

CS 4650/7650, Lecture 13:

Parsing 1

Jacob Eisenstein

September 26, 2013

So far we've explored finite-state models, which correspond to regular languages.

- **representations:** (weighted) finite state automata
- **probabilistic models:** HMMs (as a special case), CRFs
- **algorithms:** viterbi, forward-backward, $\mathcal{O}(NK^2)$ time complexity.
- **linguistic phenomena:**
 - morphology
 - language models
 - part-of-speech disambiguation
 - named entity recognition (chunking)

Is finite state enough?

1 Is English a regular language?

Regular languages are closed under intersection:

- $K \cap L$ is the set of strings in both K and L
- $K \cap L$ is regular iff K and L are regular

How to prove English is not regular:

- Let K be the set of grammatical English sentences
- Let L be some regular language
- Show that the intersection is not regular

We're going to prove this using center embedding:

1. *The cat is fat.*

2. *The cat that the dog chased is fat.*
3. **The cat that the dog is fat.*
4. *The cat that the dog that the monkey kissed chased is fat.*
5. **The cat that the dog that the monkey chased is fat.*

Proof sketch:

- K is the set of grammatical english sentences.
It excludes sentences (3) and (5).
- L is the regular language *the cat (that N)_t⁺ V _t⁺ is fat.*
- The language $L \cap K$ is *the cat (that N)_tⁿ V _tⁿ is fat.*

Note that the issue here is not just infinite repetition or productivity; FSAs can handle productive phenomena like *the big red smelly plastic figurine*.

Anyway, what do you think of this argument?

1.1 Is deep center embedding really part of English?

Karlsson (2007) searched for multiple (phrasal) center embeddings in corpora from 7 languages:

- Very few examples of double embedding
- Only 13 examples of triple embedding (none in speech)
- Zero examples of quadruple embeddings

Note that we can build an FSA to accept center-embedding up to any finite depth.

Chomsky and many linguists distinguish between

- **Competence**: the fundamental abilities of the (idealized) human language processing system
- **Performance**: real utterances produced by speakers, subject to non-linguistic factors such as cognitive limitations

Even if English *as performed* is regular, the underlying generative grammar may be context-free... **or beyond**. There is a similar proof that at least some languages are not context-free! I'll post slides with this proof idea.

1.2 How much expressiveness do we need?

- Shieber (1985) makes a similar argument with Swiss-German syntax. In response to the objection that all attested constructions are finite, Shieber writes:

Down this path lies tyranny. Acceptance of this argument opens the way to proofs of natural languages as regular, nay, **finite**.

- In practice, many real constructions are much simpler to handle in context-free rather than finite-state representations:

*The **processor has** 10 million times fewer transistors on it than today's typical microprocessors, **runs** much more slowly, and **operates** at five times the voltage...*

- The easy way:

$S \rightarrow \text{NN VP}$

$\text{VP} \rightarrow \text{VP3S} \mid \text{VPN3S} \mid \dots$

$\text{VP3S} \rightarrow \text{VP3S}, \text{VP3S}, \text{and VP3S} \mid \text{VBZ} \mid \text{VBZ NP} \mid \dots$

- The hard way: build an FST that basically replicates all of English grammar for VPs with 3S and non-3S subjects.

- Mainstream parsing focuses on CFGs, but there is some work on “mildly” context-sensitive grammars.

2 Context-Free Languages

In the Chomsky hierarchy, context-free languages (CFLs) are a strict generalization of regular languages.

regular	context-free
regular expressions	context-free grammars (CFGs)
finite-state machines	pushdown automata
paths	derivations

Context-free grammars define CFLs. They are sets of permissible *productions* which allow you to **derive** strings composed of surface symbols.

$S \rightarrow NP VP_1$

$NP \rightarrow \text{the } N \mid NP \text{ RELCLAUSE}$

$\text{RELCLAUSE} \rightarrow \text{that } NP V_t$

$V_t \rightarrow \text{ate} \mid \text{chased} \mid \text{befriended} \mid \dots$

$N \rightarrow \text{cat} \mid \text{dog} \mid \text{monkey} \mid \dots$


$VP_1 \rightarrow \text{is fat}$

An important feature of CFGs is *recursion*, in which a nonterminal can be derived from itself.


More formally , a CFG is a tuple $\langle N, \Sigma, R, S \rangle$:

- N a set of non-terminals
- Σ a set of terminals (distinct from N)
- R a set of productions, each of the form $A \rightarrow \beta$,
where $A \in N$ and $\beta \in (\Sigma \cup N)^*$
- S a designated start symbol

- Context free grammars provide rules for generating strings.
- A surface string can be **parsed** into a series of productions (a **derivation**).
- Parses can be viewed as trees or as bracketings:



$(S(NP(PRP\ She)(VP(VBZ\ eats)$
 $(NP(NP(NN\ sushi))(PP\ with(NP(NNS\ chopsticks))))))$



$(S(NP(PRP\ She)(VP(VBZ\ eats)$
 $(NP(NN\ sushi))$
 $(PP\ with(NP(NNS\ chopsticks))))))$

Semantics Ideally, each derivation will have a distinct semantic interpretation, and all possible interpretations will be represented in some derivation.

$(NP(NP\ Ban\ (PP\ on\ (NP\ nude\ dancing\)))$
 $(PP\ on\ (NP\ Governor's\ desk\)))$

$$\begin{aligned}
 & (_{NP} \textit{Ban} (_{PP} \textit{on} (_{NP} (_{NP} \textit{nude dancing}) \\
 & \qquad (_{PP} \textit{on} (_{NP} \textit{Governor's desk})))))
 \end{aligned}$$

Sadly, this is not always the case.

$$\begin{aligned}
 & (_{NP} (_{JJ} \textit{nice}) (_{JJ} \textit{little}) (_{NN} \textit{car})) \\
 & (_{NP} (_{JJ} \textit{nice}) (_{NP} (_{JJ} \textit{little}) (_{NN} \textit{car}))) \\
 & (_{NP} (_{JJ} \textit{nice}) (_{NP} (_{JJ} \textit{little}) (_{NP} (_{NN} \textit{car}))))
 \end{aligned}$$

3 Constituency

- In natural language grammars, the non-terminals should reflect syntactic categories.
- Bracketed substrings (e.g., *sushi with chopsticks*) are called **constituents**.
- There are several tests for constituency, including:
 - substitution
 - coordination
 - movement

Substitution Constituents generated by the same non-terminal should be substitutable in many contexts:

- $(_{NP} \textit{The ban}) \textit{ is on the desk.}$
- $(_{NP} \textit{The Governor's desk}) \textit{ is on the desk.}$
- $(_{NP} \textit{The ban on dancing on the desk}) \textit{ is on the desk.}$
- $* (_{PP} \textit{On the desk}) \textit{ is on the desk.}$

A more precise test for whether a set of substrings constitute a single category is whether they can be replaced by the same pronouns.

- $(_{NP} \textit{It}) \textit{ is on the desk.}$

What about verbs?

- $I (_{V} \textit{gave}) \textit{ it to Anne.}$
- $I (_{V} \textit{taught}) \textit{ it to Anne.}$
- $I (_{V} \textit{gave}) \textit{ Anne a fish}$
- $* I (_{V} \textit{taught}) \textit{ Anne a fish}$

This suggests we need nonterminals which distinguish verbs based on the arguments they can take. The technical name for this is *subcategorization*.

Coordination Constituents generated by the same non-terminal can usually be *coordinated* using words like *and* and *or*:

- *We fought* (_{PP} *on the hills*) *and* (_{PP} *in the hedges*).
- *We fought* (_{ADVP} *as well as we could*).
- **We fought* (_{ADVP} *as well as we could*) *and* (_{PP} *in the hedges*).

This too doesn't always work:

- *She* (_{VP} *went*) (_{PP} *to the store*).
- *She* (_{VP} *came*) (_{PP} *from the store*).
- *She* (*went to*) *and* (*came from*) *the store*.

Movement Valid constituents can be moved as a unit, preserving grammaticality.

- Passivization
 - (*The governor*) *banned* (*nude dancing on his desk*)
 - (*Nude dancing on his desk*) *was banned by* (*the governor*)
- Wh- movement
 - (*Nude dancing was banned*) *on* (*the desk*).
 - (*The desk*) *is where* (*nude dancing was banned*)
- Topicalization
 - (*He banned nude dancing*) *to appeal to conservatives*.
 - *To appeal to conservatives*, (*he banned nude dancing*).

4 A simple grammar of English

4.1 Noun phrases

Let's start with noun phrases:

- *She* *sleeps* (Pronoun)
- *Arlo* *sleeps* (Proper noun)
- *Fish* *sleep* (Mass noun)
- *The fish* *sleeps* (determiner + noun)
- *The blue fish* *sleeps* (DT + JJ + NN)

- *The girl from Omaha sleeps* (NP + PP)
- *The student who ate 15 donuts sleeps* (NP + RelClause)

So overall, we can summarize this fragment as

$$\begin{aligned}\text{NP} &\rightarrow \text{PRP} \mid \text{NNP} \mid \text{DT} \text{ NOM} \\ \text{NOM} &\rightarrow \text{ADJP} \text{ NOM} \mid \text{NN} \\ \text{NP} &\rightarrow \text{NP} \text{ PP} \mid \text{NP} \text{ RELCLAUSE}\end{aligned}$$

We're leaving out some detail, like pluralization and possessives, but you get the idea.

4.2 Adjectival and prepositional phrases

- *Very funny*
- *The large, blue fish*
- *The man from la mancha*

$$\begin{aligned}\text{ADJP} &\rightarrow \text{JJ} \mid \text{RB} \text{ ADJP} \mid \text{JJ} \text{ ADJP} \\ \text{PP} &\rightarrow \text{IN} \text{ NP} \mid \text{TO} \text{ NP}\end{aligned}$$

4.3 Verb phrases

- *She sleeps*
- *She sleeps restlessly*
- *She sleeps at home*
- *She eats sushi*¹
- *She gives John sushi*

$$\text{VP} \rightarrow \text{V} \mid \text{VP} \text{ RB} \mid \text{VP} \text{ PP} \mid \text{V} \text{ NP} \mid \text{V} \text{ NP} \text{ NP} \mid \text{V} \text{ NP} \text{ RB}$$

But what about **She sleeps sushi* or **She speaks John Japanese*?

- Classes of verbs can take different numbers of arguments.
- This is called **subcategorization**

$$\begin{aligned}\text{VP} &\rightarrow \text{V-INTRANS} \mid \text{V-TRANS} \text{ NP} \mid \text{V-DITRANS} \text{ NP} \text{ NP} \\ \text{VP} &\rightarrow \text{VP} \text{ RB} \mid \text{VP} \text{ PP}\end{aligned}$$

We would also need to handle modal and auxiliary verbs that allow us to create complex tenses, like *She will have eaten sushi* but not **She will have eats sushi*.

¹Sushi examples from Julia Hockenmaier

4.4 Sentences

- *She eats sushi*

$$S \rightarrow NP \ VP$$

- *Sometimes, she eats sushi*

$$S \rightarrow AdvP \ S$$

- *In Japan, she eats sushi*

$$S \rightarrow PP \ S$$

- What about **I eats sushi, *She eat sushi??*

$$S \rightarrow NP.3S \ VP.3S \mid NP.N3S \ VP.N3S$$

In general, we need **features** to capture this kind of agreement.

4.5 Conjunctions

- *She eats sushi and candy*

$$NP \rightarrow NP \ and \ NP$$

- *She eats sushi and drinks soda*

$$VP \rightarrow VP \ and \ VP$$

- *She eats sushi and he drinks soda*

$$S \rightarrow S \ and \ S$$

- *fresh and tasty sushi*

$$AdjP \rightarrow JJ \ and \ JJ$$

We'd need a little more cleverness to properly cover groups larger than two.

4.6 Odds and ends

- *I gave sushi to the girl **who eats sushi**.* This is a relative clause,

$$RELClause \rightarrow who \ VP \mid that \ VP$$

- *I took sushi from the man **offering sushi**.* This is a gerundive postmodifier.

$$NOM \rightarrow NOM \ GERUNDVP$$

$$GERUNDVP \rightarrow VBZ \mid VBZ \ NP \mid VBZ \ PP \mid \dots$$

- **Can** *she eat sushi?* (notice it's not *eats*)

$$S \rightarrow AUX \ NP \ VP$$

- ... and many more

5 Grammar design

Our goal is a grammar that avoids

- **Overgeneration**: deriving strings that are not grammatical.
- **Undergeneration**: failing to derive strings that are grammatical.

To avoid undergeneration, we would need thousands of productions.

Typically, grammars are defined in conjunction with large-scale **treebank** annotation projects.

- An annotation guideline specifies the non-terminals and how they go together.
- The annotators then apply these guidelines to data.
- The grammar rules can then be read off the data.

The Penn Treebank contains one million parsed words of Wall Street Journal text from the 1990s.


6 Grammar equivalence and normal form


- Grammars are weakly equivalent if they generate the same strings.
- Grammars are strongly equivalent if they generate the same strings **and** assign the same phrase structure to each sentence.
- In Chomsky Normal Form (CNF), all productions are either:

$$A \rightarrow BC$$

$$A \rightarrow a$$

- All CFGs can be converted into a weakly equivalent grammar in CNF.
- This is very handy for parsing algorithms.





- Binarization is easy:
group right children into new non-terminals.
- Un-binarization is important!
people with claws is not a constituent in the original parse.
- Unary productions are best handled by modifying the algorithm.

7 Parsing

Parsing is the process of determining whether a sentence is in a context-free language, by searching for a legal derivation. Some possibilities:

- **Top-down**: start with the start symbol, and see if we can derive the sentence.
- **Bottom-up**: combine the observed symbols using whatever productions we can, until we reach the start symbol
- **Left-to-right**: move through the input, incrementally building a parse tree

Before we get into these different possibilities, let's see whether exhaustive search is possible. Suppose we only have one non-terminal, X , and it has binary productions

$$\begin{aligned}
 X &\rightarrow X X \\
 X &\rightarrow \textit{the girl} \mid \textit{ate sushi} \mid \dots
 \end{aligned}$$

How many different ways could we parse a sentence? This is just equal to the number of binary bracketings of the words in the sentence, which is a Catalan number. Catalan numbers grow **super-exponentially** in the length of the sentence, $C_n = \frac{(2n)!}{(n+1)!n!}$.

8 CKY parsing

CKY is a bottom-up parsing allows us to test whether a sentence is in a context-free language, without considering all possible parses. First we form small constituents, then try to merge them into larger constituents.

Let's start with an example grammar:

$$\begin{aligned} S &\rightarrow VP\ NP \\ NP &\rightarrow NP\ PP \mid we \mid sushi \mid chopsticks \\ PP &\rightarrow P\ NP \\ P &\rightarrow with \\ VP &\rightarrow VP\ NP \mid VP\ PP \mid eat \end{aligned}$$

Suppose we encounter the sentence *We eat sushi with chopsticks*.

- The first thing that we notice is that we can apply unary productions to obtain NP VP NP P NP
- Next, we can apply a binary production to merge the first NP VP into an S.
- Or we could merge VP NP into VP
- ... and so on

Let's systematize this. Here is the CKY algorithm:

```

for j : [1,N] do
  t[j, j-1] ← {A | A → xj ∈ R}
  for i : [j-2, 0] do
    for k : [i+1, j-1] do
      t[i, j] ← t[i, j] ∪ {A | A → BC ∈ R, B ∈ t[i, k], C ∈ t[k, j]}
    end for
  end for
end for

```

To handle unary transitions, we compute the *unary closure* of each non-terminal.


- e.g., if $S \rightarrow VP$, $VP \rightarrow V$, then add $S \rightarrow V$
- At each table entry $t[i, j]$
 - For each non-terminal $A \in t[i, j]$

* Add all elements of the reflexive unary closure for A



Complexity What is the complexity of CKY?

- Space complexity: $\mathcal{O}(L^2|N|)$
- Time complexity: $\mathcal{O}(N^3|R|)$
- L is length of sentence,
 $|N|$ is the number of non-terminals,
 $|R|$ is the number of production rules
- But in practice...



~ 20K Rules

(not an
optimized
parser!)

Observed
exponent:

3.6

It's worse than worst-case! (figure from Dan Klein)

- Longer sentences “unlock” more of the grammar.

9 Ambiguity in parsing


- Syntactic ambiguity is endemic to natural language:²

[i+-i]

²Examples borrowed from Dan Klein

- Attachment ambiguity: *we eat sushi with chopsticks, I shot an elephant in my pajamas.*
 - Modifier scope: *southern food store*
 - Particle versus preposition: *The puppy tore up the staircase.*
 - Complement structure: *The tourists objected to the guide that they couldn't hear.*
 - Coordination scope: *"I see," said the blind man, as he picked up the hammer and saw.*
 - Multiple gap constructions: *The chicken is ready to eat*
- In morphology, we didn't just want to know which derivational forms are *legal*, we wanted to know which were likely.
 - Syntactic parsing is all about choosing among the many, many legal parses for a given sentence.

Here's another example, which we've seen before:



- A minimal grammar permits 36 parses!
- Real-size broad coverage grammars permit millions of parses.

Classical parsers faced a tradeoff:

- broad coverage with tons of ambiguity...
- or limited coverage in exchange for constraints on ambiguity

Consequently, deterministic parsers produced no analysis for many sentences.

9.1 Local solutions

Some ambiguity can be resolved locally:

- [*imposed* [*a ban* [*on asbestos*]]]
- [*imposed* [*a ban*]][*on asbestos*]]

- Hindle and Rooth (1990) proposed a likelihood ratio test:

$$LR(v, n, p) = \frac{P(p|v)}{P(p|n)} = \frac{P(on|imposed)}{P(on|ban)}$$

where we select VERB attachment if $LR(v, n, p) > 1$.

- But the likelihood-ratio approach ignores important information, like the phrase being attached.

- ...[*it* [*would end* [*its venture* [*with Maserati*]]]]
- ...[*it* [*would end* [*its venture*]][*with Maserati*]]

- The likelihood ratio gets this wrong

- $P(with|end) = \frac{607}{5156} = 0.118$
- $P(with|venture) = \frac{155}{1442} = 0.107$

Other features (e.g., *Maserati*) argue for noun attachment. How can we add them?

Machine learning solutions Ratnaparkhi et al (1994) propose a maximum-entropy (logistic regression) approach:

$$P(N|would\ end\ its\ venture\ with\ Maserati) = \frac{e^{\mathbf{w}^T \mathbf{f}(would\ end\ its\ venture\ with\ Maserati)}}{1 + e^{\mathbf{w}^T \mathbf{f}(would\ end\ its\ venture\ with\ Maserati)}}$$

Features include n-grams and word classes from hierarchical word clustering; accuracy is roughly 80%.

Collins and Brooks (1995) argued that attachment depends on four **heads**:

- the preposition (*with*)
- the VP attachment site (*end*)
- the NP attachment site (*venture*)
- the NP to be attached (*Maserati*)

They propose a backoff-based approach:

- First, look for counts of the tuple $\langle with, Maserati, end, venture \rangle$
- If none, try $\langle with, Maserati, end \rangle + \langle with, end, venture \rangle + \langle with, Maserati, venture \rangle$
- If none, try $\langle with, Maserati \rangle + \langle with, end \rangle + \langle with, venture \rangle$
- If none, try $\langle with \rangle$

Accuracy is roughly 84%. This approach of combining relative frequency estimation, smoothing, and backoff was very characteristic of late 1990s statistical NLP.

9.2 Beyond local solutions

Framing the problem as attachment ambiguity is limiting:

- assumes the parse is mostly done, leaving just a few attachment ambiguities to solve
- But realistic sentences have more than a few syntactic interpretations.
- Attachment decisions are interdependent:
 - *Cats scratch people with claws with knives.*
 - We may want to attach *with claws* to *scratch*.
 - But then we have nowhere to put *with knives*.

The task of statistical parsing is to produce a single analysis that resolves all syntactic ambiguities.

10 PCFGs

We want the parse τ that maximizes $P(\tau|S)$.

$$\begin{aligned}
 \arg \max_{\tau} P(\tau|S) &= \arg \max_{\tau} \frac{P(\tau, S)}{P(S)} \\
 &= \arg \max_{\tau} P(\tau, S) \\
 &= \arg \max_{\tau} P(S|\tau)P(\tau) \\
 &= \arg \max_{\tau: S = \text{yield}(\tau)} P(\tau)
 \end{aligned}$$

- The **yield** of a tree is the string of terminal symbols that can be read off the leaf nodes.
- The set $\{\tau : S = \text{yield}(\tau)\}$ is exactly the set of all derivations of S in a CFG G .


PCFGs extend the CFG by adding probability to each production:

S	$\rightarrow NP VP$	0.9
S	$\rightarrow S conj S$	0.1
NP	$\rightarrow N$	0.2
NP	$\rightarrow DT N$	0.3
NP	$\rightarrow N NP$	0.2
NP	$\rightarrow JJ NP$	0.2
NP	$\rightarrow NP PP$	0.1
VP	$\rightarrow V$	0.4
VP	$\rightarrow V NP$	0.3
VP	$\rightarrow V NP NP$	0.1
VP	$\rightarrow VP PP$	0.2
PP	$\rightarrow P NP$	1.0

The probabilities for all productions involving a single LHS must sum to 1:

$$\sum_{\alpha} P(X \rightarrow \alpha | X) = 1$$

The probability $P(\tau)$ is just the product of all the productions:



10.1 Estimation

- As in supervised HMMs, estimation is easy (for now!).
- PCFG probabilities can be estimated directly from a treebank:

$$P(VP \rightarrow VP PP) = \frac{c(VP \rightarrow VP PP)}{c(VP)}$$

10.2 Three basic problems for PCFGs

Let $\tau \in T$ be a derivation, S be a sentence, and λ a PCFG.

- **Decoding:** Find $\hat{\tau} = \arg \max_{\tau} P(\tau, S; \lambda)$
- **Likelihood:** Find $P(w; \lambda) = \sum_{\tau} P(\tau, S; \lambda)$
- **(Unsupervised) Estimation:** Find $\arg \max_{\lambda} P(S_{1...N} | \lambda)$