

A Technical Introduction to Statistical Natural Language Processing

Jacob Eisenstein

April 16, 2017

Contents

Contents	1
I Words, bags of words, and features	9
1 Linear classification and features	11
1.1 Review of basic probability	14
1.2 Naïve Bayes	21
2 Discriminative learning	31
2.1 Perceptron	32
2.2 Loss functions and large margin classification	34
2.3 Logistic regression	39
2.4 Optimization	43
2.5 *Additional topics in classification	45
2.6 Summary of learning algorithms	48
3 Linguistic applications of classification	53
3.1 Sentiment and opinion analysis	53
3.2 Word sense disambiguation	57
3.3 Design decisions for text classification	60
3.4 Evaluating text classification	62
4 Learning without supervision	65
4.1 <i>K</i> -means clustering	66
4.2 The Expectation Maximization (EM) Algorithm	67
4.3 Applications of EM	72
4.4 *Other approaches to learning with latent variables	75
5 Language models	79
5.1 N-gram language models	80

5.2	Smoothing and discounting	83
5.3	Recurrent neural network language models	88
5.4	Evaluating language models	94
5.5	Out-of-vocabulary words	96
II	Sequences and trees	99
6	Sequence labeling	101
6.1	Sequence labeling as classification	101
6.2	Sequence labeling as structure prediction	103
6.3	The Viterbi algorithm	105
6.4	Hidden Markov Models	109
6.5	Discriminative sequence labeling	116
6.6	*Unsupervised sequence labeling	127
7	Applications of sequence labeling	131
7.1	Part-of-speech tagging	131
7.2	Shallow parsing	140
7.3	Named entity recognition	140
7.4	Dialogue acts	141
7.5	Code switching	142
8	Finite-state automata	143
8.1	Automata and languages	144
8.2	Weighted Finite State Automata	149
8.3	Semirings	153
8.4	Finite state transducers	155
8.5	Weighted FSTs	157
8.6	Applications of finite state composition	159
8.7	Discriminative structure prediction	161
9	Morphology	163
9.1	Types of morphemes	166
9.2	Types of morphology	168
9.3	Computing and morphology	174
10	Context-free grammars	177
10.1	Is English a regular language?	177
10.2	Context-Free Languages	179
10.3	Constituents	182
10.4	A simple grammar of English	184

10.5 Grammar equivalence and normal form	190
11 Context-free Parsing	193
11.1 Deterministic bottom-up parsing	193
11.2 Ambiguity in parsing	196
11.3 Weighted Context-Free Grammars	200
11.4 Improving Parsing by Refined Non-terminals	205
11.5 Discriminative parsing	214
12 Dependency Parsing	217
12.1 Dependency grammar	217
12.2 Graph-based dependency parsing	222
12.3 Transition-based dependency parsing	229
12.4 Applications	232
III Meaning	235
13 Logical semantics	237
13.1 Meaning representations	237
13.2 Logical representations of meaning	240
13.3 Syntax and semantics	242
13.4 Semantic parsing	245
14 Shallow semantics	247
14.1 Predicates and arguments ¹	247
14.2 Semantic Role Labeling	251
14.3 FrameNet	255
14.4 Abstract Meaning Representation	257
15 Distributional and distributed semantics	259
15.1 The distributional hypothesis	259
15.2 Design decisions for word representations	261
15.3 Distributional semantics	263
15.4 Distributed representations	269
16 Reference Resolution	275
16.1 Forms of referring expressions	276
16.2 Learning for coreference resolution	282
16.3 Entity linking and multi-document coreference resolution	286

¹This section follows closely from J&M 2009

IV Applications	289
17 Information extraction	291
17.1 Entities	292
17.2 Relations	294
17.3 Events and processes	294
17.4 Facts, beliefs, and hypotheticals	294
18 Machine translation	295
18.1 The noisy channel model	295
18.2 Translation modeling	297
18.3 Phrase-based translation	303
18.4 Syntactic MT	304
18.5 Algorithms for SCFGs	307
V Learning	311
19 Semi-supervised learning	313
19.1 Semisupervised learning	315
19.2 Domain adaptation	323
19.3 Other learning settings	325
20 Beyond linear models	327
20.1 Representation learning	327
20.2 Convolutional neural networks	327
20.3 Recursive neural networks	327
20.4 Encoder-decoder models	327
20.5 Structure prediction	327
Bibliography	329

Preface

This text is built from the notes that I use for teaching Georgia Tech’s undergraduate and graduate courses on natural language processing, CS 4650 and 7650. There are several other good resources (e.g., Manning and Schütze, 1999; Jurafsky and Martin, 2009; Smith, 2011; Figueiredo et al., 2013; Collins, 2013), but for various reasons I wanted to create something of my own.

The text assumes familiarity with basic linear algebra, and with calculus through Lagrange multipliers. It includes a refresher on probability, but some previous exposure would be helpful. An introductory course on the analysis of algorithms is also assumed; in particular, the reader should be familiar with asymptotic analysis of the time and memory costs of algorithms, and should have seen dynamic programming. No prior background in machine learning or linguistics is assumed, and even students with background in machine learning should be sure to read the introductory chapters, since the notation used in natural language processing is different from typical machine learning presentations, due to the emphasis on structure prediction in applications of machine learning to language. Throughout the book, advanced material is marked with an asterisk, and can be safely skipped.

The notes focus on what I view as a core subset of the field of natural language processing, unified by the concepts of linear models and structure prediction. A remarkable thing about the field of natural language processing is that so many problems in language technology can be solved by a small number of methods. These notes focus on the following methods:

Search algorithms shortest path, Viterbi, CKY, minimum spanning tree, shift-reduce, integer linear programming, dual decomposition (maybe), beam search.

Learning algorithms Naïve Bayes, logistic regression, perceptron, expectation-maximization, matrix factorization, backpropagation.

The goal of this text is to teach how these methods work, and how they can be applied to problems that arise in the computer processing of natural language: document classification, word sense disambiguation, sequence labeling (part-of-speech tagging and named entity recognition), parsing, coreference resolution, relation extraction, discourse analysis,

and, to a limited degree, language modeling and machine translation. Because proper application of these techniques requires understanding the underlying linguistic phenomena, the notes also include chapters on the foundations of morphology, syntactic parts of speech, context-free grammar, semantics, and discourse; however, for a detailed understanding of these topics, a full-fledged linguistics textbook should be consulted (e.g., Akmajian et al., 2010; Fromkin et al., 2013).

-Jacob Eisenstein, April 16, 2017

Notation

w_m	word token at position m
$\mathbf{x}^{(i)}$	a (column) vector of feature counts for instance i , often word counts
$\mathbf{x}_{i:j}$	elements i through j (inclusive) of a vector \mathbf{x}
N	number of training instances
M	length of a sequence (of words or tags)
$ \mathcal{V} $	number of words in vocabulary
$y^{(i)}$	the label for instance i
\hat{y}	a predicted label
\mathbf{y}	a vector of labels
\mathcal{Y}	the set of all possible labels
K	number of possible labels $K = \mathcal{Y} $
$\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)})$	feature vector for instance i with label $y^{(i)}$
$\boldsymbol{\theta}$	a (column) vector of weights
$\Pr(A)$	probability of event A
$p_B(b)$	the marginal probability of random variable B (often implicit) taking value b
$\mathcal{Y}(\mathbf{w})$	the set of possible tag sequences for the word sequence \mathbf{w}
\diamond	the start tag
\blacklozenge	the stop tag
\square	the start token
\blacksquare	the stop token
λ	the amount of regularization

Part I

Words, bags of words, and features

Chapter 1

Linear classification and features

Suppose you want to build a spam detector, in which each document is classified as “spam” or “ham.” How would you do it, using only the text in the email?

One solution is to represent document i as a column vector of word counts: $\mathbf{x}^{(i)} = [0 \ 1 \ 1 \ 0 \ 0 \ 2 \ 0 \ 1 \ 13 \ 0 \ \dots]^T$, where $x_{i,j}$ is the count of word j in document i . Suppose the size of the vocabulary is V , so that the length of $\mathbf{x}^{(i)}$ is also V . The object $\mathbf{x}^{(i)}$ is a vector, but colloquially we call it a **bag of words**, because it includes only information about the count of each word, and not the order in which they appear.

We’ve thrown out grammar, sentence boundaries, paragraphs — everything but the words! But this could still work. If you see the word *free*, is it spam or ham? How about *Bayesian*? One approach would be to define a “spamminess” score for every word in the dictionary, and then just add them up. These scores are called **weights**, written θ , and we’ll spend a lot of time talking about where they come from.

But for now, let’s generalize: suppose we want to build a multi-way classifier to distinguish stories about sports, celebrities, music, and business. Each label $y^{(i)}$ is a member of a set of K possible labels \mathcal{Y} . Our goal is to predict a label $\hat{y}^{(i)}$, given the bag of words $\mathbf{x}^{(i)}$, using the weights θ . We’ll do this using a vector inner product between the weights θ and a **feature vector** $\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)})$. As the notation suggests, the feature vector is constructed by combining $\mathbf{x}^{(i)}$ and $y^{(i)}$. For example, feature j might be,

$$f_j(\mathbf{x}^{(i)}, y^{(i)}) = \begin{cases} 1, & \text{if } (\text{free} \in \mathbf{x}^{(i)}) \wedge (y^{(i)} = \text{SPAM}) \\ 0, & \text{otherwise} \end{cases} \quad (1.1)$$

For any pair $\langle \mathbf{x}^{(i)}, y^{(i)} \rangle$, we then define $\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)})$ as,

$$\mathbf{f}(\mathbf{x}, Y = 0) = [\mathbf{x}^\top, \underbrace{0, 0, \dots, 0}_{V \times (K-1)}]^\top \quad (1.2)$$

$$\mathbf{f}(\mathbf{x}, Y = 1) = [\underbrace{0, 0, \dots, 0}_V, \mathbf{x}^\top, \underbrace{0, 0, \dots, 0}_{V \times (K-2)}]^\top \quad (1.3)$$

$$\mathbf{f}(\mathbf{x}, Y = 2) = [\underbrace{0, 0, \dots, 0}_{2 \times V}, \mathbf{x}^\top, \underbrace{0, 0, \dots, 0}_{V \times (K-3)}]^\top \quad (1.4)$$

$$\mathbf{f}(\mathbf{x}, Y = K) = [\underbrace{0, 0, \dots, 0}_{V \times (K-1)}, \mathbf{x}^\top]^\top, \quad (1.5)$$

where $\underbrace{0, 0, \dots, 0}_{V \times (K-1)}$ is a column vector of $V \times (K - 1)$ zeros. This arrangement is shown in Figure 1.1. This notation may seem like a strange choice, but in fact it helps to keep things simple. Given a vector of weights, $\boldsymbol{\theta} \in \mathbb{R}^{V \times K}$, we can now compute the inner product $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y)$. This inner product gives a scalar measure of the score for label y , given observations \mathbf{x} . For any document $\mathbf{x}^{(i)}$, we predict the label \hat{y} as

$$\hat{y} = \operatorname{argmax}_y \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) \quad (1.6)$$

This inner product is the fundamental equation for linear classification, and it is the reason we prefer the feature function notation $\mathbf{f}(\mathbf{x}, y)$. The notation gives a clean separation between the **data** (\mathbf{x} and y) and the **parameters**, which are expressed by the single vector of weights, $\boldsymbol{\theta}$. As we will see in later chapters, this notation also generalizes nicely to **structured output spaces**, in which the space of labels \mathcal{Y} is very large, and we want to model shared substructure between labels.

Often we'll add an **offset** feature at the end of \mathbf{x} , which is always 1; we then have to also add an extra zero to each of the zero vectors. This gives the entire feature vector $\mathbf{f}(\mathbf{x}, y)$ a length of $(V + 1) \times K$. The weight associated with this offset feature can be thought of as a “bias” for each label. For example, if we expect most documents to be spam, then the weight for the offset feature for $Y = \text{spam}$ should be larger than the weight for the offset feature for $Y = \text{ham}$.

Returning to the weights $\boldsymbol{\theta}$ — where do they come from? As already suggested, we could set the weights by hand. If we wanted to distinguish, say, English from Spanish, we could use English and Spanish dictionaries, and set the weight to one for each word that



Figure 1.1: The bag-of-words and feature vector representations, for a hypothetical text classification task.

appears in the associated dictionary. For example,¹

$$\begin{aligned}
 \theta_{(E, \text{bicycle})} &= 1 & \theta_{(S, \text{bicycle})} &= 0 \\
 \theta_{(E, \text{bicicleta})} &= 0 & \theta_{(S, \text{bicicleta})} &= 1 \\
 \theta_{(E, \text{con})} &= 1 & \theta_{(S, \text{con})} &= 1 \\
 \theta_{(E, \text{ordinateur})} &= 0 & \theta_{(S, \text{ordinateur})} &= 0.
 \end{aligned}$$

Similarly, if we want to distinguish positive and negative sentiment, we could use positive and negative *sentiment lexicons*, which are defined by expert psychologists (Tausczik and Pennebaker, 2010). You'll try this in Problem Set 1.

But it is usually not easy to set classification weights by hand. Instead, we will learn them from data. For example, email users manually label thousands of messages as "spam" or "not spam"; newspapers label their own articles as "business" or "fashion." Such **instance labels** are a typical form of labeled data that we will encounter in NLP. In **supervised machine learning**, we use instance labels to automatically set the weights for a classifier. An important tool for this is probability.

¹In this notation, each tuple (language, word) indexes an element in θ , which remains a vector.

1.1 Review of basic probability

Probability theory provides a way to reason about random events. The sorts of random events that are typically used to explain probability theory include coin flips, card draws, and the weather. It may seem odd to think about the choice of a word as akin to the flip of a coin, particularly if you are the type of person to choose words carefully. But random or not, language has proven to be extremely difficult to model deterministically. Probability offers a powerful tool for modeling and manipulating linguistic data, which we will use repeatedly throughout this course.²

Probability can be thought of in terms of **random outcomes**: for example, a single coin flip has two possible outcomes, heads or tails. The set of possible outcomes is the **sample space**, and a subset of the **sample space** is an **event**. For a sequence of two coin flips, there are four possible outcomes, $\{HH, HT, TH, TT\}$, representing the ordered sequences heads-head, heads-tails, tails-heads, and tails-tails. The event of getting exactly one head includes two outcomes: $\{HT, TH\}$.

Formally, a probability is a function from events to the interval between zero and one: $\Pr : \mathcal{F} \rightarrow [0, 1]$, where \mathcal{F} is the set of possible events. An event that is certain has probability one; an event that is impossible has probability zero. For example, the probability of getting less than three heads on two coin flips is one. Each outcome is also an event (a set with exactly one element), and for two flips of a fair coin, the probability of each outcome is,

$$\Pr(\{HH\}) = \Pr(\{HT\}) = \Pr(\{TH\}) = \Pr(\{TT\}) = \frac{1}{4}. \quad (1.7)$$

Probabilities of event combinations

Because events are **sets** of outcomes, we can use set theoretic operations such as complement, intersection, and unions to reason about the probabilities of various event combinations.

For any event A , there is a **complement** $\neg A$, such that:

- The union $A \cup \neg A$ covers the entire sample space, and $\Pr(A \cup \neg A) = 1$;
- The intersection $A \cap \neg A = \emptyset$ is the empty set, and $\Pr(A \cap \neg A) = 0$.

In the coin flip example, the event of obtaining a single head on two flips corresponds to the set of outcomes $\{HT, TH\}$; the complement event includes the other two outcomes, $\{TT, HH\}$.

²A good introduction to probability theory is offered by Manning and Schütze (1999), which helped to motivate this section. For more detail, Sharon Goldwater provides another useful reference, <http://homepages.inf.ed.ac.uk/sgwater/teaching/general/probability.pdf>.

Probabilities of disjoint events

In general, when two events have an empty intersection, $A \cap B = \emptyset$, they are said to be **disjoint**. The probability of the union of two disjoint events is equal to the sum of their probabilities,

$$A \cap B = \emptyset \Rightarrow \Pr(A \cup B) = \Pr(A) + \Pr(B). \quad (1.8)$$

This is the **third axiom of probability**, and can be generalized to any countable sequence of disjoint events.

In the coin flip example, we can use this axiom to derive the probability of the event of getting a single head on two flips. This event is the set of outcomes $\{HT, TH\}$, which is the union of two simpler events, $\{HT, TH\} = \{HT\} \cup \{TH\}$. The events $\{HT\}$ and $\{TH\}$ are disjoint. Therefore,

$$\Pr(\{HT, TH\}) = \Pr(\{HT\} \cup \{TH\}) = \Pr(\{HT\}) + \Pr(\{TH\}) \quad (1.9)$$

$$= \frac{1}{4} + \frac{1}{4} = \frac{1}{2}. \quad (1.10)$$

For events that are not disjoint, it is still possible to compute the probability of their union:

$$\Pr(A \cup B) = \Pr(A) + \Pr(B) - \Pr(A \cap B). \quad (1.11)$$

This can be derived from the third axiom of probability. First, consider an event that includes all outcomes in B that are not in A , which we can write as $B - (A \cap B)$. By construction, this event is disjoint from A .³ We can therefore apply the additive rule,

$$\Pr(A \cup B) = \Pr(A) + \Pr(B - (A \cap B)) \quad (1.12)$$

$$\Pr(B) = \Pr(B - (A \cap B)) + \Pr(A \cap B) \quad (1.13)$$

$$\Pr(B - (A \cap B)) = \Pr(B) - \Pr(A \cap B) \quad (1.14)$$

$$\Pr(A \cup B) = \Pr(A) + \Pr(B) - \Pr(A \cap B). \quad (1.15)$$

Law of total probability

A set of events $\mathcal{B} = \{B_1, B_2, \dots, B_N\}$ is a **partition** of the sample space iff each pair of events is disjoint ($B_i \cap B_j = \emptyset$), and the union of the events is the entire sample space. The law of total probability states that we can **marginalize** over these events as follows,

$$\Pr(A) = \sum_{B_n \in \mathcal{B}} \Pr(A \cap B_n). \quad (1.16)$$

Note for any event B , the union $B \cup \neg B$ forms a partition of the sample space. Therefore, an important special case of the law of total probability is,

$$\Pr(A) = \Pr(A \cap B) + \Pr(A \cap \neg B). \quad (1.17)$$

³[todo: add figure]

Conditional probability and Bayes' rule

A **conditional probability** is an expression like $\Pr(A \mid B)$, which is the probability of the event A , assuming that event B happens too. For example, we may be interested in the probability of a randomly selected person answering the phone by saying *hello*, conditioned on that person being a speaker of English. We define conditional probability as the ratio,

$$\Pr(A \mid B) = \frac{\Pr(A \cap B)}{\Pr(B)} \quad (1.18)$$

The **chain rule** states that $\Pr(A \cap B) = \Pr(A \mid B) \times \Pr(B)$, which is just a simple rearrangement of terms from Equation 1.18. We can apply the chain rule repeatedly:

$$\begin{aligned} \Pr(A \cap B \cap C) &= \Pr(A \mid B \cap C) \times \Pr(B \cap C) \\ &= \Pr(A \mid B \cap C) \times \Pr(B \mid C) \times \Pr(C) \end{aligned}$$

Bayes' rule (sometimes called Bayes' law or Bayes' theorem) gives us a way to convert between $\Pr(A \mid B)$ and $\Pr(B \mid A)$. It follows from the chain rule:

$$\Pr(A \mid B) = \frac{\Pr(A \cap B)}{\Pr(B)} = \frac{\Pr(B \mid A) \times \Pr(A)}{\Pr(B)} \quad (1.19)$$

The terms in Bayes rule have specialized names, which we will occasionally use:

- $\Pr(A)$ is the **prior**, since it is the probability of event A without knowledge about whether B happens or not.
- $\Pr(B \mid A)$ is the **likelihood**, the probability of event B given that event A has occurred.
- $\Pr(A \mid B)$ is the **posterior**, since it is the probability of event A with knowledge that B has occurred.

Example Manning and Schütze (1999) have a nice example of Bayes' rule (sometimes called Bayes Law) in a linguistic setting. (This same example is usually framed in terms of tests for rare diseases.) Suppose one is interested in a rare syntactic construction, such as **parasitic gaps**, which occur on average once in 100,000 sentences. Here is an example:

(1.1) *Which class did you attend __ without registering for __?*

Lana Linguist has developed a complicated pattern matcher that attempts to identify sentences with parasitic gaps. It's pretty good, but it's not perfect:

(c) Jacob Eisenstein 2014-2017. Work in progress.

- If a sentence has a parasitic gap, the pattern matcher will find it with probability 0.95. (Skipping ahead, this is the **recall**; the **false negative rate** is defined as one minus the recall.)
- If the sentence doesn't have a parasitic gap, the pattern matcher will wrongly say it does with probability 0.005. (This is the **false positive rate**. The **precision** is defined as one minus the false positive rate.)

Suppose that Lana's pattern matcher says that a sentence contains a parasitic gap. What is the probability that this is true?

Let G be the event of a sentence having a parasitic gap, and T be the event of the test being positive. We are interested in the probability of a sentence having a parasitic gap given that the test is positive. This is the conditional probability $\Pr(G \mid T)$, and we can compute it from Bayes' rule:

$$\Pr(G \mid T) = \frac{\Pr(T \mid G) \times \Pr(G)}{\Pr(T)}. \quad (1.20)$$

We already know both terms in the numerator: $\Pr(T \mid G)$ is the recall, which is 0.95; $\Pr(G)$ is the prior, which is 10^{-5} .

We are not given the denominator, but we can compute it by using some of the tools that we have developed in this section. We first apply the law of total probability, using the partition $\{G, \neg G\}$:

$$\Pr(T) = \Pr(T \cap G) + \Pr(T \cap \neg G). \quad (1.21)$$

This says that the probability of the test being positive is the sum of the probability of a **true positive** ($T \cap G$) and the probability of a **false positive** ($T \cap \neg G$). Next, we can compute the probability of each of these events using the chain rule:

$$\Pr(T \cap G) = \Pr(T \mid G) \times \Pr(G) = 0.95 \times 10^{-5} \quad (1.22)$$

$$\Pr(T \cap \neg G) = \Pr(T \mid \neg G) \times \Pr(\neg G) = 0.005 \times (1 - 10^{-5}) \approx 0.005 \quad (1.23)$$

$$\Pr(T) = \Pr(T \cap G) + \Pr(T \cap \neg G) \quad (1.24)$$

$$= 0.95 \times 10^{-5} + 0.005 \approx 0.005. \quad (1.25)$$

We now return to Bayes' rule to compute the desired posterior probability,

$$\Pr(G \mid T) = \frac{\Pr(T \mid G) \Pr(G)}{\Pr(T)} \quad (1.26)$$

$$= \frac{0.95 \times 10^{-5}}{0.95 \times 10^{-5} + 0.005 \times (1 - 10^{-5})} \quad (1.27)$$

$$\approx 0.002. \quad (1.28)$$

Lana's pattern matcher is very accurate, with false positive and false negative rates below 5%. Yet the extreme rarity of this phenomenon means that a positive result from the detector is most likely to be wrong.

Independence

Two events are independent if the probability of their intersection is equal to the product of their probabilities: $\Pr(A \cap B) = \Pr(A) \times \Pr(B)$. For example, for two flips of a fair coin, the probability of getting heads on the first flip is independent of the probability of getting heads on the second flip. We can prove this by using the additive axiom defined above:

$$\Pr(\{HT, HH\}) = \Pr(HT) + \Pr(HH) = \frac{1}{4} + \frac{1}{4} = \frac{1}{2} \quad (1.29)$$

$$\Pr(\{HH, TH\}) = \Pr(HH) + \Pr(TH) = \frac{1}{4} + \frac{1}{4} = \frac{1}{2} \quad (1.30)$$

$$\Pr(\{HT, HH\}) \times \Pr(\{HH, TH\}) = \frac{1}{2} \times \frac{1}{2} = \frac{1}{4} \quad (1.31)$$

$$\Pr(\{HT, HH\} \cap \{HH, TH\}) = \Pr(HH) = \frac{1}{4} \quad (1.32)$$

$$= \Pr(\{HT, HH\}) \times \Pr(\{HH, TH\}). \quad (1.33)$$

Independence will play a key role in the discussion of probabilistic classification later in this chapter.

If $\Pr(A \cap B \mid C) = \Pr(A \mid C) \times \Pr(B \mid C)$, then the events A and B are **conditionally independent**, written $A \perp B \mid C$.

Random variables

Random variables are functions of events. Formally, we will treat random variables as functions from events to the space \mathbb{R}^n , where \mathbb{R} is the set of real numbers. This general notion subsumes a number of different types of random variables:

- **Indicator random variables** are functions from events to the set $\{0, 1\}$. In the coin flip example, we can define Y as an indicator random variable, for whether the coin has come up heads on at least one flip. This would include the outcomes $\{HH, HT, TH\}$. The event probability $\Pr(Y = 1)$ is the sum of the probabilities of these outcomes, $\Pr(Y = 1) = \frac{1}{4} + \frac{1}{4} + \frac{1}{4} = \frac{3}{4}$.
- A **discrete random variable** is a function from events to a countable subset of \mathbb{R} . Consider the coin flip example: the number of heads, X , can be viewed as a discrete random variable, $X \in 0, 1, 2$. The event probability $\Pr(X = 1)$ can again be computed as the sum of the probabilities of the events in which there is one head, $\{HT, TH\}$, giving $\Pr(X = 1) = \frac{1}{2}$.

Each possible value of a random variable is associated with a subset of the sample space. In the coin flip example, $X = 0$ is associated with the event $\{TT\}$, $X = 1$ is associated with the event $\{HT, TH\}$, and $X = 2$ is associated with the event $\{HH\}$.

Assuming a fair coin, the probabilities of these events are, respectively, $1/4$, $1/2$, and $1/4$. This list of numbers represents the **probability distribution** over X , written p_X , which maps from the possible values of X to the non-negative reals. For a specific value x , we write $p_X(x)$, which is equal to the event probability $\Pr(X = x)$.⁴ The function p_X is called a probability **mass** function (pmf) if X is discrete; it is called a probability **density** function (pdf) if X is continuous. In either case, we have $\int_x p_X(x)dx = 1$ and $\forall x, p_X(x) \geq 0$.

Random variables can be combined into **joint probabilities**, e.g., $p_{A,B}(a, b) = \Pr(A = a \cap B = b)$. Several ideas from event probabilities carry over to probability distributions over random variables:

- We can write a **marginal probability distribution** $p_A(a) = \sum_b p_{A,B}(a, b)$.
- We can write a **conditional probability distribution** as $p_{A|B}(a | b) = \frac{p_{A,B}(a, b)}{p_B(b)}$.
- Random variables A and B are independent iff $p_{A,B}(a, b) = p_A(a) \times p_B(b)$.

Expectations

Sometimes we want the **expectation** of a function, such as $E[g(x)] = \sum_{x \in \mathcal{X}} g(x)p(x)$. Expectations are easiest to think about in terms of probability distributions over discrete events:

- If it is sunny, Marcia will eat three ice creams.
- If it is rainy, she will eat only one ice cream.
- There's a 80% chance it will be sunny.
- The expected number of ice creams she will eat is $0.8 \times 3 + 0.2 \times 1 = 2.6$.

If the random variable X is continuous, the sum becomes an integral:

$$E[g(x)] = \int_{\mathcal{X}} g(x)p(x)dx \quad (1.34)$$

For example, a fast food restaurant in Quebec has a special offer for cold days: they give a 1% discount on poutine for every degree below zero. Assuming they use a thermometer with infinite precision, the expected price would be an integral over all possible temperatures,

$$E[\text{price}(x)] = \int_{\mathcal{X}} \min(1, 1 + x) \times \text{original-price} \times p(x)dx. \quad (1.35)$$

(Careful readers will note that the restaurant will apparently pay you for taking poutine, if the temperature falls below -100 degrees celsius.)

⁴In general, capital letters (e.g., X) refer to random variables, and lower-case letters (e.g., x) refer to specific values. I will often just write $p(x)$, when the subscript is clear from context.

Modeling and estimation

Probabilistic models give us a principled way to reason about random events and random variables, and to make predictions about the future. Let's consider the coin toss example. We can model each toss as a random event, with probability θ of the event H , and probability $1 - \theta$ of the complementary event T . If we write a random variable X as the total number of heads on three coin flips, then the distribution of X depends on θ . In this case, X is distributed as a **binomial random variable**, meaning that it is drawn from a binomial distribution, with **parameters** $(\theta, N = 3)$. We write:

$$X \sim \text{Binomial}(\theta, N = 3). \quad (1.36)$$

This is a probabilistic model of X . The binomial distribution has a number of known properties that enable us to make statements about the X , such as its expected value, the likelihood that its value will fall within some interval, etc.

Now suppose that θ is unknown, but we have run an experiment, in which we executed N trials, and obtained x heads. We can **estimate** θ by the principle of **maximum likelihood**:

$$\hat{\theta} = \underset{\theta}{\operatorname{argmax}} p_X(x; \theta, N). \quad (1.37)$$

This says that our estimate $\hat{\theta}$ should be the value that maximizes the likelihood of the data we have observed. The semicolon indicates that θ and N are parameters of the probability function. The likelihood $p_X(x; \theta, N)$ can be computed from the binomial distribution,

$$p_X(x; \theta, N) = \frac{N!}{x!(N-x)!} \theta^x (1-\theta)^{N-x}. \quad (1.38)$$

This likelihood is proportional to the product of the probability of individual outcomes: for example, the sequence T, H, H, T, H would have probability $\theta^2(1-\theta)^3$. The term $\frac{N!}{x!(N-x)!}$ arises from the many possible orderings by which we could obtain x heads on N trials. This term is constant in θ , so it can be ignored.

We can maximize likelihood by taking the derivative and setting it equal to zero. In practice, we usually maximize log-likelihood, which is a monotonic function of the likeli-

(c) Jacob Eisenstein 2014-2017. Work in progress.

hood, and is easier to manipulate mathematically.

$$\ell(\theta) = x \log \theta + (N - x) \log(1 - \theta) \quad (1.39)$$

$$\frac{\partial \ell(\theta)}{\partial \theta} = \frac{x}{\theta} - \frac{N - x}{1 - \theta} \quad (1.40)$$

$$\frac{N - x}{1 - \theta} = \frac{x}{\theta} \quad (1.41)$$

$$\frac{N - x}{x} = \frac{1 - \theta}{\theta} \quad (1.42)$$

$$\frac{N}{x} - 1 = \frac{1}{\theta} - 1 \quad (1.43)$$

$$\hat{\theta} = \frac{x}{N}. \quad (1.44)$$

In this case, the maximum likelihood estimate is equal to $\frac{x}{N}$, the fraction of trials that came up heads. This intuitive solution is also known as the **relative frequency estimate**, since it is equal to the relative frequency of the outcome.

Is maximum likelihood estimation always the right choice? Suppose you conduct one trial, and get heads — would you conclude that $\theta = 1$, so this coin is guaranteed to give heads? If not, then you must have some **prior expectation** about θ . To incorporate this prior information, we can treat θ as a random variable, and use Bayes rule:

$$p(\theta | x; N) = \frac{p(x | \theta) \times p(\theta)}{p(x)} \quad (1.45)$$

$$\propto p(x | \theta) \times p(\theta) \quad (1.46)$$

$$\hat{\theta} = \operatorname{argmax}_{\theta} p(x | \theta) \times p(\theta). \quad (1.47)$$

This is the **maximum a posteriori** (MAP) estimate. Given a form for $p(\theta)$, you can derive the MAP estimate using the same approach that was used to derive the maximum likelihood estimate.

1.2 Naïve Bayes

Back to text classification, where we were left wondering how to set the weights θ . Having just reviewed basic probability, we can now take a probabilistic approach to this problem. A **Naïve Bayes** classifier constructs the weights θ and the feature function $f(x, y)$ so that the inner product $\theta \cdot f(x, y)$ is equal to the joint log-probability $\log p(x, y)$. We can then set the weights to maximize the probability of a labeled dataset, $\{x^{(i)}, y^{(i)}\}_{i \in 1 \dots N}$, where each tuple $\langle x^{(i)}, y^{(i)} \rangle$ is a labeled instance.

(c) Jacob Eisenstein 2014-2017. Work in progress.

To carry out this strategy, We first need to define the probability $p(\{\mathbf{x}^{(i)}, y^{(i)}\}_{i \in 1 \dots N})$. We will do that through a **generative model**, which describes a hypothesized stochastic process that has generated the observed data.⁵

- For each document i ,
 - draw the label $y^{(i)} \sim \text{Categorical}(\boldsymbol{\mu})$
 - draw the vector of counts $\mathbf{x}^{(i)} \mid y^{(i)} \sim \text{Multinomial}(\boldsymbol{\phi}_{y^{(i)}})$,

The first line of this generative model is “for each document i ”, which tells us to treat each document independently: the probability of the whole dataset is equal to the product of the probabilities of each individual document. The observed word counts and document labels are **independent and identically distributed (IID)**.

$$p(\{\mathbf{x}^{(i)}, y^{(i)}\}_{i \in 1 \dots N}; \boldsymbol{\mu}, \boldsymbol{\phi}) = \prod_{i=1}^N p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\mu}, \boldsymbol{\phi}) \quad (1.48)$$

This means that the words in each document are **conditionally independent** given the parameters $\boldsymbol{\mu}$ and $\boldsymbol{\phi}$.

The second line indicates $y^{(i)} \sim \text{Categorical}(\boldsymbol{\mu})$, which means that the random variable $y^{(i)}$ is a stochastic draw from a categorical distribution with **parameter** $\boldsymbol{\mu}$. A categorical distribution is just like a weighted die: $p_{\text{cat}}(y; \boldsymbol{\mu}) = \mu_y$, where μ_y is the probability of the outcome $Y = y$. For example, if $\mathcal{Y} = \{\text{positive, negative, neutral}\}$, we might have $\boldsymbol{\mu} = [0.1, 0.7, 0.2]$. We require $\sum_y \mu_y = 1$ and $\forall_y (\mu_y \geq 0)$.

The third and final line describes the how $\mathbf{x}^{(i)}$ is sampled, conditional on $y^{(i)}$. It invokes the **multinomial distribution**, which is a probability distribution over vectors of non-negative counts. The probability mass function for this distribution is:

$$p_{\text{mult}}(\mathbf{x}; \boldsymbol{\phi}) = B(\mathbf{x}) \prod_j^V \phi_j^{x_j} \quad (1.49)$$

$$B(\mathbf{x}) = \frac{\left(\sum_j^V x_j\right)!}{\prod_j^V (x_j!)} \quad (1.50)$$

As in the categorical distribution, the parameter ϕ_j can be interpreted as a probability: specifically, the probability that any given token in the document is the word j . The multinomial distribution involves a product over words, with each term in the product

⁵We’ll see a lot of different generative models in this course. They are a helpful tool because they clearly and explicitly define the assumptions that underly the form of the probability distribution. For a very readable introduction to generative models in statistics, see Blei (2014).

equal to the probability of the word ϕ_j exponentiated by the count x_j . Words that have zero count play no role in this product, because $\phi_j^0 = 1$ for all ϕ_j .

The term $B(\mathbf{x})$ doesn't depend on ϕ , and can usually be ignored. Can you see why we need this term at all?⁶ We will return to this issue shortly.

We can write $p(\mathbf{x} \mid y; \phi)$ to indicate the conditional probability of word counts \mathbf{x} given label y , with parameter ϕ , which is equal to $p_{\text{mult}}(\mathbf{x}; \phi_y)$. By specifying the multinomial distribution for $p_{\mathbf{x}|y}$, we are working with *multinomial naïve Bayes* (MNB). Why “naïve”? Because the multinomial distribution treats each word token independently: the probability mass function factorizes across the counts.⁷ We'll see this more clearly later, when we show how MNB is an example of linear classification.

Another version of Naïve Bayes

Consider a slight modification to the generative story of NB:

- For each document i
 - Draw the label $y^{(i)} \sim \text{Categorical}(\boldsymbol{\mu})$
 - For each word $n \leq D_i$
 - * Draw the word $w_{i,n} \sim \text{Categorical}(\phi_{y^{(i)}})$

This is not quite the same model as multinomial Naive Bayes (MNB): it's a product of categorical distributions over words, instead of a multinomial distribution over word counts. This means we would generate the words in order,

$$p_W(\text{multinomial}) \times p_W(\text{Naïve}) \times p_W(\text{Bayes}). \quad (1.51)$$

Formally, this is a model for the joint probability of the word *sequence* \mathbf{w} and the label y , $p(\mathbf{w}, y)$, not the joint probability of the word *counts* \mathbf{x} and the label y , $p(\mathbf{x}, y)$.

However, as a classifier, it is identical to MNB. The final probabilities are reduced by a factor corresponding to the normalization term in the multinomial, $B(\mathbf{x})$. This means that the probability for a vector of counts \mathbf{x} is larger than the probability for a list of words \mathbf{w} that induces the same counts. But this makes sense: there can be many word sequences that correspond to a single vector of counts. For example, *man bites dog* and *dog bites man* correspond to an identical count vector, $\{\text{bites} : 1, \text{dog} : 1, \text{man} : 1\}$, and the total number of word orderings for a given count vector \mathbf{x} is exactly the ratio $B(\mathbf{x}) = \frac{(\sum_j x_j)!}{\prod_j x_j!}$.

⁶Technically, a multinomial distribution requires a second parameter, the total number of counts, which in the bag-of-words representation is equal to the number of words in the document.

⁷You can plug in any probability distribution to the generative story and it will still be naïve Bayes, as long as you are making the “naïve” assumption that your features are conditionally independent, given the label. For example, a multivariate Gaussian with diagonal covariance would be naïve in exactly the same sense.

From the perspective of classification, none of this matters, because it has nothing to do with the label y or the parameters ϕ and μ . The ratio of probabilities between any two labels y_1 and y_2 will be identical in the two models, as will the maximum likelihood estimates for the parameters μ and ϕ (which are defined below).

Prediction

The Naive Bayes prediction rule is to choose the label y which maximizes $p(\mathbf{x}, y; \mu, \phi)$:

$$\hat{y} = \underset{y}{\operatorname{argmax}} p(\mathbf{x}, y; \mu, \phi) \quad (1.52)$$

$$= \underset{y}{\operatorname{argmax}} p(\mathbf{x} \mid y; \phi) p(y; \mu) \quad (1.53)$$

$$= \underset{y}{\operatorname{argmax}} \log p(\mathbf{x} \mid y; \phi) + \log p(y; \mu) \quad (1.54)$$

Converting to logarithms makes the notation easier. It doesn't change the prediction rule because the log function is monotonically increasing.

Now we can plug in the probability distributions from the generative story.

$$\log p(\mathbf{x} \mid y; \phi) + \log p(y; \mu) = \log \left[B(\mathbf{x}) \prod_j^V \phi_{y,j}^{x_j} \right] + \log \mu_y \quad (1.55)$$

$$= \log B(\mathbf{x}) + \sum_j^V x_j \log \phi_{y,j} + \log \mu_y \quad (1.56)$$

$$= \log B(\mathbf{x}) + \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y), \quad (1.57)$$

where

$$\boldsymbol{\theta} = [\boldsymbol{\theta}^{(1)\top}, \boldsymbol{\theta}^{(2)\top}, \dots, \boldsymbol{\theta}^{(K)\top}]^\top \quad (1.58)$$

$$\boldsymbol{\theta}^{(y)} = [\log \phi_{y,1}, \log \phi_{y,2}, \dots, \log \phi_{y,V}, \log \mu_y]^\top \quad (1.59)$$

The feature function $\mathbf{f}(\mathbf{x}, y)$ is a vector of V word counts and an offset, padded by zeros for the labels not equal to y (see equations 1.2-1.5, and Figure 1.1). This construction ensures that the inner product $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y)$ only activates the features whose weights are in $\boldsymbol{\theta}^{(y)}$. These features and weights are all we need to compute the joint log-probability $\log p(\mathbf{x}, y)$ for each y . This is a key point: through this notation, we have converted the problem of computing the log-likelihood for a document-label pair $\langle \mathbf{x}^{(i)}, y^{(i)} \rangle$ into the computation of a vector inner product.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Estimation

The parameters of a multinomial distribution have a simple interpretation: they are the expected frequency for each word. Based on this interpretation, it is tempting to set the parameters empirically, as

$$\phi_{y,j} = \frac{\sum_{i: y^{(i)}=y} x_{i,j}}{\sum_{j'} \sum_{i: y^{(i)}=y} x_{i,j'}} = \frac{\text{count}(y, j)}{\sum_{j'} \text{count}(y, j')} \quad (1.60)$$

This is called a **relative frequency estimator**. It can be justified more rigorously as a **maximum likelihood estimate**.

Our prediction rule in Equation 1.52 is to choose \hat{y} to maximize the joint probability $p(\mathbf{x}, y)$. Maximum likelihood estimation proposes to choose the parameters ϕ and μ in much the same way. Specifically, we want to maximize the joint log-likelihood of some **training data**: a set of annotated examples for which we observe both the text and the true label, $\{\mathbf{x}^{(i)}, y^{(i)}\}_{i \in 1 \dots N}$. Based on the generative model that we have defined, the log-likelihood is:

$$L(\phi, \mu) = \sum_i^N \log p_{\text{mult}}(\mathbf{x}_i; \phi_{y^{(i)}}) + \log p_{\text{cat}}(y^{(i)}; \mu). \quad (1.61)$$

Let's continue to focus on the parameters ϕ . Since $p(y)$ is constant in L with respect to ϕ , we can forget it for now,

$$L(\phi) = \sum_i^N \log p_{\text{mult}}(\mathbf{x}^{(i)}; \phi_{y^{(i)}}) \quad (1.62)$$

$$= \sum_i^N \log B(\mathbf{x}) \prod_j^V \phi_{y^{(i)},j}^{x_{i,j}} \quad (1.63)$$

$$= \sum_i^N \log B(\mathbf{x}) + \sum_j^V x_{i,j} \log \phi_{y^{(i)},j}, \quad (1.64)$$

where $B(\mathbf{x})$ is constant with respect to ϕ .

We would now like to optimize L , by taking derivatives with respect to ϕ . But before we can do that, we have to deal with a set of constraints:

$$\forall y, \sum_{j=1}^V \phi_{y,j} = 1 \quad (1.65)$$

We'll do this by adding a Lagrange multiplier. Solving separately for each label y , we obtain the resulting Lagrangian,

$$\ell[\phi_y] = \sum_{i: Y^{(i)}=y} \sum_j^V x_{i,j} \log \phi_{y,j} - \lambda \left(\sum_j^V \phi_{y,j} - 1 \right) \quad (1.66)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

We can now differentiate the Lagrangian with respect to the parameter of interest, setting $\frac{\partial \ell}{\partial \phi_{y,j}} = 0$,

$$0 = \sum_{i:Y^{(i)}=y} x_{i,j} / \phi_{y,j} - \lambda \quad (1.67)$$

$$\lambda \phi_{y,j} = \sum_{i:Y^{(i)}=y} x_{i,j} \quad (1.68)$$

$$\phi_{y,j} \propto \sum_{i:Y^{(i)}=y} x_{i,j} = \sum_i \delta(Y^{(i)} = y) x_{i,j}, \quad (1.69)$$

where I use two different notations for indicating the same thing: a sum over the word counts for all documents i such that the label $Y^{(i)} = y$. This gives a solution for each ϕ_y up to a constant of proportionality. Now recall the constraint $\forall y, \sum_{j=1}^V \phi_{y,j} = 1$; this constraint arises because ϕ_y represents a vector of probabilities for each word in the vocabulary. We can exploit this constraint to obtain an exact solution,

$$\phi_{y,j} = \frac{\sum_{i:Y^{(i)}=y} x_{i,j}}{\sum_{j'=1}^V \sum_{i:Y^{(i)}=y} x_{i,j'}} \quad (1.70)$$

$$= \frac{\text{count}(y, j)}{\sum_{j'=1}^V \text{count}(y, j')}. \quad (1.71)$$

This is exactly equal to the relative frequency estimator. A similar derivation gives $\mu_y \propto \sum_i \delta(Y^{(i)} = y)$, where $\delta(Y^{(i)} = y) = 1$ if $Y^{(i)} = y$ and 0 otherwise.

Smoothing and MAP estimation

If data is sparse, you may end up with values of $\phi = 0$. For example, the word *Bayesian* may have never appeared in a spam email yet, so the relative frequency estimate $\phi_{\text{SPAM}, \text{Bayesian}} = 0$. But choosing a value of 0 would allow this single feature to completely veto a label, since $\Pr(Y = \text{SPAM} \mid \mathbf{x}) = 0$ if $\mathbf{x}_{\text{Bayesian}} > 0$.

This is undesirable, because it imposes high **variance**: depending on what data happens to be in the training set, we could get vastly different classification rules. One solution is to **smooth** the probabilities, by adding “pseudo-counts” of α to each count, and then normalizing.

$$\phi_{y,j} = \frac{\alpha + \sum_{i:Y^{(i)}=y} x_j^{(i)}}{\sum_{j'=1}^V \left(\alpha + \sum_{i:Y^{(i)}=y} x_{i,j'} \right)} = \frac{\alpha + \text{count}(y, j)}{V\alpha + \sum_{j'=1}^V \text{count}(y, j')} \quad (1.72)$$

This form of smoothing is called “Laplace smoothing”, and it has a nice Bayesian justification, in which we extend the generative story to include ϕ as a random variable (rather than as a parameter). The resulting estimate is called *maximum a posteriori*, or MAP.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Smoothing reduces **variance**, but it takes us away from the maximum likelihood estimate: it imposes a **bias**. In this case, the bias points towards uniform probabilities. Machine learning theory shows that errors on heldout data can be attributed to the sum of bias and variance. Techniques for reducing variance typically increase the bias, so there is a **bias-variance tradeoff**.⁸

- Unbiased classifiers **overfit** the training data, yielding poor performance on unseen data.
- But if we set a very large smoothing value, we can **underfit** instead. In the limit of $\alpha \rightarrow \infty$, we have zero variance: it is the same classifier no matter what data we see! But the bias of such a classifier will be high.
- Navigating this tradeoff is hard. But in general, as you have more data, variance is less of a problem, so you can just go for low bias.

The Naïvety of Naïve Bayes

Naïve Bayes is simple to work with: estimation and prediction can be done in closed form, and the nice probabilistic interpretation makes it relatively easy to extend the model. But Naïve Bayes makes assumptions which seriously limit its accuracy, especially in NLP.

- The multinomial distribution assumes that each word is generated independently of all the others (conditioned on the parameter ϕ_y). Formally, we assume conditional independence:

$$p(\text{naïve, Bayes} \mid y) = p(\text{naïve} \mid y) \times p(\text{Bayes} \mid y). \quad (1.73)$$

- But this is clearly wrong, because words “travel together.” To hone your intuitions about this, try and decide whether you believe

$$p(\text{naïve Bayes}) > p(\text{naïve}) \times p(\text{Bayes}) \quad (1.74)$$

or...

$$p(\text{naïve Bayes}) < p(\text{naïve}) \times p(\text{Bayes}). \quad (1.75)$$

Apply the chain rule!

⁸The bias-variance tradeoff is covered by Murphy (2012), but see Mohri et al. (2012) for a more formal treatment of this key concept in machine learning theory.

Traffic lights Dan Klein makes this point with an example about traffic lights. In his hometown of Pittsburgh, there is a $1/7$ chance that the lights will be broken, and both lights will be red. There is a $3/7$ chance that the lights will work, and the north-south lights will be green; there is a $3/7$ chance that the lights work and the east-west lights are green.

The *prior* probability that the lights are broken is $1/7$. If they are broken, the conditional likelihood of each light being red is 1. The prior for them not being broken is $6/7$. If they are not broken, the conditional likelihood of each individual light being red is $1/2$.

Now, suppose you see that both lights are red. According to Naïve Bayes, the probability that the lights are broken is $1/7 \times 1 \times 1 = 1/7 = 4/28$. The probability that the lights are not broken is $6/7 \times 1/2 \times 1/2 = 6/28$. So according to naive Bayes, there is a 60% chance that the lights are not broken!

What went wrong? We have made an independence assumption to factor the probability $p(R, R \mid \text{not-broken}) = p_{\text{north-south}}(R \mid \text{not-broken})p_{\text{east-west}}(R \mid \text{not-broken})$. But this independence assumption is clearly incorrect, because $p(R, R \mid \text{not-broken}) = 0$.

Less Naïve Bayes? Of course we could decide not to make the naive Bayes assumption, and model $p(R, R)$ explicitly. But this idea does not scale when the feature space is large — as it often is in NLP. The number of possible feature configurations grows exponentially, so our ability to estimate accurate parameters will suffer from high variance. With an infinite amount of data, we would be okay; but we never have that. Naïve Bayes accepts some bias, because of the incorrect modeling assumption, in exchange for lower variance.

Training, testing, and tuning (development) sets

We'll soon talk about more learning algorithms, but whichever one we apply, we will want to report its accuracy. Really, this is an educated guess about how well the algorithm will do on new data in the future.

To make an estimate of the accuracy, we need to hold out a separate “test set” from the data that we use for estimation (i.e., training, learning). Otherwise, if we measure accuracy on the same data that is used for estimation, we will badly overestimate the accuracy that we are likely to get on new data.

Recall that in addition to the parameters μ and ϕ , which are learned on training data, we also have the amount of smoothing, α . This can be considered a “tuning” parameter, and it controls the tradeoff between overfitting and underfitting the training data. Where is the best position on this tradeoff curve? It's hard to tell in advance. Sometimes it is tempting to see which tuning parameter gives the best performance on the test set, and then report that performance. Resist this temptation! It will also lead to overestimating accuracy on truly unseen future data. For that reason, this is a sure way to get your research paper rejected; in a commercial setting, this mistake may cause you to promise

much higher accuracy than you can deliver. Instead, you should split off a piece of your training data, called a “development set” (or “tuning set”).

Sometimes, people average across multiple test sets and/or multiple development sets. One way to do this is to divide your data into “folds,” and allow each fold to be the development set one time. This is called **K-fold cross-validation**. In the extreme, each fold is a single data point. This is called **leave-one-out**.

Exercises

[[todo: make exercises](#)]

Chapter 2

Discriminative learning

Naïve Bayes is a simple classifier, where both the prediction rule and the learning objective are based on the joint probability of labels and base features,

$$\log p(y^{(i)}, \mathbf{x}^{(i)}) = \log p(\mathbf{x}^{(i)} | y^{(i)}) + \log p(y^{(i)}) \quad (2.1)$$

$$= \sum_j \log p(x_{i,j} | y^{(i)}) + \log p(y^{(i)}) \quad (2.2)$$

$$= \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) \quad (2.3)$$

Equation 2.2 shows the independence assumption that makes it possible to compute this joint probability: the probability of each base feature $x_{i,j}$ is mutually independent, after conditioning on the label $y^{(i)}$.

In the equations above, we define the **feature function** $\mathbf{f}(x, y)$ so that it corresponds to “bag-of-words” features. Bag-of-words features violate the assumption of conditional independence — for example, the probability that a document will contain the word *naïve* is surely higher given that it also contains the word *Bayes* — but this violation is relatively mild. However, to get really good performance on text classification and other language processing tasks, we will need to add many other types of features. Some of these features will capture parts of words, and others will capture multi-word units. For example:

- Prefixes, such as *anti-*, *im-*, and *un-*.
- Punctuation and capitalization.
- **Bigrams**, such as *not good*, *not bad*, *least terrible*, and higher-order **n-grams**.

Many of these “rich” features violate the Naïve Bayes independence assumption more severely. Consider what happens if we add feature capturing the word prefix. Under the Naïve Bayes assumption, we make the following approximation:

$$\begin{aligned} \Pr(\text{word} = \textit{impossible}, \text{prefix} = \textit{im-} | y) &\approx \Pr(\text{prefix} = \textit{im-} | y) \\ &\times \Pr(\text{word} = \textit{impossible} | y). \end{aligned} \quad (2.4)$$

To test the quality of the approximation, we can manipulate the original probability by applying the chain rule,

$$\Pr(\text{word} = \textit{impossible}, \text{prefix} = \textit{im-} \mid y) = \Pr(\text{prefix} = \textit{im-} \mid \text{word} = \textit{impossible}, y) \times \Pr(\text{word} = \textit{impossible} \mid y) \quad (2.5)$$

But $\Pr(\text{prefix} = \textit{im-} \mid \text{word} = \textit{impossible}, y) = 1$, since *im-* is guaranteed to be the prefix for the word *impossible*. Therefore,

$$\Pr(\text{word} = \textit{impossible}, \text{prefix} = \textit{im-} \mid y) \quad (2.6)$$

$$= 1 \times \Pr(\text{word} = \textit{impossible} \mid y) \\ \gg \Pr(\text{prefix} = \textit{im-} \mid y) \times \Pr(\text{word} = \textit{impossible} \mid y). \quad (2.7)$$

The final inequality is due to the fact that the probability of any given word starting with the prefix *im-* is much less than one, and it shows that Naïve Bayes will systematically underestimate the true probabilities of conjunctions of positively correlated features. To use such features, we will need learning algorithms that do not rely on an independence assumption.

2.1 Perceptron

In Naïve Bayes, the weights can be interpreted as parameters of a probabilistic model. But this model requires an independence assumption that usually does not hold, and limits our choice of features. Why not forget about probability and learn the weights in an error-driven way? The perceptron algorithm, shown in Algorithm 1, is one way to do this.¹

What the algorithm says is this: if you make a mistake, increase the weights for features which are active with the correct label $y^{(i)}$, and decrease the weights for features which are active with the guessed label \hat{y} . This is an **online learning** algorithm, since the classifier weights change after every example. This is different from Naïve Bayes, which computes corpus statistics and then sets the weights in a single operation — Naïve Bayes is a **batch learning** algorithm.²

The perceptron algorithm may seem like a cheap heuristic: Naïve Bayes has a solid foundation in probability, but now we are just adding and subtracting constants from the weights every time there is a mistake. Will this really work? In fact, there is some nice theory for the perceptron. To understand it, we must introduce the notion of **linear separability**:

¹I have been deliberately vague about the stopping criterion; this is discussed later in the chapter.

²Later in this chapter we will encounter a third class of learning algorithm, which is **iterative**. Such algorithms perform multiple updates to the weights (like perceptron), but are also **batch**, in that they have to use all the training data to compute the update. [todo: keep this?]

Algorithm 1 Perceptron learning algorithm

```

1: procedure PERCEPTRON( $\mathbf{x}^{(1:N)}, y^{(1:N)}$ )
2:    $t \leftarrow 0$ 
3:    $\boldsymbol{\theta}_t \leftarrow \mathbf{0}$ 
4:   repeat
5:      $t \leftarrow t + 1$ 
6:     Select an instance  $i$ 
7:      $\hat{y} \leftarrow \operatorname{argmax}_y \boldsymbol{\theta}_{t-1} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y)$ 
8:     if  $\hat{y} \neq y^{(i)}$  then
9:        $\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} + \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \mathbf{f}(\mathbf{x}^{(i)}, \hat{y})$ 
10:    else
11:       $\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1}$ 
12:  until tired
13:  return  $\boldsymbol{\theta}$ 

```

Definition 1 (Linear separability). *The dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_i$ is linearly separable iff there exists some weight vector $\boldsymbol{\theta}$ and some **margin** ρ such that for every instance $\langle \mathbf{x}^{(i)}, y^{(i)} \rangle$, the inner product of $\boldsymbol{\theta}$ and the feature function for the true label, $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y^{(i)})$, is at least ρ greater than inner product of $\boldsymbol{\theta}$ and the feature function for every other possible label, $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y')$.*

$$\exists \boldsymbol{\theta}, \rho > 0 : \forall \langle \mathbf{x}^{(i)}, y^{(i)} \rangle \in \mathcal{D}, \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) \geq \rho + \max_{y' \neq y^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y'). \quad (2.8)$$

Linear separability is important because of the following guarantee: if your data is linearly separable, then the perceptron algorithm will find a separator (Novikoff, 1962).³ So while the perceptron may seem heuristic, it is guaranteed to succeed, if the learning problem is easy enough.

How useful is this proof? Minsky and Papert (1969) note that the simple logical function of *exclusive-or* is not separable, and that a perceptron is therefore incapable of learning to mimic this function. But this is not just a problem for perceptron: any linear classification algorithm, including Naïve Bayes, will fail to learn this function. In natural language, we work in very high dimensional feature spaces, with thousands or millions of features. In these high-dimensional spaces, finding a separator becomes exponentially easier. Furthermore, later theoretical work showed that if the data is not separable, it is still possible to place an upper bound on the number of errors that the perceptron algorithm will make (Freund and Schapire, 1999).

³It is also possible to prove an upper bound on the number of training iterations required to find the separator. Proofs like this are part of the field of **statistical learning theory** (Mohri et al., 2012).

Averaged perceptron

The perceptron iterates over the data repeatedly — until “tired”, as described in Algorithm 1. If the data is linearly separable, it is guaranteed that the perceptron will eventually find a separator, and then we can stop. But if the data is not separable, the algorithm can *thrash* between two or more weight settings, never converging. In this case, how do we know that we can stop training, and how should we choose the final weights? An effective practical solution is to *average* the perceptron weights across all iterations.

This procedure is shown in Algorithm 2. The learning algorithm is nearly identical to the “vanilla” perceptron, but we also maintain a vector of the weight sums, \mathbf{m} . At the end of the learning procedure, we divide this sum by the total number of updates t , to compute the averaged weights, $\bar{\theta}$. These averaged weights are then used to predict the labels of new data, such as examples in the test set. The algorithm sketch indicates that we compute the average by keeping a running sum, $\mathbf{m} \leftarrow \mathbf{m} + \theta$. However, this is inefficient, because it requires $|\theta|$ operations to update the running sum. In NLP problems, $\mathbf{f}(\mathbf{x}, y)$ is typically sparse, so $|\theta| \gg |\mathbf{f}(\mathbf{x}, y)|$ for any individual (\mathbf{x}, y) . This means that the computation of the running sum will be much more expensive than the computation of the update to θ itself, which requires only $2 \times |\mathbf{f}(\mathbf{x}, y)|$ operations. One of the exercises is to sketch a more efficient algorithm for computing the averaged weights.

Even if the data is not separable, the averaged weights will eventually converge. One possible stopping criterion is to check the difference between the average weight vectors after each pass through the data: if the norm of the difference falls below some predefined threshold, we can stop iterating. Another stopping criterion is to hold out some data, and to measure the predictive accuracy on this heldout data (this is called a **development set**, and was introduced in chapter 1). When the accuracy on the heldout data starts to decrease, the learning algorithm has begun to **overfit** the training set. At this point, it is probably best to stop; this stopping criterion is known as **early stopping**.

Generalization is the ability to make good predictions on instances that are not in the training data; it can be proved that averaging improves generalization, by computing an upper bound on the generalization error (Freund and Schapire, 1999; Collins, 2002).

2.2 Loss functions and large margin classification

Naïve Bayes chooses the weights θ by maximizing the joint likelihood $p(\{\mathbf{x}^{(i)}, y^{(i)}\}_{i \in 1 \dots N})$. This is equivalent to maximizing the log-likelihood (due to the monotonicity of the log function), and also to **minimizing** the negative log-likelihood. This negative log-likelihood

(c) Jacob Eisenstein 2014-2017. Work in progress.

Algorithm 2 Averaged perceptron learning algorithm

```

1: procedure AVG-PERCEPTRON( $\mathbf{x}^{(1:N)}, y^{(1:N)}$ )
2:    $t \leftarrow 0$ 
3:    $\boldsymbol{\theta}_0 \leftarrow \mathbf{0}$ 
4:   repeat
5:      $t \leftarrow t + 1$ 
6:     Select an instance  $i$ 
7:      $\hat{y} \leftarrow \operatorname{argmax}_y \boldsymbol{\theta}_{t-1} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y)$ 
8:     if  $\hat{y} \neq y^{(i)}$  then
9:        $\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1} + \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \mathbf{f}(\mathbf{x}^{(i)}, \hat{y})$ 
10:    else
11:       $\boldsymbol{\theta}_t \leftarrow \boldsymbol{\theta}_{t-1}$ 
12:       $\mathbf{m} \leftarrow \mathbf{m} + \boldsymbol{\theta}_t$ 
13:    until tired
14:     $\bar{\boldsymbol{\theta}} \leftarrow \frac{1}{t} \mathbf{m}$ 
15:  return  $\bar{\boldsymbol{\theta}}$ 

```

can therefore be viewed as a **loss function**,

$$\log p(\mathbf{x}, \mathbf{y}; \boldsymbol{\theta}) = \sum_{i=1}^N \log p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta}) \quad (2.9)$$

$$\ell_{\text{NB}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = -\log p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta}) \quad (2.10)$$

$$\hat{\boldsymbol{\theta}} = \operatorname{argmin}_{\boldsymbol{\theta}} \sum_{i=1}^N \ell_{\text{NB}}(\boldsymbol{\theta}, \mathbf{x}^{(i)}, y^{(i)}) \quad (2.11)$$

This minimization problem is identical to the maximum-likelihood estimation problem that we solved in the previous chapter. Framing it as minimization may seem backwards, but loss functions provide a very general framework in which to compare many approaches to machine learning. For example, an alternative loss function is the **zero-one loss**,

$$\ell_{\text{zero-one}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = \begin{cases} 0, & y^{(i)} = \operatorname{argmax}_y \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}_i, y) \\ 1, & \text{otherwise} \end{cases} \quad (2.12)$$

This loss function is closely related to accuracy, which may seem ideal. But it is **non-convex**⁴ and discontinuous, which means that it is combinatorially difficult to optimize.

⁴A function f is convex iff $\alpha f(x_i) + (1 - \alpha)f(x_j) \geq f(\alpha x_i + (1 - \alpha)x_j)$, for all $\alpha \in [0, 1]$ and for all x_i and x_j on the domain of the function. Convexity implies that any local minimum is also a global minimum, and there are effective techniques for optimizing convex functions (Boyd and Vandenberghe, 2004).

The perceptron optimizes the following loss function:

$$\ell_{\text{perceptron}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = \max_y \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}), \quad (2.13)$$

which is a **hinge loss** with the hinge point at zero. When $\hat{y} = y^{(i)}$, the loss is zero; otherwise, it increases linearly with the gap between the score for the predicted label \hat{y} and the score for the true label $y^{(i)}$. To see why this is the loss function optimized by the perceptron, just take the derivative with respect to $\boldsymbol{\theta}$,

$$\frac{\partial}{\partial \boldsymbol{\theta}} \ell_{\text{perceptron}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = \mathbf{f}(\mathbf{x}^{(i)}, \hat{y}) - \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}). \quad (2.14)$$

One way to minimize our loss is to take a step of magnitude τ in the opposite direction of this gradient,

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \tau \frac{\partial}{\partial \boldsymbol{\theta}} \ell_{\text{perceptron}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = \boldsymbol{\theta} + \tau (\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \mathbf{f}(\mathbf{x}^{(i)}, \hat{y})). \quad (2.15)$$

When the step size $\tau = 1$, this is identical to the perceptron update.

This loss function has some pros and cons with respect to the joint likelihood loss implied by Naïve Bayes.

- Both ℓ_{NB} and $\ell_{\text{perceptron}}$ are convex, making them relatively easy to optimize. However, ℓ_{NB} can be optimized in closed form, while $\ell_{\text{perceptron}}$ requires iterating over the dataset multiple times.
- ℓ_{NB} can suffer **infinite** loss on a single example, which suggests it will overemphasize some examples, and underemphasize others.
- $\ell_{\text{perceptron}}$ treats all correct answers equally. Even if $\boldsymbol{\theta}$ only gives the correct answer by a tiny margin, the loss is still zero.

This last comment suggests a potential problem. Suppose a test example is very close to a training example, but not identical. If the classifier only gets the correct answer on the training example by a small margin, then it may get the test instance wrong. To formalize this intuition, let's define the **margin** as

$$\gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \max_{y \neq y^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) \quad (2.16)$$

The margin represents the separation between the score for the correct label $y^{(i)}$, and the score for the highest-scoring label. If the instance is classified incorrectly, the margin will be negative. The intuition behind **large-margin** learning algorithms is that it is not enough just to get the training data correct — we want the correct label to be separated

from the other possible labels by a comfortable margin. We can use the margin to define a convex and continuous **margin loss**,

$$\ell_{\text{margin}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = \begin{cases} 0, & \gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) \geq 1, \\ 1 - \gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}), & \text{otherwise} \end{cases} \quad (2.17)$$

Equivalently, we can write $\ell_{\text{margin}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = (1 - \gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}))_+$, where $(x)_+ = \max(0, x)$. The margin loss is zero if we have a margin of at least 1 between the score for the true label and the best-scoring alternative, which we have written \hat{y} . It is equivalent to the hinge loss defined above, but shifted to the right on the x -axis. The margin and zero-one loss functions are shown in Figure 2.1. Note that the margin loss is a convex upper bound on the zero-one loss.

Support vector machines

We can write the weight vector $\boldsymbol{\theta} = s\mathbf{u}$, where the **norm** of \mathbf{u} is equal to one, $\|\mathbf{u}\|_2 = 1$.⁵ Think of s as the magnitude and \mathbf{u} as the direction of the vector $\boldsymbol{\theta}$. If the data is separable, there are many values of s that attain zero loss. To see this, let us redefine the margin as,

$$\gamma(\boldsymbol{\theta}, \mathbf{x}^{(i)}, y^{(i)}) = \min_{y \neq y^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) \quad (2.18)$$

$$= \min_{y \neq y^{(i)}} s(\mathbf{u}^\top (\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \mathbf{f}(\mathbf{x}^{(i)}, y))). \quad (2.19)$$

Based on this definition, if the unit vector \mathbf{u}^* satisfies $\gamma(\mathbf{u}^*, \mathbf{x}^{(i)}, y^{(i)}) > 0$, then there is some smallest value s^* such that $\forall s \geq s^*, \gamma(s\mathbf{u}^*, \mathbf{x}^{(i)}, y^{(i)}) \geq 1$. Given many possible $\boldsymbol{\theta}$ that obtain zero margin loss, we may prefer the one with the smallest norm ($s = s^*$), since this entails making the least commitment to the training data. This idea underlies the **Support Vector Machine** (SVM) classifier,⁶ which, in its most basic form, solves the following optimization problem,

$$\begin{aligned} \min_{\boldsymbol{\theta}} \quad & \|\boldsymbol{\theta}\|_2^2 \\ \text{s.t.} \quad & \forall_i \ell_{\text{margin}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = 0. \end{aligned} \quad (2.20)$$

In realistic settings, we do not know whether there is any feasible solution — that is, whether there exists any $\boldsymbol{\theta}$ so that the margin loss on every training instance is zero. We therefore introduce a set of **slack variables** $\xi_i \geq 0$, which represent a sort of “fudge factor” for each instance i — instead of requiring that the loss be exactly zero, we require that it be

⁵The norm of a vector $\|\mathbf{u}\|_2$ is defined as, $\|\mathbf{u}\|_2 = \sqrt{\sum_j u_j^2}$.

⁶Instances near the margin are called **support vectors**. In some optimization methods for this model, the support vectors play an especially important role, motivating the name.

less than ξ_i . Ideally there would not be any slack, so we add the sum of the slack variables to the objective function to be minimized:

$$\begin{aligned} \min_{\boldsymbol{\theta}} \quad & \|\boldsymbol{\theta}\|_2^2 + C \sum_i \xi_i \\ \text{s.t.} \quad & \forall_i \ell_{\text{margin}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) \leq \xi_i \\ & \forall_i \xi_i \geq 0. \end{aligned} \tag{2.21}$$

Here C is a tunable parameter that controls the penalty on the slack variables. As $C \rightarrow \infty$, slack is infinitely expensive, and we can only find a solution if the data is separable. As $C \rightarrow 0$, slack becomes free, and there is a trivial solution at $\boldsymbol{\theta} = \mathbf{0}$, regardless of the data. Thus, C plays a similar role to the smoothing parameter in Naïve Bayes (§ 1.2), trading off between a close fit to the training data and better generalization. Like the smoothing parameter of Naïve Bayes, C must be set by the user, typically by maximizing performance on a heldout development set.

To solve the constrained optimization problem defined in Equation 2.21, we can use Lagrange multipliers to convert it into the unconstrained **primal form**,⁷

$$\min_{\boldsymbol{\theta}} \quad \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2 + \sum_i \ell_{\text{margin}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}), \tag{2.22}$$

where λ is a tunable parameter that can be computed from the term C in Equation 2.21. A generic way to minimize such objective functions is **gradient descent**: moving along the gradient (obtained by differentiating with respect to $\boldsymbol{\theta}$), until the gradient is equal to zero.⁸

Let us rewrite the primal form of the SVM optimization problem as follows:

$$L_{SVM} = \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2 + \sum_i^N \ell_{\text{margin}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) \tag{2.23}$$

$$= \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2 + \sum_i^N (1 - \gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}))_+ \tag{2.24}$$

$$= \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2 + \sum_i^N (1 - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \max_{y \neq y^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y))_+ \tag{2.25}$$

⁷An alternative **dual form** is used in the formulation of the kernel-based support vector machine, which supports non-linear classification. This is described briefly at the end of the chapter.

⁸Because the margin loss is not smooth, there is not a single gradient at the point at which the loss is exactly equal to zero, but rather, a **subgradient set**. However, this is a theoretical issue that poses no difficulties in practice.

Let us define the **cost** of a misclassification as,

$$c(y^{(i)}, \hat{y}) = \begin{cases} 1, & y^{(i)} \neq \hat{y} \\ 0, & \text{otherwise.} \end{cases} \quad (2.26)$$

We can then simplify Equation 2.25,

$$L_{SVM} = \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2 + \sum_i^N (\max_y (\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) + c(y^{(i)}, y)) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}))_+, \quad (2.27)$$

where we now maximize over all $y \in \mathcal{Y}$, favoring labels that are both high-scoring (as measured by $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y)$) and wrong (as measured by $c(y^{(i)}, y)$). When the highest-scoring such label is $y = y^{(i)}$, then the margin constraint is satisfied, and the loss for this instance is zero.

Then the (sub)gradient of Equation 2.27 is:

$$\frac{\partial L_{SVM}}{\partial \boldsymbol{\theta}} = \lambda \boldsymbol{\theta} + \sum_i^N \mathbf{f}(\mathbf{x}^{(i)}, \hat{y}) - \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}), \quad (2.28)$$

where $\hat{y} = \operatorname{argmax}_y \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) + c(y^{(i)}, y)$. If we were to update $\boldsymbol{\theta}$ by adding this gradient, this would be very similar to the perceptron algorithm: the only difference is the additional cost term $c(y^{(i)}, y)$, which derives from the margin constraint.

2.3 Logistic regression

Thus far, we have seen two broad classes of learning algorithms. Naïve Bayes is a probabilistic method, where learning is equivalent to estimating a joint probability distribution. Perceptron and support-vector machines are error-driven algorithms: the learning objective is closely related to the number of errors on the training data, and will be maximized when there are zero errors. Probabilistic and error-driven approaches each have advantages: probability enables us to quantify uncertainty about the predicted labels, but error-driven learning typically leads to better performance on error-based performance metrics such as accuracy.

Logistic regression combines both of these advantages: it is error-driven like the perceptron and margin-based learning algorithms, but it is probabilistic like Naïve Bayes. To understand the motivation for logistic regression, first recall that Naïve Bayes selects weights to optimize the joint probability $p(\mathbf{x}, y)$.

- We have used the chain rule to factor this joint probability as $p(\mathbf{x}, y) = p(\mathbf{x} \mid y) \times p(y)$.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- But we could equivalently choose the alternative factorization $p(\mathbf{x}, y) = p(y \mid \mathbf{x}) \times p(\mathbf{x})$.

In classification, we always know \mathbf{x} : these are the base features from which we predict y . So there is no need to model $p(\mathbf{x})$; we really care only about the **conditional probability** $p(y \mid \mathbf{x})$ — sometimes called the **likelihood**. Logistic regression defines this probability directly, in terms of the features $\mathbf{f}(\mathbf{x}, y)$ and the weights $\boldsymbol{\theta}$.

We can think of $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y)$ as a scoring function for the compatibility of the base features \mathbf{x} and the label y . This function is an unconstrained scalar; we would like to convert it to a probability. To do this, we first **exponentiate**, obtaining $\exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y))$, which is guaranteed to be non-negative. Next, we need to **normalize**, dividing over all possible labels $y' \in \mathcal{Y}$. The resulting conditional probability is defined as,

$$p(y \mid \mathbf{x}) = \frac{\exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y))}{\sum_{y' \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y'))}. \quad (2.29)$$

Given a dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_i$, the maximum-likelihood estimator for $\boldsymbol{\theta}$ is obtained by maximizing,

$$L(\boldsymbol{\theta}) = \log p(\mathbf{y}^{(1:N)} \mid \mathbf{x}^{(1:N)}; \boldsymbol{\theta}) \quad (2.30)$$

$$= \log \prod_{i=1}^N p(y^{(i)} \mid \mathbf{x}^{(i)}; \boldsymbol{\theta}) \quad (2.31)$$

$$= \sum_{i=1}^N \log p(y^{(i)} \mid \mathbf{x}^{(i)}; \boldsymbol{\theta}) \quad (2.32)$$

$$= \sum_{i=1}^N \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \log \sum_{y' \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y')). \quad (2.33)$$

The final line is obtained by plugging in Equation 2.29 and taking the logarithm.^{9,10} Inside the sum, we have the (additive inverse of the) **logistic loss**.

- In binary classification, we can write this as

$$\ell_{\text{logistic}}(\boldsymbol{\theta}; \mathbf{x}_i, y^{(i)}) = -(y^{(i)} \boldsymbol{\theta}^\top \mathbf{x}_i - \log(1 + \exp(\boldsymbol{\theta} \cdot \mathbf{x}_i))) \quad (2.34)$$

⁹Any reasonable base will work; if it is important to you to know which one to choose, then I suggest using base 2 if you are a computer scientist, and base e otherwise.

¹⁰The log-sum-exp term is very common in machine learning. It is numerically unstable because it will underflow if the inner product is small, and overflow if the inner product is large. Scientific computing libraries usually contain special functions for computing `logsumexp`, but with some thought, you should be able to see how to create an implementation that is numerically stable.

- In multi-class classification, we have,

$$\ell_{\text{logistic}}(\boldsymbol{\theta}; \mathbf{x}_i, y^{(i)}) = -(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)})) - \log \sum_{y' \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y')) \quad (2.35)$$




Figure 2.1: Margin, zero-one, and logistic loss functions

The logistic loss is shown in Figure 2.1. Note that logistic loss is also an upper bound on the perceptron loss. A key difference from the perceptron and hinge losses is that logistic loss is never exactly zero. This means that the objective function can always be improved by choosing the correct label with more confidence.

Regularization

As with the margin-based algorithms described in § 2.2, we can obtain better generalization by penalizing the norm of $\boldsymbol{\theta}$, by adding a term of $\frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2$ to the minimization objective. This is called L_2 regularization, because it includes the L_2 norm. It can be viewed as placing a zero-mean Gaussian prior distribution on each term of $\boldsymbol{\theta}$, because the log-likelihood under a zero-mean Gaussian is,

$$\log N(\theta_j; 0, \sigma^2) \propto -\frac{1}{2\sigma^2} \theta_j^2, \quad (2.36)$$

so that $\lambda = \frac{1}{\sigma^2}$.

The effect of this regularizer will cause the estimator to trade off conditional likelihood on the training data for a smaller norm of the weights, and this can help to prevent overfitting. Indeed, regularization is generally considered to be essential to estimating high-dimensional models, as we typically do in NLP. To see why, consider what would happen to the unregularized weight for a base feature j that was active in only one instance $\mathbf{x}^{(i)}$: the conditional likelihood could always be improved by increasing the weight

for this feature, so that $\theta_{(j,y^{(i)})} \rightarrow \infty$ and $\theta_{(j,\tilde{y} \neq y^{(i)})} \rightarrow -\infty$, where (j, y) indicates the index of feature associated with $x_{i,j}$ and label y in $\mathbf{f}(\mathbf{x}^{(i)}, y)$.

Gradients

We will optimize θ through gradient descent. Specific algorithms are described in § 2.4, but because the gradient of the logistic regression objective is illustrative, it is worth working out in detail. Let us begin with the logistic loss on a single example,

$$\ell(\theta; \mathbf{x}_i, y^{(i)}) = -(\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \log \sum_{y' \in \mathcal{Y}} \exp(\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y'))) \quad (2.37)$$

$$\frac{\partial \ell}{\partial \theta} = -\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \frac{1}{\sum_{y'' \in \mathcal{Y}} \exp(\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y''))} \times \sum_{y'} \exp(\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y')) \times \mathbf{f}(\mathbf{x}^{(i)}, y') \quad (2.38)$$

$$= -\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \sum_{y'} \frac{\exp(\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y'))}{\sum_{y'' \in \mathcal{Y}} \exp(\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y''))} \times \mathbf{f}(\mathbf{x}^{(i)}, y') \quad (2.39)$$

$$= -\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \sum_{y'} p(y' | \mathbf{x}^{(i)}; \theta) \times \mathbf{f}(\mathbf{x}^{(i)}, y') \quad (2.40)$$

$$= -\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + E_{y|\mathbf{x}}[\mathbf{f}(\mathbf{x}^{(i)}, y)], \quad (2.41)$$

where the final step employs the definition of an expectation (§ 1.1). The gradient thus has the pleasing interpretation as the difference between the observed feature counts $\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)})$ and the expected counts under the current model, $E_{y|\mathbf{x}}[\mathbf{f}(\mathbf{x}^{(i)}, y)]$. When these two count vectors are equal for a single example, there is nothing more to learn from this example; when they are equal in sum over the entire dataset, there is nothing more to learn from the dataset as a whole.

As we will see shortly, a simple online approach to gradient-based optimization is to take a step along the gradient. In (unregularized) logistic regression, this gradient-based optimization is a soft version of the perceptron. Put another way, in the case that $p(y | \mathbf{x})$ is a delta function, $p(y | \mathbf{x}) = \delta(y = \hat{y})$, then the gradient step is exactly equal to the perceptron update.

If we add a regularizer $\frac{\lambda}{2} \|\theta\|_2^2$, then this contributes $\lambda\theta$ to the overall gradient:

$$L = \frac{\lambda}{2} \|\theta\|_2^2 - \sum_{i=1}^N \left(\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \log \sum_{y' \in \mathcal{Y}} \exp \theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y') \right) \quad (2.42)$$

$$\frac{\partial L}{\partial \theta} = \lambda\theta - \sum_{i=1}^N \left(\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - E_{y|\mathbf{x}}[\mathbf{f}(\mathbf{x}^{(i)}, y)] \right) \quad (2.43)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

2.4 Optimization

In Naïve Bayes, the gradient on the joint likelihood led us to a closed form solution for the parameters θ ; in passive-aggressive, we obtained a solution for each individual update from a constrained optimization problem. In logistic regression and support vector machines (SVM), we have objective functions L .

- In logistic regression, L corresponds to the regularized negative log-likelihood,

$$L_{LR} = \frac{\lambda}{2} \|\theta\|_2^2 - \sum_i \left(\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \log \sum_y \exp(\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y)) \right) \quad (2.44)$$

- In the support vector machine, L corresponds to the “primal form”,

$$L_{SVM} = \frac{\lambda}{2} \|\theta\|_2^2 + \sum_i^N (\max_y (\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) + c(y^{(i)}, y)) - \theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}))_+, \quad (2.45)$$

In both cases, the objective is convex, and there are many efficient algorithms for optimizing convex functions (Boyd and Vandenberghe, 2004). Most algorithms are based on the **gradient** $\frac{\partial L}{\partial \theta}$, or on the subgradients, in the case of non-smooth objectives in which the gradient is not unique. This section will present the most frequently-used optimization algorithms, focusing on logistic regression. However, these algorithms can also be applied to the support vector machine objective with minimal modification.

Batch optimization

In batch optimization, all the data is kept in memory and iterated over many times. The logistic loss is smooth and convex, so we can find the global optimum using gradient descent,

$$\theta_{t+1} \leftarrow \theta_t - \eta_t \frac{\partial L}{\partial \theta}, \quad (2.46)$$

where $\frac{\partial L}{\partial \theta}$ is the gradient computed over the entire training set, and η_t is some **step size**. In practice, this can be very slow to converge, as the gradient can become infinitesimally small. Second-order (Newton) optimization obtains much better convergence rates by incorporating the inverse of the Hessian matrix,

$$H_{i,j} = \frac{\partial^2}{\partial w_i \partial w_j} L. \quad (2.47)$$

Unfortunately, in NLP problems, the Hessian matrix (which is quadratic in the number of parameters) is usually too big to deal with. A typical solution is to approximate the

Hessian matrix via a **quasi-Newton optimization** technique, such as L-BFGS (Liu and Nocedal, 1989).¹¹ Quasi-Newton optimization packages are available in many scientific computing environments, and for most types of NLP practice and research, it is okay to treat them as black boxes. You will typically pass in a pointer to a function that computes the likelihood and gradient, and the solver will return a set of weights.

Online optimization

In online optimization, you consider one example (or a “mini-batch” of a few examples) at a time. **Stochastic gradient descent** (SGD) makes a stochastic online approximation to the overall gradient. Here is the SGD update for logistic regression:

$$\boldsymbol{\theta}^{(t+1)} \leftarrow \boldsymbol{\theta}^{(t)} - \eta_t \frac{\partial L_{LR}}{\partial \boldsymbol{\theta}} \quad (2.48)$$

$$= \boldsymbol{\theta}^{(t)} - \eta_t \left(\lambda \boldsymbol{\theta}^{(t)} - \sum_i^N \left(\mathbf{f}(\mathbf{x}_i, y^{(i)}) - E_{y|\mathbf{x}}[\mathbf{f}(\mathbf{x}_i, y)] \right) \right) \quad (2.49)$$

$$= (1 - \lambda \eta_t) \boldsymbol{\theta}^{(t)} + \eta_t \left(\sum_i^N \mathbf{f}(\mathbf{x}_i, y^{(i)}) - E_{y|\mathbf{x}}[\mathbf{f}(\mathbf{x}_i, y)] \right) \quad (2.50)$$

$$\approx (1 - \lambda \eta_t) \boldsymbol{\theta}^{(t)} + N \eta_t \left(\mathbf{f}(\mathbf{x}_{i(t)}, y_{i(t)}) - E_{y|\mathbf{x}}[\mathbf{f}(\mathbf{x}_{i(t)}, y)] \right) \quad (2.51)$$

where η_t is the **step size** at iteration t , and $\langle \mathbf{x}_{i(t)}, y_{i(t)} \rangle$ is an instance that is *randomly sampled* at iteration t . We can obtain a more compact form for SGD by folding the constant N into η_t and λ , so that $\tilde{\eta}_t = N \eta_t$ and $\tilde{\lambda} = \frac{\lambda}{N}$. This yields the form shown in Algorithm 3. A similar online algorithm can be derived for the SVM objective, using the subgradient in Equation 2.28.

Algorithm 3 Stochastic gradient descent for logistic regression

```

1: procedure SGD( $\mathbf{x}^{(1:N)}, y^{(1:N)}, \eta, \lambda$ )
2:    $t \leftarrow 1$ 
3:   repeat
4:     Select an instance  $i$ 
5:      $\boldsymbol{\theta}^{(t+1)} \leftarrow (1 - \tilde{\lambda} \tilde{\eta}_t) \boldsymbol{\theta}^{(t)} + \tilde{\eta}_t (\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - E_{y|\mathbf{x}}[\mathbf{f}(\mathbf{x}^{(i)}, y)])$ 
6:      $t \leftarrow t + 1$ 
7:   until tired
```

As above, the expectation is equal to a weighted sum over the labels,

$$E_{y|\mathbf{x}}[\mathbf{f}(\mathbf{x}^{(i)}, y)] = \sum_{y' \in \mathcal{Y}} p(y' | \mathbf{x}^{(i)}; \boldsymbol{\theta}) \mathbf{f}(\mathbf{x}^{(i)}, y'). \quad (2.52)$$

¹¹You can remember the order of the letters as “Large Big Friendly Giants.” Does this help you?

Again, note how similar this update is to the perceptron.

The theoretical foundation for SGD assumes that each training instance is randomly sampled (thus the name “stochastic”), but in practice, it is typical to stream through the data sequentially. It is often useful to select not a single instance, but a **mini-batch** of K instances. In this case, we would scale η_t and λ by $\frac{N}{K}$. The gradients over mini-batches will be lower variance approximations of the true gradient, and it is possible to parallelize the computation of the gradient for each instance in the mini-batch.

A key question for SGD is how to set the learning rates η_t . It can be proven that SGD will converge if $\eta_t = \eta_0 t^{-\alpha}$ for $\alpha \in [1, 2]$; however, convergence may be very slow. In practice, η_t may also be fixed to a small constant, like 10^{-3} . In either case, it is typical to try a set of different values, and see which minimizes the objective L most quickly. For more on stochastic gradient descent, as applied to a number of different learning algorithms, see (Zhang, 2004) and (Bottou, 1998). Murphy (2012) traces SGD to Nemirovski and Yudin (1978).

*AdaGrad

There are a number of ways to improve on stochastic gradient descent (Bottou et al., 2016). For NLP applications, a popular choice is use an **adaptive** step size, which can be different for every feature (Duchi et al., 2011). Features that occur frequently are likely to be updated frequently, so it is best to use a small step size; rare features will be updated infrequently, so it is better to take larger steps. The **AdaGrad** (adaptive gradient) algorithm achieves this behavior by storing the sum of the squares of the gradients for each feature, and rescaling the learning rate by its inverse:

$$\mathbf{g}_t = \lambda \boldsymbol{\theta} - \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \sum_{y' \in \mathcal{Y}} p(y' | \mathbf{x}^{(i)}) \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) \quad (2.53)$$

$$\theta_j^{(t+1)} \leftarrow \theta_j^{(t)} - \frac{\eta}{\sqrt{\sum_{t'=1}^t g_{t,j}^2}} g_{t,j}, \quad (2.54)$$

where j iterates over features in $\mathbf{f}(\mathbf{x}, y)$. AdaGrad seems to require less careful tuning of η , and Dyer (2014) reports that $\eta = 1$ works for a wide range of problems.

2.5 *Additional topics in classification

Passive-aggressive

In online learning, rather than seeking the feasible $\boldsymbol{\theta}$ with the smallest norm, we might instead prefer to make the smallest magnitude **change** to $\boldsymbol{\theta}$, while meeting the hinge loss constraint for instance $\langle \mathbf{x}^{(i)}, y^{(i)} \rangle$. Specifically, at each step t , we solve the following opti-

(c) Jacob Eisenstein 2014-2017. Work in progress.

mization problem:

$$\begin{aligned} \min w. \quad & \frac{1}{2} \|\boldsymbol{\theta} - \boldsymbol{\theta}_t\|^2 + C\xi_t \\ \text{s.t.} \quad & \ell_{\text{hinge}}(\boldsymbol{\theta}; \mathbf{x}_i, y^{(i)}) \leq \xi_t, \xi_t \geq 0 \end{aligned} \quad (2.55)$$

By forming another Lagrangian, it is possible to show that the solution to Equation 2.55 is,

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \tau_t (\mathbf{f}(y^{(i)}, \mathbf{x}^{(i)}) - \mathbf{f}(\hat{y}, \mathbf{x}^{(i)})) \quad (2.56)$$

$$\tau_t = \min \left(C, \frac{\ell(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)})}{\|\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \mathbf{f}(\mathbf{x}^{(i)}, \hat{y})\|^2} \right), \quad (2.57)$$

This algorithm is called **Passive-Aggressive** (PA; Crammer et al., 2006), because it is passive when the margin constraint is satisfied, but it aggressively changes the weights to satisfy the constraints if necessary.¹² PA is error-driven like the perceptron, and the update is nearly identical: the only difference is the learning rate τ_t , which depends on the amount of loss incurred by instance i , the norm of the difference in feature vectors between the predicted and correct labels, and the hyperparameter C , which places an upper bound on the step size. As with the perceptron, it is possible to apply weight averaging to PA, which can improve generalization. PA allows more explicit control than the Averaged Perceptron, due to the C parameter: when C is small, we make very conservative adjustments to $\boldsymbol{\theta}$ from each instance, because the slack variables aren't very expensive; when C is large, we make large adjustments to avoid using the slack variables.

Other regularizers

In Equation 2.42, we proposed to **regularize** the logistic regression estimator by penalizing the squared L_2 norm, $\|\boldsymbol{\theta}\|_2^2$. However, this is not the only way to penalize large weights; we might prefer some other norm, such as $L_0 = \|\boldsymbol{\theta}\|_0 = \sum_j \delta(\theta_j \neq 0)$, which applies a constant penalty for each non-zero weight. This norm can be thought of as a form of **feature selection**: optimizing the L_0 -regularized conditional likelihood is equivalent to trading off the log-likelihood against the number of active features. Reducing the number of active features is desirable because the resulting model will be fast, low-memory, and should generalize well, since features that are not very helpful will be pruned away. Unfortunately, the L_0 norm is non-convex and non-differentiable; optimization under L_0 regularization is NP-hard, meaning that it can be solved efficiently only if P=NP (Ge et al., 2011).

A useful alternative is the L_1 norm, which is equal to the sum of the absolute values of the weights, $\|\boldsymbol{\theta}\|_1 = \sum_j |\theta_j|$. The L_1 norm is convex, and can be used as an approximation

¹²A related algorithm without slack variables is called MIRA, for Margin-Infused Relaxed Algorithm (Crammer and Singer, 2003).

to L_0 (Tibshirani, 1996). Moreover, the L_1 norm also performs feature selection, by driving many of the coefficients to zero; it is therefore known as a **sparsity inducing regularizer**. Gao et al. (2007) compare L_1 and L_2 regularization on a suite of NLP problems, finding that L_1 regularization generally gives similar test set accuracy to L_2 regularization, but that L_1 regularization produces models that are between ten and fifty times smaller, because more than 90% of the feature weights are set to zero.

The L_1 norm does not have a gradient at $\theta_j = 0$, so we must instead optimize the L_1 -regularized objective using **subgradient** methods. The associated stochastic subgradient descent algorithms are only somewhat more complex than conventional SGD; Sra et al. (2012) survey approaches for estimation under L_1 and other regularizers.

Other views of logistic regression

Logistic regression is so named because in the binary case where $y \in \{0, 1\}$, we are performing a regression of x against y , after passing the inner product $\theta \cdot x$ through a logistic transformation to obtain a probability. However, it goes by many other names:

- Logistic regression is also called **maximum conditional likelihood** (MCL), because it is based on maximizing the conditional likelihood $p(y | x)$.
- Logistic regression can be viewed as part of a larger family of **generalized linear models** (GLMs), which include other “link functions,” such as the probit function. If you use the R software environment, you may be familiar with `glmnet`, a widely-used package for estimating GLMs.
- In the neural networks literature, the multivariate analogue of the logistic transformation is sometimes called a **softmax** layer, because it “softly” identifies the label y that maximizes the activation function $\theta \cdot f(x, y)$.

In the early NLP literature, logistic regression is frequently called **maximum entropy** (Berger et al., 1996). This is due to an alternative formulation, which tries to find the maximum entropy probability function that satisfies moment-matching constraints. The moment matching constraints specify that the empirical counts of each label-feature pair should match the expected counts:

$$\forall j, \sum_{i=1}^N f_j(\mathbf{x}^{(i)}, y^{(i)}) = \sum_{i=1}^N \sum_{y \in \mathcal{Y}} p(y | \mathbf{x}^{(i)}; \theta) f_j(\mathbf{x}^{(i)}, y) \quad (2.58)$$

Note that this constraint will be met exactly when the derivative of the likelihood function (Equation 2.41) is equal to zero. However, this constraint can be met for many values of θ , so which should we choose?

(c) Jacob Eisenstein 2014-2017. Work in progress.

The **entropy** of the conditional likelihood $p_{y|x}$ is,

$$H(p_{y|x}) = - \sum_{\mathbf{x} \in \mathcal{X}} p_{\mathbf{x}}(\mathbf{x}) \sum_{y \in \mathcal{Y}} p_{y|x}(y | \mathbf{x}) \log p_{y|x}(y | \mathbf{x}), \quad (2.59)$$

where $p_{\mathbf{x}}(\mathbf{x})$ is the probability of observing the base features \mathbf{x} . We compute an empirical estimate of the entropy by summing over all the instances in the training set,

$$\tilde{H}(p_{y|x}) = - \frac{1}{N} \sum_i \sum_{y \in \mathcal{Y}} p_{y|x}(y | \mathbf{x}^{(i)}) \log p_{y|x}(y | \mathbf{x}^{(i)}). \quad (2.60)$$

If the entropy is large, the likelihood function is smooth across possible values of y ; if it is small, the likelihood function is sharply peaked at some preferred value; in the limiting case, the entropy is zero if $p(y | \mathbf{x}) = 1$ for some y . By saying we want a maximum-entropy classifier, we are saying we want to make the weakest commitments possible, while satisfying the moment-matching constraints from Equation 2.58. The solution to this constrained optimization problem is identical to the maximum conditional likelihood (logistic-loss) formulation we considered in the previous section. This view of logistic regression is arguably a little dated, but it is useful to understand, especially when reading classic papers from the 1990s. For a tutorial on maximum entropy, see <http://www.cs.cmu.edu/afs/cs/user/aberger/www/html/tutorial/tutorial.html>.

2.6 Summary of learning algorithms

Having seen several learning algorithms, it is natural to ask which is best in various situations.

Naïve Bayes *Pros*: easy to implement; estimation is very fast, requiring only a single pass over the data; assigns probabilities to predicted labels; controls overfitting with smoothing parameter. *Cons*: the joint likelihood is arguably the wrong objective to optimize; often has poor accuracy, especially with correlated features.

Perceptron and PA *Pros*: easy to implement; online learning means it is not necessary to store all data in memory; error-driven learning means that accuracy is typically high, especially after averaging. *Cons*: not probabilistic, which can be bad in pipeline architectures, when the output of one system becomes the input for another; non-averaged perceptron performs poorly if data is not separable; hard to know when to stop learning; lack of margin can lead to overfitting.

Support vector machine *Pros*: optimizes an error-based metric, usually resulting in high accuracy; overfitting is controlled by a regularization parameter. *Cons*: not probabilistic.

Logistic regression *Pros*: error-driven and probabilistic; overfitting is controlled by a regularization parameter. *Cons*: batch learning requires black-box optimization; logistic loss sometimes gives lower accuracy than hinge loss, due to overtraining on correctly-labeled examples.

Table 2.1 summarizes some properties of Naïve Bayes, perceptron, SVM, and logistic regression. In non-probabilistic settings, I usually reach for averaged perceptron first if I am coding from scratch. If probabilities are necessary, I use logistic regression.

What about non-linear classification?

The feature spaces that we consider in NLP are usually huge, so non-linear classification can be quite difficult. Furthermore, when the number of features is much larger than the number of instances, it is usually possible to learn a linear classifier that perfectly classifies the training data. This makes selecting an appropriate **non-linear** classifier especially difficult. Nonetheless, there are some approaches to non-linear classification in NLP:

- The simplest approach is to define $f(x, y)$ to contain conjunctions or other non-linear combinations of the base features in x . For example, a bigram feature such as $\langle \text{coffee house} \rangle$ will not fire unless both base features $\langle \text{coffee} \rangle$ and $\langle \text{house} \rangle$ also fire. More generally, we can define non-linear transformations such as the element-wise product $x \odot x$ and the cross-product $x \otimes x$.
- **Kernel-based learning** is based on similarity between instances; it can be seen as a generalization of k -**nearest-neighbors**, which classifies instances by considering the label of the k most similar instances in the training set (Hastie et al., 2009). The resulting decision boundary will be non-linear in general. Kernel methods are often used in combination with the support vector machine, which has a **dual form** in which kernel functions can be inserted in place of inner products on the feature vectors.

Kernel functions can be designed to compute the similarity between structured objects, such as strings, bags-of-words, sequences, trees, and general graphs. Such methods will be discussed briefly in chapter 17.

- Boosting (Freund et al., 1999) and decision tree algorithms (Schmid, 1994) learn non-linear conjunctions of features. These methods sometimes are used less frequently in contemporary research, especially as the field increasingly emphasizes big data and simple classifiers.
- More recent work has shown how **deep learning** can perform non-linear classification, by passing the inputs through a series of non-linear transformations. Each of these transformations is learned in a supervised fashion, by **backpropagating** from a loss function. For document classification, convolutional neural networks are a

popular approach, because of their ability to capture multi-word units. These methods will be reviewed in chapter 20; surveys are offered by Goldberg (2015) and Cho (2015).

	Naive Bayes	Logistic Regression	Perceptron	SVM
Objective	Joint likelihood	Conditional likelihood	Hinge loss	Margin loss
estimation	$\max \sum_i \log \mathbf{P}(\mathbf{x}^{(i)}, y^{(i)})$	$\max \sum_i \log \mathbf{P}(y^{(i)} \mathbf{x}^{(i)})$	$\min \sum_i \delta(y^{(i)}, \hat{y})$	$\sum_i [1 - \gamma(\langle \boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)} \rangle)]_+$
tuning	$\theta_{ij} = \frac{c(x_i, y=j) + \alpha}{c(y=j) + \mathcal{V} \alpha}$	$\frac{\partial \mathcal{L}}{\partial \boldsymbol{\theta}} = \sum_i \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - E[\mathbf{f}(\mathbf{x}^{(i)}, y)]$	$\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} + \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \mathbf{f}(\mathbf{x}^{(i)}, \hat{y})$	$\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} + \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \mathbf{f}(\mathbf{x}^{(i)}, \hat{y})$
complexity	smoothing α	regularizer $\lambda \ \boldsymbol{\theta}\ _2^2$	weight averaging	slack penalty C , or regularizer λ
easy?	$\mathcal{O}(N \mathcal{V})$	$\mathcal{O}(NT \mathcal{V})$	$\mathcal{O}(NT \mathcal{V})$	$\mathcal{O}(NT \mathcal{V})$
probabilities?	very	not really	yes	yes
features?	yes	yes	no	no
	no	yes	yes	yes

Table 2.1: Comparison of classifiers. N = number of examples, $|\mathcal{V}|$ = number of features, T = number of training iterations.

Exercises

1. As noted in the discussion of averaged perceptron in § 2.1, the computation of the running sum $\mathbf{m} \leftarrow \mathbf{m} + \boldsymbol{\theta}$ is unnecessarily expensive, requiring $K \times |\mathcal{V}|$ operations. Give an alternative way to compute the averaged weights $\bar{\boldsymbol{\theta}}$, with complexity that is independent of $|\mathcal{V}|$ and linear in the sum of feature sizes $\sum_{i=1}^N |\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)})|$.
2. **[todo: reconcile with notation in this chapter]** Suppose you have two datasets $D_1 = \{x_i^{(1)}, y_i^{(1)}\}_{i \in 1 \dots N_1}$ and $D_2 = \{x_j^{(2)}, y_j^{(2)}\}_{j \in 1 \dots N_2}$, with each $y \in \{-1, 1\}$, and each $x \in \mathbb{R}^P$.
 - Let $w^{(1)}$ be the unregularized logistic regression (LR) coefficients from training on dataset D_1 , under the model, $P(y \mid x; w) = \sigma(y(x \cdot w))$, with σ indicating the sigmoid function and $x \cdot w$ indicating the dot product of the features x and the coefficients w .
 - Let $w^{(2)}$ be the unregularized LR coefficients (same model) from training on dataset D_2 .
 - Let w^* be the unregularized LR coefficients from training on the combined dataset $D_1 \cup D_2$.

Under these conditions, prove that for any feature n ,

$$w_n^* \geq \min(w_n^{(1)}, w_n^{(2)})$$

$$w_n^* \leq \max(w_n^{(1)}, w_n^{(2)}).$$

Chapter 3

Linguistic applications of classification

Having learned some techniques for classification, we will now see how they can be applied to some classical problems in natural language technology. Later in this chapter, we discuss some of the design decisions involved in text classification, as well as evaluation practices.

3.1 Sentiment and opinion analysis

A popular application of text classification is to automatically determine the **sentiment** or **opinion polarity** of documents such as product reviews and social media posts. For example, marketers are interested to know how people respond to advertisements, services, and products (Hu and Liu, 2004); social scientists are interested in how emotions are affected by phenomena such as the weather (Hannak et al., 2012), and how both opinions and emotions spread over social networks (Coviello et al., 2014; Miller et al., 2011). In the field of **digital humanities**, literary scholars track plot structures through the flow of sentiment across a novel (Jockers, 2015). A comprehensive analysis of this broad literature is beyond the scope of this chapter, but see survey manuscripts by Pang and Lee (2008) and Liu (2015).

Sentiment analysis can be framed as a fairly direct application of document classification, assuming reliable labels can be obtained. In the simplest case, sentiment analysis can be treated as a two or three-class problem, with sentiments of POSITIVE, NEGATIVE, and possibly NEUTRAL. Such annotations could be annotated by hand, or obtained automatically through a variety of means:

- Tweets containing happy emoticons can be marked as positive, sad emoticons as negative (Read, 2005; Pak and Paroubek, 2010).

- Reviews with four or more stars can be marked as positive, two or fewer stars as negative (Pang et al., 2002).
- Statements from politicians who are voting for a given bill are marked as positive (towards that bill); statements from politicians voting against the bill are marked as negative (Thomas et al., 2006).

The bag-of-words model is a good fit for sentiment analysis at the document level: if the document is long enough, we would expect the words associated with its true sentiment to overwhelm the others. Indeed, **lexicon-based sentiment analysis** avoids machine learning altogether, and classifies documents by counting words against positive and negative sentiment word lists (Taboada et al., 2011).

The problem becomes more tricky for short documents, such as single-sentence reviews or social media posts. In these documents, linguistic issues like **negation** and **irrealis** (Polanyi and Zaenen, 2006) — events that are hypothetical or otherwise non-factual — can make bag-of-words classification ineffective. Consider the following examples:

- (3.1) *That's not bad for the first day.*
 (3.2) *This is not the worst thing that can happen.*
 (3.3) *It would be nice if you acted like you understood.*
 (3.4) *There is no reason at all to believe that the polluters are suddenly going to become reasonable.* (Wilson et al., 2005)
 (3.5) *This film should be brilliant. The actors are first grade. Stallone plays a happy, wonderful man. His sweet wife is beautiful and adores him. He has a fascinating gift for living life fully. It sounds like a great plot, **however**, the film is a failure.* (Pang et al., 2002)

A minimal solution is to move from a bag-of-words model to a bag-of-**bigrams** model, where each base feature is a pair of adjacent words, e.g.,

$$\langle \text{that's}, \text{not} \rangle, \langle \text{not}, \text{bad} \rangle, \langle \text{bad}, \text{for} \rangle, \dots \quad (3.1)$$

Bigrams can handle relatively straightforward cases, such as when an adjective is immediately negated; trigrams would be required to extend to larger contexts (e.g., *not the worst*). But it should be clear that this approach will not scale to the more complex examples, such as (3.4) and (3.5). More sophisticated solutions try to account for the syntactic structure of the sentence (Wilson et al., 2005; Socher et al., 2013b), or apply more complex classifiers such as **convolutional neural networks** (Kim, 2014), which are described in chapter 20.

Related problems

Subjectivity Closely related to sentiment analysis is **subjectivity detection**, which requires identifying the parts of a text that express subjective opinions, as well as other non-factual content such speculation and hypotheticals (Riloff and Wiebe, 2003). This can be

(c) Jacob Eisenstein 2014-2017. Work in progress.

done by treating each sentence as a separate document, and then applying a bag-of-words classifier: indeed, Pang and Lee (2004) do exactly this, using a training set consisting of (mostly) subjective sentences gathered from movie reviews, and (mostly) objective sentences gathered from plot descriptions. They augment this bag-of-words model with a graph-based algorithm that encourages nearby sentences to have the same subjectivity label.

Stance classification In debates, each participant takes a side: for example, advocating for or against adopting a vegetarian lifestyle or mandating free college education. The problem of stance classification involves identifying an author’s position from the text of the argument. In some cases, there is training data available for each position, so that standard document classification techniques can be employed. In other cases, it suffices to classify each document as whether it is in support or opposition of the argument advanced by a previous document (Anand et al., 2011). In the most challenging case, there is no labeled data for any of the stances, so the only possibility is group documents that advocate the same position (Somasundaran and Wiebe, 2009). This is a form of unsupervised learning, and will be discussed in chapter 4.

Targeted and aspect based sentiment analysis The expression of sentiment is often more nuanced than a simple binary label. Consider the following examples:

(3.6) *The vodka was good, but the meat was rotten.*

(3.7) *Go to Heaven for the climate, Hell for the company.* – Mark Twain

The author of (3.6) displays a mixed overall sentiment: positive towards some entities (e.g., *the vodka*), negative towards others (e.g., *the meat*). A more ambitious goal for sentiment analysis is to determine the sentiment towards specific entities, or towards aspects of those entities. For example, **targeted sentiment analysis** seeks to identify the writer’s sentiment towards specific entities mentioned in the document (Jiang et al., 2011). This requires identifying the entities and linking them to specific sentiment words — much more than we can do with the classification-based approaches discussed thus far. For example, Kim and Hovy (2006) analyze sentence-internal structure to determine the topic of each sentiment expression.

Aspect-based opinion mining seeks to identify the sentiment of the author with respect to aspects such as PRICE and SERVICE, or, in the case of (3.7), CLIMATE and COMPANY (Hu and Liu, 2004). These aspects may be predefined, in which case it may be possible to build a separate classifier per aspect. If the aspects are not defined in advance, it may be necessary to employ **unsupervised machine learning** methods — described in the next chapter — to identify them (e.g., Branavan et al., 2009).

(c) Jacob Eisenstein 2014-2017. Work in progress.

Emotion classification While sentiment analysis is framed in terms of positive and negative categories, psychologists generally regard **emotion** as more multifaceted. For example, Ekman (1992) argues that there are six basic emotions — happiness, surprise, fear, sadness, anger, and contempt — and that they are universal across human cultures. Alm et al. (2005) build a linear classifier for recognizing the emotions expressed in children’s stories. The ultimate goal of this work was to improve text-to-speech synthesis, so that stories could be read with intonation that reflected the emotional content. They used bag-of-words features, as well as features capturing the story type (e.g., jokes, folktales), and structural features that reflect the position of each sentence in the story. The task is difficult: even human annotators frequently disagreed with each other, and the best classifiers achieved accuracy between 60-70%.

Non-classification approaches to sentiment analysis

A more challenging version of sentiment analysis is to determine not just the class of a document, but its rating on a numerical scale (Pang and Lee, 2005). If the scale is continuous, we might take a **regression** approach, identifying a set of weights θ so as to minimize the squared error of a predictor $\hat{y} = \theta \cdot x + b$, where b is an offset. This approach is called **linear regression**, and sometimes **least squares**, because the regression coefficients θ are determined by minimizing the squared error, $(y - \hat{y})^2$. If the weights are regularized using a penalty $\lambda \|\theta\|_2^2$, then the name **ridge regression** is sometimes applied. Both linear regression and ridge regression can be solved in closed form as a system of linear equations.

If the rating scale is discrete, $y \in \{1, 2, \dots, K\}$, we can take a **ranking** approach (Crammer and Singer, 2001), in which scores $\theta \cdot x$ are discretized into ranks, by also learning a set of boundaries, $b_0 = -\infty \leq b_1 \leq \dots \leq b_K$. The learning algorithm consists in making perceptron-like updates to both θ and b . This approach is ideal for settings like predicting a 1-10 rating or a grade (A - F); instead of learning one vector θ for every rank, we can learn a single θ , and then just partition the output space.

Finally, sentiment analysis is one of the only NLP tasks where hand-crafted feature weights are still widely employed. In **lexicon-based classification** (Taboada et al., 2011), the user creates a list of words for each label, and then classifies each document based on how many of the words from each list are present. In our linear classification framework, this is equivalent to choosing the following weights:

$$\theta_{y,j} = \begin{cases} 1, & j \in \mathcal{L}_y \\ 0, & \text{otherwise,} \end{cases} \quad (3.2)$$

where \mathcal{L}_y is the lexicon for label y . Compared to the machine learning classifiers discussed in the previous chapters, lexicon-based classification may seem primitive. However, supervised machine learning relies on large annotated datasets, which are time-consuming and expensive to produce. If the goal is to distinguish two or more categories in a new domain, it may be simpler to start by writing down a list of words for each category.

(c) Jacob Eisenstein 2014-2017. Work in progress.

An early lexicon was the *General Inquirer* (Stone, 1966). Today, popular sentiment lexicons include *sentiwordnet* (Esuli and Sebastiani, 2006) and an evolving set of lexicons from Liu (2015). For emotions and more fine-grained analysis, *Linguistic Inquiry and Word Count* (LIWC) provides a set of lexicons (Tausczik and Pennebaker, 2010). The MPQA lexicon indicates the polarity of some 8221 terms, as well as whether they are strongly or weakly subjective (Wiebe et al., 2005). A comprehensive comparison of sentiment lexicons is offered by Ribeiro et al. (2016). Given an initial set of small “seed” lexicons, it is possible to automatically expand the lexicon by looking for words that tend to co-occur with words in the seed sets (Hatzivassiloglou and McKeown, 1997; Qiu et al., 2011).

3.2 Word sense disambiguation

Consider the the following headlines:

(3.8) *Iraqi head seeks arms*

(3.9) *Prostitutes appeal to Pope*

(3.10) *Drunk gets nine years in violin case¹*

They are ambiguous because they contain words that have multiple meanings, or **senses**. Word sense disambiguation (WSD) is the problem of identifying the intended sense of each word token in a document. WSD is part of a larger field of research called **lexical semantics**, which is concerned with meanings of the words.

Problem definition

At a basic level, the problem of word sense disambiguation is to identify the correct sense for each word token in a document. Part-of-speech ambiguity (e.g., noun versus verb) is usually considered to be a different problem, which is solved at an earlier stage. From a linguistic perspective, senses are not really properties of words, but of **lemmas**, which are canonical forms that stand in for a set of inflected words. For example, *arm/N* is a lemma that includes the inflected form *arms/N* — the */N* indicates that it we are referring to the noun form of the word. Similarly, *arm/V* is a lemma that includes the inflected verbs (*arm/V*, *arms/V*, *armed/V*, *arming/V*). Therefore, WSD requires first identifying the correct part-of-speech and lemma for each token, and then choosing the correct sense from the inventory associated with the corresponding lemma.

How many word senses?

Words (lemmas) may have many more than two senses. For example, the word *serve* would seem to have at least the following senses:

¹These examples, and many more, can be found at <http://www.ling.upenn.edu/~beatrice/humor/headlines.html>

- [FUNCTION]: *The tree stump served as a table*
- [ENABLE]: *His evasive replies only served to heighten suspicion*
- [DISH]: *We serve only the rawest fish here*
- [ENLIST]: *She served her country in the marines*
- [JAIL]: *He served six years in Alcatraz*
- [TENNIS]: *Nobody can return his double-reverse spin serve*
- [LEGAL]: *They were served with subpoenas*²

How do we know that these senses are really different? Linguists often design tests for this purpose, and one such test is to construct a **zeugma**, which combines senses through conjunction:

(3.11) *Which flight serves breakfast?*

(3.12) *Which flights serve Tuscon?*

(3.13) *Which flights serve breakfast and Tuscon?*³

To the extent that you think that (3.13) is semantically unacceptable, you should agree that (3.11) and (3.12) refer to distinct senses of the lemma *serve*.

Standard dictionaries list senses for each word, but WSD research is dominated by a computational resource called WORDNET (<http://wordnet.princeton.edu>). WordNet is organized in terms of lemmas rather than words. An example of a wordnet entry is shown in Figure 3.1. WordNet consists of roughly 100,000 **synsets**, groups of words or phrases with an identical meaning. An example synset is $\{chump^1, fool^2, sucker^1, mark^9\}$, where the superscripts index the sense of each word that is included in the synset: for example, there are several other senses of *mark* that have very different meanings, and are not part of this synset.

A lemma is **polysemous** if it participates in multiple synsets. Besides **synonymy**, WordNet also describes many other lexical relationships, including:

- **antonymy**: x means the opposite of y , e.g. FRIEND-ENEMY;
- **hyponymy**: x is a special case of y , e.g. RED-COLOR; the inverse relationship is **hypernymy**;
- **meronymy**: x is a part of y , e.g., WHEEL-BICYCLE; the inverse relationship is **holonymy**.

WordNet has played an important role in helping WSD move from toy systems to to large-scale quantitative evaluations. However, some have argued that WordNet's sense

²Examples from Dan Klein's lecture notes, [http://www.cs.berkeley.edu/~klein/cs294-7/SP07%20cs294%20lecture%205%20--%20maximum%20entropy%20\(6pp\).pdf](http://www.cs.berkeley.edu/~klein/cs294-7/SP07%20cs294%20lecture%205%20--%20maximum%20entropy%20(6pp).pdf)

³This example is adapted from Jurafsky and Martin (2009).

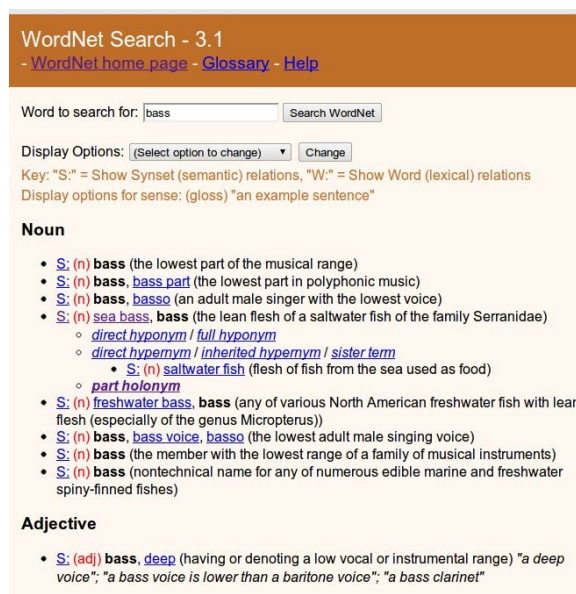


Figure 3.1: Example wordnet entry, from <http://wordnet.princeton.edu>

granularity is too fine (Ide and Wilks, 2006); more fundamentally, the premise that word senses can be differentiated in a task-neutral way has been criticized as linguistically naïve (Kilgarriff, 1997). One way of testing this question is to ask whether people tend to agree on the appropriate sense for example sentences: according to Mihalcea et al. (2004), people agree on roughly 70% of examples using WordNet senses; far better than chance, but perhaps less than we might like.

WSD as Classification

So, how can we tell living *plants* from manufacturing *plants*? The key information often lies in the context:

(3.14) *Town officials are hoping to attract new manufacturing plants through weakened environmental regulations.*

(3.15) *The endangered plants play an important role in the local ecosystem.*

The bag-of-words representation that we applied in document classification can be applied here, by treating each context as a pseudo-document. We can then construct a feature function for each potential sense y ,

$$f(\langle \text{plant}, \text{The endangered plants play an } \dots \rangle, y) = \\ \{ \langle \text{the}, y \rangle : 1, \langle \text{endangered}, y \rangle : 1, \langle \text{play}, y \rangle : 1, \langle \text{an}, y \rangle : 1, \dots \}$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

As in document classification, many of these features are irrelevant, but a few are very strong indicators. In this example, the context word *endangered* is a strong signal that the intended sense is biology rather than manufacturing. We would therefore expect a learning algorithm to assign high weight to $\langle \textit{endangered}, \text{BIOLOGY} \rangle$, and low weight to $\langle \textit{endangered}, \text{MANUFACTURING} \rangle$.

An extension of bag-of-words models is to encode the position of each context word with respect to the target, e.g.,

$$f(\langle \textit{bank}, I \textit{ went to the bank to deposit my paycheck} \rangle, y) = \\ \{ \langle i - 3, \textit{went}, y \rangle : 1, \langle i + 2, \textit{deposit}, y \rangle : 1, \langle i + 4, \textit{paycheck}, y \rangle : 1 \}$$

These **collocation features** give more information about the specific role played by each context word. This idea can be taken further by incorporating additional **syntactic** information about the grammatical role played by each context feature.

After deciding on the features, we can train a classifier to predict the right sense of each word. A **semantic concordance** is a corpus in which each open-class word (nouns, verbs, adjectives, and adverbs) is tagged with its word sense from the target dictionary or thesaurus. SEMCOR is a semantic concordance built from 234K tokens of the Brown corpus, annotated as part of the WordNet project (Fellbaum, 2010). SemCor annotations look like this:

$$(3.16) \quad \textit{As of Sunday}_n^1 \textit{ night}_n^1 \textit{ there was}_v^4 \textit{ no word}_n^2 \dots$$

As always, supervised classification is only possible if enough labeled examples can be accumulated. This is difficult, because each polysemous lemma requires its own training set: having a good classifier for *bank* is of no help at all towards disambiguating *plant*. For this reason, **unsupervised** and **semisupervised** methods are particularly important for WSD (e.g., Yarowsky, 1995). These methods will be discussed in chapter 4 and chapter 19. Unsupervised methods typically lean heavily on the heuristic of “one sense per discourse”, which means that a lemma will usually have a single, consistent sense throughout any given document. Based on this heuristic, we can propagate information from high-confidence instances to lower-confidence instances in the same document. For a survey on word sense disambiguation, see Navigli (2009).

3.3 Design decisions for text classification

Text classification involves a number of design decisions. Some of these decisions, such as smoothing or regularization, are classifier-specific; these decisions are described in the previous chapter. But even the construction of the feature vector itself involves a number of design decisions, and these decisions can often be more important than the choice of classifier.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Preprocessing

One question is whether the vocabulary should be case sensitive: do we distinguish *great*, *Great*, and *GREAT*? What about *coooooooooo!*? In social media text, this sort of **expressive lengthening** can cause the vocabulary size to explode (Brody and Diakopoulos, 2011); we might want to somehow **normalize** the text (Sproat et al., 2001) to collapse the vocabulary again.

A related issue is that suffixes may be irrelevant to the sentiment orientation of each word: for example, *love*, *loved*, and *loving* are all positive, so perhaps we should eliminate the suffix and group them together. The removal of these suffixes is called **stemming** when it is done at the character level (leaving roots like *lov-*), and is called **lemmatization** when the goal is to identify the underlying lemma (in this case, *love*). Both of these methods will be discussed in more detail in chapter 8.

Still another preprocessing decision involves **tokenization**: breaking the text into tokens. This is more complicated than simply looking for whitespace, since we may want to tokenize items such as *well-bred* into $\langle \text{well}, \text{bred} \rangle$, *isn't* into $\langle \text{is}, \text{n't} \rangle$; at the same time, we would like to keep *U.S.* as a single token. This too will be discussed in chapter 8.

Vocabulary

Regularization makes it possible to use large feature spaces without overfitting. However, in some cases it is still preferable to reduce the size of the feature vector by limiting the vocabulary. For example, words such as *the*, *to*, and *and* seem intuitively to play little role in expressing the topic, sentiment, or stance, yet they are very frequent; removing these **stopwords** may therefore improve the classifier. This is typically done by creating a list and then ignoring all terms that match the list. However, the social psychologist Jamie Pennebaker has shown that seemingly inconsequential words — especially personal pronouns like *I*, *you*, and *we* — can sometimes offer surprising insights about the author (Chung and Pennebaker, 2007).

More aggressively, we might assume that sentiment is typically carried by word classes such as adjectives and adverbs (see Chapter 7), and therefore we could focus on these words (Hatzivassiloglou and McKeown, 1997; Turney, 2002). However, Pang et al. (2002) find that in their case, eliminating non-adjectives causes the performance of the classifier to decrease.

Finally, if the goal is simply to create small models that can be stored and loaded efficiently, an alternative approach is to use **feature hashing** (Weinberger et al., 2009). Each feature is assigned an index using a hash function. If a hash function that permits collisions is chosen (typically by taking the hash output modulo some integer), then the model can be made arbitrarily small, as multiple features share a single weight. Because most features are rare, accuracy is surprisingly robust to such collisions (Ganchev and Dredze, 2008).

Count or binary?

Finally, we may consider whether we want our feature vector to include the **count** of each word, or its mere **presence**. This gets at a subtle limitation of linear classification: two *failures* may be worse than one, but is it really twice as bad? A more flexible classifier could assign diminishing weight to each additional instance, but this is hard to do in the linear classification framework, and it's hard to see how much the weight should diminish. Pang et al. (2002) take a simpler approach, using binary presence/absence indicators in the feature vector: $f_i(\mathbf{x}, y) \in \{0, 1\}, \forall i$. They find that classifiers trained on these binary feature vectors outperform classifiers trained on count-based features.

3.4 Evaluating text classification

In any text classification setting, it is critical to reserve a held-out test set, and use this data for only one purpose: to evaluate the overall accuracy of a single classifier. Using this data more than once would cause your estimated accuracy to be overly optimistic. Since it is typically necessary to set hyperparameters or perform feature selection, you may need to construct various “tuning” or “development” sets, but these should not intersect with the test data. For more details, see § 1.2.

There are a number of ways to evaluate classifier performance. The simplest is **accuracy**: the number of correct predictions, divided by the total number of instances. Why isn't this always the right choice? Suppose we were building a classifier to detect whether an essay receives a passing grade. Due perhaps to grade inflation, 95% of all essays receive a passing grade. This means that a classifier that always says “pass” will get 95% accuracy. But this classifier isn't telling us anything useful at all.

Precision, recall, and F -measure

Another way to evaluate this classifier is in terms of its **precision** and **recall**. For each label $y \in \mathcal{Y}$, we define a **positive** instance as one that the classifier labels as $Y_i = y$, and a **negative** instance as one that the classifier labels as $Y_i \neq y$. We can then define four quantities:

- **True positive**: positive and correct, TP
- **False positive**: positive but incorrect, FP
- **True negative**: negative and correct, TN
- **False negative**: negative and incorrect, FN .

(c) Jacob Eisenstein 2014-2017. Work in progress.

From these quantities, we can then define the **recall** and **precision**:

$$r = \frac{TP}{TP + FN} \quad (3.3)$$

$$p = \frac{TP}{TP + FP} \quad (3.4)$$

The recall is the proportion of positive labels among those that **should** have been labeled as positive (for some label y). The precision is the proportion of positive labels among those that **were** labeled as positive. Our “always pass” classifier above would have 100% recall for the positive label, but 95% precision. It would have 0% recall for the negative label, and undefined precision.

The **F -measure** is the harmonic mean of recall and precision,

$$F = \frac{2 \times r \times p}{r + p}. \quad (3.5)$$

F -measure is a classic measure of classifier performance for binary classification problems with unbalanced class distribution. Sometimes it is called F_1 , as there are generalizations of F -measure in which the precision is multiplied by some constant β^2 .

Macro- F_1 is the average F -measure across several classes. In a multi-class problem with unbalanced class distributions, the macro- F_1 is a balanced measure of how well the classifier recognizes each class. In **micro- F_1** , we compute true positives, false positives, and false negatives for each class, and then add them up before computing a single F -measure. This metric is balanced across instances rather than classes, so will weight each class in proportion to how frequently it appears.

Chapter 4

Learning without supervision

So far we've assumed the following setup:

- A **training set** where you get observations x_i and labels y_i
- A **test set** where you only get observations x_i

What if you never get labels y_i ? For example, suppose you are trying to do word sense disambiguation. You get a bunch of text, and you suspect that there are at least two different meanings for the word *concern*. But you don't have any labels for specific instances in which this word is used. What can you?

As described in chapter 3, in supervised word sense disambiguation, we often build feature vectors from the words that appear in the context of the word that we are trying to disambiguate. For example, for the word *concern*, the immediate context might typically include words from one of the following two groups:

1. *services, produces, banking, pharmaceutical, energy, electronics*
2. *about, said, that, over, in, with, had*

Now suppose we were to scatterplot each instance of *concern* on a graph, so that the x-axis is the density of words in group 1, and the y-axis is the density of words in group 2. In such a graph, shown in Figure 4.1, two or more blobs might emerge. These blobs would correspond to the different sense of *concern*.

But in reality, we don't know the word groupings in advance.¹ We have to try to apply the same idea in a very high dimensional space, where every word gets its own dimension — and most dimensions are irrelevant!

¹One approach, which we do not consider here, would be to get them from some existing resource, such as the dictionary definition (Lesk, 1986).

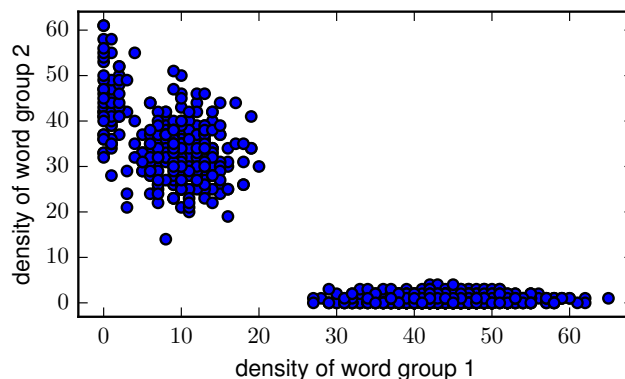


Figure 4.1: Counts of words from two different context groups

Now here’s a related scenario, from a different problem. Suppose you download thousands of news articles, and make a scatterplot, where each point corresponds to a document: the x-axis is the frequency of the word *hurricane*, and the y-axis is the frequency of the word *election*. Again, three clumps might emerge: one for documents that are largely about the hurricane, another for documents largely about the election, and a third clump for documents about neither topic.

These examples are intended to show that we can find structure in data, even without labels — just look for clumps in the scatterplot of features. But again, in reality we cannot make scatterplots of just two words; we may have to consider hundreds or thousands of words. It would be impossible to visualize such a high-dimensional scatterplot, so we will need to design algorithmic approaches to finding these groups.

4.1 *K*-means clustering

You might know about classic clustering algorithms like *K*-means. These algorithms maintain a cluster assignment for each instance, and a central location for each cluster. They then repeatedly update the cluster assignments and the locations, until convergence. Pseudocode for *K*-means is shown in Algorithm 4.

K-means can be used to find coherent clusters of documents in high-dimensional data. When we assign each point to its nearest center, we are choosing which cluster it is in; when we re-estimate the location of the centers, we are determining the defining characteristic of each cluster. *K*-means is a classic algorithmic that has been used and modified in thousands of papers (Jain, 2010); for an application of *K*-means to word sense induction, see Pantel and Lin (2002).

Of the many variants of *K*-means, one that is particularly relevant for our purposes is called **soft *K*-means**. The key difference is that instead of directly assigning each point x_i

Algorithm 4 K -means clustering algorithm

```

1: procedure  $K$ -MEANS( $\mathbf{x}_{1:N}$ )
2:   Initialize cluster centers  $\mu_k \leftarrow \text{Random}()$ 
3:   repeat
4:     for all  $i$  do
5:       Assign each point to the nearest cluster:  $z_i \leftarrow \min_k \text{Distance}(\mathbf{x}_i, \mu_k)$ 
6:     for all  $k$  do
7:       Recompute each cluster center from the points in the cluster:  $\mu_k \leftarrow$ 
          $\frac{1}{\sum_i \delta(z_i=k)} \sum_i \delta(z_i=k) \mathbf{x}_i$ 
8:   until converged

```

to a specific cluster z_i , soft K -means assigns each point a **distribution** over clusters $q_i(z_i)$, so that $\sum_k q_i(k) = 1$, and $\forall_k 0 \leq q_i(k) \leq 1$. The centroid of each cluster is then computed from a **weighted average** of the points in the cluster, where the weights are taken from the q distribution.

We will now explore a more principled, statistical version of soft K -means, called **expectation maximization** (EM) clustering. By understanding the statistical principles underlying the algorithm, we can extend it in a number of ways.

4.2 The Expectation Maximization (EM) Algorithm

Let's go back to the Naïve Bayes model:

$$\log p(\mathbf{x}, \mathbf{y}; \phi, \mu) = \sum_i \log p(\mathbf{x}_i | y_i; \phi) p(y_i; \mu) \quad (4.1)$$

For example, \mathbf{x} can describe the documents that we see today, and \mathbf{y} can correspond to their labels. But suppose we never observe y_i ? Can we still do anything with this model?

Since we don't know \mathbf{y} , let's marginalize it:

$$\log p(\mathbf{x}) = \sum_i^N \log p(\mathbf{x}_i) \quad (4.2)$$

$$= \sum_i \log \sum_{y_i} p(\mathbf{x}_i | y_i; \phi) p(y_i; \mu) \quad (4.3)$$

$$(4.4)$$

We will estimate the parameters ϕ and μ by maximizing the log-likelihood of $\mathbf{x}_{1:N}$, which is our (unlabeled) observed data. Why is this a good thing to maximize? If we

don't have labels, discriminative learning is impossible (there's nothing to discriminate), so maximum likelihood is all we have.

Unfortunately, maximizing $\log P(\mathbf{x})$ directly is intractable. So to estimate this model, we must employ approximation. We do this by introducing an **auxiliary variable** q_i , for each y_i . We want q_i to be a **distribution**, so we have the usual constraints: $\sum_y q_i(y) = 1$ and $\forall y, q_i(y) \geq 0$. In other words, q_i defines a probability distribution over \mathcal{Y} , for each instance i .

Now since $\frac{q_i(y)}{q_i(y)} = 1$, we can multiply the right side by this ratio and preserve the equality,

$$\log p(\mathbf{x}) = \sum_i \log \sum_{y_i} p(\mathbf{x}_i | y_i; \phi) p(y_i; \mu) \frac{q_i(y)}{q_i(y)} \quad (4.5)$$

$$= \sum_i \log E_q \left[\frac{p(\mathbf{x}_i | y; \phi) p(y; \mu)}{q_i(y)} \right], \quad (4.6)$$

by the definition of expectation, $E_q[f(x)] = \sum_x q(x)f(x)$. Note that $E_q[\cdot]$ just means the expectation under the distribution q .

Now we apply **Jensen's inequality**, which says that because \log is a concave function, we can push it inside the expectation, and obtain a lower bound.

$$\log p(\mathbf{x}) \geq \sum_i E_q \left[\log \frac{p(\mathbf{x}_i | y; \phi) p(y; \mu)}{q_i(y)} \right] \quad (4.7)$$

$$\mathcal{J} = \sum_i E_q [\log p(\mathbf{x}_i | y; \phi)] + E_q [\log p(y; \mu)] - E_q [\log q_i(y)] \quad (4.8)$$

By maximizing \mathcal{J} , we are maximizing a lower bound on the joint log-likelihood $\log p(\mathbf{x})$. Now, \mathcal{J} is a function of two sets of arguments:

- the distributions q_i for each i
- the parameters μ and ϕ

We'll optimize with respect to each of these in turn, holding the other one fixed.

Step 1: the E-step

First, we expand the expectation in the lower bound as:

$$\mathcal{J} = \sum_i E_q [\log p(\mathbf{x}_i | y; \phi)] + E_q [\log p(y; \mu)] - E_q [\log q_i(y)] \quad (4.9)$$

$$= \sum_i \sum_y q_i(y) (\log p(\mathbf{x}_i | y; \phi) + \log p(y; \mu) - \log q_i(y)) \quad (4.10)$$

As in Naïve Bayes, we have a “sum-to-one” constraint: in this case, $\sum_y q_i(y) = 1$. Once again, we incorporate this constraint into a Lagrangian:

$$\mathcal{J}_q = \sum_i^N \sum_{y \in \mathcal{Y}} q_i(y) (\log p(\mathbf{x}_i | y; \phi) + \log p(y; \mu) - \log q_i(y)) + \lambda_i (1 - \sum_y q_i(y)) \quad (4.11)$$

We then optimize by taking the derivative and setting it equal to zero:

$$\frac{\partial \mathcal{J}_q}{\partial q_i(y)} = \log p(\mathbf{x}_i | y; \phi) + \log p(y; \mu) - \log q_i(y) - 1 - \lambda_i \quad (4.12)$$

$$\log q_i(y) = \log p(\mathbf{x}_i | y; \phi) + \log p(y; \mu) - 1 - \lambda_i \quad (4.13)$$

$$q_i(y) \propto p(\mathbf{x}_i | y; \phi) p(y; \mu) = p(\mathbf{x}_i, y; \phi, \mu) \quad (4.14)$$

Since q_i is defined over the labels \mathcal{Y} , we normalize it as,

$$q_i(y) = \frac{p(\mathbf{x}_i, y; \phi, \mu)}{\sum_{y' \in \mathcal{Y}} p(\mathbf{x}_i, y'; \phi, \mu)} = p(y | \mathbf{x}_i; \phi, \mu) \quad (4.15)$$

After normalizing, each $q_i(y)$ — which is the soft distribution over clusters for data \mathbf{x}_i — is set to the posterior probability $p(y | \mathbf{x}_i)$ under the current parameters μ, ϕ . This is called the E-step, or “expectation step,” because it is derived from updating the bound on the expected likelihood under $q(y)$. Note that although we introduced the Lagrange multipliers λ_i as additional parameters, we were able to drop these parameters because we solved for $q_i(y)$ to a constant of proportionality.

Step 2: the M-step

Next, we hold $q(y)$ fixed and maximize the bound with respect to the parameters, ϕ and μ . Lets focus on ϕ , which parametrizes the likelihood, $p(\mathbf{x} | y; \phi)$. Again, we have a constraint that $\sum_j^V \phi_{y,j} = 1$, so we start by forming a Lagrangian,

$$\mathcal{J}_\phi = \sum_i^N \sum_{y \in \mathcal{Y}} q_i(y) (\log p(\mathbf{x}_i | y; \phi) + \log p(y; \mu) - \log q_i(y)) + \sum_{y \in \mathcal{Y}} \lambda_y (1 - \sum_j^V \phi_{y,j}). \quad (4.16)$$

Again, we solve by setting the derivative equal to zero:

$$\frac{\partial \mathcal{J}_\phi}{\partial \phi_{y,j}} = \sum_i^N q_i(y) \frac{x_{i,j}}{\phi_{y,j}} - \lambda_y \quad (4.17)$$

$$\lambda_y \phi_{y,j} = \sum_i^N q_i(y) x_{i,j} \quad (4.18)$$

$$\phi_{y,j} \propto \sum_i^N q_i(y) x_{i,j}. \quad (4.19)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

Now because $\sum_j \phi_{y,j} = 1$, we can normalize as follows,

$$\phi_{y,j} = \frac{\sum_i q_i(y) x_{i,j}}{\sum_{j' < V} \sum_i q_i(y) x_{i,j'}} \quad (4.20)$$

$$= \frac{E_q[\text{count}(y, j)]}{E_q[\text{count}(y)]}, \quad (4.21)$$

where $j \in \{1, 2, \dots, V\}$ indexes base features, such as words.

So ϕ_y is now equal to the relative frequency estimate of the **expected counts** under the distribution $q(y)$.

- As in supervised Naïve Bayes, we can apply smoothing to add α to all these counts.
- The update for μ is identical: $\mu_y \propto \sum_i q_i(y)$, the expected proportion of cluster $Y = y$. If needed, we can add smoothing here too.
- So, everything in the M-step is just like Naïve Bayes, except that we use expected counts rather than observed counts.

This is the M -step for a model in which the likelihood $P(\mathbf{x} \mid \mathbf{y})$ is multinomial. For other likelihoods, there may be no closed-form solution for the parameters in the M -step. We may therefore run gradient-based optimization at each M-step, or we may simply take a single step along the gradient step and then return to the E-step (Berg-Kirkpatrick et al., 2010).

Coordinate ascent

Algorithms that alternate between updating various subsets of the parameters are called “coordinate ascent” algorithms.

The objective function \mathcal{J} is **biconvex**, meaning that it is separately convex in $q(\mathbf{y})$ and $\langle \mu, \phi \rangle$, but it is not jointly convex in all terms. In the coordinate ascent algorithm that we have defined, each step is guaranteed not to decrease \mathcal{J} . This is sometimes called “hill climbing”, because you never go down. Specifically, EM is guaranteed to converge to a **local optima** — a point which is as good or better than any of its immediate neighbors. But there may be many such points, and the overall procedure is **not** guaranteed to find a global maximum. Figure 4.2 shows the objective function for EM with ten different random initializations: while the objective function increases monotonically in each run, it converges to several different values.

The fact that there is no guarantee of global optimality means that initialization is important: where you start can determine where you finish. This is not true in the supervised learning algorithms that we have considered, such as logistic regression — although deep learning algorithms do suffer from this problem. But for logistic regression, and for many other supervised learning algorithms, we don’t need to worry about initialization,

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 4.2: Sensitivity of expectation maximization to initialization

because it won't affect our ultimate solution: we are guaranteed to reach the global minimum. Recent work on **spectral learning** has sought to obtain similar guarantees for “latent variable” models, such as the case we are considering now, where x is observed and y is latent. This work is briefly touched on in § 4.4.

Variants In **hard EM**, each q_i distribution assigns probability of 1 to a single \hat{y}_i , and probability of 0 to all others (Neal and Hinton, 1998). This is similar in spirit to K -means clustering. In problems where the space \mathcal{Y} is large, it may be easier to find the maximum likelihood value \hat{y} than it is to compute the entire distribution $q_i(y)$. Spitzkovsky et al. (2010) show that hard EM can outperform standard EM in some cases.

Another variant of the coordinate ascent procedure combines EM with stochastic gradient descent (SGD). In this case, we can do a local E-step at each instance i , and then immediately make an gradient update to the parameters $\langle \mu, \phi \rangle$. This is particularly relevant in cases where there is no closed form solution for the parameters, so that gradient ascent will be necessary in any case. This algorithm is called “incremental EM” by Neal and Hinton (1998), and online EM by Sato and Ishii (2000) and Cappé and Moulines (2009). Liang and Klein (2009) apply a range of different online EM variants to NLP problems, obtaining better results than standard EM in many cases.

How many clusters?

All along, we have assumed that the number of clusters $K = \#|\mathcal{Y}|$ is given. In some cases, this assumption is valid. For example, the dictionary or WordNet might tell us the number of senses for a word. In other cases, the number of clusters should be a tunable

parameter: some readers may want a coarse-grained clustering of news stories into three or four clusters, while others may want a fine-grained clusterings into twenty or more. But in many cases, we will have to choose K ourselves, with little outside guidance.

One solution is to choose the number of clusters to maximize some computable quantity of the clustering. First, note that the likelihood of the training data will always increase with K . For example, if a good solution is available for $K = 2$, then we can always obtain that same solution at $K > 2$; usually we can find an even better solution by fitting the data more closely. The Akaike Information Criterion (AIC; Akaike, 1974) solves this problem by minimizing a linear combination of the log-likelihood and the number of model parameters, $AIC = 2m - 2\mathcal{L}$, where m is the number of parameters and \mathcal{L} is the log-likelihood. Since the number of parameters increases with the number of clusters K , the AIC may prefer more parsimonious models, even if they do not fit the data quite as well.

Another choice is to maximize the **predictive likelihood** on heldout data $\mathbf{x}_{1:N_h}^{(h)}$. This data is not used to estimate the model parameters ϕ and μ ; we can compute the predictive likelihood on this data by keeping the parameters ϕ and μ fixed, and running a single iteration of the E-step. In document clustering or **topic modeling** (Blei, 2012), a typical approach is to split each instance (document) in half. We use the first half to estimate $q_i(z_i)$, and then on the second half we compute the expected log-likelihood,

$$\ell_i = \sum_z q_i(z) (\log p(\mathbf{x}_i | z; \phi) + \log p(z; \mu)). \quad (4.22)$$

On heldout data, this quantity will not necessarily increase with the number of clusters K , because for high enough K , we are likely to overfit the training data. Thus, choosing K to maximize the predictive likelihood on heldout data will limit the extent of overfitting. Note that in general we cannot analytically find the K that maximizes either AIC or the predictive likelihood, so we must resort to grid search: trying a range of possible values of K , and choosing the best one.

Finally, it is worth mentioning an alternative approach, called **Bayesian nonparametrics**, in which the number of clusters K is treated as another latent variable. This enables statistical inference over a set of models with a variable number of clusters; this is not possible with EM, but there are several alternative inference procedures that are suitable for this case (Murphy, 2012), including MCMC (§ 4.4). Reisinger and Mooney (2010) provide a nice example of Bayesian nonparametrics in NLP, applying it to unsupervised word sense induction.

4.3 Applications of EM

EM is not really an “algorithm” like, say, quicksort. Rather, it is a framework for learning with missing data. The recipe for using EM on a problem of interest is:

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Introduce latent variables z , such that it is easy to write the probability $P(\mathcal{D}, z)$, where \mathcal{D} is your observed data; it should also be easy to estimate the associated parameters, given knowledge of z .
- Derive the E-step updates for $q(z)$, which is typically factored as $q(z) = \prod_i q_{z_i}(z_i)$, where i is an index over instances.
- The M-step updates typically correspond to the soft version of some supervised learning algorithm, like Naïve Bayes.

Some more applications of this basic setup are presented here.

Word sense clustering

In the “demos” folder, you can find a demonstration of expectation maximization for word sense clustering. I assume we know that there are two senses, and that the senses can be distinguished by the contextual information in the document. The basic framework is identical to the clustering model of EM as presented above.

Semi-supervised learning

Nigam et al. (2000) offer another application of EM: **semi-supervised learning**. They apply this idea to document classification in the classic “20 Newsgroup” dataset, in which each document is a post from one of twenty newsgroups from the early days of the internet.

In the setting considered by Nigam et al. (2000), we have labels for some of the instances, $\langle \mathbf{x}^{(\ell)}, \mathbf{y}^{(\ell)} \rangle$, but not for others, $\langle \mathbf{x}^{(u)} \rangle$. The question they pose is: can unlabeled data improve learning? If so, then we might be able to get good performance from a smaller number of labeled instances, simply by incorporating a large number of unlabeled instances. This idea is called **semi-supervised learning**, because we are learning from a combination of labeled and unlabeled data; the setting is described in much more detail in chapter 19.

As in Naïve Bayes, the learning objective is to maximize the joint likelihood,

$$\log p(\mathbf{x}^{(\ell)}, \mathbf{x}^{(u)}, \mathbf{y}^{(\ell)}) = \log p(\mathbf{x}^{(\ell)}, \mathbf{y}^{(\ell)}) + \log p(\mathbf{x}^{(u)}) \quad (4.23)$$

We treat the labels of the unlabeled documents as missing data — in other words, as a latent variable. In the E-step we impute $q(y)$ for the unlabeled documents only. The M-step computes estimates of μ and ϕ from the sum of the observed counts from $\langle \mathbf{x}^{(\ell)}, \mathbf{y}^{(\ell)} \rangle$ and the expected counts from $\langle \mathbf{x}^{(u)} \rangle$ and $q(\mathbf{y})$.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Nigam et al. (2000) further parametrize this approach by weighting the unlabeled documents by a scalar λ , which is a tuning parameter. The resulting criterion is:

$$\mathcal{L} = \log p(\mathbf{x}^{(\ell)}, \mathbf{y}^{(\ell)}) + \lambda \log p(\mathbf{x}^{(u)}) \quad (4.24)$$

$$\geq \log p(\mathbf{x}^{(\ell)}, \mathbf{y}^{(\ell)}) + \lambda E_q[\log p(\mathbf{x}^{(u)}, y)] \quad (4.25)$$

The scaling factor does not really have a probabilistic justification, but it can be important to getting good performance, especially when the amount of labeled data is small in comparison to the amount of unlabeled data. In that scenario, the risk is that the unlabeled data will dominate, causing the parameters to drift towards a “natural clustering” that may be a bad fit for the labeled data. Nigam et al. (2000) show that this approach can give substantial improvements in classification performance when the amount of labeled data is small.

Multi-component modeling

Now let us consider an alternative application of EM to supervised classification. One of the classes in 20 newsgroups is `comp.sys.mac.hardware`; suppose that within this newsgroup there are two kinds of posts: reviews of new hardware, and question-answer posts about hardware problems. The language in these **components** of the `mac.hardware` class might have little in common. So we might do better if we model these components separately. Nigam et al. (2000) show that EM can be applied to this setting as well.

Recall that Naïve Bayes is based on a generative process, which provides a stochastic explanation for the observed data. For multi-component modeling, we envision a slightly different generative process, incorporating both the observed label y_i and the latent component z_i :

- For each document i ,
 - draw the label $y_i \sim \text{Categorical}(\mu)$
 - draw the component $z_i \mid y_i \sim \text{Categorical}(\beta_{y_i})$, where $z_i \in 1, 2, \dots, K_z$.
 - draw the vector of counts $\mathbf{x}_i \mid z_i \sim \text{Multinomial}(\phi_{z_i})$

Our labeled data includes $\langle \mathbf{x}_i, y_i \rangle$, but not z_i , so this is another case of missing data. Again, we sum over the missing data, applying Jensen’s inequality to as to obtain a lower bound on the log-likelihood,

$$\log p(\mathbf{x}_i, y_i) = \log \sum_z^{K_z} p(\mathbf{x}_i, y_i, z) \quad (4.26)$$

$$\geq \log p(y_i; \mu) + E_q[\log p(\mathbf{x}_i \mid z; \phi) + \log p(z \mid y_i; \psi) - \log q_i(z)]. \quad (4.27)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

We are now ready to apply expectation maximization. As usual, the distribution over the missing data — the component z_i — $q_i(z)$ is updated in the E-step. Then during the m-step, we compute:

$$\beta_{y,z} = \frac{E_q [\text{count}(y, z)]}{\sum_{z'} E_q [\text{count}(y, z')]} \quad (4.28)$$

$$\phi_{z,j} = \frac{E_q [\text{count}(z, j)]}{\sum_{j'} E_q [\text{count}(z, j')]} \quad (4.29)$$

Suppose we assume each class y is associated with K components, \mathcal{Z}_y . We can then add a constraint to the E-step so that $q_i(z) = 0$ if $z \notin \mathcal{Z}_y \wedge Y_i = y$.

4.4 *Other approaches to learning with latent variables

Expectation maximization is a very general way to think about learning with latent variables, but it has some limitations. One is the sensitivity to initialization, which means that we cannot simply run EM once and expect to get a good solution. Indeed, in practical applications of EM, quite a lot of attention may be devoted to finding a good initialization. A second issue is that EM tends to be easiest to apply in cases where the latent variables have a clear decomposition (in the cases we have considered, they decompose across the instances). For these reasons, it is worth briefly considering some alternatives to EM.

Sampling

Recall that in EM, we set $q(\mathbf{z}) = \prod_i q_i(z_i)$, factoring the q distribution into conditionally independent q_i distributions. In sampling-based algorithms, rather than maintaining a distribution over each latent variable, we draw random samples of the latent variables. If the sampling algorithm is designed correctly, this procedure will eventually converge to drawing samples from the true posterior, $p(\mathbf{z}_{1:N} \mid \mathbf{x}_{1:N})$. For example, in the case of clustering, we will draw samples from the distribution over clusterings of the data. If a single clustering is required, we can select the one with the highest joint likelihood, $p(\mathbf{z}_{1:N}, \mathbf{x}_{1:N})$.

This general family of algorithms is called **Markov Chain Monte Carlo** (MCMC): “Monte Carlo” because it is based on a series of random draws; “Markov Chain” because the sampling procedure must be designed such that each sample depends only on the previous sample, and not on the entire sampling history. Gibbs Sampling is a particularly simple and effective MCMC algorithm, in which we sample each latent variable from its posterior distribution,

$$z_i \mid \mathbf{x}, \mathbf{z}_{-i} \sim p(z_i \mid \mathbf{x}, \mathbf{z}_{-i}), \quad (4.30)$$

where \mathbf{z}_{-i} indicates $\{\mathbf{z} \setminus z_i\}$, the set of all latent variables except for z_i .

(c) Jacob Eisenstein 2014-2017. Work in progress.

What about the parameters, ϕ and μ ? One possibility is to turn them into latent variables too, by adding them to the generative story. This requires specifying a prior distribution; the Dirichlet is a typical choice of prior for the parameters of a multinomial, since it has support over vectors of non-negative numbers that sum to one, which is exactly the set of permissible parameters for a multinomial. For example,

$$\phi_y \sim \text{Dirichlet}(\alpha), \forall y \quad (4.31)$$

We can then sample $\phi_y \mid \mathbf{x}, \mathbf{z} \sim p(\phi_y \mid \mathbf{x}, \mathbf{z}, \alpha)$; this posterior distribution will also be Dirichlet, with parameters $\alpha + \sum_{i: y_i=y} \mathbf{x}_i$. Alternatively, we can analytically marginalize these parameters, as in **Collapsed Gibbs Sampling**; this is usually preferable if possible. Finally, we might maintain ϕ and μ as parameters rather than latent variables. We can employ sampling in the E-step of the EM algorithm, obtaining a hybrid algorithm called Monte Carlo Expectation Maximization (MCEM; Wei and Tanner, 1990).

In principle, these algorithms will eventually converge to the true posterior distribution. However, there is no way to know how long this will take; there is not even any way to check on whether the algorithm has converged. In practice, convergence again depends on initialization, since it might take ages to recover from a poor initialization. Thus, while Gibbs Sampling and other MCMC algorithms provide a powerful and flexible array of techniques for statistical inference in latent variable models, they are not a panacea for the problems experienced by EM.

Murphy (2012) includes an excellent chapter on MCMC; for a more comprehensive treatment, see Robert and Casella (2013).

Spectral learning

A more recent approach to learning with latent variables is based on the **method of moments**. In these approaches, we avoid the problem of non-convex log-likelihood by using a different estimation criterion. Let us write $\bar{\mathbf{x}}_i$ for the normalized vector of word counts in document i , so that $\bar{\mathbf{x}}_i = \mathbf{x}_i / \sum_j x_{ij}$. Then we can form a matrix of word-word co-occurrence counts,

$$\mathbf{C} = \sum_i \mathbf{x}_i \mathbf{x}_i^\top. \quad (4.32)$$

We can also compute the expected value of this matrix under $p(\mathbf{x} \mid \phi, \mu)$, as

$$E[\mathbf{C}] = \sum_i \sum_k P(Z_i = k \mid \mu) \phi_k \phi_k^\top \quad (4.33)$$

$$= \sum_k N \mu_k \phi_k \phi_k^\top \quad (4.34)$$

$$= \Phi \text{Diag}(N\mu) \Phi^\top, \quad (4.35)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

where Φ is formed by horizontally concatenating $\phi_1 \dots \phi_K$, and $\text{Diag}(N\mu)$ indicates a diagonal matrix with values $N\mu_k$ at position (k, k) . Now, by setting \mathbf{C} equal to its expectation, we obtain,

$$\mathbf{C} = \Phi \text{Diag}(N\mu) \Phi^\top, \quad (4.36)$$

which is very similar to the eigendecomposition $\mathbf{C} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^\top$. This suggests that simply by finding the eigenvectors and eigenvalues of \mathbf{C} , we could obtain the parameters ϕ and μ , and this is what motivates the name **spectral learning**.

However, there is a key difference in the constraints on the solutions to the two problems. In eigendecomposition, we require orthonormality, so that $\mathbf{Q}\mathbf{Q}^\top = \mathbb{I}$. But in estimating the parameters of a mixture model, we require the columns of Φ represents probability vectors, $\forall k, j, \phi_{k,j} \geq 0, \sum_j \phi_{k,j} = 1$, and that the entries of μ correspond to the probabilities over components. Thus, spectral learning algorithms must include a procedure for converting the solution into vectors of probabilities. One approach is to replace eigendecomposition (or the related singular value decomposition) with non-negative matrix factorization (Xu et al., 2003), which guarantees that the solutions are non-negative (Arora et al., 2013).

After obtaining the parameters ϕ and μ , we can obtain the distribution over clusters for each document by simply computing $p(z_i \mid \mathbf{x}_i; \phi, \mu) \propto p(\mathbf{x}_i \mid z_i; \phi)p(z_i; \mu)$. The advantages of spectral learning are that it obtains (provably) good solutions without regard to initialization, and that it can be quite fast in practice. Anandkumar et al. (2014) describe how similar matrix and tensor factorizations can be applied to statistical estimation in many other forms of latent variable models.

Chapter 5

Language models

In probabilistic classification, we are interested in computing the probability of a label, conditioned on the text. Let us now consider something like the inverse problem: computing the probability of text itself. Specifically, we will consider models that assign probability to a sequence of word tokens,¹ $p(w_1, w_2, \dots, w_M)$, with $w_m \in \mathcal{V}$. The set \mathcal{V} is a discrete vocabulary,

$$\mathcal{V} = \{\text{aardvark}, \text{abacus}, \dots, \text{zither}\}. \quad (5.1)$$

Why would we want to compute the probability of a word sequence? In many applications, our goal is to produce word sequences as output:

- In **machine translation**, we convert from text in a source language to text in a target language.
- In **speech recognition**, we convert from audio signal to text.
- In **summarization**, we convert from long texts into short texts.
- In **dialogue systems**, we convert from the user's input (and perhaps an external knowledge base) into a text response.

In each of these cases, a key subcomponent is to compute the probability of the output text. By choosing high-probability output, we hope to generate texts that are more **fluent**. For example, suppose we want to translate a sentence from Spanish to English.

$$(5.1) \quad \textit{El cafe negro me gusta mucho}.$$

¹The linguistic term “word” does not cover everything we might want to model, such as names, numbers, and emoticons. Instead, we prefer the term **token**, which refers to anything that can appear in a sequence of linguistic data. **Tokenizers** are programs for segmenting strings of characters or bytes into tokens. In standard written English, tokenization is relatively straightforward, and can be performed using a regular expression. But in languages like Chinese, tokens are not usually separated by spaces, so tokenization can be considerably more challenging. For more on tokenization algorithms, see Manning et al. (2008), chapter 2.

A literal word-for-word translation (sometimes called a **gloss**) is,

(5.2) *The coffee black me pleases much.*

A good language model of English will tell us that the probability of this translation is low, in comparison with more grammatical alternatives, such as,

$$p(\text{The coffee black me pleases much}) < p(\text{I love dark coffee}). \quad (5.2)$$

How can we use this fact? Warren Weaver, one of the early leaders in machine translation, viewed it as a problem of breaking a secret code (Weaver, 1955):

When I look at an article in Russian, I say: ‘This is really written in English, but it has been coded in some strange symbols. I will now proceed to decode.’

This observation motivates a generative model (like Naïve Bayes):

- The English sentence $w^{(e)}$ is generated from a **language model**, $p_e(w^{(e)})$.
- The Spanish sentence $w^{(s)}$ is then generated from a **translation model**, $p_{s|e}(w^{(s)} | w^{(e)})$.

Given these two distributions, we can then perform translation by Bayes rule:

$$p_{e|s}(w^{(e)} | w^{(s)}) \propto p_{e,s}(w^{(e)}, w^{(s)}) \quad (5.3)$$

$$= p_{s|e}(w^{(s)} | w^{(e)}) \times p_e(w^{(e)}). \quad (5.4)$$

This is sometimes called the **noisy channel model**, because it envisions English text turning into Spanish by passing through a noisy channel, $p_{s|e}$. What is the advantage of modeling translation this way, as opposed to modeling $p_{e|s}$ directly? The crucial point is that the two distributions $p_{s|e}$ (the translation model) and p_e (the language model) can be estimated from separate data. The translation model requires **bitext** — examples of correct translations. But the language model requires only text in English. Such monolingual data is much more widely available, which means that the fluency of the output translation can be improved simply by scraping more webpages. Furthermore, once estimated, the language model p_e can be reused in any application that involves generating English text, from summarization to speech recognition.

5.1 N-gram language models

How can we estimate the probability of a sequence of word tokens? The simplest idea would be to apply a **relative frequency estimator**. For example, consider the quote, attributed to Picasso, “computers are useless, they can only give you answers.” We can estimate

(c) Jacob Eisenstein 2014-2017. Work in progress.

the probability of this sentence as follows:

$$\begin{aligned} & p(\text{Computers are useless, they can only give you answers}) \\ &= \frac{\text{count}(\text{Computers are useless, they can only give you answers})}{\text{count}(\text{all sentences ever spoken})} \end{aligned} \quad (5.5)$$

This estimator is **unbiased**: in the theoretical limit of infinite data, the estimate will be correct. But in practice, we are asking for accurate counts over an infinite number of events, since sequences of words can be arbitrarily long. Even if we set an aggressive upper bound of, say, $n = 20$ tokens in the sequence, the number of possible sequences is $|\mathcal{V}|^{20}$. A small vocabulary for English would have $|\mathcal{V}| = 10^4$, so we would have 10^{80} possible sequences. Clearly, this estimator is very data-hungry, and suffers from high variance: even grammatical sentences will have probability zero if they happen not to have occurred in the training data.² We therefore need to introduce bias to have a chance of making reliable estimates from finite training data. The language models that follow in this chapter introduce bias in various ways.

We begin with n -gram language models, which compute the probability of a sequence as the product of probabilities of subsequences. The probability of a sequence $p(\mathbf{w}) = p(w_1, w_2, \dots, w_M)$ can be refactored using the chain rule:

$$p(\mathbf{w}) = p(w_1, w_2, \dots, w_M) \quad (5.6)$$

$$= p(w_1) \times p(w_2 | w_1) \times p(w_3 | w_2, w_1) \times \dots \times p(w_M | w_{M-1}, \dots, w_1) \quad (5.7)$$

Each element in the product is the probability of a word given all its predecessors. We can think of this as a *word prediction* task: given the context *Computers are*, we want to compute a probability over the next token. The relative frequency estimate of the probability of the word *useless* in this context is,

$$\begin{aligned} p(\text{useless} | \text{computers are}) &= \frac{\text{count}(\text{computers are useless})}{\sum_{x \in \mathcal{V}} \text{count}(\text{computers are } x)} \\ &= \frac{\text{count}(\text{computers are useless})}{\text{count}(\text{computers are})}. \end{aligned}$$

Note that we haven't made any approximations yet, and we could have just as well applied the chain rule in reverse order,

$$p(\mathbf{w}) = p(w_M) \times p(w_{M-1} | w_M) \times \dots \times p(w_1 | w_2, \dots, w_M), \quad (5.8)$$

²Chomsky has famously argued that this is evidence against the very concept of probabilistic language models: no such model could distinguish the grammatical sentence *colorless green ideas sleep furiously* from the ungrammatical permutation *furiously sleep ideas green colorless*. Indeed, even the bigrams in these two examples are unlikely to occur — at least, not in texts written before Chomsky proposed this example.

or in any other order. But this means that we also haven't really improved anything either: to compute the conditional probability $p(w_M | w_{M-1}, w_{M-2}, \dots, w_1)$, we need to model $|\mathcal{V}|^{M-1}$ contexts. We cannot estimate such a distribution from any reasonable finite sample.

To solve this problem, n -gram models make a crucial simplifying approximation: condition on only the past $n - 1$ words.

$$p(w_m | w_{m-1} \dots w_1) \approx p(w_m | w_{m-1}, \dots, w_{m-n+1}) \quad (5.9)$$

This means that the probability of a sentence w can be computed as

$$p(w_1, \dots, w_M) \approx \prod_m^M p(w_m | w_{m-1}, \dots, w_{m-n+1}) \quad (5.10)$$

To compute the probability of an entire sentence, it is convenient to pad the beginning and end with special symbols \diamond and \blacklozenge . Then the bigram ($n = 2$) approximation to the probability of *I like black coffee* is:

$$p(I \text{ like black coffee}) = p(I | \diamond) \times p(\text{like} | I) \times p(\text{black} | \text{like}) \times p(\text{coffee} | \text{black}) \times p(\blacklozenge | \text{coffee}). \quad (5.11)$$

In this model, we have to estimate and store the probability of only $|\mathcal{V}|^n$ events, which is exponential in the order of the n -gram, and not $|\mathcal{V}|^M$, which is exponential in the length of the sentence. The n -gram probabilities can be computed by relative frequency estimation,

$$\Pr(W_m = i | W_{m-1} = j, W_{m-2} = k) = \frac{\text{count}(i, j, k)}{\sum_{i'} \text{count}(i', j, k)} = \frac{\text{count}(i, j, k)}{\text{count}(j, k)} \quad (5.12)$$

A key design question is how to set the hyperparameter n , which controls the size of the context used in each conditional probability. If this is misspecified, the language model will sacrifice accuracy. Let's consider the potential problems concretely.

When n is too small. Consider the following sentences:

(5.3) ***Gorillas** always like to groom **THEIR** friends.*

(5.4) *The **computer** that's on the 3rd floor of our office building **CRASHED**.*

The uppercase bolded words depend crucially on their predecessors in lowercase bold: the likelihood of *their* depends on knowing that *gorillas* is plural, and the likelihood of *crashed* depends on knowing that the subject is a *computer*. If the n -grams are not big enough to capture this context, then the resulting language model would offer probabilities that are too low for these sentences, and too high for sentences that fail basic linguistic tests like number agreement.

(c) Jacob Eisenstein 2014-2017. Work in progress.

When n is too big. In this case, we cannot make good estimates of the n -gram parameters from our dataset, because of data sparsity. To handle the *gorilla* example, we would need to model 6-grams; which means accounting for $|\mathcal{V}|^6$ events. Under a very small vocabulary of $|\mathcal{V}| = 10^4$, this means estimating the probability of 10^{24} distinct events.

These two problems point to another **bias-variance** tradeoff. A small n -gram size introduces high bias with respect to the true distribution, and a large n -gram size introduces high variance due to the huge number of possible events. But in reality the situation is even worse, because we often have both problems at the same time! Language is full of long-range dependencies that we cannot capture because n is too small; at the same time, language datasets are full of rare phenomena, whose probabilities we fail to estimate accurately because n is too large.

We will seek approaches to keep n large, while still making low-variance estimates of the underlying parameters. To do this, we will introduce a different sort of bias: **smoothing**. But first, let's take a digression to discuss how to evaluate language models.

5.2 Smoothing and discounting

Limited data is a persistent problem in estimating language models. In § 5.1, we presented n -grams as a partial solution. But as we saw, sparse data can be a problem even for low-order n -grams; at the same time, many linguistic phenomena, like subject-verb agreement, cannot be incorporated into language models without higher-order n -grams. It is therefore necessary to add additional inductive biases to n -gram language models. This section covers some of the most intuitive and common approaches, but there are many more (Chen and Goodman, 1999).

Smoothing

A major concern in language modeling is to avoid the situation $p(w) = 0$, which could arise as a result of a single unseen n -gram. A similar problem arose in Naïve Bayes, and there we solved it by **smoothing**: adding imaginary “pseudo” counts. The same idea can be applied to n -gram language models, as shown here in the bigram case,

$$p_{\text{smooth}}(w_m \mid w_{m-1}) = \frac{\text{count}(w_{m-1}, w_m) + \alpha}{\sum_{w' \in \mathcal{V}} \text{count}(w_{m-1}, w') + |\mathcal{V}| \alpha}. \quad (5.13)$$

This basic framework is called **Lidstone smoothing**, but special cases have other names:

- **Laplace smoothing** corresponds to the case $\alpha = 1$.
- **Jeffreys-Perks law** corresponds to the case $\alpha = 0.5$. Manning and Schütze (1999) offer more insight on the justifications for this setting.

(c) Jacob Eisenstein 2014-2017. Work in progress.

			Lidstone smoothing, $\alpha = 0.1$		Discounting, $d = 0.1$	
	counts	unsmoothed probability	effective counts	smoothed probability	effective counts	smoothed probability
<i>impropriety</i>	8	0.4	7.826	0.391	7.9	0.395
<i>offense</i>	5	0.25	4.928	0.246	4.9	0.245
<i>damage</i>	4	0.2	3.961	0.198	3.9	0.195
<i>deficiencies</i>	2	0.1	2.029	0.101	1.9	0.095
<i>outbreak</i>	1	0.05	1.063	0.053	0.9	0.045
<i>infirmity</i>	0	0	0.097	0.005	0.25	0.013
<i>cephalopods</i>	0	0	0.097	0.005	0.25	0.013

Table 5.1: Example of Lidstone smoothing and absolute discounting in a bigram language model, for the context (*alleged*, *-*), for a toy corpus with a total of twenty counts over the seven words shown. Note that discounting decreases the probability for all but the unseen words, while Lidstone smoothing increases the effective counts and probabilities for *deficiencies* and *outbreak*.

To maintain normalization, anything that we add to the numerator (α) must also appear in the denominator ($|\mathcal{V}|\alpha$). This idea is reflected in the concept of **effective counts**:

$$c_i^* = (c_i + \alpha) \frac{M}{M + |\mathcal{V}|\alpha}, \quad (5.14)$$

where c_i is the count of event i , c_i^* is the effective count, and $M = \sum_i^{|\mathcal{V}|} c_i$ is the total number of terms in the dataset (w_1, w_2, \dots, w_M). This term ensures that $\sum_i^{|\mathcal{V}|} c_i^* = \sum_i^{|\mathcal{V}|} c_i = M$. The **discount** for each n-gram is then computed as,

$$d_i = \frac{c_i^*}{c_i} = \frac{(c_i + \alpha)}{c_i} \frac{M}{(M + |\mathcal{V}|\alpha)}.$$

Discounting and backoff

Discounting “borrows” probability mass from observed n-grams and redistributes it. In Lidstone smoothing, we borrow probability mass by increasing the denominator of the relative frequency estimates, and redistribute it by increasing the numerator for all n-grams. But instead, we could borrow the same amount of probability mass from all observed counts, and redistribute it among only the unobserved counts. This is called **absolute discounting**. For example, suppose we set an absolute discount $d = 0.1$ in a bigram model, and then redistribute this probability mass equally over the unseen words. The resulting probabilities are shown in Table 5.1.

Discounting reserves some probability mass from the observed data, and we need not redistribute this probability mass equally. Instead, we can **backoff** to a lower-order

(c) Jacob Eisenstein 2014-2017. Work in progress.

language model. In other words, if you have trigrams, use trigrams; if you don't have trigrams, use bigrams; if you don't even have bigrams, use unigrams. This is called **Katz backoff**. In this smoothing model, bigram probabilities are computed as,

$$c^*(i, j) = c(i, j) - d \quad (5.15)$$

$$P_{\text{Katz}}(i | j) = \begin{cases} \frac{c^*(i, j)}{c(j)} & \text{if } c(i, j) > 0 \\ \alpha(j) \times \frac{P_{\text{unigram}}(i)}{\sum_{i': c(i', j)=0} P_{\text{unigram}}(i')} & \text{if } c(i, j) = 0. \end{cases} \quad (5.16)$$

The term $\alpha(j)$ indicates the amount of probability mass that has been discounted for context j . This probability mass is then divided across all the unseen events, $\{i' : c(i', j) = 0\}$, proportional to the unigram probability of each word i' . The discount parameter d can be optimized to maximize performance (typically held-out log-likelihood) on a development set.

*Interpolation

Backoff is one way to combine n -gram models across various values of n . An alternative approach is **interpolation**: setting the probability of a word in context to a weighted sum of its probabilities across progressively shorter contexts.

Instead of choosing a single n for the size of the n -gram, we can take the weighted average across several n -gram probabilities. For example, for an interpolated trigram model,

$$\begin{aligned} P_{\text{Interpolation}}(i | j, k) &= \lambda_3 p_3^*(i | j, k) \\ &\quad + \lambda_2 p_2^*(i | j) \\ &\quad + \lambda_1 p_1^*(i). \end{aligned}$$

In this equation, p_n^* is the unsmoothed empirical probability given by an n -gram language model, and λ_n is the weight assigned to this model. To ensure that the interpolated $p(w)$ is still a valid probability distribution, we must obey the constraint, $\sum_n \lambda_n = 1$. But how to find the specific values of λ ?

An elegant solution is **expectation maximization**. Recall from chapter 4 that we can think about EM as learning with **missing data**: we just need to choose missing data such that learning would be easy if it weren't missing. What's missing in this case? We can think of each word w_m as drawn from an n -gram of unknown size, $z_m \in \{1 \dots n_{\text{max}}\}$. This z_m is the missing data that we are looking for. Therefore, the application of EM to this problem involves the following **generative process**:

- For each token $m \in \{1, 2, \dots, M\}$:
 - draw $z_m \sim \text{Categorical}(\lambda)$,

(c) Jacob Eisenstein 2014-2017. Work in progress.

– draw $w_m \sim p_{z_m}^*(w_m \mid w_{m-1}, \dots, w_{m-z_m})$.

If the missing data $\{Z_m\}$ were known, then we could estimate λ from relative frequency estimation,

$$\lambda_z = \frac{\text{count}(Z_m = z)}{M} \quad (5.17)$$

$$\propto \sum_{m=1}^M \delta(Z_m = z). \quad (5.18)$$

But since we do not know the values of the latent variables Z_m , we impute a distribution q_m in the E-step, which represents the degree of belief that word token w_m was generated from a n -gram of order z_m ,

$$q_m(z) \triangleq \Pr(Z_m = z \mid \mathbf{w}_{1:m}; \lambda) \quad (5.19)$$

$$= \frac{p(w_m \mid \mathbf{w}_{1:m-1}, Z_m = z) \times p(z)}{\sum_{z'} p(w_m \mid \mathbf{w}_{1:m-1}, Z_m = z') \times p(z')} \quad (5.20)$$

$$\propto p_z^*(w_m \mid \mathbf{w}_{1:m-1}) \times \lambda_z. \quad (5.21)$$

In the M-step, we can compute λ by summing the expected counts under q ,

$$\lambda_z \propto \sum_{m=1}^M q_m(z). \quad (5.22)$$

By iterating between updates to q and λ , we will ultimately converge at a solution. The complete algorithm is shown in Algorithm 5.

*Kneser-Ney smoothing

Kneser-Ney smoothing is based on absolute discounting, but it redistributes the resulting probability mass in a different way from Katz backoff. Empirical evidence points to Kneser-Ney smoothing as the state-of-art for n -gram language modeling ?.

To motivate Kneser-Ney smoothing, consider the example: *I recently visited* .. Which of the following is more likely?

- *Francisco*
- *Duluth*

Now suppose that both bigrams *visited Duluth* and *visited Francisco* are unobserved in our training data, and furthermore, that the unigram probability $p^*(\text{Francisco})$ is greater than $p^*(\text{Duluth})$. Nonetheless we would still guess that $p(\text{visited Duluth}) > p(\text{visited Francisco})$,

(c) Jacob Eisenstein 2014-2017. Work in progress.

Algorithm 5 Expectation-maximization for interpolated language modeling

```

1: procedure ESTIMATE INTERPOLATED  $n$ -GRAM ( $\mathbf{w}_{1:M}, \{\mathbf{p}_n^*\}_{n \in 1:n_{\max}}$ )
2:   for  $z \in \{1, 2, \dots, n_{\max}\}$  do ▷ Initialization
3:      $\lambda_z \leftarrow \frac{1}{n_{\max}}$ 
4:   repeat
5:     for  $m \in \{1, 2, \dots, M\}$  do ▷ E-step
6:       for  $z \in \{1, 2, \dots, n_{\max}\}$  do
7:          $q_m(z) \leftarrow \mathbf{p}_z^*(w_m \mid \mathbf{w}_{1:m-}) \times \lambda_z$ 
8:        $\mathbf{q}_m \leftarrow \text{Normalize}(\mathbf{q}_m)$ 
9:       for  $z \in \{1, 2, \dots, n_{\max}\}$  do ▷ M-step
10:         $\lambda_z \leftarrow \frac{1}{M} \sum_{m=1}^M q_m(z)$ 
11:   until tired
12:   return  $\lambda$ 

```

because *Duluth* is a more **versatile** word: it occurs in many contexts, while *Francisco* usually occurs in a single context, following the word *San*. This notion of versatility is the key to Kneser-Ney smoothing.

Writing u for a context of undefined length, and $\text{count}(w, u)$ as the count of word w in context u , we define the Kneser-Ney bigram probability as

$$\mathbf{p}_{KN}(w \mid u) = \begin{cases} \frac{\text{count}(w, u) - d}{\text{count}(u)}, & \text{count}(w, u) > 0 \\ \alpha(u) \times \mathbf{p}_{\text{continuation}}(w), & \text{otherwise} \end{cases} \quad (5.23)$$

$$\mathbf{p}_{\text{continuation}}(w) = \frac{|u : \text{count}(w, u) > 0|}{\sum_{w' \in \mathcal{V}} |u' : \text{count}(w', u') > 0|}. \quad (5.24)$$

First, note that we reserve probability mass using absolute discounting d , which is taken from all unobserved n -grams. The total amount of discounting in context u is $d \times |w : \text{count}(w, u) > 0|$, and we divide this probability mass equally among the unseen n -grams,

$$\alpha(u) = |w : \text{count}(w, u) > 0| \times \frac{d}{\text{count}(u)}. \quad (5.25)$$

This is the amount of probability mass left to account for versatility, which we define via the *continuation probability* $\mathbf{p}_{\text{continuation}}(w)$ as proportional to the number of observed contexts in which w appears. In the numerator of the continuation probability we have the number of contexts u in which w appears, and in the denominator, we normalize by summing the same quantity over all words w' .

The idea of modeling versatility by counting contexts may seem heuristic, but there is an elegant theoretical justification from Bayesian nonparametrics (Teh, 2006). Kneser-Ney

smoothing on n -grams was the dominant language modeling technique — widely used in speech recognition and machine translation — before the arrival of neural language models.

5.3 Recurrent neural network language models

Until this decade, n -grams were the dominant language modeling approach. But in a few years, they have been almost completely supplanted by a new family of language models based on **neural networks**. These models do not make the n -gram assumption of restricted context; indeed, they can incorporate arbitrarily distant contextual information, while remaining computationally and statistically tractable.

The first insight is to treat word prediction as a **discriminative** learning task: rather than directly estimating the distribution $p(w | u)$ from (smoothed) relative frequencies, we now treat language modeling as a machine learning problem, and estimate parameters that maximize the log conditional probability of a corpus.³

The second insight is to reparametrize the probability distribution $p(w | u)$ as a function of two dense K -dimensional numerical vectors, $\beta_w \in \mathbb{R}^K$, and $v_u \in \mathbb{R}^K$,

$$p(w | u) = \frac{\exp(\beta_w \cdot v_u)}{\sum_{w' \in \mathcal{V}} \exp(\beta_{w'} \cdot v_u)}, \quad (5.26)$$

where $\beta_w \cdot v_u$ represents a dot product. Note that the denominator ensures that it is a properly normalized probability distribution. In the neural networks literature, this function is sometimes known as a **softmax** layer, written

$$(\text{SoftMax}(\mathbf{a}))_i = \frac{\exp(a_i)}{\sum_j \exp(a_j)}, \quad (5.27)$$

where \mathbf{a} is a vector of scores and $\text{SoftMax}(\mathbf{a})_i$ is a normalized distribution.⁴

The word vectors β_w are parameters of the model, and are estimated directly. As we will see in chapter 15, these vectors carry useful information about word meaning, and semantically similar words tend to have highly correlated vectors.

The context vectors v_u can be computed in various ways, depending on the model. Here we will consider a relatively simple — but effective — neural language model, the **recurrent neural network** (RNN; Mikolov et al., 2010). The basic idea is to recurrently update the context vectors as we move through the sequence. Let us write \mathbf{h}_m for the

³This idea is not in itself new; for example, Rosenfeld (1996) applies logistic regression to language modeling, and Roark et al. (2007) apply perceptrons and conditional random fields (§ 6.5).

⁴The logistic regression classifier can be viewed as an application of the softmax transformation to the vector constructed by computing the inner products of weights and features for all possible labels.




Figure 5.1: The recurrent neural network language model, viewed as an “unrolled” computation graph. Solid lines indicate direct computation, and dotted blue lines indicate probabilistic dependencies. Circles indicate variables; the left column of nodes are parameters.

contextual information at position m in the sequence. RNNs employ the following recurrence:

$$\mathbf{x}_m \triangleq \phi_{w_m} \quad (5.28)$$

$$\mathbf{h}_m = g(\Theta \mathbf{h}_{m-1} + \mathbf{x}_m) \quad (5.29)$$

$$p(w_{m+1} \mid w_1, w_2, \dots, w_m) = \frac{\exp(\beta_{w_{m+1}} \cdot \mathbf{h}_m)}{\sum_{w' \in \mathcal{V}} \exp(\beta_{w'} \cdot \mathbf{h}_m)}, \quad (5.30)$$

where ϕ is a matrix of **input word embeddings**, and \mathbf{x}_m denotes the embedding for word w_m . The function g is an element-wise nonlinear **activation function**. Typical choices are:

- $\tanh(x)$, the hyperbolic tangent;
- $\sigma(x)$, the **sigmoid function** $\frac{1}{1+\exp(-x)}$;
- $(x)_+$, the **rectified linear unit**, $(x)_+ = \max(x, 0)$, also called **ReLU**.

These activation functions are shown in Figure 5.2. The sigmoid and tanh functions “squash” their inputs into a fixed range: $[0, 1]$ for the sigmoid, $[-1, 1]$ for tanh. This makes it possible to chain together many instances of these functions without numerical instability.

A key point about the RNN language model is that although each w_m depends only on the context vector \mathbf{h}_{m-1} , this vector is in turn influenced by **all** previous tokens, w_1, w_2, \dots, w_{m-1} , through the recurrence operation: w_1 affects \mathbf{h}_1 , which affects \mathbf{h}_2 , and so on, until the information is propagated all the way to \mathbf{h}_{m-1} , and then on to w_m (see Figure 5.1). This is an important distinction from n -gram language models, where any information outside




Figure 5.2: Nonlinear activation functions for neural networks

the n -word window is ignored. Thus, in principle, the RNN language model can handle long-range dependencies, such as number agreement over long spans of text — although it would be difficult to know where exactly in the vector \mathbf{h}_m this information is represented. The main limitation is that information is attenuated by repeated application of the nonlinearity g . **Long short-term memories** (LSTMs), described below, are a variant of RNNs that address this issue, using memory cells to propagate information through the sequence without applying non-linearities (Hochreiter and Schmidhuber, 1997).

The denominator in Equation 5.30 is a computational bottleneck, because it involves a sum over the entire vocabulary. One solution is to use a **hierarchical softmax** function, which computes the sum more efficiently by organizing the vocabulary into a tree (Mikolov et al., 2011). Another strategy is to optimize an alternative metric, such as **noise-contrastive estimation** (Gutmann and Hyvärinen, 2012), which learns by distinguishing observed instances from artificial instances generated from a noise distribution (Mnih and Teh, 2012).

Estimation by backpropagation

The recurrent neural network language model has the following parameters:

- $\phi_i \in \mathbb{R}^K$, the “input” word vectors (these are sometimes called **word embeddings**, since each word is embedded in a K -dimensional space);
- $\beta_i \in \mathbb{R}^K$, the “output” word vectors;
- $\Theta \in \mathbb{R}^{K \times K}$, the recurrence operator.

Each of these parameters must be estimated. We do this by formulating an objective function over the training corpus, $\ell(w)$, and then employ **backpropagation** to incrementally update the parameters after encountering each training example. Backpropagation is a term from the neural network literature, which means that we use the chain rule of

differentiation to obtain gradients on each parameter. After obtaining these gradients, we can apply an online learning algorithm such as stochastic gradient descent or adagrad, as discussed in § 2.4.

For example, suppose we want to obtain the gradient of the log-likelihood with respect to a single row of the recurrence operator, θ_k . Let us first define the total objective as a sum over local error functions e_m ,

$$\ell(\mathbf{w}) = \sum_{m=1}^M e_m(\mathbf{h}_{m-1}) \quad (5.31)$$

$$e_m(\mathbf{h}_{m-1}) \triangleq -\log p(w_m \mid w_1, w_2, \dots, w_{m-1}) \quad (5.32)$$

$$= -\beta_{w_m} \cdot \mathbf{h}_{m-1} + \log \sum_{w' \in \mathcal{V}} \exp(\beta_{w'} \cdot \mathbf{h}_{m-1}). \quad (5.33)$$

We can now differentiate the objective with respect to θ_k :

$$\frac{\partial}{\partial \theta_k} \ell(\mathbf{w}) = \sum_m \frac{\partial e_m(\mathbf{h}_{m-1})}{\partial \theta_k} \quad (5.34)$$

$$= \sum_{m=1}^M (\nabla_{\mathbf{h}_{m-1}} e_m) \frac{\partial}{\partial \theta_k} \mathbf{h}_{m-1}. \quad (5.35)$$

In the first line, we simply distribute the derivative across the sum. In the second line, we apply the chain rule of calculus. The term $\nabla_{\mathbf{h}_{m-1}} e_m$ refers to the gradient of the error e_m evaluated at \mathbf{h}_{m-1} , and is equal to,

$$\nabla_{\mathbf{h}_{m-1}} e_m = -\beta_{w_m} + B \text{SoftMax}(B \mathbf{h}_{m-1}), \quad (5.36)$$

where B is a matrix with all word output embeddings stacked vertically, $B = (\beta_1^\top, \beta_2^\top, \dots, \beta_{|\mathcal{V}|}^\top)$.

Next we compute the derivative of \mathbf{h}_{m-1} , first noting that within the vector \mathbf{h}_{m-1} , only the element $h_{m-1,k}$ depends on θ_k .

$$\frac{\partial}{\partial \theta_k} \mathbf{h}_{m-1} = \frac{\partial}{\partial \theta_k} h_{m-1,k} \quad (5.37)$$

$$= \frac{\partial}{\partial \theta_k} g(\theta_k \cdot \mathbf{h}_{m-2} + x_{m-1,k}) \quad (5.38)$$

$$= (\nabla_{\theta_k \cdot \mathbf{h}_{m-2} + x_{m-1,k}} g) \times \frac{\partial}{\partial \theta_k} (\theta_k \cdot \mathbf{h}_{m-2} + x_{m-1,k}) \quad (5.39)$$

$$= (\nabla_{\theta_k \cdot \mathbf{h}_{m-2} + x_{m-1,k}} g) \times (\mathbf{h}_{m-2} + \theta_k \odot \frac{\partial}{\partial \theta_k} \mathbf{h}_{m-2}), \quad (5.40)$$

where \odot is an elementwise (Hadamard) vector product, and $(\nabla_{\theta_k \cdot \mathbf{h}_{m-2} + x_{m-1,k}} g)$ is the elementwise gradient of the non-linear activation function for \mathbf{h}_{m-1} evaluated at the scalar

$\theta_k \cdot \mathbf{h}_{m-2} + x_{m-1,k}$. For example, if g is the elementwise hyperbolic tangent, then its gradient is,

$$(\nabla_{\theta_k \cdot \mathbf{h}_{m-2} + x_{m-1,k}} g) = (1 - \tanh^2(\theta_k \cdot \mathbf{h}_{m-2} + x_{m-1,k})). \quad (5.41)$$

The application of backpropagation to sequence models such as recurrent neural networks is known as **backpropagation through time**. A key point is that the derivative $\frac{\partial \mathbf{h}_{m-1}}{\partial \theta_k}$ depends recurrently on $\frac{\partial \mathbf{h}_{m-2}}{\partial \theta_k}$, and on all $\frac{\partial \mathbf{h}_n}{\partial \theta_k}$ for $n < m$. Furthermore, we will need to compute $\frac{\partial \mathbf{h}_{m-2}}{\partial \theta_k}$ **again**, to account for the error term $e_{m-1}(\mathbf{h}_{m-2})$. To avoid redoing work, it is best to cache such derivatives, so that they can be reused during backpropagation.

Backpropagation is implemented by neural network toolkits such as TensorFlow (Abadi et al., 2016), Torch (Collobert et al., 2011), and DyNet (Neubig et al., 2017). In these toolkits, the user defines a **computation graph** representing the neural network structure, which culminates in a scalar loss function. The toolkit then automatically computes the gradient of the loss function with respect to all model parameters, by applying the chain rule of differentiation across the computation graph. Unlike the classification objectives considered in chapter 2, neural network objectives are usually non-convex function of the parameters, so there is no learning procedure that is guaranteed to converge to the global optimum. Nonetheless, gradient-based optimization often yields parameter estimates that are very effective in practice.

Hyperparameters

The RNN language model has several hyperparameters that must be tuned to ensure good performance. The model capacity is controlled by the size of the word and context vectors K , which play a role that is somewhat analogous to the size of the n -gram context. For datasets that are large with respect to the vocabulary (i.e., there is a large token-to-type ratio), we can afford to estimate a model with a large K , which enables more subtle distinctions between words and contexts. When the dataset is relatively small, then K must be smaller too. However, this general advice has not yet been formalized into any concrete formula for choosing K , and trial-and-error is still necessary. Overfitting can also be prevented by **dropout**, which involves randomly setting some elements of the computation to zero (Srivastava et al., 2014), forcing the learner not to rely too much on any particular dimension of the word or context vectors. (The dropout rate must also be tuned by the user.) Other design decisions include: the nature of the nonlinear activation function g , the size of the vocabulary, and the parametrization of the learning algorithm, such as the learning rate.

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 5.3: The long short-term memory (LSTM) architecture. For clarity, only the variables (and not the parameters) are shown. Gates are shown in shaded boxes. In an LSTM language model, each h_m would be used to predict the next word w_{m+1} .

Alternative neural language models

A well known problem with RNNs is that backpropagation across long chains tends to lead to “vanishing” or “exploding” gradients (Bengio et al., 1994). For example, the input embedding of word w_1 affects the likelihood of a distant word such as w_{29} , but this impact may be attenuated by backpropagation through the intervening time steps. One solution is to rescale the gradients, or to clip them at some maximum value (Pascanu et al., 2013). An alternative is to change the model architecture itself.

A popular variant of RNNs, which is more robust to these problems, is the **long short-term memory (LSTM)** (Hochreiter and Schmidhuber, 1997; Sundermeyer et al., 2012). This model augments the hidden state h_m with a “memory cell” c_m . The value of the memory cell at each time m is a linear interpolation between two quantities: its previous value c_{m-1} , and an “update” \tilde{c}_m , which is computed from the current input x_m and the previous hidden state h_{m-1} . The next state h_m is then computed from the memory cell. Because the memory cell is never passed through the non-linear function g , it is possible for information to propagate through the network over long distances.

The interpolation weights are controlled by a set of gates, which are themselves functions of the input and previous hidden state. The gates are computed from sigmoid activations, ensuring that their values will be in the range $[0, 1]$. They can therefore be viewed as soft, differentiable logic gates. The LSTM architecture is shown in Figure 5.3, and the

complete update equations are:

$$\mathbf{f}_m = \sigma(\Theta^{(h \rightarrow f)} \cdot \mathbf{h}_{m-1} + \Theta^{(x \rightarrow f)} \cdot \mathbf{x}_m) \quad \text{forget gate} \quad (5.42)$$

$$\mathbf{i}_m = \sigma(\Theta^{(h \rightarrow i)} \cdot \mathbf{h}_{m-1} + \Theta^{(x \rightarrow i)} \cdot \mathbf{x}_m) \quad \text{input gate} \quad (5.43)$$

$$\tilde{\mathbf{c}}_m = \tanh(\Theta^{(h \rightarrow c)} \cdot \mathbf{h}_{m-1} + \Theta^{(w \rightarrow c)} \cdot \mathbf{x}_m) \quad \text{update candidate} \quad (5.44)$$

$$\mathbf{c}_m = \mathbf{f}_m \odot \mathbf{c}_{m-1} + \mathbf{i}_m \odot \tilde{\mathbf{c}}_m \quad \text{memory cell update} \quad (5.45)$$

$$\mathbf{o}_m = \sigma(\Theta^{(h \rightarrow o)} \cdot \mathbf{h}_{m-1} + \Theta^{(x \rightarrow o)} \cdot \mathbf{x}_m) \quad \text{output gate} \quad (5.46)$$

$$\mathbf{h}_m = \mathbf{o}_m \odot \mathbf{c}_m \quad \text{output.} \quad (5.47)$$

As above, \odot refers to an elementwise (Hadamard) product. The LSTM model has been shown to outperform standard recurrent neural networks across a wide range of problems (it was first used for language modeling by Sundermeyer et al. (2012)), and is now widely used for sequence modeling tasks. There are several LSTM variants, of which the Gated Recurrent Unit (Cho et al., 2014) is presently one of the more well known. Many software packages implement a variety of RNN architectures, so choosing between them is simple from a user’s perspective. Jozefowicz et al. (2015) provide an empirical comparison of various modeling choices circa 2015. Notable earlier non-recurrent architectures include the neural probabilistic language model (Bengio et al., 2003) and the log-bilinear language model (Mnih and Hinton, 2007). Much more detail on these models can be found in the text by Goodfellow et al. (2016).

5.4 Evaluating language models

Because language models are typically components of larger systems — language modeling is not usually an application itself — we would prefer **extrinsic evaluation**. This means evaluating whether the language model improves performance on the application task, such as machine translation or speech recognition. But this is often hard to do, and depends on details of the overall system which may be irrelevant to language modeling. In contrast, **intrinsic evaluation** is task-neutral. Better performance on intrinsic metrics may be expected to improve extrinsic metrics across a variety of tasks, unless we are over-optimizing the intrinsic metric. We will discuss intrinsic metrics here, but bear in mind that it is important to also perform extrinsic evaluations to ensure that the improvements obtained on these intrinsic metrics really carry over to the applications that we care about.

Held-out likelihood

The goal of probabilistic language models is to accurately measure the probability of sequences of word tokens. Therefore, an intrinsic evaluation metric is the likelihood that the language model assigns to **held-out data**, which is not used during training. Specifically,

(c) Jacob Eisenstein 2014-2017. Work in progress.

we compute,

$$\ell(\mathbf{w}) = \sum_{m=1}^M \log p(w_m \mid w_{m-1}, \dots), \quad (5.48)$$

treating the entire held-out corpus as a single stream of tokens.

Typically, unknown words are mapped to the $\langle \text{UNK} \rangle$ token. This means that we have to estimate some probability for $\langle \text{UNK} \rangle$ on the training data. One way to do this is to fix the vocabulary \mathcal{V} to the $|\mathcal{V}| - 1$ words with the highest counts in the training data, and then convert all other tokens to $\langle \text{UNK} \rangle$. Other strategies for dealing with out-of-vocabulary terms are discussed in § 5.5.

Perplexity

Held-out likelihood is usually presented as **perplexity**, which is a deterministic transformation of the log-likelihood into an information-theoretic quantity,

$$\text{Perplex}(\mathbf{w}) = 2^{-\frac{\ell(\mathbf{w})}{M}}, \quad (5.49)$$

where M is the total number of tokens in the held-out corpus.

Lower perplexities correspond to higher likelihoods, so lower scores are better on this metric. (How to remember: lower perplexity is better, because your language model is less perplexed.) To understand perplexity, here are some special cases:

- In the limit of a perfect language model, probability 1 is assigned to the held-out corpus, with $\text{Perplex}(\mathbf{w}) = 2^{-\frac{1}{M} \log_2 1} = 2^0 = 1$.
- In the opposite limit, probability zero is assigned to the held-out corpus, which corresponds to an infinite perplexity, $\text{Perplex}(\mathbf{w}) = 2^{-\frac{1}{M} \log_2 0} = 2^\infty = \infty$.
- Assume a uniform, unigram model in which $p(w_i) = \frac{1}{|\mathcal{V}|}$ for all words in the vocabulary. Then,

$$\begin{aligned} \log_2(\mathbf{w}) &= \sum_{m=1}^M \log_2 \frac{1}{|\mathcal{V}|} = - \sum_{m=1}^M \log_2 |\mathcal{V}| = -M \log_2 |\mathcal{V}| \\ \text{Perplex}(\mathbf{w}) &= 2^{\frac{1}{M} M \log_2 |\mathcal{V}|} \\ &= 2^{\log_2 |\mathcal{V}|} \\ &= |\mathcal{V}|. \end{aligned}$$

This is the “worst reasonable case” scenario, since you could build such a language model without even looking at the data.

(c) Jacob Eisenstein 2014-2017. Work in progress.

In practice, n -gram language models tend to give perplexities in the range between 1 and $|\mathcal{V}|$. For example, Jurafsky and Martin estimate a language model over a vocabulary of roughly 20,000 words, on 38 million tokens of text from the Wall Street Journal (Jurafsky and Martin, 2009, page 97). They report the following perplexities on a held-out set of 1.5 million tokens:

- Unigram ($n = 1$): 962
- Bigram ($n = 2$): 170
- Trigram ($n = 3$): 109

Will this trend continue?

5.5 Out-of-vocabulary words

Through this chapter, we have assumed a **closed-vocabulary** setting — the vocabulary \mathcal{V} is assumed to be a finite set. In realistic application scenarios, this assumption may not hold. Consider, for example, the problem of translating newspaper articles. The following sentence appeared in a Reuters article on January 6, 2017:⁵

The report said U.S. intelligence agencies believe Russian military intelligence, the **GRU**, used intermediaries such as **WikiLeaks**, **DCLeaks.com** and the **Guccifer 2.0** “persona” to release emails...

Suppose that you trained a language model on the Gigaword corpus,⁶ which was released in 2003. The bolded terms either did not exist at this date, or were not widely known; they are unlikely to be in the vocabulary. The same problem can occur for a variety of other terms: new technologies, previously unknown individuals, new words (e.g., *hashtag*), and numbers.

One solution is to simply mark all such terms with a special token, $\langle \text{UNK} \rangle$. While training the language model, we decide in advance on the vocabulary (often the K most common terms), and mark all other terms in the training data as $\langle \text{UNK} \rangle$. If we do not want to determine the vocabulary size in advance, an alternative approach is to simply mark the first occurrence of each word type as $\langle \text{UNK} \rangle$.

In some scenarios, we may prefer to make distinctions about the likelihood of various unknown words. This is particularly important in languages that have rich morphological systems, with many inflections for each word. For example, Spanish is only moderately complex from a morphological perspective, yet each verb has dozens of inflected forms.

⁵Bayoumy, Y. and Strobel, W. (2017, January 6). U.S. intel report: Putin directed cyber campaign to help Trump. *Reuters*. Retrieved from <http://www.reuters.com/article/us-usa-russia-cyber-idUSKBN14Q1T8> on January 7, 2017.

⁶<https://catalog.ldc.upenn.edu/LDC2003T05>

In such languages, there will necessarily be many word types that we do not encounter in a corpus, which are nonetheless predictable from the morphological rules of the language. To use a somewhat contrived English example, if *transfenestrate* is in the vocabulary, our language model should assign a non-zero probability to the past tense *transfenestrated*, even if it does not appear in the training data.

One way to accomplish this is to supplement word-level language models with **character-level language models**. Such models can use n -grams or RNNs, but with a fixed vocabulary equal to the set of ASCII or Unicode characters. For example Ling et al. (2015b) propose an LSTM model over characters, and Kim (2014) employ a **convolutional neural network** (LeCun and Bengio, 1995). A more linguistically motivated approach is to segment words into meaningful subword units, known as **morphemes** (see chapter 9). For example, Botha and Blunsom (2014) induce vector representations for morphemes, which they build into a log-bilinear language model; Bhatia et al. (2016) incorporate morpheme vectors into an LSTM.

Exercises

1. exercise tk

Part II

Sequences and trees

Chapter 6

Sequence labeling

In sequence labeling, we want to assign tags to words, or more generally, we want to assign discrete labels to elements in a sequence. There are many applications of sequence labeling in natural language processing, and chapter 7 presents an overview. One of the most classic application of sequence labeling is **part-of-speech tagging**, which involves tagging each word by its grammatical category. Coarse-grained grammatical categories include **NOUNs**, which describe things, properties, or ideas, and **VERBs**, which describe actions and events. Given a simple sentence like,

(6.1) They can fish.

we would like to produce the tag sequence N V V, with the modal verb *can* labeled as a verb in this simplified example.

6.1 Sequence labeling as classification

One way to solve tagging problems is to treat them as classification. We can write $f((\mathbf{w}, m), y)$ to indicate the feature function for applying tag y to word w_m in the sequence w_1, w_2, \dots, w_M . A simple tagging model would have a single base feature, the word itself:

$$f((\mathbf{w} = \text{they can fish}, m = 1), N) = \langle \text{they}, N \rangle \quad (6.1)$$

$$f((\mathbf{w} = \text{they can fish}, m = 2), V) = \langle \text{can}, V \rangle \quad (6.2)$$

$$f((\mathbf{w} = \text{they can fish}, m = 3), V) = \langle \text{fish}, V \rangle. \quad (6.3)$$

Here the feature function takes three arguments as input: the sentence to be tagged (*they can fish* in all cases), the proposed tag (e.g., N or V), and the word token to which this tag is applied. This simple feature function then returns a single feature: a tuple including the word to be tagged and the tag that has been proposed. If the vocabulary size is V and the number of tags is K , then there are $V \times K$ features. Each of these features must

be assigned a weight. These weights can be learned from a labeled dataset using a classification algorithm such as perceptron, but this isn't necessary in this case: it would be equivalent to define the classification weights directly, with $\theta_{w,y} = 1$ for the tag y most frequently associated with word w , and $\theta_{w,y} = 0$ for all other tags.

However, it is easy to see that this simple classification approach can go wrong. Consider the word *fish*, which often describes an animal rather than an activity; in these cases, *fish* should be tagged as a noun. To tag ambiguous words correctly, the tagger must rely on context, such as the surrounding words. We can build this context into the feature set by incorporating the surrounding words as additional features:

$$\begin{aligned} f((w = \text{they can fish}, 1), N) = \{ & \langle w_i = \text{they}, y_i = N \rangle, \\ & \langle w_{i-1} = \diamond, y_i = N \rangle, \\ & \langle w_{i+1} = \text{can}, y_i = N \rangle \} \end{aligned} \quad (6.4)$$

$$\begin{aligned} f((w = \text{they can fish}, 2), V) = \{ & \langle w_i = \text{can}, y_i = V \rangle, \\ & \langle w_{i-1} = \text{they}, y_i = V \rangle, \\ & \langle w_{i+1} = \text{fish}, y_i = V \rangle \} \end{aligned} \quad (6.5)$$

$$\begin{aligned} f((w = \text{they can fish}, 3), V) = \{ & \langle w_i = \text{fish}, y_i = V \rangle, \\ & \langle w_{i-1} = \text{can}, y_i = V \rangle, \\ & \langle w_{i+1} = \blacklozenge, y_i = V \rangle \}. \end{aligned} \quad (6.6)$$

These features contain enough information that a tagger should be able to choose the right label for the word *fish*: words that follow the modal verb *can* are likely to be verbs themselves, so the feature $\langle w_{i-1} = \text{can}, y_i = V \rangle$ should have a large positive weight.

However, even with this enhanced feature set, it may be difficult to tag some sequences correctly. One reason is that there are often relationships between the tags themselves. For example, in English it is relatively rare for a verb to follow another verb — particularly if we differentiate MODAL verbs like *can* and *should* from more typical verbs, like *give*, *transcend*, and *befuddle*. We would like to incorporate preferences **against** such tag sequences, and preferences **for** other tag sequences, such as NOUN-VERB.

The need for such preferences is best illustrated by a **garden path sentence**:

(6.2) The old man the boat.

Grammatically, the word *the* is a DETERMINER. When you read the sentence, what part of speech did you first assign to *old*? Typically, this word is an ADJECTIVE — abbreviated as J — which is a class of words that modify nouns. Similarly, *man* is usually a noun. The resulting sequence of tags is D J N D N. But this is a mistaken “garden path” interpretation, which ends up leading nowhere. It is unlikely that a determiner would directly follow a noun,¹ and particularly unlikely that the entire sentence would lack a verb. The only possible verb in the sentence is the word *man*, which can refer to the act of maintaining and piloting something — often boats. But if *man* is tagged as a verb, then *old* is seated

¹The main exception is the double object construction, as in *I gave the child a toy*.

between a determiner and a verb, and must be a noun. And indeed, adjectives can often have a second interpretation as nouns when used in this way (e.g., *the young*, *the restless*). This reasoning, in which the labeling decisions are intertwined, cannot be applied in a setting where each tag is produced by an independent classification decision.

6.2 Sequence labeling as structure prediction

As an alternative, we can think of the entire sequence of tags as a label itself. For a given sequence of words $\mathbf{w}_{1:M} = (w_1, w_2, \dots, w_M)$, there is a set of possible taggings $\mathcal{Y}(\mathbf{w}_{1:M}) = \mathcal{Y}^M$, where $\mathcal{Y} = \{N, V, D, \dots\}$ refers to the set of individual tags, and \mathcal{Y}^M refers to the set of tag sequences of length M . We can then treat the sequence labeling problem as a classification problem in the label space $\mathcal{Y}(\mathbf{w}_{1:M})$,

$$\hat{\mathbf{y}}_{1:M} = \underset{\mathbf{y}_{1:M} \in \mathcal{Y}(\mathbf{w}_{1:M})}{\operatorname{argmax}} \quad \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}_{1:M}, \mathbf{y}_{1:M}), \quad (6.7)$$

where $\mathbf{y}_{1:M} = (y_1, y_2, \dots, y_M)$ is a sequence of M tags. Note that in this formulation, we have a feature function that consider the entire tag sequence $\mathbf{y}_{1:M}$. Such a feature function can therefore include features that capture the relationships between tagging decisions, such as the preference that determiners not follow nouns, or that all sentences have verbs.

Given that the label space is exponentially large in the length of the sequence w_1, \dots, w_M , can it ever be practical to perform tagging in this way? The problem of making a series of interconnected labeling decisions is known as **inference**. Because natural language is full of interrelated grammatical structures, inference is a crucial aspect of contemporary natural language processing. In English, it is not unusual to have sentences of length $M = 20$; part-of-speech tag sets vary in size from 10 to several hundred. Taking the low end of this range, we have $|\mathcal{Y}(\mathbf{w}_{1:M})| \approx 10^{20}$, one hundred billion billion possible tag sequences. Enumerating and scoring each of these sequences would require an amount of work that is exponential in the sequence length, so inference is intractable.

However, the situation changes when we restrict the feature function. Suppose we choose features that never consider more than one tag. We can indicate this restriction as,

$$\mathbf{f}(\mathbf{w}, \mathbf{y}) = \sum_{m=1}^M \mathbf{f}(\mathbf{w}, y_m, m), \quad (6.8)$$

where we use the shorthand $\mathbf{w} \triangleq \mathbf{w}_{1:M}$. The summation in (6.8) means that the overall feature vector is the sum of feature vectors associated with each individual tagging decision. These features are not capable of capturing the intuitions that might help us solve garden path sentences, such as the insight that determiners rarely follow nouns in English. But this restriction does make it possible to find the globally optimal tagging, by

(c) Jacob Eisenstein 2014-2017. Work in progress.

making a sequence of individual tagging decisions.

$$\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \mathbf{y}) = \boldsymbol{\theta} \cdot \sum_{m=1}^M \mathbf{f}(\mathbf{w}, y_m, m) \quad (6.9)$$

$$= \sum_{m=1}^M \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, m) \quad (6.10)$$

$$\hat{y}_m = \operatorname{argmax}_{y_m} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, m) \quad (6.11)$$

$$\hat{\mathbf{y}} = (\hat{y}_1, \hat{y}_2, \dots, \hat{y}_M) \quad (6.12)$$

Note that we are still searching over an exponentially large set of tag sequences! But the feature set restriction results in decoupling the labeling decisions that were previous interconnected. As a result, it is not necessary to score every one of the $|\mathcal{Y}|^M$ tag sequences individually — we can find the optimal sequence by scoring the local parts of these decisions.

Now let's consider a slightly less restrictive feature function: rather than considering only individual tags, we will consider adjacent tags too. This means that we can have negative weights for infelicitous tag pairs, such as noun-determiner, and positive weights for typical tag pairs, such as determiner-noun and noun-verb. We define this feature function as,

$$\mathbf{f}(\mathbf{w}, \mathbf{y}) = \sum_{m=1}^M \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m). \quad (6.13)$$

Let's apply this feature function to the shorter example, *they can fish*, using features for word-tag and tag-tag pairs:

$$\mathbf{f}(\mathbf{w} = \text{they can fish}, \mathbf{y} = \text{N V V}) = \sum_{m=1}^M \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m) \quad (6.14)$$

$$\begin{aligned} &= \mathbf{f}(\mathbf{w}, \text{N}, \diamond, 1) \\ &\quad + \mathbf{f}(\mathbf{w}, \text{V}, \text{N}, 2) \\ &\quad + \mathbf{f}(\mathbf{w}, \text{V}, \text{V}, 3) \end{aligned} \quad (6.15)$$

$$\begin{aligned} &= \langle w_m = \text{they}, y_m = \text{N} \rangle + \langle y_m = \text{N}, y_{m-1} = \diamond \rangle \\ &\quad + \langle w_m = \text{can}, y_m = \text{V} \rangle + \langle y_m = \text{V}, y_{m-1} = \text{N} \rangle \\ &\quad + \langle w_m = \text{fish}, y_m = \text{V} \rangle + \langle y_m = \text{V}, y_{m-1} = \text{V} \rangle \\ &\quad + \langle y_m = \diamond, y_{m-1} = \text{V} \rangle. \end{aligned} \quad (6.16)$$

We end up with seven active features: one for each word-tag pair, and one for each tag-tag pair (this includes a final tag $y_{M+1} = \diamond$). These features capture what are arguably the two main sources of information for part-of-speech tagging: which tags are appropriate

(c) Jacob Eisenstein 2014-2017. Work in progress.

for each word, and which tags tend to follow each other in sequence. Given appropriate weights for these features, we can expect to make the right tagging decisions, even for difficult cases like *the old man the boat*.

The example shows that even with the restriction to the feature set shown in Equation 6.13, it is still possible to construct expressive features that are capable of solving many sequence labeling problems. But the key question is: does this restriction make it possible to perform efficient inference? The answer is yes, and the solution is the **Viterbi algorithm** (Viterbi, 1967).

6.3 The Viterbi algorithm

We now consider the inference problem,

$$\hat{y} = \operatorname{argmax}_y \theta \cdot f(w, y) \quad (6.17)$$

$$f(w, y) = \sum_{m=1}^M f(w, y_m, y_{m-1}, m). \quad (6.18)$$

Given this restriction on the feature function, we can solve this inference problem using **dynamic programming**, a algorithmic technique for reusing work in recurrent computations. As is often the case in dynamic programming, we begin by solving an auxiliary problem: rather than finding the best tag sequence, we simply try to compute the **score** of the best tag sequence,

$$\max_y \theta \cdot f(w, y) = \max_{y_{1:M}} \sum_{m=1}^M \theta \cdot f(w, y_m, y_{m-1}, m) \quad (6.19)$$

$$= \max_{y_{1:M}} \theta \cdot f(w, y_M, y_{M-1}, M) + \sum_{m=1}^{M-1} \theta \cdot f(w, y_m, y_{m-1}, m) \quad (6.20)$$

$$= \max_{y_M} \max_{y_{M-1}} \theta \cdot f(w, y_M, y_{M-1}, M) + \max_{y_{1:M-2}} \sum_{m=1}^{M-1} \theta \cdot f(w, y_m, y_{m-1}, m). \quad (6.21)$$

In this derivation, we first removed the final element $\theta \cdot f(w, y_M, y_{M-1}, M)$ from the sum over the sequence, and then we adjusted the scope of the the max operation, since the elements $(y_1 \dots y_{M-2})$ are irrelevant to the final term.

Let us now define the **Viterbi variable**,

$$v_m(k) \triangleq \max_{y_{1:m-1}} \theta \cdot f(w, k, y_{m-1}, m) + \sum_{n=1}^{m-1} \theta \cdot f(w, y_n, y_{n-1}, n), \quad (6.22)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

Algorithm 6 The Viterbi algorithm.

```

for  $k \in \{0, \dots, K\}$  do
   $v_1(k) = \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, k, \diamond, m)$ 
for  $m \in \{2, \dots, M\}$  do
  for  $k \in \{0, \dots, K\}$  do
     $v_m(k) = \max_{k'} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, k, k', m) + v_{m-1}(k')$ 
     $b_m(k) = \operatorname{argmax}_{k'} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, k, k', m) + v_{m-1}(k')$ 
   $y_M = \operatorname{argmax}_k v_M(k) + \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \diamond, k, M + 1)$ 
for  $m \in \{M - 1, \dots, 1\}$  do
   $y_m = b_m(y_{m+1})$ 
return  $y_{1:M}$ 

```

where lower-case m indicates any position in the sequence, and $k \in \mathcal{Y}$ indicates a tag for that position. The variable $v_m(k)$ represents the score of the best tag sequence $(\hat{y}_1, \hat{y}_2, \dots, \hat{y}_m)$ that terminates in $\hat{y}_m = k$. From this definition, we can compute the score of the best tagging of the sequence by plugging the Viterbi variables $v_M(\cdot)$ into Equation 6.21,

$$\max_{\mathbf{y}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \mathbf{y}) = \max_k v_M(k). \quad (6.23)$$

Now, let us look more closely at how we can compute these Viterbi variables.

$$v_m(k) \triangleq \max_{\mathbf{y}_{1:m-1}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, k, y_{m-1}, m) + \sum_{n=1}^{m-1} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_n, y_{n-1}, n) \quad (6.24)$$

$$= \max_{y_{m-1}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, k, y_{m-1}, m) + \max_{\mathbf{y}_{1:m-2}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_{m-1}, y_{m-2}) + \sum_{n=1}^{m-2} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_n, y_{n-1}, n) \quad (6.25)$$

$$= \max_{y_{m-1}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, k, y_{m-1}, m) + v_{m-1}(y_{m-1}) \quad (6.26)$$

$$v_1(y) = \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y, \diamond, 1). \quad (6.27)$$

Equation 6.26 is a **recurrence** for computing the Viterbi variables: each $v_m(k)$ can be computed in terms of $v_{m-1}(\cdot)$, and so on. We can therefore step forward through the sequence, computing first all variables $v_1(\cdot)$ from Equation 6.27, and then computing all variables $v_2(\cdot)$, $v_3(\cdot)$, and so on, until we reach the final set of variables $v_M(\cdot)$.

Graphically, it is customary to arrange these variables in a matrix, with the sequence index m on the columns, and the tag index k on the rows. In this representation, each $v_{m-1}(k)$ is connected to each $v_m(k')$, forming a **trellis**, as shown in Figure 6.1. As shown in the figure, special nodes are set aside for the start and end states.

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 6.1: The trellis representation of the Viterbi variables, for the example *they can fish*, using the weights shown in Table 6.1.

Our real goal is to find the best scoring sequence, not simply to compute its score. But as is often the case in dynamic programming, solving the auxiliary problem gets us almost all the way to our original goal. Recall that each $v_m(k)$ represents the score of the best tag sequence ending in that tag k in position m . To compute this, we maximize over possible values of y_{m-1} . If we keep track of the tag that maximizes this choice at each step, then we can walk backwards from the final tag, and recover the optimal tag sequence. This is indicated in Figure 6.1 by the solid blue lines, which we trace back from the final position. These “back-pointers” are written $b_m(k)$, indicating the optimal tag y_{m-1} on the path to $Y_m = k$.

Why does this work? We can make an inductive argument. Suppose k is indeed the optimal tag for word m , and we now need to decide on the tag y_{m-1} . Because we make the inductive assumption that we know $y_m = k$, and because the feature function is restricted to adjacent tags, we need not consider any of the tags $y_{m+1:M}$; these tags, and the features that describe them, are irrelevant to the inference of y_{m-1} , given that we have $y_m = k$. Thus, we are looking for the tag \hat{y}_{m-1} that maximizes,

$$\hat{y}_{m-1} = \operatorname{argmax}_{y_{m-1}} \theta \cdot \mathbf{f}(\mathbf{w}, k, y_{m-1}, m) + \max_{\mathbf{y}_{1:m-2}} \sum_{n=1}^{m-1} \theta \cdot \mathbf{f}(\mathbf{w}, y_n, y_{n-1}, n) \quad (6.28)$$

$$= \operatorname{argmax}_{y_{m-1}} \theta \cdot \mathbf{f}(\mathbf{w}, k, y_{m-1}, m) + v_{m-1}(y_{m-1}), \quad (6.29)$$

which we obtain by plugging in the definition of the Viterbi variable. The value \hat{y}_{m-1} was identified during forward pass, when computing the value of the Viterbi variable $v_m(k)$.

The complete Viterbi algorithm is shown in Algorithm 6. This formalizes the recurrences that were described in the previous paragraphs, and handles the boundary conditions at the start and end of the sequence. Specifically, when computing the initial Viterbi variables $v_1(\cdot)$, we use a special tag, \diamond , to indicate the start of the sequence. When com-

	<i>they</i>	<i>can</i>	<i>fish</i>
N	-2	-3	-3
V	-10	-1	-3

(a) Weights for emission features.

	N	V	◆
◇	-1	-2	$-\infty$
N	-3	-1	-12
V	-1	-3	-1

(b) Weights for transition features. The “from” tags are on the columns, and the “to” tags are on the rows.

Table 6.1: Feature weights for the example trellis shown in Figure 6.1. Emission weights from \diamond and \blacklozenge are implicitly set to $-\infty$.

putting the final tag Y_M , we use another special tag, \blacklozenge , to indicate the end of the sequence. These special tags enable the use of transition features for the tags that begin and end the sequence: for example, conjunctions are unlikely to end sentences in English, so we would like a large negative weight for the feature $\langle CC, \blacklozenge \rangle$; nouns are relatively likely to appear at the beginning of sentences, so we would like a more positive (or less negative) weight for the feature $\langle \diamond, N \rangle$.

What is the complexity of this algorithm? If there are K tags and M positions in the sequence, then there are $M \times K$ Viterbi variables to compute. Computing each variable requires finding a maximum over K possible predecessor tags. The total computation cost of populating the trellis is therefore $\mathcal{O}(MK^2)$, with an additional factor for the number of active features at each position. After completing the trellis, we simply trace the backwards pointers to the beginning of the sequence, which takes $\mathcal{O}(M)$ operations.

Example

To illustrate the Viterbi algorithm with an example, let us consider the minimal tagset $\{N, V\}$, corresponding to nouns and verbs. Even in this tagset, there is considerable ambiguity: for example, the words *can* and *fish* can each take both tags. Of the $2 \times 2 \times 2 = 8$ possible taggings for the sentence *they can fish*, four are possible given these possible tags, and two are grammatical. (The tagging *they*/N *can*/V *fish*/N corresponds to the scenario of putting fish into cans.)

To begin, we use the feature weights defined in Table 6.1. These weights are used to incrementally fill in the trellis. As described in Algorithm 6, we fill in the cells from left to right, with each column corresponding to a word in the sequence. As we fill in the cells, we must keep track of the back-pointers $b_m(k)$ — the previous cell that maximizes the score of tag k at word m . These are represented in the figure with the thick blue lines. At the end of the algorithm, we recover the optimal tag sequence by tracing back the optimal path from the final position, $(M + 1, \blacklozenge)$.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Higher-order features

The Viterbi algorithm was made possible by a restriction of the features to consider only pairs of adjacent tags. In a sense, we can think of this as a bigram language model, at the tag level. A natural question is how to generalize Viterbi to tag trigrams, which would involve the following feature decomposition:

$$f(\mathbf{w}, \mathbf{y}) = \sum_m^M f(\mathbf{w}, y_m, y_{m-1}, y_{m-2}, m). \quad (6.30)$$

One possibility is to take the Cartesian product of the tagset with itself, $\mathcal{Y}^{(2)} = \mathcal{Y} \times \mathcal{Y}$. The tags in this product space are ordered pairs, representing adjacent tags at the token level: for example, the tag $\langle N, V \rangle$ would represent a noun followed by a verb. Transitions between such tags must be consistent: we can have a transition from $\langle N, V \rangle$ to $\langle V, N \rangle$ (corresponding to the token-level tag sequence $N V N$), but not from $\langle N, V \rangle$ to $\langle N, N \rangle$, which would not correspond to any token-level tag sequence. This constraint can be enforced in the feature weights, with $\theta_{\langle\langle a,b \rangle, \langle c,d \rangle\rangle} = -\infty$ if $b \neq c$. The remaining feature weights can encode preferences for and against various tag trigrams.

In the Cartesian product tag space, there are K^2 tags, suggesting that the time complexity will increase to $\mathcal{O}(MK^4)$. However, it is unnecessary to max over predecessor tag bigrams that are incompatible with the current tag bigram. By exploiting these constraints, it is possible to limit the time complexity to $\mathcal{O}(MK^3)$. The space complexity is $\mathcal{O}(MK^2)$. In general, the time and space complexity of higher-order Viterbi grows exponentially with the order of the tag n -grams that are considered in the feature decomposition.

6.4 Hidden Markov Models

We now consider how to learn the weights θ that parametrize the Viterbi sequence labeling algorithm. We begin with a probabilistic approach. Recall that the probabilistic Naïve Bayes classifier selects the label y to maximize $p(y | \mathbf{x}) \propto p(y, \mathbf{x})$. In probabilistic sequence labeling, our goal is similar: select the tag sequence that maximizes $p(\mathbf{y} | \mathbf{w}) \propto p(\mathbf{y}, \mathbf{w})$. Just as Naïve Bayes could be cast as a linear classifier maximizing $\theta \cdot \mathbf{f}(\mathbf{x}, y)$, we can cast our probabilistic classifier as a linear decision rule. Furthermore, the feature restriction in (6.13) can be viewed as a conditional independence assumption on the random variables \mathbf{y} . Thanks to this assumption, it is possible to perform inference using the Viterbi algorithm.

Naïve Bayes was introduced as a generative model — a probabilistic story that explains the observed data as well as the hidden label. A similar story can be constructed for probabilistic sequence labeling: first, we draw the tags from a prior distribution, $\mathbf{y} \sim p(\mathbf{y})$; next, we draw the tokens from a conditional likelihood distribution, $\mathbf{w} | \mathbf{y} \sim p(\mathbf{w} | \mathbf{y})$.

(c) Jacob Eisenstein 2014-2017. Work in progress.

However, for inference to be tractable, additional independence assumptions are required. Here we make two assumptions. First, the probability of each token depends only on its tag, and not on any other element in the sequence:

$$p(\mathbf{w} | \mathbf{y}) = \prod_{m=1}^M p(w_m | y_m). \quad (6.31)$$

Next, we introduce an independence assumption on the form of the prior distribution over labels: each label y_m depends only on its predecessor,

$$p(\mathbf{y}) = \prod_{m=1}^M p(y_m | y_{m-1}), \quad (6.32)$$

where $y_0 = \diamond$ in all cases. Due to this **Markov** assumption, probabilistic sequence labeling models are known as **hidden Markov models** (HMMs). We now state the generative model under these independence assumptions,

- For $m \in (1, 2, \dots, M)$,
 - draw $y_m | y_{m-1} \sim \text{Categorical}(\lambda_{y_{m-1}})$;
 - draw $w_m | y_m \sim \text{Categorical}(\phi_{y_m})$

This generative story formalizes the hidden Markov model. Given the parameters λ and ϕ , we can compute $p(\mathbf{w}, \mathbf{y})$ for any token sequence \mathbf{w} and tag sequence \mathbf{y} . The HMM is often represented as a **graphical model** (Wainwright and Jordan, 2008), as shown in Figure 6.2. This representation makes the independence assumptions explicit: if a variable v_1 is probabilistically conditioned on another variable v_2 , then there is an arrow $v_2 \rightarrow v_1$ in the diagram. If there are no arrows between v_1 and v_2 , they are **conditionally independent**, given each variable’s **Markov blanket**. In the hidden Markov model, the Markov blanket for each tag y_m includes the “parent” y_{m-1} , and the “children” y_{m+1} and w_m .²

It is important to reflect on the implications of the HMM independence assumptions. A non-adjacent pair of tags y_m and y_n are conditionally independent; if $m < n$ and we are given y_{n-1} , then y_m offers no additional information about y_n . However, if we are not given any information about the tags in a sequence, then all tags are probabilistically coupled.

Estimation

The hidden Markov model has two groups of parameters:

²In general graphical models, a variable’s Markov blanket includes its parents, children, and its children’s other parents (Murphy, 2012).




Figure 6.2: Graphical representation of the hidden Markov model. Arrows indicate probabilistic dependencies.

Emission probabilities. The probability $p_e(w_m | y_m; \phi)$ is the emission probability, since the words are treated as probabilistically “emitted”, conditioned on the tags.

Transition probabilities. The probability $p_t(y_m | y_{m-1}; \lambda)$ is the transition probability, since it assigns probability to each possible tag-to-tag transition.

Both of these groups of parameters are typically computed from relative frequency estimation on a labeled corpus,

$$\phi_{k,i} \triangleq \Pr(W_m = i | Y_m = k) = \frac{\text{count}(W_m = i, Y_m = k)}{\text{count}(Y_m = k)}$$

$$\lambda_{k,k'} \triangleq \Pr(Y_m = k' | Y_{m-1} = k) = \frac{\text{count}(Y_m = k', Y_{m-1} = k)}{\text{count}(Y_{m-1} = k)}.$$

Smoothing is more important for the emission probability than the transition probability, because the event space is much larger. Smoothing techniques such as additive smoothing, interpolation, and backoff (see chapter 5) can all be applied here.

Inference

The goal of inference in the hidden Markov model is to find the highest probability tag sequence,

$$\hat{\mathbf{y}} = \underset{\mathbf{y}}{\operatorname{argmax}} p(\mathbf{y} | \mathbf{w}). \quad (6.33)$$

As in Naïve Bayes, it is equivalent to find the tag sequence with the highest **log**-probability, since the log function is monotonically increasing. It is furthermore equivalent to maximize the joint probability $p(\mathbf{y}, \mathbf{w}) = p(\mathbf{y} | \mathbf{w}) \times p(\mathbf{w}) \propto p(\mathbf{y} | \mathbf{w})$, which is proportional to the conditional probability. Therefore, we can reformulate the inference problem as,

$$\hat{\mathbf{y}} = \underset{\mathbf{y}}{\operatorname{argmax}} \log p(\mathbf{y}, \mathbf{w}). \quad (6.34)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

We can now apply the HMM independence assumptions:

$$\log p(\mathbf{y}, \mathbf{w}) = \log p(\mathbf{y}) + \log p(\mathbf{w} | \mathbf{y}) \quad (6.35)$$

$$= \sum_{m=1}^M \log p_y(y_m | y_{m-1}) + \log p_{w|y}(w_m | y_m) \quad (6.36)$$

$$= \sum_{m=1}^M \log \lambda_{y_m, y_{m-1}} + \log \phi_{y_m, w_m}. \quad (6.37)$$

This log probability can be rewritten as a dot product of weights and features,

$$\log p(\mathbf{y}, \mathbf{w}) = \sum_{m=1}^M \log \lambda_{y_m, y_{m-1}} + \log \phi_{y_m, w_m} \quad (6.38)$$

$$= \sum_{m=1}^M \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m), \quad (6.39)$$

where the feature function is defined,

$$\mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m) = \{\langle y_m, y_{m-1} \rangle, \langle y_m, w_m \rangle\}, \quad (6.40)$$

and the weight vector $\boldsymbol{\theta}$ encodes the log-parameters $\log \lambda$ and $\log \phi$.

This derivation shows that HMM inference can be viewed as an application of the Viterbi decoding algorithm, given an appropriately defined feature function and weight vector. The local product $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m)$ can be interpreted probabilistically,

$$\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m) = \log p_y(y_m | y_{m-1}) + \log p_{w|y}(w_m | y_m) \quad (6.41)$$

$$= \log p(y_m, w_m | y_{m-1}). \quad (6.42)$$

Now recall the definition of the Viterbi variables,

$$v_m(k) = \max_{\mathbf{y}_{1:m-1}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, Y_m = k, y_{m-1}, m) + \sum_{n=1}^{m-1} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_n, y_{n-1}, n) \quad (6.43)$$

$$= \max_{\mathbf{y}_{1:m-1}} \log p_{y_m, w_m | y_{m-1}}(k, w_m | y_{m-1}) + \sum_{n=1}^{m-1} \log p(y_n, w_n | y_{n-1}) \quad (6.44)$$

$$= \max_{\mathbf{y}_{1:m-1}} \log p(\mathbf{y}_{1:m-1}, Y_m = k, \mathbf{w}_{1:m}). \quad (6.45)$$

In words, the Viterbi variable $v_m(k)$ is the log probability of the best tag sequence ending in $Y_m = k$, joint with the word sequence $\mathbf{w}_{1:m}$. The log probability of the best complete tag sequence is therefore,

$$\max_{\mathbf{y}_{1:M}} \log p(\mathbf{y}_{1:M}, \mathbf{w}_{1:M}) = \max_{y_M} \log p_y(\diamond | y_M) + v_M(y_M). \quad (6.46)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

The Viterbi algorithm can also be implemented using probabilities, rather than log probabilities. In this case, each $v_m(k)$ is equal to,

$$v_m(k) = \max_{\mathbf{y}_{1:m-1}} p(\mathbf{y}_{1:m-1}, Y_m = k, \mathbf{w}_{1:m}) \quad (6.47)$$

$$= \max_{y_{m-1}} p(Y_m = k, w_m \mid y_{m-1}) \times \max_{\mathbf{y}_{1:m-2}} p(\mathbf{y}_{1:m-2}, y_{m-1}, \mathbf{w}_{1:m-1}) \quad (6.48)$$

$$= \max_{y_{m-1}} p(Y_m = k, w_m \mid y_{m-1}) \times v_{m-1}(y_{m-1}) \quad (6.49)$$

$$= p_E(w_m \mid Y_m = k) \times \max_{y_{m-1}} p_T(y_m \mid y_{m-1}) \times v_{m-1}(y_{m-1}) \quad (6.50)$$

$$= \max_{y_{m-1}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, Y_m = k, y_{m-1}, m)) \times v_{m-1}(y_{m-1}). \quad (6.51)$$

In the final line, we use the fact that $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m) = \log p(y_m, w_m \mid y_{m-1})$, and exponentiate the dot product to obtain the probability.

In practice, the probabilities tend towards zero over long sequences, so the log-probability version of Viterbi is more practical from the standpoint of numerical stability. However, this version connects to a broader literature on inference in graphical models. Each Viterbi variable is computed by **maximizing** over a set of **products**. Thus, the Viterbi algorithm is a special case of the **max-product algorithm** for inference in graphical models (Wainwright and Jordan, 2008).

The Forward Algorithm

In an influential survey, Rabiner (1989) defines three problems for hidden Markov models:

Decoding Find the best tags \mathbf{y} for a sequence \mathbf{w} .

Likelihood Compute the marginal probability $p(\mathbf{w}) = \sum_{\mathbf{y}} p(\mathbf{w}, \mathbf{y})$.

Learning Given only unlabeled data $\{\mathbf{w}^{(1)}, \mathbf{w}^{(2)}, \dots, \mathbf{w}^{(N)}\}$, estimate the transition and emission distributions.

The Viterbi algorithm solves the decoding problem. We'll talk about the learning problem in § 6.6. Let's now consider how to compute the marginal likelihood $p(\mathbf{w}) = \sum_{\mathbf{y}} p(\mathbf{w}, \mathbf{y})$, which involves summing over all possible tag sequences. There are at least two reasons we might want to do this:

Language modeling Note that the probability $p(\mathbf{w})$ is also computed by the language models that were discussed in chapter 5. In those language models, we used only unlabeled corpora, conditioning each token w_m on previous tokens. An HMM-based language model would leverage a corpus of part-of-speech annotations, and therefore might be expected to generalize better than an n-gram language model — for example, it would be more likely to assign positive probability to a nonsense grammatical sentence like *colorless green ideas sleep furiously*.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Tag marginals It is often important to compute marginal probabilities of individual tags, $p(y_m \mid \mathbf{w}_{1:M})$. This is the probability distribution over tags for token m , conditioned on all of the words $\mathbf{w}_{1:M}$. For example, we might like to know the probability that a given word is tagged as a verb, regardless of how all the other words are tagged. We will discuss how to compute this probability in § 6.5, but as a preview, we will use the following form,

$$p(y_m \mid \mathbf{w}_{1:M}) = \frac{p(y_m, \mathbf{w}_{1:M})}{p(\mathbf{w}_{1:M})}, \quad (6.52)$$

which involves the marginal likelihood in the denominator.

We can compute the marginal likelihood using a dynamic program that is nearly identical to the Viterbi algorithm. We will use probabilities for now, and show the conversion to log-probabilities later. The core of the algorithm is to compute a set of **forward variables**,

$$\alpha_m(k) \triangleq p(Y_m = k, \mathbf{w}_{1:m}). \quad (6.53)$$

From this definition, we can compute the marginal likelihood by summing over the final forward variables,

$$p(\mathbf{w}) = \sum_{k \in \mathcal{Y}} p(Y_M = k, \mathbf{w}_{1:M}) \quad (6.54)$$

$$= \sum_{k \in \mathcal{Y}} \alpha_M(k). \quad (6.55)$$

To capture the probability of terminating the sequence on each possible tag Y_M , we can pad the end of \mathbf{w} with an extra token \blacksquare , which can only be emitted from the stop tag \blacklozenge .

The forward variables themselves can be computed recursively,

$$\alpha_m(k) = p(Y_m = k, \mathbf{w}_{1:m}) \quad (6.56)$$

$$= p(w_m \mid Y_m = k) \times \Pr(Y_m = k, \mathbf{w}_{1:m-1}) \quad (6.57)$$

$$= p(w_m \mid Y_m = k) \times \sum_{k' \in \mathcal{Y}} \Pr(Y_m = k, Y_{m-1} = k', \mathbf{w}_{1:m-1}) \quad (6.58)$$

$$= p(w_m \mid Y_m = k) \times \sum_{k' \in \mathcal{Y}} \Pr(Y_m = k \mid Y_{m-1} = k') \times \Pr(Y_{m-1} = k', \mathbf{w}_{1:m-1}) \quad (6.59)$$

$$= p(w_m \mid Y_m = k) \times \sum_{k' \in \mathcal{Y}} \Pr(Y_m = k \mid Y_{m-1} = k') \times \alpha_{m-1}(k') \quad (6.60)$$

$$= \sum_{k' \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}_{1:M}, Y_m = k, Y_{m-1} = k', m)) \times \alpha_{m-1}(k'). \quad (6.61)$$

The derivation relies on the independence assumptions in the hidden Markov model: W_m depends only on Y_m , and Y_m is conditionally independent from $W_{1:m-1}$ and all tags, given

(c) Jacob Eisenstein 2014-2017. Work in progress.

Y_{m-1} . We complete the derivation by introducing Y_{m-1} and summing over all possible values, and by then applying the chain rule to obtain the final recursive form.

Procedurally, we compute the forward variables in just the same way as we compute the Viterbi variables: we first compute all $\alpha_1(\cdot)$, then all $\alpha_2(\cdot)$, and so on. We initialize each $\alpha_0(k) = p(Y_m = k \mid Y_{m-1} = \diamond)$, to capture the transition probability from the start symbol. Comparing Equation 6.60 to Equation 6.50, the sole difference is that instead of maximizing over possible values of Y_{m-1} , we sum. Just as the Viterbi algorithm is a special case of the max-product algorithm for inference in graphical models, the forward algorithm is a special case of the **sum-product** algorithm for computing marginal likelihoods.

In practice, it is numerically more stable to compute the marginal log-probability. In the log domain, the forward recurrence is,

$$\alpha_m(k) \triangleq \log p(\mathbf{w}_{1:m}, Y_m = k) \quad (6.62)$$

$$\begin{aligned} &= \log \sum_{k' \in \mathcal{Y}} \exp(\log p(w_m \mid Y_m = k) + \log \Pr(Y_m = k \mid Y_{m-1} = k')) \\ &\quad + \log p(\mathbf{w}_{1:m-1}, Y_{m-1} = k') \end{aligned} \quad (6.63)$$

$$\begin{aligned} &= \log \sum_{k' \in \mathcal{Y}} \exp(\log p(w_m \mid Y_m = k) + \log \Pr(Y_m = k \mid Y_{m-1} = k') + \alpha_{m-1}(k')) \\ &\quad (6.64) \end{aligned}$$

$$= \log \sum_{k' \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}_{1:M}, Y_m = k, Y_{m-1} = k', m) + \alpha_{m-1}(k')). \quad (6.65)$$

Scientific programming libraries provide numerically robust implementations of the log-sum-exp function, which should prevent overflow and underflow from exponentiation.

Semiring Notation and the Generalized Viterbi Algorithm

We have now seen the Viterbi and Forward recurrences, each of which can be performed over probabilities or log probabilities. These four recurrences are closely related, and can in fact be expressed as a single recurrence in a more general notation, known as **semiring algebra**. We use the symbol \oplus to represent generalized addition, and the symbol \otimes to represent generalized multiplication.³ Given these operators, we can denote a generalized Viterbi recurrence as,

$$v_m(k) = \bigoplus_{k' \in \mathcal{Y}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, Y_m = k, Y_{m-1} = k', m) \otimes v_{m-1}(k'). \quad (6.66)$$

³In a semiring, the addition and multiplication operators must both obey associativity, and multiplication must distribute across addition; the addition operator must be commutative; there must be additive and multiplicative identities $\bar{0}$ and $\bar{1}$, such that $a \oplus \bar{0} = a$ and $a \otimes \bar{1} = a$; and there must be a multiplicative annihilator $\bar{0}$, such that $a \otimes \bar{0} = \bar{0}$.

Each recurrence that we have seen so far is a special case of this generalized Viterbi recurrence:

- In the max-product Viterbi recurrence over probabilities, the \oplus operation corresponds to maximization, and the \otimes operation corresponds to multiplication.
- In the forward recurrence over probabilities, the \oplus operation corresponds to addition, and the \otimes operation corresponds to multiplication.
- In the max-product Viterbi recurrence over log-probabilities, the \oplus operation corresponds to maximization, and the \otimes operation corresponds to addition. (This is sometimes called the **tropical semiring**, in honor of the Brazilian mathematician Imre Simon.)
- In the forward recurrence over log-probabilities, the \oplus operation corresponds to log-addition, $a \oplus b = \log(e^a + e^b)$. The \otimes operation corresponds to addition.

The mathematical abstraction offered by semiring notation can be applied to the software implementations of these algorithms, yielding concise and modular implementations. The OPENFST library (Allauzen et al., 2007) is an example of a software package in which the algorithms are parametrized by the choice of semiring.

6.5 Discriminative sequence labeling

Today, hidden Markov models are rarely used for supervised sequence labeling. This is because HMMs are limited to only two phenomena:

- Word-tag probabilities, via the emission probability $p_E(w_m | y_n)$;
- local context, via the transition probability $p_T(y_m | y_{m-1})$.

However, as we have seen, the Viterbi algorithm can be applied to much more general feature sets, as long as the decomposition $f(\mathbf{w}, \mathbf{y}) = \sum_{m=1}^M f(\mathbf{w}, y_m, y_{m-1}, m)$ is observed. In this section, we discuss methods for learning the weights on such features. However, let's first pause to ask what additional features might be needed.

Word affix features. Consider the problem of part-of-speech tagging on the first four lines of the poem *Jabberwocky* (Carroll, 1917):

- (6.3) 'Twas brillig, and the slithy toves
 Did gyre and gimble in the wabe:
 All mimsy were the borogoves,
 And the mome raths outgrabe.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Many of these words are made up, so you would have no information about their probabilities of being associated with any particular part of speech. Yet it is not so hard to see what their grammatical roles might be in this passage. Context helps: for example, the word *slithy* follows the determiner *the*, and therefore is likely to be a noun or adjective. Which do you think is more likely? The suffix *-thy* is found in a number of adjectives — e.g., *frothy, healthy, pithy, worthy*. The suffix is also found in a handful of nouns — e.g., *apathy, sympathy* — but nearly all of these nouns contain *-pathy*, unlike *slithy*. The suffix gives some evidence that *slithy* is an adjective, and indeed it is: later in the text we find that it is a combination of the adjectives *lithe* and *slimy*.⁴

Fine-grained context. Another useful source of information is fine-grained context — that is, contextual information that is more specific than the previous tag. For example, consider the noun phrases *this fish* and *these fish*. Many part-of-speech tagsets distinguish between singular and plural nouns, but do not distinguish between singular and plural determiners; for example, the Penn Treebank tagset follows these conventions. A hidden Markov model would be unable to correctly label *fish* as singular or plural in both of these cases, because it only has access to two features: the preceding tag (determiner in both cases) and the word (*fish* in both cases). The classification-based tagger discussed in § 6.1 had the ability to use preceding and succeeding words as features, and we would like to incorporate this information into a sequence labeling algorithm.

Example Suppose we have the tagging D J N (determiner, adjective, noun) for the sequence *the slithy toves* in Jabberwocky, so that

$$\begin{aligned} w &= \text{the slithy toves} \\ y &= \text{D J N}. \end{aligned}$$

We now create the feature vector for this example, assuming that we have word-tag features (indicated by prefix *W*), tag-tag features (indicated by prefix *T*), and suffix features (indicated by prefix *M*). We assume access to a method for extracting the suffix *-thy* from *slithy*, *-es* from *toves*, and \emptyset from *the*, indicating that this word has no suffix. The resulting feature vector is,

$$\begin{aligned} f(\text{the slithy toves}, \text{D J N}) = \{ &\langle W : \text{the}, \text{D} \rangle, \langle M : \emptyset, \text{D} \rangle, \langle T : \diamond, \text{D} \rangle \\ &\langle W : \text{slithy}, \text{J} \rangle, \langle M : \text{-thy}, \text{J} \rangle, \langle T : \text{D}, \text{J} \rangle \\ &\langle W : \text{toves}, \text{N} \rangle, \langle M : \text{-es}, \text{N} \rangle, \langle T : \text{J}, \text{N} \rangle \\ &\langle T : \text{N}, \blacklozenge \rangle \}. \end{aligned}$$

⁴**Morphology** is the study of how words are formed from smaller linguistic units. Computational approaches to morphological analysis are touched on in chapter 8; Bender (2013) provides a good overview of the underlying linguistic principles.

We now consider several discriminative methods for learning feature weights in sequence labeling. In chapter 2, we considered three types of discriminative classifiers: perceptron, support vector machine, and logistic regression. Each of these classifiers has a structured equivalent, enabling it to be trained from labeled sequences rather than individual tokens.

Structured perceptron

The perceptron classifier updates its weights by increasing the weights for features that are associated with the correct label, and decreasing the weights for features that are associated with incorrectly predicted labels:

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y) \quad (6.67)$$

$$\boldsymbol{\theta}^{(t+1)} \leftarrow \boldsymbol{\theta}^{(t)} + \mathbf{f}(\mathbf{x}, y) - \mathbf{f}(\mathbf{x}, \hat{y}). \quad (6.68)$$

We can apply exactly the same update in the case of structure prediction,

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \mathbf{y}) \quad (6.69)$$

$$\boldsymbol{\theta}^{(t+1)} \leftarrow \boldsymbol{\theta}^{(t)} + \mathbf{f}(\mathbf{w}, \mathbf{y}) - \mathbf{f}(\mathbf{w}, \hat{\mathbf{y}}). \quad (6.70)$$

This learning algorithm is called **structured perceptron**, because it learns to predict the structured output \mathbf{y} . The key difference is that instead of computing $\hat{\mathbf{y}}$ by enumerating the entire set \mathcal{Y} , we use the Viterbi algorithm to search this set efficiently. In this case, the output structure is the sequence of tags (y_1, y_2, \dots, y_M) ; the algorithm can be applied to other structured outputs as long as efficient inference is possible. As in perceptron classification, weight averaging is crucial to get good performance (see § 2.1).

Example For the example *They can fish*, suppose the reference tag sequence is N V V, but our tagger incorrectly returns the tag sequence N V N. Given **feature templates** $\langle w_m, y_m \rangle$ and $\langle y_{m-1}, y_m \rangle$, the corresponding structured perceptron update is:

$$\theta_{\langle \text{fish}, \text{V} \rangle} \leftarrow \theta_{\langle \text{fish}, \text{V} \rangle} + 1 \quad (6.71)$$

$$\theta_{\langle \text{fish}, \text{N} \rangle} \leftarrow \theta_{\langle \text{fish}, \text{N} \rangle} - 1 \quad (6.72)$$

$$\theta_{\langle \text{V}, \text{V} \rangle} \leftarrow \theta_{\langle \text{V}, \text{V} \rangle} + 1 \quad (6.73)$$

$$\theta_{\langle \text{V}, \text{N} \rangle} \leftarrow \theta_{\langle \text{V}, \text{N} \rangle} - 1 \quad (6.74)$$

$$\theta_{\langle \text{V}, \blacklozenge \rangle} \leftarrow \theta_{\langle \text{V}, \blacklozenge \rangle} + 1 \quad (6.75)$$

$$\theta_{\langle \text{N}, \blacklozenge \rangle} \leftarrow \theta_{\langle \text{N}, \blacklozenge \rangle} - 1. \quad (6.76)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

Structured Support Vector Machines

Large-margin classifiers such as the support vector machine improve on the perceptron by learning weights that push the classification boundary away from the training instances. In many cases, large-margin classifiers outperform the perceptron, so we would like to apply similar ideas to sequence labeling. A support vector machine in which the output is a structured object, such as a sequence, is called a **structured support vector machine** (Tsochantaridis et al., 2004).⁵

In classification, we formalized the large-margin constraint as,

$$\forall y \neq y^{(i)}, \theta \cdot f(x, y^{(i)}) - \theta \cdot f(x, y) \geq 1, \quad (6.77)$$

which says that we require a margin of at least 1 between the scores for all labels y that are not equal to the correct label $y^{(i)}$. The weights θ are then learned by constrained optimization (see § 2.2).

We can apply this idea to sequence labeling by formulating an equivalent set of constraints for all possible labelings $\mathcal{Y}(w)$ for an input w . However, there are two problems with this idea. First, in sequence labeling, some predictions are more wrong than others: we may miss only one tag out of fifty, or we may get all fifty wrong. We would like our learning algorithm to be sensitive to this difference. Second, the number of constraints is equal to the number of possible labelings, which is exponentially large in the length of the sequence.

The first problem can be addressed by adjusting the constraint to require larger margins for more serious errors. Let $c(y^{(i)}, \hat{y}) \geq 0$ represent the **cost** of predicting label \hat{y} when the true label is $y^{(i)}$. We can then generalize the margin constraint,

$$\forall y \neq y^{(i)}, \theta \cdot f(w^{(i)}, y^{(i)}) - \theta \cdot f(w^{(i)}, y) \geq c(y^{(i)}, y). \quad (6.78)$$

This cost-augmented margin constraint specializes to the constraint in Equation 6.77 if we choose the delta function $c(y^{(i)}, y) = \delta(y^{(i)} \neq y)$. For sequence labeling, we can instead use a structured cost function, such as the **Hamming cost**,

$$c(y^{(i)}, y) = \sum_{m=1}^M \delta(y_m^{(i)} \neq y_m). \quad (6.79)$$

With this cost function, we require that the true labeling be separated from the alternatives by a margin that is proportional to the number of incorrect tags in each alternative labeling. Other cost functions are possible as well.

The second problem is that the number of constraints is exponential in the length of the sequence. This can be addressed by focusing on the prediction \hat{y} that *maximally* violates

⁵This model is also known as a **max-margin Markov network** (Taskar et al., 2003), emphasizing that the scoring function is constructed from a sum of components, which are Markov independent.

the margin constraint. We find this prediction by solving the following **cost-augmented decoding** problem:

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y} \neq \mathbf{y}^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}) + c(\mathbf{y}^{(i)}, \mathbf{y}) \quad (6.80)$$

$$= \operatorname{argmax}_{\mathbf{y} \neq \mathbf{y}^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}) + c(\mathbf{y}^{(i)}, \mathbf{y}), \quad (6.81)$$

where in the second line we drop the term $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}^{(i)})$, which is constant in \mathbf{y} .

We can now formulate the margin constraint for sequence labeling,

$$\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}) - \max_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} \left(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}) + c(\mathbf{y}^{(i)}, \mathbf{y}) \right) \geq 0. \quad (6.82)$$

If the score for $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}^{(i)})$ is greater than the cost-augmented score for all alternatives, then the constraint will be met. Therefore we can maximize over the entire set $\mathcal{Y}(\mathbf{w})$, meaning that we can apply Viterbi directly.⁶

The name “cost-augmented decoding” is due to the fact that the objective includes the standard decoding problem, $\max_{\hat{\mathbf{y}}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \hat{\mathbf{y}})$, plus an additional term for the cost. Essentially, we want to train against predictions that are strong and wrong: they should score highly according to the model, yet incur a large loss with respect to the ground truth. We can then adjust the weights to reduce the score of these predictions.

For cost-augmented decoding to be tractable, the cost function must decompose into local parts, just as the feature function $\mathbf{f}(\cdot)$ does. The Hamming cost, defined above, obeys this property. To solve this cost-augmented decoding problem using the Hamming cost, we can simply add features $f_m(y_m) = \delta(y_m \neq y_m^{(i)})$, and assign a weight of 1 to these features. Decoding can then be performed using the Viterbi algorithm. Are there cost functions that do not decompose into local parts? Suppose we want to assign a constant loss c to any prediction $\hat{\mathbf{y}}$ in which k or more predicted tags are incorrect, and zero loss otherwise. This loss function is combinatorial over the predictions, and thus we cannot decompose it into parts.

As with large-margin classifiers, it is possible to formulate the learning problem in an unconstrained form, by combining a regularization term on the weights and a Lagrangian for the constraints:

$$\min_{\boldsymbol{\theta}} \frac{1}{2} \|\boldsymbol{\theta}\|_2^2 - C \left(\sum_i \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}) - \max_{\mathbf{y} \in \mathcal{Y}(\mathbf{w}^{(i)})} \left[\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}) + c(\mathbf{y}^{(i)}, \mathbf{y}) \right] \right), \quad (6.83)$$

In this formulation, C is a parameter that controls the tradeoff between the regularization term and the margin constraints. A number of optimization algorithms have been

⁶To maximize over the set $\mathcal{Y}(\mathbf{w}) \setminus \mathbf{y}^{(i)}$ we would need an alternative version of Viterbi that returns the k -best predictions. K -best Viterbi may be useful for other reasons — for example, in interactive applications, it can be helpful to show the user multiple possible taggings. The design of k -best Viterbi is left an exercise.

proposed for structured support vector machines, some of which are discussed in § 2.2. An empirical comparison by Kummerfeld et al. (2015) shows that stochastic subgradient descent — which is relatively easy to implement — is highly competitive, especially on the sequence labeling task of named entity recognition.

Conditional random fields

Structured perceptron is easy to implement, and structured support vector machines give excellent performance. However, sometimes we need to compute probabilities over labelings, $p(\mathbf{y} \mid \mathbf{w})$, and we would like to do this in a discriminative way. The **Conditional Random Field** (CRF; Lafferty et al., 2001) is a conditional probabilistic model for sequence labeling; just as structured perceptron is built on the perceptron classifier, conditional random fields are built on the logistic regression classifier.⁷ The basic probability model is,

$$p(\mathbf{y} \mid \mathbf{w}) = \frac{\exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \mathbf{y}))}{\sum_{\mathbf{y}' \in \mathcal{Y}(\mathbf{w})} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \mathbf{y}'))}. \quad (6.84)$$

This is almost identical to logistic regression, but because the label space is now tag sequences, we require efficient algorithms for both **decoding** (searching for the best tag sequence given a sequence of words \mathbf{w} and a model $\boldsymbol{\theta}$) and for **normalizing** (summing over all tag sequences). These algorithms will be based on the usual locality assumption on the feature function, $\mathbf{f}(\mathbf{w}, \mathbf{y}) = \sum_{m=1}^M \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m)$.

Decoding in CRFs

Decoding — finding the tag sequence $\hat{\mathbf{y}}$ that maximizes $p(\mathbf{y} \mid \mathbf{w})$ — is a direct application of the Viterbi algorithm. The key observation is that the decoding problem does not depend on the denominator of $p(\mathbf{y} \mid \mathbf{w})$,

$$\begin{aligned} \hat{\mathbf{y}} &= \operatorname{argmax}_{\mathbf{y}} p(\mathbf{y} \mid \mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{y}} \log p(\mathbf{y} \mid \mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{y}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{y}, \mathbf{w}) - \log \sum_{\mathbf{y}' \in \mathcal{Y}(\mathbf{w})} e^{\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{y}', \mathbf{w})} \\ &= \operatorname{argmax}_{\mathbf{y}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{y}, \mathbf{w}). \end{aligned}$$

This is identical to the decoding problem for structured perceptron, so the same Viterbi recurrence as defined in Equation 6.26 can be used.

⁷The name “Conditional Random Field” is derived from **Markov random fields**, a general class of models in which the probability of a configuration of variables is proportional to a product of scores across pairs (or more generally, cliques) of variables in a **factor graph**. In sequence labeling, the pairs of variables include all adjacent tags $\langle y_m, y_{m-1} \rangle$. The probability is **conditioned** on the words $\mathbf{w}_{1:M}$, which are always observed, motivating the term “conditional” in the name.

Learning in CRFs

As with logistic regression, we learn the weights θ by minimizing the regularized negative log conditional probability,

$$\ell = \frac{\lambda}{2} \|\theta\|^2 - \sum_{i=1}^N \log p(\mathbf{y}^{(i)} \mid \mathbf{w}^{(i)}; \theta) \quad (6.85)$$

$$= \frac{\lambda}{2} \|\theta\|^2 - \sum_{i=1}^N \theta \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}) + \log \sum_{\mathbf{y}' \in \mathcal{Y}(\mathbf{w}^{(i)})} \exp(\theta \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}')), \quad (6.86)$$

where λ controls the amount of regularization. We will optimize θ by moving along the gradient of this loss. Probabilistic programming environments, such as THEANO (Bergstra et al., 2010) and TORCH (Collobert et al., 2011), can compute this gradient using automatic differentiation. However, it is worth deriving the gradient to understand how this model works, and why learning is computationally tractable.

As in logistic regression, the gradient includes a difference between observed and expected feature counts:

$$\frac{d\ell}{d\theta_j} = \lambda \theta_j + \sum_{i=1}^N E[f_j(\mathbf{w}^{(i)}, \mathbf{y})] - f_j(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}), \quad (6.87)$$

where $f_j(\mathbf{w}^{(i)}, \mathbf{y}^{(i)})$ refers to the count of feature j for token sequence $\mathbf{w}^{(i)}$ and tag sequence $\mathbf{y}^{(i)}$.

The expected feature counts are computed by summing over all possible labelings of the word sequence,

$$E[f_j(\mathbf{w}^{(i)}, \mathbf{y})] = \sum_{\mathbf{y} \in \mathcal{Y}(\mathbf{w}^{(i)})} p(\mathbf{y} \mid \mathbf{w}^{(i)}; \theta) f_j(\mathbf{w}^{(i)}, \mathbf{y}) \quad (6.88)$$

This looks bad: it is a sum over an exponential number of labelings. To solve this problem, we again rely on the assumption that the overall feature vector decomposes into a sum of local feature vectors, which we exploit to compute the expected feature counts as a sum

(c) Jacob Eisenstein 2014-2017. Work in progress.

across the sequence:

$$E[f_j(\mathbf{w}, \mathbf{y})] = \sum_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} p(\mathbf{y} | \mathbf{w}; \boldsymbol{\theta}) f_j(\mathbf{w}, \mathbf{y}) \quad (6.89)$$

$$= \sum_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} p(\mathbf{y} | \mathbf{w}; \boldsymbol{\theta}) \sum_{m=1}^M f_j(\mathbf{w}, y_m, y_{m-1}, m) \quad (6.90)$$

$$= \sum_{m=1}^M \sum_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} p(\mathbf{y} | \mathbf{w}; \boldsymbol{\theta}) f_j(\mathbf{w}, y_m, y_{m-1}, m) \quad (6.91)$$

$$= \sum_{m=1}^M \sum_{k, k' \in \mathcal{Y}} \sum_{\mathbf{y}: y_{m-1}=k', y_m=k} p(\mathbf{y} | \mathbf{w}; \boldsymbol{\theta}) f_j(\mathbf{w}, k, k', m) \quad (6.92)$$

$$= \sum_{m=1}^M \sum_{k, k' \in \mathcal{Y}} f_j(\mathbf{w}, k, k', m) \sum_{\mathbf{y}: y_{m-1}=k', y_m=k} p(\mathbf{y} | \mathbf{w}; \boldsymbol{\theta}) \quad (6.93)$$

$$= \sum_{m=1}^M \sum_{k, k' \in \mathcal{Y}} f_j(\mathbf{w}, k', k, m) \Pr(Y_{m-1} = k', Y_m = k | \mathbf{w}; \boldsymbol{\theta}). \quad (6.94)$$

This derivation works by interchanging the sum over tag sequences with the sum over indices m . At each position in the sequence, the locality restriction on features ensures that we need only the marginal probability of the tag bigram, $\Pr(Y_{m-1} = k', Y_m = k | \mathbf{w}; \boldsymbol{\theta})$. These tag bigram marginals are also used in unsupervised approaches to sequence labeling. In principle, these marginals still require a sum over the exponentially many label sequences in which $Y_{m-1} = k'$ and $Y_m = k$. However, the marginals can be computed efficiently using the **forward-backward algorithm**.

*Forward-backward algorithm

Recall that in the hidden Markov model, it was possible to use the forward algorithm to compute marginal probabilities $p(y_m, \mathbf{w}_{1:m})$. We now derive a more general version of the forward algorithm, in which label sequences are scored in terms of **potentials** $\psi_m(k, k')$:

$$\psi_m(k, k') \triangleq \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, k, k', m)) \quad (6.95)$$

$$p(\mathbf{y} | \mathbf{w}) = \frac{\prod_{m=1}^M \psi_m(y_m, y_{m-1})}{\sum_{\mathbf{y}'} \prod_{m=1}^M \psi_m(y'_m, y'_{m-1})}. \quad (6.96)$$

Equation 6.96 simply expresses the CRF conditional likelihood, under a change of notation.

(c) Jacob Eisenstein 2014-2017. Work in progress.

The tag bigram marginal probabilities can be written as,

$$\Pr(Y_{m-1} = k', Y_m = k \mid \mathbf{w}; \boldsymbol{\theta}) = \frac{\sum_{\mathbf{y}: y_m=k, y_{m-1}=k'} \prod_{n=1}^M \psi_n(y_n, y_{n-1})}{\sum_{\mathbf{y}'} \prod_{n=1}^M \psi_n(y'_n, y'_{n-1})}. \quad (6.97)$$

where the denominator is the marginal $p(\mathbf{w}_{1:M}) = \sum_{\mathbf{y}} p(\mathbf{w}, \mathbf{y}_{1:M})$, sometimes known as the **partition function**.⁸ Let us now consider how to compute each of these terms efficiently.

Computing the numerator In Equation 6.97, we sum over all tag sequences that include the transition $(Y_{m-1} = k') \rightarrow (Y_m = k)$. Because we are only interested in sequences that include this arc, we can decompose this sum into three parts: the sum over **prefixes** $\mathbf{y}_{1:m-1}$, the transition $(Y_{m-1} = k') \rightarrow (Y_m = k)$, and the sum over **suffixes** $\mathbf{y}_{m:M}$,

$$\begin{aligned} \sum_{\mathbf{y}: Y_m=k, Y_{m-1}=k'} \prod_{n=1}^M \psi_n(y_n, y_{n-1}) &= \sum_{\mathbf{y}_{1:m-1}: y_{m-1}=k'} \prod_{n=1}^{m-1} \psi_n(y_n, y_{n-1}) \\ &\quad \times \psi_m(k, k') \\ &\quad \times \sum_{\mathbf{y}_{m:M}: y_m=k} \prod_{n=m+1}^M \psi_n(y_n, y_{n-1}). \end{aligned} \quad (6.98)$$

The result is product of three terms: a score for getting to the position $(Y_{m-1} = k')$, a score for the transition from k' to k , and a score for finishing the sequence from $(Y_m = k)$. Let us define the first term as a **forward variable**,

$$\alpha_m(k) \triangleq \sum_{\mathbf{y}_{1:m}: y_m=k} \prod_{n=1}^m \psi_n(y_n, y_{n-1}) \quad (6.99)$$

$$= \sum_{k' \in \mathcal{Y}} \psi_m(k, k') \sum_{\mathbf{y}_{1:m-1}: y_{m-1}=k'} \prod_{n=1}^{m-1} \psi_n(y_n, y_{n-1}) \quad (6.100)$$

$$= \sum_{k' \in \mathcal{Y}} \psi_m(k, k') \times \alpha_{m-1}(k'). \quad (6.101)$$

Thus, we compute the forward variables while moving from left to right over the trellis. This forward recurrence is a generalization of the forward recurrence defined in § 6.4. If $\psi_m(k, k') = p_E(w_m \mid Y_m = k) \times p_T(k \mid k')$, then we exactly recover the Hidden Markov Model forward variable $\alpha_m(k) = p(\mathbf{w}_{1:m}, Y_m = k)$ as computed in § 6.4.

⁸The terminology of “potentials” and “partition functions” comes from statistical mechanics (Bishop, 2006).

The third term of Equation 6.98 can also be defined recursively, this time moving over the trellis from right to left. The resulting recurrence is called the **backward algorithm**:

$$\beta_{m-1}(k) \triangleq \sum_{\mathbf{y}_{m-1:M}: y_{m-1}=k} \prod_{n=m}^M \psi_n(y_n, y_{n-1}) \quad (6.102)$$

$$= \sum_{k' \in \mathcal{Y}} \psi_m(k', k) \sum_{\mathbf{y}_{m:M}: y_m=k'} \prod_{n=m+1}^M \psi_n(y_n, y_{n-1}) \quad (6.103)$$

$$= \sum_{k' \in \mathcal{Y}} \psi_m(k', k) \times \beta_m(k'). \quad (6.104)$$

In practice, numerical stability requires that we use log-potentials rather than potentials, $\log \psi_m(y_m, y_{m-1}) = \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m)$. Then the sums must be replaced with log-sum-exp:

$$\log \alpha_m(k) = \log \sum_{k' \in \mathcal{Y}} \exp(\log \psi_m(k, k') + \log \alpha_{m-1}(k')) \quad (6.105)$$

$$\log \beta_{m-1}(k) = \log \sum_{k' \in \mathcal{Y}} \exp(\log \psi_m(k', k) + \log \beta_m(k')). \quad (6.106)$$

Both the forward and backward algorithm operate on the trellis, which implies a space complexity $\mathcal{O}(MK)$. Because they require computing a sum over K terms at each node in the trellis, their time complexity is $\mathcal{O}(MK^2)$.

Computing the normalization term The normalization term (partition function), sometimes abbreviated as Z , can be written as,

$$Z \triangleq \sum_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} p(\mathbf{w}, \mathbf{y}) \quad (6.107)$$

$$= \sum_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \mathbf{y})) \quad (6.108)$$

$$= \sum_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} \prod_{m=1}^M \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m)) \quad (6.109)$$

$$= \sum_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} \prod_{m=1}^M \psi_m(y_m, y_{m-1}). \quad (6.110)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

This term can be computed directly from either the forward or backward probabilities:

$$Z = \sum_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} \prod_{m=1}^M \psi_m(y_m, y_{m-1}) \quad (6.111)$$

$$= \alpha_{M+1}(\diamond) \quad (6.112)$$

$$= \beta_0(\diamond). \quad (6.113)$$

CRF learning: wrapup Having computed the forward and backward variables, we can compute the desired marginal probability as,

$$P(Y_{m-1} = k', Y_m = k \mid \mathbf{w}_{1:M}) = \frac{\alpha_{m-1}(k') \psi_m(k, k') \beta_m(k)}{Z}. \quad (6.114)$$

This computation is known as the **forward-backward algorithm**. From the resulting marginals, we can compute the feature expectations $E[f_j(\mathbf{w}, \mathbf{y})]$; from these expectations, we compute a gradient on the weights $\frac{\partial \mathcal{L}}{\partial \theta}$. Stochastic gradient descent or quasi-Newton optimization can then be applied. As the optimization algorithm changes the weights, the potentials change, and therefore so do the marginals. Each iteration of the optimization algorithm therefore requires recomputing the forward and backward variables for each training instance.⁹

Neural sequence labeling

Recently, neural network methods have been applied to sequence labeling. These methods can be employed in tandem with structure prediction algorithms such as Viterbi, although they are effective on their own. A relatively straightforward approach is to train a recurrent neural network or LSTM, as described in chapter 5; however, rather than predicting the next word in sequence, the model can be trained to predict the tag of the current word. A particularly effective approach is to train a **bidirectional recurrent neural network** (Graves and Schmidhuber, 2005), which estimates two hidden vectors, $\mathbf{h}_m = (\vec{\mathbf{h}}_m, \overleftarrow{\mathbf{h}}_m)$, with $\vec{\mathbf{h}}_m$ indicating the hidden state in a standard left-to-right RNN or LSTM, and $\overleftarrow{\mathbf{h}}_m$ indicating the hidden state in a left-to-right model. In this way, information from the entire sentence is brought to bear on the tagging decision for y_m . This approach was employed by Ling et al. (2015b), who find that bi-LSTMs perform considerably better than bi-RNNs on part-of-speech tagging, and that bidirectional variants of both models perform slightly but consistently better than their unidirectional counterparts. Note that LSTM-based tagging is a classification approach, so dynamic programming is not required to find the best tag sequence.

⁹The `CRFsuite` package implements several learning algorithms for CRFs (<http://www.chokkan.org/software/crfsuite/>).

It is also possible in neural sequence labeling to couple the tagging decisions, by introducing additional parameters for tag-to-tag transitions. Lample et al. (2016) dub this the **LSTM-CRF**, due to its combination of aspects of the long short-term memory and conditional random field models. They find that it improves performance on the task of **named entity recognition**, a sequence labeling task that is described in detail in the next chapter. This task has particularly strong dependencies between adjacent tags, so it is not surprising to see an advantage for structure predictions here.

Both Ling et al. (2015b) and Lample et al. (2016) find that it is advantageous to model unseen and rare words through their character-level representations. They train nested bi-LSTMs for each word, and take the concatenation of the introductory and final hidden states, $(\vec{h}_M, \overleftarrow{h}_0)$ as the word embeddings, which is then used as input in the tagging bi-LSTM. Lample et al. (2016) combine these character-level embeddings with **pre-trained** word embeddings, which were estimated from an unlabeled dataset that is many orders of magnitude larger than the labeled data. They do not backpropagate into the word embeddings, and learn only the transition and output parameters of the model.

6.6 *Unsupervised sequence labeling

In unsupervised sequence labeling, we want to induce a Hidden Markov Model from a corpus of unannotated text $w^{(1)}, w^{(2)}, \dots, w^{(N)}$. This is an example of the general problem of **structure induction**, which is the unsupervised version of **structure prediction**. The tags that result from unsupervised sequence labeling might be useful for some downstream task, or they might help us to better understand the language's inherent structure.

Unsupervised learning in hidden Markov models can be performed using the **Baum-Welch algorithm**, which combines forward-backward with expectation-maximization (EM). In the M-step, we compute the HMM parameters from expected counts:

$$\begin{aligned} \Pr(W = i \mid Y = k) &= \phi_{k,i} = \frac{E[\text{count}(W = i, Y = k)]}{E[\text{count}(Y = k)]} \\ \Pr(Y_m = k \mid Y_{m-1} = k') &= \lambda_{k',k} = \frac{E[\text{count}(Y_m = k, Y_{m-1} = k')]}{E[\text{count}(Y_{m-1} = k')]} \end{aligned}$$

The expected counts are computed in the E-step, using the forward and backward variables as defined in Equation 6.101 and Equation 6.104. Because we are working in a hidden Markov model, we define the potentials as,

$$\psi_m(k, k') = p_E(w_m \mid Y_m = k; \phi) \times p_T(Y_m = k \mid Y_{m-1} = k'; \lambda). \quad (6.115)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

The expected counts are then,

$$E[\text{count}(W = i, Y = k)] = \sum_{m=1}^M \Pr(Y_m = k \mid \mathbf{w}_{1:M}) \delta(W_m = i) \quad (6.116)$$

$$= \sum_{m=1}^M \frac{\Pr(Y_m = k, \mathbf{w}_{1:m}) p(\mathbf{w}_{m+1:M} \mid Y_m = k)}{p(\mathbf{w}_{1:M})} \delta(W_m = i) \quad (6.117)$$

$$= \frac{1}{\alpha_{M+1}(\diamond)} \sum_{m=1}^M \alpha_m(k) \beta_m(k) \delta(W_m = i) \quad (6.118)$$

We use the chain rule to separate $\mathbf{w}_{1:m}$ and $\mathbf{w}_{m+1:M}$, and then use the definitions of the forward and backward variables. In the final step, we normalize by $p(\mathbf{w}_{1:M}) = \alpha_{M+1}(\diamond) = \beta_0(\diamond)$.

The expected transition counts can be computed in a similar manner:

$$E[\text{count}(Y_m = k, Y_{m-1} = k')] = \sum_{m=1}^M \Pr(Y_m = k, Y_{m-1} = k' \mid \mathbf{w}_{1:M}) \quad (6.119)$$

$$\propto \sum_{m=1}^M p(Y_{m-1} = k', \mathbf{w}_{1:m-1}) p(w_{m+1:M} \mid Y_m = k) \\ \times p(w_m, Y_m = k \mid Y_{m-1} = k') \quad (6.120)$$

$$= \sum_{m=1}^M p(Y_{m-1} = k', \mathbf{w}_{1:m-1}) p(w_{m+1:M} \mid Y_m = k) \\ \times p(w_m \mid Y_m = k) p(Y_m = k \mid Y_{m-1} = k') \quad (6.121)$$

$$= \sum_{m=1}^M \alpha_{m-1}(k') \beta_m(k) \phi_{k,w_m} \lambda_{k' \rightarrow k}. \quad (6.122)$$

Again, we use the chain rule to separate out $\mathbf{w}_{1:m-1}$ and $\mathbf{w}_{m+1:M}$, and use the definitions of the forward and backward variables. The final computation also includes the parameters ϕ and λ , which govern (respectively) the emission and transition properties between w_m, y_m , and y_{m-1} . Note that the derivation only shows how to compute this to a constant of proportionality; we would divide by $p(\mathbf{w}_{1:M})$ to go from the joint probability $p(Y_{m-1} = k', Y_m = k, \mathbf{w}_{1:M})$ to the desired conditional $\Pr(Y_{m-1} = k', Y_m = k \mid \mathbf{w}_{1:M})$.

Linear dynamical systems

The forward-backward algorithm can be viewed as Bayesian state estimation in a discrete state space. In a continuous state space, $y_m \in \mathbb{R}$, the equivalent algorithm is the **Kalman Smoother**. It also computes marginals $p(y_m \mid \mathbf{x}_{1:M})$, using a similar two-step algorithm

(c) Jacob Eisenstein 2014-2017. Work in progress.

of forward and backward passes. Instead of computing a trellis of values at each step, we would compute a probability density function $q_{y_m}(y_m; \mu_m, \Sigma_m)$, characterized by a mean μ_m and a covariance Σ_m around the latent state. Connections between the Kalman Smoother and the forward-backward algorithm are elucidated by Minka (1999) and Murphy (2012).

Alternative unsupervised learning methods

As noted in § 4.4, expectation-maximization is just one of many techniques for structure induction. One alternative is to use a family of randomized algorithms called **Markov Chain Monte Carlo (MCMC)**. In these algorithms, we compute a marginal distribution over the latent variable \mathbf{y} **empirically**, by drawing random samples. The randomness explains the “Monte Carlo” part of the name; typically, we employ a Markov Chain sampling procedure, meaning that each sample is drawn from a distribution that depends only on the previous sample (and not on the entire sampling history). A simple MCMC algorithm is **Gibbs Sampling**, in which we iteratively sample each y_m conditioned on all the others (Finkel et al., 2005):

$$p(y_m \mid \mathbf{y}_{-m}, \mathbf{w}_{1:M}) \propto p(w_m \mid y_m) p(y_m \mid \mathbf{y}_{-m}). \quad (6.123)$$

Gibbs Sampling has been applied to unsupervised part-of-speech tagging by Goldwater and Griffiths (2007). *Beam sampling* is a more sophisticated sampling algorithm, which randomly draws entire sequences $\mathbf{y}_{1:M}$, rather than individual tags y_m ; this algorithm was applied to unsupervised part-of-speech tagging by Van Gael et al. (2009).

EM is guaranteed to find only a local optimum; MCMC algorithms will converge to the true posterior distribution $p(\mathbf{y}_{1:M} \mid \mathbf{w}_{1:M})$, but this is only guaranteed in the limit of infinite samples. Recent work has explored the use of **spectral learning** for latent variable models, which use matrix and tensor decompositions to provide guaranteed convergence under mild assumptions (Song et al., 2010; Hsu et al., 2012).

Exercises

1. Consider the garden path sentence, *The old man the boat*. Given word-tag and tag-tag features, what inequality in the weights must hold for the correct tag sequence to outscore the garden path tag sequence for this example?
2. Sketch out an algorithm for a variant of Viterbi that returns the top- n label sequences. What is the time and space complexity of this algorithm?
3. Show how to compute the marginal probability $\Pr(Y_{m-2} = k, Y_m = k')$, in terms of the forwards and backward variables, and the potentials ψ_m .
4. more tk

Chapter 7

Applications of sequence labeling

Sequence labeling has applications throughout computational linguistics. This chapter focuses on the classical applications of part-of-speech tagging, shallow parsing, and named entity recognition, and touches briefly on two applications to interactive settings: dialogue act recognition and the detection of code-switching points between languages.

7.1 Part-of-speech tagging

The **syntax** of a language is the collection of principles under which sequences of words are judged to be grammatically acceptable by fluent speakers. One of the most basic syntactic concepts is the **part-of-speech** (POS), which refers to the syntactic role of each word in a sentence. We have already referred to this concept informally in the previous chapter, and you may have some intuitions from your own study of English. For example, in the sentence *Akanksha likes vegetarian sandwiches*, you may already know that *Akanksha* and *sandwiches* are nouns, *likes* is a verb, and *vegetarian* is an adjective.

Parts-of-speech can help to disentangle or explain various linguistic problems. Recall Chomsky's proposed distinction in chapter 5:

(7.1) Colorless green ideas sleep furiously.

(7.2) *Ideas colorless furiously green sleep.

Why is the first grammatical and the second not? One explanation is that the first example contains part-of-speech transitions that are typical in English: adjective to adjective, adjective to noun, noun to verb, and verb to adverb. In contrast, the second sentence contains transitions that are unusual: noun to adjective and adjective to verb. The ambiguity in sentences like *teacher strikes idle children* can also be explained in terms of parts of speech: in the interpretation that was likely intended, *strikes* is a noun and *idle* is a verb; in the alternative explanation, *strikes* is a verb and *idle* is an adjective.

Part-of-speech tagging is often taken as a early step in a natural language processing pipeline. Indeed, parts-of-speech provide features that can be useful for many of the tasks that we will encounter later, such as parsing, coreference resolution, and relation extraction. [todo: say more here]

Part-of-speech inventories

We have discussed a few parts-of-speech already: noun, verb, adjective, adverb. Jurafsky and Martin (2009) describe these as the four major **open word classes**, meaning that these are classes in which new words can be created.¹ These four open classes can be divided into more fine-grained subcategories, such as proper and common nouns; in addition, there are a number of closed classes. This section provides an overview of part-of-speech categories for English, but see Bender (2013) for a deeper linguistic perspective.

When creating a part-of-speech tagged dataset, the tagset inventory must be explicitly defined. The best known POS dataset for English is the Penn Treebank (PTB; Marcus et al., 1993), which includes a set of 45 tags. This chapter describes the linguistics of part-of-speech categories, and lists the corresponding PTB tags. The next section contrasts the PTB tagset with other tagsets for English and for other languages.

- **Nouns** describe entities and concepts.
 - **Proper nouns** name specific people and entities: e.g., *Georgia Tech*, *Janet*, *Buddhism*. In English, proper nouns are usually capitalized. The Penn Treebank (PTB) tags for proper nouns are: NNP (singular), NNPS (plural).
 - **Common nouns** cover all other nouns. In English, they are often preceded by determiners, e.g. *the book*, *a university*, *some people*. Common nouns decompose into two main types:
 - * **Count nouns** have a plural and need an article in the singular, *dogs*, *the dog*;
 - * **Mass nouns** don't have a plural and don't need an article in the singular:

(7.3) *Fire is dangerous.*

(7.4) *Skiing is an expensive hobby.*

In the Penn Treebank, singular and mass nouns are tagged NN, and plural nouns are tagged NNS.

- **Pronouns** refer to specific entities or events that are already known to the reader or listener.
 - * **Personal pronouns** refer to people or entities: *you*, *she*, *I*, *it*, *me*. The PTB tag is PRP.

¹Languages need not have all four classes: for example, it has been argued that Korean does not have adjectives Kim (2002).

- * **Possessive pronouns** are pronouns that indicate possession: *your, her, my, its, one's, our*. The PTB tag is PRP\$.
- * **Wh-pronouns** (WP) are used in question forms, and as relative pronouns:

(7.5) *Where are you going?*

(7.6) *The man **who** wasn't there.*

Possessive wh-pronouns (e.g., *The guy **whose** email got hacked*) are distinguished in the PTB with the tag WP\$.

Pronouns are considered a **closed class**, because unlike other nouns, it is generally not possible to introduce new pronouns.²

- **Verbs** describe activities, processes, and events. For example, *eat, write, sleep* are all verbs.

- Just as nouns can be differentiated by number, verbs are differentiated by properties of the action they describe, and by their form. For English, the Penn Treebank differentiates the following types of verbs: VB (infinitive, e.g. *to shake*), VBD (past, e.g. *shook*), VBG (present participle, e.g. *shaking*), VBN (past participle, e.g. *shaken*), VBZ (present third-person singular, e.g. *shakes*), VBP (present, non-third-person singular, e.g. *shake*).

Note that these verb classes include properties of the event like tense, as well as subject-verb agreement (third-person singular). Many languages have much more complex inflectional systems than English. For example, French verbs have unique inflections for every combination of person and number; when combined with features such as tense, mood, and aspect, there are several dozen possible inflections for each verb.

- **Modals** are a closed subclass of verbs; they give additional information about the event described by the main verb of the sentence, e.g., *someone **should** do something*. In the PTB, their tag is MD. The PTB distinguishes modal verbs as all verbs which do not take a -s suffix in the third person singular present (Santorini, 1990).
- The verb *to be* requires special treatment, as it must appear with a **predicative adjective** or noun, e.g.

(7.7) *She is hungry.*

(7.8) *We are Georgians.*

²Recent efforts to create gender-neutral singular pronouns in English are the exception that proves the rule: while battles over *s/he* and singular *they* have raged for decades, legions of new common nouns have been introduced (e.g., *selfie, emoji*) without much controversy.

The verbs *is* and *are* in these cases are called **copula**. The PTB does not distinguish copula from other verbs, but other tagsets do. More generally, in **light verb** constructions, the meaning is largely shaped by a predicative adjective or noun phrase, e.g. *he got fired*, *we took a walk*.

- **Auxiliary verbs** include *be*, *have*, *will*, which form complex tenses, negations, and questions.

(7.9) *They **had** spoken.*

(7.10) *She **did** not know.*

(7.11) ***Did** she know?*

(7.12) *We **will have** done it.*

Auxiliary verbs are not distinguished with special tags in the PTB; these cases are tagged as verbs, according to their person, number, and tense.

- **Adjectives** describe properties of entities. In English, adjectives can be used in two ways:

- **Attributive**: *an **antique** land*;
- **Predicative**: *the land was **antique**.*

Adjectives (tagged JJ) may be **gradable**, meaning that they have a comparative form (e.g., *bigger*, *smellier*; tagged JJR) and superlative form (*biggest*, *smelliest*; tagged JJS). Adjectives like *antique* are not gradable.

- **Adverbs** describe properties of events. In the following examples, the bolded words are all adverbs.

(7.13) *He spoke **carefully**.*

(7.14) *She lives **downstairs**.*

(7.15) *I study **here**.*

(7.16) *Go **left** at the first traffic light.*

Adverbs are generally tagged RB, but **comparative adverbs** (e.g., *They played **harder***) are tagged RBR. Adverbs can describe a range of details about events:

- The **manner** of the event, e.g., *slowly*, *slower*, *fast*, *hesitantly*.
- The **degree** of the event, e.g., *extremely*, *very*, *highly*.
- Adverbs also include temporal information, such as *yesterday*, *Monday*, *soon*, and spatial information, such as *here*.

Adverbs do not only modify verbs; they may also modify sentences, adjectives, or other adverbs.

(7.17) *Apparently, the very ill man walks extremely slowly.*

In this example, *very* modifies the adjective *ill*, *slowly* modifies the verb *walks*, *extremely* modifies the adverb *slowly*, and *apparently* modifies the entire sentence that follows it.

- **Prepositions** are a closed class of words that can come before noun phrases, forming a prepositional phrase that relates the noun phrase to something else in the sentence.
 - *I eat sushi with soy sauce.* The prepositional phrase **attaches** to the noun *sushi*.
 - *I eat sushi with chopsticks.* The prepositional phrase here attaches to the verb *eat*.

The preposition *To* gets its own tag TO, because it forms the **infinitive** with bare form verbs (VB), e.g. *I want to eat*. All other prepositions are tagged IN in the PTB.

- **Coordinating conjunctions** (PTB tag: CC) join two elements,

(7.18) *vast and trunkless legs*

(7.19) *She plays backgammon or she does homework.*

(7.20) *She eats and drinks quickly.*

(7.21) *Sandeep lives north of Midtown and south of Buckhead.*

(7.22) *Max cooked, and Abigail ate, all the pizza.*

- **Subordinating conjunctions** introduce a subordinate clause, e.g.

(7.23) *She thinks that Chomsky is wrong about language models.*

In the PTB, subordinating conjunctions are grouped with prepositions in the tag IN.

- **Particles** are words that travel combine with verbs to create **phrasal verbs** with meaning that is distinct from the verb alone, e.g.,

(7.24) *Come on.*

(7.25) *He brushed himself off.*

(7.26) *Let's check out that new restaurant.*

Particles are a closed class, and are tagged RP in the PTB.

- **Determiners** (PTB tag: DT) are a closed class of words that precede noun phrases.
 - **Articles:** *the, an, a*. These words describe the **information status** and **unique-ness** of the noun phrase that follows. For example, *the book* refers to a unique entity, which is already known to the listener; *a book* can refer to the existence of an entity (*Tahira loaned him a book.*)

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Demonstratives: *this, these, that*
- Quantifiers: *some, every, few*
- **Wh-determiners**: e.g., *Which bagel should I choose?, Do you know **when** it will be ready?*
- **Pre-determiners** appear in constructions like ***both** my children*, where *my* is the central determiner.

Of these determiners, only wh-determiners and pre-determiners are distinguished with special tags, WDT and PDT. All others are tagged DT.

- Finally, there are a few oddball categories.
 - **Existential there**, e.g. *There is no way out of here*, gets its own tag, EX.
 - The possessive ending *'s* is annotated POS. This means that we assume we have separated this ending into a separate token.
 - Other special tags are reserved for numbers (CD), foreign words (FW), list items (LS), interjections (e.g., *uh, wow*, tagged UH), and a range of non-alphabetic symbols (commas, dollar signs, quotation marks, etc.)

Tagset granularity

The previous section illustrates some of the design decisions that were made in creating the Penn Treebank tagset. In general, there is a tradeoff between capturing important linguistic details, and choosing a tagset that can annotated efficiently. The PTB strikes one balance; other tagsets make other decisions. In English, the Brown corpus favored a more granular tagset, with 87 part-of-speech tags (Francis, 1964)³. It has:

- specific tags for the *be, do, and have* verbs, which the other two tagsets just lump in with other verbs;
- distinct tags for possessive determiners (*my name*) and possessive pronouns (*mine*);
- distinct tags for the third-person singular pronouns (e.g., *it, he*) and other pronouns (e.g., *they, we, I*).

In the other direction, Petrov et al. (2012) propose a “universal” set of twelve part-of-speech tags, which apply across many languages (Petrov et al., 2012). The Universal tagset aggressively groups categories that are distinguished in the other tagsets:

- all nouns are grouped, ignoring number and the proper/common distinction (see below);
- all verbs are grouped, ignoring inflection;

³See <http://www.comp.leeds.ac.uk/ccalas/tagsets/brown.html> for the tagset details.

- preposition and postpositions are grouped as “adpositions”;
- all punctuation is grouped;
- coordinating and subordinating conjunctions (e.g. *and* versus *that*) are grouped.

In the Universal Dependency Treebank (Nivre et al., 2016), the coarse-grained universal part-of-speech tags are augmented with **lexical features**, such as gender, number, case, and tense. These features are annotated only for languages in which they apply: for example, gender would be annotated for determiners in Spanish, but not in English.

How to decide among these tagging strategies? Each has its own advantages. The Brown tags can be useful for certain applications, and they may have strong tag-to-tag relations that make tagging easier, as described in the next chapter). But they are more expensive to annotate. The Universal tags are intended to generalize across many languages and many types of text, and should be easier to annotate.




Figure 7.1: [todo: attribution?]

Example

To understand the linguistic differences between these tagsets, let's look at an example:

- (7.27) My name is Ozymandias, king of kings:
 Look on my works, ye Mighty, and despair!

(c) Jacob Eisenstein 2014-2017. Work in progress.

The part-of-speech tags for this couplet from Ozymandias are shown in Table 7.1.

	Brown		PTB	Universal
My	possessive (DD\$)	determiner	possessive pronoun (PRP\$)	pronoun (PRON)
name	noun, singular, common (NN)		NN	NOUN
is	verb “to be” 3rd person, singular (BEZ)		verb 3rd person, singular (VBZ)	VERB
Ozymandias	proper noun, singular (NP)		proper noun, singular (NNP)	NOUN
,	comma (,)		comma (,)	punctuation (.)
king	NN		NN	NOUN
of	preposition (IN)		preposition (IN)	adposition (ADP)
kings	noun, plural, common (NNS)		NNS	NOUN
:	colon (:)		mid-sentence punc (:)	.
Look	verb, base: uninflected present, imperative, or infinite (VB)		VB	VERB
on	IN		IN	ADP
my	DD\$		PRP\$	PRON
works	NNS		NNS	NOUN
ye	personal pronoun, nominative, non 3S (PPSS)		personal pronoun, nominative (PRP)	PRON
mighty	adjective (JJ)		JJ	adjective (ADJ)
,	comma (,)		comma (,)	punctuation (.)
and	coordinating conjunction (CC)		CC	conjunction (CONJ)
despair	VB		VB	VERB

Table 7.1: Part-of-speech annotations from three tagsets for the first couplet of the poem Ozymandias.

Accurate part-of-speech tagging

Part-of-speech tagging is the problem of selecting the correct tag for each word in a sentence. Success is typically measured by accuracy on an annotated test set, which is simply

(c) Jacob Eisenstein 2014-2017. Work in progress.

the fraction of tokens that were tagged correctly. POS tagging has been a benchmark problem in natural language processing since the mid-1990s.

Baselines

A simple baseline for part-of-speech tagging is to choose the most common tag for each word. For example, in the Universal Dependency treebank, the word *talk* appears 96 times, and 85 of those times it is labeled as a verb: therefore, we will always predict verb when we see it in the test set. For words that do not appear in the training corpus, the baseline simply guesses the most common tag overall, which is noun. In the Penn Treebank, this simple baseline obtains accuracy above 92%. So the interesting question is how many of the remaining 8% of instances can be classified correctly.

Hidden Markov models 6.4 were a popular choice for part-of-speech tagging: the emission distribution links tags to words, and the transition distribution captures tag ordering constraints. However, hidden Markov models can struggle with rare words, which may not appear in the training data. For these words, it is important to use word-internal information, such as suffixes and capitalization; despite some heroic efforts (Brants, 2000), these features are difficult to incorporate into hidden Markov models.

Structure prediction

Conditional random fields and structured perceptron perform at or near the state-of-the-art for part-of-speech tagging in English. For example, (Collins, 2002) achieved 97.1% accuracy on the Penn Treebank, using a structured perceptron, using the following base features (originally introduced by Ratnaparkhi (1996)):

- current word, w_m
- previous words, w_{m-1}, w_{m-2}
- next words, w_{m+1}, w_{m+2}
- previous tag, y_{m-1}
- previous two tags, (y_{m-1}, y_{m-2})
- for rare words:
 - first k characters, up to $k = 4$
 - last k characters, up to $k = 4$
 - whether w_m contains a number, uppercase character, or hyphen.

A group from Stanford used a similar set of features in a model that is closely related to conditional random fields (Toutanova et al., 2003), attaining 97.3% accuracy. Eight years later, Chris Manning (one of the authors on the Stanford paper), remarked on how difficult it had been to obtain further improvements from either better features or machine learning models (Manning, 2011). Manning notes that despite the seemingly impressive 97% accuracy on the token level, only a little more than half of all sentences are tagged

(c) Jacob Eisenstein 2014-2017. Work in progress.

The	U.S.	Army	captured	Atlanta	on	May	14	,	1864	.
B-ORG	I-ORG	I-ORG	O	B-LOC	O	B-DATE	I-DATE	I-DATE	I-DATE	O

Table 7.2: BIO notation for named entity recognition

completely correctly, suggesting that the task is far from solved. He suggests that better annotations may be necessary to drive further improvements.

Neural part-of-speech tagging

[todo: remind readers about LSTMs from § 5.3. Or maybe put the modeling stuff in the previous chapter?]

- Simple LSTM tagging, using pre-trained word embeddings Huang et al. (2015); Ma and Hovy (2016)
- Bi-directional LSTM (Graves et al., 2013)
- LSTM-CRF (Huang et al., 2015; Ma and Hovy, 2016). Improvements on PTB are small (error rate from 2.7% to 2.4%). Bigger improvements for NER

7.2 Shallow parsing

7.3 Named entity recognition

A standard approach to tagging named entity spans is to use discriminative sequence labeling methods such as conditional random fields and structured perceptron. As described in chapter 6, these methods use the Viterbi algorithm to search over all possible label sequences, while scoring each sequence using a feature function that decomposes across adjacent tags. Named entity recognition is formulated as a tagging problem by assigning each word token to a tag from a tagset. However, there is a major difference from part-of-speech tagging: in NER we need to recover **spans** of tokens, such as *The United States Army*. To do this, the tagset must distinguish tokens that are at the **beginning** of a span from tokens that are inside a span.

BIO notation This is accomplished by the **BIO notation**, shown in Table 17.1. Each token at the beginning of a name span is labeled with a B- prefix; each token within a name span is labeled with an I- prefix. Tokens that are not parts of name spans are labeled as O. From this representation, it is unambiguous to recover the entity name spans within a labeled text. Another advantage is from the perspective of learning: tokens at the beginning of name spans may have different properties than tokens within the name, and the learner can exploit this. This insight can be taken even further, with special labels for

(c) Jacob Eisenstein 2014-2017. Work in progress.

the last tokens of a name span, and for **unique** tokens in name spans, such as *Atlanta* in the example in Table 17.1. This is called **BILOU** notation, and has been shown to yield improvements in supervised named entity recognition (Ratinov and Roth, 2009).

Entity types The number of possible entity types depends on the labeled data. An early dataset was released as part of a shared task in the Conference on Natural Language Learning (CoNLL), containing entity types **LOC** (location), **ORG** (organization), and **PER** (person). Later work has distinguished additional entity types, such as dates, **[todo: etc]**. **[todo: find cites]** Special purpose corpora have been built for domains such as biomedical text, where entities include protein types **[todo: etc]**.

Features The use of Viterbi decoding restricts the feature function $f(w, y)$ to $\sum_m f(w, y_m, y_{m-1}, m)$, so that each feature can consider only local adjacent tags. Typical features include tag transitions, word features for w_m and its neighbors, character-level features for prefixes and suffixes, and “word shape” features to capture capitalization. As an example, base features for the word *Army* in the example in Table 17.1 include:

<CURR-WORD:*Army*,
 PREV-WORD:*U.S.*,
 NEXT-WORD:*captured*,
 PREFIX-1:*A*-,
 PREFIX-2:*Ar*-,
 SUFFIX-1:*-y*,
 SUFFIX-2:*-my*,
 SHAPE:*Xxxx*>

Another source of features is to use **gazetteers**: lists of known entity names. For example, it is possible to obtain from the U.S. Social Security Administration a list of **[todo: hundreds of thousands]** of frequently used American names — more than could be observed in any reasonable annotated corpus. Tokens or spans that match an entry in a gazetteer can receive special features; this provides a way to incorporate hand-crafted resources such as name lists in a learning-driven framework.

Features in recent state-of-the-art systems are summarized in papers by ? and Ratinov and Roth (2009).

7.4 Dialogue acts

[todo: define problem] (Stolcke et al., 2000; Galley, 2006)

(c) Jacob Eisenstein 2014-2017. Work in progress.

7.5 Code switching

[**todo: define problem**] Solorio and Liu (2008) use a classification-based approach. Nguyen and Dogruöz²³ (2013) use a conditional random field.

Chapter 8

Finite-state automata

Consider the following problems:

- Segment a word into its stem and affixes: *impossibility* \rightarrow *im+possibl+ity*.
- Convert a sequence of morphemes like *im+possible+ity* into the correct sequence of characters (*impossibility*).
- Decide whether a given word is morphotactically correct, or more generally, rank all the possible realizations for a morphological expression like NEGATION + *possible*: *impossible*, *inpossible*, *nonpossible*, *unpossible*, etc.
- Given a speech utterance and a large set of potential text transcriptions, choose the one with the highest probability according to an n-gram language model.
- Perform context-sensitive spelling correction, so as to correct examples like *their at piece* to *they're at peace*.

All of these problems relate to the content of the previous two chapters — language models and morphology — but none of them seem easily solved by supervised classifiers. This chapter presents a new tool for language technology: finite state automata. Finite-state automata are particularly suited for scoring strings (sequences of characters, words, morphemes, or phonemes), and for converting one string into another. A key advantage of finite state automata is their modularity: the output of one finite-state transducer can be the input for another, allowing the combination of simple components into cascades with rich and complex behaviors.

Finite-state automata are a formalism for representing a subset of formal languages, the **regular** languages; these are languages that can be defined with regular expressions. While there is strong evidence that natural language is not regular — that is, the question of whether a given sentence is grammatical cannot be answered with any regular expres-

sion — finite state automata can be used as the building block for a surprisingly wide range of applications in language technology.¹

8.1 Automata and languages

Finite state automata emerge from formal language theory. Here are some basic formalisms that will be used throughout this chapter:

- An **alphabet** Σ is a set of symbols, e.g. $\{a, b, c, \dots, z\}$, or $\{aardvark, abacus, \dots, zyxt\}$.
- A **string** ω is a sequence of symbols, $\omega \in \Sigma^*$. The empty string ϵ contains zero symbols.
- A **language** $L \subseteq \Sigma^*$ is a set of strings.
- An **automaton** is an abstract model of a computer, which reads a string $\omega \in \Sigma^*$, and determines whether or not $\omega \in L$.

This seems a very different notion of “language” than English or Hindi. But could we think of these natural languages in the same way as formal languages? If *impossible* is acceptable as an English word but *unpossible* is not, might it be possible to build an automaton that formalizes the underlying linguistic distinction?

Finite-state automata

A finite-state **acceptor** (FSA) is a special type of automaton, which is capable of modeling some, but not all languages. Formally, finite-state automata are defined by a tuple $M = \langle Q, \Sigma, q_0, F, \delta \rangle$, consisting of:

- a finite alphabet Σ of input symbols;
- a finite set of **states** $Q = \{q_0, q_1, \dots, q_n\}$;
- a **start state** $q_0 \in Q$;
- a set of **final states** $F \subseteq Q$;
- a **transition function** $\delta : Q \times \Sigma \rightarrow 2^Q$. The transition function maps from a state and an input symbol to a **set** of possible resulting states.

Given this definition, M accepts a string ω if there is a path from q_0 to any state $q_i \in F$ that consumes all of the symbols in ω . If M accepts ω , this means that ω is in the formal language L defined by M .

¹A more formal treatment of finite state automata and their applications to language is offered by Mohri et al. (2002). Knight and May (2009) show how finite-state automata can be composed together to create impressive applications, focusing on **transliteration** of words and names between languages with different scripts. Here, we’ll build the formalism from the ground up, starting with finite-state acceptors, then adding weights, and then adding transduction, finally arriving at the same sorts of applications.

Example Consider the following FSA, M_1 .

$$\Sigma = \{a, b\} \quad (8.1)$$


$$Q = \{q_0, q_1\} \quad (8.2)$$

$$F = \{q_1\} \quad (8.3)$$

$$\delta = \{ \{ (q_0, a) \rightarrow q_0 \}, \quad (8.4)$$

$$\{ (q_0, b) \rightarrow q_1 \},$$

$$\{ (q_1, b) \rightarrow q_1 \} \}$$



This FSA defines a language over an alphabet of two symbols, a and b . The transition function δ is written as a set of tuples: the tuple $\{(q_0, a) \rightarrow q_0\}$ says that if you are in state q_0 and you see symbol a , you can consume it and stay in q_0 . Because each pair of initial state and symbol has at most one resulting state, this FSA is **deterministic**: each string ω induces at most one path. Note that δ does not contain any information about what to do if you encounter the symbol a while in state q_1 . In this case, you are stuck, and cannot accept the input string.

What strings does this FSA accept? We begin in q_0 , but we have to get to q_1 , since this is the only final state. We can accept any number of a symbols while in q_0 , but we require a b symbol to transition to q_1 . Once there, we can accept any number of b symbols, but if we see an a symbol, there is nothing we can do. So the regular expression corresponding to the language defined by M_1 is a^*bb^* . To see this, consider what M_1 would do if it were fed each of the following strings: $aaabb$; aa ; $abbba$; bb .

Regular languages* Can every formal language be recognized by some finite state automata? No. Finite state automata can only recognize **regular languages**. The classic example of a non-regular language is $a^n b^n$; this language includes only those strings that contain n copies of symbol a , followed by n copies of symbol b . The **pumping lemma** demonstrates that this language cannot be accepted by any FSA. The proof is by contradiction. Suppose M is an FSA that accepts the language $a^n b^n$. By definition M must have a finite number of states; if we choose a string $a^m b^m$ such that m is bigger than the number of states in M , then the path through M must contain a cycle, and the transitions on this cycle must accept only the symbol a . But if there is a cycle, then we can repeat the cycle any number of times, “pumping up” the number of a symbols in the string. The automaton M must therefore also accept strings $a^{m'} b^m$, with $m' > m$. But these strings are not in

the language $a^n b^n$, so we arrive at a contradiction. The proof will be covered in detail by any textbook on theory of computation (e.g., Sipser, 2012).

Determinism

- In a deterministic (D)FSA, the transition function is defined so that $\delta : Q \times \Sigma \rightarrow Q$. This means that every pair of initial state and symbol can transition to at most one resulting state.
- In a nondeterministic (N)FSA, $\delta : Q \times \Sigma \rightarrow 2^Q$. This means that a pair of initial state and symbol can transition to multiple resulting states. As a consequence, an NFSA may have multiple paths to accept a given string.
- We can determinize any NFSA using the powerset construction, but the number of states in the resulting DFSA may be exponential in the size of the original NFSA.
- Any **regular expression** can be converted into an NFSA, and thus into a DFSA.

The English Dictionary as an FSA We can build a simple “chain” FSA which accepts any single word. So, we can define the English dictionary with an FSA. However, we can make this FSA much more compact. (see slides)

- Begin by taking the **union** of all of the chain FSAs by defining **epsilon transitions** (transitions that do not consume an input symbol) from the start state to chain FSAs for each word (5303 states / 5302 arcs using a 850 word dictionary of “basic English”).
- Eliminate the epsilon transitions by pushing the first letter to the front (4454 states / 4453 arcs)
- **Determinize** (2609 / 2608)
- **Minimize** (744 / 1535). The cost of minimizing an acyclic FSA is $O(E)$. This data structure is called a **trie**.

Operations In discussing talked about three operations: union, determinization and minimization. Other important operations are:

Intersection only accept strings in both FSAs: $\omega \in (M_1 \cap M_2)$ iff $\omega \in M_1 \cap \omega \in M_2$.

Negation only accept strings not accepted by FSA M : $\omega \in (\neg M)$ iff $\omega \notin M$.

concatenation accept strings of the form $\omega = [\omega_1 \omega_2]$, where $\omega_1 \in M_1$ and $\omega_2 \in M_2$.

FSAs are **closed** under all these operations, meaning that resulting automaton is still an FSA (and therefore still defines a regular language).

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 8.1: First try at modeling English morphology

FSAs for Morphology

Now for some morphology. Suppose that we want to write a program that accepts only those words that are constructed in accordance with English derivational morphology:

- *grace, graceful, gracefully*
- *disgrace, disgraceful, disgracefully, ...*
- *Google, Googler, Googleology, ...*
- **gracelyful, *disungracefully, ...*

As we saw in the English dictionary example, we could just make a list, and then take the union of the list using ϵ -transitions. The list would get very long, and it would not account for productivity (our ability to make new words like *antiwordificationist*). So let's try to use finite state machines instead. Our FSA will have to encode rules about morpheme ordering, called *morphotactics*.

Every word must have a stem, so we do not want to accept proposed words like *dis-* or *-ly*. This suggests that we should have at least two states: one for before we have seen a stem, and one for after. Assuming the alphabet Σ consists of all English morphemes, we can define a transition function so that it is only possible to transition from q_0 to q_1 by consuming a stem morpheme; by defining $F = \{q_1\}$, we can ensure that every word has a stem. For prefixes, we can allow self-transitions in q_0 on prefix morphemes; we can do the same in q_1 for suffix morphemes.

The resulting FSA is shown in Figure 8.1 will accept *grace, disgrace, graceful disgraceful*, and even *disgracefully* (with two self-transitions in q_1). However, it will also accept **gracelyful* and **gracerly*. To deal with these cases, we need to think about what the suffixes are doing. The suffix *-ful* converts the noun *grace* into an adjective *graceful*; it does the same for words like *thoughtful* and *sinful*. The suffix *-ly* converts the adjective *graceful* to the adverb *gracefully* (to see the difference, compare *the ballet was graceful* to *the ballerina moved gracefully*.) These examples suggest that we need additional states in our FSA, such as




Figure 8.2: Second try at modeling English morphology, this time distinguishing parts-of-speech

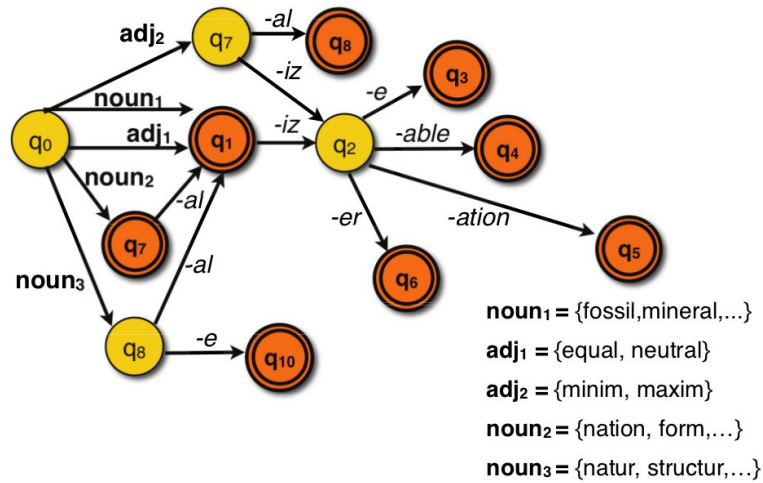


Figure 8.3: A fragment of a finite-state acceptor for derivational morphology. From Julia Hockenmaier's slides.

q_{noun} , $q_{\text{adjective}}$, and q_{adverb} . Each of these is a potential final state, and the suffixes allow transitions between them. This FSA is shown in Figure 8.2.

However, with a little more thought, we see that this approach is still too simple. First, not every noun can be made into an adjective: **chairful* and **monkeyful* are perhaps suggestive of some kind of poetic meaning, but would not be recognized as standard English. Second, many nouns are made into adjectives using different suffixes, such as *music+al*, *fish+y*, and *elv+ish*. We need to create additional noun states to distinguish these noun groups, so as to avoid accepting ill-formed words like **musicky* and **fishful*. We could continue to refine the FSA, coming ever-closer to an accurate model of English morphotactics. A fragment of such an FSA is shown in Figure 8.3.

This approach makes a key assumption: every word is either in or out of the language, with no wiggle room. Perhaps you agreed that *musicky* and *fishful* were not valid English words; but if forced to choose, you probably find *a fishful stew* or *a musicky tribute* prefer-

able to *behaving disgracefully*. To take the argument further, here are some Google counts for various derivational forms:

- *superfast*: 70M; *ultrafast*: 16M; *hyperfast*: 350K; *megafast*: 87K
- *suckitude*: 426K; *suckiness*: 378K
- *nonobvious*: 1.1M; *unobvious*: 826K; *disobvious*: 5K

Given this diversity of possible realizations of the same idea, rather than asking whether a word is **acceptable**, we might like to ask how acceptable it is. But finite state acceptors gives us no way to express *preferences* among technically valid choices. We will need to augment the formalism for this.

8.2 Weighted Finite State Automata

A weighted finite-state automaton $M = \langle Q, \Sigma, \pi, \xi, \delta \rangle$ consists of:

- A finite set of states $Q = \{q_0, q_1, \dots, q_n\}$
- A finite alphabet Σ of input symbols
- Initial weight function, $\pi : Q \rightarrow \mathbb{R}$
- Final weight function $\xi : Q \rightarrow \mathbb{R}$
- A transition function $\delta : Q \times \Sigma \times Q \rightarrow \mathbb{R}$

We have departed from the FSA formalism in three ways:

- Every state can be a start state, with score π_q .
- Every state can be an end state, with score ξ_q .
- Transitions are possible between any pair of states on any input, with a score $\delta_{q_i, \omega, q_j}$.

Now, we can score every path through a weighted finite state acceptor (WFSA) by the sum of the weights for the transitions, plus the scores for the initial and final states. The **shortest path algorithm** finds the minimum-cost path through a WFSA for a string ω , with time complexity $\mathcal{O}(E + V \log V)$, where E is the number of edges and V is the number of vertices (Cormen et al., 2009).

Weighted finite state automata (WFSAs) are a generalization of unweighted FSAs: for any FSA M we can build an equivalent WFSA by setting $\pi_q = \infty$ for all $q \neq q_0$, $\xi_q = \infty$ for all $q \notin F$, and $\delta_{q_i, \omega, q_j}$ for all transitions $\{(q_1, \omega) \rightarrow q_2\}$ that are not permitted by the transition function of M .




Figure 8.4: A weighted finite state acceptor for computing edit distance from the word *edit*.

Applications of WFSAs

We can use WFSAs to score derivational morphology as suggested above. But let's start with some simpler examples.

Edit distance

An **edit distance** is a function of two strings, which quantifies their similarity: for example, *she* and *he* differ by only the addition of a single letter, while *you* and *me* differ on every letter. There are a huge number of ways to compute edit distance (Manning et al., 2008), with applications in information retrieval, bioinformatics, and beyond.

Here we consider a simple edit distance, which computes the minimum number of character insertions, deletions, and substitutions required to get from one word to another. Insertions and deletions are penalized by a cost of one; substitutions have a cost of two. To compute this cost, we build a WFSa with one state for every letter in the word, plus an initial state q_0 : for example, for the word *edit*, we build a machine with states q_0, q_e, q_d, q_i, q_t .

- The initial cost for q_0 is zero; for every other state, the initial cost is infinite.
- The final cost for q_t is zero; for every other state, the final cost is infinite.
- We define the transition function as follows:
 - The cost for “correct” symbols and rightward moves is zero: for example, $\delta_{q_0,e,q_e} = 0$, and $\delta_{q_i,t,q_t} = 0$.
 - The cost for self-transitions is one, regardless of the symbol: for example, $\delta_{q_d,*,q_d} = 1$. These self-transitions correspond to **insertions**.
 - The cost for epsilon transitions to the right is one: for example, $\delta_{q_e,\epsilon,q_d} = 1$. These transitions correspond to **deletions**.
 - The cost of all other transitions is ∞ .

(c) Jacob Eisenstein 2014-2017. Work in progress.

The machine is shown in Figure 8.4. The total edit distance for a string is the *sum* of costs across the best path through machine. Note that we did not define a cost for **substitutions** (e.g., from *him* to *ham*), because substitutions can be performed by a combination of insertion and deletion, for a total cost of two. However, some edit distances assign a cost of one to substitutions; can you see how to modify the WFSA to compute such an edit distance?

N-gram language models

Weighted finite state acceptors can also be used to compute probabilities of sequences — for example, the probability of a word sequence from an n-gram language model. To do this, we define the states and transitions so that each transition is equal to a condition probability, $\delta_{q_i, \omega_m, q_j} = p(q_i, \omega_m \mid q_j)$, so that the product is equal to the joint probability of the state sequence and the string,

$$p(q_{1:M}, \omega_{1:M}) = \prod_m^M p(q_m, \omega_m \mid q_{m-1}). \quad (8.5)$$

For example, to construct a unigram language model over a vocabulary \mathcal{V} of size V , we need just a single state. All transitions are self-transitions, with probability equal to the unigram word probability, $\delta_{q_0, w, q_0} = p_1(w)$.

To construct a bigram language model, we need to model the conditional probability $p(w_m \mid w_{m-1})$. To do this in a WFSA, we must create V different states: one for each context. Then we define the transition function as,

$$\delta_{q_i, w_m, q_j} = \begin{cases} p(w_m \mid w_{m-1} = i), & j = m \\ 0, & \text{otherwise.} \end{cases} \quad (8.6)$$

Because each state represents a context, we require the transition function to ensure that we are in the right state after observing w_m : thus, we assign zero probability to all other transitions. The start function captures the probability $p(w \mid \diamond)$, and the final state function captures the probability $p(\blacklozenge \mid w)$. Thus, the bigram probability of any string is computed by the product of transition scores,

$$p_2(w_{1:M}) = p(w_1 \mid \diamond) \times \left(\prod_{m=2}^M p(w_m \mid w_{m-1}) \right) \times p(\blacklozenge \mid w_M) \quad (8.7)$$

$$= \pi_{w_1} \times \left(\prod_{m=2}^M \delta_{q_{w_{m-1}}, w_m, q_{w_m}} \right) \times \xi_{w_M}. \quad (8.8)$$

Can you see how to construct a trigram language model in the same way?

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 8.5: WFSA implementing an interpolated bigram/unigram language model (Knight and May, 2009). [todo: maybe redraw this for clarity?]

Interpolated n-gram language model

Knight and May (2009) show how to implement an interpolated bigram/unigram language model using a WFSA. Recall that an interpolated bigram language model computes probability,

$$\hat{p}(w_m | w_{m-1}) = \lambda p_1(w_m) + (1 - \lambda) p_2(w_m | w_{m-1}), \quad (8.9)$$

with \hat{p} indicating the interpolated probability, p_2 indicating the bigram probability, and p_1 indicating the unigram probability.

Note that Equation 8.9 involves both the multiplication and addition of probabilities. Knight and May (2009) achieve this through the use of **non-determinism**. The basic idea is shown in Figure 8.5. At each of the top row of states in Figure 8.5, there are two possible ϵ -transitions, which consume no input. With score λ , we transition to the generic state U , which “forgets” the local context; transitions out of U are scored according to the unigram probability model p_1 . With score $1 - \lambda$, we transition to one of the context-remembering states, S' , T' , H' , E' . Each of these states encodes the bigram context, and outgoing transitions are scored according to the bigram probability model p_2 .

Any given path through this WFSA will have a score that multiplies together the probabilities of generating the words in the input, as well as the decisions about whether to use the unigram or bigram probability models. However, due to the non-determinism, each input string will have many possible paths to acceptance. Let’s write these paths as sequences z_1, z_2, \dots, z_M , with each $z_m \in \{1, 2\}$, indicating whether the unigram or bigram model was chosen to generating w_m . Then the string b, a will have the following paths and

(c) Jacob Eisenstein 2014-2017. Work in progress.

scores:

$$\text{score}(1, 1, 1) = \lambda \times p_1(b) \times \lambda \times p_1(a) \times \lambda \times p_1(\diamond) \quad (8.10)$$

$$= \lambda^3 p_1(a) p_1(b) p_1(\diamond) \quad (8.11)$$

$$\text{score}(1, 1, 2) = \lambda^2 (1 - \lambda) p_1(b) p_1(a) p_2(\diamond | a) \quad (8.12)$$

$$\text{score}(1, 2, 1) = \lambda^2 (1 - \lambda) p_1(b) p_2(a | b) p_1(\diamond) \quad (8.13)$$

$$\text{score}(1, 2, 2) = \lambda (1 - \lambda)^2 p_1(b) p_2(a | b) p_2(\diamond | a) \quad (8.14)$$

$$\text{score}(2, 1, 1) = \lambda^2 (1 - \lambda) p_2(b | \diamond) p_1(a) p_1(\diamond) \quad (8.15)$$

$$\text{score}(2, 1, 2) = \lambda^2 (1 - \lambda) p_2(b | \diamond) p_1(a) p_2(\diamond | a) \quad (8.16)$$

$$\text{score}(2, 2, 1) = \lambda^2 (1 - \lambda) p_2(b | \diamond) p_2(a | b) p_1(\diamond) \quad (8.17)$$

$$\text{score}(2, 2, 2) = (1 - \lambda)^3 p_2(b | \diamond) p_2(a | b) p_2(\diamond | a), \quad (8.18)$$

where \diamond is the special start symbol and \blacklozenge is the special stop symbol. Each of these scores is a joint probability $p(\mathbf{w}_{1:M}, \mathbf{z}_{1:M})$; summing over them gives $\sum_{\mathbf{z}_{1:M}} p(\mathbf{w}_{1:M}, \mathbf{z}_{1:M}) = p(\mathbf{w}_{1:M})$, which is the desired marginal probability under the interpolated language model. Thus, in this case, we want not the score of the single best path, but the sum of the scores of **all** paths that accept a given input string.

8.3 Semirings

We have now seen three examples: an FSA for derivational morphology, and WFSA for edit distance and language modeling. Several things are different across these examples.

Scoring

- In the derivational morphology FSA, we wanted a boolean “score”: is the input a valid word or not?
- In the edit distance WFSA, we wanted a numerical (integer) score, with lower being better.
- In the interpolated language model, we wanted a numerical (real) score, with higher being better.

Nondeterminism

- In the derivational morphology FSA, we accept if there is any path to a terminating state.
- In the edit distance WFSA, we want the score of the single best path.
- In the interpolated language model, we want to sum over non-deterministic choices.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Semiring notation allows us to combine all of these different possibilities into a single formalism.

Formal definition

A semiring is a system $(\mathbb{K}, \oplus, \otimes, \bar{0}, \bar{1})$

- \mathbb{K} is the set of possible values, e.g. $\{\mathbb{R}_+ \cup \infty\}$, the non-negative reals union with infinity
- \oplus is an addition operator
- \otimes is a multiplication operator
- $\bar{0}$ is the additive identity
- $\bar{1}$ is the multiplicative identity

A semiring must meet the following requirements:

- $(a \oplus b) \oplus c = a \oplus (b \oplus c)$, $(\bar{0} \oplus a) = a$, $a \oplus b = b \oplus a$
- $(a \otimes b) \otimes c = a \otimes (b \otimes c)$, $a \otimes \bar{1} = \bar{1} \otimes a = a$
- $a \otimes (b \oplus c) = (a \otimes b) \oplus (a \otimes c)$, $(a \oplus b) \otimes c = (a \otimes c) \oplus (b \otimes c)$
- $a \otimes \bar{0} = \bar{0} \otimes a = \bar{0}$

Semirings of interest :

Name	\mathbb{K}	\oplus	\otimes	$\bar{0}$	$\bar{1}$	Applications
Boolean	$\{0, 1\}$	\vee	\wedge	0	1	identical to an unweighted FSA
Probability	\mathbb{R}_+	+	\times	0	1	sum of probabilities of all paths
Log-probability	$\mathbb{R} \cup -\infty \cup \infty$	\oplus_{\log}	+	$-\infty$	0	log marginal probability
Tropical	$\mathbb{R} \cup -\infty \cup \infty$	\min	+	∞	0	best single path

where $\oplus_{\log}(a, b)$ is defined as $\log(e^a + e^b)$.

Semirings allow us to compute a more general notion of the “shortest path” for a WFSA.

- Our initial score is $\bar{1}$
- When we take a step, we use \otimes to combine the score for the step with the running total.
- When nondeterminism lets us take multiple possible steps, we combine their scores using \oplus .

(c) Jacob Eisenstein 2014-2017. Work in progress.

Example Let's see how this works out for our language model example.

$$\begin{aligned} \text{score}(\{a, b, a\}) &= \bar{1} \otimes (\lambda \otimes p_2(a|*) \oplus (1 - \lambda) \otimes p_1(a)) \\ &\quad \otimes (\lambda \otimes p_2(b|a) \oplus (1 - \lambda) \otimes p_1(b)) \\ &\quad \otimes (\lambda \otimes p_2(a|b) \oplus (1 - \lambda) \otimes p_1(a)) \end{aligned}$$

Now if we plug in the **probability semiring**, we get

$$\begin{aligned} \text{score}(\{a, b, a\}) &= 1 \times (\lambda p_2(a|*) + (1 - \lambda)p_1(a)) \\ &\quad \times (\lambda p_2(b|a) + (1 - \lambda)p_1(b)) \\ &\quad \times (\lambda p_2(a|b) + (1 - \lambda)p_1(a)) \end{aligned}$$

But if we plug in the **log probability semiring**, we need the edge weights to be equal to $\log p_1$, $\log p_2$, $\log \lambda$, and $\log(1 - \lambda)$. Then we get:

$$\begin{aligned} \text{score}(\{a, b, a\}) &= 0 + \log(\exp(\log \lambda + \log p_2(a|*)) + \exp(\log(1 - \lambda) + \log p_1(a))) \\ &\quad + \log(\exp(\log \lambda + \log p_2(b|a)) + \exp(\log(1 - \lambda) + \log p_1(b))) \\ &\quad + \log(\exp(\log \lambda + \log p_2(a|b)) + \exp(\log(1 - \lambda) + \log p_1(a))) \\ &= 0 + \log(\lambda p_2(a|*) + (1 - \lambda)p_1(a)) \\ &\quad + \log(\lambda p_2(b|a) + (1 - \lambda)p_1(b)) \\ &\quad + \log(\lambda p_2(a|b) + (1 - \lambda)p_1(a)), \end{aligned}$$

which is exactly equal to the log of the score from the probability semiring.

- The score on any specific path will be the semiring **product** of all steps along the path.
- The score of any input will be the semiring **sum** of the scores of all paths that successfully process the input.
- What happens if we use the tropical semiring?

8.4 Finite state transducers

Finite state acceptors can determine whether a string is in a language, and weighted finite state acceptors can compute a score for every string from a given alphabet. We now consider a family of automata which can **transduce** one string into another. Formally, finite state transducers (FSTs) define **regular relations** over pairs of strings. We can think of them in two different ways:

(c) Jacob Eisenstein 2014-2017. Work in progress.

- **Recognizer:** An FST accepts a pair of strings (input and output) if the pair is in the regular relation defined by the transducer.
- **Translator:** An FST takes an input string, and returns an output, such that the input/output pair is in the regular relation.

Like FSAs, finite-state transducers are defined as tuples. In this case, we define $M = \langle Q, \Sigma, \Delta, q_0, F, \delta, \sigma \rangle$, including:

- a finite set of states $Q = \{q_0, q_1, \dots, q_n\}$;
- the finite alphabets Σ for input symbols and Δ for output symbols;
- an initial state $q_0 \in Q$, and a set of final states $F \subseteq Q$;
- a transition function $\delta : \langle Q \times \Sigma^* \rangle \rightarrow \langle Q \times \Delta^* \rangle$.

Example Consider the following FST, shown in Figure 8.6, which performs **pluralization** of some English words:

$$Q = \{q_0, q_{\text{regular}}, q_{\text{needs-e}}, q_{\text{pluralized}}\} \quad (8.19)$$

$$N = \{aardvark, \dots, wish, wit, \dots, zyzzzyva^2\} \text{ (the set of all English nouns)} \quad (8.20)$$

$$\Sigma = N \cup \{+PL\} \quad (8.21)$$

$$\Delta = N \cup \{+s, +es\} \quad (8.22)$$

$$q_0 = q_0 \quad (8.23)$$

$$F = \{q_{\text{regular}}, q_{\text{needs-e}}, q_{\text{pluralized}}\} \quad (8.24)$$

$$\begin{aligned} \delta = \{ & \langle \langle q_0, aardvark \rangle \rightarrow \langle q_{\text{regular}}, aardvark \rangle \rangle, \\ & \langle \langle q_0, wish \rangle \rightarrow \langle q_{\text{needs-e}}, wish \rangle \rangle, \\ & \langle \langle q_0, wit \rangle \rightarrow \langle q_{\text{regular}}, wit \rangle \rangle, \\ & \dots \\ & \langle \langle q_{\text{regular}}, +PL \rangle \rightarrow \langle q_{\text{pluralized}}, +s \rangle \rangle \\ & \langle \langle q_{\text{needs-e}}, +PL \rangle \rightarrow \langle q_{\text{pluralized}}, +es \rangle \rangle \} \end{aligned} \quad (8.25)$$

This machine will accept the pairs $\langle wit+PL, wits \rangle$, $\langle wish+PL, wishes \rangle$, $\langle wit, wit \rangle$, but not the pairs $\langle wit+PL, wites \rangle$, $\langle wish+PL, wists \rangle$, $\langle wish+PL, wish \rangle$. Thus, it correctly handles a small part of English orthography for pluralization; with a different word list, it could also be used to conjugate verbs to third-person singular. Consider how you might modify this FST to perform lemmatization.

Non-determinism Unlike non-deterministic finite state acceptors, not all non-deterministic finite state transducers (NFSTs) can be determinized. However, special subsets of NFSTs called **subsequential** transducers can be determinized efficiently (see 3.4.1 in Jurafsky and Martin (2009)).




Figure 8.6: A finite state transducer for pluralizing English words.

8.5 Weighted FSTs

Weights can be added to FSTs in much the same way as they are added to FSAs. For any pair $\langle q \in Q, s \in \Sigma^* \rangle$, we have a set of possible transitions, $\langle q \in Q, t \in \Delta^*, \omega \in \mathbb{K} \rangle$, with a weight ω in the domain defined by the semiring. Table 8.1 shows the relationship between FSAs, FSTs, and their weighted generalizations.

	acceptor	transducer
unweighted	FSA: $\Sigma^* \rightarrow \{0, 1\}$	FST: $\Sigma^* \rightarrow \Sigma^*$
weighted	WFSA: $\Sigma^* \rightarrow \mathbb{K}$	WFST: $\Sigma^* \rightarrow \langle \Sigma^*, \mathbb{K} \rangle$

Table 8.1: A unified view of finite state automata

Example In § 8.2, we saw how to build an FSA that would compute the edit distance from any single word. With WFSTs, we can build a general edit distance computer, which computes the edit distance between any **pair** of words.

- $Q_0 \xrightarrow[a]{a} Q_0 : 0$
- $Q_0 \xrightarrow[\epsilon]{a} Q_0 : 1$
- $Q_0 \xrightarrow[a]{\epsilon} Q_0 : 1$

The shortest path for a pair of strings $\langle s, t \rangle$ in this transducer has a score equal to the minimum edit distance between the strings (in the tropical semiring). We can think of each path as defining a potential **alignment** between s and t . That is, there are many ways to transduce *she* into *he*; in the minimum edit distance path, we have the alignment $\langle s, \epsilon \rangle, \langle h, h \rangle, \langle e, e \rangle$.

Operations on FSTs

FSTs are:

- Closed under **union**. If T_1 recognizes the relation R_1 and T_2 recognizes the relation R_2 , then there exists an FST that recognizes the relation $R_1 \cup R_2$.
- Closed under **inversion**. If T_1 recognizes the relation $R_1 = \{s_i, t_i\}_i$, then there exists an FST that recognizes the relation defined by $\{t_i, s_i\}_i$, effectively switching the inputs and outputs.
- Closed under **projection**. If T_1 recognizes the relation $R_1 = \{s_i, t_i\}_i$, then there exist FSTs that recognize the relations defined by $\{s_i, \epsilon\}_i$ and $\{\epsilon, t_i\}_i$. Note that these relations ignore either the input or the output, and so are equivalent to finite state acceptors (FSAs).
- Not closed under **difference**, **complementation**, and **intersection**;
- Closed under **composition**, as described below.

FST composition is the basis for implementing the noisy channel model in FSTs, and can be used to support dozens of cool applications. Through composition, we can create finite state **cascades** that link together several simple models; closure guarantees that the resulting model is still a WFST.

Finite state composition

Suppose we have a transducer T_1 from Σ^* to Γ^* , and another transducer T_2 from Γ^* to Δ^* . Then the composition $T_1 \circ T_2$ is an FST from Σ^* to Δ^* . More formally,

Unweighted definition iff $\langle x, z \rangle \in T_1$ and $\langle z, y \rangle \in T_2$, then $\langle x, y \rangle \in T_1 \circ T_2$.

Weighted definition

$$(T_1 \circ T_2)(x, y) = \bigoplus_{z \in \Sigma^*} T_1(x, z) \otimes T_2(z, y) \quad (8.26)$$

Note that weighted composition in the Boolean semiring is identical to unweighted composition.

Designing algorithms for automatic FST composition is relatively straightforward if there are no epsilon transitions; otherwise it's more challenging (Allauzen et al., 2009). Luckily, software toolkits like OpenFST take care of this for you.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Example

- $T_1 : Q_0 \xrightarrow[a]{x} Q_0, Q_0 \xrightarrow[b]{y} Q_0$
- $T_2 : Q_1 \xrightarrow{a} Q_1, Q_1 \xrightarrow{b} Q_2, Q_2 \xrightarrow{b} Q_2$
- $T_1 \circ T_2 : Q_1 \xrightarrow{x} Q_1, Q_1 \xrightarrow{y} Q_2, Q_2 \xrightarrow{y} Q_2$

For simplicity T_2 is written as a finite-state acceptor, not a transducer. Acceptors are a special case of transducers, where the output alphabet is $\Delta = \{\epsilon\}$.

8.6 Applications of finite state composition**Edit distance**

Consider the general edit distance computer developed in section 8.5. It assigns scores to pairs of strings. If we compose it with an FSA for a given string (e.g., *tech*), we get a WFSA, who assigns score equal to the minimum edit distance from *tech* for the input string.

- Composing an FST with a FSA yields a FSA.
- A very useful design pattern is to build a **decoding** WFSA by composing a general-purpose WFST with an unweighted FSA representing the input.
- The best path through the resulting WFSA will be the minimum cost / maximum likelihood decoding.

Transliteration

English is written in a Roman script, but many languages are not. **Transliteration** is the problem of converting strings between scripts. It is especially important for names, which don't have agreed-upon translations.

A simple transliteration system can be implemented through the noisy-channel model.

- T_1 is an English character model, implemented as a transducer so that strings are scored as $\log p_r(c_1, c_2, \dots, c_M)$.
- T_2 is a character-to-character transliteration model. This can be based on explicit rules,³ or on conditional probabilities $\log p_t(c^{(f)} \mid c^{(r)})$.
- T_3 is an acceptor for a given string that is to be transliterated.

³http://en.wikipedia.org/wiki/Romanization_of_Russian

The machine $T_1 \circ T_2 \circ T_3$ scores English character strings based on their orthographic fluency (T_1) and adequacy (T_2).

Suppose you were given an Roman-script character model and a set of foreign-script strings, but no equivalent Roman-script strings. How would you use EM to learn a transliteration model?

Knight and May (2009) provide a more complex transliteration model, which transliterates between Roman and Katakana scripts, using a deep cascade that includes models of the underlying phonology. In their model,

Word-based translation

Machine translation can be implemented as a finite-state cascade. A simple approach is to compose three automata:

- T_1 is a language model, implemented as a transducer, where every path inputs and outputs the same string, with a score equal to $\log p(w_1, w_2, \dots, w_M)$. This model's responsibility is to tell us that $p(\text{Coffee black me pleases much}) \ll p(\text{I like black coffee a lot})$.
- T_2 is the translation machine. It contains a single state, and every transition takes a word from the source language and outputs a word in the target language. The weights are typically set to $p(w^{(t)} \mid w^{(s)})$. This model should assign a high probability to $p(\text{cafe} \mid \text{coffee})$, and a low probability to $p(\text{cafe} \mid \text{tea})$.

Suppose we are translating Spanish to English. Then T_1 maps from English to English, since it is a language model in English; T_2 maps from English to Spanish. By the definition of finite state composition (Equation 8.26), the scores of the paths through these two transducers will be combined with the \otimes operator; in the probability semiring, this means we will compute $p(w^{(e)})p(w^{(s)} \mid w^{(e)}) = p(w^{(s)}, w^{(e)})$.

- T_3 is a deterministic finite-state acceptor, which accepts only the sentence to be translated. By composing $T_1 \circ T_2 \circ T_3$, we get a weighted finite-state acceptor for sentences in the target language (in our example, English).

Recall that the composition $T_1 \circ T_2$ represents the joint probability $p(w^{(s)}, w^{(e)})$. The effect of T_3 is to “lock” $w^{(s)}$ to the sentence to be translated. The shortest path in the composed machine $T_1 \circ T_2 \circ T_3$ thus computes,

$$\hat{w}^{(e)} = \operatorname{argmax} w^{(e)} p(w^{(s)}, w^{(e)}) \quad (8.27)$$

$$= \operatorname{argmax} w^{(e)} p(w^{(e)} \mid w^{(s)}), \quad (8.28)$$

which is the maximum-likelihood translation.

- Finally, note that we will need to allow ϵ -transitions in the translation model to handle cases like the translation of *mucho* to *a lot*. This introduces non-determinism to the finite-state cascade; again, we can think of this in terms of possible **alignments**

(c) Jacob Eisenstein 2014-2017. Work in progress.

between the source and target languages. The shortest-path algorithm computes the maximum likelihood translation while implicitly summing over all alignments.

8.7 Discriminative structure prediction

Now suppose we would like to use perceptron to learn to perform morphological segmentation. Imagine we are given a set of words $x_{1:N}$ and their true segmentations $y_{1:N}$. We would like to use perceptron to learn the weights of a WFST. How can we do it?

Recall that perceptron relies on computing a feature function $f(x, y)$. We will make this feature vector exactly equal to the finite-state transitions taken in the shortest-path transduction of x to y . That is, each potential transition $(Q_i, \omega) \rightarrow Q_o$ corresponds to some entry j in the vector $f(x, y)$, and the value $f_j(x, y)$ is equal to the number of times that transition was taken. Although FSTs can manipulate arbitrarily long strings, there will still be only a finite number of possible transitions, since both the state space and the alphabet are finite. The scores for these transitions can then be formed into the vector of weights θ , so that the score of the best path from x to y can be represented as the inner product $\theta \cdot f(x, y)$.

Let these transitions be represented in the weighted FST T . Given an instance x , we build a chain acceptor A_x . By composing T and A_x , we obtain a WFSA in which the shortest path corresponds to the prediction \hat{y} , and the transitions on this path are the feature vector $f(x, \hat{y})$. We then compute the score of the best scoring path for accepting the true y segmentation in this machine; the transitions on this path form the feature vector $f(x, y)$. Given these two feature vectors, the perceptron update is as usual: $\theta^{(t+1)} \leftarrow \theta^{(t)} + f(x, y) - f(x, \hat{y})$. Weight averaging and passive-aggressive can be applied here, just as they were applicable in straightforward classification.

But unlike classification, we have now learned a function for making predictions over an **infinite set of labels**: all possible morphological segmentations for all possible words. We were able to do this by designing a feature function that shares features across different labels: if y and \hat{y} are nearly the same, then they will involve many of the same finite-state transitions, and so the feature vector $f(x, y)$ and $f(x, \hat{y})$ will be nearly the same too. This is a powerful idea that will enable us to apply the tools of classification to a huge range of problems in language technology, including part-of-speech tagging, parsing, and even machine translation.

Chapter 9

Morphology

So far we have been focusing on NLP at the word level. Now we will explore meaning **inside of words**. We've already hinted at a morphological problem by introducing the idea of **lemmas**, where *serve/served/serving* all have the lemma *serve*.

From the perspective of document classification, these multiple forms may just seem like an annoyance, which we can get rid of by lemmatization or stemming (more on this later). But morphology conveys information which can be crucial for some applications.

Information retrieval With a search query like *bagel*, we want to get hits for the **inflected** form *bagels*; the same goes for irregular inflections like *corpus/corpora*, *goose/geese*. In **query expansion**, the search query is expanded to include all inflections of the search terms. Note that this isn't always what we want: for example, given a query for *Apple*, we may not want hits for *apples*.

Information extraction A major goal of information extraction is to capture references to events, and their properties. Event timing is conveyed in morphology: in English, we have suffixes for past tense (*she talked*), the past participle (*she had spoken*), and the present participle (*she is speaking*). Other languages can indicate many more details about event timing through morphology; for example, Romance languages like French have a much larger inventory of verb endings:

<i>J'achete un velo</i>	I buy a bicycle (now)
<i>J'acheterai un velo</i>	I will buy a bicycle
<i>J'achetais un velo</i>	I was buying a bicycle
<i>J'ai acheté un velo</i>	I bought a bicycle
<i>J'acheterais un velo</i>	I would buy a bicycle

In English, this function is mostly filled by auxiliary verbs like *will*, *was*, *had*, and *would*. This makes morphological analysis relatively less important for English, as we can get a

long way with carefully constructed n-gram patterns (Riloff, 1996). But in languages like French and Spanish — where second-language learners are tormented by conjugation tables with dozens of different inflections — there seems little alternative to morphological analysis if language technology is to generalize across many verbs.

Document classification Even document classification tasks, such as sentiment analysis, are potentially impacted by morphology. For example, suppose you are doing sentiment analysis, and you encounter the out-of-vocabulary words *unfriended*, *antichrist*, *unputdownable*, or *disenchanted*. As unknown words, they would make no contribution to the overall sentiment polarity in a bag-of-words system. But with some morphological reasoning, we can see that they are indeed strongly subjective.

Translation In addition to recognizing morphology, there are applications in which we need to produce it. Translation is a classic case, especially when translating from morphologically simple languages like English and Chinese to morphologically rich languages, like French, Czech, German, and Swahili. Here again, a purely word-based approach would suffer from data sparsity: relatively rare words would be unlikely to be seen in every inflection, and thus the translation system would be unable to produce them.

Morphology, Orthography, and Phonology

Morphology interacts closely with two related systems: orthography and phonology. The **surface form** of a word is the form that is written down or spoken. This form results from the interactions between morphology and the orthographic and phonological systems. More specifically:

- **Morphology** describes how meaning is constructed from combining affixes. For example, it is a morphological fact of English that adding the affix +S to many nouns creates a plural.

$$\text{berry} + \text{PLURAL} \rightarrow \text{berry+s}$$

Morphological rules may also include stem changes, such as *goose* + PLURAL \rightarrow *geese*.

- **Orthography** specifically relates to writing. For example,

$$\text{berry+s} \rightarrow \text{berries}$$

is an orthographic rule. We have lots of these in English, which is one reason English spelling is difficult.

- **Phonology** describes how sounds combine. For example, the different pronunciations of the final *s* in *cats* (s) and *dogs* (z) follow from a phonological rule (Bender, 2013, example 25, page 30).

(c) Jacob Eisenstein 2014-2017. Work in progress.

Surface form	lemma	features
<i>duck</i>	<i>duck</i>	NOUN+SINGULAR
<i>ducks</i>	<i>duck</i>	NOUN+PLULAR
<i>duck</i>	<i>duck</i>	VERB+PRESENT
<i>ducks</i>	<i>duck</i>	VERB+THIRDPERSON+PRESENT

Table 9.1: Fragment of a morphologically-aware dictionary

In English, morphologically distinct words may be pronounced differently even when they are spelled the same, and this can reflect morphological differences. *read*+PRESENT vs. *read*+PAST. Conversely, morphological variants may be spelled differently even when they sound the same, like *The Champions' league* versus *The Champion's league* versus *The Champions league*.

Productivity

One idea for dealing with morphology is to build a morphologically-aware dictionary. The keys in this dictionary would correspond to **surface forms**, such as *served*. The values would include both the underlying **lemma** as well as any morphological features: in this case, the lemma is **serve**, and the feature is PAST. Given such a dictionary, we simply look up each surface form that we encounter.

As shown in the example in Table 9.1, we may need multiple entries for the same surface form; this means that there is ambiguity, so simple lookup will not suffice. Still another problem is that morphology is **productive**, meaning that it applies to new words. If you only know the words *Google* or *iPad*, you can immediately understand their inflected forms.

- Have you Googled that yet?
- I have broken all three iPads.

Derivational morphology (more on this later) is productive in another way: you can produce new words by applying morphological changes to existing words. hyper+un+desire+able+ity

In some languages, derivational morphology can create extremely complicated words. Jurafsky and Martin (2009) have a fun example from Turkish:

In the homework, you'll see examples from Swahili, which also has complex morphology. A dictionary of all possible surface forms in such languages would be gargantuan. So instead of building a static dictionary, we will try to model the underlying morphological and orthographic rules.

(c) Jacob Eisenstein 2014-2017. Work in progress.

A Turkish word

uygarlaştıramadıklarımızdanmışsınızcasına

uygar_laş_tır_ama_dık_lar_ımız_dan_mış_sınız_casına

“as if you are among those whom we were not able to civilize (=cause to become civilized)”

uygar: *civilized*

_laş: *become*

_tır: *cause somebody to do something*

_ama: *not able*

_dık: *past participle*

_lar: *plural*

_ımız: *1st person plural possessive (our)*

_dan: *among (ablative case)*

_mış: *past*

_sınız: *2nd person plural (you)*

K. Oflazer pc to J&M

Figure 9.1: From (Jurafsky and Martin, 2009)

9.1 Types of morphemes

There are two broad classes of morphemes: **stems** and **affixes**. Intuitively, stems are the “main” part of meaning, and affixes are the modifiers. Typically, **stems** can appear on their own (they are **free**) and affixes cannot (they are **bound**).

Affixes can be categorized by where they appear with respect to the stem.

- **Prefixes:** *un+learn, pre+view*.

- These examples are **derivational**, in that they form new words, rather than forming grammatical variants of the same word (*inflectional* morphology; more on this in § 9.2).
- English has no inflectional prefixes, but other languages do. For example, in Swahili, *u-na-kata* means *you are cutting*, while *u-me-kata* means *you have cut*. In this example, *na* and *me* are prefixes, *kata* is the root.¹

- **Suffixes** are the typical way of inflecting words in English, and in other languages in the Indo-European family. For example, in English: *I learn+ed, She learn+s, three ap-*

¹Would it be better to think about *u*, *na*, and *me* as words? This example suggests that the word/affix distinction is not always clear-cut.

ple+s, *four fox+es*. English suffixes can also be derivational: for example: *modern+ity*, *fix+able*, and *deriv+ation+al*.

- **Circumfixes** go around the stem.
 - German has a circumfix for the past participle: *sagen* (say) → *ge+sag+t* (said)
 - English has a very small number of circumfix examples: *bold* → *em+bold+en*, and, arguably, *light* → *en+light+en*. Both of these examples are derivational.
 - French negation can be seen as a circumfix: *Je mange+NEG* → *Je ne mange pas* (I do not eat).²
 - More generally, morphemes can be non-contiguous, e.g. (Bender, 2013, example 7, page 12):

(7)	Root	Pattern	Part of Speech	Phonological Form	Orthographic Form	Gloss
	ktb	CaCaC	(v)	katav	כתב	'wrote'
	ktb	hiCCiC	(v)	hixtiv	הכתוב	'dictated'
	ktb	miCCaC	(n)	mixtav	מכתב	'a letter'
	ktb	CCaC	(n)	ktav	כתב	'writing, alphabet'

[heb]

In this example, the root *ktb* (related to writing) is combined with patterns that indicate where to insert vowels to produce different parts-of-speech and meanings.

- **Infixes** go inside the stem.
 - In Tagalog (spoken in the Philippines), the root *hingi* indicates a request, and the infix *um* creates *humingi*, as in *I asked*.
 - English, *absolutely+fucking* →

(9.1) *absofuckinglutely*

(9.2) *?absfuckingsolutely*

where the '?' prefix indicates questionable linguistic acceptability.

- Morphology may be **non-segmental**, meaning that it doesn't involve any affix at all. For example, the pluralization of *goose* to *geese* is not accomplished through any affix, but through vowel alteration; the past tense marking of *eat* → *ate* is another example

²In spoken French, the *ne* is gradually disappearing, so that *Je mange pas* is now acceptable.

of this phenomenon, known as *apophony*. Languages in which morphemes are represented by affixes that are “glued together” (like *talk+ed* or *think+ing*) are known as **agglutinative**; languages in which morphemes are represented by changes to spelling and sound are known as **fusional**.

- What about words like *fish*, which have the same form in both singular and plural? We say that this word has a **zero** plural.

9.2 Types of morphology

Morphology serves a variety of linguistic functions, and acts in a variety of ways. Inflectional and derivational morphology are distinguished by their function; other forms of morphology, such as cliticization and compounding are distinguished by how they work. In this section, we will focus mainly on inflectional and derivational morphology, describing their roles in English, and in other languages when there is no adequate example in English.

Inflectional morphology

Inflectional morphology adds information about the stem, typically grammatical properties such as tense, number, and case. English has a relatively simple system of inflectional morphology, compared to many other languages.

Affix	Syntactic/semantic effect	Examples
-s	NUMBER: plural	<i>cats</i>
-'s	possessive	<i>cat's</i>
-s	TENSE: present, SUBJ: 3sg	<i>jumps</i>
-ed	TENSE: past	<i>jumped</i>
-ed/-en	ASPECT: perfective	<i>eaten</i>
-ing	ASPECT: progressive	<i>jumping</i>
-er	comparative	<i>smaller</i>
-est	superlative	<i>smallest</i>

Figure 9.2: From (Bender, 2013)

Nouns

English nouns are marked for **number** and **possession**. Number is typically marked by the suffix +s, e.g., *hat* + PLURAL → *hat+s*, but some words are pluralized differently, e.g.,

(c) Jacob Eisenstein 2014-2017. Work in progress.

geese, children, and fish. Number is binary in English (singular versus plural), but many languages, such as Arabic and Sanskrit, include an additional **dual** number for groups of two. English has residual traces of the dual number, with *both* versus *all* and *either* versus *any*. Some Austronesian languages even have a **trial** number, for groups of three, and languages such as Arabic have a **paucal** number, for small groups. Conversely, nouns are not marked for number at all in Japanese and Indonesian.

Many languages mark nouns for **case**, which is the syntactic role that the noun plays in the sentence. In English, we do distinguish the case of some pronouns:

- *He* (NOMINATIVE) *gave her* (OBLIQUE) *his* (GENITIVE) *guitar*.
- *She gave him her guitar*.
- *I gave you our guitar*.
- *You gave me your guitar*.

The third person masculine pronoun appears as *he* in the nominative case, *him* in the oblique case, and *his* in the genitive case. English distinguishes these cases for all personal pronouns except for the second person, where the nominative and oblique cases are both *you*.

Other languages — such as Latin, Russian, Sanskrit, and Tamil — mark the case of all nouns. These languages have additional cases, such as dative (indirect object), accusative (direct object), and vocative (address). In German, noun is not inflected for case, but the articles and adjectives are, as shown in example 49 from Bender (2013):

- (9.3) *Der alte Mann gab dem kleinen Affen die grosse Banane.*
 The old man (NOM) gave the little monkey (DATIVE) the big banana (ACCUSATIVE)

Notice how *der*, *dem*, and *die* all mean the same thing (*the*), but they are spelled differently due to the case marking. The adjectives (*alte*, *kleinen*, *grosse*) are also marked for case.

Many languages — such as Romance languages — mark the gender and number of nouns by inflecting the article and adjective. e.g., Spanish:

- (9.4) *El coche rojo pasó la luz roja.*
 The red car ran the red light.
- (9.5) *Los coches rojos pasó las luces rojas.*
 The red cars ran the red lights.

Here, *la* is the feminine article and *el* is the masculine article; the adjective for *red* is inflected to *roja* when describing a feminine noun (*luz*, meaning light), and *rojo* when describing a masculine noun (*coche*, meaning car). The article and adjective must **agree** with noun for the sentence to be grammatical. The following examples are ungrammatical for this reason:

- (9.6) **Los coches rojo pasó la luce rojas*

(c) Jacob Eisenstein 2014-2017. Work in progress.

(9.7) **Los coches rojas pasó las luces rojos*

In English, demonstrative determiners mark number: e.g., *this book* vs *these books*, and the determiner and noun must agree, e.g. **this books*. Agreement is also required between subject and verb, as we will see shortly.

Romance languages like Spanish and French mark gender as masculine and feminine, but it need not be binary:

- English pronouns include neuter *it*; German, Sanskrit, and Latin do this for all nouns.
- Danish and Dutch distinguish **neuter** from **common** gender.[**todo: example**]
- Other languages distinguish **animate** and **inanimate** genders.

Verbs

English verbs are inflected for tense and number distinguishing past (*she acted*), present (*you act*), and third person singular (*she acts*). As with nouns, these inflections may change the orthography (*plan+ed* → *planned*), and there are many irregular patterns, e.g. *they eat* / *she eats* / *we ate*. English verbs are also inflected for aspect, distinguishing the perfective (*I had eaten*) and progressive (*I am eating*). The perfective and the past tense are identical for regular verbs, e.g. *we had talked*, *we talked*.

Many languages (e.g., Chinese and Indonesian), do not mark tense with morphology. For example, Indonesian uses function words rather than morphology to distinguish tense (Table 9.2).

<i>Saya makan apel</i>	I eat an apple
<i>Saya sedang makan apel</i>	I am eating an apple
<i>Saya telah makan apel</i>	I already ate an apple
<i>Saya akan makan apel</i>	I will eat an apple

Table 9.2: Indonesian uses function words (*sedang*, *telah*, *makan*) rather than morphology to distinguish verb tense. [**todo: switch to exe**]

Romance languages distinguish many more tenses than English with morphology. For example, Spanish has multiple past tenses: **preterite** and **imperfect**, distinguishing events that occurred at a specific past point in time from a continuous or repeated past state:

(9.8) *I ate onions yesterday*

Comí cebollas ayer

(9.9) *I ate onions every day*

Comía cebollas cada día

(c) Jacob Eisenstein 2014-2017. Work in progress.

Spanish and French also have endings for conditional (*comería cebollas, I would eat onions*) and future (*comeré cebollas, I will eat onions*). In English, these differences are marked with time signals rather than morphology. In French and Spanish, time signals are also an option, e.g. *voy a comer cebollas*, which literally translates to *I am going to eat onions*.

Romance languages also have separate verb forms for every combination of number and person, while in English, only the third-person singular is distinguished:

- English: *I speak / you speak / she speaks / we speak / you (pl) speak / they speak*
- Spanish: *Yo hablo / tu hablas / ella habla / nosotros hablamos / vosotros hablais / ellas hablan*
- French: *Je parle / tu parles / elle parle / nous parlons / vous parlez / ils parlent*

In Spanish and in many other Romance languages (but not French), the verb morphology is sufficiently descriptive that the subject is often omitted, since it can often be easily recovered from the verb ending and the context.

Other things can be marked with affixes, such as **evidentiality** – how the speaker came to know the information. In Eastern Pomo (a California language), there are verb suffixes for four evidential categories (McLendon, 2003):

-ink'e	nonvisual sensory
-ine	inferential
-le	hearsay
-ya	direct knowledge

Adjectives and adverbs

Adjectives in English mark comparative and superlative (*taller, tallest*). Adverbs can mark comparative and superlative too: *Yangfeng paddles fast, Yi paddles faster, Uma paddles fastest*. As we have seen, adjectives can mark gender and number in languages like French and Spanish, where they are required to agree with the noun and determiner; adjectives also mark case in languages like German and Latin.

Synthetic and isolating languages

Languages with complex morphology are called **synthetic**; languages with simple morphology are called **isolating** or **analytic**. The **index of synthesis** quantifies this property by measuring the ratio of the number of morphemes in a given text to the number of words. On this index, English is relatively, but not extremely, analytic.

An approximation of the index of synthesis is the type-token ratio. Can you see why? If you count the number of unique surface forms in 10K *parallel* sentences from a corpus of European Parliament transcripts, you get:

(c) Jacob Eisenstein 2014-2017. Work in progress.

Language	Index of synthesis
Vietnamese	1.06
Yoruba	1.09
English	1.68
Old English	2.12
Swahili	2.55
Turkish	2.86
Russian	3.33
Inuit (Eskimo)	3.72

Figure 9.3: From Bender (2013)

- English: 16k distinct word types
- French: 22k
- German: 32k
- Finnish: 55k

Derivational Morphology

Derivational morphology is a way to create new words and change part-of-speech.

- **nominalization**
 - *V + -ation: computerization*
 - *V + -er: walker*
 - *Adj + -ness: fussiness*
 - *Adj + -ity: obesity*
- **negation:** *undo, unseen, misnomer*
- **adjectivization:** *V + -able : doable, thinkable, N + -al : tonal, national, N + -ous: famous, glamorous*
- **abverbization:** *ADJ + -ily: clumsily*
- **lots more:** *rewrite, phallocentrism, ...*

You can create totally new words this way.

word → *wordify* → *wordification* → *wordificationism* → *antiwordificationism* → *hyperantiwordificationism*

As with inflection, derivational morphology can require orthographic changes, e.g. *true+ly* → *truly* and *fussy+ness* → *fussiness*. It can also cause phonological changes, such

(c) Jacob Eisenstein 2014-2017. Work in progress.



Figure 9.4: The written space in watermelon disappeared as the word became more frequent over the 19th century. From Google ngrams.

as the change emphasis from *imPOSSible* to *impossiBILity*, and the change in vowel from *ferTILE* to *ferTILity*.

Other types of morphology

Cliticization combines *Georgia*+*'s* into *Georgia's*; the possessive clitic *'s* is syntactically independent but phonologically dependent. This syntactic independence can be seen in examples like (Bender, 2013, example 21):

(9.10) Jesse met the president of the university's cousin

In this example, the possessive modifies the *president*, but it attaches to the right edge of the entire noun phrase.

- Pronouns appear as clitics in French, e.g., *j'accuse* (I accuse), as does negation *Je n'accuse personne* (I don't accuse anyone).
- Another example is from Hebrew: *l'shana tova* (literally for year good, meaning happy new year); the preposition *for* appears as a clitic.

Compounding combines two words into a new word:

(9.11) *cream* → *ice cream*

We can think of *ice cream* as a word since it is a non-compositional combination of *ice* and *cream*. Perhaps someday the written space will be dropped, as it has been in *watermelon* (Figure 9.4).

Portmanteaus combine words, truncating one or both.

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 9.5: Google n-grams plots for inflections of *show*. While the past participle *had shown* is decreasing, this does not seem to be due to competition from the more regular *had showed*; rather, there appears to be a broader decrease in frequency of the past participle, shown by the parallel pattern for *had known*.

(9.12) *smoke + fog* → *smog*

(9.13) *glass + asshole* → *glasshole*

Urban Dictionary is a fun source of contemporary portmanteaus.

Irregularities

English morphology contains a lot of irregularities: *know/knew/known*, *foot/feet*, *go/went*, etc. if you are not a native speaker, learning these was probably a pain in the neck. The good news is that there are fewer of these all the time! English is undergoing a process in which these irregular forms are gradually being replaced: for example, the past tense of *show* used to be *shew*, just as the past tense of *know* is still *knew* (Figure 9.5a). This transformation remains incomplete, as the past participle of *show* is still *shown*, and not *showed* (Figure 9.5b). However, this example points to the bad news for language learners: the most frequently-occurring words, like *know*, will be the last to change — if ever!

9.3 Computing and morphology

In this section, we will briefly overview some of the computational problems related to morphology. We don't yet have many tools to solve these problems, but we will soon: chapter 8 presents finite-state automata, which are the workhorse of morphological analysis in NLP. For now, we will simply state the problem definitions, and discuss some of the challenges involved.

Lemmatization

[todo: write]

(c) Jacob Eisenstein 2014-2017. Work in progress.

Stemming

[**todo: write**]

Generation

[**todo: write**]

Normalization

[**todo: write**]

Chapter 10

Context-free grammars

So far we've explored finite-state models, which are capable of defining regular languages (and regular relations).

- **representations:** (weighted) finite state automata
- **probabilistic models:** HMMs (as a special case), CRFs
- **algorithms:** Viterbi, Forward-Backward, $\mathcal{O}(MK^2)$ time complexity.
- **linguistic phenomena:**
 - morphology
 - language models
 - part-of-speech disambiguation
 - named entity recognition (chunking)

Clearly there are formal languages that are not describable using finite-state machinery, such as the classic $a^n b^n$. But is the finite-state representation enough for natural language?

10.1 Is English a regular language?

In this section, we consider a proof that English is not regular, and therefore, no finite-state automaton could perfectly model English syntax. The proof begins by noting that regular languages are closed under **intersection**.

- $K \cap L$ is the set of strings in both K and L
- $K \cap L$ is regular iff K and L are regular

The proof strategy is as follows:

- Let K be the set of grammatical English sentences
- Let L be some regular language
- Show that the intersection is not regular

We're going to prove this using **center embedding**, as shown in the examples below:

(10.1) *The cat is fat.*

(10.2) *The cat that the dog chased is fat.*

(10.3) **The cat that the dog is fat.*

(10.4) *The cat that the dog that the monkey kissed chased is fat.*

(10.5) **The cat that the dog that the monkey chased is fat.*

Proof sketch:

- K is the set of grammatical english sentences.
It excludes examples (10.3) and (10.5).
- L is the regular language *the cat (that N)⁺ V_t ⁺ is fat*. It is crucial to see that this language is itself regular, and could be recognized with a finite-state acceptor.
- The language $L \cap K$ is *the cat (that N)ⁿ V_t ⁿ is fat*. This language is homomorphic to $a^n b^n$, which is known not to be regular. Since L is regular and $L \cap K$ is not regular, it follows that K cannot be regular.

It is important to understand that the issue is not just infinite repetition or productivity; FSAs can handle productive phenomena like *the big red smelly plastic figurine*. It is specifically the center-embedding phenomenon, because this leads to the same structure as the classic $a^n b^n$ language. What do you think of this argument?

Is deep center embedding really part of English?

Karlsson (2007) searched for multiple (phrasal) center embeddings in corpora from 7 languages:

- Very few examples of double embedding
- Only 13 examples of triple embedding (none in speech)
- Zero examples of quadruple embeddings

Note that we can build an FSA to accept center-embedding up to any finite depth. So in practice, we could build an FSA that accepts any center-embedded sentence that has ever been written. Does that defeat the proof? Chomsky and many linguists distinguish between

(c) Jacob Eisenstein 2014-2017. Work in progress.

Competence the fundamental abilities of the (idealized) human language processing system;

Performance real utterances produced by speakers, subject to non-linguistic factors such as cognitive limitations.

Even if English *as performed* is regular, the underlying generative grammar may be context-free... **or beyond**.

How much expressiveness do we need?

Shieber (1985) makes a similar argument, showing that case agreement in Swiss-German cross-serial constructions is homomorphic to a formal language $wa^mb^nc^md^ny$, which is weakly non-context free. In response to the objection that all attested constructions are finite, Shieber writes:

Down this path lies tyranny. Acceptance of this argument opens the way to proofs of natural languages as regular, nay, **finite**.

Regardless of what we think of these theoretical arguments, the fact is that in practice, many real constructions appear to be much simpler to handle in context-free rather than finite-state representations. For example,

(10.6) *The **processor** has 10 million times fewer transistors on it than today's typical micro-processors, **runs** much more slowly, and **operates** at five times the voltage...*

The verbs *has*, *runs*, and *operates* agree with the subject *the processor*; we want to accept this sentence, but reject all sentences in which this subject-verb agreement is lost. Handling this in a finite state representation would building separate components for third-person singular and non-third-person singular forms, and then replicating essentially all of verb-related syntax in each component. A **grammar** — formally defined in the next section — would vastly simplify things:

$$\begin{aligned} S &\rightarrow \text{NN VP} \\ \text{VP} &\rightarrow \text{VP3S} \mid \text{VPN3S} \mid \dots \\ \text{VP3S} &\rightarrow \text{VP3S}, \text{VP3S}, \text{and VP3S} \mid \text{VBZ} \mid \text{VBZ NP} \mid \dots \end{aligned}$$

10.2 Context-Free Languages

The Chomsky Hierarchy Every automaton defines a language, and different classes of automata define different classes of languages. The Chomsky hierarchy formalizes this set of relationships:

- **finite-state automata** define **regular** languages;

(c) Jacob Eisenstein 2014-2017. Work in progress.

- **pushdown automata** define **context-free** languages;
- **Turing machines** define **recursively-enumerable** languages.

In the Chomsky hierarchy, context-free languages (CFLs) are a strict generalization of regular languages.

regular languages	context-free languages
regular expressions	context-free grammars (CFGs)
finite-state machines	pushdown automata
paths	derivations

Context-free grammars define CFLs. They are sets of permissible *productions* which allow you to **derive** strings composed of surface symbols. An important feature of CFGs is *recursion*, in which a nonterminal can be derived from itself.

More formally, a CFG is a tuple $\langle N, \Sigma, R, S \rangle$:


- N a set of non-terminals
- Σ a set of terminals (distinct from N)
- R a set of productions, each of the form $A \rightarrow \beta$,
where $A \in N$ and $\beta \in (\Sigma \cup N)^*$
- S a designated start symbol

Context free grammars provide rules for generating strings.

- The left-hand side (LHS) of each production is a non-terminal $\in N$
- The right-hand side (RHS) of each production is a sequence of terminals or non-terminals, $\{n, \sigma\}^*$, $n \in N$, $\sigma \in \Sigma$.

A **derivation** t is a sequence of steps from S to a surface string $w \in \Sigma^*$, which is the **yield** of the derivation. A derivation can be viewed as trees or as bracketings, as shown in Figure 10.1.

If there is some derivation t in grammar G such that w is the yield of t , then w is in the language defined by the grammar. Equivalently, for grammar G , we can write that $|\mathcal{T}_G(w)| \geq 1$. When there are multiple derivations of w in grammar G , this is a case of derivational **ambiguity**; if any such w exists, then we can say that the grammar itself is ambiguous.



Example The grammar below handles the case of center embedding:

$$S \rightarrow NP VP_1 \quad (10.1)$$

$$NP \rightarrow the\ NP \mid NP\ RELCLAUSE \quad (10.2)$$

$$RELCLAUSE \rightarrow that\ NP\ V_t \quad (10.3)$$

$$V_t \rightarrow ate \mid chased \mid befriended \mid \dots \quad (10.4)$$

$$N \rightarrow cat \mid dog \mid monkey \mid \dots \quad (10.5)$$

$$VP_1 \rightarrow is\ fat \quad (10.6)$$

Here we are using a shorthand, where $\alpha \rightarrow \beta \mid \gamma$ implies two productions, $\alpha \rightarrow \beta$ and $\alpha \rightarrow \gamma$.

Semantics Ideally, each derivation will have a distinct semantic interpretation, and all possible interpretations will be represented in some derivation.

$$\begin{aligned} & (NP(NP\ Ban\ (PP\ on\ (NP\ nude\ dancing\))) \\ & \quad (PP\ on\ (NP\ Governor's\ desk\))) \end{aligned}$$

$$\begin{aligned} & (NP\ Ban\ (PP\ on\ (NP(NP\ nude\ dancing\) \\ & \quad (PP\ on\ (NP\ Governor's\ desk\))))) \end{aligned}$$

In practice, this is quite hard to achieve with context-free grammars. For example, Johnson (1998) notes that there are three possible derivations for the verb phrase *ate dinner on the table with a fork*:

“flat” (*ate dinner (on the table) (with a fork)*)

“two-level” (*((ate dinner) (on the table) (with a fork))*)

“adjunction” (*((((ate dinner) (on the table)) (with a fork))*)

In this case, there doesn't seem to be any meaningful difference between these derivations. The grammar could avoid this problem by limiting its set of productions, but this change might cause problems in other cases.

10.3 Constituents

Our goal in using context-free grammars is usually not to determine whether a string is in the language defined by the grammar, but to acquire the derivation itself, which should explain the organization of the text and give some clue to its meaning. Therefore, a key question in grammar design is how to define the non-terminals.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Every non-terminal production **yields** a contiguous portion of the input string. For example, the VP non-terminal in Figure 10.1 (both parses) yields the substring *eats sushi with chopsticks*, and the PP non-terminal yields *with chopsticks*. These substrings, which are bracketed in the figure, are known as **constituents**. The main difference between the two parses in Figure 10.1 is that the second parse includes *sushi with chopsticks* as a constituent, and the first parse does not.

In a given string, which substrings should be constituents? Linguistics offers several tests for constituency, including: substitution, coordination, and movement.

Substitution

Constituents generated by the same non-terminal should be substitutable in many contexts:

- (10.7) (NP *The ban*) *is on the desk.*
- (10.8) (NP *The Governor's desk*) *is on the desk.*
- (10.9) (NP *The ban on dancing on the desk*) *is on the desk.*
- (10.10) *(PP *On the desk*) *is on the desk.*

A more precise test for whether a set of substrings constitute a single category is whether they can be replaced by the same pronouns.

- (10.11) (NP *It*) *is on the desk.*

What about verbs?

- (10.12) *I* (V *gave*) *it to Anne.*
- (10.13) *I* (V *taught*) *it to Anne.*
- (10.14) *I* (V *gave*) *Anne a fish*
- (10.15) **I* (V *taught*) *Anne a fish*

This suggests that *gave* and *taught* are not substitutable. We might therefore need non-terminals that distinguish verbs based on the arguments they can take. The technical name for this is **subcategorization**. [todo: more details]

Coordination

Constituents generated by the same non-terminal can usually be *coordinated* using words like *and* and *or*:

- (10.16) *We fought* (PP *on the hills*) *and* (PP *in the hedges*).
- (10.17) *We fought* (ADV_P *as well as we could*).
- (10.18) **We fought* (ADV_P *as well as we could*) *and* (PP *in the hedges*).

(c) Jacob Eisenstein 2014-2017. Work in progress.

Like all such tests, coordination does not always work:

(10.19) *She* (_{VP} *went*) (_{PP} *to the store*).

(10.20) *She* (_{VP} *came*) (_{PP} *from the store*).

(10.21) *She* (_? *went to*) *and* (_? *came from*) *the store*.

Typically we would not think of *went to* and *came from* as constituents, but they can be coordinated.

Movement Valid constituents can be moved as a unit, preserving grammaticality. There are a number of ways in which such movement can occur in English.

Passivization (10.22) *(The governor) banned (nude dancing on his desk)*

(10.23) *(Nude dancing on his desk) was banned by (the governor)*

Wh- movement (10.24) *(Nude dancing was banned) on (the desk).*

(10.25) *(The desk) is where (nude dancing was banned)*

Topicalization (10.26) *(He banned nude dancing) to appeal to conservatives.*

(10.27) *To appeal to conservatives, (he banned nude dancing).*

10.4 A simple grammar of English

A goal of grammar design is to thread the line between two potential problems:

Overgeneration deriving strings that are not grammatical.

Undergeneration failing to derive strings that are grammatical.

To avoid undergeneration in a real language, we would need thousands of productions. Designing such a large grammar without overgeneration is extremely difficult.

Typically, grammars are defined in conjunction with large-scale **treebank** annotation projects.

- An annotation guideline specifies the non-terminals and how they go together.
- The annotators then apply these guidelines to data.
- The grammar rules can then be read off the data.

The Penn Treebank (PTB) contains one million parsed words of Wall Street Journal text (Marcus et al., 1993).

In the remainder of this section, we consider a small grammar of English.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Noun phrases

Let's start with noun phrases:

(10.28) *She sleeps* (Pronoun)

(10.29) *Arlo sleeps* (Proper noun)

These examples suggest that pronouns and proper nouns are substitutable, so we can define a production,

$$\text{NP} \rightarrow \text{PRP} \mid \text{NNP}, \quad (10.7)$$

where NP stands for **noun phrase**. In this grammar, we will treat part-of-speech tags as the terminal vocabulary, but we could easily extend this to words by defining productions,

$$\text{PRP} \rightarrow \textit{she} \mid \textit{he} \mid \textit{I} \mid \textit{you} \dots \quad (10.8)$$

$$\text{NNP} \rightarrow \textit{Arlo} \mid \textit{Abigail} \dots \quad (10.9)$$

What else could be a noun phrase?

(10.30) *A lobster sleeps*

(10.31) *The lobster sleeps*

(10.32) *Lobsters sleep*

(10.33) **Lobster sleeps*

The first two examples show that we can have common nouns (NN) as long as they are preceded by determiners (DT). We can also have plural nouns (NNS). But we cannot have common nouns **without** determiners — the final example doesn't work unless *Lobster* is a proper name.

We can handle these cases by defining a new nonterminal, NOM, which stands for **nominal**. A nominal is a constituent that cannot be a noun phrase by itself, but requires a determiner. We then add two productions:

$$\text{NP} \rightarrow \text{DT NOM} \mid \text{NNS} \quad (10.10)$$

$$\text{NOM} \rightarrow \text{NN} \mid \text{NNS} \quad (10.11)$$

Notice that these productions also allow *The lobsters sleep*, using the $\text{NOM} \rightarrow \text{NNS}$ production.

Noun phrases may also contain various **modifiers**.

(10.34) *The blue fish sleeps* (adjective)

(10.35) *The four crabs sleep* (cardinality)

(c) Jacob Eisenstein 2014-2017. Work in progress.

We could try to handle these cases by adding to the nominal productions,

$$\text{NOM} \rightarrow \text{JJ NOM} \mid \text{CD NOM} \quad (10.12)$$

where JJ is an adjective and CD is a **cardinality**. Note that these productions are **recursive**, because NOM appears on the right-hand side. This means we can use the production to create a nominal with an infinite number of modifiers. This works for adjectives (*the angry blue plastic lobster*), but not for cardinals: **the four three crabs* is ungrammatical, so this grammar now **overgenerates**. We would need to further refine the grammar to handle this case properly, as well as to avoid **undergenerating** cases like *four crabs sleep*.

Modifiers can also come at the end of the noun phrase:

(10.36) *The girl from Omaha sleeps* (prepositional phrase)

(10.37) *Cats in Catalonia cry* (prepositional phrase)

(10.38) *The student who ate 15 donuts sleeps* (relative clause)

(10.39) *Mary from Omaha sleeps*

(10.40) *Cats who are in Catalonia cry*

(10.41) *?Mary who ate 15 donuts sleeps*

These examples suggest that **prepositional phrases** (*from Omaha, in Catalonia*) can be attached to the end of any noun phrase. For **relative clauses** (*...who ate 15 donuts*), the situation is somewhat less clear. If we accept examples like (10.41), then we can handle both of these cases by adding the following NP productions,

$$\text{NP} \rightarrow \text{NP PP} \mid \text{NP RELCLAUSE} \quad (10.13)$$

We again have recursion: because the NP tag appears on the right side of the production, it is possible generate infinitely long noun phrases, like *the student from the city in the state below the river ...*

So overall, we can summarize the NP fragment of the grammar as,

$$\begin{aligned} \text{NP} &\rightarrow \text{PRP} \mid \text{NNP} \mid \text{DT NOM} \mid \text{NP PP} \mid \text{NP RELCLAUSE} \\ \text{NOM} &\rightarrow \text{NN} \mid \text{ADJP NOM} \mid \text{CD NNS} \mid \text{NNS} \end{aligned}$$

Are we done? Not close. We still haven't handled cardinal numbers in satisfactory way, and we are leaving out important details like number agreement, causing the grammar to overgenerate examples like *Mary sleep*. The process of grammar design would involve continuing to probe at the grammar with these sorts of examples until we handled as many as possible.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Adjectival and prepositional phrases

The noun phrase grammar mentioned prepositional phrases, such as

(10.42) *cats from Catalonia*

(10.43) *pizza in the refrigerator*

(10.44) *pizza in the old, broken refrigerator*

(10.45) *the red switch under the panel next to the radiator*

These examples suggest that prepositional phrases are formed by placing a preposition before any noun phrase — including noun phrases that already contain prepositional phrases, as in (10.45). This suggests the simple production,

$$PP \rightarrow P \text{ NP}. \quad (10.14)$$

The noun phrase fragment also includes adjective modifiers, like *the blue lobster*. But in fact, adjectives can combine into phrases.

(10.46) *the large blue fish*

(10.47) *the very funny hat*

The first example, we have two adjectives; in the second, we have an adverb followed by an adjective. This suggests the following productions:

$$ADJP \rightarrow JJ \mid RB \text{ ADJP} \mid JJ \text{ ADJP} \quad (10.15)$$

$$NOM \rightarrow ADJP \text{ NN} \mid ADJP \text{ NNS} \quad (10.16)$$

Notice that if we instead added $NOM \rightarrow ADJP \text{ NOM}$, we would be introducing a considerable amount of ambiguity to the grammar. This would give us two different ways of generating multiple adjectives: by a series of NOM productions, or a series of ADJP productions. The proposed solution here increases the number of production rules, but decreases the number of ways to derive the same string.

Verb phrases

Let's now consider the verb and its modifiers.

(10.48) *She sleeps*

(10.49) *She sleeps restlessly*

(10.50) *She sleeps at home*

(10.51) *She eats sushi*

(10.52) *She gives John sushi*

Each of these examples requires a production,

$$VP \rightarrow V \mid VP \text{ RB} \mid VP \text{ PP} \mid V \text{ NP} \mid V \text{ NP NP} \quad (10.17)$$

But what about **She sleeps sushi* or **She speaks John Japanese*? We need a more fine-grained verb non-terminal to handle these cases.

$$VP \rightarrow VP \text{ RB} \mid VP \text{ PP} \quad (10.18)$$

$$VP \rightarrow V\text{-INTRANS} \mid V\text{-TRANS NP} \mid V\text{-DITRANS NP NP} \quad (10.19)$$

$$V\text{-INTRANS} \rightarrow \textit{sleeps} \mid \textit{talks} \mid \textit{eats} \mid \dots \quad (10.20)$$

$$V\text{-TRANS} \rightarrow \textit{eats} \mid \textit{knows} \mid \textit{gives} \mid \dots \quad (10.21)$$

$$V\text{-DITRANS} \rightarrow \textit{gives} \mid \textit{tells} \mid \dots \quad (10.22)$$

Notice that many verbs can be produced by multiple non-terminals: because we could have *Mary eats* and *Mary eats sushi*, we have to be able to derive *eats* from both V-INTRANS and V-TRANS.

To complete this fragment, we would also need to handle modal and auxiliary verbs that create complex tenses, like *She will have eaten sushi* but not **She will have eats sushi*.

Sentences

We can now define the part of the grammar that deals with entire sentences. Perhaps the simplest type of sentence includes a subject and a predicate,

$$(10.53) \quad \textit{She eats sushi}$$

To handle this we simply need,

$$S \rightarrow NP \text{ VP}. \quad (10.23)$$

This rule can handle a number of other examples, like *she gives Alice the sushi*, *she eats*, etc. But things get more complex when we consider that sentences can be embedded inside other sentences:

$$(10.54) \quad \textit{Sometimes, she eats sushi}$$

$$(10.55) \quad \textit{In Japan, she eats sushi}$$

We therefore add two more productions,

$$S \rightarrow \text{ADVP} \text{ S} \quad (10.24)$$

$$S \rightarrow \text{PP} \text{ S} \quad (10.25)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

What about **I eats sushi*, **She eat sushi*? To handle these, we need additional productions that enforce subject-verb agreement:

$$S \rightarrow NP.3S \ VP.3S \mid NP.N3S \ VP.N3S$$

In some languages, there are many other forms of agreement. **Feature grammars** provide a notation that can capture this kind of agreement, while remaining in the context-free class of languages.

Coordination

As mentioned above, one test for constituency is whether constituents of the same proposed type can be **coordinated** using words like *and* and *or*. For example,

(10.56) *She eats (sushi) and (candy)*

(10.57) *She (eats sushi) and (drinks soda)*

(10.58) *(She eats sushi) and (he drinks soda)*

(10.59) *(fresh) and (tasty) sushi*

These examples motivate, respectively, the following productions,

$$NP \rightarrow NP \ Cc \ NP \quad (10.26)$$

$$VP \rightarrow VP \ Cc \ VP \quad (10.27)$$

$$S \rightarrow S \ Cc \ S \quad (10.28)$$

$$ADJP \rightarrow ADJP \ Cc \ ADJP \quad (10.29)$$

$$Cc \rightarrow and \mid or \mid \dots \quad (10.30)$$

We would need a little more cleverness to properly cover coordinations of more than two elements.

Odds and ends

Consider the example,

(10.60) *I gave sushi to the girl **who eats sushi**.*

This is a relative clause, which we already hinted at in the section on noun phrases. It requires its own non-terminal.

$$RELCLAUSE \rightarrow WP \ VP \quad (10.31)$$

$$WP \rightarrow who \mid that \mid which \mid \dots \quad (10.32)$$

Here are some related examples:

(c) Jacob Eisenstein 2014-2017. Work in progress.

(10.61) *I took sushi from the man **offering** sushi.*

(10.62) *I gave sushi to the woman **working at home**.*

This is a gerundive postmodifier, which again requires its own non-terminal.

$$\text{NOM} \rightarrow \text{NOM GERUNDVP} \quad (10.33)$$

$$\text{GERUNDVP} \rightarrow \text{VBG} \mid \text{VBG NP} \mid \text{VBG PP} \mid \dots \quad (10.34)$$

$$\text{VBG} \rightarrow \text{offering} \mid \text{working} \mid \text{talking} \mid \dots \quad (10.35)$$

Finally, we need to deal with questions, such as *can she eat sushi?* (and notice it's not *can she **eats** sushi*).

$$\text{S} \rightarrow \text{AUX NP VP} \quad (10.36)$$

$$\text{AUX} \rightarrow \text{can} \mid \text{did} \mid \dots \quad (10.37)$$

Clearly this is just a small fragment of all the productions and non-terminals we would need to generate all observed English sentences. And as we will see, even this grammar fragment suffers from significant ambiguity. It is this issue that we will tackle in chapter 11.

10.5 Grammar equivalence and normal form

There may be many grammars that express the same context-free language.

- Grammars are **weakly equivalent** if they generate the same strings.
- Grammars are **strongly equivalent** if they generate the same strings **and** assign the same phrase structure to each string.

In Chomsky Normal Form (CNF), all productions are either:

$$A \rightarrow BC$$

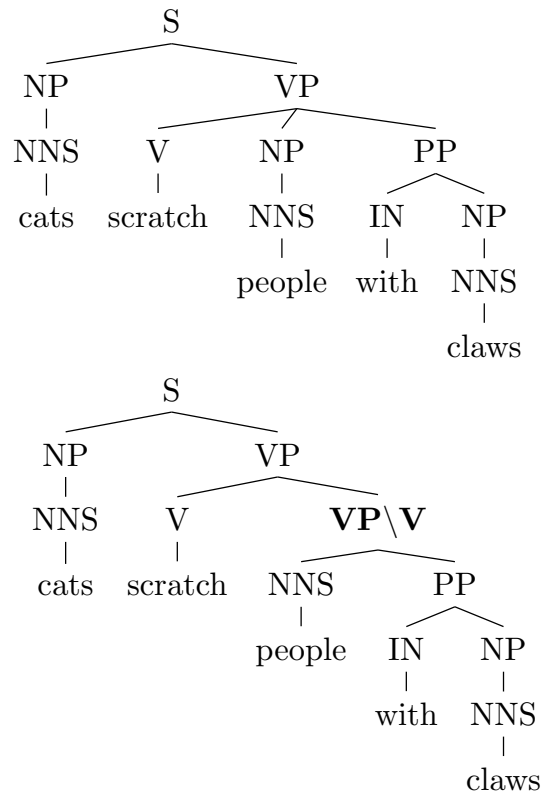
$$A \rightarrow a$$

All CFGs can be converted into a CNF grammar that is weakly equivalent — meaning that it generates exactly the same set of strings. As we will soon see, this conversion is very useful for parsing algorithms.

In CNF, all productions have either two or zero non-terminals on the right-hand side. To deal with productions that have more than two non-terminals on the RHS, we create new “dummy” non-terminals. For example, if we have the production,

$$W \rightarrow X Y Z, \quad (10.38)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

Figure 10.2: Binarization of the $VP \rightarrow V NP PP$ production

we can replace it with two productions,

$$W \rightarrow X W \backslash X \quad (10.39)$$

$$W \backslash X \rightarrow Y Z. \quad (10.40)$$

In these productions, $W \backslash X$ is a new dummy non-terminal, indicating a phrase that would be W if its left neighbor is an X . Figure 10.2 conveys this idea in a real example, with the non-terminal $VP \backslash V$ indicating a verb phrase that requires its left neighbor to be a verb.

Note that *people with claws* was not a constituent in the original grammar, but it is a constituent in the binarized grammar. Therefore, after parsing it is important to take care to “un-binarize” the resulting parse.

What about unary productions, such as $NP \rightarrow NNS$? While we could easily deal with this in the grammar, as we will see, in practice it is best dealt with by modifying the parsing algorithm itself.

Chapter 11

Context-free Parsing

Parsing is the task of determining whether a string can be produced by a given context-free grammar, and if so, how. The “how” question involves obtaining a hierarchical structure, as shown in Figure 10.1 in the previous chapter. Before we discuss specific parsing algorithms, let us consider whether exhaustive search is possible. Suppose we only have one non-terminal, X , and it has the following productions:

$$\begin{aligned} X &\rightarrow X X \\ X &\rightarrow \text{aardvark} \mid \text{abacus} \mid \dots \mid \text{zyther} \end{aligned}$$

In this grammar, the number of possible derivations for each string is equal to the number of binary bracketings, which is a **Catalan number**. Catalan numbers grow **super-exponentially** in the length of the sentence, $C_n = \frac{(2n)!}{(n+1)!n!}$. Clearly we cannot search the space of possible derivations naïvely; as with sequence labeling, we will make independence assumptions that allow us to search efficiently by reusing shared substructures with dynamic programming. This chapter will focus on a bottom-up algorithm called **CKY**; chapter 12 will describe a left-to-right algorithm called **shift-reduce**.

11.1 Deterministic bottom-up parsing

The CKY algorithm¹ is a bottom-up approach to parsing in a context free grammar. It efficiently tests whether a string is in a language, without considering all possible parses. The algorithm first forms small constituents, and then tries to merge them into larger constituents.

¹The name is for Cocke-Kasami-Younger, the inventors of the algorithm. It is sometimes called **chart parsing**, because of its chart-like data structure.

Let's start with a simple example grammar:

$$\begin{aligned} S &\rightarrow NP VP \\ NP &\rightarrow NP PP \mid we \mid sushi \mid chopsticks \\ PP &\rightarrow P NP \\ P &\rightarrow with \\ VP &\rightarrow V NP \mid V PP \\ V &\rightarrow eat \end{aligned}$$

Suppose we encounter the sentence *we eat sushi with chopsticks*.

- The first thing that we notice is that we can apply unary terminal productions to obtain the part-of-speech sequence NP V NP P NP.
- Next, we can apply a binary production to merge the first NP VP into an S.
- Or we could merge VP NP into VP ...
- ... and so on.

The CKY algorithm systematizes this approach, incrementally constructing a table t in which each cell $t[i, j]$ contains the set of nonterminals that can derive the span $w_{i:j-1}$. The algorithm fills in the upper right triangle of the table; it begins with the diagonal, which corresponds to substrings of length 1, and then computes derivations for progressively larger substrings, until reaching the upper right corner $t[0, M]$, which corresponds to the entire input. If the start symbol S is in $t[0, M]$, then the string w is in the language defined by the grammar.

- We begin by filling in the diagonal: the entries $t[m, m+1]$ for all $m \in \{0 \dots M-1\}$. These are filled with terminal productions that yield the individual tokens; for the word $w_2 = sushi$, we fill in $t[2, 3] = \{NP\}$, and so on.
- Then we fill in the next diagonal, in which each cell corresponds to a subsequence of length two: $t[0, 2], t[1, 3], \dots, t[M-2, M]$. These are filled in by looking for binary productions capable of producing at least one entry in each of the cells corresponding to left and right children. For example, the cell $t[1, 3]$ includes VP because the grammar includes the production $VP \rightarrow V NP$, and we have $V \in t[1, 2]$ and $NP \in t[2, 3]$.
- When we move to the next diagonal, there is an additional decision to make: where to split the left and right children. The cell $t[i, j]$ corresponds to the subsequence $w_{i:j-1}$, and we must choose some **split point** $i < k < j$, so that $w_{i:k-1}$ is the left child and $w_{k:j-1}$ is the right child. We do this by looping over all possible k , and then looking for productions that generate elements in $t[i, k]$ and $t[k, j]$; the left-hand side of all such productions can be added to $t[i, j]$. When it is time to compute $t[i, j]$,

(c) Jacob Eisenstein 2014-2017. Work in progress.

Algorithm 7 The CKY algorithm for parsing with context-free grammars

```

1: for  $m \in \{0 \dots M - 1\}$  do
2:    $t[m, m + 1] \leftarrow \{X : X \rightarrow w_m \in R\}$ 
3:   for  $n \in \{m + 1 \dots M\}$  do
4:      $t[m, n] \leftarrow \emptyset$ 
5:   for  $\ell \in \{2 \dots M\}$  do
6:     for  $m \in \{0 \dots M - \ell\}$  do
7:       for  $k \in \{m + 1 \dots m + \ell - 1\}$  do
8:          $t[m, m + \ell] \leftarrow t[m, m + \ell] \cup \{X : (X \rightarrow Y Z) \in R \wedge Y \in t[m, k] \wedge Z \in t[k, m + \ell]\}$ 
9:   if  $S \in t[0, M]$  then
10:    return True
11:  else
12:    return False

```

the cells $t[i, k]$ and $t[k, j]$ have already been filled in, since these cells correspond to shorter sub-strings of the input.

- The process continues until we reach $t[0, M]$.

Algorithm 7 further formalizes this process, and Figure 11.1 shows the chart that arises from parsing the sentence *we eat sushi with chopsticks* using the grammar defined above.

An important detail about the CKY algorithm is that it assumes that all productions with non-terminals on the right-hand side (RHS) are binary. What do we do when this is not true?

- For productions with more than two elements on the right-hand side, we **binarize**, creating additional non-terminals (see § 10.5). For example, if we have the production $VP \rightarrow V NP NP$ (for ditransitive verbs), we might convert to $VP \rightarrow VP_{ditrans}/NP NP$, and then add the production $VP_{ditrans}/NP \rightarrow V NP$.
- What about unary productions like $VP \rightarrow V$? In practice, this is handled by making a second pass on each diagonal, in which each cell $t[i, j]$ is augmented with all possible unary productions capable of generating each item already in the cell. Suppose our example grammar is extended to include the production $VP \rightarrow V$. Then the cell $t[1, 2]$ — corresponding to the word *eat* — would first include the set $\{V\}$, and would be augmented to the set $\{V, VP\}$ during this second pass. This would then make it possible to parse sentences like *We eat*.

Computing the parse tree We are usually interested not only in whether a sentence is in a grammar, but in what syntactic structure is revealed by parsing. As with the Viterbi algorithm, we can compute this structure by keeping a set of back-pointers while populating the CKY table. If we add an entry X to cell $t[i, j]$ by using the production $X \rightarrow YZ$ and




Figure 11.1: An example completed CKY chart. There are two paths to VP in position $t[1, 5]$, one in solid black and another in dashed blue.

the split point k , then we keep back-pointers $(i, j, X) \rightarrow (i, k, Y)$ and $(i, j, X) \rightarrow (k, j, Z)$. Once the table is constructed, we select back-pointers from $(0, M, S)$, and recursively follow them until they ground out at individual words.

For ambiguous sentences, there will be multiple paths to reach $S \in t[0, M]$. For example, in Figure 11.1, we reach $t[0, M]$ through a production that includes $VP_{int}[1, 5]$. Parsing is ambiguous for this example because there are two different ways to reach $VP \in t[1, 5]$: one with $(eat\ sushi)$ and $(with\ chopsticks)$ as children, and another with (eat) and $(sushi\ with\ chopsticks)$ as children. The presence of multiple paths indicates that the input could have been generated by the grammar in more than one way. In § 11.2, we consider different ways to resolve this ambiguity.

Complexity The space complexity of the CKY algorithm is $\mathcal{O}(M^2 \#|N|)$. We are building a table of size M^2 , and each cell must hold up to $\#|N|$ elements, where $\#|N|$ is the number of non-terminals. The time complexity is $\mathcal{O}(M^3 \#|R|)$. At each cell, we search over $\mathcal{O}(M)$ split points, and $\#|R|$ productions, where $\#|R|$ is the number of production rules in the grammar. Notice that these are considerably worse than the finite-state algorithms of Viterbi and forward-backward, which are linear time; generic shortest-path for finite-state automata has complexity $\mathcal{O}(M \log M)$.

11.2 Ambiguity in parsing

Unfortunately, parsing ambiguity is endemic to natural language:

- **Attachment ambiguity:** *we eat sushi with chopsticks, I shot an elephant in my pajamas.* In each of these examples, the prepositions (*with, in*) can attach to either the verb or the direct object.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- **Modifier scope:** *southern food store, plastic cup holder*. In these examples, the first word could be modifying the subsequent adjective, or the final noun.
- **Particle versus preposition:** *The puppy tore up the staircase*. Phrasal verbs like *tore up* often include particles which could also act as prepositions.
- **Complement structure:** *The tourists objected to the guide that they couldn't hear*. This is another form of attachment ambiguity, where the complement *that they couldn't hear* could attach to the main verb (*objected*), or to the indirect object (*the guide*).
- **Coordination scope:** *"I see," said the blind man, as he picked up the hammer and saw*. In this example, the lexical ambiguity for *saw* enables it to be coordinated either with the noun *hammer* or the verb *picked up*.

These forms of ambiguity can combine, so that a seemingly simple headlines like *Fed raises interest rates* can have dozens of possible analyses, even in a minimal grammar. Broad coverage grammars permit millions of parses of typical sentences. Faced with this ambiguity, classical deterministic parsers faced a tradeoff:

- achieve broad coverage but admit a huge amount of ambiguity;
- or settle for limited coverage in exchange for constraints on ambiguity.

Rather than attempting to design a grammar that achieves broad coverage and low ambiguity, contemporary methods use labeled data to learn models capable of selecting the correct syntactic analysis.

Parser evaluation

Before continuing to parsing algorithms that are able to handle ambiguity, we need to consider how to measure parsing performance. Suppose we have a set of **reference parses** — the ground truth — and a set of **system parses** that we would like to score. A simple solution would be **per-sentence accuracy**: the parser is scored by the proportion of sentences on which the system and reference parses exactly match.² But we would like to assign *partial credit* for correctly matching parts of the reference parse. The PARSEval metrics (Grishman et al., 1992) do that, scoring each system parse via:

Precision, the fraction of brackets in the system parse that match a bracket in the reference parse.

Recall, the fraction of brackets in the reference parse that match a bracket in the system parse.

²Most parsing papers do not report results on this metric, but Finkel et al. (2008) find that a near-state-of-the-art parser finds the exact correct parse on 35% of sentences of length ≤ 40 , and on 62% of parses of length ≤ 15 in the Penn Treebank.




Figure 11.2: Suppose that the left parse is the system output, and the right parse is the ground truth; the precision is 0.75 and the recall is 1.0.

As in chapter 3, the F-measure is the harmonic mean of precision and recall, $F = \frac{2*P*R}{R+P}$.

In **labeled** precision and recall, the system must also match the non-terminals for each bracket; in **unlabeled** precision and recall, it is only required to match the bracketing structure.

In Figure 11.2, suppose that the left tree is the system parse and the right tree is the reference parse. We have the following spans:

- $S \rightarrow w_{0:5}$ is **true positive**, because it appears in both trees.
- $VP \rightarrow w_{1:5}$ is **true positive** as well.
- $NP \rightarrow w_{2:5}$ is **false positive**, because it appears only in the system output.
- $PP \rightarrow w_{3:5}$ is **true positive**, because it appears in both trees.

So for this parse, we have a (labeled and unlabeled) precision of $\frac{3}{4} = 0.75$, and a recall of $\frac{3}{3} = 1.0$, for an F-measure of 0.86.

Local solutions

Some ambiguity can be resolved locally. Consider the following examples,

(11.1) [*imposed* [*a ban* [*on asbestos*]]]

(11.2) [*imposed* [*a ban*]] [*on asbestos*]]

This is a case of attachment ambiguity: do we attach the prepositional phrase *on asbestos* to the verb *imposed*, or the noun phrase *a ban*? To solve this problem, Hindle and Rooth (1990) proposed a likelihood ratio test:

$$LR(v, n, p) = \frac{p(p | v)}{p(p | n)} = \frac{p(\text{on} | \text{imposed})}{p(\text{on} | \text{ban})} \quad (11.1)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

where they select VERB attachment if $LR(v, n, p) > 1$. The probabilities are estimated from annotated training data.

This approach is capable of modeling which prepositions tend to attach to which verbs and nouns. However, it ignores any information about the object of the prepositional phrase, which might also factor into this decision. For example:

(11.3) ...[it [would end [its venture [with Maserati]]]]

(11.4) ...[it [would end [its venture][with Maserati]]]

The first analysis attaches *with Maserati* to the *venture*, which is almost surely preferred. Yet the likelihood ratio test ignores *Maserati*, and prefers to link the preposition *with* to the verb *end* (as in *It will end with a bang*).

- $p(\textit{with} \mid \textit{end}) = \frac{607}{5156} = 0.118$
- $p(\textit{with} \mid \textit{venture}) = \frac{155}{1442} = 0.107$

A richer probabilistic approach is undertaken by Collins and Brooks (1995), who model attachment as depending on four **heads**:

- the preposition (e.g., *with*)
- the VP attachment site (e.g., *end*)
- the NP attachment site (e.g., *venture*)
- the NP to be attached (e.g., *Maserati*)

They propose a backoff-based approach:

- First, look for counts of the tuple $\langle \textit{with}, \textit{Maserati}, \textit{end}, \textit{venture} \rangle$, and see how often VP and NP attachment were preferred for this tuple.
- If there are no counts for the full tuple, back off to the triples, $\langle \textit{with}, \textit{Maserati}, \textit{end} \rangle + \langle \textit{with}, \textit{end}, \textit{venture} \rangle + \langle \textit{with}, \textit{Maserati}, \textit{venture} \rangle$, and count how often VP and NP attachment were preferred in each case.
- If there are no counts even for these triples, then try $\langle \textit{with}, \textit{Maserati} \rangle + \langle \textit{with}, \textit{end} \rangle + \langle \textit{with}, \textit{venture} \rangle$.
- Finally, if there are no counts for even these tuples, simply compute how often the preposition preferred NP or VP attachment. Since prepositions are a closed class, we can expect to have sufficient data for each preposition.

Accuracy of this method is roughly 84%. This approach of combining relative frequency estimation, smoothing, and backoff was very characteristic of 1990s statistical natural language processing. More conventional classification-based techniques can also be used for this problem; for example, Ratnaparkhi et al. (1994) designed a set of features and then trained a logistic regression classifier.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Beyond local solutions

Framing the problem as attachment ambiguity is limiting. It assumes the parse is mostly done, leaving just a few attachment ambiguities to solve. But realistic sentences have more than a few syntactic interpretations, and attachment decisions are interdependent. For example, consider the garden-path sentence,

(11.5) *Cats scratch people with claws with knives.*

We may want to attach *with claws* to *scratch*, as would be correct in the shorter sentence in *cats scratch people with claws*. But then we have nowhere to attach *with knives*. Only by considering these decisions jointly can we make the right choice. The task of statistical parsing is to produce a single analysis that resolves all syntactic ambiguities.

11.3 Weighted Context-Free Grammars

In a **weighted context-free grammar** (WCFG), each production $X \rightarrow \alpha$ is associated with a score $\psi_{X \rightarrow \alpha}$. The score of a derivation is simply the combination (sum or product) of the scores of all the productions. For any given string w , the “best” parse is the one corresponding to the highest-scoring derivation. More formally, for a given sequence w , we want to select the parse τ that maximizes the score,

$$\hat{\tau} = \operatorname{argmax}_{\tau: \text{yield}(\tau)=w} \sum_{(X \rightarrow \alpha) \in \tau} \psi_{X \rightarrow \alpha},$$

where we model a derivation τ as a set of productions of the form $X \rightarrow \alpha$. As in CFGs, the **yield** of a tree is the string of terminal symbols that can be read off the leaf nodes. The set $\{\tau : w = \text{yield}(\tau)\}$ is exactly the set of all derivations of w in a CFG G .

Probabilistic context-free grammars

An important special case of WCFGs is a **probabilistic context-free grammar** (PCFG), in which the weight for each production $X \rightarrow \alpha$ corresponds to a log-probability, $\psi_{X \rightarrow \alpha} = \log p(\alpha \mid X)$. These probabilities are conditioned on the left-hand side, so they must normalize to one over possible right-hand sides, $\sum_{\alpha'} p(\alpha' \mid X) = 1$. For example, for the verb phrase productions, we might have,

VP \rightarrow V	0.3
VP \rightarrow V NP	0.6
VP \rightarrow V NP NP	0.1

which would indicate that transitive verbs are twice as common as intransitive verbs, which in turn are three times more common than ditransitive verbs.

(c) Jacob Eisenstein 2014-2017. Work in progress.

S	→ NP VP	0.9
S	→ S Cc S	0.1
NP	→ N	0.2
NP	→ Dt N	0.3
NP	→ N NP	0.2
NP	→ Jj NP	0.2
NP	→ NP PP	0.1
VP	→ V	0.4
VP	→ V NP	0.3
VP	→ V NP NP	0.1
VP	→ VP PP	0.2
PP	→ P NP	1.0

Table 11.1: A fragment of an example probabilistic context-free grammar (PCFG)

Given probabilities on the productions, we can then score the probability of a derivation as a **product** of the probabilities of all of the productions. Consider the PCFG in Table 11.1 and the parse in Figure 11.3. The probability of this parse is:³

$$\begin{aligned}
p(\tau, w) = & p(S \rightarrow NP VP) \\
& \times p(NP \rightarrow N) \times p(N \rightarrow \textit{they}) \\
& \times p(VP \rightarrow VP PP) \\
& \times p(VP \rightarrow V NP) \times p(V \rightarrow \textit{eat}) \\
& \times p(NP \rightarrow N) \times p(N \rightarrow \textit{sushi}) \\
& \times p(PP \rightarrow P NP) \times p(P \rightarrow \textit{with}) \\
& \times p(NP \rightarrow N) \times p(N \rightarrow \textit{chopsticks}) \tag{11.2}
\end{aligned}$$

$$\begin{aligned}
= & 0.9 \times 0.2 \times 0.2 \times 0.3 \times 0.2 \times 1.0 \times 0.2 \\
& \times \text{probability of terminal productions} \tag{11.3}
\end{aligned}$$

Now if we consider the alternative parse in which the prepositional phrase attaches to the noun, all of these probabilities are the same, with one exception: instead of the production $VP \rightarrow VP PP$, we would have the production $NP \rightarrow NP PP$. Since $p(VP \rightarrow VP PP) > p(NP \rightarrow NP PP)$ in the PCFG, the verb phrase attachment would be preferred.⁴

³In the remainder of the chapter, we will use the notation $p(X \rightarrow YZ)$ for the probability of producing Y and Z , conditioned on the left-hand side being X .

⁴This example hints at a big problem with PCFG parsing on non-terminals such as NP, VP, and PP: we will **always** prefer either VP or PP attachment, without regard to what is being attached! This problem is addressed later in the chapter.




Figure 11.3: An example derivation

Feature-based parsing

The scores for each production can also be computed as an inner product of weights and features,

$$\psi_{X \rightarrow \alpha} = \theta \cdot \mathbf{f}(X, \alpha), \quad (11.4)$$

where the feature vector $\mathbf{f}(X, \alpha, \mathbf{w})$ is a function of the left-hand side X and the right-hand side α . More generally, we can estimate weights for productions covering specific parts of the input,

$$\psi_{X \rightarrow \alpha, i, j, k} = \theta \cdot \mathbf{f}(X, \alpha, \mathbf{w}, i, j, k), \quad (11.5)$$

where the feature vector is now a function of the details of the production (X and α), as well as the text \mathbf{w} and the indices of the spans to derive: the parent $\mathbf{w}_{i:j-1}$, the left child $\mathbf{w}_{i:k-1}$, and the right child $\mathbf{w}_{k:j-1}$. This is equivalent to characterizing the entire parse τ in terms of a locally-decomposable feature vector,

$$\mathbf{f}(\tau, \mathbf{w}) = \sum_{(X \rightarrow \alpha, i, j, k) \in \tau} \mathbf{f}(X, \alpha, \mathbf{w}, i, j, k) \quad (11.6)$$

$$\hat{\tau} = \operatorname{argmax}_{\tau} \theta \cdot \mathbf{f}(\tau, \mathbf{w}) \quad (11.7)$$

$$= \operatorname{argmax}_{\tau} \sum_{(X \rightarrow \alpha, i, j, k) \in \tau} \theta \cdot \mathbf{f}(X, \alpha, \mathbf{w}, i, j, k). \quad (11.8)$$

This enables the use of richer features, such as the words that border the span $\mathbf{w}_{i:j-1}$, the specific word at the split point w_k , the presence of a verb or noun in the left child span $\mathbf{w}_{i:j-1}$, etc. The use of such features does not affect the applicability of the CKY parsing algorithm: we can still compute each element of the table $t[i, j]$ recursively, and therefore we can still find the best parse in polynomial time. The only restriction is that the features for each production $X \rightarrow \alpha$ cannot consider other non-terminals besides the parent X

and the children α . This is analogous to the Viterbi restriction that features consider only adjacent tags.

Estimation

Probabilistic context free grammars are similar to hidden Markov models, in that they are generative models of text. The parameters in hidden Markov models can be estimated from labeled data by relative frequency, and the same approach can be applied in PCFGs. In this case, the parameters of interest correspond to probabilities of productions, conditional on the left hand side. The relative frequency estimate is therefore,

$$p(X \rightarrow \alpha) = \frac{\text{count}(X \rightarrow \alpha)}{\text{count}(X)}. \quad (11.9)$$

For example, the probability of the production $\text{NP} \rightarrow \text{DT NN}$ is equal to the count of this production divided by the count of the non-terminal NP. This applies to terminal productions as well: the probability of $\text{NN} \rightarrow \text{centipede}$ is the count of how often *centipede* appears in the corpus as generated from an NN tag, divided by the total count of the NN tag. These counts can be obtained from an annotated dataset, such as the **Penn Treebank**, which includes syntactic annotations over one million words of English text (Marcus et al., 1993). Even with one million words, it will be difficult to compute probabilities of relatively rare events, such as $\text{NN} \rightarrow \text{centipede}$. Therefore, smoothing techniques will again be critical for making PCFGs effective.

Feature-based parsing models can be estimated using either structure perceptron or maximum conditional likelihood. For structure perceptron, we would compute,

$$\hat{\tau} = \underset{\tau: \text{yield}(\tau) = \mathbf{w}^{(i)}}{\text{argmax}} \quad \boldsymbol{\theta} \cdot \mathbf{f}(\tau, \mathbf{w}^{(i)}) \quad (11.10)$$

$$\boldsymbol{\theta} \leftarrow \mathbf{f}(\tau^{(i)}, \mathbf{w}^{(i)}) - \mathbf{f}(\hat{\tau}, \mathbf{w}^{(i)}). \quad (11.11)$$

Alternatively, we can estimate the weights $\boldsymbol{\theta}$ by maximizing the conditional log-likelihood, $\sum_{i=1}^N \log p(\tau^{(i)} \mid \mathbf{w}^{(i)})$, which is analogous to the conditional random field (CRF) model for sequence labeling. Finkel et al. (2008) present a CRF-based parsing model. Carreras et al. (2008) present a structure perceptron model, although they perform parsing in alternative syntactic formalism known as **Tree-Adjoining Grammar** (Joshi and Schabes, 1997). § 11.5 gives more details on these discriminative parsing models.

Parsing with weighted context-free grammars

It is not difficult to extend the CKY algorithm to include probabilities or other weights. Let us write $\psi_{X \rightarrow Y Z}$ for the score for the production $X \rightarrow Y Z$. In the original CKY algorithm for deterministic parsing, each cell $t[i, j]$ stored a set of non-terminals capable of deriving

Algorithm 8 CKY algorithm with weighted productions

```

for  $m \in \{0, \dots, M-1\}$  do
  for all  $X \in \text{tags}(w_j)$  do
     $t[m, m+1, X] \leftarrow P(X \rightarrow w_m)$ 
  for  $\ell \in \{2 \dots M\}$  do
    for  $m \in \{0, \dots, M-\ell\}$  do
      for  $k \in \{m+1, \dots, m+\ell-1\}$  do
        for all  $(X \rightarrow YZ) \in R$  do
           $t[m, m+\ell, X] \leftarrow t[m, m+\ell, X] + \max_{k, X \rightarrow YZ} \psi_{X \rightarrow YZ} + t[m, k, Y] + t[k, m+\ell, Z]$ 

```

the span $w_{i:j-1}$. We now augment the table to be indexed by the tuple (i, j, X) , with $t[i, j, X]$ indicating the score of the best possible derivation of $w_{i:j-1}$ from non-terminal X . Algorithm 8 shows how to perform CKY parsing using such a table.

We also keep the back-pointers corresponding to the best path to $t[i, j, X]$; the best-scoring derivation can be obtained by tracing these pointers from $t[0, M, S]$ back to each terminal, just as the best sequence of labels in the Viterbi algorithm can be computed by tracing pointers backwards from the end of the trellis. Note that we need only store back-pointers for the **best** path to $t[i, j, X]$; this follows from the locality assumption that the score for a parse is a combination of the scores of each production in the parse.

Semiring CKY As with hidden Markov models, we can generalize weighted CKY parsing using semiring notation. The basic recurrence becomes,

$$t[m, m+\ell, X] = t[m, m+\ell, X] \otimes \left(\bigoplus_{k, X \rightarrow YZ} \psi_{X \rightarrow YZ} \otimes t[m, k, Y] \otimes t[k, m+\ell, Z] \right). \quad (11.12)$$

This notation makes it possible to capture a number of different parsing algorithms compactly. If the scores ψ correspond to log-probabilities or feature-weight inner products, then \otimes is addition, as in Algorithm 8. If ψ are probabilities, then \otimes is multiplication. By setting \oplus equal to the max operation, the CKY algorithm computes the score of the best parse for a given sentence.

If the scores ψ correspond to probabilities, we can set \oplus to summation and \otimes to multiplication. Then the value at $t[0, M]$ corresponds to the total probability of all derivations of the input. This is known as the **inside algorithm**, and it is the tree-structured version of the **forward algorithm** from § 6.4. The inside algorithm can be used for unsupervised or semi-supervised parsing (Pereira and Schabes, 1992).

Finally, if we set \otimes to be the boolean “and” operation, and \oplus to be the boolean “or” operation, then $t[0, M] = \text{True}$ if and only if there is at least one derivation of the in-

put from the grammar. Thus, the semiring notation generalizes across the weighted and unweighted CKY algorithms.

11.4 Improving Parsing by Refined Non-terminals

PCFG parsing on the Penn Treebank dataset does not perform well: Johnson (1998) shows that a PCFG estimated from treebank production counts obtains an F-measure of only $F = 0.72$. There are several problems with the use of weighted context free grammars on the Penn Treebank dataset:

- One problem is that the context-free assumption is too strict: for example, the probability of the production $NP \rightarrow NP PP$ is much higher if the parent of the noun phrase is a verb phrase (indicating that the NP is a direct object) than if the parent is a sentence (indicating that the NP is the subject of the sentence). Accurately modeling this “vertical” context is essential for accurate parsing.
- Another problem is that the Penn Treebank non-terminals are simply too coarse: there are many kinds of noun phrases and verb phrases, and accurate parsing sometimes requires knowing the difference. As we have already seen, when faced with prepositional phrase attachment ambiguity, a weighted CFG will either always choose NP attachment (if $\psi_{NP \rightarrow NP PP} > \psi_{VP \rightarrow VP PP}$), or it will always choose VP attachment. To get more nuanced behavior, more fine-grained non-terminals are needed.
- More generally, accurate parsing requires some amount of **semantics** — understanding the meaning of the text to be parsed. Consider the example *cats scratch people with claws*: knowledge of about *cats*, *claws*, and scratching is necessary to correctly resolve the attachment ambiguity.
- As a more extreme case, consider the example shown in Figure 11.4. The analysis on the left is preferred because of the conjunction of similar entities *France* and *Italy*. But given the non-terminals shown in the analyses, there is no way to differentiate these two parses, since they include exactly the same productions.

In all cases, what is needed seems to be more precise non-terminals. One possibility would be to rethink the linguistics behind the Penn Treebank, and ask the annotators to try again. But the original annotation effort took five years, and a more fine-grained set of non-terminals would only make things worse. We will therefore focus on automated techniques.

Parent annotations and other tree transformations

The key assumption underlying weighted context-free parsing is that productions depend only on the identity of the non-terminal on the left-hand side, and not on its ancestors

(c) Jacob Eisenstein 2014-2017. Work in progress.

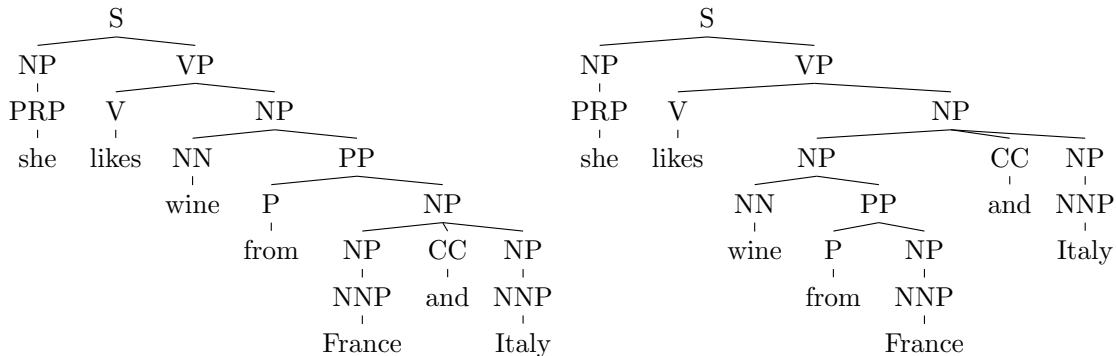


Figure 11.4: The left parse is preferable because of the conjunction of phrases headed by *France* and *Italy*.

in the parse. The validity of this assumption is an empirical question, and it depends on the non-terminals themselves: ideally, every noun phrase would be distributionally identical, so the assumption would hold. But in PTB-style analysis of English grammar, the observed probability of productions often depends on the parent of the left-hand side. For example, noun phrases are more likely to be modified by prepositional phrases when they are in the object position (e.g., *they amused the students from Georgia*) than in the subject position (e.g., *the students from Georgia amused them*). This means that the $\text{NP} \rightarrow \text{NP PP}$ production is more likely if the entire constituent is the child of a VP than if it is the child of S.

$$P(\text{NP} \rightarrow \text{NP PP}) = 11\% \quad (11.13)$$

$$P(\text{NP UNDER S} \rightarrow \text{NP PP}) = 9\% \quad (11.14)$$

$$P(\text{NP UNDER VP} \rightarrow \text{NP PP}) = 23\%. \quad (11.15)$$

Johnson (1998) proposes to capture this phenomenon via **parent annotation**. Each non-terminal is augmented with the identity of its parent, as shown in Figure 11.5). This is sometimes called **vertical Markovization**, since we introduce a Markov dependency between each node and its parent (Klein and Manning, 2003).

Using this transformation and a number of related heuristics, Johnson (1998) was able to improve the accuracy of PCFG-based parsing from 72% to 80%, at the cost of increasing the number of production rules from 14,962 to 22,773. (Recall that the number of production rules is a constant factor in the time complexity of WCFG parsing.) This increase in the number of rules is relatively modest, considering that parent annotation squares the number of non-terminals.

Parent annotation weakens the PCFG independence assumptions. This could improve accuracy by enabling the parser to make more fine-grained distinctions, which




Figure 11.5: Parent annotation in a CFG derivation

Non-terminal	Direction	Priority
S	right	VP SBAR ADJP UCP NP
VP	left	VBD VBN MD VBZ TO VB VP VBG VBP ADJP NP
NP	right	N* EX \$ CD QP PRP ...
PP	left	IN TO FW

Table 11.2: A fragment of head percolation rules

better capture real linguistic phenomena. However, each production is more rare (since the non-terminals are more specific), so the more careful smoothing is required to dampen the variance over production probabilities.

Lexicalization

Recall that some of the problems with PCFG parsing that were suggested above have to do with **meaning** — for example, preferring to coordinate constituents that are of the same type, like *cats* and *dogs* rather than *cats* and *houses*. A simple way to capture semantics is through the words themselves: we can annotate each non-terminal with **head** word of the phrase.

Head words are deterministically assigned according to a set of rules, sometimes called **head percolation rules**. In many cases, these rules are straightforward: the head of a $\text{NP} \rightarrow \text{DT N}$ production is the noun, the head of a $\text{S} \rightarrow \text{NP VP}$ production is the head of the VP, etc. A fragment of the head percolation rules used in many parsing systems are found in Table 11.2.⁵

The meaning of the first rule is that to find the head of an S constituent, we first look for the rightmost VP child; if we don't find one, we look for the rightmost SBAR child, and so on down the list. Verb phrases are headed by left verbs (the head of *can walk home* is *walk*, since the modal verb *can* is tagged MD), noun phrases are headed by the rightmost noun-like non-terminal (so the head of *the red cat* is *cat*),⁶ and prepositional phrases are headed

⁵From <http://www.cs.columbia.edu/~mcollins/papers/heads>

⁶The noun phrase non-terminal is sometimes treated as a special case. Collins (1997) uses a heuristic that

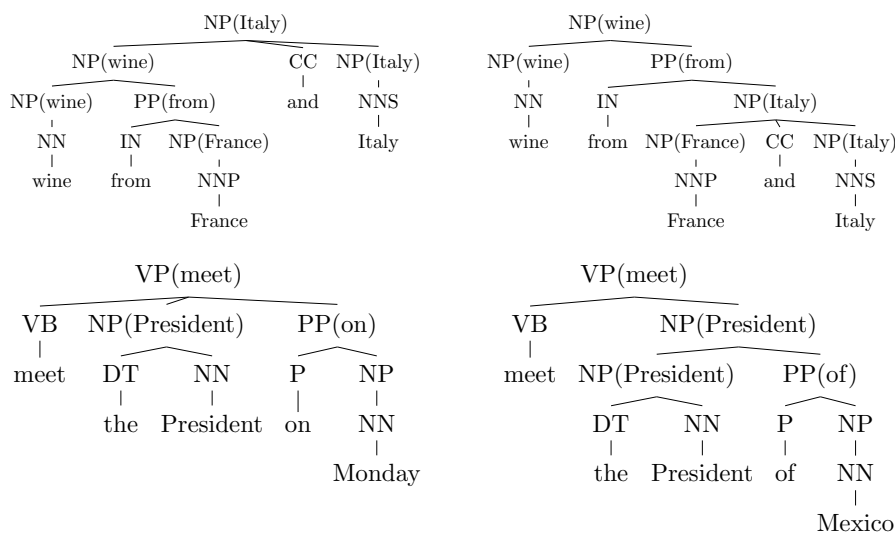


Figure 11.6: Lexicalization can address ambiguity on coordination scope (upper) and PP attachment (lower)

by the preposition (the head of *at Georgia Tech* is *at*). Some of these rules are somewhat arbitrary — there’s no particular reason why the head of *cats and dogs* should be *dogs* — but the point here is just to get some lexical information that can support parsing, not to make any deep claims about syntax.

Given these rules, we can lexicalize the parse trees for some of our examples, as shown in Figure 11.6.

- In the upper part of Figure 11.6, we see how lexicalization can help solve coordination scope ambiguity. We will correctly coordinate *France* and *Italy* if,

$$p(\text{NP}(\text{Italy}) \rightarrow \text{NP}(\text{France}) \text{ CC } \text{NP}(\text{Italy})) > p(\text{NP}(\text{Italy}) \rightarrow \text{NP}(\text{wine}) \text{ CC } \text{NP}(\text{Italy})). \quad (11.16)$$

- In the lower part of Figure 11.6, we see how lexicalization can help solve attachment ambiguity. Here we assume that,

$$p(\text{VP}(\text{meet}) \rightarrow \alpha \text{ PP}(\text{on})) \gg p(\text{NP}(\text{President}) \rightarrow \beta \text{ PP}(\text{on})) \quad (11.17)$$

$$p(\text{VP}(\text{meet}) \rightarrow \alpha \text{ PP}(\text{of})) \ll p(\text{NP}(\text{President}) \rightarrow \beta \text{ PP}(\text{of})) \quad (11.18)$$

In plain English: *meetings* are usually *on* things; *Presidents* are *of* things.

looks for the rightmost child which is a noun-like part-of-speech (e.g., *Nn*, *Nnp*), a possessive marker, or a superlative adjective (e.g., *the greatest*). If no such child is found, the heuristic then looks for the **leftmost** NP. If there is no child with tag NP, the heuristic then applies another priority list, this time from right to left.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Recall that verbs may be intransitive, transitive, or ditransitive. Lexicalization can help distinguish these cases, as shown by the lexicalized PCFG probabilities for the ditransitive VP production,

$$p(\text{VP} \rightarrow \text{V NP NP}) = 0.00151 \quad (11.19)$$

$$p(\text{VP}(\text{said}) \rightarrow \text{V}(\text{said}) \text{ NP NP}) = 0.00001 \quad (11.20)$$

$$p(\text{VP}(\text{gave}) \rightarrow \text{V}(\text{gave}) \text{ NP NP}) = 0.01980. \quad (11.21)$$

Overall, lexicalization had a major impact on parsing accuracy, which the best lexicalized parsers attaining accuracies in the range of 87-89% (Collins, 1997, 2003; Charniak, 1997). However, lexicalized parsing introduces significant technical challenges. First, the CKY parsing algorithm must keep track of the heads of each phrase, which adds algorithmic complexity. Second, the set of possible lexicalized productions is vastly larger, since it is quadratic in the size of the vocabulary (the left and right children each have a head, and one of these heads is chosen to head the unified phrase). We now briefly overview solutions to these problems.

Algorithms for lexicalized parsing

In weighted CFG parsing, the table element $t[i, j, X]$ keeps the score of the best derivation of the span $w_{i:j-1}$ from the non-terminal X . However, for this constituent to participate on the right-hand side of higher-level productions, we must also know its head. One solution is to expand the table to include cells of the form $t[i, j, h, X]$, where h is the index of the head of the non-terminal X for the span $w_{i:j-1}$, with $i \leq h < j$.

We can compute each element in the table by first computing the score of the best production in which the head comes from the left child, $t_\ell[i, j, h, X]$, then computing the score of the best production in which the head comes from the right child, $t_r[i, j, h, X]$, and finally taking the max over these two possibilities.

$$t_\ell[i, j, h, X] = \max_{X \rightarrow YZ} \max_{k > h} \max_{k \leq h' < j} t[i, k, h, Y] + t[k, j, h', Z] + \psi_{X(h) \rightarrow Y(h)Z(h')} \quad (11.22)$$

$$t_r[i, j, h, X] = \max_{X \rightarrow YZ} \max_{k \leq h} \max_{i \leq h' < j} t[i, k, h', Y] + t[k, j, h, Z] + \psi_{X(h) \rightarrow Y(h')Z(h)} \quad (11.23)$$

$$t[i, j, h, X] = \max(t_\ell[i, j, h, X], t_r[i, j, h, X]). \quad (11.24)$$

To compute t_ℓ , we maximize over all split points $k > h$, since the head word must be in the left child. We then maximize again over possible head words h' for the right child. An analogous computation is performed for t_r . The size of the table is now $\mathcal{O}(M^3 \#|N|)$, where M is the length of the input and $\#|N|$ is the number of non-terminals. Furthermore, each cell is computed by performing $\mathcal{O}(M^2)$ operations, since we maximize over both the split point k and the head h' . The time complexity of the algorithm is therefore $\mathcal{O}(M^5 \#|N|)$, which is impractical. Fortunately, the Eisner (1996) algorithm reduces

this complexity back to $\mathcal{O}(M^3)$, using a more complex algorithm that maintains multiple tables.

The Charniak Parser

We now approach the problem of how to estimate weights for lexicalized productions $X(i) \rightarrow Y(j) Z(k)$. These productions are said to be **bilexical**, because they involve scores over pairs of words: in the example *...meet the President of Mexico*, we hope to choose the right attachment point by modeling the bilexical affinities of *(meet, of)* and *(President, of)*. The number of such word pairs is of course quadratic in the size of the vocabulary, making it difficult to estimate them directly from data.

The Charniak (1997) parser addresses this issue in the context of probabilistic parsing, so that $\psi_{X(i) \rightarrow Y(j) Z(k)}$ is equal to the (log) probability of the lexicalized production. This probability is then decomposed into a product of: a **rule probability**, which is the probability of the unlexicalized production $X \rightarrow YZ$, conditioned on the head word and parent of X (the same idea as parent annotation); a **head probability**, which is the probability of the head of X conditioned on X , the parent of X , and the head of the parent of X .

Recall the example from Figure 11.6, focusing on the bottom right example, *...meet the President of Mexico*. In the case of the production $PP(of) \rightarrow P(of) NP(Mexico)$, the rule probability is $p_{\text{rule}}(PP \rightarrow P NP \mid PP, NP, of)$, since the parent is a noun phrase and the head word is *of*. The head probability is $p_{\text{head}}(of \mid PP, NP, President)$, since the parent is a noun phrase and the head of the parent is *President*. This captures the bilexical affinity between *President* and *of*, which is key to accurately parsing this example.

Even with this decomposition, it is necessary to smooth the rule and head probabilities to reduce the variance of the probability estimates. This is done by interpolating the full probabilities with simplified probabilities that condition on less information.

The Collins Parser

The Charniak parser focuses on lexical relationships between children and parents. Motivated by the linguistic theory of **lexicalized tree-adjoining grammar** (Joshi and Schabes, 1997), the Collins (2003) parser focuses on relationships between adjacent children of the same parent. We can write each production as,

$$X \rightarrow L_m L_{m-1} \dots L_1 H R_1 \dots R_{n-1} R_n,$$

where H is the child containing the head word, each L_i is a child element to the left of the head, and each R_j is a child element to the right of the head. In the Collins parser, these elements are generated probabilistically from the head outward. The outermost elements of L and R are special symbols, written \blacklozenge .

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 11.7: Example verb phrase for understanding the Collins parser

For example, consider the verb phrase *dumped sacks into the river*, shown in Figure 11.7. To model this rule, we would compute:

$$p(\text{VP}(\textit{dumped}, \text{VBD}) \rightarrow [\blacklozenge, \text{VBD}(\textit{dumped}, \text{VBD}), \text{NP}(\textit{sacks}, \text{NNS}), \text{PP}(\textit{into}, \text{P}), \blacklozenge]),$$

with each phrase augmented by its head word and the head word’s part of speech (e.g., VBD for the head word *dumped*).

This probability is computed through a generative model, in which the head is generated first (conditioned on the parent), and then each dependent is conditioned on the parent non-terminal and the head word. In this way, we do not directly estimate the full probability of a lexicalized production rule, but rather, we compute it from simpler probabilities involving the head and parent. Nonetheless, it is still necessary to smooth these probabilities by interpolating them with less expressive probability functions.

The Collins parser models bilexical dependencies between the head and its siblings. Bilexical probabilities require counts over pairs of words, a space of $\mathcal{O}(|\mathcal{V}|^2)$ events. It is this large event space that makes these probabilities difficult to estimate, necessitating smoothing. Is it worth it? Bikel (2004) evaluates the importance of bilexical probabilities to the performance of the Collins parser. He found that bilexical probabilities are rarely available — because most of the possible bilexical pairs in the test data are unobserved in the training data — but that bilexical probabilities are indeed active in 29% of the rules in the **top-scoring** parses. Still, bilexical probabilities play a relatively small role in accuracy: an equivalent parser which conditions on only a single head suffers only 0.3% decrease in F-measure. A completely unlexicalized parser performs considerably worse, indicating that some amount of lexicalization is still necessary for top performance.

Refinement grammars

Lexicalization improves on pure PCFG parsing by adding detailed information in the form of lexical heads. However, estimating the probabilities of lexicalized parsing rules is difficult, requiring additional independence assumptions and complex smoothing. Klein and Manning (2003) argue that the right level of linguistic detail is somewhere between treebank categories and individual words. For example:

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Some parts-of-speech and non-terminals are truly substitutable: for example, *cat*/N and *dog*/N.
- But others are not: for example, *on*/PP behaves differently from *of*/PP. This is an example of **subcategorization**.
- Similarly, the words *and* and *but* should be distinguished from other coordinating conjunctions.

Figure 11.8 shows an example of an error that is corrected through the introduction of a new NP-TMP subcategory for temporal noun phrases. Klein and Manning (2003) show how the introduction of a number of such categories can make unlexicalized PCFG parsing competitive with lexicalized methods.




Figure 11.8: State-splitting creates a new non-terminal called NP-TMP, for temporal noun phrases. This corrects the PCFG parsing error in (a), resulting in the correct parse in (b).

***Automated state-splitting** Klein and Manning (2003) use linguistic insight and error analysis to manually split PTB non-terminals so as to make parsing easier. Later work by Klein and his students automated this state-splitting process, by treating the “refined” non-terminals as latent variables. For example, we might split the noun phrase non-terminal into NP1, NP2, NP3, ..., without defining in advance what each refined non-terminal corresponds to.

Petrov et al. (2006) employ expectation-maximization to solve this problem. In the E-step, we estimate a marginal distribution q over the refinement type of each non-terminal. Note that this E-step is subject to the constraints of the original Penn Treebank annotation: an NP can be reannotated as NP4, but not as VP3. Now, the marginals are defined as $p(X \rightsquigarrow w_{i:j} \mid w_{1:M})$, which is the probability that the span $w_{i:j}$ is derived from the non-terminal X , conditioning on the entire sentence $w_{1:M}$ and marginalizing over all

Proper nouns			
NNP-14	<i>Oct.</i>	<i>Nov.</i>	<i>Sept.</i>
NNP-12	<i>John</i>	<i>Robert</i>	<i>James</i>
NNP-2	<i>J.</i>	<i>E.</i>	<i>L.</i>
NNP-1	<i>Bush</i>	<i>Noriega</i>	<i>Peters</i>
NNP-15	<i>New</i>	<i>San</i>	<i>Wall</i>
NNP-3	<i>York</i>	<i>Francisco</i>	<i>Street</i>
Personal Pronouns			
PRP-0	<i>It</i>	<i>He</i>	<i>I</i>
PRP-1	<i>it</i>	<i>he</i>	<i>they</i>
PRP-2	<i>it</i>	<i>them</i>	<i>him</i>

Table 11.3: Examples of automatically refined non-terminals and some of the words that they generate (Petrov et al., 2006).

other parts of the derivation. Such marginals can be computed by a two-pass recursive algorithm called the **inside-outside algorithm** (Lari and Young, 1990).

- In the **inside** step, we compute the likelihood $p(\mathbf{w}_{i:j} \mid X)$, which is simply the probability of deriving the span $\mathbf{w}_{i:j}$ from the non-terminal X ; this probability depends only on the grammar, and not on any other parts of the sentence.
- In the **outside** step, we compute the probability $p(X \mid \mathbf{w}_{1:i-1}, \mathbf{w}_{j+1:M})$, which is the probability of a non-terminal X governing the span $\mathbf{w}_{i:j}$, conditioned on the “outside” parts of the sentence, $\mathbf{w}_{1:i-1}$ and $\mathbf{w}_{j+1:M}$.

Each of these probabilities can be computed recursively. The marginal is then computed from the product of the inside and outside probabilities. The inside-outside algorithmic is a direct analogue of the forward-backward algorithm, which we used to compute the marginals necessary for training conditional random fields over sequence models.

In the M-step, we recompute the parameters of the grammar, based on the expected counts from the E-step. As usual, this process can be iterated to convergence. To determine the number of refinement types for each tag, Petrov et al. (2006) apply a split-merge heuristic; Liang et al. (2007) and Finkel et al. (2007) apply Bayesian nonparametrics.

This approach yielded state-of-the-art accuracy at the time. Some examples of refined non-terminals are shown in Table 11.3. The proper nouns differentiate months, first names, middle initials, last names, first names of places, and second names of places; each of these will tend to appear in different parts of grammatical productions. The personal pronouns differentiate grammatical role, with PRP-0 appearing in subject position at the beginning of the sentence (note the capitalization), PRP-1 appearing in subject position but not at the beginning of the sentence, and PRP-2 appearing in object position.

(c) Jacob Eisenstein 2014-2017. Work in progress.

11.5 Discriminative parsing

The methods described in the previous section are all based on generative parsing models, in which the probability of a parse is a product of the probabilities of the individual productions. As we have seen, these models can be improved by using finer-grained non-terminals, via parent-annotation, lexicalization, and state-splitting. An alternative path to making parsing more accurate is to use techniques from discriminative machine learning. With the exception of reranking (discussed below), the introduction of discriminative methods to parsing came relatively late. The main reason is that these learning algorithms require multiple passes over the data, applying the parser repeatedly. Unlike sequence labeling, where the time complexity of inference is linear in the size of the input, the cost of inference for parsing is non-trivial — cubic in the length of the input. These limitations prevented well-known discriminative learning techniques, such as structured perceptron, from being applied sooner.

Reranking

An inexpensive way to get the benefits of discriminative learning is through **reranking** (Charniak and Johnson, 2005; Collins and Koo, 2005). First, a generative model — such as the Collins or Charniak parser — is used to identify the K -best parses for a sentence. (A modified version of CKY can compute the K -best parses efficiently.) Then a discriminative learning algorithm is trained to select the best of these parses. The discriminative model does not need to search over all parses, it only needs to consider the best K identified by the “generator.” This means that the discriminator can use arbitrary features, such as structural features that capture parallelism and right-branching, which could not be easily incorporated into a bottom-up parsing model. Because learning is discriminative, rerankers can also use rich lexicalized features, relying on regularization to combat overfitting. Overall, this approach yields substantial improvements in accuracy on the Penn Treebank, and can be applied to improve any generative parsing model. The main limitation is that reranking can only find the best parse among the K -best offered by the generator, so it is inherently limited by the ability of the generator to find high-quality parse candidates.

Discriminative parsing

As shown in § 11.3, the weights on productions need not correspond to probabilities; the CKY algorithm can apply to **any** set of weights, as long as they are context-free. Discriminative learning can therefore be applied by setting $\psi_{X \rightarrow YZ} = \theta \cdot f(X \rightarrow YZ, w, i, j, k)$, with the indices i, j, k indicating the boundaries of the parent ($w_{i:j-1}$) and its left and right children ($w_{i:k-1}$ and $w_{k:j-1}$). Such features could incorporate lexical information, so that we learn weights for non-terminal productions as well as for lexicalized forms. For example:

(c) Jacob Eisenstein 2014-2017. Work in progress.

- $f1: NP(*) \rightarrow NP(*) PP(*)$
- $f2: NP(cats) \rightarrow NP(cats) PP(*)$
- $f3: NP(*) \rightarrow NP(*) PP(claws)$
- $f4: NP(cats) \rightarrow NP(cats) PP(claws)$

Through regularization, we can find weights that strike a good balance between frequently-observed features ($f1$) and more discriminative features ($f4$).

This approach was implemented by Finkel et al. (2008) in the context of weighted CFG parsing with conditional random fields. They used stochastic gradient descent for training, with the inside-outside algorithm (analogous to forward-backward, but for trees) to compute expected feature counts. Like CKY, the runtime of the inside-outside algorithm is cubic in the length of the input. Because each instance must be visited and parsed many times during stochastic gradient descent, efficiency is critical. One solution is to “prefilter” the CKY parsing chart, identifying and eliminating productions which cannot be part of any complete parse.

Carreras et al. (2008) use the averaged perceptron to perform conditional parsing, employing an alternative feature decomposition based on tree-adjoining grammar (TAG; Carreras et al., 2008). They use features that capture “grandparent” dependencies between words and the heads of their parents’ parents. These second-order dependency features make the time complexity $\mathcal{O}(M^4)$ in the length of the input, so pruning is again required to make parsing efficient enough to train accurately.

Neural parsing

Recent work has applied neural representations to parsing, representing units of text with dense numerical vectors (Socher et al., 2013a; Durrett and Klein, 2015). Neural approaches to natural language processing will be surveyed in chapter 20. [todo: say a little more about durrett and klein]

Chapter 12

Dependency Parsing

The previous chapter discussed algorithms for analyzing sentences in terms of nested **constituents**, such as noun phrases and verb phrases. The combination of constituency structure and head-percolation rules yields a set of **dependencies** between individual words. These dependencies are a more “bare-bones” version of syntax, leaving out information that is present in the full constituent parse. Nonetheless, the dependency representation is still capable of capturing important linguistic phenomena, such as the prepositional phrase attachment and coordination scope. For this reason, dependency parsing is increasingly used in applications that require syntactic analysis. While dependency structures can be obtained as a byproduct of constituent parsing, it is more efficient to extract them directly. Indeed, accurate dependency parses can be obtained by algorithms with time complexity that is linear in the length of the sentence. This chapter begins by overviewing dependency grammar, and then presents the two dominant approaches to dependency parsing, graph-based and transition-based dependency parsing.

12.1 Dependency grammar

In lexicalized parsing, non-terminals such as NP are augmented with **head words**, as shown in Figure 12.1a. In this sentence, the head of the S constituent is the main verb, *scratch*; this non-terminal then produces the noun phrase *the cats*, whose head word is *cats*, and from which we finally derive the word *the*. Thus, the word *scratch* occupies the central position for the sentence, with the word *cats* playing a supporting role. In turn, *cats* occupies the central position for the noun phrase, with the word *the* playing a supporting role.

These relationships, which hold between the words in the sentence, can be formalized in a directed graph structure. In this graph, there is an edge from word i to word j iff word i is the head of the first branching node above a node headed by j . Thus, in our example, we would have $scratch \rightarrow cats$ and $cats \rightarrow the$. We would not have the edge




Figure 12.1: Dependency grammar is closely linked to lexicalized context free grammars: each lexical head has a dependency path to every other word in the constituent.

scratch \rightarrow *the*, because although *scratch* dominates *the* in the graph, it is not the head of a node that produces a node headed by *the*. These edges describe syntactic **dependencies**, a bilocal relationship between a **head** and a **dependent**, which is at the heart of **dependency grammar** (Tesnière, 1966).

If we continue to build out this **dependency graph**, we will eventually reach every word in the sentence, as shown in Figure 12.1b. In this graph — and in all graphs constructed in this way — every word will have exactly one incoming edge, except for the root word, which is indicated by a special incoming arrow from above. Another feature of this graph is that it is **weakly connected**, in the sense that if we replaced the directed edges with undirected edges, there would be a path between all pairs of nodes. From these properties, it can be shown that there are no cycles in the graph (or else at least one node would have to have more than one incoming edge), and therefore, the graph is a **tree**.

Although we have begun by motivating dependency grammar in terms of lexicalized constituent parsing, there is a rich literature on dependency grammar as a model of syntax in its own right (Tesnière, 1966). Kübler et al. (2009) provides a comprehensive overview of this literature.

What do the edges mean?

A dependency edge implies an asymmetric syntactic relationship between the head and dependent words. For a pair like *the cats* or *cats scratch*, how do we decide which is the head? Here are some possible criteria:

- The head sets the syntactic category of the construction: for example, nouns are the heads of noun phrases, and verbs are the heads of verb phrases.

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 12.2: A labeled dependency parse

- The modifier may be optional while the head is mandatory: for example, in the sentence *cats scratch people with claws*, the substrings *cats scratch* and *cats scratch people* are grammatical sentences, but *with claws* is not.
- The head determines the morphological form of the modifier: for example, in languages that require gender agreement, the gender of the noun determines the gender of the adjectives and determiners.

As always, these guidelines sometimes conflict, but it is possible to use these basic principles to define fairly consistent conventions at the level of part-of-speech tags, similar to the head percolation rules from lexicalized constituent parsing.

Edges may be **labeled** to indicate the nature of the syntactic relation that holds between the two elements. An example is shown in Figure 12.2. The edge between *scratch* and *cats* is labeled NSUBJ, with *scratch* as the head; this indicates that the noun subject of the predicate verb *scratch* is headed by the word *cats*. The edge from *scratch* to *people* is labeled with DOBJ; this indicates that the word *people* is the head of the direct object. The Stanford typed dependencies have become a standard inventory of dependency types for English (De Marneffe and Manning, 2008). De Marneffe et al. (2014) propose a more minimal “universal” set of dependencies that is suitable for many languages.

Ambiguity and difficult cases

[**todo: update this section with current standards from universal dependency treebank (Nivre et al., 2016)**] The attachment ambiguity in the sentence shown in Figure 12.2 can be represented by a single change: replacing the edge from *scratch* to *with* by an edge from *people* to *with*. This should give you an idea of why labeled dependency trees are useful: they tell us who did what to whom.

However, dependency trees are less structurally expressive than lexicalized CFG derivations. That means they hide information that would be present in a CFG parse. Often this “information” is in fact irrelevant for any conceivable linguistic purpose: for example, Figure 12.3 shows three different ways of representing prepositional phrase adjuncts to

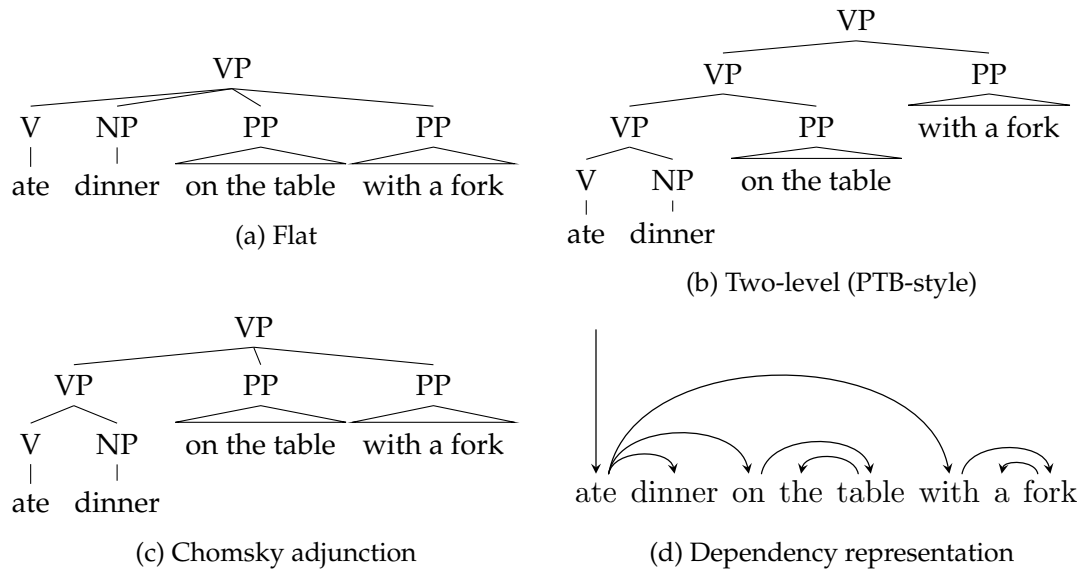


Figure 12.3: The three different CFG analyses of this verb phrase all correspond to a single dependency structure.

the verb *ate*. Because there is apparently no meaningful difference between these analyses, the Penn Treebank decides by convention to use the two-level representation. As shown in Figure 12.3d, these three cases all look the same in a dependency parse. So if you didn't think there was any meaningful difference between these three constituent representations, you may view this as an advantage of the dependency representation.

Dependency grammar still leaves open some tricky representational decisions. For example, coordination is a challenge: in the sentence, *Abigail and Max like kimchi* (Figure 12.4), which word is the immediate dependent of the main verb *likes*? Choosing either *Abigail* or *Max* seems arbitrary; for fairness we might choose *and*, but this seems in some ways to be the least important word in the noun phrase. One typical solution is to simply choose the left-most item in the coordinated structure — in this case, *Abigail*. Another alternative, as shown in Figure 12.4c, is a **collapsed** dependency grammar in which conjunctions are not included as nodes in the graph, but are instead used to label the edges (De Marneffe et al., 2006). Popel et al. (2013) survey alternatives for handling this phenomenon across several dependency treebanks.

The same logic that makes us reluctant to accept *and* as the head of a coordinated noun phrase may also make us reluctant to accept a preposition as the head of a prepositional phrase. In the sentence *cats scratch people with claws*, surely the word *claws* is more central than the word *with* — and it is precisely the bilexical relations between *scratch*, *claws*, and *people* that help guide us to the correct syntactic interpretation. Yet there are also arguments for preferring the preposition as the head — as we saw in § 11.4, the preposition




Figure 12.4: Three alternatives for representing coordination in a dependency parse

itself is what helps us to choose verb attachment in *meet the President **on** Monday* and noun attachment in *meet the President **of** Mexico*. Collapsed dependency grammar is again a possible solution: we can collapse out the prepositions so that the dependency chain,

$$President \rightarrow_{prep} of \rightarrow_{pobj} Mexico$$

would be replaced by $President \rightarrow_{PREP:of} Mexico$. Dependency annotation is an active area of research due to the ongoing development of the universal dependency treebank, which has produced dependency-annotated corpora in 47 languages at the time of this writing (Nivre et al., 2016).¹

Projectivity

The dependency graphs that can be built from all possible lexicalized constituent parses of a sentence with M words are a proper subset of the spanning trees over M nodes. In other words, there exist spanning trees that do not correspond to any lexicalized constituent parse. This is because syntactic constituents are **contiguous** spans of text, so that the head h of the constituent that spans the nodes from i to j must have a path to every node in this span. This property is known as **projectivity**. Informally, it means that “crossing edges” are prohibited. The formal definition follows:

Definition 2 (Projectivity). *An edge from i to j is projective iff all k between i and j are descendants of i . A dependency parse is projective iff all its edges are projective.*

If we were to annotate a dependency parse directly — rather than deriving it from a lexicalized constituent parse — such non-projective edges would occur. Figure 12.5 gives an example of a non-projective dependency graph in English. This dependency graph does not correspond to any constituent parse. In languages where non-projectivity is common, such as Czech and German, it is better to annotate dependency trees directly, rather than deriving them from constituent parses. An example is the Prague Dependency Treebank (Böhmová et al., 2003), which contains 1.5 million words of Czech, with

¹<http://universaldependencies.org/>

	% non-projective edges	% non-projective sentences
Czech	1.86%	22.42%
English	0.39%	7.63%
German	2.33%	28.19%

Table 12.1: Frequency of non-projective dependencies in three languages (Kuhlmann and Nivre, 2010)




Figure 12.5: An example of a non-projective dependency parse in English

approximately 12,000 non-projective edges (see Table 12.1). Even though relatively few dependencies are non-projective in Czech and German, many sentences have at least one such dependency.

As we will see in the next section, projectivity has important consequences for the sorts of algorithms that can perform dependency parsing.

12.2 Graph-based dependency parsing

Let $\mathbf{y} = \{\langle i, j, r \rangle\}$ indicate a dependency graph with relation r from head word w_i to dependent word w_j . We would like to define a scoring function $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{y}, \mathbf{w})$, where $\mathbf{f}(\mathbf{y}, \mathbf{w})$ is a vector of features on the dependency graph and sentence, and $\boldsymbol{\theta}$ is a vector of weights. The dependency parsing problem is then the structure prediction problem,

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{y}, \mathbf{w}). \quad (12.1)$$

As usual, the number of possible labelings $\mathcal{Y}(\mathbf{w})$ is exponential in the length of the input. In the case of non-projective dependency parsing, the set $\mathcal{Y}(\mathbf{w})$ includes all possible spanning trees over a complete graph with M nodes, where M is the length of the sentence \mathbf{w} . The size of this set is M^{M-2} (Wu and Chao, 2004). Algorithms that search over this space of possible graphs are known as **graph-based dependency parsers**.

In sequence labeling and constituent parsing, it was possible to search efficiently over an exponential space by choosing a feature function that decomposes into a sum of local

feature vectors. A similar approach is possible for dependency parsing, by requiring the feature function to decompose across dependency arcs $i \rightarrow j$:

$$f(\mathbf{y}, \mathbf{w}) = \sum_{\langle i, j, r \rangle \in \mathbf{y}} f(\mathbf{w}, i, j, r) \quad (12.2)$$

$$\theta \cdot f(\mathbf{y}, \mathbf{w}) = \sum_{\langle i, j, r \rangle \in \mathbf{y}} \theta \cdot f(\mathbf{w}, i, j, r). \quad (12.3)$$

Dependency parsers that operate under this assumption are known as **arc-factored**, since the overall (exponentiated) score is a product of scores over all arcs. As described later in this section, the arc-factored assumption enables efficient algorithms for dependency parsing.

Features

Typical features for arc-factored dependency parsing are similar to those used in sequence labeling and discriminative constituent parsing. They include: the length and direction of the dependency arc; the words linked by the dependency relation; their prefixes, suffixes, and part-of-speech tags (as produced by an automatic tagger); and their neighbors in the sentence. In labeled dependency parsing, each of these features are also conjoined with the relation type r .

Bilexical features, which include both the head and the dependent, will be helpful for common words, but will be extremely sparse for rare words. It is therefore necessary to include features at various levels of detail, such as: word-word, word-tag, tag-word, and tag-tag. For example, for the arc *scratch* \rightarrow *cats*, we might have the features,

$$\begin{array}{ll} \{w_i \rightarrow w_j : & \text{scratch} \rightarrow \text{cats}, \\ w_i \rightarrow t_j : & \text{scratch} \rightarrow \text{NNS}, \\ t_i \rightarrow w_j : & \text{VBP} \rightarrow \text{cats}, \\ t_i \rightarrow t_j : & \text{VBP} \rightarrow \text{NNS} \} \end{array}$$

Regularized discriminative learning algorithms can then learn to trade off between features that are rare but highly predictive, and features that are common but less informative.

As with sequence labeling, it is possible to include features on neighboring words without breaking the locality restriction: we can consider features such as the identity, part-of-speech, and shape of the preceding and succeeding words, $w_{i-1}, w_{i+1}, w_{j-1}, w_{j+1}$. What we cannot do (yet) is consider other parts of the graph \mathbf{y} , such as the parent of i (which I will denote $w_{\Gamma(i)}$) or the siblings of j , the set $\{w_j : \Gamma(j) = i\}$. This requires higher-order dependency parsing, discussed in § 12.2.

To give a concrete example, the seminal paper by McDonald et al. (2005a) includes the following features for an arc between words w_i and w_j , with part-of-speech tags t_i and t_j :

(c) Jacob Eisenstein 2014-2017. Work in progress.

Unigram features $\langle w_i \rangle; \langle t_i \rangle; \langle w_i, t_i \rangle; \langle w_j \rangle; \langle t_j \rangle; \langle w_j, t_j \rangle$.

Bigram features $\langle w_i, t_i, w_j, t_j \rangle; \langle w_i, w_j, t_j \rangle; \langle t_i, w_j, t_j \rangle; \langle w_i, t_i, t_j \rangle; \langle w_i, t_i, w_j \rangle; \langle w_i, w_j \rangle; \langle t_i, t_j \rangle$.

“In-between” features $\langle t_i, t_k, t_j \rangle$ for all k between i and j .

Neighbor features $\langle t_i, t_{i+1}, t_{j-1}, t_j \rangle; \langle t_{i-1}, t_i, t_{j-1}, t_j \rangle; \langle t_i, t_{i+1}, t_j, t_{j+1} \rangle; \langle t_{i-1}, t_i, t_j, t_{j+1} \rangle$

In addition, all the word features are supplemented with the five-character prefixes for all words longer than five characters (e.g., *unconscionable* \rightarrow *uncon*). The bigram features include several varieties of backoff from the most detailed 4-tuple feature; McDonald et al. (2005a) note that these backoff features were particularly helpful, presumably because they improve generalization. The “in-between” features activate for all part-of-speech tags between positions i and j in the sentence. This feature group helps to “rule out situations when a noun would attach to another noun with a verb in between, which is a very uncommon phenomenon.”

Learning

Having formulated graph-based dependency parsing as a structure prediction problem, we can apply similar learning algorithms to those used in sequence labeling. The most direct application is structured perceptron,

$$\hat{y} = \operatorname{argmax}_{y' \in \mathcal{Y}(w)} \theta \cdot f(w, y') \quad (12.4)$$

$$\theta = \theta + f(w, y) - f(w, \hat{y}) \quad (12.5)$$

This is just like sequence labeling, but now the argmax requires a maximization over all dependency trees for the sentence. Algorithms for performing this search efficiently are described below. We can apply all the usual tricks from chapter 2: weight averaging, large-margin, and regularization. McDonald et al. (2005a,b) were the first to treat dependency parsing as a structure prediction problem, using MIRA (a close relative of the passive-aggressive algorithm we saw in chapter 2) to obtain high accuracy parses in both projective and non-projective settings.

Conditional random fields (CRFs) are globally-normalized conditional models (see chapter 6), and they can be applied to any graphical model in which we can efficiently compute marginal probabilities over individual random variables — in this case, we need marginals over the edges. The marginals are required because the unregularized log-likelihood has a gradient that sums over all possible edges, taking the difference between the features in the observed dependency parses and the expected feature counts under $p(y | w)$:

$$\frac{\partial \mathcal{L}}{\partial \theta} = \sum_{(i,j) \in \mathcal{Y}} f(w, i, j) - \sum_{i,j} p(i \rightarrow j | w) f(w, i, j) \quad (12.6)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

Algorithm 9 Chu-Liu-Edmonds algorithm for unlabeled dependency parsing

```

1: procedure CHU-LIU-EDMONDS( $\{\psi(i \rightarrow j)\}_{i,j \in \{1 \dots M\}}$ )
2:   for  $j \in 1 \dots M$  do
3:      $h_j \leftarrow \operatorname{argmax}_i \psi(i \rightarrow j)$  ▷ Find the best incoming edge for each node
4:    $\tau \leftarrow \{j, h_j\}_{j \in 1 \dots M}$  ▷  $\tau$  is the graph of the best incoming edges
5:    $\mathcal{C} \leftarrow \text{FINDCYCLES}(\tau)$ 
6:   if  $\mathcal{C} = \emptyset$  then
7:     return  $\tau$  ▷ If  $\tau$  has no cycles, it is the best tree
8:   else ▷ Otherwise, collapse each cycle
9:     for each cycle  $c \in \mathcal{C}$  do
10:       Remove all nodes in the cycle from the graph
11:       Add a “super-node” representing the cycle
12:     Let  $G$  be the resulting graph
13:     return CHU-LIU-EDMONDS( $G$ ) ▷ Call recursively on the collapsed graph

```

For projective dependency trees, the marginal probabilities can be computed in cubic time, using a variant of the inside-outside algorithm (Lari and Young, 1990). For non-projective dependency parsing, marginals can also be computed in cubic time, using the **matrix-tree theorem** (Koo et al., 2007; McDonald et al., 2007; Smith and Smith, 2007). Details of these methods are described by Kübler et al. (2009).

Algorithms for non-projective dependency parsing

In **non-projective dependency parsing**, the goal is to identify the highest-scoring spanning tree over the words in the sentence. The arc-factored assumption ensures that the score for each spanning tree will be computed as a sum over scores for the edges. We can precompute these scores, $\psi(i \rightarrow j, r) = \theta \cdot \mathbf{f}(\mathbf{w}, i, j, r)$, before applying a parsing algorithm. (We must compute $\mathcal{O}(M^2 R)$ such scores, where M is the length of the sentence and R is the number of dependency relation types, so this is a lower bound on the time complexity of any exact algorithm for dependency parsing.)

Based on these scores, we build a weighted connected graph. Arc-factored non-projective dependency parsing is then equivalent to finding the the spanning tree that achieves the maximum total score, $\sum_{\langle i, j, r \rangle \in \mathbf{y}} \psi(i \rightarrow j, r)$. The **Chu-Liu-Edmonds algorithm** (Chu and Liu, 1965; Edmonds, 1967) computes this spanning tree in time $\mathcal{O}(M^3 R)$. The algorithm, which is sketched in Algorithm 9, operates recursively. It first identifies the highest scoring incoming edge for each node, and then checks the graph for cycles. If there are no cycles, the resulting graph is a spanning tree, and moreover, it is the maximum spanning tree, because there is no better-scoring incoming edge for each node. If there is a cycle, the cycle is collapsed into a “super-node”, whose incoming edges have scores equal to

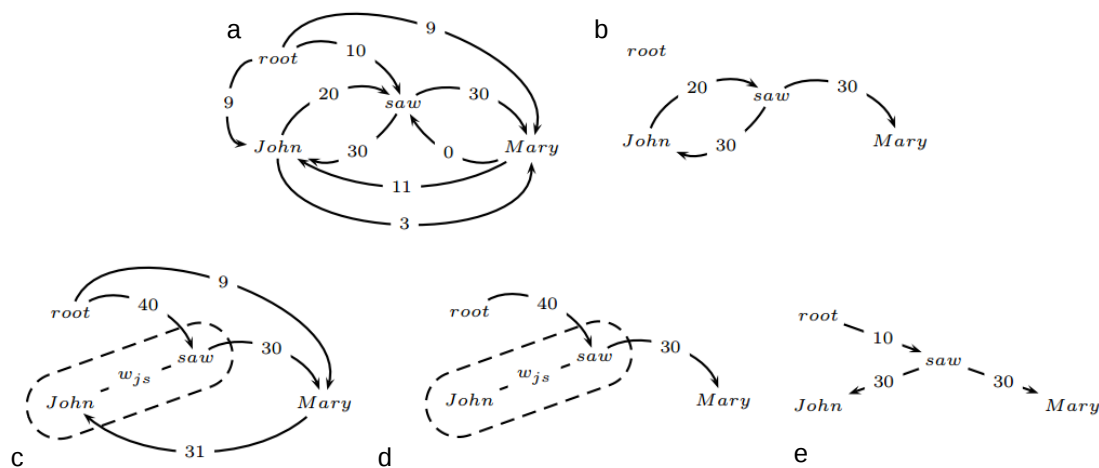


Figure 12.6: An illustration of the MST algorithm on a simple example. Figure borrowed from McDonald et al. (2005b).

the scores of the best spanning tree that includes both the edge and all nodes in the cycle. [todo: more detail on what happens when you collapse cycles].

The algorithm works because it can be proved that the maximum spanning tree on the contracted graph is equivalent to the maximum spanning tree on the original graph. The basic process is illustrated in Figure 12.6. In part (a), we see the complete graph, which includes all edge scores $\psi(i \rightarrow j)$. In (b), we see the highest scoring incoming edge for each node. In (c), the cycle between *John* and *saw* is contracted, creating new incoming edges with weight 40 from the root, and weight 31 from *Mary*. In (d), we find the highest-scoring incoming edge in the new graph. There are no remaining cycles, so we recover the maximum spanning tree.

Let us consider the time complexity of unlabeled dependency parsing first. For each of the M words in the sentence, one must search all $M - 1$ other words for the highest-scoring incoming edge, for a time complexity of $\mathcal{O}(M^2)$. In the worst case, it is necessary

(c) Jacob Eisenstein 2014-2017. Work in progress.

to contract the graph M times. If we redo the search within each contraction, we face a total cost of $\mathcal{O}(M^3)$. Recall that the CKY constituent parsing algorithm is also cubic time complexity in the length of the sentence. However, further optimizations are possible, resulting in a complexity of $\mathcal{O}(M^2)$ (Tarjan, 1977). To generalize the algorithm to labeled dependency parsing, it is necessary only to compute the best scoring label for each possible edge. Because of the arc-factoring assumption, the edge labels are decoupled from each other, so this can be done as a preprocessing step, with total complexity of $\mathcal{O}(M^2 R)$.

Algorithms for projective dependency parsing

The Chu-Liu-Edmonds algorithm finds the best scoring dependency tree, but it does not enforce the projectivity constraint. For languages in which we expect projectivity — such as English — we may prefer to ensure that the parsing algorithm returns only projective trees. Note that the arc-factored assumption makes it impossible to **learn** to produce projective trees, since projectivity cannot be encoded in a feature that decomposes over individual arcs.

Recall that it is possible to convert any lexicalized constituent parse directly into a projective dependency parse. This means that any algorithm for lexicalized constituent parsing is also an algorithm for projective dependency parsing. One such algorithm is presented in § 11.4, in which we built a table where the cell $t[i, j, h, X]$ contains the score of the best derivation of the substring $w_{i:j}$ from non-terminal X , in which the head is w_h . For unlabeled projective dependency parsing, we can apply a very similar algorithm:

$$t_\ell[i, j, h] = \max_{k > h} \max_{k \leq h' < j} t[i, k, h] + t[k, j, h'] + \psi(h \rightarrow h') \quad (12.7)$$

$$t_r[i, j, h] = \max_{k \leq h} \max_{i \leq h' < k} t[i, k, h'] + t[k, j, h] + \psi(h \rightarrow h') \quad (12.8)$$

$$t[i, j, h] = \max(t_\ell[i, j, h], t_r[i, j, h]). \quad (12.9)$$

The goal is for $t[i, j, h]$ to contain the score of the best-scoring projective dependency tree for $w_{i:j}$, headed by w_h . We must first maximize over all h' , which is the location of an immediate dependent of w_h . Projectivity guarantees that the subtree headed by h' will extend to one of the endpoints of the entire span: either from the left endpoint i to some midpoint k , or from some midpoint k to the right endpoint j . We compute the best score for each of these possibilities separately in Equation 12.7 and Equation 12.8. Computing each of these scores also involves maximizing over all possible midpoints k .

We construct the table t from the bottom up: first compute scores for all subtrees of size 2, then size 3, and so on. The total size of the table is $\mathcal{O}(M^3)$, and to complete each cell we must search over $\mathcal{O}(M)$ dependents and $\mathcal{O}(M)$ split points. Thus, the overall complexity is $\mathcal{O}(M^5)$. The Eisner (1996) algorithm reduces this complexity to $\mathcal{O}(M^3)$ by maintaining multiple tables. For a detailed description of this algorithm, see Kübler et al. (2009). As with the Chu-Liu-Edmonds algorithm, the best-scoring label for each edge can be computed as a preprocessing step. [todo: write up formal algorithm description]

Higher-order dependency parsing

Arc-factored dependency parsers can only score dependency graphs as a product across their edges. However, it can be useful to consider higher-order features, which consider pairs or triples of edges, as shown in Figure 12.7. Second-order features consider **siblings** and **grandchildren**; third-order features consider **grand-siblings** (siblings and grandparents together) and **tri-siblings**.




Figure 12.7: Feature templates for higher-order dependency parsing (Koo and Collins, 2010)

Why might we need higher-order dependency features? Consider the example *cats scratch people with claws*, where the preposition *with* could attach to either *scratch* or *people*. In a lexicalized first-order arc-factored dependency parser, we would have the following feature sets for the two possible parses:

- $\langle \text{ROOT} \rightarrow \text{scratch} \rangle, \langle \text{scratch} \rightarrow \text{cats} \rangle, \langle \text{scratch} \rightarrow \text{people} \rangle, \langle \text{scratch} \rightarrow \text{with} \rangle, \langle \text{with} \rightarrow \text{claws} \rangle$
- $\langle \text{ROOT} \rightarrow \text{scratch} \rangle, \langle \text{scratch} \rightarrow \text{cats} \rangle, \langle \text{scratch} \rightarrow \text{people} \rangle, \langle \text{people} \rightarrow \text{with} \rangle, \langle \text{with} \rightarrow \text{claws} \rangle$

The only difference between the feature vectors are the features $\langle \text{scratch} \rightarrow \text{with} \rangle$ and $\langle \text{people} \rightarrow \text{with} \rangle$, but both are reasonable features, both syntactically and semantically. A first-order arc-factored dependency parsing model would therefore struggle to find the right solution to this sentence. However, if we add grandchild features, then our feature sets include:

- $\langle \text{scratch} \rightarrow \text{with} \rightarrow \text{claws} \rangle$
- $\langle \text{people} \rightarrow \text{with} \rightarrow \text{claws} \rangle$,

The first feature is preferable, so a second-order dependency parser would have a better chance of correctly parsing this sentence. In general, higher-order features can yield substantial improvements in dependency parsing accuracy (e.g., Koo and Collins, 2010).

Projective second-order parsing can still be performed in $\mathcal{O}(M^3)$ time (and $\mathcal{O}(M^2)$ space), using a modified version of the Eisner algorithm. Projective third-order parsing can be performed in $\mathcal{O}(M^4)$ time and $\mathcal{O}(M^3)$ space (Koo and Collins, 2010). Approximate pruning algorithms can reduce this cost significantly by filtering out unpromising edges (Rush and Petrov, 2012).

Given the tractability of higher-order projective dependency parsing, you may be surprised to learn that non-projective second-order dependency parsing is NP-Hard! This can be proved by reduction from the vertex cover problem (Neuhaus and Bröker, 1997). One heuristic solution is to do projective parsing first, and then post-process the projective dependency parse to add non-projective edges (Nivre and Nilsson, 2005). More recent work has applied advanced techniques for approximate inference in graphical models, including belief propagation (Smith and Eisner, 2008), integer linear programming (Martins et al., 2009), variational inference (Martins et al., 2010), and Markov Chain Monte Carlo (Zhang et al., 2014).

12.3 Transition-based dependency parsing

Graph-based dependency parsing offers exact inference, meaning that it is possible to recover the best-scoring parse. But this exactness comes at a price: we can use only a limited set of features. These limitations are felt more keenly in dependency parsing than in sequence labeling, because second-order dependency features are critical to correctly identify certain types of attachments. Graph-based dependency parsing may also be criticized on the basis of intuitions about human language processing: people read and listen to sentences *sequentially*, incrementally building mental models of the sentence structure and meaning before getting to the end (Jurafsky, 1996). This seems hard to reconcile with graph-based algorithms, which perform bottom-up operations on the entire sentence, requiring the parser to keep every word in memory.

Transition-based algorithms address both of these objections. They work by moving through the sentence sequentially, while incrementally updating a stored representation of what has been read thus far. After processing the entire sentence, they return an analysis of its syntactic structure. A simple transition-based parser is the **arc-standard** parsing algorithm, which is similar to the LR algorithm that is used to parse programming languages. Transition-based parsing algorithms maintain a configuration state, which includes a stack where elements can be pushed and popped. They update the state incrementally through a series of actions, until the input is consumed and the stack is empty.

In the arc-standard parser, the configuration c is a tuple $c = (\sigma, \beta, A)$, where σ is a stack, β is the input buffer, and A is a set of dependency arcs. The initial state is $(\sigma = [0], \beta = w_{1:M}, A = \emptyset)$, where: $\sigma = [0]$ indicates that the stack begins with the root node; $\beta = w_{1:M}$ indicates that the buffer begins with the entire input string (indexed from 1); and $A = \emptyset$ means that there are not yet any arcs. We can then apply three possible

actions:

- **SHIFT**: moves the first item from the input buffer on to the top of the stack,

$$(\sigma, i|\beta, A) \Rightarrow (\sigma|i, \beta, A), \quad (12.10)$$

where we write $i|\beta$ to indicate that i is the leftmost item in the input buffer, and $\sigma|i$ to indicate the result of pushing i on to stack σ .

- **ARC-LEFT**: creates a new left-facing arc between the item on the top of the stack and the first item in the input buffer. This item is then “popped” to the front of the input buffer, and the arc is added to A .

$$(\sigma|i, j|\beta, A) \Rightarrow (\sigma, j|\beta, A \cup (j, \ell, i)), \quad (12.11)$$

where ℓ is the (optional) label of the dependency arc.

- **ARC-RIGHT**: creates a new right-facing arc between the item on the top of the stack and the first item in the input buffer; this item is then “popped” to the front of the input buffer, and the arc is added to A .

$$(\sigma|i, j|\beta, A) \Rightarrow (\sigma, i|\beta, A \cup (i, \ell, j)), \quad (12.12)$$

where again ℓ is the label of the dependency arc.

The **ARC-LEFT** action cannot be performed when the root node 0 is on top of the stack, since this node must be the root of the entire tree. Neither **ARC-LEFT** nor **ARC-RIGHT** can be performed if the result would create a second incoming edge for any word. When the stack σ and the input buffer β are empty, parsing is complete.

Learning transition-based parsers

Transition-based parsing requires selecting a series of actions. In parsing programming languages, shift-reduce parsers can choose the appropriate action deterministically, because programming languages are unambiguous by design. For natural language, we use machine learning classification to determine the best series of actions; for example, Yamada and Matsumoto (2003) use a support vector machine classifier (see § 2.1) to decide whether to shift or create a dependency arc at each stage in parsing.

To train a transition-based dependency parser, we can treat each parsing decision as a separate training instance. However, our ground truth input is not a list of parsing decisions, but rather, a dependency tree. We therefore require an **oracle** to convert the ground truth dependency tree into a list of parsing decisions, which can then be used as training data (Nivre, 2008).²

²**Spurious ambiguity** occurs when there are multiple derivations for the same dependency structure. This is the case in arc-standard dependency parsing: the structure $1 \leftarrow 2 \rightarrow 3$ can be obtained from two different action sequences (Cohen et al., 2012).

Typical features for transition-based dependency parsing include: the word and part-of-speech of the top element on the stack; the word and part-of-speech of the first, second, and third elements on the input buffer; pairs and triples of words and parts-of-speech from the top of the stack and the front of the buffer; the distance (in tokens) between the element on the top of the stack and the element in the front of the input buffer; the number of modifiers of each of these elements; and higher-order dependency features as described above in the section on graph-based dependency parsing. Zhang and Nivre (2011) describe a transition-based parser with rich features, which gave state-of-the-art performance (at the time) in both English and Chinese.

Pros and cons of transition-based dependency parsing

A key advantage of transition-based parsing is that it is much faster than graph-based methods. Since every word can be shifted once and every arc-creation action eliminates a word from the stack, the time complexity is linear in the length of the input. In contrast, graph-based parsing algorithms have quadratic or cubic time complexity.

Transition-based parsing can suffer from search errors, since an early mistake can make it impossible to find the best parse. This means that there could be an action sequence that would be preferred by the current parsing model, but is nonetheless not chosen because the first few actions in the sequence score badly. Put another way, transition-based parsing is **greedy** — unlike graph-based algorithms, which are guaranteed to find the best-scoring overall analysis. Solutions to this problem are discussed below.

Nonetheless, transition-based parsing achieves comparable accuracy to graph-based methods, in far less time (Nivre, 2004; Nivre et al., 2007). One reason is that in exchange for giving up on global inference (and thereby accepting the possibility of search errors), we free ourselves from any restrictions on the features that can be used in the classifier that selects each parsing action. For example, features may consider any number of previous parsing decisions, any aspect of the current stack, and any part of the input.

Alternative transition-based parsing algorithms

Arc-eager dependency parsing changes the ARC-RIGHT action so that right dependents can be attached before all of **their** dependents have been found. In arc-eager parsing, the ARC-RIGHT action creates an arc, and then pushes both the parent and child elements on to the stack. To remove these elements, it adds an addition REDUCE action, which can be applied to elements on the stack for whom an incoming edge has already been identified. Arc-eager parsing is arguably more cognitively plausible, because it constructs larger connected components incrementally, rather than having a deep stack with lots of disconnected elements (Abney and Johnson, 1991; Nivre, 2004).

(c) Jacob Eisenstein 2014-2017. Work in progress.

Beam search A drawback of transition-based parsing is the possibility for search errors, in which a poor decision early in the parse will lead to cascading errors. **Beam search** is an improvement on greedy transition-based parsing, with the goal of eliminating search errors. As we move through the sentence, we keep a beam of possible hypotheses. For each element on the beam, we consider possible actions, and obtain a list of the top- k possibilities across all such actions. We then update the beam with the results of these top- k actions, and proceed (Zhang and Clark, 2008). Huang et al. (2012b) offer alternative perceptron learning rules that are specifically designed for learning in the beam search setting.

Shift-reduce parsing for CFGs Transition-based parsing can also be used for parsing in (binarized) context-free grammars. Here we use a shift-reduce parser, where each reduce operation creates a new non-terminal that produces the top two elements in the stack. When the input is consumed and the only element on the stack is a tree derived from the start symbol S , the input has been completely parsed.

Neural transition-based parsing

[todo: 2-3 paragraphs about the following papers:]

- Each shift-reduce decision is made by a locally-trained neural network (Chen and Manning, 2014). See also (Dyer et al., 2015)
- Shift-reduce decisions are made by neural network trained on global conditional likelihood (Andor et al., 2016), using beam search.

12.4 Applications

Dependency parsing is used in many real-world applications: any time you want to know about pairs of words which might not be adjacent, you can use dependency links instead of typical regular expression search patterns. For example, we may want to match strings like *delicious pastries*, *delicious French pastries*, and *the pastries are delicious*³

It is now possible to search Google n-grams by dependency edges; for example, finding the trend in how often a dependency edge has appeared over time. For example, we might be interested in knowing when people started talking about *writing code*, but we also want *write some code*, *write the code*, *write all the code*, etc. By searching on dependency edges, we can recover this information, as shown in Figure 12.8. This capability has implications for research in digital humanities, as shown by the analysis of Shakespeare performed by Muralidharan and Hearst (2013).

³Recall that the copula *is* is collapsed in many dependency grammars, such as the Universal Dependency treebank Nivre et al. (2016).




Figure 12.8: Google n-grams results for the bigram *write code* and the dependency arc *write => code* (and their morphological variants)

A classic application of dependency parsing is **relation extraction**, which is described in chapter 17. The goal of relation extraction is to identify entity pairs, such as

⟨TOLSTOY, WAR AND PEACE⟩
 ⟨MARQUÉZ, 100 YEARS OF SOLITUDE⟩
 ⟨SHAKESPEARE, A MIDSUMMER NIGHT’S DREAM⟩,

which stand in some relation to each other (in this case, the relation is authorship). Such entity pairs are often referenced via consistent chains of dependency relations. Therefore, dependency paths are often a useful feature in supervised systems which learn to detect new instances of a relation, based on labeled examples of other instances of the same relation type (Culotta and Sorensen, 2004; Fundel et al., 2007; Mintz et al., 2009).

Cui et al. (2005) show how dependency parsing can improve question answering. For example, you might ask,

(12.1) *What % of the nation’s cheese does Wisconsin produce?*

Now suppose your corpus contains this sentence:

(12.2) *In Wisconsin, where farmers produce 28% of the nation’s cheese, ...*

The location of *Wisconsin* in the surface form of this string might make it a poor match for the query. However, in the dependency graph, there is an edge from *produce* to *Wisconsin* in both the question and the potential answer, raising the likelihood that this span of text is relevant to the question.

A final example comes from sentiment analysis. As discussed in chapter 3, the polarity of a sentence can be reversed by negation, e.g.

(12.3) *There is no reason at all to believe the polluters will suddenly become reasonable.*

(c) Jacob Eisenstein 2014-2017. Work in progress.

By tracking the sentiment polarity through the dependency parse, we can better identify the overall polarity of the sentence, determining when key sentiment words are reversed (Wilson et al., 2005; Nakagawa et al., 2010).

Exercises

1. The dependency structure $1 \leftarrow 2 \rightarrow 3$, with 2 as the root, can be obtained from more than one set of actions in arc-standard parsing. List both sets of actions that can obtain this parse.

Part III

Meaning

Chapter 13

Logical semantics

A grand ambition of natural language processing, and indeed, all of artificial intelligence, is to convert natural language into a representation that supports **semantic inferences**.¹ Many applications of language technology involve some level of semantic understanding:

- Answering questions. This includes “real-life” questions like *where can a guy find a decent cup of coffee around here?*, and also “quiz show” questions like *what’s the middle name of the mother of the 44th President of the United States?*
- Translating a sentence from one language into another, while preserving the underlying meaning.
- Building a robot that can follow natural language instructions and execute useful tasks.
- Fact-checking an article by searching the web for contradictory evidence.
- Logic-checking an argument by trying to identify contradictions or unsupported assertions.

Most approaches towards achieving this level of semantic understanding involve converting natural language to some form of **meaning representation**. Jurafsky and Martin (2009) compare several alternative representations, showing parallels between several representations that are superficially distinct. Therefore, we will focus on logical representations: **boolean logic**, **first-order logic**, and the **lambda calculus**.

13.1 Meaning representations

The goal of a meaning representation is to provide a way to express **propositions**, while abstracting over the ambiguity and vagueness of natural language. There are several

¹Alternative readings on this topic include the chapter from Jurafsky and Martin (2009), a more involved “informal” reading from Levy and Manning (2009), and a yet more involved introduction from Briscoe (2011).

criteria that a meaning representation should meet:

Verifiability It should be possible to test the truth of assertions in the meaning representation. Indeed, in **truth-conditional semantics**, the meaning of a sentence is said to be identical to its truth conditions: that is, to the set of facts that must hold in the world for the sentence to be true.

We might imagine that verifiability should be tested against the real world: for example, if faced with the proposition *Alice hates apples*, we could verify it by finding Alice and asking her. However, it is better still to be able to reason about **possible worlds**, such as fictional worlds in which *Alice* (or *apples*) might refer to arbitrarily different entities. In **model-theoretic semantics**, each proposition has a **denotation** in a model of the world, enabling propositions to be verified against specific models corresponding to possible worlds. Why is this useful? Consider that Lois Lane is unaware that Superman and Clark Kent are the same person — that is, **SUPERMAN** and **CLARKKENT** have different denotations in her model. Model-theoretic semantics makes it possible to interpret statements from her perspective, so that, for example, it would not be absurd for her to ask Clark to speak with Superman.²

Truth-conditional semantics allows us to define additional concepts of **equivalence** and **entailment**. A statement *P* is entailed by statement *Q* iff the truth conditions for *Q* imply the truth conditions for *P*. For example, the statement *Alice gives Bob a book about calculus* entails the statements *Alice gives Bob a book*, *Alice gives someone a book*, *Someone gives Bob a book*, etc. Iff *P* entails *Q* and *Q* entails *P*, then we can say that *P* and *Q* are logically equivalent.

No ambiguity Each sentence in the meaning representation should have exactly one meaning. In truth conditional semantics, this means that each sentence in the meaning representation has exactly one corresponding set of truth conditions.

Clearly this criterion is not met by natural language. Many of the syntactic ambiguities that we encountered in previous sections have corresponding semantic ambiguities: consider the truth conditions for the two possible PP attachments in our example *cats scratch people with claws*, or the example *she fed her dog biscuits*. Natural language also has ambiguity at the lexical level: the sentence *Dong bought a plant* would have distinct truth conditions depending on whether *plant* refers to something like a shrub, or a factory for producing widgets.

Jurafsky and Martin (2009) mention a converse criterion, **canonical form**, which requires that each meaning (set of truth conditions) has a single representation. For

²Example from Percy Liang's slides on semantics, <http://icml.cc/2015/tutorials/icml2015-nlu-tutorial.pdf>

example, if we consider the database query language SQL as a meaning representation, then it is easy to design superficially distinct queries that will return the same results regardless of what database they are applied to:

- (13.1) `SELECT RestaurantID, City FROM Restaurants WHERE City = 'Atlanta' OR City = 'New York'`
- (13.2) `SELECT RestaurantID, City FROM Restaurants WHERE City = 'New York' OR City = 'Atlanta'`

In general, it is difficult to design meaning representations in which every meaning has a single canonical form. However, removing unnecessary flexibility can vastly simplify the computation associated with verifying statements and performing **inference** (described below).

Expressiveness Meaning representation is useful only to the extent that it enables us to talk about a wide range of different things. This is partly a matter of the **non-logical vocabulary** that the representation includes: the set of entities (e.g., *Alice*, *Bob*) and relations (e.g., *likes*, *brother-of*) that can be included in sentences. However, there are also deeper structural limits on expressiveness. Consider the following possibilities:

- (13.3) *Alice admires Bob*
- (13.4) *Alice admires Bob and Bob trusts Alice*
- (13.5) *Alice admires someone*
- (13.6) *Alice admires someone who trusts her*
- (13.7) *Everyone whom Alice admires trusts someone*
- (13.8) *Not everyone whom Alice admires trusts Bob*

To handle all of these cases, we must have an appropriate **logical vocabulary**, including boolean connectives and quantifiers. More on this in § 13.2.

Inference We would like to be able to combine assertions in our meaning representation to infer new facts about the world. For example, given the assertion *Bart is Lisa's brother*, we should be able to infer that *Someone is Lisa's brother*. Given the additional information that Lisa is female, we should be able to infer that *Lisa is Bart's sister* — although this inference is of a different type, since it requires additional knowledge about the relations **BROTHER** and **SISTER**.

How do natural languages like English do on these criteria? They are infinitely expressive, but highly ambiguous. Because we cannot establish the truth conditions of natural language expressions without ambiguity, it is difficult to speak of verifying their meaning or drawing further inferences.

(c) Jacob Eisenstein 2014-2017. Work in progress.

But if natural language is not itself a meaning representation, we would still like to be able to find the most likely meaning, or the set of possible meanings, for a given natural language sentence. This task is known as **semantic parsing**, and it typically rests on the assumption that meaning is determined **compositionally**, with the meaning of a sentence determined by the meanings of its constituent expressions, and the operations that are used to combine them. In particular, we will assume that the relevant substrings of a sentence correspond to the syntactic constituents identified during CFG-style parsing, and that each parsing production corresponds to some semantic operation. More on this in § 13.3.

13.2 Logical representations of meaning

We will build a meaning representation on logical semantics, which does a pretty good job of meeting the criteria established in the previous section.

Propositional logic

The bare bones of logical meaning representation are boolean operations on propositions:

Propositional symbols We use the symbols P, Q, \dots to represent propositions; for example, P may correspond to the proposition, *bagels are delicious*.

Boolean operators We can evaluate the truth of more complex statements through boolean operators: negation ($\neg P$, which is true if P is false), conjunction ($P \wedge Q$, which is true if both P and Q are true), and disjunction ($P \vee Q$, which is true if at least one of P and Q is true). Other operators can be derived from these: for example, implication ($P \Rightarrow Q$) has identical truth conditions to $\neg P \vee Q$; equivalence ($P \Leftrightarrow Q$) has identical truth conditions to $(P \wedge Q) \vee (\neg P \wedge \neg Q)$. In fact, if we have \neg , then only one of \wedge and \vee is needed; we can derive the other.

We can define axioms or inference rules in terms of these boolean connectives (commutativity, associativity, etc), and then derive further equivalences, which can support some inferences. For example, suppose $P = \textit{The music is loud}$ and $Q = \textit{Max can't sleep}$. Then if we have $P \Rightarrow Q$ (*If the music is loud, Max can't sleep*) and P (*the music is loud*), then we have Q (*Max can't sleep*). However, there are other inferences that we cannot perform with propositional logic alone. For example, let $R = \textit{The music is quiet}$; then we might hope that $R \Rightarrow \neg P$, but this is not supported without knowing more about the propositions themselves. For this, we turn to predicate logic.

Predicate logic

Predicate logic extends our meaning representation with several additional classes of terms:

Constants These are elements that name individual entities in the model, such as MAX and THEMUSIC. We say that the **denotation** of each constant in a model \mathcal{M} is an element in the model, e.g., $\llbracket \text{MAX} \rrbracket = d$ and $\llbracket \text{THEMUSIC} \rrbracket = m$.

Predicates Predicates can be thought of as sets of objects, or equivalently, as functions from objects to truth values. For example CANSLEEP is a predicate, and we may have $\llbracket \text{CANSLEEP} \rrbracket = \{d, e, \dots\}$, denoting the set of individuals who can sleep. We can then test the proposition CANSLEEP(MAX) by asking whether $\llbracket \text{MAX} \rrbracket \in \llbracket \text{CANSLEEP} \rrbracket$.

Functions Functions can be thought of as sets of pairs of objects, or equivalently, as functions from one object to another. For example BROTHER-OF is a function, so that $\llbracket \text{BROTHER-OF}(\text{LISA}) \rrbracket = \llbracket \text{BART} \rrbracket$.

We can now express statements like

$$\text{ISQUIET}(\text{THEMUSIC}) \Leftrightarrow \neg \text{ISLOUD}(\text{THEMUSIC}), \quad (13.1)$$

but this only applies to a specific constant, THEMUSIC, and not more generally. For example, we might prefer to say that *anything* that is quiet is not loud. To make such general statements, we will need two additional elements in our meaning representation:

Variables These are mechanisms for referring to objectives, which are not locally specified. We can then write BROTHER-OF(x) or ISLOUD(x), using x here as an **unbound variable**.

Quantifiers To bind variables, we use quantifiers. Variables can be used to refer to some particular unspecified object, or to all possible objectives. Correspondingly, we have two connectives, \exists and \forall . The statement,

$$\exists x : \text{BROTHER-OF}(\text{LISA}) = x, \quad (13.2)$$

uses the **existential quantifier** \exists to assert that there is at least one object which is the brother of Lisa in the model. The statement,

$$\forall x : \text{ISLOUD}(x) \Leftrightarrow \neg \text{ISQUIET}(x) \quad (13.3)$$

uses the **universal quantifier** \forall to generalize the relationship between the predicates ISLOUD and ISQUIET; for this sentence to be true, it must be the case that for all entities in the model, the predicate ISLOUD only holds in exactly those cases in which the predicate ISQUIET does not hold.

Lambda calculus

Predicate logic is verifiable, unambiguous, expressive enough for a wide range of statements, and supports inferences; it does a good job meeting all of the criteria listed at the beginning of the chapter. But we still need a few more pieces before we can build logical meanings from natural language sentences.

Recall the assumption of **compositionality**, which states that the meaning of a natural language sentence is composed from the meaning of its constituents. Now, a simple sentence like *Max likes dragons* has two top-level constituents in a CFG parse: the NP *Max*, and the VP *likes dragons*. The meaning of *Max* is the constant MAX, and the meaning of the entire sentence might be $\text{LIKES}(\text{MAX}, \text{DRAGONS})$. But what is the meaning of the VP constituent *likes dragons*?

We will think of the meaning of VPs such as *likes dragons* as **functions** which require additional arguments to form a sentence in predicate logic. The notation for describing such functions is called **lambda calculus**, and it involves expressions such as $\lambda x.P(x)$, which indicates a function that takes an argument x and then has value $P(x)$. The application of a function $\lambda x.P(x)$ to an argument A is written

$$\lambda x.P(\dots, x, \dots)(A) \quad (13.4)$$

$$P(\dots, A, \dots), \quad (13.5)$$

indicating that A is playing the role occupied by the variable x , which is bound here by the lambda expression. It is crucial to note that P itself may be a lambda expression, so that application can be performed multiple times.

13.3 Syntax and semantics

We will now extend CFG products to include the meaning of each constituent, using rules of the form,

$$X : \alpha \rightarrow Y : \beta \quad Z : \gamma, \quad (13.6)$$

where X, Y, Z are syntactic non-terminals and α, β, γ are the meanings associated with each constituent.

For example, consider the very simple fragment,

$$S : \beta(\alpha) \rightarrow \text{NP} : \alpha \quad \text{VP} : \beta \quad (13.7)$$

$$\text{VP} : \beta(\alpha) \rightarrow \text{V} : \beta \quad \text{NP} : \alpha \quad (13.8)$$

$$\text{Abigail, NP} : \text{ABIGAIL} \quad (13.9)$$

$$\text{Max, NP} : \text{MAX} \quad (13.10)$$

$$\text{likes, V} : \lambda y. \lambda x. \text{LIKE}(x, y) \quad (13.11)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

Lines 13.9-13.11 describe the **lexicon**, listing the syntactic categories and semantic meanings of individual words. Words may have multiple entries in the lexicon, depending on their semantics; for example, the verb *eats* may be intransitive (*Abigail eats*) or transitive (*Abigail eats kimchi*), so we need two lexical entries:

$$\text{eats}, V : \lambda x. \text{EAT}(x) \quad (13.12)$$

$$\text{eats}, V : \lambda y. \lambda x. \text{EAT}(x, y). \quad (13.13)$$

Now, given the sentence *Max likes Abigail*, we get the following analysis,

$$P = \lambda y. \lambda x. \text{LIKES}(x, y)(\text{MAX})(\text{ABIGAIL}) \quad (13.14)$$

$$= \lambda x. \text{LIKES}(x, \text{ABIGAIL})(\text{MAX}) \quad (13.15)$$

$$= \text{LIKES}(\text{MAX}, \text{ABIGAIL}) \quad (13.16)$$

Noun phrases

What about sentences with more complex noun phrases like *Max has a red bear* or *Abigail eats all the spicy snacks*? To handle these cases, we'll need to deal with determiners, adjectives, and general nouns. Let's start with a relatively simple case,

(13.9) *A dog likes Max.*

The desired analysis is,

$$(\text{A dog likes Max.}) . \text{sem} = \exists x. \text{DOG}(x) \wedge \text{LIKES}(x, \text{MAX}), \quad (13.17)$$

where $(\text{text}) . \text{sem}$ indicates the semantics of *text*.

We already know that the meaning of the verb phrase *likes Max* is $\lambda x. \text{LIKES}(x, \text{MAX})$, and we would like to apply this function to the argument specified by the noun phrase. But somehow we have to get to a solution where the outermost term is the existential quantifier $\exists x$, and not the predicate LIKES . How can we do it?

The solution is to introduce some additional operations for **type-shifting**. The semantic type of the verb phrase *likes Max* was a function mapping from entities to truth values, $\lambda x. \text{LIKES}(x, \text{MAX})$. We now introduce the **type-raising** operation $\alpha \rightarrow \lambda P. P(\alpha)$, indicating that the semantics α can be replaced with a function that takes P as an argument, and returns $P(\alpha)$. Applying type-raising to the verb phrase *likes Max*, we obtain, $\lambda P. P(\lambda x. \text{LIKES}(x, \text{MAX}))$.

Now, how should we think of the noun phrase *a dog*? The determiner implies an existential quantifier (there exists some dog...) over all dogs, $\exists x. \text{DOG}(x)$. Moreover, we are planning to apply some additional functions to explain what this dog is doing. So the semantics we want is $\lambda P. \exists (x) \text{DOG}(x) \wedge P(x)$. We can get there by appropriately defining

(c) Jacob Eisenstein 2014-2017. Work in progress.

the determiner a , and the production $\text{NP} \rightarrow \text{DET NN}$.

$$\text{NP} : \beta(\alpha) \rightarrow \text{DET} : \beta \quad \text{NN} : \alpha \quad (13.18)$$

$$a, \text{DET} : \lambda P. \lambda Q. \exists x. P(x) \wedge Q(x) \quad (13.19)$$

$$\text{dog}, \text{NN} : \lambda x. \text{DOG}(x) \quad (13.20)$$

Note that although we have typically treated the noun as the head of a noun phrase, it is the determiner whose semantics takes precedence in Equation 13.18. This enables us to properly assess the meaning of the phrase $a \text{ dog}$,

$$(a \text{ dog}).\text{sem} = (\lambda P. \lambda Q. \exists x. P(x) \wedge Q(x))(\lambda x. \text{DOG}(x)) \quad (13.21)$$

$$= \lambda Q. \exists (x). \text{DOG}(x) \wedge Q(x) \quad (13.22)$$

So now we have the two pieces,

$$(a \text{ dog}).\text{sem} = \lambda Q. \exists (x). \text{DOG}(x) \wedge Q(x) \quad (13.23)$$

$$(\text{likes Max}).\text{sem} = \lambda x. \text{LIKES}(x, \text{MAX}) \quad (13.24)$$

$$= \lambda P. P(\lambda x. \text{LIKES}(x, \text{MAX})), \quad (13.25)$$

using type-raising on the verb phrase. We can now combine the pieces, using the verb phrase semantics as a function on the noun phrase,

$$(a \text{ dog likes Max}).\text{sem} = (\lambda P. P(\lambda x. \text{LIKES}(x, \text{MAX}))) (\lambda Q. \exists (x). \text{DOG}(x) \wedge Q(x)) \quad (13.26)$$

$$= (\lambda Q. \exists (x). \text{DOG}(x) \wedge Q(x)) (\lambda x. \text{LIKES}(x, \text{MAX})) \quad (13.27)$$

$$= \exists (x). \text{DOG}(x) \wedge \text{LIKES}(x, \text{MAX}), \quad (13.28)$$

which is the desired semantics that we identified above for this sentence. A useful exercise is to try to do the same kind of analysis for the sentence *Max likes a dog*.

$$(a \text{ dog}).\text{sem} = \lambda P. \exists x. P(x) \wedge \text{DOG}(x) \quad (13.29)$$

$$(\text{likes}).\text{sem} = \lambda y. \lambda z. \text{LIKES}(z, y) \quad (13.30)$$

$$= \lambda Q. Q(\lambda y. \lambda z. \text{LIKES}(z, y)) \quad (13.31)$$

$$(\text{likes a dog}).\text{sem} = (\lambda Q. Q(\lambda y. \lambda z. \text{LIKES}(z, y))) (\lambda P. \exists x. P(x) \wedge \text{DOG}(x)) \quad (13.32)$$

$$= (\lambda P. \exists x. P(x) \wedge \text{DOG}(x)) (\lambda y. \lambda z. \text{LIKES}(z, y)) \quad (13.33)$$

$$= \exists x. (\lambda y. \lambda z. \text{LIKES}(z, y))(x) \wedge \text{DOG}(x) \quad (13.34)$$

$$= \exists x. \lambda z. \text{LIKES}(z, x) \wedge \text{DOG}(x) \quad (13.35)$$

$$(\text{Max likes a dog}).\text{sem} = \exists x. \text{LIKES}(\text{MAX}, x) \wedge \text{DOG}(x) \quad (13.36)$$

[todo: double-check this]

(c) Jacob Eisenstein 2014-2017. Work in progress.

Full semantic analysis of natural language requires handling many more phenomena, but the basic strategy of function application and type-shifting covers much of what is needed. Jurafsky and Martin (2009) provide more details than presented here, and a book-length treatment is offered by Blackburn and Bos (2005).

13.4 Semantic parsing

The goal of **semantic parsing** is to convert natural language statements to a representation such as predicate logic with lambda calculus. Zettlemoyer and Collins (2005) show that it is possible to train such a system, using labeled data of natural language sentences and their associated logical meanings. They use a linear model, in which each syntactic-semantic production has an associated feature weight, which is learned from labeled data.³ A key point is that a sentence may have analyses that produce the same logical interpretation, which is known as **spurious ambiguity**. They do not have labeled data for the specific productions, so they treat this as a latent variable, and learn using a latent variable perceptron, where

$$z^* = \operatorname{argmax}_z \theta \cdot f(w, y, z) \quad (13.37)$$

$$\hat{y}, \hat{z} = \operatorname{argmax}_{y, z} \theta \cdot f(w, y, z) \quad (13.38)$$

$$\theta^{(t+1)} \leftarrow \theta^{(t)} + f(w, y, z^*) - f(w, \hat{y}, \hat{z}), \quad (13.39)$$

with y indicating the logical interpretation and z indicating the derivation of that interpretation from the input w .

A more ambitious approach is to train a semantic parser not from sentences annotated by their logical forms, but rather, from question-answer pairs, e.g., $\langle \textit{Where is Georgia Tech?}, \textit{Atlanta} \rangle$. There are now two latent variables: the logical form y , and the derivation of that logical form, z . We constrain the logical form y such that its denotation $\llbracket y \rrbracket$ is identical to the denotation of the logical form of the answer, e.g.,

$$\llbracket \lambda x. \text{LOCATED-IN}(\text{GEORGIA TECH}, x) \rrbracket = \llbracket \text{ATLANTA} \rrbracket. \quad (13.40)$$

This idea has been implemented by Clarke et al. (2010) and Liang et al. (2013), yielding systems that can answer questions about geographical relationships with above 90% accuracy.

³Zettlemoyer and Collins (2005) do not use context-free grammar, but instead use a mildly context-sensitive formalism called **Combinatory Categorical Grammar** (CCG). Semantic parsing is considerably easier to explain in CCG, but would require introducing a new syntactic formalism.

Chapter 14

Shallow semantics

“Full” compositional semantics requires representations at least as expressive as first-order logic. Machine learning approaches have improved robustness, and recent work has driven down the requirements for manually-created resources. But coverage is still relatively limited, with best performance in narrow domains like travel and geography.

Shallow semantics comprises a set of alternative approaches, which trade the expressiveness of representations like first-order logic for shallower representations which can be parsed more robustly, with broader coverage.

14.1 Predicates and arguments¹

Shallow semantics focuses on predicate-argument relations. For example, the sentence *Abigail trusts Max* can be interpreted as `trusts(ABIGAIL, MAX)`, where `trusts` is a predicate and `ABIGAIL` and `MAX` are its arguments. This is exactly the sort of relation that we saw in first-order logical semantics too, but in shallow semantics we will typically work without variables and quantification. (Recent explorations of intermediate representations between FOL and shallow predicate-argument relations are described in Section 14.4.)

To see how shallow semantics can represent meaning, consider these four sentences (borrowed from the slides of a tutorial by Kristina Toutanova and Scott Yih).

(14.1) [Yesterday]₃, [Kristina]₀ hit [Scott]₁ [with a baseball]₂

(14.2) [Scott]₁ was hit by [Kristina]₀ [yesterday]₃ [with a baseball]₂

(14.3) [Yesterday]₃, [Scott]₁ was hit [with a baseball]₂ by [Kristina]₀

(14.4) [Kristina]₀ hit [Scott]₁ [with a baseball]₂ [yesterday]₃

We don’t need first-order logic to realize that these sentences are semantically identical. Shallow semantics will suffice: the *roles* in each sentence are filled by the same text.

¹This section follows closely from J&M 2009

- [Hitter]₀: *Kristina*
- [Person hit]₁: *Scott*
- [Instrument of hitting]₂: *with a baseball*
- [Time of hitting]₃: *yesterday*

The event semantics representation for the sentence *Scott was hit by Kristina yesterday* (and all of the other examples) is:

$$\begin{aligned} \exists e. \text{Hitting}(e) \wedge \text{Hitter}(e, \text{Kristina}) \wedge \text{PersonHit}(e, \text{Scott}) \\ \wedge \text{TimeOfHitting}(e, \text{Yesterday}) \end{aligned}$$

In this example, *Hitter*, *PersonHit*, and *TimeOfHitting* are roles. We use these specific roles because of the **predicate verb** *hit*. Roles that relate to a specific predicate are called **deep roles**.

Thematic roles

Without knowing more about deep roles like *Hitter*, we cannot do much inference. But building classifiers for every role of every predicate would be a lot of work, and we would struggle to get enough training data to accomplish this. Is there a shortcut?

Consider the example *Scott was paid by Kristina yesterday*. Clearly *yesterday* is filling the same role in this example as in Examples § 14.1-item 14.4, describing the time at which the events occur — regardless of whether the event is *hitting* or *paying*. But arguably, the role-fillers *Scott*, *Kristina* and *yesterday* also have similar thematic functions as in the earlier sentence about baseballs.

- *Kristina* is causing the event by performing an action, which she does volitionally (on purpose); we can generalize her **thematic role** in these examples as the AGENT of the event.
- *Scott* is the primary experiencer of the effects of the event. We can generalize his thematic role as the PATIENT.

AGENT and PATIENT are the two best-known examples of **thematic roles** (Fillmore, 1968),² which attempt to generalize across predicates. They are also among the least controversial (Dowty, 1991); other thematic roles are shown in Table 14.1, but it is important to emphasize that this particular role inventory is not universally accepted, or even accepted to the same extent as, say, the Penn Treebank syntactic categories.

²The idea of thematic roles can be traced to the Sanskrit linguist Pāṇini (7th-4th century BCE!).

AGENT	The volitional causer <i>The waiter spilled the soup</i>
EXPERIENCER	The experiencer <i>The soup gave all three of us a headache.</i>
FORCE	The non-volitional causer <i>The wind blew my soup off the table.</i>
THEME	The participant most directly affected <i>The wind blew my my soup off the table.</i>
RESULT	The end product <i>The cook has prepared a cold duck soup.</i>
CONTENT	The proposition or content of a propositional event <i>The waiter assured me that the soup is vegetarian.</i>
INSTRUMENT	An instrument used in an event <i>It's hard to eat soup with chopsticks.</i>
BENEFICIARY	The beneficiary <i>The waiter brought me some soup.</i>
SOURCE	The origin of the object of a transfer event <i>The stack of canned soup comes from Pittsburgh.</i>
GOAL	The destination of the object of a transfer event <i>He brought the bowl of soup to our table.</i>

Table 14.1: Definitions and examples of thematic roles (Jurafsky and Martin, 2009)

Case frames Different verbs take different thematic roles as arguments. The possible arguments for a verb is the **case frame** or **thematic grid**. For example, for *break*:

- AGENT: Subject, THEME: Object
John broke the window.
- AGENT: Subject, THEME: Object, INSTRUMENT: PP (with)
John broke the window with a rock.
- INSTRUMENT: Subject, THEME: Object
The rock broke the window.
- THEME: Subject
The window broke.

(c) Jacob Eisenstein 2014-2017. Work in progress.

When two verbs have similar case frames, this is a clue that they might be semantically related: (e.g., *break*, *shatter*, *smash*).

Many verbs permit multiple orderings of the same arguments. These are known as **diathesis alternations**. For example, *give* permits the dative alternation,

(14.5) [AGENT *Doris*] *gave* [GOAL *Cary*] [THEME *the book*].

(14.6) [AGENT *Doris*] *gave* [THEME *the book*] [GOAL *to Cary*].

Again, similar alternation patterns suggest semantic similarity. For example, verbs that display the dative alternation include some broad classes:

- “verbs of future having” (advance, allocate, offer, owe)
- “verbs of sending” (forward, hand, mail)
- “verbs of throwing” (kick, pass, throw)

The purpose of thematic roles is to abstract above verb-specific roles. But it is usually possible to construct examples in which thematic roles are insufficiently specific.

- *Intermediary instruments* can act as subjects:
 1. *The cook opened the jar with the new gadget.*
 2. *The new gadget opened the jar.*
- *Enabling instruments* cannot:
 1. *Shelly ate the pizza with the fork.*
 2. **The fork ate the pizza.*

Thematic roles are bundles of semantic properties, but it’s not clear how many properties are necessary. For example, AGENTS are usually animate, volitional, sentient, causal, but any of these properties may be missing occasionally. The distinction between agents and patients is explored in detail by Dowty (1991).

The Proposition Bank

In the Proposition Bank (**PropBank**), roles are verb-specific, with some sharing (Palmer et al., 2005).

- ARG0: proto-agent (has agent-like properties)
- ARG1: proto-patient (has patient-like properties)
- ARG2 ... ARGN: verb-specific
- 13 universal adjunct-like arguments: temporal, manner, location, cause, negation, ...

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 14.1: Examples of PropBank-style annotations, borrowed from the slides of Toutanova and Yih

PropBank contains two main resources:⁴ “frame files” describing the roles for each verbal predicate (3,324 such files are included), and labeled sentences, built on the Penn TreeBank (113,000 such propositions are annotated). Some example PropBank-style sentence annotations are shown in Figure 14.1. The overlap with the Penn Treebank makes it possible to test the relationship between semantic roles and syntactic constituents. Similar PropBanks have been created for other languages, including Arabic, Chinese, Hindi, and Korean. PropBank is used as the standard dataset for popular shared tasks on Semantic Role Labeling (SRL); some of the main approaches are described in § 14.2.

PropBank describes the predicate-argument structure of verbs, but words belong to other syntactic categories may have argument structures of their own. A related resource is NomBank (Meyers et al., 2004), which annotates the arguments of noun phrases, such as:

(14.7) [ARG0 students'] [REL knowledge] of [ARG1 two-letter consonant sounds]

In this example, the syntactic head is *knowledge*, and this is also the word that defines the semantic relation (REL). The “proto-agent” in this case is *students'*, and the “proto-patient” is *two-letter consonant sounds*.

14.2 Semantic Role Labeling

Semantic role labeling (SRL) is the task of assigning semantic labels to spans of text. Labels describe the role of the phrase with respect to the *predicate verb*. In practice, this usually means PropBank labels, e.g. Arg0, Arg1, etc, so our goal is to produce labelings such as those shown in Figure 14.1.

⁴<https://catalog.ldc.upenn.edu/LDC2004T14>; <http://verbs.colorado.edu/propbank/framesets-english/scratch-v.html>

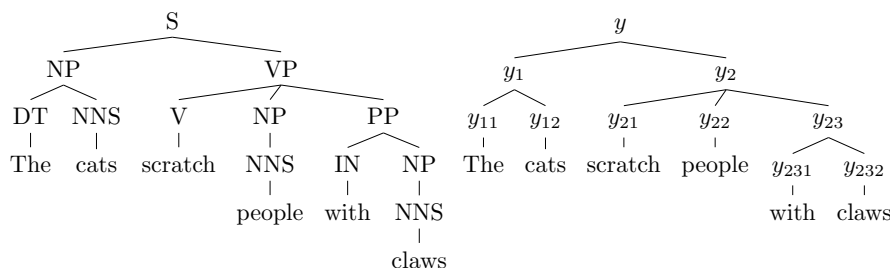


Figure 14.2: Conversion of a constituent parse tree to variables for semantic role labeling

While there are many possible approaches to Semantic Role Labeling (SRL), an effective solution is to treat it as another case of structured prediction. The problem has a few components:

1. identify all predicates in the sentence;
2. identify all argument spans;
3. label the argument spans.

Early approaches treated these problems in isolation, but more recent work has shown that it is best to treat them jointly. Assuming for the moment that we have identified the predicate, the remaining problem can be viewed as simply **tagging** the remaining words in a tagset $\mathcal{T} = \{A0, A1, A2, \dots, A_{M-TMP}, \dots, \emptyset\}$. Thus, the output of an SRL system might be written,

(14.8) *Kristina* / A0 *hit* / PRED *Scott* / A1 *with* / A2 *a* / A2 *baseball* / A2

This would suggest that SRL can be solved by applying a sequence labeling algorithm such as structured perceptron with Viterbi. But recall that Viterbi is based on sequential features, $f(\mathbf{w}, \mathbf{y}) = \sum_m f(\mathbf{w}, y_m, y_{m-1}, m)$; these features are not particularly useful in SRL, because sequential constraints and preferences are less important here than they are in tasks such as part-of-speech tagging and named-entity recognition — recall examples (§ 14.1-item 14.4). In fact, it is better to consider the tree structure offered by a constituent parse of the sentence: in PropBank, 96% of the arguments correspond to a “gold” constituent (from the manual annotation), and 90% correspond to a constituent from an automatic parser (Punyakanok et al., 2008). Therefore we will treat the problem of SRL as a problem of **labeling constituents**, rather than labeling words. This transformation is illustrated in Figure 14.2.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Given a sentence w and a parse tree τ , our goal is now to assign each y_i to a value in the set \mathcal{T} . We optimize a scoring function,

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(w, \tau)} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{y}, w, \tau) \quad (14.1)$$

$$\mathbf{f}(\mathbf{y}, w, \tau) = \sum_{i \in \text{constituents}(\tau)} \mathbf{f}(y_i, w, \tau), \quad (14.2)$$

where we assume that the features decompose across labels y_i . Notice that the features may consider any part of the parse tree, since we are not searching over parse trees. Useful features for this problem include: the predicate verb (which is given); the syntactic type (e.g., NP, VP), head word, first word, and last word of the constituent; whether the constituents comes before or after the predicate; and the **syntactic path** from the constituent to the predicate. This last feature describes a series of steps up and down the parse tree: in the example shown in Figure 14.2, the path from *the cats* (y_1) to the predicate *scratch* (y_{21}) is written NP \uparrow S \downarrow VP \downarrow V. The syntactic path feature captures regularities in the syntactic positions of constituent arguments. For more discussion of features, see Gildea and Jurafsky (2002) and Surdeanu et al. (2007).

The inference problem defined in Equation 14.1 specifies a search over $\mathcal{Y}(w, \tau)$, which is all permissible labelings of the parse tree τ for the sentence w . How should we define this set? If every constituent is allowed to have any label in \mathcal{T} , then we have $\mathcal{Y}(w, \tau) = \mathcal{T}^{\#\tau|}$. But this seems too permissive: it would allow a single argument to appear in multiple places (for example, both *cats* and *claws* labeled as A0), and would also allow multi-word constituents like *the cats* to realize a different argument from their children, like *cats*.

Rather than explicitly defining the set $\mathcal{Y}(w, \tau)$, it is useful to think of **constraints** that a labeling \mathbf{y} must obey. To do this, we will redefine \mathbf{y} slightly, so that it includes a set of indicator features,

$$Y_{i,t} = \begin{cases} 1, & \text{argument } i \text{ takes tag } t \\ 0, & \text{otherwise} \end{cases} \quad (14.3)$$

Now, we can define $\mathcal{Y}(w, \tau)$ to include only those labelings that obey a set of **constraints**. For example:

- All arguments get at most one label, $\forall i \sum_t y_{i,t} = 1$. Note we use equality, because you can always have the \emptyset label.
- No duplicate argument classes, $\forall t \neq \emptyset, \sum_i y_{i,t} \leq 1$
- Overlapping arguments get at most one non-null label:

$$\forall \langle i, j \rangle : i \rightsquigarrow_{\tau} j, y_{i,\emptyset} + y_{j,\emptyset} \geq 1 \quad (14.4)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Some arguments are forbidden, e.g. $\sum_i y_{i,A2} = 0$. Many predicates cannot take all types of arguments: for example, the verb *dream* can only take *A0* and *A1*, so we would add this constraint to make it impossible to label anything as *A2* or *A3*.

All of the constraints are linear, meaning we can write them as a matrix-vector product, $\mathbf{A}\mathbf{y} \leq \mathbf{b}$. Moreover, we can redefine the feature function as,

$$\mathbf{f}(y_i, \mathbf{w}, \tau, i) = \sum_t y_{i,t} \times \mathbf{f}(\mathbf{w}, \tau, i, t). \quad (14.5)$$

After this redefinition, the scoring function is,

$$\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{y}, \mathbf{w}, \tau) = \sum_i \boldsymbol{\theta} \cdot \mathbf{f}(y_i, \mathbf{w}, \tau, i) \quad (14.6)$$

$$= \sum_i \sum_t (\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \tau, i, t)) \times y_{i,t}. \quad (14.7)$$

We can therefore reframe the overall optimization problem as,

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y} \in \mathcal{T}^{\#|\tau|}} \sum_i \sum_t (\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \tau, i, t)) \times y_{i,t} \quad (14.8)$$

$$s.t. \mathbf{A}\mathbf{y} \leq \mathbf{b} \quad (14.9)$$

The objective function is linear in \mathbf{y} , the constraints are linear inequalities, and each $y_{i,t} \in \{0, 1\}$. This optimization problem is therefore a case of **integer linear programming** (ILP). Unfortunately, ILP is known to be NP-hard, including in the binary special case. However, because ILP has many commercial applications, it is a well-studied problem, with heuristic approximations that work well in the overwhelming majority of practical cases. One such algorithm is implemented in the free software GNU Linear Programming Kit (GLPK); Gurobi and CPLEX provide commercial implementations. Integer linear programming is an example of a **combinatorial optimization** problem, with alternative solutions such as **dual decomposition**. Das et al. (2012) develop an “augmented” dual decomposition algorithm which obtains identical accuracy to CPLEX, while running roughly ten times faster.

A final note about constrained optimization approaches to SRL is that you might be uncomfortable about committing to a single syntactic parse, given that even the best parsers have a 10% error rate. Punyakanok et al. (2008) show that you can do better by considering the constituents of five different parsers at the same time! The trick is simple: add constraints preventing the optimizer from selecting constituents that overlap across parses.

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 14.3: Using semantic role labeling to align questions and answers

Applications of SRL Why might we want to do this? One application is to automatic question answering systems like IBM Watson. Consider the example question, *Who discovered prions?*. Somewhere in our database, we have the statement 1997: *Stanley B. Prusiner, United States, discovery of prions....* How can we link them up? Shen and Lapata (2007) use semantic roles to align questions against the content of factual sentences, as shown in Figure 14.3.

[**todo: more applications**]

14.3 FrameNet

PropBank does not attempt to group related predicates, such as BUY/SELL, GIVE/RECEIVE, and RISE/FALL. FrameNet provides a richer model of shallow semantics by grouping predicates and arguments into a predefined **frame** ontology. To see how this works, consider the following examples from Jurafsky and Martin (2009):

(14.9) [_{A1} *The price of bananas*] increased [_{A2} 5%].

(14.10) [_{A1} *The price of bananas*] rose [_{A2} 5%].

(c) Jacob Eisenstein 2014-2017. Work in progress.

FRAMENET ANNOTATION:

[Buyer Chuck] *bought* [Goods a car] [Seller from Jerry] [Payment for \$1000].

[Seller Jerry] *sold* [Goods a car] [Buyer to Chuck] [Payment for \$1000].

PROPBANK ANNOTATION:

[Arg0 Chuck] *bought* [Arg1 a car] [Arg2 from Jerry] [Arg3 for \$1000].

[Arg0 Jerry] *sold* [Arg1 a car] [Arg2 to Chuck] [Arg3 for \$1000].

Figure 14.4: A comparison of framenet and propbank, from Toutanova and Yih [todo: I think]




Figure 14.5: FrameNet annotation, figure from Fleischman et al, 2003

(14.11) *There has been a* [A₂ 5%] *increase* [A₁ in the price of bananas].

The first two sentences involve different verbs; the second sentence conveys same semantics with a noun. Nonetheless, the meaning is the same.

A frame defines a set of *lexical units* and a set of *frame elements*, as shown in Figure 14.5. The relationship between Framenet and PropBank annotation is shown in Figure 14.4. The FrameNet corpus is publicly available online,⁵ and of this writing, annotation is still ongoing.

Unlike PropBank, Framenet is not based on TreeBank parses, and example sentences are chosen by hand. Shi and Mihalcea (2004) present a deterministic algorithm for FrameNet

⁵<https://framenet.icsi.berkeley.edu/fndrupal/about>

parsing, and Das et al. (2010, 2014) provide a structured prediction approach. But compared to PropBank, there is much less work on parsing to the Framenet representation.

14.4 Abstract Meaning Representation

Recent work has focused on a new form of shallow semantics, the **abstract meaning representation** (AMR), which is more structured than PropBank-style semantics, but less ambitious than first-order logic. A major gap in semantic role labeling is the inability to link arguments that refer to a single entity: for example:

(14.12) Abby told Max she would visit him in San Quentin.

In this example, there are three entities, *Abby*, *Max*, and *San Quentin*, and two predicates, *told* and *visit*. The associated AMR structure is shown in Figure 14.6. This graph includes a number of pieces of information about the semantics. The **coreference** relations between *Abby* and *she*, and *Max* and *he* are indicated by having multiple incoming arrows to the nodes representing these entities. In addition, the types of the entities are represented with special nodes: *Abby* and *Max* are of type PER, person; *San Quentin* is of type LOC. The graph also indicates the sense of each predicate, the relationship between the predicates, and the role of each argument.

[todo: Talk a little about AMR parsing] [todo: talk about applications of AMR]

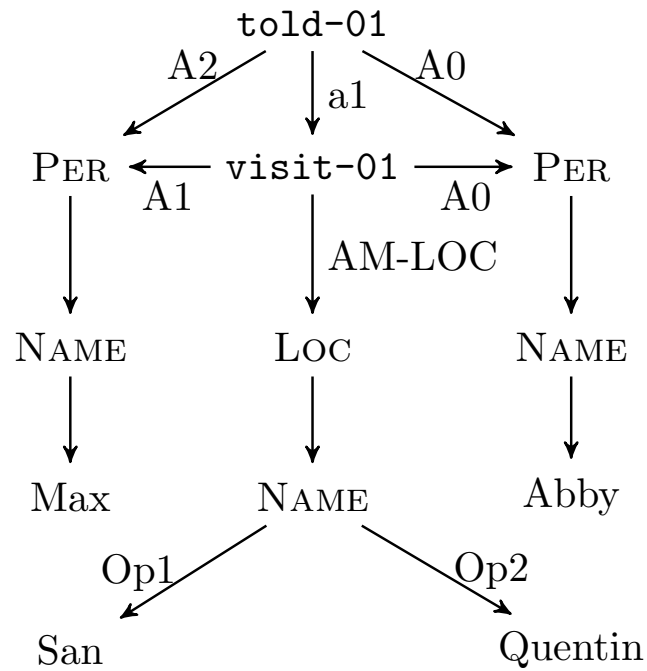


Figure 14.6: An example parse in the Abstract Meaning Representation (AMR)

Chapter 15

Distributional and distributed semantics

A recurring theme in natural language processing is that the mapping from words to meaning is complex. In chapter 3, we saw that a single word form, like *bank*, can have multiple meanings. Conversely, a single meaning may be created by multiple surface forms, a lexical semantic relationship known as **synonymy**. Other lexical semantic relationships include **antonymy** (opposite meaning), **hyponymy** (instance-of), and **meronymy** (part-whole), each of which have semantic consequences for the interpretation of a sentence, utterance, or text.

Despite this complex mapping between words and meaning, the text analytic methods that we have considered thus far tend to consider words as the basic unit of analysis. All of the classifiers, sequence labelers, and parsers from the first chapters rely heavily on lexical features. The logical and frame semantic methods from chapter 13 and chapter 14 rely on hand-crafted lexicons that map from words to semantic predicates. But how can we analyze texts that contain words that we haven't seen before?

This chapter describes methods that create representations of word meaning by analyzing unlabeled data. The goal is for words with similar meanings to have similar representations; if this can be achieved, then it is possible for natural language processing systems to generalize from words that appear in costly resources such as labeled data or semantic lexicons to the broader vocabulary. The primary theory that makes it possible to achieve such representations from unlabeled data is the **distributional hypothesis**.

15.1 The distributional hypothesis

Here's a word you may not know: *tezgüino*.¹ Our position in deciding how to interpret sentences containing this word is similar to that of a natural language processing system

¹The example is from Lin (1998).

	C1	C2	C3	C4	...
<i>tezgüino</i>	1	1	1	1	
<i>loud</i>	0	0	0	0	
<i>motor oil</i>	1	0	0	1	
<i>tortillas</i>	0	1	0	1	
<i>choices</i>	0	1	0	0	
<i>wine</i>	1	1	1	1	

Table 15.1: Distributional statistics for *tezgüino* and five related terms

when encountering a word that does not appear in the labeled training data.

Suppose we see that *tezgüino* is used in the following contexts:

- **C1:** *A bottle of _____ is on the table.*
- **C2:** *Everybody likes _____.*
- **C3:** *Don't have _____ before you drive.*
- **C4:** *We make _____ out of corn.*

What other words fit into these contexts? How about: *loud*, *motor oil*, *tortillas*, *choices*, *wine*? Each row of Table 15.1 is a vector that summarizes the contextual properties for each word, with a value of one for contexts in which the word can appear, and a value of zero for contexts in which it cannot. Based on these vectors, we can conclude:

- *wine* is very similar to *tezgüino*;
- *motor oil* and *tortillas* are fairly similar to *tezgüino*;
- *loud* is quite different.

These vectors, which we will call **word representations**, describe the **distributional** properties of each word. Does vector similarity imply semantic similarity? This is the **distributional hypothesis**, stated by Firth (1957) as: “You shall know a word by the company it keeps.” This hypothesis has been stood the test of time: distributional statistics are a core part of language technology today, mainly because they make it possible to leverage large amounts of unlabeled data to learn about rare words that do not appear in labeled training data.

A striking demonstration of the power of distributional statistics is in their ability to represent lexical semantic relationships such as analogies. Figure 15.1 shows three examples. Distributional statistics are converted into vector **word embeddings**, using the GloVe algorithm, discussed later in this chapter (Pennington et al., 2014). These vectors




Figure 15.1: Lexical semantic relationships have regular linear structures in two dimensional projections of distributional statistics. From <http://nlp.stanford.edu/projects/glove/>.

are then projected into a two dimensional space. In each case, word-pair relationships correspond to regular linear patterns in this two dimensional space. No labeled data about the nature of these relationships was required to identify this underlying structure.

15.2 Design decisions for word representations

There are many approaches for computing word representations. To provide some structure to this space, it is useful to consider three dimensions along which word representations can be compared.

Representation

The most critical question is how words are to be represented. At present, the dominant choice is to represent words as k -dimensional vectors of real numbers, known as **word embeddings**. This representation dates back at least to the late 1980s (Deerwester et al., 1990), and is still used in currently popular techniques such as word2vec (Mikolov et al., 2013a). In word embeddings, similar words typically have high cosine similarity,

$$\cos(\mathbf{u}_i, \mathbf{u}_j) = \frac{\mathbf{u}_i \cdot \mathbf{u}_j}{\|\mathbf{u}_i\| \|\mathbf{u}_j\|}. \quad (15.1)$$

Dense vector word embeddings are well-suited for neural network architectures; they can also be applied in linear classifiers and structure prediction models (Turian et al., 2010), although some authors report that it can be difficult to learn linear models using real-valued features (Kummerfeld et al., 2015). A popular alternative is bit-string representations, such as **Brown clusters**, in which each word is represented by a variable-length sequence of zeros and ones (Brown et al., 1992). Another representational question is whether to estimate one embedding per word surface form or per word stem [**todo: find cite**], or

(c) Jacob Eisenstein 2014-2017. Work in progress.

<i>The moment one learns English, complications set in.</i>		
Brown Clusters (Brown et al., 1992)		$\{\text{learns}\}$
WORD2VEC (Mikolov et al., 2013a) ($c = 2$)		$\{\text{one, learns, complications, set}\}$
Structured WORD2VEC (Ling et al., 2015a) ($c = 2$)		$\{(\text{one}, -2), (\text{learns}, -1), (\text{complications}, +1), (\text{set}, +2)\}$
Dependency texts (Levy and Goldberg, 2014a)	con-	$\{\text{learns}/\text{DOBJ}^{-1}\}$

Table 15.2: Contexts for the word *English*, according to various word representations. For dependency context, $\text{learns}/\text{DOBJ}^{-1}$ means that there is a relation of type DOBJ from the word *learns*.

whether to estimate distinct embeddings for each word sense (Huang et al., 2012a; Li and Jurafsky, 2015).

Context

The distributional hypothesis says that word meaning is related to the contexts in which the word appears, but the notion of context can be defined in a number of different ways. In the *tezgüino* example, contexts are entire sentences, but in practice there are far too many sentences for this to work — the resulting vectors would be too sparse. At the opposite end of the spectrum, the immediately preceding word could be taken as the context, and this is indeed the context considered in Brown clusters, which are discussed in ???. WORD2VEC takes an intermediate approach, using local neighborhoods of words (e.g., $c = 5$) as contexts (Mikolov et al., 2013a). Contexts can also be much larger: for example, in **latent semantic analysis**, each word’s context vector includes an entry per document, with a value of one if the word appears in the document (Deerwester et al., 1990); in **explicit semantic analysis**, these documents are Wikipedia pages (Gabrilovich and Markovitch, 2007).

Words in context can be labeled by their position with respect to the target word w_m (e.g., two words before, one word after), and this appears to make the resulting word representations more sensitive to syntactic differences (Ling et al., 2015a). Another way to incorporate syntax is to perform parsing as a preprocessing step, and then form context vectors from the set of dependency edges (Levy and Goldberg, 2014a) or predicate-argument relations (Lin, 1998). The resulting context vectors for several of these methods are shown in ??.

(c) Jacob Eisenstein 2014-2017. Work in progress.

The choice of context has a profound effect on the resulting representations, which can be viewed in terms of word similarity. Applying latent semantic analysis (??) to contexts of length two and length 30, one obtains the following nearest-neighbors for the word *dog*:²

- ($c = 2$): *cat, horse, fox, pet, rabbit, pig, animal, mongrel, sheep, pigeon*
- ($c = 30$): *kennel, puppy, pet, bitch, terrier, rottweiler, canine, cat, to bark, Alsatian*

Which word list is better? Every item on the $c = 2$ is an animal, while the $c = 30$ list starts with *kennel* and even includes the verb *to bark*.

Estimation procedure

15.3 Distributional semantics

Local distributional statistics: Brown clusters

One way to use context is to perform word clustering. This can improve the performance of downstream (supervised learning) tasks, because even if a word is not observed in any labeled instances, other members of its clusters might be. The Brown et al. (1992) clustering algorithm provides one way to do this. The algorithm is over 20 years old and is still widely used in NLP; for example, Owoputi et al. (2012) use it to obtain large improvements in Twitter part-of-speech tagging.³

In Brown clustering, the context is just the immediately adjacent words. The similarity metric is built on a generative probability model:

- Assume each word w has a class $C(w)$
- Assume a generative model $\log p(w) = \sum_i \log p(w_i | c_i) + \log p(c_i | c_{i-1})$
(What does this remind you of?)

The word clusters $C(w)$ are not observed; our goal is to infer them from data. Now, in this model, we assume that,

$$p(w_i | c_i) = \begin{cases} \frac{\text{count}(w_i)}{\text{count}(c_i)}, & c_i = C(w_i) \\ 0, & \text{otherwise.} \end{cases} \quad (15.2)$$

This means that each word **type** has a single cluster — unlike in hidden Markov models, where a given word might be generated from multiple tags. Due to this constraint, we

²The example is from lecture slides by Marco Baroni, Alessandro Lenci, and Stefan Evert, who applied latent semantic analysis to the British National Corpus. You can play with an online demo here: <http://clac.cimc.unitn.it/infomap-query/>

³You can download Brown clusters at <http://metaoptimize.com/projects/wordreprs/>.

Algorithm 10 The bottom-up Brown et al. (1992) clustering algorithm

$\forall w, C(w) = w$ (start with every word in its own cluster)

while all clusters not merged **do**

 merge the c_i and c_j to maximize clustering quality.

Each word is described by a bitstring representation of its merge path




Figure 15.2: A small subtree produced by bottom-up Brown clustering

will not apply the expectation maximization algorithm which was used in unsupervised hidden markov model learning (§ 6.6). Instead, Brown et al. (1992) use a hierarchical clustering algorithm, shown in Algorithm 10. This is a **bottom-up** clustering algorithm, in that every word begins in its own cluster, and then clusters are merged until everything is clustered together. The series of merges taken by the algorithm is called a **dendrogram**, and it looks like a tree. For example, if the words *bike* and *bicycle* are first merged with each other, and then the cluster was merged with another cluster containing just the word *tricycle*, we would have the small tree shown in Figure 15.2.

For any desired number of clusters K , we can get a clustering by “cutting” the tree at some height. But in Brown clustering, we are usually interested not only in the resulting clusters from some cut of the merge tree, but also in the bitstrings that represent the series of mergers that led to the final clustering. A classical approach to semi-supervised learning is to use Brown bitstring prefixes in place of (or in addition to) lexical features, thus generalizing to words that are unseen in labeled data. The bitstrings for Figure 15.2 would be 0 for *tricycle*, 10 for *bicycle*, and 11 for *bike*. Subtrees from Brown clustering on a larger dataset are shown in Figure 15.3. The examples are drawn from a paper by Miller et al. (2004), who use Brown cluster bitstring prefixes as features for named entity recognition; this approach has also been used in dependency parsing (Koo et al., 2008) and in Twitter part-of-speech tagging (Owoputi et al., 2012).

The complexity of Algorithm 10 is $\mathcal{O}(V^3)$, where V is the size of the vocabulary. We are merging V clusters, since we start off with each word in its own cluster; each merger involves searching over $\mathcal{O}(V^2)$ pairs of clusters, to find the pair that maximizes the improvement in clustering quality. Cubic complexity is too slow for practical purposes, so we will explore a faster approximate algorithm later.

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 15.3: Brown subtrees from Miller et al. (2004)

Brown clusters and mutual information

We now explore the Brown clustering algorithm more mathematically, and then derive a more efficient clustering algorithm. First, some notation:

- \mathcal{V} is the set of all words.
- N is number of observed word tokens.
- $C : \mathcal{V} \rightarrow \{1, 2, \dots, k\}$ defines a partition of words into k classes.
- $\text{count}(w)$ is the number of times we see word $w \in \mathcal{V}$. This function can also be used to count classes.
- $\text{count}(w, v)$ is the number of times w immediately precedes v . This function can also be used to count class bigrams.

$$p(w_1, w_2, \dots, w_N; C) = \prod_m p(w_m | C(w_m)) p(C(w_m) | C(w_{m-1}))$$

$$\log p(w_1, w_2, \dots, w_N; C) = \sum_m \log p(w_m | C(w_m)) \times p(C(w_m) | C(w_{m-1}))$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

This is kind of like a hidden Markov model, but each word can only be produced by a single cluster. Now let's define the "quality" of a clustering as the average log-likelihood:

$$\begin{aligned}
J(C) &= \frac{1}{N} \sum_m \log (p(w_m | C(w_m)) \times p(C(w_m) | C(w_{m-1}))) \\
&= \sum_{w, w'} \frac{n(w, w')}{N} \log (p(w' | C(w')) \times p(C(w') | C(w))) && \text{sum over word types instead} \\
&= \sum_{w, w'} \frac{n(w, w')}{N} \log \left(\frac{n(w')}{n(C(w'))} \times \frac{n(C(w), C(w'))}{n(C(w))} \right) && \text{definition of probabilities} \\
&= \sum_{w, w'} \frac{n(w, w')}{N} \log \left(\frac{n(w')}{1} \times \frac{n(C(w), C(w'))}{n(C(w)) \times n(C(w'))} \times \frac{N}{N} \right) && \text{re-arrange, multiply by one} \\
&= \sum_{w, w'} \frac{n(w, w')}{N} \log \left(\frac{n(w')}{N} \times \frac{n(C(w), C(w')) \times N}{n(C(w)) \times n(C(w'))} \right) && \text{re-arrange terms} \\
&= \sum_{w, w'} \frac{n(w, w')}{N} \log \frac{n(w')}{N} + \frac{n(w, w')}{N} \log \left(\frac{n(C(w), C(w')) \times N}{n(C(w)) \times n(C(w'))} \right) && \text{distribution through log} \\
&= \sum_{w'} \frac{n(w')}{N} \log \frac{n(w')}{N} + \sum_{c, c'} \frac{n(c, c')}{N} \log \left(\frac{n(c, c') \times N}{n(c) \times n(c')} \right) && \text{sum across bigrams and classes} \\
&= \sum_{w'} p(w') \log p(w') + \sum_{c, c'} p(c, c') \log \frac{p(c, c')}{p(c) \times p(c')} && \text{multiply by } \frac{N^{-2}}{N^{-2}} \text{ inside log} \\
&= -H(W) + I(C)
\end{aligned}$$

The last step uses the following definitions from information theory:

Entropy The entropy of a discrete random variable is the expected negative log-likelihood,

$$H(X) = -E[\log P(X)] = -\sum_x P(X = x) \log P(X = x). \quad (15.3)$$

For example, for a fair coin we have $H(X) = \frac{1}{2} \log \frac{1}{2} + \frac{1}{2} \log \frac{1}{2} = -\log 2$; for a (virtually) certain outcome, we have $H(x) = 1 \times \log 1 + 0 \times \log 0 = 0$. We have already seen entropy in a few other contexts.

Mutual information The information shared by two random variables is the mutual information,

$$I(X; Y) = \sum_{y \in Y} \sum_{x \in X} p_{X,Y}(x, y) \log \left(\frac{p_{X,Y}(x, y)}{p_X(x) p_Y(y)} \right). \quad (15.4)$$

For example, if X and Y are independent, then $p_{X,Y}(x, y) = p_X(x) p_Y(y)$, so the mutual information is $\log 1 = 0$. In

(c) Jacob Eisenstein 2014-2017. Work in progress.

Algorithm 11 Exchange clustering algorithm

For K most frequent words, set $C_i = i$.
for $i = (m + 1) : V$ **do**
 Set $C_i = K + 1$
 Let $\langle c, c' \rangle$ be the two clusters whose merger minimizes the decrease in $I(C)$
 Merge c and c'

By $I(C)$, we are using a shorthand for the mutual information of adjacent word classes, $\langle C_{m-1}, C_m \rangle$,

$$I(C) = \sum_{C_m=c, C_{m-1}=c'} P(C_m = c, C_{m-1} = c') \log \left(\frac{P(C_m = c, C_{m-1} = c')}{P(C_m = c)P(C_{m-1} = c')} \right) \quad (15.5)$$

The entropy $H(W)$ does not depend on the clustering, so this term is constant; choosing a clustering with maximum mutual information $I(C)$ is equivalent to maximizing the log-likelihood. Now let's see how to do that efficiently.

 $V \log V$ approximate algorithm

With this model in hand, we can now define a more efficient algorithm, shown in Algorithm 11. The algorithm keeps exactly K clusters at every point in time, so the merger operation requires considering only $\mathcal{O}(K^2)$ clusters. We have to pass over the entire vocabulary once for a cost of $\mathcal{O}(V)$, but more importantly, we must sort the words by frequency, for a cost of $\mathcal{O}(V \log V)$, giving a total cost of $\mathcal{O}(V \log V + VK^2)$.

Syntactic distributional statistics

Local context is contingent on syntactic decisions that may have little to do with semantics:

(15.1) *I gave Tim the ball.*

(15.2) *I gave the ball to Tim.*

(You may recall from § 14.1 that this is the **dative alternation**.) Using the syntactic structure of the sentence might give us a more meaningful context, yielding better clusters.

There are several examples of this idea in practice. Pereira et al. (1993) cluster nouns based on the verbs for which they are the direct object: the context vector for each noun is **the count of occurrences as a direct object of each verb**. As with Brown clustering, they

(c) Jacob Eisenstein 2014-2017. Work in progress.

employ a class-based probability model:

$$\hat{p}(n, v) = \sum_{c \in \mathcal{C}} p(v | c) \times p(c, n) \quad (15.6)$$

$$= \sum_{c \in \mathcal{C}} p(v | c) \times p(n | c) \times p(c), \quad (15.7)$$

where n is the noun, v is the verb, and c is the class of the noun. They maximize the likelihood under this model using an iterative algorithm similar to expectation maximization (chapter 4).

Lin (1998) extends this idea from nouns to all words, using context statistics based on the incoming dependency edges. For any pair of words i and j and relation r , we can compute:

$$p(i, j | r) = \frac{n(i, j, r)}{\sum_{i', j'} n(i', j', r)} \quad (15.8)$$

$$p(i | r) = \sum_j p(i, j | r) \quad (15.9)$$

Now, let $T(i)$ be the set of pairs $\langle j, r \rangle$ such that $p(i, j | r) > p(i | r) \times p(j | r)$: then $T(i)$ contains words j that are especially likely to be joined with word i in relation r . Similarity between u and v can be defined through $T(u)$ and $T(v)$.

Lin considers several similarity measures for $T(u)$ and $T(v)$. Many of these are used widely in other contexts (usually for comparing clusterings or other sets), and are worth knowing about:

Cosine similarity $\frac{|T(u) \cap T(v)|}{\sqrt{|T(u)| |T(v)|}}$

Dice similarity $\frac{2 \times |T(u) \cap T(v)|}{|T(u)| + |T(v)|}$

Jaccard similarity $\frac{|T(u) \cap T(v)|}{|T(u)| + |T(v)| - |T(u) \cap T(v)|}$

However, Lin's chosen metric is more complex than any of these well-known alternatives:

$$\frac{\sum_{\langle r, w \rangle \in T(u) \cup T(v)} I(u, r, w) + I(v, r, w)}{\sum_{\langle r, w \rangle \in T(u)} I(u, r, w) + \sum_{\langle r, w \rangle \in T(v)} I(v, r, w)}, \quad (15.10)$$

where $I(u, r, w)$ is the mutual information between u and w , conditioned on r .

Results of the algorithm are shown in Figure 15.4. An interesting point in these results is that while many of the pairs are indeed synonyms, some have the **opposite** meaning. This is particularly evident for the adjectives, with pairs like *good/bad* and *high/low* at the top. It's useful to think about why this might be the case, and how you might fix it.

Lin's algorithm was also evaluated on its ability to match synonym pairs in human-generated thesauri. Its measure of text similarity was a better matched to WordNet than was the (human-written) Roget thesaurus!

(c) Jacob Eisenstein 2014-2017. Work in progress.

Nouns			Adjective/Adverbs		
Rank	Respective Nearest Neighbors	Similarity	Rank	Respective Nearest Neighbors	Similarity
1	earnings profit	0.572525	1	high low	0.580408
11	plan proposal	0.47475	11	bad good	0.376744
21	employee worker	0.413936	21	extremely very	0.357606
31	battle fight	0.389776	31	deteriorating improving	0.332664
41	airline carrier	0.370589	41	alleged suspected	0.317163
51	share stock	0.351294	51	clerical salaried	0.305448
61	rumor speculation	0.327266	61	often sometimes	0.281444
71	outlay spending	0.320535	71	bleak gloomy	0.275557
81	accident incident	0.310121	81	adequate inadequate	0.263136
91	facility plant	0.284845	91	affiliated merged	0.257666
101	charge count	0.278339	101	stormy turbulent	0.252846
111	baby infant	0.268093	111	paramilitary uniformed	0.246638
121	actor actress	0.255098	121	sharp steep	0.240788
131	chance likelihood	0.248942	131	communist leftist	0.232518
141	catastrophe disaster	0.241986	141	indoor outdoor	0.224183
151	fine penalty	0.237606	151	changed changing	0.219697
161	legislature parliament	0.231528	161	defensive offensive	0.211062
171	oil petroleum	0.227277	171	sad tragic	0.206688
181	strength weakness	0.218027	181	enormously tremendously	0.199936
191	radio television	0.215043	191	defective faulty	0.193863
201	coupe sedan	0.209631	201	concerned worried	0.186899

Figure 15.4: Similar word pairs from the clustering method of Lin (1998)

15.4 Distributed representations

Distributional semantics are computed from context statistics. **Distributed** semantics are a related but distinct idea: that meaning is best represented by numerical vectors rather than discrete combinatoric structures. Distributed representations are often distributional: this section will focus on latent semantic analysis and word2vec, both of which are distributed representations that are based on distributional statistics. However, distributed representations need not be distributional: for example, they can be learned in a supervised fashion from labeled data, as in the sentiment analysis work of Socher et al. (2013b).

Latent semantic analysis

Thus far, we have considered context vectors that are large and sparse. We can arrange these vectors into a matrix $\mathbf{X} \in \mathbb{R}^{V \times N}$, where rows correspond to words and columns correspond to contexts. However, for rare words i and j , we might have $\mathbf{x}_i^\top \mathbf{x}_j = 0$, indicating zero counts of shared contexts. So we'd like to have a more robust representation.

We can obtain this by factoring $\mathbf{X} \approx \mathbf{U}_K \mathbf{S}_K \mathbf{V}_K^\top$, where

$$\mathbf{U}_K \in \mathbb{R}^{V \times K}, \quad \mathbf{U}_K \mathbf{U}_K^\top = \mathbb{I} \quad (15.11)$$

$$\mathbf{S}_K \in \mathbb{R}^{K \times K}, \quad \mathbf{S}_K \text{ is diagonal, non-negative} \quad (15.12)$$

$$\mathbf{V}_K \in \mathbb{R}^{D \times K}, \quad \mathbf{V}_K \mathbf{V}_K^\top = \mathbb{I} \quad (15.13)$$

Here K is a parameter that determines the fidelity of the factorization; if $K = \min(V, N)$, then $\mathbf{X} = \mathbf{U}_K \mathbf{S}_K \mathbf{V}_K^\top$. Otherwise, we have

$$\mathbf{U}_K, \mathbf{S}_K, \mathbf{V}_K = \underset{\mathbf{U}_K, \mathbf{S}_K, \mathbf{V}_K}{\operatorname{argmin}} \|\mathbf{X} - \mathbf{U}_K \mathbf{S}_K \mathbf{V}_K^\top\|_F, \quad (15.14)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

subject to the constraints above. This means that $\mathbf{U}_K, \mathbf{S}_K, \mathbf{V}_K$ give the rank- K matrix $\tilde{\mathbf{X}}$ that minimizes the Frobenius norm, $\sqrt{\sum_{i,j} (x_{i,j} - \tilde{x}_{i,j})^2}$.

This factorization is called the **Truncated Singular Value Decomposition**, and is closely related to eigenvalue decomposition of the matrices $\mathbf{X}\mathbf{X}^\top$ and $\mathbf{X}^\top\mathbf{X}$. In general, the complexity of SVD is $\min(\mathcal{O}(D^2V), \mathcal{O}(V^2N))$. The standard library LAPACK (Linear Algebra PACKage) includes an iterative optimization solution for SVD, and (I think) this what is called by Matlab and Numpy.

However, for large sparse matrices it is often more efficient to take a stochastic gradient approach. Each word-context observation $\langle w, c \rangle$ gives a gradient on $\mathbf{u}_w, \mathbf{v}_c$, and \mathbf{S} , so we can take a gradient step. This is part of the algorithm that was used to win the Netflix challenge for predicting movie recommendation — in that case, the matrix includes raters and movies (Koren et al., 2009).

Return to NLP applications, the slides provide a nice example from Deerwester et al. (1990), using the titles of computer science research papers. In the example, the context-vector representations of the terms *user* and *human* have negative correlations, yet their distributional representations have high correlation, which is appropriate since these terms have roughly the same meaning in this dataset.

Word vectors and neural word embeddings

Discriminatively-trained word embeddings very hot area in NLP. The idea is to replace factorization approaches with discriminative training, where the task may be to predict the word given the context, or the context given the word.

Suppose we have the word w and the context c , and we define

$$u_\theta(w, c) = \exp(\mathbf{a}_w^\top \mathbf{b}_c) \quad (15.15)$$

$$(15.16)$$

with $\mathbf{a}_w \in \mathbb{R}^K$ and $\mathbf{b}_c \in \mathbb{R}^K$. The vector \mathbf{a}_w is then an **embedding** of the word w , representing its properties. We are usually less interested in the context vector \mathbf{b} ; the context can include surrounding words, and the vector \mathbf{b}_c is often formed as a sum of context embeddings for each word in a window around the current word. Mikolov et al. (2013a) draw the size of this context as a random number r .

The popular word2vec software⁴ uses these ideas in two different types of models:

Skipgram model In the skip-gram model (Mikolov et al., 2013a), we try to maximize the log-probability of the context,

⁴<https://code.google.com/p/word2vec/>

$$J = \frac{1}{M} \sum_m \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{m+j} \mid w_m) \quad (15.17)$$

$$p(w_{m+j} \mid w_m) = \frac{u_\theta(w_{m+j}, w_m)}{\sum_{w'} u_\theta(w', w_m)} \quad (15.18)$$

$$= \frac{u_\theta(w_{m+j}, w_m)}{Z(w_m)} \quad (15.19)$$

This model is considered to be slower to train, but better for rare words.

CBOW The continuous bag-of-words (CBOW) (Mikolov et al., 2013b,c) is more like a language model, since we predict the probability of words given context.

$$J = \frac{1}{M} \sum_m \log p(w_m \mid c) \quad (15.20)$$

$$= \frac{1}{M} \sum_m \log u_\theta(w_m, c) - \log Z(c) \quad (15.21)$$

$$u_\theta(w_m, c) = \exp \left(\sum_{-c \leq j \leq c, j \neq 0} \mathbf{a}_{w_m}^\top \mathbf{b}_{w_{m+j}} \right) \quad (15.22)$$

The CBOW model is faster to train (Mikolov et al., 2013a). One efficiency improvement is build a Huffman tree over the vocabulary, so that we can compute a hierarchical version of the softmax function with time complexity $\mathcal{O}(\log V)$ rather than $\mathcal{O}(V)$. Mikolov et al. (2013a) report two-fold speedups with this approach.

The recurrent neural network language model (§ 5.3) is still another way to compute word representations. In this model, the context is summarized by a recurrently-updated state vector $\mathbf{c}_m = f(\Theta \mathbf{c}_{m-1} + \mathbf{U} \mathbf{x}_m)$, where $\Theta \in \mathbb{R}^{K \times K}$ defines a the recurrent dynamics, $\mathbf{U} \in \mathbb{R}^{K \times V}$ defines “input embeddings” for each word, and $f(\cdot)$ is a non-linear function such as tanh or sigmoid. The word distribution is then,

$$P(W_{m+1} = i \mid \mathbf{c}_m) = \frac{\exp(\mathbf{c}_m^\top \mathbf{v}_i)}{\sum_{i'} \exp(\mathbf{c}_m^\top \mathbf{v}_{i'})}, \quad (15.23)$$

where \mathbf{v}_i is the “output embedding” of word i .

(c) Jacob Eisenstein 2014-2017. Work in progress.

*Estimating word embeddings

[[todo: link to rnnlm, show pictures](#)] Training word embedding models can be challenging, because they require probabilities that need to be normalized over the entire vocabulary. This implies a training time complexity of $\mathcal{O}(VK)$ for each instance. Since these models are often trained on hundreds of billions of words, with $V \approx 10^6$ and $K \approx 10^3$, this cost is too high. Estimation techniques eliminate the factor V by making approximations.

One such approximation is negative sampling, which is a heuristic variant of noise-contrastive estimation (Gutmann and Hyvärinen, 2012).

We introduce an auxiliary variable D , where

$$D = \begin{cases} 1, & w \text{ is drawn from the empirical distribution } \hat{p}(w | c) \\ 0, & w \text{ is drawn from the noise distribution } q(w) \end{cases} \quad (15.24)$$

Now we will optimize the objective

$$\sum_{(w,c) \in \mathcal{D}} \log P(D = 1 | c, w) + \sum_{i=1, w' \sim q}^k \log P(D = 0, | c, w'), \quad (15.25)$$

setting

$$P(D = 1 | c, w) = \frac{u_\theta(w, c)}{u_\theta(w, c) + k \times q(w)} \quad (15.26)$$

$$P(D = 0 | c, w) = 1 - P(D = 1 | c, w) \quad (15.27)$$

$$= \frac{k \times q(w)}{u_\theta(w, c) + k \times q(w)}, \quad (15.28)$$

where k is the number of noise samples. Note that we have dropped the normalization term $\sum_{w'} u_\theta(w', c)$. Gutmann and Hyvärinen (2012) show that it is possible to treat the normalization term as an additional parameter z_c , which can be directly estimated (see also Vaswani et al., 2013). Andreas and Klein (2015) go one step further, setting $z_c = 1$, in what has been called a “self-normalizing” probability distribution. This might be trouble if we were trying to directly maximize $\log p(w | c)$, but this is where the auxiliary variable formulation helps us out: if we set θ such that $\sum_{w'} u_\theta(w' | c) \gg 1$, we will get a very low probability for $P(D = 0)$. [[todo: needs a little more explanation](#)]

We can further simplify by setting $k = 1$ and $q(w)$ to a uniform distribution, arriving at

$$P(D = 1 | c, w) = \frac{u_\theta(w, c)}{u_\theta(w, c) + 1} \quad (15.29)$$

$$P(D = 0 | c, w) = \frac{1}{u_\theta(w, c) + 1} \quad (15.30)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

The derivative with respect to a is obtained from the objective

$$L = \sum_m \log p(D = 1 \mid c_m, w_m) + \log p(D = 0 \mid c, w') \quad (15.31)$$

$$= \sum_m \log u_\theta(w_m, c_m) - \log(1 + u_\theta(w_m, c_m)) - \log(1 + u_\theta(w', c_m)) \quad (15.32)$$

$$\frac{\partial L}{\partial \mathbf{a}_i} = \sum_{m:w_m=i} \mathbf{b}_{c_m} - \frac{1}{1 + u_\theta(w_m, c_m)} \frac{\partial u_\theta(i, c_m)}{\partial \mathbf{a}_i} + \sum_m \frac{q(i)}{1 + u_\theta(i, c_m)} \frac{\partial u_\theta(i, c_m)}{\partial \mathbf{a}_i} \quad (15.33)$$

$$= \sum_{m:w_m=i} \mathbf{b}_{c_m} - P(D = 1 \mid w_m = i, c_m) \mathbf{b}_{c_m} - \sum_m q(i) P(D = 0 \mid i, c_m) \mathbf{b}_{c_m} \quad (15.34)$$

$$= \sum_m (\delta(w_m = i) - q(i)) P(D = 0 \mid w_m = i, c_m) \mathbf{b}_{c_m}. \quad (15.35)$$

The gradient with respect to \mathbf{b} is similar. In practice, we simply sample w' at each instance and compute the update with respect to \mathbf{a}_{w_m} and $\mathbf{a}_{w'}$. In practice, AdaGrad performs well for this optimization.

*Connection to matrix factorization

Recent work has drawn connections between this procedure for training the skip-gram model and weighted matrix factorization approaches (Pennington et al., 2014; Levy and Goldberg, 2014b). For example, Levy and Goldberg (2014b) show that skip-gram with negative sampling is equivalent to factoring a matrix X , where

$$X_{i,j} = PMI(W = i, C = j) - \log k, \quad (15.36)$$

where k is a constant offset equal to the number of negative samples drawn in Equation 15.25, and PMI is the **pointwise mutual information** of the events of the word $W = i$ and the context $C = j$,

$$PMI(W = i, C = j) = \log \frac{P(W = i, C = j)}{P(W = i)P(C = j)} \quad (15.37)$$

$$= \log \frac{n(W = i, C = j)}{M} \frac{M}{n(W = i)} \frac{M}{n(C = j)} \quad (15.38)$$

$$= \log \frac{n(W = i, C = j)}{n(W = i)} \frac{M}{n(C = j)}. \quad (15.39)$$

Word embeddings can be obtained by solving the truncated singular value decomposition $\mathbf{U}\Sigma\mathbf{V}^\top = \mathbf{X}$, setting the embedding of word i to $\mathbf{u}_i\sqrt{(\Sigma_{i,i})}$.

This connection suggests that the differences between recent work on neural word embeddings and much older work on Latent Semantic Analysis may be smaller than they initially seemed! Online learning approaches such as negative sampling stream over

(c) Jacob Eisenstein 2014-2017. Work in progress.

data, and require hyperparameter tuning to set the appropriate learning rate. On the other hand, PMI is undefined for word-context pairs that are unobserved (due to the logarithm of zero), requiring a heuristic solution such as positive PMI, $PPMI(i, j) = \max(0, PMI(i, j))$, or shifted positive PPMI $SPPMI_k(i, j) = \max(0, PMI(i, j) - \log k)$. Levy and Goldberg (2014b) find that singular value decomposition on shifted positive PMI does better than skipgram negative sampling on some lexical semantic tasks, but worse on others.

Chapter 16

Reference Resolution

References are one of the most noticeable forms of linguistic ambiguity, afflicting not just automated natural language processing systems, but also fluent human readers. For this reason, warnings to avoid “ambiguous pronouns” are ubiquitous in manuals and tutorials on writing style. But referential ambiguity is not limited to pronouns, as shown in the following text:

(16.1) *Apple Inc Chief Executive Tim Cook has jetted into China for talks with government officials as **he**₁ seeks to clear up a pile of problems in **[[the firm’s]₂ biggest growth market]**₃... **[Cook]**₄ is on **[his]**₅ first trip to **[the country]**₆ since taking over...*¹

Each of the bolded substrings in the passage refers to an entity that is introduced earlier in the story. These references include the pronouns *he* and *his*, but also the shortened name *Cook*, and most challengingly, **nominals** such as *the firm* and *the firm’s biggest growth market*. Only by resolving several of these references can we reach the (correct) inference that China is Apple’s biggest growth market.

The task of reference resolution is often broken into two components:

- **Coreference resolution**, which is the task of linking spans of text such as the *the firm* to other spans, such as *Apple Inc*. A subset of coreference resolution is the task of **anaphora resolution**, which typically involves resolving only pronoun anaphora such as *he* and *her*.
- **Entity linking**,² which is the task of linking spans of text to entities in a knowledge base. This step is a prerequisite for the model-based semantic parsing that was considered in chapter 13.

¹<http://www.reuters.com/article/us-apple-china-idUSBRE82Q06420120327>, retrieved on March 26, 2017

²Amusingly, there are many names for this task: deduplication, approximate string match, entity clustering, record linking, multidocument coreference resolution, etc.

These tasks have traditionally been distinguished because they seem to require different sorts of knowledge to perform, and different resources to evaluate. As we will see, coreference resolution — especially anaphora resolution — is constrained by syntax and by compatibility of attributes such as gender, number, and animacy. Coreference resolution can be evaluated by comparing against a ground truth that is specified at the document level, without reference to any external information. In contrast, solving entity linking requires making inferences about name compatibility and about semantic properties of each entity — although these inferences are sometimes necessary for coreference resolution too. Evaluation of entity linking requires linking textual references to some predefined external ontology. Of the two tasks, research on coreference resolution is more mature, and will therefore be the focus of this chapter. Approaches to entity linking and related tasks are summarized in ??.

16.1 Forms of referring expressions

The three main forms of referring expressions — pronouns, names, and nominals — each pose unique challenges for the reader. As a coarse-grained summary, pronouns are constrained by syntax and semantic attributes; name references constrained by rules for matching; nominals are linked by world knowledge.

Pronouns

Pronouns are a closed class of words that are used for references. A natural way to think about pronoun resolution is what Kehler (2007) calls the SMASH approach:


- Search for candidate referents;
- Match against hard agreement constraints;
- And Select using Heuristics, which are “soft” constraints such as recency and parallelism.

In the search step, candidates are identified from the preceding text or speech.³ In models such as **centering theory**, any entity that has previously been **evoked** can be **accessed** in any subsequent unit of text (Grosz et al., 1995). However, cognitive constraints may imply that entities which have not been mentioned recently are unlikely to be accessed without be re-introduced, and correspondingly, computational constraints may encourage algorithms to consider only referents that are relatively recent.

³Pronouns whose referents come later are known as **cataphora**, e.g.,

(16.1) *Many years later, as **he** faced the firing squad, Colonel Aureliano Buendía was to remember that distant afternoon when his father took him to discover ice.*

(This is the first sentence of *One Hundred Years of Solitude*, by Gabriel García Márquez.)

Figure 16.1: *Mary* c-commands *her/herself*

Matching constraints for pronouns

Semantic constraints include morphologically marked information such as number, person, gender, and animacy.

- (16.2) Tim Cook has jetted in for talks with officials as **he** seeks to clear up a pile of problems...

We can identify the following features of the pronoun and possible referents:

- Number(*he*) = singular
- Number(*officials*) = plural
- Number(*Tim Cook*) = singular

Since there are no other possible referents, *he* almost certainly refers back to *Tim Cook*.

Other features include person, gender, and animacy, as in the following examples:

- (16.3) *Sally met my brother. He charmed her.*
 (16.4) *Sally met my brother. She charmed him.*
 (16.5) **We₁ told them₁ not to go.*
 (16.6) *Putin brought a bottle of vodka. It was from Russia.*

Another source of constraints comes from syntax. To understand these constraints, it is helpful to introduce some linguistic terminology:

- *x* **c-commands** *y* iff the first branching node above *x* also dominates *y*;
- *x* **binds** *y* iff *x* and *y* are co-indexed and *x* c-commands *y*;
- if *y* is not bound, it is **free**.

For example, consider the tree in Figure 16.1. In this example, *Mary* c-commands *her/herself*, because the first branching node above *Mary* also dominates *her/herself*. However, *her/herself* does not c-command *Mary*. Thus, the pronoun *her* **cannot** refer to *Mary*,




Figure 16.2: *Mary* does not c-command *her*, but *Mary's mom* does.




Figure 16.3: The scope of *Abigail* is limited by the S non-terminal. Either *she* or *her* (but not both) can bind to *Abigail*.

because non-reflexive pronouns cannot refer to antecedents that c-command them. On the other hand, the reflexive *herself* **must** refer to *Mary*.

Now consider the example, shown in Figure 16.2. Here, *Mary* does **not** c-command *her*, but *Mary's mom* c-commands *her*. Thus, *her* **can** refer to *Mary* — and we cannot use reflexive *herself* in this context, unless we are talking about *Mary's mom*. However, *her* does not have to refer to *Mary*.

A more complex example is shown in Figure 16.3. This indicates how the constraints defined here have a limited domain. [todo: explain how this is limited] The pronoun *she* can refer to *Abigail*, because *Abigail* is outside the domain of *she*. Similarly, *her* can also refer to *Abigail*. But *she* and *her* cannot be coreferent.

Heuristics

After applying constraints, there will be a number of candidate referents for each pronoun. In the SMASH paradigm, heuristics are then applied to compare among these possibilities.

Recency is a particularly strong heuristic. All things equal, readers will prefer the more recent referent for a given pronoun, particularly when comparing referents that occur in different sentences. Jurafsky and Martin (2009) offer the following example:

(c) Jacob Eisenstein 2014-2017. Work in progress.

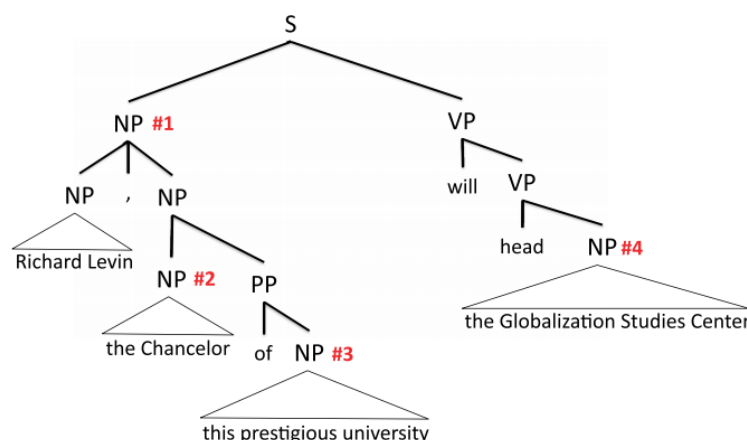


Figure 16.4: Left-to-right breadth-first tree traversal, proposed by Hobbs (1978), as implemented by Lee et al. (2013)

- (16.7) *The doctor found an old map in the captain's chest. Jim found an even older map hidden on the shelf. **It** described an island.*

Readers are expected to prefer the second, older map as the referent for the pronoun *it*.

However, subjects are often preferred over objects, and this can contradict the preference for recency when two candidate referents are in the same sentence. For example,

- (16.8) *Asha loaned Mei a book on Spanish. **She** is always trying to help people.*

Here, we may prefer to link *she* to *Asha* rather than *Mei*, because of *Asha*'s position in the subject role of the preceding sentence. (Arguably, this preference would be reversed if the second sentence were *She is visiting Argentina next month.*)

A third heuristic is parallelism:

- (16.9) *Asha loaned Mei a book on Spanish. Olya loaned **her** a book on Portuguese.*

Here *Mei* is preferred as the referent for *her*, contradicting the preference for the subject *Asha* in the preceding sentence.

Hobbs (1978) unifies recency and subject-role heuristics by traversing the document in a syntax-driven fashion: each preceding sentence is traversed breadth-first, left-to-right (Figure 16.4). In this way, *Asha* would be preferred as the referent for *she* in (16.8), while the older map would be preferred as the referent for *it* in *ex:coref-recency*. **Centering theory** offers an alternative unification of recency and syntactic prominence, maintaining ordered lists of candidates referents throughout the text or discourse (Grosz et al., 1995). Centering can also be viewed as a generative model, in that it predicts the form of the referring expression that will be used for each entity reference in a sentence.

In early work on reference resolution, Lappin and Leass (1994) set weights for a half-dozen syntactic preferences by hand, choosing the referent with the highest overall weight. More recent work uses machine learning approaches to quantify the importance of each of these factors, as discussed in ???. However, pronoun resolution often cannot be performed successfully using syntactic heuristics alone. This is shown by the classic example pair:

(16.10) *The **city council** denied the protesters a permit because **they** feared violence.*

(16.11) *The city council denied **the protesters** a permit because **they** advocated violence.*⁴

Non-referential pronouns

While pronouns are generally used to refer to things, they need not refer to entities, as shown in the following examples:

(16.12) *They told me that I was too ugly, but I didn't believe **it**.*

(16.13) *Alice saw Bob get angry, and I saw **it** too.*

(16.14) *They told me that I was too ugly, but **that** was a lie.*

(16.15) *Jess said she worked in security.
I suppose **that's** one way to put it.*

Pronouns may also have **generic** referents, meaning that they do not refer to entities in any model, but rather, to possible entities:

(16.16) *A good father takes care of **his** kids.*

(16.17) *On the moon, **you** have to carry **your** own oxygen.*

(16.18) *Every farmer who owns a donkey beats it. (Geach, 1962)*

Finally, pronouns need not refer to anything at all:

(16.19) ***It's** raining.*

(16.20) ***It's** crazy out there.*

(16.21) ***It's** money that she's really after.*

(16.22) ***It** sucks that we have to work so hard.*

In the first two examples above, *it* is **pleonastic**; the third and fourth examples are **cleft** and **extraposition**. How can we automatically distinguish these usages of *it* from referential pronouns? Bergsma et al. (2008) propose a substitutability test. Consider the difference between the following two examples:

(16.23) *You can make **it** in advance.*

(16.24) *You can make **it** in showbiz.*

⁴This pair is attributed to Winograd (1972), but I downloaded that article and didn't find it.

In the second example, the pronoun **it** is non-referential. One way to see this is by substituting another pronoun, like **them**, into these examples:

(16.25) *You can make them in advance.*

(16.26) *?You can make them in showbiz.*

The questionable grammaticality of the second example suggests that **it** cannot be referential. Bergsma et al. (2008) operationalize this idea by comparing distributional statistics for 5-grams around the word *it*, testing how often other pronouns or nouns ever appear in the same position as *it*. In cases where other pronouns are frequent, the *it* is likely referential.

Proper Names

If a proper name is used as a referring expression, it often refers to another proper name, so that the coreference problem is simply to determine whether the two names match. Subsequent proper name references often use a shortened form, as in the running example:

(16.27) *Apple Inc Chief Executive **Tim Cook** has jetted into China ... **Cook** is on his first business trip to the country since taking over ...*

In this news article, the family name *Cook* is used as a referring expression; in informal conversation, it might be more typical to use the given name *Tim*, while more formal venues, such as *The Economist*, would use the title form *Mr Cook*.

Thus, exact match is unlikely to identify many proper name references. A typical solution is to match the syntactic **head words** of the reference with the referent. Recall that the head word of a phrase can be identified by applying head percolation rules to the phrasal parse tree (chapter 11); alternatively, the head can be identified as the root of the dependency subtree covering the name. For sequences of proper nouns, the head word will be the final word, which in the example is *Cook*.

While useful, there are a number of caveats to the practice of matching head words of proper names.

- In the European tradition, family names tend to be more specific than given names, and family names usually come last. However, other traditions have other practices: for example, in Chinese names, the family name typically comes first; in Japanese, honorifics come after the name, as in *Nobu-San* (*Mr. Nobu*).
- In many organization names, it is also the case that the head word is not the most informative: for example, *Georgia Tech* and *Virginia Tech* are distinguished by the modifiers *Virginia* and *Georgia*, and not the heads. This concern applies even when the referring expression is a substring of the candidate referent: *Lebanon* does not refer to the same entity as *Southern Lebanon*, and Lee et al. (2011) add a rule to deal with the specific case of geographical modifiers.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Finally, proper names can be nested, as in [*the CEO of [Microsoft]*]. Haghighi and Klein (2009) introduce a constraint to prevent nested noun phrases from being marked as coreferential.

Despite these difficulties, proper names are the easiest category of references to resolve (Stoyanov et al., 2009). In machine learning systems, one solution is to include a range of matching features, including exact match, head match, string inclusion, and even matching on “bags” of tokens, so that, e.g., *Tim Cook* matches *Cook, Tim* (Bontcheva et al., 2002). In addition to matching features, competitive systems include large lists, or **gazetteers**, of acronyms (e.g., *the National Basketball Association / NBA*), demonyms (e.g., *the Israelis / Israel*), and other aliases (e.g., *the Georgia Institute of Technology / Georgia Tech*).⁵ The learning algorithm can then determine the appropriate weights for each matching feature.

Nominals

In coreference resolution, noun phrases that are neither pronouns nor names are referred to as **nominals**. In the running example, nominal references include:

- *the firm (Apple Inc)*
- *the firm’s biggest growth market (China)*
- *the country (China).*

Nominals are generally more difficult to resolve than pronouns and names (Durrett and Klein, 2013, e.g.), and the examples above suggest why this may be the case: world knowledge is required to identify *Apple Inc* as a *firm*, and *China* as a *growth market*. Other difficult examples include the use of colloquial expressions, such as coreference between *Clinton transition officials* and *the Clinton camp* (Soon et al., 2001). But there are also cases that can be handled by surface features such as head word match: for example, *the tax cut bill* may be referenced later by *the Republican bill* or even *the bill*.

Attempts to use semantics to improve nominal coreference have met with limited success. Durrett and Klein (2013) employ WordNet synonymy and hypernymy relations on head words, named entity types (e.g., person, organization), and unsupervised clustering over nominal heads. These features give only limited improvement over simple baseline using surface features such as string match.

16.2 Learning for coreference resolution

Coreference resolution is a non-traditional learning problem, because it is not obvious what constitutes an “instance.” A number of proposals have been put forward. In dis-

⁵Lists of aliases were used heavily in the Message Understanding Conference (MUC) systems of the 1990s, which helped to define the coreference resolution task (Grishman and Sundheim, 1996). They are still used in some of the most competitive systems at the time of this writing (e.g. Martschat and Strube, 2015).

cussing these approaches, references and candidate referents are called **mentions**; chains of references are called **entities**.⁶ [todo: add figure] Ground truth annotations identify the entities. In our running example, this would be:

- *Apple Inc Chief Executive Tim Cook, he, Cook*
- *Apple Inc, the firm*
- *China, the firm's biggest growth market, the country.*

“Singleton” mentions (e.g., *government officials*) are annotated in the ACE Corpus (?), but not in the OntoNotes corpus (Hovy et al., 2006).

Coreference resolution can be viewed as a structure prediction problem, where the goal is to identify a set of coreference chains c among all possible coreference structures $\mathcal{C}(w)$,

$$\operatorname{argmax}_{c \in \mathcal{C}(w)} \theta \cdot f(c, w). \quad (16.1)$$

Each chain c_i consists of a set of mentions $\{m_j\}$. Typically it is the coreference resolution system’s job to identify the mentions from unannotated text, although systems are sometimes evaluated with “gold” mentions from the annotators.

The main approaches to coreference resolution can be distinguished by how they decompose the feature function $f(c, w)$. In mention-based models, features are defined over pairs of entities. This can facilitate inference, but mention-based models can suffer from incoherent entity chains, such as $\{Hillary Clinton \leftarrow Clinton \leftarrow Mr Clinton\}$. In entity-based models, the goal is to ensure that the entire entity chain is coherent. This can make inference more difficult, since the number of possible entity groupings is exponential in the number of mentions. A second distinction is whether the training instances are pairs (mention-mention pairs or mention-entity pairs), or whether learning is performed by ranking all possible candidates (mentions or entities) for a given mention.

Mention-pair and mention ranking models

In the **mention-pair model**, a label $y_{ij} \in \{0, 1\}$ is assigned to each pair of mentions $\langle i, j \rangle, i < j$. If i and j corefer, then $y_{ij} = 1$, and we say that i is the **antecedent** of j ; otherwise, $y_{ij} = 0$. Thus, the mention-pair model reduces coreference resolution to binary classification, enabling the application of off-the-shelf machine learning algorithms: Soon et al. (2001) use decision trees, and Bengtson and Roth (2008) use the averaged perceptron.

Under the constraint that each mention has at most one antecedent, the **antecedent structure** $\{y_{ij}\}$ induces a unique set of entities c . However, the converse is not true: a

⁶In many annotations, the term **markable** is used to refer to spans of text that can **potentially** mention an entity. The set of markables includes non-referential pronouns such as pleonastic *it*, which does not mention any entity. Part of the job of the coreference system is to avoid incorrectly linking these non-referential markables to any mention chains.

single set of entities c may be compatible with multiple antecedent structures. Since the ground truth annotations give c but not y , additional heuristics must be employed to convert the labeled data into training examples. Furthermore, we must impose the constraint that each mention have at most one antecedent. One solution is to pair the classifier with a search heuristic, based on SMASH: search backwards from j until finding an antecedent i which corefers with j with high confidence, and then stop searching. During training, for each reference j with antecedent i , we include negative examples $y_{i'j} = 0$ for all $i < i' < j$.

In **mention ranking**, the classifier learns to identify a single antecedent $a_i \in \{1, 2, \dots, i-1, i\}$ for each referring expression i , where $a_i = i$ indicates that the mention i does not refer to any previously-introduced entity (Denis and Baldridge, 2007). Specifically, the model chooses,

$$\hat{a}_i = \operatorname{argmax}_{a \in \{1, 2, \dots, i\}} \boldsymbol{\theta} \cdot \mathbf{f}(i, a, \mathbf{w}), \quad (16.2)$$

where $\mathbf{f}(i, a, \mathbf{w})$ defines a set of features on the mention pair $\langle i, a \rangle$. A special set of features can be employed for the case $a_i = i$, although later work on mention ranking has employed a two-stage model, in which an “anaphoricity” classifier determines whether the mention i refers to a previously defined entity; if so, then the ranking decision is performed over the set $1, 2, \dots, i-1$ (Denis and Baldridge, 2008).

As with the binary coreference variables $\{y_{ij}\}$, the antecedent variables $\{a_i\}$ relate to the entity chains in a many-to-one mapping: each set of assignment variables induces a single entity clustering, but an entity clustering can correspond to many different settings of assignment variables. When mention i has multiple possible antecedents in the clustering c , a typical approach is to select the most recent compatible antecedent. However, by using a probabilistic ranking model, $p(\{a_i\} \mid \mathbf{w})$, Durrett and Klein (2013) are able to sum over the set of all antecedent structures $\mathcal{A}(c)$ that are compatible with the gold coreference clustering c ,

$$p(a_i \mid i, \mathbf{w}) = \frac{\exp(\boldsymbol{\theta} \cdot \mathbf{f}(i, a_i, \mathbf{w}))}{\sum_{a' \leq i} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(i, a', \mathbf{w}))} \quad (16.3)$$

$$p(\mathbf{a} \mid \mathbf{w}) = \prod_i^M p(a_i \mid i, \mathbf{w}) \quad (16.4)$$

$$p(c \mid \mathbf{w}) = \sum_{\mathbf{a} \in \mathcal{A}(c)} p(\mathbf{a} \mid \mathbf{w}). \quad (16.5)$$

In this way, Durrett and Klein learn a model that tries to assign high scores for all valid antecedent structures.

Entity-based models

Many of the practical difficulties with mention-based models stem from the fact that they treat coreference resolution like a classification or ranking problem, when in fact it is a

(c) Jacob Eisenstein 2014-2017. Work in progress.

clustering problem: the goal is to group the mentions together into clusters that correspond to the underlying entities. Entity-based approaches attempt to identify these clusters directly. Such methods require defining features at the entity level, measuring whether the set of mentions are internally consistent. Cardie and Wagstaff (1999) provide an early example of entity-based coreference, incrementally merging clusters of mentions under the constraint that all pairs of mentions in the entity are compatible in number, gender, animacy, etc. They define a set of soft preferences for merging when there are multiple clusters that are compatible. More recent methods for entity-based coreference resolution have applied machine learning in the context of incremental search over the space of coreference clusterings (e.g., Clark and Manning, 2015).

The gap between entity-based and mention-pair models can be partially bridged by enforcing transitivity on the mention-pair variables: if $y_{ij} = 1$ and $y_{jk} = 1$, then $y_{ik} = 1$. This constraint can be written as a linear inequality,

$$\forall i < j < k, y_{ik} \geq y_{ij} + y_{jk} - 1. \quad (16.6)$$

The transitivity constraint ensures that each mention is linked to all antecedents in its cluster. We can then formulate the inference problem as,

$$\max \sum_{i,j} \theta \cdot f(y_{ij}, w) \quad (16.7)$$

$$s.t. \forall i < j < k, y_{ik} \geq y_{ij} + y_{jk} - 1. \quad (16.8)$$

In this formulation, features are still defined over mention pairs — rather than over entire entities — but transitivity ensures that all pairs in the cluster are compatible, avoiding incoherent clusters like $\{Hillary\ Clinton \leftarrow Clinton \leftarrow Mr\ Clinton\}$. However, this coherence comes at a computational price: the constrained optimization problem is NP-hard. **Integer linear programming** (ILP), which we saw in chapter 14, is one solution (Klenner, 2007; Finkel and Manning, 2008); correlational clustering is another (McCallum and Wellner, 2004).

[todo: discuss recent methods for using deep learning to acquire entity representations (??)]

Deterministic methods

Unlike many other areas of natural language processing, it is possible to build competitive systems for coreference resolution without machine learning (Haghighi and Klein, 2009). One such architecture is shown in Figure 16.5. The basic idea is to apply a series of rule-based methods, or “sieves”, starting with high-precision rules and progressively increasing recall. Each sieve builds on the output of its predecessor, so that it is possible to consider entity-level information. For example, in the case of $\{Hillary\ Clinton \leftarrow$

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 16.5: Architecture of Stanford's deterministic "multi-pass sieve" coreference system (Lee et al., 2013)

Clinton \leftarrow *she*}, the name-matching sieve would link *Clinton* and *Hillary Clinton*, and the pronoun-matching sieve would then link *she* to the combined cluster.

The Stanford deterministic system made a strong showing at 2011 CoNLL shared task on coreference, winning nearly every track in the competition (Pradhan et al., 2011). This was particularly surprising, given the dominance of non-deterministic methods based on machine learning in virtually all other areas of natural language processing. While learning-based systems have regained the upper hand in recent years (e.g., Björkelund and Kuhn, 2014; ?), the resurgence of deterministic, rule-based artificial intelligence in coreference resolution tells us that it may differ in important ways from other tasks, such as tagging and dependency parsing.

16.3 Entity linking and multi-document coreference resolution

[todo: later]

Exercises

1. The size of the largest entity typically grows linearly with the number of mentions in a document. Using this assumption, give an asymptotic estimate of the num-

(c) Jacob Eisenstein 2014-2017. Work in progress.

ber of antecedent structures that are compatible with a coreference clustering in a document with M mentions.

Part IV

Applications

Chapter 17

Information extraction

A fundamental challenge for artificial intelligence (AI) is **knowledge acquisition**: how to give computers enough knowledge so as to make their inferential capabilities useful (?). From an AI perspective, one of the major motivations for natural language processing is to provide a solution to this problem — acquiring knowledge in the way that people often do, by reading. This problem is sometimes called **information extraction**; in contrast to **information retrieval**, where the goal is to retrieve informative documents for a human reader, the goal of information extraction is to synthesize these documents into structured knowledge representations, such as database entries.

This chapter distinguishes information extraction from **question answering**, where the goal is to provide natural language answers to natural language questions. The tasks are closely related: a question answering system might proceed by first parsing the question (determining what information is required), then identifying relevant records in the knowledge base, and then crafting a natural language response. In many scenarios — such as the IBM question answering system “Watson” — the required knowledge base is too large to create by hand, so it must be created by information extraction techniques, similar to those discussed here.

A large part of information extraction can be unified in terms of **entities**, **relations**, and **events**. Entities are uniquely specified objects in the world, such as people, places, organizations, and times. Relations link pairs of entities, as in `sibling(LUKE, LEIA)`. We can think of each relation type as defining a table, in which each row contains two entities. Events link arbitrary numbers of arguments, as in the following example:

```
battle : ⟨location : ATLANTA,  
         date : 1864,  
         victor : UNITED STATES ARMY,  
         defeated : CONFEDERATE ARMY⟩.
```

We can think of each event type as defining a table, in which the rows define various

“slots” pertaining to the event. The task of **knowledge base population** is closely related to information extraction, and the goal is to fill in relevant slots in just such a table.

The attentive reader will notice a close kinship between information extraction, as defined here, and the task of shallow semantic parsing defined in chapter 14. For example, in semantic role labeling, the goal was to identify predicates and their arguments; we may think of predicates as corresponding to events, and the arguments as defining slots in the event representation. The key difference is that semantic role labeling and related tasks require correctly analyzing each sentence — a goal sometimes described as **micro-reading**. In information extraction, we need only correctly identify the relations and events that are referred to in a corpus. Many relations and events may be mentioned multiple times, but in information extraction and knowledge base population, we need only identify them once — thus the goal here is sometimes described as **macro-reading**. While macro-reading is a more forgiving task than micro-reading, it requires reasoning over an entire corpus, posing additional problems of computational tractability. It may also be necessary to provide **information provenance** [todo: good term?], linking the extracted knowledge back to the original source or sources.

17.1 Entities

The starting point for information extraction is to identify mentions of entities in text. For example, consider the following text.

(17.1) *The United States Army captured a hill overlooking Atlanta on May 14, 1864.*

Given this text, we have two goals:

1. **Identify** the spans *United States Army*, *Atlanta*, and *May 14, 1864* as entity mentions. We may also want to recognize the **named entity types**: organization, location, and date. This task is known as **named entity recognition**.
2. **Link** these spans to known entities in a knowledge base, U.S. ARMY, ATLANTA, and MAY 14, 1864. This task is known as **entity linking**.

Named entity recognition (NER)

A standard approach to tagging named entity spans is to use discriminative sequence labeling methods such as conditional random fields and structured perceptrons. As described in chapter 6, these methods use the Viterbi algorithm to search over all possible label sequences, while scoring each sequence using a feature function that decomposes across adjacent tags. Named entity recognition is formulated as a tagging problem by assigning each word token to a tag from a tagset. However, there is a major difference from part-of-speech tagging: in NER we need to recover **spans** of tokens, such as *The*

(c) Jacob Eisenstein 2014-2017. Work in progress.

The	U.S.	Army	captured	Atlanta	on	May	14	,	1864	.
B-ORG	I-ORG	I-ORG	O	B-LOC	O	B-DATE	I-DATE	I-DATE	I-DATE	O

Table 17.1: BIO notation for named entity recognition

United States Army. To do this, the tagset must distinguish tokens that are at the **beginning** of a span from tokens that are **inside** a span.

BIO notation This is accomplished by the “BIO notation”, shown in Table 17.1. Each token at the beginning of a name span is labeled with a B- prefix; each token within a name span is labeled with an I- prefix. Tokens that are not parts of name spans are labeled as O. From this representation, it is unambiguous to recover the entity name spans within a labeled text. Another advantage is from the perspective of learning: tokens at the beginning of name spans may have different properties than tokens within the name, and the learner can exploit this. This insight can be taken even further, with special labels for the last tokens of a name span, and for **unique** tokens in name spans, such as *Atlanta* in the example in Table 17.1. This is called **BILOU** notation, and has been shown to yield improvements in supervised named entity recognition Ratinov and Roth (2009).[\[todo: check this cite\]](#)

Entity types The number of possible entity types depends on the labeled data. An early dataset was released as part of a shared task in the Conference on Natural Language Learning (CoNLL), containing entity types LOC (location), ORG (organization), and PER (person). Later work has distinguished additional entity types, such as dates, [\[todo: etc\]](#). [\[todo: find cites\]](#) Special purpose corpora have been built for domains such as biomedical text, where entities include protein types [\[todo: etc\]](#).

Features The use of Viterbi decoding restricts the feature function $f(\mathbf{w}, \mathbf{y})$ to $\sum_m f(\mathbf{w}, y_m, y_{m-1}, m)$, so that each feature can consider only local adjacent tags. Typical features include tag transitions, word features for w_m and its neighbors, character-level features for prefixes and suffixes, and “word shape” features to capture capitalization. As an example, base

(c) Jacob Eisenstein 2014-2017. Work in progress.

features for the word *Army* in the example in Table 17.1 include:

```

(CURR-WORD:Army,
 PREV-WORD:U.S.,
 NEXT-WORD:captured,
  PREFIX-1:A-,
  PREFIX-2:Ar-,
  SUFFIX-1:-y,
  SUFFIX-2:-my,
   SHAPE:Xxxx)

```

Another source of features is to use **gazetteers**: lists of known entity names. For example, it is possible to obtain from the U.S. Social Security Administration a list of [**todo: hundreds of thousands**] of frequently used American names — more than could be observed in any reasonable annotated corpus. Tokens or spans that match an entry in a gazetteer can receive special features; this provides a way to incorporate hand-crafted resources such as name lists in a learning-driven framework.

Features in recent state-of-the-art systems are summarized in papers by ? and Ratinov and Roth (2009).

Alternative modeling frameworks*

Apart from sequence labeling, there are other formulations for named entity recognition, which are arguably better customized for the task.

17.2 Relations

Knowledge-base population

Distant supervision

17.3 Events and processes

17.4 Facts, beliefs, and hypotheticals

Chapter 18

Machine translation

Machine translation (MT) is one of the “holy grail” problems in natural language processing. Solving it would be a major advance in facilitating communication between people all over the world, and so it has received a lot of attention and funding since the early 1950s. However, it has proved incredibly challenging, and while there has been substantial progress towards usable MT systems — especially for so-called “high resource” languages like English and French — we are still far from automatically producing translations that capture the nuance and depth of human language.

18.1 The noisy channel model

Throughout the course, we’ve been working with the general formulation,

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} \theta \cdot f(x, y). \quad (18.1)$$

Now suppose we make x a sentence in a foreign (**source**) language, and $y \in \mathcal{Y}$ a sentence in the **target language**. We can thus view translation in the same linear structure-prediction formalism that we have used for tagging, parsing, etc. Will this work?

There are two major criteria for a translation:

- **Adequacy:** The translation \hat{y} should adequately reflect the linguistic content of w . For example, if $x = A \text{ Vinay le gusta Python}$, the gloss¹ $y = To \text{ Vinay it like Python}$ is considered adequate because it contains all the relevant content. The output $y = Vinay \text{ debugs memory leaks}$ will score poorly.
- **Fluency:** The translation \hat{y} should read like fluent text in the target language. By this criterion, the gloss $y = To \text{ Vinay it like Python}$ will score poorly, and $y = Vinay \text{ likes Python}$ will be preferred.

	Adequate?	Fluent?
<i>To Vinay it like Python</i>	yes	no
<i>Vinay debugs memory leaks</i>	no	yes
<i>Vinay likes Python</i>	yes	yes

Table 18.1: Adequacy and fluency for translations of the Spanish *A Vinay le gusta Python*

An early insight in machine translation was that the scoring function for a translation can decompose across these criteria:

$$\theta \cdot f(x, y) = \theta_t \cdot f_t(x, y) + \theta_\ell \cdot f_\ell(y) \quad (18.2)$$

The features f_t represent the translation model, which corresponds to the adequacy criterion; the features f_ℓ represent the language model, which corresponds to the fluency criterion.

The advantage of this decomposition is that we can estimate θ_ℓ from unlabeled data in the target language. Because unlabeled text data is widely available, in principle we can easily improve the fluency of our translations by estimating very high-order language models from ample unlabeled text. In this case, we can express these features as

$$f_\ell(y) = \bigcup_i \{y_{i:(i+k)}\} \quad (18.3)$$

$$\theta_\ell(\{y_i, y_{i+1}, \dots, y_{i+k}\}) = \log p(y_{i+k} \mid y_i, y_{i+1}, \dots, y_{i+k-1}) \quad (18.4)$$

When estimating these probabilities, we will naturally want to apply all the smoothing tricks that we learned in Chapter 5. Note that we will also have to add padding of K “buffer” words at the beginning and end of the input.

This approach is indeed a component of the many MT systems, but there is a catch: as the size of the N-gram features increases, the problem of **decoding** — selecting the best scoring translation \hat{y} — becomes exponentially more difficult. We will consider this issue later. For now, just note that this formulation ensures that,

$$\theta_\ell \cdot f_\ell(y) = \log p(y). \quad (18.5)$$

Now let’s consider the translation component. If we can set

$$\theta_t \cdot f_t(y, x) = \log p(x \mid y), \quad (18.6)$$

then the sum of these two scores yields,

$$\theta_t \cdot f_t(y, x) + \theta_\ell \cdot f_\ell(y) = \log p(x \mid y) + \log p(y) \quad (18.7)$$

$$= \log p(x, y). \quad (18.8)$$

¹A “gloss” is a word-for-word translation.

In other words, we can obtain the translation \hat{y} which has the maximum joint log-likelihood $\log p(y, x)$. We want the translation with the highest conditional probability,

$$\operatorname{argmax}_y p(y | x) = \operatorname{argmax}_y \frac{p(y, x)}{p(x)}, \quad (18.9)$$

but since x is given, we can ignore the denominator $p(x)$ and just select the y that maximizes the joint probability.

This approach is called the **noisy channel model**, and was pioneered by researchers who were experts in cryptography. They proposed to view translation as *decoding* the output of a stochastic cipher.

- Imagine that the original text y was written in English (or whatever is the target language), and is modeled as drawn from a source language model $y \sim p_\ell(y)$
- The original text was then stochastically encoded, according to the translation model, $x | y \sim p_t(x | y)$.
- If we can estimate the stochastic processes p_ℓ and p_t , we can reverse the cipher and obtain the original text.

18.2 Translation modeling

Language modeling is covered in Chapter 5, so this chapter will mainly focus on the translation model, $p_t(x | y)$. To estimate this model, we will need a parallel corpus, which contains sentences in both languages.

- Parallel corpora are often available from national and international governments. **The Hansards corpus** contains aligned English and French sentences from the Canadian parliament. **The EuroParl corpus** contains sentences for 21 languages, aligned with their English translations.
- More recent work has explored the use of web documents (Kilgariff and Grefenstette, 2003; Resnik and Smith, 2003) and crowdsourcing for MT (Zaidan and Callison-Burch, 2011).

Once a parallel corpus is obtained, we can consider how to characterize the translation model, f_t . The sets \mathcal{X} and \mathcal{Y} are far too huge for us to directly estimate the adequacy of every possible translation pair. So we need to decompose this problem into smaller units.

The **Vauquois Pyramid** is a theory of how translation should be modeled. At the lowest level, we translate individual words, but the distance here is far, because languages express ideas differently. If we can move up the triangle to syntactic structure, the distance for translation is reduced; we then need only produce target-language text from the

(c) Jacob Eisenstein 2014-2017. Work in progress.

syntactic representation, which can be as simple as reading off a tree. Further up the triangle lies semantics; translating between semantic representations should be easier still, but mapping between semantics and surface text is a difficult, unsolved problem. At the top of the triangle is **interlingua**, a semantic representation that is so generic, it is identical across all human languages. Philosophers may debate whether such a thing as interlingua is really possible (Derrida, 1985), but the idea of linking translation and semantic understanding can be viewed as a grand challenge for natural language technology.




Figure 18.1: The Vauquois Pyramid (“Direct translation and transfer translation pyramid”). Licensed under Creative Commons Attribution-Share Alike 3.0 via Wikimedia Commons.)

Returning to earth, the simplest decomposition of the translation model is a word-based translation: each word in the source string should be aligned to a word in the translation. In this approach, we need an **alignment** $\mathcal{A}(x, y)$, which contains a list of pairs of source and target tokens. For example, given $x = A \text{ Vinay le gusta Python}$ and $y = \text{Vinay likes Python}$, one possible word-to-word alignment is,

$$\mathcal{A}(x, y) = \{\langle \text{Vinay}, \text{Vinay} \rangle, \langle \text{gusta}, \text{likes} \rangle, \langle \text{Python}, \text{Python} \rangle, \langle A, \emptyset \rangle, \langle \text{le}, \emptyset \rangle\}. \quad (18.10)$$

Another, less promising, possibility is:

$$\mathcal{A}(x, y) = \{\langle A, \text{Vinay} \rangle, \langle \text{Vinay}, \text{likes} \rangle, \langle \text{le}, \text{Python} \rangle, \langle \text{gusta}, \emptyset \rangle, \langle \text{Python}, \emptyset \rangle\}. \quad (18.11)$$

Given the alignment, we can define the translation probability as,

$$p_t(x, \mathcal{A} \mid y) = \prod_i p(x_i, a_i \mid y_{a_i}) \quad (18.12)$$

$$= \prod_i p_a(a_i \mid i, N_x, N_y) \times p_{x|y}(x_i \mid y_{a_i}). \quad (18.13)$$

This probability model makes some assumptions that we now state explicitly:

(c) Jacob Eisenstein 2014-2017. Work in progress.

- The alignment probability decomposes as $p(\mathcal{A} \mid \mathbf{x}, \mathbf{y}) = \prod_i p_a(a_i \mid i, N_x, N_y)$. This means that each alignment decision is independent of the others, and depends only on the index i , and the sentence lengths N_x and N_y .
- The translation probability decomposes as $p(\mathbf{x} \mid \mathbf{y}, \mathcal{A}) = \prod_i p_{x|y}(x_i \mid y_{a_i})$, which means that each word in \mathbf{x} depends only on its aligned word in \mathbf{y} . This means that we are doing word-based translation only, ignoring context. The hope is that the language model will correct any disfluencies that arise from word-to-word translation.

A series of translation models with increasingly relaxed independence assumptions was produced by researchers at IBM in the 1980s and 1990s, known as IBM Models 1-6 (Och and Ney, 2003). IBM model 1 makes the strongest independence assumption:

$$p_a(a_i \mid i, N_x, N_y) = \frac{1}{N_y} \quad (18.14)$$

In this model every alignment is equally likely! This is almost surely wrong, but it makes learning easy.

Let's consider how to translate with IBM model 1. The key idea is to treat the alignment as a **hidden variable**. If we knew the alignment, we could easily estimate a translation model, and we could find the optimal translation as

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y}} p(\mathbf{x}, \mathbf{y}) \quad (18.15)$$

$$= \operatorname{argmax}_{\mathbf{y}} \sum_{\mathcal{A}} p(\mathbf{x}, \mathbf{y}, \mathcal{A}) \quad (18.16)$$

$$= \operatorname{argmax}_{\mathbf{y}} p_{\ell}(\mathbf{y}) \sum_{\mathcal{A}} p_t(\mathbf{x}, \mathcal{A} \mid \mathbf{y}) \quad (18.17)$$

Conversely, if we had an accurate translation model, we could estimate beliefs about each alignment decision,

$$q(a_i \mid \mathbf{x}, \mathbf{y}) \propto p_a(a_i \mid i, N_x, N_y) \times p_{x|y}(x_i \mid y_{a_i}). \quad (18.18)$$

We therefore have a classic chicken-and-egg problem, which we can solve using the iterative expectation-maximization (EM) algorithm.

E-step Update beliefs about word alignment,

$$q_i(a_i) \propto p_a(a_i \mid i, N_x, N_y) p_{x|y}(x_i \mid y_{a_i}) \quad (18.19)$$

M-step Update the translation model,

$$\theta_{u \rightarrow v} = \log \frac{\sum_i \sum_j q_i(a_i = j) \delta(y_j = u \wedge x_i = v)}{\sum_i \sum_j q_i(a_i = j) \delta(y_j = u)} \quad (18.20)$$

(c) Jacob Eisenstein 2014-2017. Work in progress.

Example for IBM Model 1

Suppose we have an English/French bilingual text (**bitext**) with two sentence pairs:

(18.1) *The coffee*
Le cafe

(18.2) *My coffee*
Mon cafe

We start with the following translation probabilities:

	<i>le</i>	<i>mon</i>	<i>cafe</i>
<i>the</i>	$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$
<i>my</i>	$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$
<i>coffee</i>	$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$

In the E-step, we compute alignment probabilities for each sentence.

$$q_0(0) \propto p_a(0) \times p(\textit{le} \mid \textit{the}) = \frac{1}{2} \times \frac{1}{3} \quad (18.21)$$

$$q_0(1) \propto p_a(1) \times p(\textit{le} \mid \textit{coffee}) = \frac{1}{2} \times \frac{1}{3} \quad (18.22)$$

$$q_0(\cdot) = \left[\frac{1}{2}, \frac{1}{2} \right] \quad (18.23)$$

The same logic applies to all the alignment decisions: we begin with $q_i(j) = \frac{1}{N}$ in every case. Now we move to the M-step, where we will plug in these (apparently uninformative) alignment probabilities:

$$p_{x|y}(\textit{le} \mid \textit{the}) = \frac{\sum_{i,j} q_i(j) \delta(y_i = \textit{le} \wedge x_j = \textit{the})}{\sum_{i,j} q_i(j) \delta(x_j = \textit{the})} = \frac{\frac{1}{2}}{\frac{1}{2} + \frac{1}{2}} = \frac{1}{2} \quad (18.24)$$

$$p_{x|y}(\textit{cafe} \mid \textit{the}) = \frac{\sum_{i,j} q_i(j) \delta(y_i = \textit{le} \wedge x_j = \textit{the})}{\sum_{i,j} q_i(j) \delta(x_j = \textit{the})} = \frac{\frac{1}{2}}{\frac{1}{2} + \frac{1}{2}} = \frac{1}{2} \quad (18.25)$$

$$p_{x|y}(\textit{mon} \mid \textit{the}) = \frac{\sum_{i,j} q_i(j) \delta(y_i = \textit{le} \wedge x_j = \textit{the})}{\sum_{i,j} q_i(j) \delta(x_j = \textit{the})} = \frac{0}{\frac{1}{2} + \frac{1}{2}} = 0 \quad (18.26)$$

The math works out similarly for $p(\cdot \mid \textit{my})$. But the English word *coffee* appears in both

sentence pairs, so:

$$p_{x|y}(le \mid cafe) = \frac{\frac{1}{2}}{4 \times \frac{1}{2}} = \frac{1}{4} \quad (18.27)$$

$$p_{x|y}(coffee \mid cafe) = \frac{2 \times \frac{1}{2}}{4 \times \frac{1}{2}} = \frac{1}{2} \quad (18.28)$$

$$p_{x|y}(mon \mid cafe) = \frac{\frac{1}{2}}{4 \times \frac{1}{2}} = \frac{1}{4} \quad (18.29)$$

$$(18.30)$$

To summarize the new translation probabilities:

	<i>le</i>	<i>mon</i>	<i>cafe</i>
<i>the</i>	$\frac{1}{2}$	0	$\frac{1}{2}$
<i>my</i>	0	$\frac{1}{2}$	$\frac{1}{2}$
<i>coffee</i>	$\frac{1}{4}$	$\frac{1}{4}$	$\frac{1}{2}$

We now go back to the E-step and compute the alignments again.

$$q_0(0) \propto p_a(0) \times p(le \mid the) = \frac{1}{2} \times \frac{1}{2} \quad (18.31)$$

$$q_0(1) \propto p_a(1) \times p(le \mid coffee) = \frac{1}{2} \times \frac{1}{4} \quad (18.32)$$

$$q_0(\cdot) = \left[\frac{2}{3}, \frac{1}{3} \right] \quad (18.33)$$

$$q_1(0) \propto p_a(0) \times p(le \mid coffee) = \frac{1}{2} \times \frac{1}{4} \quad (18.34)$$

$$q_1(1) \propto p_a(1) \times p(cafe \mid coffee) = \frac{1}{2} \times \frac{1}{2} \quad (18.35)$$

$$q_1(\cdot) = \left[\frac{1}{3}, \frac{2}{3} \right] \quad (18.36)$$

Having learned something about the translation model, the alignments are no longer uniform. The situation for the second sentence is identical, so is not shown here.

(c) Jacob Eisenstein 2014-2017. Work in progress.

	<i>le</i>	<i>mon</i>	<i>café</i>
<i>the</i>	$\frac{2}{3}$	0	$\frac{1}{3}$
<i>my</i>	0	$\frac{2}{3}$	$\frac{1}{3}$
<i>coffee</i>	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{2}{3}$

If we return to the M-step, we end up with sharper translation probabilities:

$$p_{x|y}(le | the) = \frac{\sum_{i,j} q_i(j) \delta(y_i = le \wedge x_j = the)}{\sum_{i,j} q_i(j) \delta(x_j = the)} = \frac{\frac{2}{3}}{\frac{2}{3} + \frac{1}{3}} = \frac{2}{3} \quad (18.37)$$

$$p_{x|y}(café | the) = \frac{\sum_{i,j} q_i(j) \delta(y_i = le \wedge x_j = the)}{\sum_{i,j} q_i(j) \delta(x_j = the)} = \frac{\frac{1}{3}}{\frac{1}{3} + \frac{2}{3}} = \frac{1}{3} \quad (18.38)$$

$$p_{x|y}(mon | the) = 0 \quad (18.39)$$

$$p_{x|y}(le | café) = \frac{\frac{1}{3}}{\frac{1}{3} + \frac{2}{3} + \frac{1}{3} + \frac{2}{3}} = \frac{1}{6} \quad (18.40)$$

$$p_{x|y}(coffee | café) = \frac{2 \times \frac{2}{3}}{2} = \frac{2}{3} \quad (18.41)$$

$$p_{x|y}(mon | café) = \frac{\frac{1}{3}}{2} = \frac{1}{6} \quad (18.42)$$

The process will eventually converge to assign all of the probability mass for each English word to its correct French translation. Note that we have made no assumptions about the word alignments at all! The only information that we have exploited is the co-occurrence of words across sentence pairs. But we can do even better in models that make reasonable assumptions about alignment — for example, that alignments tend to be monotonic ($i > j \rightarrow a_i > a_j$), etc.

Better alignment models IBM Model 2 tries to learn the prior distribution from data,

$$p_a(a_i; i, N_x, N_y) = \phi_{a_i, i, N_x, N_y} \quad (18.43)$$

$$s.t. \forall i, N_x, N_y, \sum_a \phi_{a, i, N_x, N_y} = 1. \quad (18.44)$$

The variables a_i and i are integer indices, so ϕ_{a_i, i, N_x, N_y} represents the probability that token i is aligned to token a_i in sentence pairs with length N_x and N_y . We compute this probability by the relative frequency estimate,

$$\phi_{a, i, N_x, N_y} = \frac{\sum_{\mathbf{y}, \mathbf{x}: \#|\mathbf{y}|=N_y, \#|\mathbf{x}|=N_x} q_i(a)}{\sum_{\mathbf{y}, \mathbf{x}: \#|\mathbf{y}|=N_y, \#|\mathbf{x}|=N_x} \delta(\#|\mathbf{y}| < i)}, \quad (18.45)$$

where we are summing only over sentence pairs with lengths N_x, N_y .

(c) Jacob Eisenstein 2014-2017. Work in progress.

Adding a parameter for the alignment model makes the overall objective function non-convex (see chapter 4 for a review of convexity). The practical consequence of this is that initialization matters; it is no longer sufficient to just initialize the translation model to uniform probabilities and hope that everything works out. A good solution is to first run IBM Model 1, and then use the resulting translation model as the initialization for IBM Model 2.

IBM model 3 adds a term for the “fertility” of each word — that is, the number of words that typically align to it. For example, some English verbs are translated as multiword phrases in Spanish:

- (18.3) *Mary did not **slap** the green witch.*
*Maria no **daba una bofetada** a la bruja verde.*

By learning these fertility probabilities from data, the alignment model has a better chance of learning the correct translation rules for such multiword phrases. But note that even in the best case, we would have to model the translation of *slap* into *daba una bofetada* as,

$$p_{x|y,A}(daba\ una\ bofetada\ |\ slap) \quad (18.46)$$

$$= p_{x|y}(daba\ |\ slap) \times p_{x|y}(una\ |\ slap) \times p_{x|y}(bofetada\ |\ slap). \quad (18.47)$$

This seems wrong, since the word *una* is just an indefinite article — the Spanish feminine for the English word *a*. We therefore turn to models that go beyond word-based translation.

18.3 Phrase-based translation

The problem identified with the example *daba una bofetada* is an instance of a more general issue: translation is often not a matter of word to word substitutions. Multiword expressions are often not translated literally:

- (18.4) *clean up*
faire (make) le (the) menage (home)

Handling this in a word-to-word translation model seems unnecessarily difficult. Furthermore, phrases tend to move together:

- (18.5) *i like the food a lot*
la (the) comida (food) me (I) gusta (like) mucho (a lot)

We would therefore have to learn that the alignment decisions for *la* and *comida* should be made jointly.

Phrase-based translation generalizes on word-based models by building translation tables and alignments between multiword spans of text. The generalization from word-based translation is surprisingly straightforward: the translation tables can now condition


on multi-word units, and can assign probabilities to multi-word units; alignments are mappings from spans to spans, $\langle (i, j), (k, \ell) \rangle$, so that

$$p(\mathbf{x} \mid \mathbf{y}, \mathcal{A}) = \prod_{\langle (i,j), (k,\ell) \rangle \in \mathcal{A}} p_{\mathbf{x}|\mathbf{y}}(\{x_i, x_{i+1}, \dots, x_j\} \mid \{y_k, y_{k+1}, \dots, y_\ell\}), \quad (18.48)$$

where we require that the alignment set \mathcal{A} cover both sentences with non-overlapping spans, as shown in ??. [todo: add figure]

18.4 Syntactic MT

Consider the English sentence, *The green witch eats the hot soup.*



Where NPB is a “bare NP,” without the determiner. We might get this non-terminal from binarizing a CFG.


We can view the CFG as a process for **generating** English sentences.

Synchronous CFGs are a generalization of CFGs. They generate text in two different languages simultaneously. Each RHS has two components, one for each language. Subscripts show the mapping between non-terminals in the RHS. For example:

$$\begin{array}{ll}
 S \rightarrow NP_1 VP_2, & NP_1 VP_2 \\
 VP \rightarrow V_1 NP_2, & V_1 NP_2 \\
 NP \rightarrow DT_1 NPB_2, & DT_1 NPB_2 \\
 NPB \rightarrow JJ_1 NPB_2, & NPB_2 JJ_1
 \end{array}$$

The key production is the fourth one, which handles the re-ordering of adjectives and nouns. Let’s use this SCFG to generate the English and Spanish versions of this sentence.

(c) Jacob Eisenstein 2014-2017. Work in progress.



- On the slides there is another example, in Japanese. Since Japanese is a SOV language (subject-object-verb), we need a production: $VP \rightarrow V_1 NP_2, NP_2 V_1$.
- As with CFGs, we can attach a probability to each production, and compute the joint probability of the derivation and the text as the product of these productions.

Binarization

Let's define a rank- n CFG as a grammar with at most n elements on a right-hand side.

- CFGs can always be binarized.
 - e.g. $NP \rightarrow DT [JJ NN]$ becomes

$$NP \rightarrow DT NPB$$

$$NPB \rightarrow JJ NN$$

- Therefore, the set of languages that can be defined by a 2-CFG is identical to the set that can be defined by 3-CFG, 4-CFG, etc...
- What about SCFGs?
 - Rank 3:

$$\begin{array}{ll}
 A \rightarrow B [C D], & [C D] B \\
 A \rightarrow B V, & V B \\
 V \rightarrow C D, & C D
 \end{array}$$

Yes, we can. 2-SCFG = 3-SCFG.

(c) Jacob Eisenstein 2014-2017. Work in progress.

– Rank 4:


$$\begin{array}{ll}
 A \rightarrow B C D E, & C E B D \\
 A \rightarrow [B C] D E, & [C E B] D \\
 A \rightarrow B [C D] E, & [C E B D] \\
 A \rightarrow B C [D E], & C [E B D]
 \end{array}$$

In each chunk that we might want to replace in the first language, we have one or more intervening symbols in the second language. Therefore, $3\text{-SCFG} \subsetneq 4\text{-SCFG}$.

- The subset of $2\text{-SCFG} = 3\text{-SCFG}$ is equivalently called **inversion transduction grammar**. The notation is slightly different, we write $A \rightarrow [B C]$ when the order is preserved and $A \rightarrow \langle B C \rangle$ when it is inverted.


No raising or lowering

SCFGs can only reorder sibling nodes. Is that enough? Not always.



SCFGs cannot swap the subject and object, because they aren't siblings in the original grammar.

We could solve this by changing the grammar,



By including the verb *misses/manque à* directly into the rule, we ensure that it doesn't apply to other verbs.

With other syntactic translation models (synchronous tree substitution grammar or tree adjoining grammars), this case can be handled without flattening.

18.5 Algorithms for SCFGs

Translation

In principle, translation in SCFGs is nearly identical to parsing. Suppose we have the Spanish phrase *la razón principal*, and the synchronous grammar

$NP \rightarrow D \ NPB,$	$D \ NPB$	1.0
$NPB \rightarrow N_1 \ J_2,$	$J_2 \ N_1$	0.8
$NPB \rightarrow N_1 \ N_2,$	$N_1 \ N_2$	0.2
$D \rightarrow la,$	the	0.5
$N \rightarrow razon,$	$reason$	0.5
$N \rightarrow principal,$	$principal$	0.5
$J \rightarrow principal,$	$main$	1.0

Now we can apply CKY, building the translation on the English side. We should get two possible translations, *the reason principal* ($p(e, f, \tau) = 0.05$) and *the main reason* ($p(e, f, \tau) = 0.4$).

What is the complexity of translation with binarizable SCFGs? It's just like CFG parsing: $\mathcal{O}(n^3)$.

Bitext parsing

To learn a translation model, we might need to synchronously parse the **bitext**: both the source and target side language.

We can do this with a dynamic program.

Assuming we are dealing with 2-SCFG or 3-SCFG, here's what we need to keep track of:

- The non-terminals that we have derived
- Their spans in the source language (start and end)
- Their spans in the target language (start and end)

Suppose we are given spans $\langle i, j \rangle$ in the source and $\langle i', j' \rangle$ in the target. Then we are looking for split points k and k' and a production that can derive the subspans $\langle i, k \rangle, \langle k, j \rangle$ and $\langle i', k' \rangle, \langle k', j' \rangle$.

What is the space complexity of bitext parsing? $\mathcal{O}(|S|n^4)$, where $|S|$ is the number of non-terminals.

What is the time complexity of bitext parsing? $\mathcal{O}(|R|n^6)$, where $|R|$ is the number of production rules.

Specificially, we have the recurrence

$$\begin{aligned} \psi(X, i, j, i', j') = & \max_{k, k', A, B} P(S \rightarrow A B, A B) \otimes \psi(A, i, k, i', k') \otimes \psi(B, k, j, k', j') \\ & \oplus P(S \rightarrow A B, B A) \otimes \psi(A, i, k, k', j') \otimes \psi(B, k, j, i', k') \end{aligned}$$

Note: in general, bitext parsing is exponential in the rank of the SCFG (unless $P = NP$).

Intersection with language model

For fluent translations, we typically want to multiply in the language model probability on the target side.

- This (usually) corresponds **intersection** of an SCFG with a finite state machine.
- Sidenote: **what about context-free language models?**
 - $A = \{a^m b^m c^n\}$
 - $B = \{a^m b^n c^n\}$
 - $A \cap B = \{a^n b^n c^n\}$, not a CFL!
 - CFLs are not closed under intersection.
 - Determining if $s \in A \cap B$ is in PSPACE
- There are exact dynamic programming algorithms for intersecting an SCFG and an FSA, but they are very slow. One solution is **cube pruning**.
- We can equivalently view this as an ILP

$$\begin{aligned} \min. \quad & \sum_v \theta_v y_v + \sum_e \theta_e y_e + \sum_{\langle v, w \rangle \in \mathcal{B}} \theta(v, w) y(v, w) \\ \text{s.t.} \quad & C0 : y_v, y_e \text{ form a derivation} \\ & C1 : y_v = \sum_{w: \langle w, v \rangle \in \mathcal{B}} y(w, v) \\ & C2 : y_v = \sum_{w: \langle v, w \rangle \in \mathcal{B}} y(v, w) \end{aligned}$$

- Here y_e and y_v are indicator variables that define what words and hyperedges appear in the derivation.
- We can solve this optimization with Lagrangian relaxation.
 - Replace the outgoing constraints C2 with multipliers $u(v)$
 - At first, $u(v) = 0, \forall v$

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Without the outgoing constraints, we can optimize efficiently
 - If the outgoing constraints happen to be met, we are done
 - Otherwise, update $u(v)$ and try again.
- Lagrangian relaxation finds the exact solution 97% of the time, is many times faster than ILP.

Part V

Learning

Chapter 19

Semi-supervised learning

So far we have focused on learning a classifier — typically represented by a set of weights θ — from a set of labeled examples $\{(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})\}_{i=1}^N$. As we have seen, it is possible to use this framework to formulate tasks ranging from document classification to semantic parsing. But what if you don't have those labeled examples for the domain or task that you want to solve?

This scenario happens all the time — class projects, interdisciplinary collaborations, and commercial applications. Digital text is available from an increasingly diverse set of **domains**: social media, electronic health records, e-government, etc. But this text is usually unlabeled, so we cannot train NLP systems within these domains. Lack of labeled data in the target domains and tasks is the main limitation to language technology being applicable more broadly.

There are two “simple” solutions that one might undertake:

1. Use some other labeled data and hope it works.

Unfortunately, it probably won't. For example, in applying parsers trained on the Penn Treebank to social media texts, researchers have observed massive decreases in accuracy (Foster et al., 2011; Gimpel et al., 2011).

2. Label data yourself.

This is a lot of work. For example:

- **The Switchboard corpus** contains phoneme annotations of telephone conversations, e.g.

film → F IH_N UH_GL_N M
be all → BCL B IY IY_TR AO_TR AO L_DL

This took 400 hours of annotation time per hour of speech.

- **The Penn Chinese Treebank** is a set of CFG annotations for Chinese. It took 2 years to get 4000 sentences annotated.

Crowd-sourcing has recently become popular as a means to annotate large amounts of data quickly. This can work well, but effort and expertise are needed to get good annotations for linguistically complex tasks (Snow et al., 2008; Zaidan and Callison-Burch, 2011).

In this chapter, we will explore an alternative to either of these approaches: harnessing data that is unlabeled, or is labeled in a different domain or task. We will think of our annotated data as a *sample* from some underlying distribution.

$$\{(\mathbf{x}^{(i)}, \mathbf{y}_i^{(i)})\}_{i=1}^N \sim \mathcal{D} \quad (19.1)$$

This allows us to formulate various learning scenarios:

Semisupervised learning Imagine that N is small, so that it is hard to learn a model that generalizes well. We would like to leverage **unlabeled** data,

$$\{\mathbf{x}_i\}_{i=1}^{N_u} \sim \mathcal{D}, \quad (19.2)$$

which is drawn from the same underlying distribution \mathcal{D} , but for which labels are unavailable. Since this data is not labeled, it is usually available in very large quantities, so $N_u \gg N$.

We have already seen two examples of semi-supervised learning. The first was the use of **expectation-maximization** in document classification in chapter 4; in the E-step, we impute beliefs about the labels of unlabeled documents, and then use these beliefs to update our model in the M-step (Nigam and Ghani, 2000). Another example of semi-supervised learning was given in chapter 15. There we saw how to use unlabeled data to build **Brown clusters**. These clusters then act as features, generalizing over individual words by capturing lexical similarity (Miller et al., 2004; Koo et al., 2008).

While these techniques are effective, they are limited. Expectation-maximization requires a generative model, which may be a less effective classifier than a discriminative alternative such as logistic regression or support vector machines. Brown clusters are useful features, but they are learned separately from the main label prediction task. In § 19.1, we will explore additional techniques for semisupervised learning, such as bootstrapping and multiview learning.

Active learning This setting is similar to semi-supervised learning, but with a twist: we can iteratively query a user for labels for a small number of unlabeled instances. This is relevant in commercial settings, where a company can pay a small staff of

(c) Jacob Eisenstein 2014-2017. Work in progress.

annotators to label examples until performance is good enough. The key question is deciding which examples to label next. Settles (2010) surveys a number of alternatives; we will not explore the issue here.

Supervised domain adaptation Now imagine that we have a large amount of labeled data in some **source** domain, but a much smaller amount of information in the **target** domain. For example, the source domain could be 20th century newspaper articles (as in the Penn Treebank), and the target domain could be something like social media posts or patient medical records. We don't have enough target domain data to learn a good model. But if we simply combine all the data from the two domains, the source domain instances will dominate, and we will suffer from the resulting **domain shift**. We will consider various techniques for learning effectively from both domains.

Multitask (transfer) learning Similar to supervised domain adaptation, but rather than assuming that the underlying distribution $p(X, Y)$ shifts across domains, we assume that only the label distribution $p(Y | X)$ shifts. For example, we are working in the newstext domain, and we have a large amount of labeled data for part-of-speech tagging, and a small amount of labeled data for named-entity recognition. The goal is then to learn a better model using both labeled datasets.

Unsupervised domain adaptation This setting combines features of semisupervised learning and supervised domain adaptation: we have labeled data in the source domain, but no labeled data in the target domain. The prototypical example of this situation is in sentiment polarity analysis of product reviews: you are given annotated reviews of, say, coffee machines, but you want to predict the sentiment for reviews of bicycles (Blitzer et al., 2007). Another relevant setting is the application of syntactic analyzers such as part-of-speech taggers to historical texts (Yang and Eisenstein, 2015).

19.1 Semisupervised learning

Let's first consider the question of why would unlabeled data might help in a supervised classification task. Suppose you want to do sentiment analysis in French. I give you two labeled examples:

(19.1) ☺ *émouvant avec grâce et **style***

(19.2) ☹ *fastidieusement inauthentique et **banale***

You have a bunch of unlabeled examples too:

(19.3) *pleine de **style** et d'**intrigue***

(c) Jacob Eisenstein 2014-2017. Work in progress.

(19.4) *la banalité n'est dépassée que par sa **prétention***

(19.5) ***prétentieux**, de la première minute au rideau final*

(19.6) *imprégné d'un air d'**intrigue***

If we just learn from the labeled data, we might conclude that *style* is positive and that *banale* is negative. This isn't much. However, we can propagate this information to the unlabeled data, and potentially learn more.

- If we are confident about *style* being positive, then we can guess that (19.3) is also positive.
- That suggests that *intrigue* is also positive.
- We can then propagate this information to (19.6), and learn more.
- Similarly, we can propagate from the labeled data to (19.4), which we guess to be negative. This suggests that *pretention* is also negative, which we propagate to (19.5).

What happened here? Instances (19.3) and (19.4) were “similar” to our labeled examples for positivity and negativity, respectively. We used them to expand those concepts, which allowed us to correctly label instances (19.5) and (19.6), which didn't share any important features with our original labeled data. In doing this, we made a key assumption: that similar instances will have similar labels. (Is this assumption reasonable? Keep this question in mind.) In this case, we defined similarity in terms of sharing some key words (non-stopwords).

To see how this can help conceptually, think about similarity just in terms of 1D space. If you have only the two labeled instances, your decision boundary should be right in between. (Do you remember what criterion justifies this choice?) But if you have a bunch of unlabeled instances, you might want to draw this boundary in a different place. Let's now see how we can operationalize this idea in an algorithm.

Semi-supervised learning with EM

We've already seen one way to do this: use expectation-maximization (EM) to marginalize over the labels of the unseen data. So we are maximizing

$$p(X^\ell, Y^\ell, X^U) = p(X^\ell, Y^\ell) \sum_{Y^U} p(X^U, Y^U). \quad (19.3)$$

Expectation-maximization maximizes a bound on the joint probability defined above, by iterating between two steps:

E-step Fit a distribution $Q(y_i)$ for all unlabeled i ;

M-step Maximize the expected likelihood under this distribution.

(c) Jacob Eisenstein 2014-2017. Work in progress.

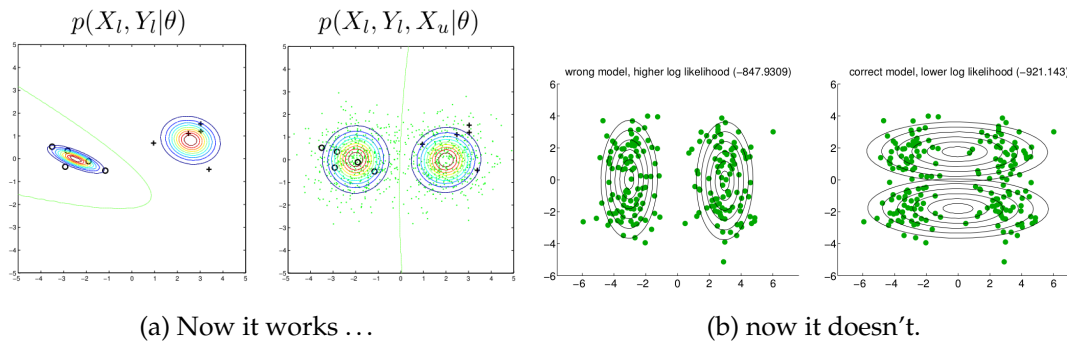


Figure 19.1: Expectation-maximization for semi-supervised learning on Gaussian data [todo: find credits for these images; Jerry Zhu?]

You can see why this can work in the example shown in Figure 19.1a: by incorporating unlabeled data, we get a much more reasonable decision boundary.

However, things can also go wrong, as shown in Figure 19.1b. In this example, the correct model (left) has a lower log-likelihood than the incorrect model (right). The basic problem here is that the model is wrong. The label is related to the observations, but not in the simplistic, Gaussian way that we had assumed. In chapter 4, we discussed a heuristic to try to deal with this problem: downweighting the contribution of the unseen data to the likelihood function. But this requires setting the weight parameter, which depends on a host of problem-specific characteristics, such as the underlying variance of the data. We will now consider some alternatives that often work better.

Bootstrapping and co-training

EM is sort of like self-training or **bootstrapping**: we use our current model to estimate $Q(y_i)$, and then update the model as if $Q(y_i)$ is correct.

- The probabilistic nature of this is nice, but it limits us to relatively weak classifiers.
- If we are willing to give up on probability, we can bootstrap **any** classifier.

Rather than imputing beliefs about all unlabeled instances $Q(y_i)$, we can add just a few, highly confident instances at each step. This is similar to how we proceeded in the French sentiment labeling example above. The simplest version of this algorithm is 1-nearest-neighbor: for each unlabeled data point, if its nearest neighbor has a label, then propagate that label. This approach does not make the parametric assumptions that doomed us in Figure 19.1b; instead, it relies on the similarity graph over instances. For some types of data, this is more reasonable, but it can also fail, as shown in the slides [todo: add these figures here].

There is some “folk wisdom” about when bootstrapping works:

(c) Jacob Eisenstein 2014-2017. Work in progress.

- It works better for generative models (e.g., Naive Bayes) than for discriminative models (e.g., perceptron)
- It works better when the Naive Bayes assumption is stronger.
 - Suppose we want to classify NEs as PERSON or LOCATION
 - Features: string and context
 - * *located on Peachtree Street*
 - * *Dr. Walker said ...*

$$P(W_{m+1} = \text{street}, W_{m-1} = \text{on} \mid Y_m = \text{LOC}) \\ \approx P(W_{m+1} = \text{street} \mid Y_m = \text{LOC})P(W_{m-1} = \text{on} \mid Y_m = \text{LOC})$$

Cotraining makes the bootstrapping assumptions explicit (Blum and Mitchell, 1998).

- Assume two, **conditionally independent**, views of a problem.
- Assume each view is sufficient to do good classification.

Sketch of learning algorithm:

- On labeled data, minimize error.
- On unlabeled data, **constrain** the models from different views to agree with each other.

Co-training example See the slides for an animated version of this. Assume we want to do named entity classification: determine whether an NE is a Location or Person. We have two views: the name itself, and its context.

	$x^{(1)}$	$x^{(2)}$	y
1.	Peachtree Street	located on	LOC
2.	Dr. Walker	said	PER
3.	Zanzibar	located in	? \rightarrow LOC
4.	Zanzibar	flew to	? \rightarrow LOC
5.	Dr. Robert	recommended	? \rightarrow PER
6.	Oprah	recommended	? \rightarrow PER

Algorithm

- Use classifier 1 to label example 5.
- Use classifier 2 to label example 3.

(c) Jacob Eisenstein 2014-2017. Work in progress.




Figure 19.2: Semi-supervised sentiment analysis as a graph

- Retrain both classifiers, using newly labeled data.
- Use classifier 1 to label example 4.
- Use classifier 2 to label example 6.

Multiview Learning is another approach in this style. Cotraining treats the output of each view's classifier as a labeled instance for the other view. In multiview learning, we add a **co-regularizer** that penalizes disagreement between the views on the unlabeled instances. This allows us to define a single objective function. In the case of two-view linear regression, the function is

$$\begin{aligned} \min_{w,v} \sum_i^L (y_i - w^\top x_i^{(1)})^2 + (y_i - v^\top x_i^{(2)})^2 + \lambda_1 \|w\|^2 + \lambda_1 \|v\|^2 \\ + \lambda_2 \sum_{i=L+1}^{L+U} (w^\top x_i^{(1)} - v^\top x_i^{(2)})^2 \end{aligned} \quad (19.4)$$

The only difference from standard regression is the co-regularizer, which penalizes disagreement on the unlabeled data.

An early version of this idea is **co-boosting** (Collins and Singer, 1999), where each view is a boosting classifier, and features are added incrementally to each view.

Graph-based approaches

Let's go back to sentiment analysis in French. We can view this data as a **graph**, with edges between similar instances, as shown in Figure 19.2. Unlabeled instances propagate information through the graph.

Where does the graph come from?

- Sometimes there is a natural similarity metric (time, position in the document).
- Otherwise, we can compute similarity from features. If the features are Gaussian, we could say:

$$\text{sim}(i, j) = \exp \left(-\frac{\|x_i - x_j\|^2}{2\sigma^2} \right)$$

If the features are discrete, we might use KL-divergence.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Then we add an edge between i and j when $\text{sim}(i, j) > \tau$

Given a graph with edge weights s_{ij} , we can formulate semi-supervised learning as an optimization problem:


$$\begin{aligned} \min_{\mathbf{z}} \quad & \sum_{i,j} s_{ij} (z_i - z_j)^2 \\ \text{s.t.} \quad & \forall_{i \in \{1 \dots N_\ell\}} z_i = y_i \\ & \forall_i z_i \in \{0, 1\} \end{aligned} \tag{19.5}$$

This looks like a combinatorial problem. Specifically, it looks like (binary) integer linear programming, which is NP-complete. But assuming $s_{ij} \geq 0$, this specific problem can be reformulated as maximum-flow, with polynomial time solutions. Rao and Ravichandran (2009) apply this idea to expanding polarity lexicons. In their graph:


- Nodes are words
- Edges are wordnet relations
- They label a few nodes for sentiment polarity, and propagate those labels to other parts of the graph.
- However, they use a slightly modified version of the mincut idea: randomized min-cut (Blum et al., 2004).

Randomized min-cut

Suppose we have this initial graph:



What is the solution? Actually, the following solutions are all equivalent:



Another problem with mincuts is that it doesn't distinguish high-confidence and low-confidence predictions. Both of these problems can be dealt with by randomization:

- Add random noise to adjacency matrix.
- Rerun mincuts multiple times.
- Deduce the final classification by voting.

(c) Jacob Eisenstein 2014-2017. Work in progress.

Label propagation

A related approach is **label propagation** (Zhu and Ghahramani, 2002), which Rao and Ravichandran also consider. The basic idea is that we relax y_i from an integer $\{0, 1\}$ to a real number \mathbb{R} . Then we solve the optimization problem,

$$\begin{aligned} \min_Y \sum_{i,j} s_{ij}(y_i - y_j)^2 \\ \text{s.t. } Y_L \text{ is clamped to initial values} \end{aligned}$$

The advantages are:

- a unique global optimum
- a natural notion of confidence: distance of y_i from 0.5

Let's look at the objective:

$$\begin{aligned} J &= \frac{1}{2} \sum_{i,j} s_{ij}(y_i - y_j)^2 \\ &= \frac{1}{2} \sum_{i,j} s_{ij}(y_i^2 + y_j^2 - 2y_i y_j) \\ &= \sum_i y_i^2 \sum_j s_{i,j} - \sum_{i,j} s_{ij} y_i y_j \\ &= \mathbf{y}^\top \mathbf{D} \mathbf{y} - \mathbf{y}^\top \mathbf{S} \mathbf{y} \\ &= \mathbf{y}^\top \mathbf{L} \mathbf{y} \end{aligned}$$

We have introduced three matrices

- Let \mathbf{S} be the $n \times n$ similarity matrix.
- Let \mathbf{D} be the **degree matrix**, $d_{ii} = \sum_j s_{ij}$. \mathbf{D} is diagonal.
- Let \mathbf{L} be the unnormalized **graph Laplacian** $\mathbf{L} = \mathbf{D} - \mathbf{S}$
- So we want to minimize $\mathbf{y}^\top \mathbf{L} \mathbf{y}$ with respect to \mathbf{y}_u , the labels of the unannotated instances.

In principle, this is easily solveable:

- Partition the Laplacian $\mathbf{L} = \begin{bmatrix} \mathbf{L}_{\ell\ell} & \mathbf{L}_{\ell u} \\ \mathbf{L}_{u\ell} & \mathbf{L}_{uu} \end{bmatrix}$
- Then the closed form solution is $\mathbf{y}_u = -\mathbf{L}_{uu}^{-1} \mathbf{L}_{u\ell} \mathbf{y}_\ell$

(c) Jacob Eisenstein 2014-2017. Work in progress.

- This is great ... if we can invert \mathbf{L}_{uu} .

In practice, $\mathbf{L}_{u,u}$ is huge, so we can't invert it unless it has special structure. Zhu and Ghahramani (2002) propose an iterative solution called **label propagation**.

- Let $\mathbf{T}_{ij} = \frac{s_{ij}}{\sum_k s_{kj}}$, row-normalizing \mathbf{S} .
- Let \mathbf{Y} be an $n \times C$ matrix of labels, where C is the number of classes.
- Until tired,
 - Set $\mathbf{Y} = \mathbf{T}\mathbf{Y}$
 - Row-normalize \mathbf{Y}
 - Clamp the seed examples in \mathbf{Y} to their original values
- There's a flavor of EM here, with \mathbf{Y} representing our belief $q_i(y_i)$. But there's no M-step in which we update model parameters. That's because we're in a graph-based framework, and we're assuming the graph is correct.

Both mincut and label propagation are **transductive** learning algorithms: they learn jointly over the training and test data. This is fine in some settings, but not if you want to train a system and then apply it to new test data later — you'd have to retrain it all over again.

Manifold regularization (Belkin et al., 2006) addresses this issue, by learning functions that are smooth on the “graph manifold.” In practice, this means that they give similar labels to nearby datapoints in the unlabeled data. Suppose we are interested in learning a classification function f . Then we can optimize:

$$\operatorname{argmin}_f \frac{1}{\ell} \sum_i \ell(f(\mathbf{x}_i), y_i) + \lambda_1 \|\mathbf{f}\|^2 + \lambda_2 \sum_{i,j} s_{ij} (f(\mathbf{x}_i) - f(\mathbf{x}_j))^2$$

- The first term corresponds to the loss on the labeled training data; we can use any convex loss functions, such as logistic or hinge loss.
- The second term corresponds to the smoothness, akin to regularizing the weights in a linear classifier.
- The third term penalizes making different predictions for similar instances in the unlabeled data

The representer theorem guarantees that we can solve for f as long as ℓ is convex. We can then apply f to any new unlabeled test data.

(c) Jacob Eisenstein 2014-2017. Work in progress.

19.2 Domain adaptation

In domain adaptation, we have a lot of labeled data, but it's in the wrong domain. Some features will be shared across domains. For example, if we are classifying movies or toasters, *good* is a good word, and *sucks* is a bad word. But as we've seen, real review text is usually more subtle. What about a word like *unpredictable*? This is a good word for a movie, but not such a good word for a kitchen appliance.

Supervised domain adaptation

In supervised domain adaptation (transfer learning), we have:

- Lots of labeled data in a “source” domain, $\{(x_i, y_i)\}_{i=1}^{\ell_S} \sim \mathcal{D}_S$ (e.g., reviews of restaurants)
- A little labeled data in a “target” domain, $\{(x_i, y_i)\}_{i=1}^{\ell_T} \sim \mathcal{D}_T$ (e.g., reviews of chess stores)

Here are some (surprisingly-competitive) baselines (see slides)

- Source-only: train on the source data, apply it to the target data.
- Target-only: forget the source data, just train on the limited target data.
- Big blob: merge the source and target data into a single training set. Optionally downweight the source data.
- Prediction: train a classifier on the source data, use its prediction as a feature in the target data.
- Interpolation: train two classifiers, combine their outputs

Here are two less-obvious approaches:

Priors :

Train a (logistic-regression) classifier on the source data. Treat its weights as the priors on the target data, and regularize towards these weights rather than towards zero (Chelba and Acero 2004).

Feature augmentation Create **copies** of each feature, for each domain and for the cross-domain setting.

- The copies fire in the appropriate domains, and the learning algorithm decides whether a feature is general or domain-specific.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- For example, suppose we have domains for Appliances and Movies, and features *outstanding* and *sturdy*. We replicate the features, obtaining

$$\begin{aligned} &\langle \textit{outstanding}, \text{APP.} \rangle, \langle \textit{outstanding}, \text{MOV.} \rangle, \langle \textit{outstanding}, \text{ALL} \rangle \\ &\langle \textit{sturdy}, \text{APP.} \rangle, \langle \textit{sturdy}, \text{MOV.} \rangle, \langle \textit{sturdy}, \text{ALL} \rangle \end{aligned}$$

- Ideally, we will learn a positive weight for $\langle \textit{outstanding}, \text{ALL} \rangle$, because the feature works in both domains, and a small weight for the domain-specific copies of the *outstanding* feature.
- We will also learn a positive weight for $\langle \textit{sturdy}, \text{APP.} \rangle$, because the feature works only in the Appliance domain.

See slides for a diagram of how this works.

Unsupervised domain adaptation

Without labeled data in the target domain, can we learn anything? If the source and target domain are somewhat related, then we can. A very popular approach is structural correspondence learning (SCL) (Blitzer et al., 2007).

- Suppose there are a few words that are good predictors in both domains; we'll call these **pivot features**
- Pivot features can be selected by finding words that are
 - Popular in both domains
 - High mutual-information with the label in the source domain
- The label is unknown in the target domain, so we can't learn to predict it. Instead we'll predict the pivots. We train a linear classifier for each pivot, obtaining weights θ_n for pivot n .
- For example, we can learn that the domain-specific feature *fast-multicore* is a good predictor of the pivot *excellent*.
- We can horizontally concatenate the pivot predictor weights, forming

$$\Theta = [\theta_1, \theta_2, \dots, \theta_N] \quad (19.6)$$

- The matrix Θ is large, and contains redundant information (since many pivots are closely related to each other). We factor $\Theta \approx USV^T$ using singular value decomposition (SVD).
- We use U to **project** features from both domains into a shared space, $U^\top x$.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- We then learn to predict the label in the source domain, using the augmented instance $\langle \mathbf{x}, U^\top \mathbf{x} \rangle$. In U contains meaningful correspondences between the domains, then the weights learned on these features will work for the target domain instances too.
- This idea yields substantial improvements in adapting sentiment classifiers across product domains, e.g., books, movies, and appliances (Blitzer et al., 2007).

See the slides for a graphical explanation of these ideas, with slightly different notation.

19.3 Other learning settings

There are many other settings in which we learn from something other than in-domain labeled data:

- **Active learning.** The model can query the annotator for labels (see above)
- **Feature labeling.** Annotators label *features* rather than instances. For example, you provide a list of five prototype words for each POS tag (Haghighi and Klein, 2006).
- **Feature expectations.** Learn from *constraints* on feature-label relationships; for example, the word “the” is a determiner at least 90% of the time. In EMNLP 2013, this idea was applied to multilingual learning (which I’ll discuss in the final lecture). The basic idea of this paper is to align words between sentences and insist that aligned words have the same tag most of the time.
- **Multi-instance learning.** The learner gets a “bag” of instances, and a label. If the label is positive, then at least one instance in the bag is positive, but you don’t know which one.

This idea is often related to **distant supervision**. The learner gets a label indicating that there is a relationship, such as BORN-IN(OBAMA, HAWAII), and a set of instances containing sentences that mention the two arguments, *Obama* and *Hawaii*. Many of these sentences do not actually instantiate the desired relation (e.g., *Obama visited Hawaii in 2008...*), but we assume that at least one does, and we must learn from this.

Chapter 20

Beyond linear models

20.1 Representation learning

20.2 Convolutional neural networks

20.3 Recursive neural networks

20.4 Encoder-decoder models

20.5 Structure prediction

Recently, several researchers have applied neural networks and other distributed representations to dependency parsing. These methods diverge from the approach of scoring edges by the inner product of a weight vector with a large, sparse feature vector. Instead, each word is represented by a small, dense **embedding** vector, which may be estimated from unlabeled data in a preprocessing step. These embeddings are typically used in combination with transition-based dependency parsers, either as features (Bansal et al., 2014), or as part of an integrated neural network parsing model (Henderson et al., 2008; Chen and Manning, 2014; Dyer et al., 2015). These models are described in more detail in chapter 20. Embeddings can also be learned for features (rather than for words) in a graph-based parsing algorithm (Lei et al., 2014).

Bibliography

- Abadi, M., Agarwal, A., Barham, P., Brevdo, E., Chen, Z., Citro, C., Corrado, G. S., Davis, A., Dean, J., Devin, M., Ghemawat, S., Goodfellow, I. J., Harp, A., Irving, G., Isard, M., Jia, Y., Józefowicz, R., Kaiser, L., Kudlur, M., Levenberg, J., Mané, D., Monga, R., Moore, S., Murray, D. G., Olah, C., Schuster, M., Shlens, J., Steiner, B., Sutskever, I., Talwar, K., Tucker, P. A., Vanhoucke, V., Vasudevan, V., Viégas, F. B., Vinyals, O., Warden, P., Wattenberg, M., Wicke, M., Yu, Y., and Zheng, X. (2016). Tensorflow: Large-scale machine learning on heterogeneous distributed systems. *CoRR*, abs/1603.04467.
- Abney, S. P. and Johnson, M. (1991). Memory requirements and local ambiguities of parsing strategies. *Journal of Psycholinguistic Research*, 20(3):233–250.
- Akaike, H. (1974). A new look at the statistical model identification. *Automatic Control, IEEE Transactions on*, 19(6):716–723.
- Akmajian, A., Demers, R. A., Farmer, A. K., and Harnish, R. M. (2010). *Linguistics: An introduction to language and communication*. MIT press, Cambridge, MA, sixth edition.
- Allauzen, C., Riley, M., and Schalkwyk, J. (2009). A generalized composition algorithm for weighted finite-state transducers. In *INTERSPEECH*.
- Allauzen, C., Riley, M., Schalkwyk, J., Skut, W., and Mohri, M. (2007). Openfst: A general and efficient weighted finite-state transducer library. In *International Conference on Implementation and Application of Automata*, pages 11–23. Springer.
- Alm, C. O., Roth, D., and Sproat, R. (2005). Emotions from text: machine learning for text-based emotion prediction. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 579–586.
- Anand, P., Walker, M., Abbott, R., Fox Tree, J. E., Bowmani, R., and Minor, M. (2011). Cats rule and dogs drool!: Classifying stance in online debate. In *Proceedings of the 2nd Workshop on Computational Approaches to Subjectivity and Sentiment Analysis (WASSA 2.011)*, pages 1–9, Portland, Oregon. Association for Computational Linguistics.

- Anandkumar, A., Ge, R., Hsu, D., Kakade, S. M., and Telgarsky, M. (2014). Tensor decompositions for learning latent variable models. *The Journal of Machine Learning Research*, 15(1):2773–2832.
- Andor, D., Alberti, C., Weiss, D., Severyn, A., Presta, A., Ganchev, K., Petrov, S., and Collins, M. (2016). Globally normalized transition-based neural networks. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 2442–2452, Berlin.
- Andreas, J. and Klein, D. (2015). When and why are log-linear models self-normalizing? In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 244–249, Denver, CO.
- Arora, S., Ge, R., Halpern, Y., Mimno, D., Moitra, A., Sontag, D., Wu, Y., and Zhu, M. (2013). A practical algorithm for topic modeling with provable guarantees. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 280–288.
- Bansal, M., Gimpel, K., and Livescu, K. (2014). Tailoring continuous word representations for dependency parsing. In *Proceedings of the Association for Computational Linguistics (ACL)*, Baltimore, MD.
- Belkin, M., Niyogi, P., and Sindhvani, V. (2006). Manifold regularization: A geometric framework for learning from labeled and unlabeled examples. *The Journal of Machine Learning Research*, 7:2399–2434.
- Bender, E. M. (2013). *Linguistic Fundamentals for Natural Language Processing: 100 Essentials from Morphology and Syntax*, volume 6 of *Synthesis Lectures on Human Language Technologies*. Morgan & Claypool Publishers.
- Bengio, Y., Ducharme, R., Vincent, P., and Janvin, C. (2003). A neural probabilistic language model. *The Journal of Machine Learning Research*, 3:1137–1155.
- Bengio, Y., Simard, P., and Frasconi, P. (1994). Learning long-term dependencies with gradient descent is difficult. *Neural Networks, IEEE Transactions on*, 5(2):157–166.
- Bengtson, E. and Roth, D. (2008). Understanding the value of features for coreference resolution. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 294–303, Honolulu, HI.
- Berg-Kirkpatrick, T., Bouchard-Côté, A., DeNero, J., and Klein, D. (2010). Painless unsupervised learning with features. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 582–590, Los Angeles, CA.
- Berger, A. L., Pietra, V. J. D., and Pietra, S. A. D. (1996). A maximum entropy approach to natural language processing. *Computational linguistics*, 22(1):39–71.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Bergsma, S., Lin, D., and Goebel, R. (2008). Distributional identification of non-referential pronouns. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 10–18, Columbus, OH.
- Bergstra, J., Breuleux, O., Bastien, F., Lamblin, P., Pascanu, R., Desjardins, G., Turian, J., Warde-Farley, D., and Bengio, Y. (2010). Theano: A cpu and gpu math compiler in python. In *Proceedings of the 9th Python in Science Conf*, pages 1–7.
- Bhatia, P., Guthrie, R., and Eisenstein, J. (2016). Morphological priors for probabilistic neural word embeddings. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- Bikel, D. M. (2004). Intricacies of Collins’ parsing model. *Computational Linguistics*, 30(4):479–511.
- Bishop, C. M. (2006). *Pattern recognition and machine learning*. springer.
- Björkelund, A. and Kuhn, J. (2014). Learning structured perceptrons for coreference resolution with latent antecedents and non-local features. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 47–57, Baltimore, MD.
- Blackburn, P. and Bos, J. (2005). *Representation and inference for natural language: A first course in computational semantics*. CSLI.
- Blei, D. M. (2012). Probabilistic topic models. *Communications of the ACM*, 55(4):77–84.
- Blei, D. M. (2014). Build, compute, critique, repeat: Data analysis with latent variable models. *Annual Review of Statistics and Its Application*, 1:203–232.
- Blitzer, J., Dredze, M., and Pereira, F. (2007). Biographies, bollywood, boom-boxes and blenders: Domain adaptation for sentiment classification. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 440–447, Prague.
- Blum, A., Lafferty, J., Rwebangira, M. R., and Reddy, R. (2004). Semi-supervised learning using randomized mincuts. In *icml*, page 13.
- Blum, A. and Mitchell, T. (1998). Combining labeled and unlabeled data with co-training. In *Proceedings of the eleventh annual conference on Computational learning theory*, pages 92–100. ACM.
- Böhmová, A., Hajič, J., Hajičová, E., and Hladká, B. (2003). The prague dependency treebank. In *Treebanks*, pages 103–127. Springer.
- Bontcheva, K., Dimitrov, M., Maynard, D., Tablan, V., and Cunningham, H. (2002). Shallow methods for named entity coreference resolution. In *Chames de références et résolveurs danaphores, workshop TALN*.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Charniak, E. and Johnson, M. (2005). Coarse-to-fine n-best parsing and maxent discriminative reranking. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 173–180, Ann Arbor, Michigan.
- Chen, D. and Manning, C. D. (2014). A fast and accurate dependency parser using neural networks. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 740–750.
- Chen, S. F. and Goodman, J. (1999). An empirical study of smoothing techniques for language modeling. *Computer Speech & Language*, 13(4):359–393.
- Cho, K. (2015). Natural language understanding with distributed representation. *CoRR*, abs/1511.07916.
- Cho, K., Van Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H., and Bengio, Y. (2014). Learning phrase representations using rnn encoder-decoder for statistical machine translation. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- Chu, Y.-J. and Liu, T.-H. (1965). On shortest arborescence of a directed graph. *Scientia Sinica*, 14(10):1396.
- Chung, C. and Pennebaker, J. W. (2007). The psychological functions of function words. *Social communication*, pages 343–359.
- Clark, K. and Manning, C. D. (2015). Entity-centric coreference resolution with model stacking. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 1405–1415, Beijing.
- Clarke, J., Goldwasser, D., Chang, M.-W., and Roth, D. (2010). Driving semantic parsing from the world’s response. In *CONLL*, pages 18–27. Association for Computational Linguistics.
- Cohen, S. B., Gómez-Rodríguez, C., and Satta, G. (2012). Elimination of spurious ambiguity in transition-based dependency parsing. *CoRR*, abs/1206.6735.
- Collins, M. (1997). Three generative, lexicalised models for statistical parsing. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 16–23.
- Collins, M. (2002). Discriminative training methods for hidden markov models: theory and experiments with perceptron algorithms. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 1–8.
- Collins, M. (2003). Head-driven statistical models for natural language parsing. *Computational linguistics*, 29(4):589–637.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Collins, M. (2013). Notes on natural language processing. <http://www.cs.columbia.edu/~mcollins/notes-spring2013.html>.
- Collins, M. and Brooks, J. (1995). Prepositional phrase attachment through a backed-off model. In *Workshop on Very Large Corpora*.
- Collins, M. and Koo, T. (2005). Discriminative reranking for natural language parsing. *Computational Linguistics*, 31(1):25–70.
- Collins, M. and Singer, Y. (1999). Unsupervised models for named entity classification. In *emnlp*, pages 189–196.
- Collobert, R., Kavukcuoglu, K., and Farabet, C. (2011). Torch7: A matlab-like environment for machine learning. Technical Report EPFL-CONF-192376, EPFL.
- Cormen, T. H., Leiserson, C. E., Rivest, R. L., and Stein, C. (2009). *Introduction to algorithms*. MIT press, third edition.
- Coviello, L., Sohn, Y., Kramer, A. D., Marlow, C., Franceschetti, M., Christakis, N. A., and Fowler, J. H. (2014). Detecting emotional contagion in massive social networks. *PloS one*, 9(3):e90315.
- Crammer, K., Dekel, O., Keshet, J., Shalev-Shwartz, S., and Singer, Y. (2006). Online passive-aggressive algorithms. *The Journal of Machine Learning Research*, 7:551–585.
- Crammer, K. and Singer, Y. (2001). Pranking with ranking. In *Neural Information Processing Systems (NIPS)*, pages 641–647, Vancouver.
- Crammer, K. and Singer, Y. (2003). Ultraconservative online algorithms for multiclass problems. *The Journal of Machine Learning Research*, 3:951–991.
- Cui, H., Sun, R., Li, K., Kan, M.-Y., and Chua, T.-S. (2005). Question answering passage retrieval using dependency relations. In *Proceedings of the 28th annual international ACM SIGIR conference on Research and development in information retrieval*, pages 400–407. ACM.
- Culotta, A. and Sorensen, J. (2004). Dependency tree kernels for relation extraction. In *Proceedings of the Association for Computational Linguistics (ACL)*.
- Das, D., Chen, D., Martins, A. F., Schneider, N., and Smith, N. A. (2014). Frame-semantic parsing. *Computational Linguistics*, 40(1):9–56.
- Das, D., Martins, A. F., and Smith, N. A. (2012). An exact dual decomposition algorithm for shallow semantic parsing with constraints. In *Proceedings of SEMEVAL*, pages 209–217.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Das, D., Schneider, N., Chen, D., and Smith, N. A. (2010). Probabilistic frame-semantic parsing. In *Human language technologies: The 2010 annual conference of the North American chapter of the association for computational linguistics*, pages 948–956. Association for Computational Linguistics.
- De Marneffe, M.-C., Dozat, T., Silveira, N., Haverinen, K., Ginter, F., Nivre, J., and Manning, C. D. (2014). Universal stanford dependencies: A cross-linguistic typology. In *LREC*, pages 4585–4592.
- De Marneffe, M.-C., MacCartney, B., and Manning, C. D. (2006). Generating typed dependency parses from phrase structure parses. In *LREC*, pages 449–454.
- De Marneffe, M.-C. and Manning, C. D. (2008). The stanford typed dependencies representation. In *Coling 2008: Proceedings of the workshop on Cross-Framework and Cross-Domain Parser Evaluation*, pages 1–8. Association for Computational Linguistics.
- Deerwester, S. C., Dumais, S. T., Landauer, T. K., Furnas, G. W., and Harshman, R. A. (1990). Indexing by latent semantic analysis. *JASIS*, 41(6):391–407.
- Denis, P. and Baldridge, J. (2007). A ranking approach to pronoun resolution. In *IJCAI*.
- Denis, P. and Baldridge, J. (2008). Specialized models and ranking for coreference resolution. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing, EMNLP '08*, pages 660–669, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Derrida, J. (1985). Des tours de babel. In Graham, J., editor, *Difference in translation*. Cornell University Press, Ithaca, NY.
- Dowty, D. (1991). Thematic proto-roles and argument selection. *Language*, pages 547–619.
- Duchi, J., Hazan, E., and Singer, Y. (2011). Adaptive subgradient methods for online learning and stochastic optimization. *The Journal of Machine Learning Research*, 12:2121–2159.
- Durrett, G. and Klein, D. (2013). Easy victories and uphill battles in coreference resolution. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*.
- Durrett, G. and Klein, D. (2015). Neural crf parsing. In *Proceedings of the Association for Computational Linguistics (ACL)*, Beijing.
- Dyer, C. (2014). Notes on adagrad. www.ark.cs.cmu.edu/cdyer/adagrad.pdf.
- Dyer, C., Ballesteros, M., Ling, W., Matthews, A., and Smith, N. A. (2015). Transition-based dependency parsing with stack long short-term memory. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 334–343, Beijing.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Edmonds, J. (1967). Optimum branchings. *Journal of Research of the National Bureau of Standards B*, 71(4):233–240.
- Eisner, J. M. (1996). Three new probabilistic models for dependency parsing: An exploration. In *COLING*, pages 340–345.
- Ekman, P. (1992). Are there basic emotions?
- Esuli, A. and Sebastiani, F. (2006). Sentiwordnet: A publicly available lexical resource for opinion mining. In *LREC*, volume 6, pages 417–422. Citeseer.
- Fellbaum, C. (2010). *WordNet*. Springer.
- Figueiredo, M., Graça, J., Martins, A., Almeida, M., and Coelho, L. P. (2013). LXMLS lab guide. <http://lxmls.it.pt/2013/guide.pdf>.
- Fillmore, C. J. (1968). The case for case. In Bach, E. and Harms, R., editors, *Universals in linguistic theory*. Holt, Rinehart, and Winston.
- Finkel, J. R., Grenager, T., and Manning, C. (2005). Incorporating non-local information into information extraction systems by gibbs sampling. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 363–370, Ann Arbor, Michigan.
- Finkel, J. R., Grenager, T., and Manning, C. D. (2007). The infinite tree. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 272–279, Prague.
- Finkel, J. R., Kleeman, A., and Manning, C. D. (2008). Efficient, feature-based, conditional random field parsing. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 959–967, Columbus, OH.
- Finkel, J. R. and Manning, C. D. (2008). Enforcing transitivity in coreference resolution. In *Proceedings of the 46th Annual Meeting of the Association for Computational Linguistics on Human Language Technologies: Short Papers*, pages 45–48. Association for Computational Linguistics.
- Firth, J. R. (1957). *Papers in Linguistics 1934-1951*. Oxford University Press.
- Foster, J., Cetinoglu, O., Wagner, J., Le Roux, J., Nivre, J., Hogan, D., and van Genabith, J. (2011). From news to comment: Resources and benchmarks for parsing the language of web 2.0. In *Proceedings of the International Joint Conference on Natural Language Processing (IJCNLP)*, pages 893–901, Chiang Mai, Thailand. Asian Federation of Natural Language Processing.
- Francis, W. N. (1964). A standard sample of present-day english for use with digital computers. Report to the U.S Office of Education on Cooperative Research Project No. E-007.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Freund, Y., Schapire, R., and Abe, N. (1999). A short introduction to boosting. *Journal-Japanese Society For Artificial Intelligence*, 14(771-780):1612.
- Freund, Y. and Schapire, R. E. (1999). Large margin classification using the perceptron algorithm. *Machine learning*, 37(3):277–296.
- Fromkin, V., Rodman, R., and Hyams, N. (2013). *An introduction to language*. Cengage Learning.
- Fundel, K., Küffner, R., and Zimmer, R. (2007). Rellexrelation extraction using dependency parse trees. *Bioinformatics*, 23(3):365–371.
- Gabrilovich, E. and Markovitch, S. (2007). Computing semantic relatedness using wikipedia-based explicit semantic analysis. In *Proceedings of the International Joint Conference on Artificial Intelligence (IJCAI)*, volume 7, pages 1606–1611.
- Galley, M. (2006). A skip-chain conditional random field for ranking meeting utterances by importance. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 364–372.
- Ganchev, K. and Dredze, M. (2008). Small statistical models by random feature mixing. In *Proceedings of the ACL08 HLT Workshop on Mobile Language Processing*, pages 19–20.
- Gao, J., Andrew, G., Johnson, M., and Toutanova, K. (2007). A comparative study of parameter estimation methods for statistical natural language processing. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 824–831, Prague.
- Ge, D., Jiang, X., and Ye, Y. (2011). A note on the complexity of l_p minimization. *Mathematical programming*, 129(2):285–299.
- Geach, P. T. (1962). *Reference and generality: An examination of some medieval and modern theories*. Cornell University Press.
- Gildea, D. and Jurafsky, D. (2002). Automatic labeling of semantic roles. *Computational linguistics*, 28(3):245–288.
- Gimpel, K., Schneider, N., O’Connor, B., Das, D., Mills, D., Eisenstein, J., Heilman, M., Yogatama, D., Flanigan, J., and Smith, N. A. (2011). Part-of-speech tagging for Twitter: annotation, features, and experiments. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 42–47, Portland, OR.
- Goldberg, Y. (2015). A primer on neural network models for natural language processing. *arXiv preprint arXiv:1510.00726*.
- Goldwater, S. and Griffiths, T. (2007). A fully bayesian approach to unsupervised part-of-speech tagging. In *Annual meeting-association for computational linguistics*, volume 45.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Goodfellow, I., Bengio, Y., and Courville, A. (2016). *Deep learning*. MIT Press.
- Graves, A., Mohamed, A.-r., and Hinton, G. (2013). Speech recognition with deep recurrent neural networks. In *2013 IEEE international conference on acoustics, speech and signal processing*, pages 6645–6649. IEEE.
- Graves, A. and Schmidhuber, J. (2005). Framewise phoneme classification with bidirectional lstm and other neural network architectures. *Neural Networks*, 18(5):602–610.
- Grishman, R., Macleod, C., and Sterling, J. (1992). Evaluating parsing strategies using standardized parse files. In *Proceedings of the third conference on Applied natural language processing*, pages 156–161. Association for Computational Linguistics.
- Grishman, R. and Sundheim, B. (1996). Message understanding conference-6: A brief history. In *Proceedings of the International Conference on Computational Linguistics (COLING)*, volume 96, pages 466–471.
- Grosz, B. J., Weinstein, S., and Joshi, A. K. (1995). Centering: A framework for modeling the local coherence of discourse. *Computational linguistics*, 21(2):203–225.
- Gutmann, M. U. and Hyvärinen, A. (2012). Noise-contrastive estimation of unnormalized statistical models, with applications to natural image statistics. *The Journal of Machine Learning Research*, 13(1):307–361.
- Haghighi, A. and Klein, D. (2006). Prototype-driven learning for sequence models. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 320–327, New York, NY.
- Haghighi, A. and Klein, D. (2009). Simple coreference resolution with rich syntactic and semantic features. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 1152–1161, Singapore.
- Hannak, A., Anderson, E., Barrett, L. F., Lehmann, S., Mislove, A., and Riedewald, M. (2012). Tweetin’ in the rain: Exploring societal-scale effects of weather on mood. In *Proceedings of the International Conference on Web and Social Media (ICWSM)*.
- Hastie, T., Tibshirani, R., and Friedman, J. (2009). *The elements of statistical learning*. Springer, New York, second edition.
- Hatzivassiloglou, V. and McKeown, K. R. (1997). Predicting the semantic orientation of adjectives. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 174–181, Madrid, Spain.
- Henderson, J., Merlo, P., Musillo, G., and Titov, I. (2008). A latent variable model of synchronous parsing for syntactic and semantic dependencies. In *CONLL*, pages 178–182.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Hindle, D. and Rooth, M. (1990). Structural ambiguity and lexical relations. In *Proceedings of the Workshop on Speech and Natural Language*.
- Hobbs, J. R. (1978). Resolving pronoun references. *Lingua*, 44(4):311–338.
- Hochreiter, S. and Schmidhuber, J. (1997). Long short-term memory. *Neural computation*, 9(8):1735–1780.
- Hovy, E., Marcus, M., Palmer, M., Ramshaw, L., and Weischedel, R. (2006). Ontonotes: the 90% solution. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 57–60, New York, NY.
- Hsu, D., Kakade, S. M., and Zhang, T. (2012). A spectral algorithm for learning hidden markov models. *Journal of Computer and System Sciences*, 78(5):1460–1480.
- Hu, M. and Liu, B. (2004). Mining and summarizing customer reviews. In *Proceedings of Knowledge Discovery and Data Mining (KDD)*, pages 168–177.
- Huang, E. H., Socher, R., Manning, C. D., and Ng, A. Y. (2012a). Improving word representations via global context and multiple word prototypes. In *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics: Long Papers-Volume 1*, pages 873–882. Association for Computational Linguistics.
- Huang, L., Fayong, S., and Guo, Y. (2012b). Structured perceptron with inexact search. In *Proceedings of the 2012 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 142–151, Montréal, Canada. Association for Computational Linguistics.
- Huang, Z., Xu, W., and Yu, K. (2015). Bidirectional lstm-crf models for sequence tagging. *arXiv preprint arXiv:1508.01991*.
- Ide, N. and Wilks, Y. (2006). Making sense about sense. In *Word sense disambiguation*, pages 47–73. Springer.
- Jain, A. K. (2010). Data clustering: 50 years beyond k-means. *Pattern recognition letters*, 31(8):651–666.
- Jiang, L., Yu, M., Zhou, M., Liu, X., and Zhao, T. (2011). Target-dependent twitter sentiment classification. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 151–160, Portland, OR.
- Jockers, M. L. (2015). Szuzhet? <http://bla.bla.com>.
- Johnson, M. (1998). Pcfg models of linguistic tree representations. *Computational Linguistics*, 24(4):613–632.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Joshi, A. K. and Schabes, Y. (1997). Tree-adjoining grammars. In *Handbook of formal languages*, pages 69–123. Springer.
- Jozefowicz, R., Zaremba, W., and Sutskever, I. (2015). An empirical exploration of recurrent network architectures. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 2342–2350.
- Jurafsky, D. (1996). A probabilistic model of lexical and syntactic access and disambiguation. *Cognitive Science*, 20(2):137–194.
- Jurafsky, D. and Martin, J. H. (2009). *Speech and Language Processing*. Prentice Hall, 2 edition.
- Karlsson, F. (2007). Constraints on multiple center-embedding of clauses. *Journal of Linguistics*, 43(02):365–392.
- Kehler, A. (2007). Rethinking the SMASH approach to pronoun interpretation. In *Interdisciplinary perspectives on reference processing*, New Directions in Cognitive Science Series, pages 95–122. Oxford University Press.
- Kilgarrieff, A. (1997). I don’t believe in word senses. *CoRR*, cmp-lg/9712006.
- Kilgarrieff, A. and Grefenstette, G. (2003). Introduction to the special issue on the web as corpus. *Computational linguistics*, 29(3):333–347.
- Kim, M.-J. (2002). Does korean have adjectives? *MIT Working Papers in Linguistics*, 43:71–89.
- Kim, S.-M. and Hovy, E. (2006). Extracting opinions, opinion holders, and topics expressed in online news media text. In *Proceedings of the Workshop on Sentiment and Subjectivity in Text*, pages 1–8, Sydney, Australia. Association for Computational Linguistics.
- Kim, Y. (2014). Convolutional neural networks for sentence classification. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 1746–1751.
- Klein, D. and Manning, C. D. (2003). Accurate unlexicalized parsing. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 423–430.
- Klenner, M. (2007). Enforcing consistency on coreference sets. In *Recent Advances in Natural Language Processing (RANLP)*, pages 323–328.
- Knight, K. and May, J. (2009). Applications of weighted automata in natural language processing. In *Handbook of Weighted Automata*, pages 571–596. Springer.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Koo, T., Carreras, X., and Collins, M. (2008). Simple semi-supervised dependency parsing. In *Proceedings of ACL-08: HLT*, pages 595–603, Columbus, Ohio. Association for Computational Linguistics.
- Koo, T. and Collins, M. (2010). Efficient third-order dependency parsers. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 1–11, Uppsala, Sweden.
- Koo, T., Globerson, A., Carreras, X., and Collins, M. (2007). Structured prediction models via the matrix-tree theorem. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 141–150.
- Koren, Y., Bell, R., and Volinsky, C. (2009). Matrix factorization techniques for recommender systems. *Computer*, 42(8):30–37.
- Kübler, S., McDonald, R., and Nivre, J. (2009). Dependency parsing. *Synthesis Lectures on Human Language Technologies*, 1(1):1–127.
- Kuhlmann, M. and Nivre, J. (2010). Transition-based techniques for non-projective dependency parsing. *Northern European Journal of Language Technology (NEJLT)*, 2(1):1–19.
- Kummerfeld, J. K., Berg-Kirkpatrick, T., and Klein, D. (2015). An empirical analysis of optimization for max-margin nlp. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- Lafferty, J., McCallum, A., and Pereira, F. (2001). Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *icml*.
- Lample, G., Ballesteros, M., Subramanian, S., Kawakami, K., and Dyer, C. (2016). Neural architectures for named entity recognition. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 260–270, San Diego, CA.
- Lappin, S. and Leass, H. J. (1994). An algorithm for pronominal anaphora resolution. *Computational linguistics*, 20(4):535–561.
- Lari, K. and Young, S. J. (1990). The estimation of stochastic context-free grammars using the inside-outside algorithm. *Computer speech & language*, 4(1):35–56.
- LeCun, Y. and Bengio, Y. (1995). Convolutional networks for images, speech, and time series. *The handbook of brain theory and neural networks*, 3361.
- Lee, H., Chang, A., Peirsman, Y., Chambers, N., Surdeanu, M., and Jurafsky, D. (2013). Deterministic coreference resolution based on entity-centric, precision-ranked rules. *Computational Linguistics*, 39(4):885–916.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Lee, H., Peirsman, Y., Chang, A., Chambers, N., Surdeanu, M., and Jurafsky, D. (2011). Stanford’s multi-pass sieve coreference resolution system at the conll-2011 shared task. In *Proceedings of the Conference on Natural Language Learning (CoNLL)*, pages 28–34. Association for Computational Linguistics.
- Lei, T., Xin, Y., Zhang, Y., Barzilay, R., and Jaakkola, T. (2014). Low-rank tensors for scoring dependency structures. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 1381–1391, Baltimore, MD.
- Lesk, M. (1986). Automatic sense disambiguation using machine readable dictionaries: how to tell a pine cone from an ice cream cone. In *Proceedings of the 5th annual international conference on Systems documentation*, pages 24–26. ACM.
- Levy, O. and Goldberg, Y. (2014a). Dependency-based word embeddings. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 302–308, Baltimore, MD.
- Levy, O. and Goldberg, Y. (2014b). Neural word embedding as implicit matrix factorization. In *Neural Information Processing Systems (NIPS)*, Montréal.
- Levy, R. and Manning, C. (2009). An informal introduction to computational semantics.
- Li, J. and Jurafsky, D. (2015). Do multi-sense embeddings improve natural language understanding? In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 1722–1732.
- Liang, P., Jordan, M. I., and Klein, D. (2013). Learning dependency-based compositional semantics. *Computational Linguistics*, 39(2):389–446.
- Liang, P. and Klein, D. (2009). Online em for unsupervised models. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 611–619, Boulder, CO.
- Liang, P., Petrov, S., Jordan, M. I., and Klein, D. (2007). The infinite pcfg using hierarchical dirichlet processes. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 688–697.
- Lin, D. (1998). Automatic retrieval and clustering of similar words. In *Proceedings of the 17th international conference on Computational linguistics-Volume 2*, pages 768–774. Association for Computational Linguistics.
- Ling, W., Dyer, C., Black, A., and Trancoso, I. (2015a). Two/too simple adaptations of word2vec for syntax problems. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, Denver, CO.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Ling, W., Luís, T., Marujo, L., Astudillo, R. F., Amir, S., Dyer, C., Black, A. W., and Trancoso, I. (2015b). Finding function in form: Compositional character models for open vocabulary word representation. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- Liu, B. (2015). *Sentiment Analysis: Mining Opinions, Sentiments, and Emotions*. Cambridge University Press.
- Liu, D. C. and Nocedal, J. (1989). On the limited memory BFGS method for large scale optimization. *Mathematical programming*, 45(1-3):503–528.
- Ma, X. and Hovy, E. (2016). End-to-end sequence labeling via bi-directional lstm-cnns-crf. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 1064–1074, Berlin.
- Manning, C. D. (2011). Part-of-speech tagging from 97% to 100%: is it time for some linguistics? In *International Conference on Intelligent Text Processing and Computational Linguistics*, pages 171–189. Springer.
- Manning, C. D., Raghavan, P., Schütze, H., et al. (2008). *Introduction to information retrieval*, volume 1. Cambridge university press.
- Manning, C. D. and Schütze, H. (1999). *Foundations of Statistical Natural Language Processing*. MIT press, Cambridge, Massachusetts.
- Marcus, M. P., Marcinkiewicz, M. A., and Santorini, B. (1993). Building a large annotated corpus of English: The Penn Treebank. *Computational Linguistics*, 19(2):313–330.
- Martins, A. F. T., Smith, N. A., and Xing, E. P. (2009). Concise integer linear programming formulations for dependency parsing. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 342–350, Suntec, Singapore.
- Martins, A. F. T., Smith, N. A., Xing, E. P., Aguiar, P. M. Q., and Figueiredo, M. A. T. (2010). Turbo parsers: Dependency parsing by approximate variational inference. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 34–44.
- Martschat, S. and Strube, M. (2015). Latent structures for coreference resolution. *Transactions of the Association for Computational Linguistics*, 3:405–418.
- McCallum, A. and Wellner, B. (2004). Conditional models of identity uncertainty with application to noun coreference. In *NIPS*, pages 905–912.
- McDonald, R., Crammer, K., and Pereira, F. (2005a). Online large-margin training of dependency parsers. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 91–98, Ann Arbor, Michigan.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- McDonald, R., Hannan, K., Neylon, T., Wells, M., and Reynar, J. (2007). Structured models for fine-to-coarse sentiment analysis. In *Proceedings of ACL*.
- McDonald, R., Pereira, F., Ribarov, K., and Hajič, J. (2005b). Non-projective dependency parsing using spanning tree algorithms. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 523–530.
- McLendon, S. (2003). Evidentials in eastern pomo with a comparative survey of the category in other pomoan languages. *TYPOLOGICAL STUDIES IN LANGUAGE*, 54:101–130.
- Meyers, A., Reeves, R., Macleod, C., Szekely, R., Zielinska, V., Young, B., and Grishman, R. (2004). The nombank project: An interim report. In *HLT-NAACL 2004 workshop: Frontiers in corpus annotation*, pages 24–31.
- Mihalcea, R., Chklovski, T. A., and Kilgarrieff, A. (2004). The senseval-3 english lexical sample task. In *Proceedings of SENSEVAL-3*, pages 25–28, Barcelona, Spain. Association for Computational Linguistics.
- Mikolov, T., Chen, K., Corrado, G., and Dean, J. (2013a). Efficient estimation of word representations in vector space. In *Proceedings of International Conference on Learning Representations*.
- Mikolov, T., Deoras, A., Povey, D., Burget, L., and Cernocky, J. (2011). Strategies for training large scale neural network language models. In *Automatic Speech Recognition and Understanding (ASRU), 2011 IEEE Workshop on*, pages 196–201. IEEE.
- Mikolov, T., Karafiát, M., Burget, L., Cernocký, J., and Khudanpur, S. (2010). Recurrent neural network based language model. In *INTERSPEECH*, pages 1045–1048.
- Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., and Dean, J. (2013b). Distributed representations of words and phrases and their compositionality. In *Advances in Neural Information Processing Systems*, pages 3111–3119.
- Mikolov, T., Yih, W.-t., and Zweig, G. (2013c). Linguistic regularities in continuous space word representations. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 746–751.
- Miller, M., Sathi, C., Wiesenenthal, D., Leskovec, J., and Potts, C. (2011). Sentiment flow through hyperlink networks. In *Proceedings of the International Conference on Web and Social Media (ICWSM)*.
- Miller, S., Guinness, J., and Zamanian, A. (2004). Name tagging with word clusters and discriminative training. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 337–342, Boston, MA.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Minka, T. P. (1999). From hidden markov models to linear dynamical systems. Tech. Rep. 531, Vision and Modeling Group of Media Lab, MIT.
- Minsky, M. and Papert, S. (1969). *Perceptrons*. MIT press.
- Mintz, M., Bills, S., Snow, R., and Jurafsky, D. (2009). Distant supervision for relation extraction without labeled data. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 1003–1011, Suntec, Singapore.
- Mnih, A. and Hinton, G. (2007). Three new graphical models for statistical language modelling. In *Proceedings of the 24th international conference on Machine learning, ICML '07*, pages 641–648, New York, NY, USA. ACM.
- Mnih, A. and Teh, Y. W. (2012). A fast and simple algorithm for training neural probabilistic language models. In *Proceedings of the International Conference on Machine Learning (ICML)*.
- Mohri, M., Pereira, F., and Riley, M. (2002). Weighted finite-state transducers in speech recognition. *Computer Speech & Language*, 16(1):69–88.
- Mohri, M., Rostamizadeh, A., and Talwalkar, A. (2012). *Foundations of machine learning*. MIT press.
- Muralidharan, A. and Hearst, M. A. (2013). Supporting exploratory text analysis in literature study. *Literary and linguistic computing*, 28(2):283–295.
- Murphy, K. P. (2012). *Machine Learning: A Probabilistic Perspective*. The MIT Press.
- Nakagawa, T., Inui, K., and Kurohashi, S. (2010). Dependency tree-based sentiment classification using crfs with hidden variables. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 786–794, Los Angeles, CA.
- Navigli, R. (2009). Word sense disambiguation: A survey. *ACM Computing Surveys (CSUR)*, 41(2):10.
- Neal, R. M. and Hinton, G. E. (1998). A view of the em algorithm that justifies incremental, sparse, and other variants. In *Learning in graphical models*, pages 355–368. Springer.
- Nemirovski, A. and Yudin, D. (1978). On Cezari’s convergence of the steepest descent method for approximating saddle points of convex-concave functions. *Soviet Math. Dokl.*
- Neubig, G., Dyer, C., Goldberg, Y., Matthews, A., Ammar, W., Anastasopoulos, A., Balles-teros, M., Chiang, D., Clothiaux, D., Cohn, T., Duh, K., Faruqui, M., Gan, C., Garrette, D., Ji, Y., Kong, L., Kuncoro, A., Kumar, G., Malaviya, C., Michel, P., Oda, Y., Richardson, M., Saphra, N., Swayamdipta, S., and Yin, P. (2017). Dynet: The dynamic neural network toolkit.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Neuhaus, P. and Bröker, N. (1997). The complexity of recognition of linguistically adequate dependency grammars. In *eacl*, pages 337–343.
- Nguyen, D. and Dogruöz, A. S. (2013). Word level language identification in online multilingual communication. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- Nigam, K. and Ghani, R. (2000). Analyzing the effectiveness and applicability of co-training. In *Proceedings of the ninth international conference on Information and knowledge management*, pages 86–93. ACM.
- Nigam, K., McCallum, A. K., Thrun, S., and Mitchell, T. (2000). Text classification from labeled and unlabeled documents using em. *Machine learning*, 39(2-3):103–134.
- Nivre, J. (2004). Incrementality in deterministic dependency parsing. In *Proceedings of the Workshop on Incremental Parsing: Bringing Engineering and Cognition Together*, pages 50–57. Association for Computational Linguistics.
- Nivre, J. (2008). Algorithms for deterministic incremental dependency parsing. *Computational Linguistics*, 34(4):513–553.
- Nivre, J., Agić, Ž., Ahrenberg, L., Aranzabe, M. J., Asahara, M., Atutxa, A., Ballesteros, M., Bauer, J., Bengoetxea, K., Berzak, Y., Bhat, R. A., Bick, E., Börstell, C., Bosco, C., Bouma, G., Bowman, S., Cebirolu Eryiit, G., Celano, G. G. A., Chalub, F., Çöltekin, Ç., Connor, M., Davidson, E., de Marneffe, M.-C., Diaz de Ilarraza, A., Dobrovoljc, K., Dozat, T., Droganova, K., Dwivedi, P., Eli, M., Erjavec, T., Farkas, R., Foster, J., Freitas, C., Gajdošová, K., Galbraith, D., Garcia, M., Gärdenfors, M., Garza, S., Ginter, F., Goenaga, I., Gojenola, K., Gökrmak, M., Goldberg, Y., Gómez Guinovart, X., Gonzáles Saavedra, B., Grioni, M., Grūžitis, N., Guillaume, B., Hajič, J., Hà M, L., Haug, D., Hladká, B., Ion, R., Irimia, E., Johannsen, A., Jørgensen, F., Kaşkara, H., Kanayama, H., Kanerva, J., Katz, B., Kenney, J., Kotsyba, N., Krek, S., Laippala, V., Lam, L., Lê Hng, P., Lenci, A., Ljubešić, N., Ljashevskaya, O., Lynn, T., Makazhanov, A., Manning, C., Măranduc, C., Mareček, D., Martínez Alonso, H., Martins, A., Mašek, J., Matsumoto, Y., McDonald, R., Missilä, A., Mititelu, V., Miyao, Y., Montemagni, S., Mori, K. S., Mori, S., Moskalevskiy, B., Muischnek, K., Mustafina, N., Müürisep, K., Nguyn Th, L., Nguyn Th Minh, H., Nikolaev, V., Nurmi, H., Osenova, P., Östling, R., Øvrelid, L., Paiva, V., Pascual, E., Passarotti, M., Perez, C.-A., Petrov, S., Piitulainen, J., Plank, B., Popel, M., Pretkalnia, L., Prokopidis, P., Puolakainen, T., Pyysalo, S., Rademaker, A., Ramasamy, L., Real, L., Rituma, L., Rosa, R., Saleh, S., Saulite, B., Schuster, S., Seeker, W., Seraji, M., Shakurova, L., Shen, M., Silveira, N., Simi, M., Simionescu, R., Simkó, K., Šimková, M., Simov, K., Smith, A., Spadine, C., Suhr, A., Sulubacak, U., Szántó, Z., Tanaka, T., Tsarfaty, R., Tyers, F., Uematsu, S., Uria, L., van Noord, G., Varga, V., Vincze, V., Wallin, L., Wang, J. X., Washington, J. N., Wirén, M., Žabokrtský, Z., Zeldes, A., Zeman, D., and Zhu, H. (2016).

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Universal dependencies 1.4. LINDAT/CLARIN digital library at the Institute of Formal and Applied Linguistics, Charles University in Prague.
- Nivre, J., Hall, J., Nilsson, J., Chanev, A., Eryigit, G., Kübler, S., Marinov, S., and Marsi, E. (2007). Maltparser: A language-independent system for data-driven dependency parsing. *Natural Language Engineering*, 13(2):95–135.
- Nivre, J. and Nilsson, J. (2005). Pseudo-projective dependency parsing. In *Proceedings of the 43rd Annual Meeting on Association for Computational Linguistics*, pages 99–106. Association for Computational Linguistics.
- Novikoff, A. B. (1962). On convergence proofs on perceptrons. In *Proceedings of the Symposium on the Mathematical Theory of Automata*, volume 12, pages 615–622.
- Och, F. J. and Ney, H. (2003). A systematic comparison of various statistical alignment models. *Computational linguistics*, 29(1):19–51.
- Owoputi, O., OConnor, B., Dyer, C., Gimpel, K., and Schneider, N. (2012). Part-of-speech tagging for twitter: Word clusters and other advances. Technical Report CMU-ML-12-107, Carnegie Mellon University.
- Pak, A. and Paroubek, P. (2010). Twitter as a corpus for sentiment analysis and opinion mining. In *LREC*, volume 10, pages 1320–1326.
- Palmer, M., Gildea, D., and Kingsbury, P. (2005). The proposition bank: An annotated corpus of semantic roles. *Computational linguistics*, 31(1):71–106.
- Pang, B. and Lee, L. (2004). A sentimental education: Sentiment analysis using subjectivity summarization based on minimum cuts. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 271–278.
- Pang, B. and Lee, L. (2005). Seeing stars: Exploiting class relationships for sentiment categorization with respect to rating scales. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 115–124, Ann Arbor, Michigan.
- Pang, B. and Lee, L. (2008). Opinion mining and sentiment analysis. *Foundations and trends in information retrieval*, 2(1-2):1–135.
- Pang, B., Lee, L., and Vaithyanathan, S. (2002). Thumbs up?: sentiment classification using machine learning techniques. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 79–86.
- Pantel, P. and Lin, D. (2002). Discovering word senses from text. In *KDD*, pages 613–619. ACM.
- (c) Jacob Eisenstein 2014-2017. Work in progress.

- Pascanu, R., Mikolov, T., and Bengio, Y. (2013). On the difficulty of training recurrent neural networks. In *Proceedings of the 30th International Conference on Machine Learning (ICML-13)*, pages 1310–1318.
- Pennington, J., Socher, R., and Manning, C. (2014). Glove: Global vectors for word representation. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 1532–1543.
- Pereira, F. and Schabes, Y. (1992). Inside-outside reestimation from partially bracketed corpora. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 128–135.
- Pereira, F., Tishby, N., and Lee, L. (1993). Distributional clustering of english words. In *Proceedings of the 31st annual meeting on Association for Computational Linguistics*, pages 183–190. Association for Computational Linguistics.
- Petrov, S., Barrett, L., Thibaux, R., and Klein, D. (2006). Learning accurate, compact, and interpretable tree annotation. In *Proceedings of the Association for Computational Linguistics (ACL)*.
- Petrov, S., Das, D., and McDonald, R. (2012). A universal part-of-speech tagset. In *Proceedings of LREC*.
- Polanyi, L. and Zaenen, A. (2006). Contextual valence shifters. In *Computing attitude and affect in text: Theory and applications*. Springer.
- Popel, M., Marecek, D., Stepánek, J., Zeman, D., and Zabokrtský, Z. (2013). Coordination structures in dependency treebanks. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 517–527, Sophia, Bulgaria.
- Pradhan, S., Ramshaw, L., Marcus, M., Palmer, M., Weischedel, R., and Xue, N. (2011). Conll-2011 shared task: Modeling unrestricted coreference in ontonotes. In *Proceedings of the Fifteenth Conference on Computational Natural Language Learning: Shared Task*, pages 1–27. Association for Computational Linguistics.
- Punyakanok, V., Roth, D., and Yih, W.-t. (2008). The importance of syntactic parsing and inference in semantic role labeling. *Computational Linguistics*, 34(2):257–287.
- Qiu, G., Liu, B., Bu, J., and Chen, C. (2011). Opinion word expansion and target extraction through double propagation. *Computational linguistics*, 37(1):9–27.
- Rabiner, L. (1989). A tutorial on hidden markov models and selected applications in speech recognition. *Proceedings of the IEEE*, 77(2):257–286.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Rao, D. and Ravichandran, D. (2009). Semi-supervised polarity lexicon induction. In *Proceedings of the European Chapter of the Association for Computational Linguistics (EACL)*, pages 675–682.
- Ratinov, L. and Roth, D. (2009). Design challenges and misconceptions in named entity recognition. In *Proceedings of the Thirteenth Conference on Computational Natural Language Learning*, pages 147–155. Association for Computational Linguistics.
- Ratnaparkhi, A. (1996). A maximum entropy model for part-of-speech tagging. In *emnlp*, pages 133–142.
- Ratnaparkhi, A., Reynar, J., and Roukos, S. (1994). A maximum entropy model for prepositional phrase attachment. In *Proceedings of the workshop on Human Language Technology*, pages 250–255. Association for Computational Linguistics.
- Read, J. (2005). Using emoticons to reduce dependency in machine learning techniques for sentiment classification. In *Proceedings of the ACL student research workshop*, pages 43–48. Association for Computational Linguistics.
- Reisinger, J. and Mooney, R. J. (2010). Multi-prototype vector-space models of word meaning. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 109–117, Los Angeles, CA.
- Resnik, P. and Smith, N. A. (2003). The web as a parallel corpus. *Computational Linguistics*, 29(3):349–380.
- Ribeiro, F. N., Araújo, M., Gonçalves, P., Gonçalves, M. A., and Benevenuto, F. (2016). Sentibench—a benchmark comparison of state-of-the-practice sentiment analysis methods. *EPJ Data Science*, 5(1):1–29.
- Riloff, E. (1996). Automatically generating extraction patterns from untagged text. In *Proceedings of the national conference on artificial intelligence*, pages 1044–1049.
- Riloff, E. and Wiebe, J. (2003). Learning extraction patterns for subjective expressions. In *Proceedings of the 2003 conference on Empirical methods in natural language processing*, pages 105–112. Association for Computational Linguistics.
- Roark, B., Saraclar, M., and Collins, M. (2007). Discriminative n_i/i_i -gram language modeling. *Computer Speech & Language*, 21(2):373–392.
- Robert, C. and Casella, G. (2013). *Monte Carlo statistical methods*. Springer Science & Business Media.
- Rosenfeld, R. (1996). A maximum entropy approach to adaptive statistical language modelling. *Computer Speech & Language*, 10(3):187–228.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Rush, A. M. and Petrov, S. (2012). Vine pruning for efficient multi-pass dependency parsing. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 498–507.
- Santorini, B. (1990). Part-of-speech tagging guidelines for the penn treebank project (3rd revision). Technical Report MS-CIS-90-47, University of Pennsylvania.
- Sato, M.-A. and Ishii, S. (2000). On-line em algorithm for the normalized gaussian network. *Neural computation*, 12(2):407–432.
- Schmid, H. (1994). Probabilistic part-of-speech tagging using decision trees. In *Proceedings of the international conference on new methods in language processing*, volume 12, pages 44–49. Manchester, UK.
- Settles, B. (2010). Active learning literature survey. *University of Wisconsin, Madison*, 52(55-66):11.
- Shen, D. and Lapata, M. (2007). Using semantic roles to improve question answering. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 12–21.
- Shi, L. and Mihalcea, R. (2004). An algorithm for open text semantic parsing. In *Proceedings of the 3rd Workshop on RObust Methods in Analysis of Natural Language Data*, pages 59–67. Association for Computational Linguistics.
- Shieber, S. M. (1985). Evidence against the context-freeness of natural language. *Linguistics and Philosophy*, 8(3):333–343.
- Sipser, M. (2012). *Introduction to the Theory of Computation*. Cengage Learning.
- Smith, D. A. and Eisner, J. (2008). Dependency parsing by belief propagation. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 145–156, Honolulu, HI.
- Smith, D. A. and Smith, N. A. (2007). Probabilistic models of nonprojective dependency trees. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 132–140.
- Smith, N. A. (2011). Linguistic structure prediction. *Synthesis Lectures on Human Language Technologies*, 4(2):1–274.
- Snow, R., O’Connor, B., Jurafsky, D., and Ng, A. Y. (2008). Cheap and fast—but is it good?: evaluating non-expert annotations for natural language tasks. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 254–263, Honolulu, HI.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Socher, R., Bauer, J., Manning, C. D., and Ng, A. Y. (2013a). Parsing with compositional vector grammars. In *Proceedings of the Association for Computational Linguistics (ACL)*, Sophia, Bulgaria.
- Socher, R., Perelygin, A., Wu, J. Y., Chuang, J., Manning, C. D., Ng, A. Y., and Potts, C. (2013b). Recursive deep models for semantic compositionality over a sentiment tree-bank. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- Solorio, T. and Liu, Y. (2008). Learning to predict code-switching points. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 973–981, Honolulu, HI. Association for Computational Linguistics.
- Somasundaran, S. and Wiebe, J. (2009). Recognizing stances in online debates. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 226–234, Suntec, Singapore.
- Song, L., Boots, B., Siddiqi, S. M., Gordon, G. J., and Smola, A. J. (2010). Hilbert space embeddings of hidden markov models. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 991–998.
- Soon, W. M., Ng, H. T., and Lim, D. C. Y. (2001). A machine learning approach to coreference resolution of noun phrases. *Computational linguistics*, 27(4):521–544.
- Spitkovsky, V. I., Alshawi, H., Jurafsky, D., and Manning, C. D. (2010). Viterbi training improves unsupervised dependency parsing. In *CONLL*, pages 9–17.
- Sproat, R., Black, A., Chen, S., Kumar, S., Ostendorf, M., and Richards, C. (2001). Normalization of non-standard words. *Computer Speech & Language*, 15(3):287–333.
- Sra, S., Nowozin, S., and Wright, S. J. (2012). *Optimization for machine learning*. MIT Press.
- Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., and Salakhutdinov, R. (2014). Dropout: A simple way to prevent neural networks from overfitting. *The Journal of Machine Learning Research*, 15(1):1929–1958.
- Stolcke, A., Ries, K., Coccaro, N., Shriberg, E., Bates, R., Jurafsky, D., Taylor, P., Martin, R., Van Ess-Dykema, C., and Meteer, M. (2000). Dialogue act modeling for automatic tagging and recognition of conversational speech. *Computational linguistics*, 26(3):339–373.
- Stone, P. J. (1966). *The General Inquirer: A Computer Approach to Content Analysis*. The MIT Press.
- Stoyanov, V., Gilbert, N., Cardie, C., and Riloff, E. (2009). Conundrums in noun phrase coreference resolution: Making sense of the state-of-the-art. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 656–664, Suntec, Singapore.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Sundermeyer, M., Schlüter, R., and Ney, H. (2012). Lstm neural networks for language modeling. In *INTERSPEECH*.
- Surdeanu, M., Màrquez, L., Carreras, X., and Comas, P. R. (2007). Combination strategies for semantic role labeling. *Journal of Artificial Intelligence Research*, pages 105–151.
- Taboada, M., Brooke, J., Tofiloski, M., Voll, K., and Stede, M. (2011). Lexicon-based methods for sentiment analysis. *Computational linguistics*, 37(2):267–307.
- Tarjan, R. E. (1977). Finding optimum branchings. *Networks*, 7(1):25–35.
- Taskar, B., Guestrin, C., and Koller, D. (2003). Max-margin markov networks. In *Neural Information Processing Systems (NIPS)*.
- Tausczik, Y. R. and Pennebaker, J. W. (2010). The psychological meaning of words: LIWC and computerized text analysis methods. *Journal of Language and Social Psychology*, 29(1):24–54.
- Teh, Y. W. (2006). A hierarchical bayesian language model based on pitman-yor processes. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 985–992.
- Tesnière, L. (1966). *Éléments de syntaxe structurale*. Klincksieck, Paris, second edition.
- Thomas, M., Pang, B., and Lee, L. (2006). Get out the vote: Determining support or opposition from Congressional floor-debate transcripts. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 327–335.
- Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society. Series B (Methodological)*, pages 267–288.
- Toutanova, K., Klein, D., Manning, C. D., and Singer, Y. (2003). Feature-rich part-of-speech tagging with a cyclic dependency network. In *Proceedings of NAACL*.
- Tsochantaridis, I., Hofmann, T., Joachims, T., and Altun, Y. (2004). Support vector machine learning for interdependent and structured output spaces. In *Proceedings of the twenty-first international conference on Machine learning*, page 104. ACM.
- Turian, J., Ratinov, L., and Bengio, Y. (2010). Word representations: A simple and general method for semi-supervised learning. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 384–394, Uppsala, Sweden.
- Turney, P. (2002). Thumbs up or thumbs down? semantic orientation applied to unsupervised classification of reviews. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 417–424.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Van Gael, J., Vlachos, A., and Ghahramani, Z. (2009). The infinite hmm for unsupervised pos tagging. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 678–687, Singapore.
- Vaswani, A., Zhao, Y., Fossum, V., and Chiang, D. (2013). Decoding with large-scale neural language models improves translation. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 1387–1392.
- Viterbi, A. (1967). Error bounds for convolutional codes and an asymptotically optimum decoding algorithm. *IEEE transactions on Information Theory*, 13(2):260–269.
- Wainwright, M. J. and Jordan, M. I. (2008). Graphical models, exponential families, and variational inference. *Foundations and Trends® in Machine Learning*, 1(1-2):1–305.
- Weaver, W. (1955). Translation. *Machine translation of languages*, 14:15–23.
- Wei, G. C. and Tanner, M. A. (1990). A monte carlo implementation of the em algorithm and the poor man’s data augmentation algorithms. *Journal of the American Statistical Association*, 85(411):699–704.
- Weinberger, K., Dasgupta, A., Langford, J., Smola, A., and Attenberg, J. (2009). Feature hashing for large scale multitask learning. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 1113–1120.
- Wiebe, J., Wilson, T., and Cardie, C. (2005). Annotating expressions of opinions and emotions in language. *Language resources and evaluation*, 39(2):165–210.
- Wilson, T., Wiebe, J., and Hoffmann, P. (2005). Recognizing contextual polarity in phrase-level sentiment analysis. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 347–354.
- Winograd, T. (1972). Understanding natural language. *Cognitive psychology*, 3(1):1–191.
- Wu, B. Y. and Chao, K.-M. (2004). *Spanning trees and optimization problems*. CRC Press.
- Xu, W., Liu, X., and Gong, Y. (2003). Document clustering based on non-negative matrix factorization. In *SIGIR*, pages 267–273. ACM.
- Yamada, H. and Matsumoto, Y. (2003). Statistical dependency analysis with support vector machines. In *Proceedings of IWPT*, volume 3, pages 195–206.
- Yang, Y. and Eisenstein, J. (2015). Unsupervised multi-domain adaptation with feature embeddings. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, Denver, CO.

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Yarowsky, D. (1995). Unsupervised word sense disambiguation rivaling supervised methods. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 189–196. Association for Computational Linguistics.
- Zaidan, O. F. and Callison-Burch, C. (2011). Crowdsourcing translation: Professional quality from non-professionals. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 1220–1229, Portland, OR.
- Zettlemoyer, L. S. and Collins, M. (2005). Learning to map sentences to logical form: Structured classification with probabilistic categorial grammars. In *Proceedings of UAI*.
- Zhang, T. (2004). Solving large scale linear prediction problems using stochastic gradient descent algorithms. In *Proceedings of the twenty-first international conference on Machine learning*, page 116. ACM.
- Zhang, Y. and Clark, S. (2008). A tale of two parsers: investigating and combining graph-based and transition-based dependency parsing using beam-search. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 562–571, Honolulu, HI.
- Zhang, Y., Lei, T., Barzilay, R., Jaakkola, T., and Globerson, A. (2014). Steps to excellence: Simple inference with refined scoring of dependency trees. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 197–207, Baltimore, MD.
- Zhang, Y. and Nivre, J. (2011). Transition-based dependency parsing with rich non-local features. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 188–193, Portland, OR.
- Zhu, X. and Ghahramani, Z. (2002). Learning from labeled and unlabeled data with label propagation. Technical report, Technical Report CMU-CALD-02-107, Carnegie Mellon University.

Index

- F*-measure, 63
- K*-means, 66
- k*-nearest-neighbors, 49
- access (an entity), 276
- accuracy, 62
- activation function, 89
- AdaGrad, 45
- adequacy (translation), 295
- adjectives, 134
- adverbs, 134
- alignment, 298
- anaphora resolution, 275
- antecedent (coreference), 283
- antecedent structure, 283
- antonymy, 58, 259
- Arc-eager dependency parsing, 231
- arc-factored, 223
- arc-standard, 229
- Attachment ambiguity, 196
- attributive adjectives, 134
- auxiliary verbs, 134
- backpropagation, 49, 90
- backpropagation through time, 92
- Baum-Welch algorithm, 127
- Bayes' rule, 16
- Beam search, 232
- bias, 81
- bias-variance, 83
- bidirectional recurrent neural network, 126
- Bigrams, 31
- bigrams, 54
- bilexical, 210
- Bilexical features, 223
- binarize, 195
- binding (anaphora), 277
- binomial random variable, 20
- BIO notation, 140
- bitext, 80, 300
- Brown clusters, 261
- c-command, 277
- Catalan number, 193
- cataphora, 276
- centering theory, 276, 279
- chain rule, 16
- character-level language models, 97
- CKY, 193
- cleft, 280
- closed class, 133
- closed-vocabulary, 96
- collapsed dependency grammar, 220
- collocation features, 60
- common nouns, 132
- comparative adverbs, 134
- complement event (probability), 14
- Complement structure, 197
- computation graph, 92
- conditional independence, 18
- conditional probability, 16
- conditional probability distribution, 19

- Conditional Random Field, 121
- conditionally independent, 110
- convexity, 35
- convolutional neural network, 97
- convolutional neural networks, 54
- coordinating conjunctions, 135
- Coordination scope, 197
- copula, 134
- Coreference resolution, 275
- cost-augmented decoding, 39, 120
- count, 62
- count nouns, 132
- deep learning, 49
- degree (of events), 134
- dependency grammar, 218
- dependency graph, 218
- dependent, 218
- determiners, 135
- development set, 34
- digital humanities, 53
- discount, 84
- discrete random variable, 18
- disjoint events, 15
- distributional, 260
- distributional hypothesis, 259, 260
- domains, 313
- dropout, 92
- dual form, 49
- dynamic programming, 105
- early stopping, 34
- effective counts, 84
- emotion, 56
- entity (coreference), 283
- Entity linking, 275
- entity linking, 292
- entropy, 48
- estimation, 20
- event (probability), 14
- evoke (an entity), 276
- existential there, 136
- expectation, 19
- expectation maximization, 67, 85
- explicit semantic analysis, 262
- expressive lengthening, 61
- extraposition, 280
- false positive, 17
- feature function, 31
- feature hashing, 61
- feature templates, 118
- fluent, 79
- forward algorithm, 204
- forward variables, 114
- forward-backward algorithm, 123
- free (anaphora), 277
- garden path sentence, 102
- gazetteers, 282
- gazeteers, 141, 294
- Generalization, 34
- generative model, 22
- generative process, 85
- Gibbs Sampling, 129
- gloss, 80
- gradable adjectives, 134
- gradient descent, 38
- gradient-based optimization, 43
- graph-based dependency parsing, 222
- graphical model, 110
- Hamming cost, 119
- hard EM, 71
- head, 218
- head word, 281
- head words, 217
- held-out data, 94
- hidden Markov models, 110
- hierarchical softmax, 90
- hinge loss, 36
- hyponymy, 58, 259

(c) Jacob Eisenstein 2014-2017. Work in progress.

- Indicator random variables, 18
- inference, 103
- infinitive, 135
- information provenance, 292
- information status, 135
- input word embeddings, 89
- inside algorithm, 204
- inside-outside algorithm, 213
- instance labels, 13
- Integer linear programming, 285
- interpolation, 85
- irrealis, 54

- Jeffreys-Perks law, 83
- joint probabilities, 19

- Kalman Smoother, 128
- Katz backoff, 85
- kernel methods, 49
- knowledge base population, 292

- labeled dependencies, 219
- language model, 80
- Laplace smoothing, 83
- large-margin, 36
- latent semantic analysis, 262
- least squares, 56
- lemma, 57
- lemmatization, 61
- lexical features (linguistics), 137
- lexicalized tree-adjoining grammar, 210
- lexicon-based classification, 56
- lexicon-based sentiment analysis, 54
- Lidstone smoothing, 83
- light verb, 134
- likelihood, 16, 40
- linear regression, 56
- linear separability, 32
- Logistic regression, 39
- Long short-term memories, 90
- long short-term memory, 93
- loss function, 35

- LSTM, 93
- LSTM-CRF, 127

- macro-reading, 292
- manner (of events), 134
- margin, 33, 36
- margin loss, 37
- marginal probability distribution, 19
- marginalize, 15
- markable, 283
- Markov blanket, 110
- Markov Chain Monte Carlo (MCMC), 129
- mass nouns, 132
- matrix-tree theorem, 225
- max-margin Markov network, 119
- max-product algorithm, 113
- maximum a posteriori, 21
- maximum conditional likelihood, 47
- maximum entropy, 47
- maximum likelihood, 20
- maximum likelihood estimate, 25
- mention (coreference), 283
- mention ranking, 284
- mention-pair model, 283
- meronymy, 58, 259
- micro-reading, 292
- mini-batch, 45
- modals, 133
- Modifier scope, 197
- morphemes, 97
- Morphology, 117
- multinomial distribution, 22

- n-grams, 31
- Naïve Bayes, 21
- named entity recognition, 127, 292
- named entity types, 292
- negation, 54
- neural networks, 88
- noise-contrastive estimation, 90

(c) Jacob Eisenstein 2014-2017. Work in progress.

- noisy channel model, 80, 297
- nominals, 275
- nominals (coreference), 282
- non-projective dependency parsing, 225
- norm, 37
- normalize, 61
- nouns, 132

- online learning, 32
- open word classes, 132
- opinion polarity, 53
- overfitting, 34

- parameters, 20
- parent annotation, 206
- part-of-speech, 131
- part-of-speech tagging, 101
- Particle versus preposition, 197
- particles (part-of-speech), 135
- partition, 15
- partition function, 124
- Passive-Aggressive, 46
- Penn Treebank, 203
- perplexity, 95
- personal pronouns, 132
- pleonastic, 280
- possessive pronouns, 133
- posterior, 16
- potentials, 123
- pre-determiners, 136
- precision, 62, 63
- predicative adjective, 133
- predicative adjectives, 134
- prepositions, 135
- presence, 62
- prior, 16
- prior expectation, 21
- probabilistic context-free grammar, 200
- probabilistic models, 20
- probability distribution, 19
- projectivity, 221

- pronouns, 132
- proper nouns, 132

- random outcomes, 14
- ranking, 56
- recall, 62, 63
- rectified linear unit, 89
- recurrent neural network, 88
- regression, 56
- relative frequency estimate, 21
- relative frequency estimator, 25
- ReLU, 89
- reranking, 214
- ridge regression, 56

- sample space, 14
- semantic concordance, 60
- semantics, 205
- semiring algebra, 115
- semisupervised, 60
- sentiment, 53
- shift-reduce, 193
- sigmoid function, 89
- slack variables, 37
- smoothing, 83
- soft K -means, 66
- softmax, 47, 88
- source language, 295
- sparsity, 47
- spectral learning, 71
- Spurious ambiguity, 230
- statistical learning theory, 33
- stemming, 61
- Stochastic gradient descent, 44
- stopwords, 61
- structured perceptron, 118
- structured support vector machine, 119
- subcategorization, 183
- subgradient, 47
- subgradient set, 38
- subjectivity detection, 54

(c) Jacob Eisenstein 2014-2017. Work in progress.

- subordinating conjunctions, 135
- sum-product, 115
- supervised machine learning, 13
- Support Vector Machine, 37
- synonymy, 259
- syntactic, 60
- syntactic dependency, 217
- syntax, 131
- target language, 295
- targeted sentiment analysis, 55
- third axiom of probability, 15
- token, 79
- tokenization, 61
- Tokenizers, 79
- training data, 25
- translation model, 80
- Tree-Adjoining Grammar, 203
- trellis, 106
- tropical semiring, 116
- true positive, 17
- uniqueness, 135
- unsupervised, 60
- unsupervised machine learning, 55
- verbs, 133
- vertical Markovization, 206
- Viterbi algorithm, 105
- Viterbi variable, 105
- weighted context-free grammar, 200
- wh-determiners, 136
- wh-pronouns, 133
- word embeddings, 90, 260, 261
- word representations, 260
- word sense disambiguation, 57
- word senses, 57
- zero-one loss, 35