Linear Models
for
Statistical Natural Language Processing

Jacob Eisenstein

September 25, 2014

Chapter 1

Introduction

This is a collection of notes that I use for teaching Georgia Tech Computer Science 4650 and 7650, "Natural Language." The notes focus on what I view as a core subset of the field of natural language processing, unified by the concept of linear models. This includes approaches to document classification, word sense disambiguation, sequence labeling (part-of-speech tagging and named entity recognition), parsing, coreference resolution, relation extraction, discourse analysis, and, to a limited degree, language modeling and machine translation. The theme was inspired by Fernando Pereira's EMNLP 2008 keynote, "Are linear models right for language."[1] The notes are heavily influenced by several other good resources (e.g., Manning and Schütze, 1999; Jurafsky and Martin, 2009; Figueiredo et al., 2013; Collins, 2013), but for various reasons I wanted to create something of my own.

---

[1]You can see a version of this talk — not the one I saw — online at `vimeo.com/30676245`

Chapter 2

Notation

w_n	word token at position n
$\boldsymbol{x}_i$	a vector of feature counts for instance i, often word counts
N	number of training instances
V	number of words in vocabulary
$\boldsymbol{\theta}$	a vector of weights
y_i	the label for instance i
$\boldsymbol{y}$	vector of labels across all instances
$\mathcal{Y}$	set of all possible labels
K	number of possible labels $K = \#|\mathcal{Y}|$
$\boldsymbol{f}(\boldsymbol{x}_i, y_i)$	feature vector for instance i with label y_i
$P(A)$	probability function of event A
$\mathrm{p}_B(b)$	the marginal probability of random variable B taking value b
M	length of a sequence (of words or tags)
$\mathcal{T}(\boldsymbol{w})$	the set of possible tag sequences for the word sequence $\boldsymbol{w}$
$\Diamond$	the start symbol
$\square$	the stop symbol
λ	the amount of regularization

Chapter 3

Linear classification and features

Suppose you want to build a spam detector. Spam vs. Ham. How would you do it, using only the text in the email?

One solution is to represent document i as a column vector of word counts: $\boldsymbol{x}_i = [0\ 1\ 1\ 0\ 0\ 2\ 0\ 1\ 13\ 0\ldots]^\mathsf{T}$, where $x_{i,j}$ is the count of word j in document i. Suppose the size of the vocabulary is V, so that the length of $\boldsymbol{x}_i$ is also V.

We've thrown out grammar, sentence boundaries, paragraphs — everything but the words! But this could still work. If you see the word *free*, is it spam or ham? How about *c1al1s*? How about *Bayesian*? One approach would be to define a "spamminess" score for every word in the dictionary, and then just add them up. This is also a commonly-used approach to sentiment analysis, where each word is scored as one of $\{1, 0, -1\}$, with 1 indicating positive sentiment and -1 indicating negative sentiment.

These scores are called **weights**, written $\boldsymbol{\theta}$, and we'll spend a lot of time later talking about where they come from. But for now, let's generalize: suppose we want to build a multi-way classifier to distinguish stories about sports, celebrities, music, and business. Each label is an element y_i in a set of K possible labels $\mathcal{Y}$. Then for any pair $\langle \boldsymbol{x}_i, y_i \rangle$, we can define a *feature vector* $\boldsymbol{f}(\boldsymbol{x}_i, y_i)$, such that:

$$\boldsymbol{f}(\boldsymbol{x}, y = 0) = [\boldsymbol{x}_i^\mathsf{T}\ \boldsymbol{0}_{V(K-1)}^\mathsf{T}]^\mathsf{T} \tag{3.1}$$

$$\boldsymbol{f}(\boldsymbol{x}, y = 1) = [\boldsymbol{0}_V^\mathsf{T}\ \boldsymbol{x}_i^\mathsf{T}\ \boldsymbol{0}_{V(K-2)}^\mathsf{T}]^\mathsf{T} \tag{3.2}$$

$$\boldsymbol{f}(\boldsymbol{x}, y = 2) = [\boldsymbol{0}_{2V}^\mathsf{T}\ \boldsymbol{x}_i^\mathsf{T}\ \boldsymbol{0}_{V(K-3)}^\mathsf{T}]^\mathsf{T} \tag{3.3}$$

$$\ldots \tag{3.4}$$

$$\boldsymbol{f}(\boldsymbol{x}, K) = [\boldsymbol{0}_{V(K-1)}^\mathsf{T}\ \boldsymbol{x}_i^\mathsf{T}]^\mathsf{T}, \tag{3.5}$$

where $\boldsymbol{0}_{VK}$ is a column vector of VK zeros. Often we'll add an **offset** feature at

the end of $\boldsymbol{x}$, which is always 1; we then have to also add an extra zero to each of the zero vectors. This gives the entire feature vector $\boldsymbol{f}(\boldsymbol{x}, y)$ a length of $(V + 1)K$.

Now, given a vector of weights, $\boldsymbol{\theta} \in \mathcal{R}^{(V+1)K}$, we can compute the inner product $\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(x, y)$. Then for any document $\boldsymbol{x}_i$, we can predict a label $\hat{y}$ as

$$\hat{y} = \arg \max_{y} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{x}_i, y) \tag{3.6}$$

We could just set the weights by hand. If we wanted to distinguish, say, English from Spanish, we could just use English and Spanish dictionaries, and set each weight to 1. For example,

$$\begin{aligned}
\theta_{\text{english,bicycle}} &= 1 & \theta_{\text{spanish,bicycle}} &= 0 \\
\theta_{\text{english,bicicleta}} &= 0 & \theta_{\text{spanish,bicicleta}} &= 1 \\
\theta_{\text{english,con}} &= 1 & \theta_{\text{spanish,con}} &= 1 \\
\theta_{\text{english,ordinateur}} &= 0 & \theta_{\text{spanish,ordinateur}} &= 0
\end{aligned}$$

Similarly, if we want to distinguish positive and negative sentiment, we could use positive and negative *sentiment lexicons*, which are defined by expert psychologists (Tausczik and Pennebaker, 2010). You'll try this in Project 1.

But it's usually not easy to set the weights by hand. Instead, we will learn them from data. For example, suppose that an email user has manually labeled thousands of messages as "spam" or "not spam"; or a newspaper may label its own articles as "business" or "fashion." Such **instance labels** are a typical form of labeled data that we will encounter in NLP. In **supervised machine learning**, we use instance labels to automatically set the weights for a classifier. An important tool for this is probability.

3.1 Review of basic probability

This section is inspired/borrowed from Manning and Schütze (1999).

- **Formally**: When we write $P()$, this denotes a function $P : \mathcal{F} \to [0, 1]$ from an **event space** $\mathcal{F}$ to a **probability**. A probability is a real number between zero and one, with zero representing impossibility and one representing certainty.

- The probabilities of disjoint event sets are additive: $A_i \cap A_j = \varnothing \Rightarrow P(A_i \cup A_j) = P(A_i) + P(A_j)$. This is a restatement of the Third Axiom of probability.

- For example, you might ask what is the probability of two heads on three coin flips. There are eight possible series of three flips $HHH, HHT, \ldots$, and each is an equally likely event. Of these events, three meet the criterion, HHT, HTH, THH. So the probability is $\frac{3}{8}$.

- More generally, $P(A_i \cup A_j) = P(A_i) + P(A_j) - P(A_i \cap A_j)$. This can be derived from the third axiom.

$$P(A_i \cup A_j) = P(A_i) + P(A_j - (A_i \cap A_j)) \tag{3.7}$$
$$P(A_j) = P(A_j - (A_i \cap A_j)) + P(A_i \cap A_j) \tag{3.8}$$
$$P(A_j - (A_i \cap A_j)) = P(A_j) - P(A_i \cap A_j) \tag{3.9}$$
$$P(A_i \cup A_j) = P(A_i) + P(A_j) - P(A_i \cap A_j) \tag{3.10}$$

- If the probability $P(A \cap B) = P(A)P(B)$, then the events A and B are *independent*, written $A \perp B$.

Conditional probability and Bayes' Rule

A conditional probability is an expression like $P(A \mid B)$, where we are interested in the probability of A conditioned on B happening.

- Conditional probability: $P(A \mid B) = P(A \cap B)/P(B)$

- If $P(A \cap B \mid C) = P(A \mid C)P(B \mid C)$, then the events A and B are **conditionally independent**, written $A \perp B \mid C$.

- Chain rule: $P(A \cap B) = P(A \mid B)P(B)$, which is just a rearrangement of terms.

- We can apply the chain rule multiple times:

$$P(A \cap B \cap C) = P(A \mid B \cap C)P(B \cap C)$$
$$= P(A \mid B \cap C)P(B \mid C)P(C)$$

We'll do this a lot later in the course.

- Bayes' rule follows from the Chain rule: $P(A \mid B) = P(A \cap B)/P(B) = P(B \mid A)P(A)/P(B)$

Often we want the maximum a posteriori (MAP) estimate

$$\hat{B} = \arg\max_B P(B \mid A)$$
$$= \arg\max_B P(A \mid B)P(B)/P(A)$$
$$\propto \arg\max_B P(A \mid B)P(B)$$

- We don't need to normalize the probability because $P(A)$ is the same for all values of B.

- If we do need to compute the conditional $P(A \mid B)$, we can compute $P(A)$ by summing over $P(A \cap B) + P(A \cap \overline{B})$, where $B \cap \overline{B} = \varnothing$ and $B \cup \overline{B} = \Omega$, the entire sample space (such that $P(\Omega) = 1$).

- More generally, if $\bigcup_i B_i = \Omega$ and $\forall_{i,j}, B_i \cap B_j = \varnothing$, then $P(A) = \sum_i P(A \mid B_i)P(B_i)$.

Example   Manning and Schütze (1999) have a nice example of Bayes Rule (Bayes Law) in a linguistic setting.

- Suppose one is interested in a rare syntactic construction, perhaps parasitic gaps, which occurs on average once in 100,000 sentences.

    - (An example of a sentence with a parasitic gap is *Which class did you attend __ without registering for __?* -JE)

- Lana Linguist has developed a complicated pattern matcher that attempts to identify sentences with parasitic gaps. Its pretty good, but it's not perfect:

    - If a sentence has a parasitic gap, it will say so with probability 0.95 (this is the **recall** -JE).

    - If it doesn't, it will wrongly say it does with probability 0.005 (this is the **false positive rate**, the additive inverse of **precision** -JE).

- Suppose the test says that a sentence contains a parasitic gap. What is the probability that this is true?

- (This example is usually framed in terms of tests for rare diseases. -JE)

Solution: Let G be the event of a sentence having a parasitic gap, and T be the event of the test being positive.

$$P(G \mid T) = \frac{P(G \mid T)P(T)}{P(G \mid T)P(T) + P(G \mid \overline{T})P(\overline{T})} \tag{3.11}$$

$$= \frac{0.95 \times 0.00001}{0.95 \times 0.00001 + 0.005 \times 0.99999} \approx 0.002 \tag{3.12}$$

Random variables

A random variable takes on a specific value in $\mathbb{R}^n$, typically with $n = 1$, but not always. Discrete random variables can take values only in some countable subset of $\mathbb{R}$.

- Recall the coin flip example. The number of heads, H, can be viewed as a discrete random variable, $H \in 0, 1, 2, 3$.

- The probability mass associated with each number is $\{\frac{1}{8}, \frac{3}{8}, \frac{3}{8}, \frac{1}{8}\}$.

- This set of numbers represents the **probability distribution** over H, written $P(H = h) = \mathrm{p}(h)$.

- To indicate that the RV H is distributed as $\mathrm{p}(h)$, we write $H \sim \mathrm{p}(h)$.

- The function $\mathrm{p}(h)$ is called a probability **mass** function (pmf) if h is discrete, and a probability **density** function (pdf) if h is continuous.

- If we have more than one variable, we can write a joint probability $\mathrm{p}(a, b) = P(A = a, B = b)$.

- We can write a **marginal** probability $\mathrm{p}_A(a) = \sum_b \mathrm{p}(a, b)$.

- Random variables are independent iff $\mathrm{p}_{A,B}(a, b) = \mathrm{p}_A(a)\mathrm{p}_B(b)$.

- We can write a conditional probability as $\mathrm{p}(a \mid b) = \frac{\mathrm{p}(a,b)}{\mathrm{p}_B(b)}$.

Expectations

Sometimes we want the **expectation** of a function, such as $E[g(x)] = \sum_{x \in \mathcal{X}} g(x)\mathrm{p}(x)$.
Expectations are easiest to think about in terms of probability distributions over discrete events:

- If it is sunny, Marcia will eat three ice creams.

- If it is rainy, she will eat only one ice cream.

- There's a 80% chance it will be sunny.

- The expected number of ice creams she will eat is $0.8 \times 3 + 0.2 \times 1 = 2.6$.

If the random variable X is continuous, the sum becomes an integral:

$$E[g(x)] = \int_{\mathcal{X}} g(x)\mathrm{p}(x)dx \tag{3.13}$$

For example, a fast food restaurant in Quebec gives a 1% discount on french fries for every degree below zero. Assuming they used a thermometer with infinite precision, the expected price would be an integral over all possible temperatures.

3.2   Naïve Bayes

Back to classification! A Naïve Bayes classifier chooses the weights $\boldsymbol{\theta}$ to maximize the *joint* probability of a labeled dataset, $\mathrm{p}(\boldsymbol{x}_{1:N}, \boldsymbol{y}_{1:N})$, where $\langle \boldsymbol{x}_i, y_i \rangle$ is a labeled instance.

We first need to define the probability $\mathrm{p}(\boldsymbol{x}, y)$. We'll do that through a "generative model," which describes a hypothesized stochastic process that has generated the observed data.[1]

- For each document i,

  - draw the label $y_i \sim \mathrm{Categorical}(\mu)$
  - draw the vector of counts $\boldsymbol{x}_i \sim \mathrm{Multinomial}(\phi_{y_i})$,

---

[1]We'll see a lot of different generative models in this course. They are a helpful tool because they clearly and explicitly define the assumptions that underly the form of the probability distribution.

The first thing this generative model tells us is that we can treat each document independently: the probability of the whole dataset is equal to the product of the probabilities of each individual document. The observed word counts and document labels are independent and identically distributed (IID).

$$p(\boldsymbol{x}, \boldsymbol{y}; \mu, \phi) = \prod_i p(\boldsymbol{x}_i, y_i; \mu, \phi) \tag{3.14}$$

This means that the words in each document are **conditionally independent** given the parameters μ and ϕ.

When we write $y_i \sim \text{Categorical}(\mu)$, that means y_i is a stochastic draw from a categorical distribution with **parameter** μ. A categorical distribution is just like a weighted die: $p_{\text{cat}}(y; \mu) = \mu_y$, where μ_y is the probability of the outcome $Y = y$. We require $\sum_y \mu_y = 1$ and $\forall_y, \mu_y \geq 0$.

A multinomial distribution is only slightly more complex:

$$p_{\text{mult}}(\boldsymbol{x}; \phi) = \frac{\left(\sum_j x_j\right)!}{\prod_j x_j!} \prod_j \phi_j^{x_j} \tag{3.15}$$

We again require that $\sum_j \phi_j = 1$ and $\forall j, \phi_j \geq 0$. The first part of the equation doesn't depend on ϕ, and can usually be ignored. Can you see why we need the first part at all?[2]

We can write $p(\boldsymbol{x}_i \mid y_i; \phi)$ to indicate the conditional probability of word counts $\boldsymbol{x}_i$ given label y_i, with parameter ϕ, which is equal to $p_{\text{mult}}(\boldsymbol{x}_i; \phi_{y_i})$.

By specifying the multinomial distribution, we are working with *multinomial naïve Bayes* (MNB). Why "naïve"? Because the multinomial distribution treats each word token independently: the probability mass function factorizes across the counts.[3] We'll see this more clearly later, when we show how MNB is an example of linear classification.

Another version of Naı̈ve Bayes

Consider a slight modification to the generative story of NB:

---

[2]Technically, a multinomial distribution requires a second parameter, the total number of counts (the number of words in the document). Even more technically, that number should be treated as a random variable, and drawn from some other distribution. But none of that matters for classification.

[3]You can plug in any probability distribution to the generative story and it will still be naïve Bayes, as long as you are making the "naïve" assumption that your features are generated independently.

- For each document i

    - Draw the label $y_i \sim \text{Categorical}(\mu)$

    - For each word $n \leq D_i$

        * Draw the word $w_{i,n} \sim \text{Categorical}(\phi_{y_i})$

This is not quite the same model as multinomial Naive Bayes (MNB): it's a product of categorical distributions over words, instead of a multinomial distribution over word counts. This means we would generate the words in order, like $p_W(\textit{multinomial})p_W(\textit{Naive})p_W(\textit{Bayes})$. Formally, this is a model for the joint probability $p(\boldsymbol{w}, y)$, not $p(\boldsymbol{x}, y)$.

However, as a classifier, it is identical to MNB. The final probabilities are reduced by a factor corresponding to the normalization term in the multinomial, $\frac{(\sum_j x_j)!}{\prod_j x_j!}$. This means that the resulting probabilities for a given $\boldsymbol{x}$ are different. However, none of this has anything to do with the label y or the parameters ϕ. The ratio of probabilities between any two labels y_1 and y_2 will be identical, as will the maximum likelihood estimates for the parameters μ and ϕ (defined later).

Prediction

The Naive Bayes prediction rule is to choose the label y which maximizes $p(\boldsymbol{x}, y; \phi, \mu)$:

$$
\begin{aligned}
\hat{y} &= \arg \max_y p(\boldsymbol{x}, y; \mu, \phi) \\
&= \arg \max_y p(\boldsymbol{x} \mid y; \phi)p(y; \mu) \\
&= \arg \max_y \log p(\boldsymbol{x} \mid y; \phi) + \log p(y; \mu)
\end{aligned}
$$

Converting to logarithms makes the notation easier. It doesn't change the prediction rule because the log function is monotonically increasing.

Now we can plug in the probability distributions from the generative story.

$$\log \mathrm{p}(\boldsymbol{x}, y; \mu, \phi) = \arg\max_y \log \mathrm{p}(\boldsymbol{x} \mid y; \phi) + \log \mathrm{p}(y; \mu)$$

$$= \log \left[\frac{\left(\sum_j x_j\right)!}{\prod_j x_j!} \prod_j \phi_{y,j}^{x_j} \right] + \log \mu_y$$

$$= \log \frac{\left(\sum_j x_j\right)!}{\prod_j x_j!} + \sum_j x_j \log \phi_{y,j} + \log \mu_y$$

$$\propto \sum_j x_j \log \phi_{y,j} + \log \mu_y$$

$$= \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}, y),$$

where

$$\boldsymbol{\theta} = [\boldsymbol{\theta}^{(1)\mathsf{T}}, \boldsymbol{\theta}^{(2)\mathsf{T}}, \dots, \boldsymbol{\theta}^{(K)\mathsf{T}}]^\mathsf{T}$$

$$\boldsymbol{\theta}^{(y)} = [\log \phi_{y,1} \ \log \phi_{y,2} \ \dots \ \log \phi_{y,M} \ \log \mu_y]^\mathsf{T}$$

and $\boldsymbol{f}(\boldsymbol{x}, y)$ is a vector of word counts and an offset, padded by zeros for the labels not equal to y (see equations 3.1-3.5). This ensures that the inner product $\boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}, y)$ only activates the features in $\boldsymbol{\theta}^{(y)}$, which are what we need to compute the joint log-probability $\log \mathrm{p}(\boldsymbol{x}, y)$ for each y.

Estimation

The parameters of a multinomial distribution have a simple interpretation: they're the expected frequency for each word. Based on this interpretation, it's tempting to set the parameters empirically, as

$$\phi_{y,j} = \frac{\sum_{i:Y_i=y} x_{i,j}}{\sum_{j'} \sum_{i:Y_i=y} x_{i,j'}} = \frac{\mathrm{count}(y, j)}{\sum_{j'} \mathrm{count}(y, j')} \tag{3.16}$$

In NLP this is called a *relative frequency estimator*. It can be justified more rigorously as a *maximum likelihood estimate*.

As in prediction, we want to maximize the joint likelihood of the data,

$$L = \sum_i \log \mathrm{p}_{\mathrm{mult}}(\boldsymbol{x}_i; \boldsymbol{\phi}_{y_i}) + \log \mathrm{p}_{\mathrm{cat}}(y_i; \mu) \tag{3.17}$$

Since $p(y)$ is unrelated to ϕ, we can forget about it for now. But before we can just optimize L, we have to deal with a constraint:

$$\sum_j \phi_{y,j} = 1 \tag{3.18}$$

We'll do this by adding a Lagrange multiplier. Here's the resulting Lagrangian:

$$\ell[\boldsymbol{\phi}_y] = \sum_{i:Y_i=y} \sum_j x_{ij} \log \phi_{y,j} + \lambda(\sum_j \phi_{y,j} - 1) \tag{3.19}$$

We solve by setting $\frac{\partial \ell}{\phi_j} = 0$.

$$0 = \sum_{i:Y_i=y} x_{i,j}/\phi_{y,j} - \lambda$$

$$\lambda\phi_{y,j} = \sum_{i:Y_i=y} x_{i,j}$$

$$\phi_{y,j} \propto \sum_{i:Y_i=y} x_{i,j} = \sum_i \delta(Y_i = y)x_{i,j}$$

$$= \frac{\sum_{i:Y_i=y} x_{i,j}}{\sum_{j'} \sum_{i:Y_i=y} x_{i,j'}}$$

Similarly, $\mu_y \propto \sum_i \delta(Y_i = y)$, where $\delta(Y_i = y) = 1$ if $Y_i = y$ and 0 otherwise.

Smoothing and MAP estimation

If data is sparse, you can end up with values of $\phi = 0$, allowing a single feature to completely veto a label. This is undesirable, because it imposes high **variance**: depending on what data happens to be in the training set, we could get vastly different classification rules.

One solution is Laplace smoothing: adding "pseudo-counts" of α to each estimate, and then normalize.

$$\phi_{y,j} = \frac{\alpha + \sum_{i:Y_i=y} x_{i,j}}{\sum_{j'} \alpha + \sum_{i:Y_i=y} x_{i,j'}} = \frac{\alpha + \text{count}(i,j)}{V\alpha + \sum_{j'} \text{count}(i,j')} \tag{3.20}$$

Laplace smoothing has a nice Bayesian justification, in which we extend the generative story to include ϕ as a random variable (rather than as a parameter). The resulting estimate is called *maximum a posteriori*, or MAP.

Smoothing reduces **variance**, but it takes us away from the maximum-likelihood estimate: it imposes a **bias** (towards uniform probabilities). Machine learning theory shows that errors on held out data result from the sum of bias and variance. Techniques for reducing variance typically increase the bias, so there is a **bias-variance tradeoff**.

- Unbiased classifiers **overfit** the training data, yielding poor performance on unseen data.

- But if we set a very large smoothing value, we can **underfit** instead. In the limit of $\alpha \to \infty$, we have zero variance: it is the same classifier no matter what data we see! But the bias of such a classifier will be high.

- Navigating this tradeoff is hard. But in general, as you have more data, variance is less of a problem, so you just go for low bias.

Training, testing, and tuning (development) sets

We'll soon talk about more learning algorithms, but whichever one we apply, we will want to report its accuracy. Really, this is an educated guess about how well the algorithm will do on new data in the future.

To do this, we need to hold out a separate "test set" from the data that we use for estimation (i.e., training, learning). Otherwise, if we measure accuracy on the same data that is used for estimation, we will badly overestimate the accuracy we're likely to get on new data. See `http://xkcd.com/1122/` for a cartoon related to this idea.

Many learning algorithms also have "tuning" parameters:

- the smoothing pseudo-counts α in Naive Bayes

- the regularization λ in logistic regression

- the slack weight C in the support-vector machine

All of these tuning parameters really do the same thing: they navigate the bias-variance tradeoff. Where is the best position on this tradeoff curve? It's hard to tell in advance. Sometimes it is tempting to see which tuning parameter gives the best performance on the test set, and then report that performance. Resist this temptation! It will also lead to overestimating accuracy on truly unseen future

data. For that reason, this is a sure way to get your research paper rejected. Instead, you should split off a piece of your training data, called a "development set" (or "tuning set").

Sometimes, people average across multiple test sets and/or multiple development sets. One way to do this is to divide your data into "folds," and allow each fold to be the development set one time. This is called **K-fold cross-validation**. In the extreme, each fold is a single data point. This is called **leave-one-out**.

The Naivety of Naive Bayes

Naive Bayes is very simple to work with. Estimation and prediction can done in closed form, and the nice probabilistic interpretation makes it relatively easy to extend the model in various ways.

But Naive Bayes makes assumptions which seriously limit its accuracy, especially in NLP.

- The multinomial distribution assumes that each word is generated independently of all the others (conditioned on the parameter ϕ_y). Formally, we assume conditional independence:

$$\mathrm{p}(\textit{naïve}, \textit{Bayes}; \phi) = \mathrm{p}(\textit{naïve}; \phi)\mathrm{p}(\textit{Bayes}; \phi). \tag{3.21}$$

- But this is clearly wrong, because words "travel together." Question for you, is it:

$$\mathrm{p}(\textit{naïve Bayes}) > \mathrm{p}(\textit{naïve})\mathrm{p}(\textit{Bayes}) \tag{3.22}$$

or...

$$\mathrm{p}(\textit{naïve Bayes}) < \mathrm{p}(\textit{naïve})\mathrm{p}(\textit{Bayes}) \tag{3.23}$$

Apply the chain rule!

Traffic lights   Dan Klein makes this point with an example about traffic lights. In his hometown of Pittsburgh, there is a 1/7 chance that the lights will be broken, and both lights will be red. There is a 3/7 chance that the lights will work, and the north-south lights will be green; there is a 3/7 chance that the lights work and the east-west lights are green.

The *prior* probability that the lights are broken is 1/7. If they are broken, the conditional likelihood of each light being red is 1. The prior for them not being broken is 6/7. If they are not broken, the conditional likelihood of each being light being red is 1/2.

Now, suppose you see that both lights are red. According to Naive Bayes, the probability that the lights are broken is $1/7 \times 1 \times 1 = 1/7 = 4/28$. The probability that the lights are not broken is $6/7 \times 1/2 \times 1/2 = 6/28$. So according to naive Bayes, there is a 60% chance that the lights are not broken!

What went wrong? We have made an independence assumption to factor the probability $P(R, R \mid \text{not-broken}) = P_{\text{north-south}}(R \mid \text{not-broken})P_{\text{east-west}}(R \mid \text{not-broken})$. But this independence assumption is clearly incorrect, because $P(R, R \mid \text{not-broken}) = 0$.

Less Naive Bayes? Of course we could decide not to make the naive Bayes assumption, and model $P(R, R)$ explicitly. But this idea does not scale when the feature space is large (as it often is in NLP). The number of possible feature configurations grows exponentially, so our ability to estimate accurate parameters will suffer from high variance. With an infinite amount of data, we'd be fine (in theory, maybe not in practice); but we never have that. Naive Bayes accepts some bias (because of the incorrect modeling assumption) in exchange for lower variance.

3.3 Recap

- Bag-of-words representation $\boldsymbol{f}(\boldsymbol{x}, y)$

- Classification as a dot-product $\boldsymbol{\theta}^{\mathsf{T}}\boldsymbol{f}(\boldsymbol{x}, y)$

- Naive Bayes

    - Define $\mathrm{p}(\boldsymbol{x}, \boldsymbol{y})$ via a *generative model*
    - Prediction: $\hat{y} = \arg\max_y \mathrm{p}(\boldsymbol{x}_i, y)$
    - Learning:

$$\boldsymbol{\theta} = \arg\max_{\boldsymbol{\theta}} \mathrm{p}(\boldsymbol{x}, \boldsymbol{y}; \boldsymbol{\theta})$$

$$\mathrm{p}(\boldsymbol{x}, \boldsymbol{y}; \boldsymbol{\theta}) = \prod_i \mathrm{p}(\boldsymbol{x}_i, y_i; \boldsymbol{\theta}) = \prod_i \mathrm{p}(\boldsymbol{x}_i | y_i)\mathrm{p}(y_i)$$

$$\phi_{y,j} = \frac{\sum_{i:Y_i=y} x_{ij}}{\sum_{i:Y_i=y} \sum_j x_{ij}}$$

$$\mu_y = \frac{\text{count}(Y = y)}{N}$$

This gives the maximum-likelihood estimator (MLE; same as relative frequency estimator)

- Bias-variance tradeoff: MLE is high-variance, so add smoothing pseudo counts α. This reduces variance but adds bias.

Chapter 4

Sentiment analysis

Todo: add notes about sentiment analysis here

Chapter 5

Discriminative learning

5.1 Features

Naive Bayes is a simple classifier, where the weights are learned based on the joint probability of labels and words. It includes an independence assumption: all features are mutually independent, conditioned on the label.

- We have defined a **feature function** $f(x, y)$, which corresponds to "bag-of-words" features. While these features do violate the independence assumption, the violation is relatively mild.

- We may be interested in other features, which violate independence more severely. Can you think of any?

  - Prefixes, e.g. *anti-*, *im-*, *un-*
  - Punctuation and capitalization
  - Bigrams, e.g. *not good*, *not bad*, *least terrible*, ...

Rich feature sets generally cannot be combined with Naive Bayes because the distortions resulting from violations of the independence assumption overwhelm the additional power of better features.

$$\text{p}(\textit{not bad food}|y) \approx \text{p}(\textit{not}|y)\text{p}(\textit{bad}|y)\text{p}(\textit{food}|y) \tag{5.1}$$
$$\text{p}(\textit{not bad food}|y) \not\approx \text{p}(\textit{not}|y)\text{p}(\textit{bad}|y)\text{p}(\textit{not bad}|y)\text{p}(\textit{food}|y) \tag{5.2}$$

To use these features, we will need learning algorithms that do not rely on an independence assumption.

5.2   Perceptron

In NB, the weights can be interpreted as parameters of a probabilistic model. But this model requires an independence assumption that usually does not hold, and limits our choice of features.

Why not forget about probability and learn the weights in an error-driven way?

- Until converged, at each iteration t

    - Select an instance i
    - Let $\hat{y} = \arg\max_y \boldsymbol{\theta}_t^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}_i, y)$
    - If $\hat{y} = y_i$, do nothing
    - If $\hat{y} \neq y_i$, set $\boldsymbol{\theta}_{t+1} \leftarrow \boldsymbol{\theta}_t + \boldsymbol{f}(\boldsymbol{x}_i, y_i) - \boldsymbol{f}(\boldsymbol{x}_i, \hat{y})$

Basically we are saying: if you make a mistake, increase the weights for features which are active with the correct label y_i, and decrease the weights for features which are active with the guessed label $\hat{y}$.

This seems like a cheap heuristic, right? Will it really work? In fact, there is some nice theory for the perceptron.

- If there is a set of weights that correctly separates your data, then your data is **separable**.

- Formally, your data is (linearly) separable if there exists a set of weights $\boldsymbol{\theta}$ such that

$$\forall \boldsymbol{x}_i, y_i, \ \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}_i, y_i) > \max_{y' \neq y_i} \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}_i, y') \tag{5.3}$$

- If your data is linearly separable, it can be proven that the perceptron algorithm will eventually find a separator.

- What if your data is not separable?

    - the number of errors is bounded...
    - but the algorithm will thrash. That is, the weights will cycle between different values, and will never converge.

The perceptron is an **online** learning algorithm.

- This means that it adjusts the weights after every example.

- This is different from Naïve Bayes, which computes corpus statistics and then sets the weights in a single operation. This is a **batch learning** algorithm.

- Other algorithms are **iterative**, in that they perform multiple updates to the weights, but are also **batch**, in that they have to use all the training data to compute the update. We'll mention two of those algorithms later.

Voted (averaged) perceptron

One solution to thrashing is to average the weights across all iterations:

$$\overline{\boldsymbol{\theta}} = \frac{1}{T} \sum_{t=1}^{T} \boldsymbol{\theta}_t$$

$$y = \arg\max_{y} \overline{\boldsymbol{\theta}}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{x}, y)$$

There is some analysis showing that voting can improve generalization (Freund and Schapire, 1999; Collins, 2002). However, this rule as described here is not practical. Can you see why not, and how to fix it?

5.3   Loss functions and large-margin classification

Naive Bayes chooses the weights $\boldsymbol{\theta}$ by maximizing the likelihood $p(\boldsymbol{x}, \boldsymbol{y})$. This can be seen, equivalently, as maximizing the log-likelihood (due to the monotonicity of the log function), and as **minimizing** the negative log-likelihood. This negative log-likelihood can therefore be viewed as a **loss function**, which is minimized:

$$\log p(\boldsymbol{x}, \boldsymbol{y}; \boldsymbol{\theta}) = \sum_{i=1}^{N} \log p(\boldsymbol{x}_i, y_i; \boldsymbol{\theta}) \tag{5.4}$$

$$\ell_{\mathrm{NB}}(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i) = -\log p(\boldsymbol{x}_i, y_i; \boldsymbol{\theta}) \tag{5.5}$$

$$\hat{\boldsymbol{\theta}} = \arg\min_{\boldsymbol{\theta}} \sum_{i=1}^{N} \ell_{\mathrm{NB}}(\boldsymbol{\theta}, \boldsymbol{x}_i, y_i) \tag{5.6}$$

This may seem confusing and backwards, but loss functions provide a very general framework in which to compare many approaches to machine learning.

For example, even though the perceptron is not a probabilistic model, it is also trying to minimize a **loss function**:

$$\ell_{\text{perceptron}}(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i) = \begin{cases} 0, & y_i = \arg\max_y \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(x_i, y) \\ 1, & \text{otherwise} \end{cases} \tag{5.7}$$

This loss function has some pros and cons in comparison with Naive Bayes.

- ℓ_{NB} can suffer **infinite** loss on a single example, which suggests it will overemphasize some examples, and underemphasize others.

- $\ell_{\text{perceptron}}$ treats all errors equally. It only cares if the example is correct, and not about how confident the classifier was. Since we usually evaluate on accuracy, this is a better match.

- $\ell_{\text{perceptron}}$ is non-convex[1] and discontinuous. Finding the global optimum is intractable when the data is not separable.

We can fix this last problem by defining a loss function that behaves more nicely. To do this, let's define the *margin* as

$$\gamma(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i) = \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}_i, y_i) - \max_{y \neq y_i} \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}_i, y) \tag{5.8}$$

Then we can write a convex and continuous "hinge loss" as

$$\ell_{\text{hinge}}(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i) = \begin{cases} 0, & \gamma(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i) \geq 1, \\ 1 - \gamma(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i), & \text{otherwise} \end{cases} \tag{5.9}$$

Equivalently, we can write $\ell_{\text{hinge}}(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i) = (1 - \gamma(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i))_+$, where $(x)_+$ indicates the positive part of x.

Essentially, we want a *margin* of at least 1 between the score for the true label and the best-scoring alternative, which we have written $\hat{y}$.

The hinge and perceptron loss functions are shown in Figure 5.1.

---

[1]As a reminder, a function f is convex iff $\alpha f(x_i) + (1 - \alpha)f(x_j) \geq f(\alpha x_i + (1 - \alpha)x_j)$, for all $\alpha \in [0, 1]$ and for all x_i and x_j on the domain of the function. Convexity implies that any local minimum is also a global minimum, and there are a wide array of techniques for optimizing convex functions (Boyd and Vandenberghe, 2004)

Figure 5.1: Hinge and perceptron loss functions

Large-margin online classification

Note that we can write $\boldsymbol{\theta} = s\boldsymbol{u}$, where $||\boldsymbol{u}||_2 = 1$. Think of s as the magnitude and $\boldsymbol{u}$ as the direction of the vector $\boldsymbol{\theta}$. If the data is separable, there are many values of s which attain zero hinge loss. For generality, we will try to make the smallest magnitude change to $\boldsymbol{\theta}$ possible.[2]

At step t, we optimize:

$$\boldsymbol{\theta}_{t+1} = \arg\min_{\boldsymbol{\theta}} \frac{1}{2}||\boldsymbol{\theta} - \boldsymbol{\theta}_t||^2 \ s.t. \ \ell_{\text{hinge}}(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i) = 0 \tag{5.10}$$

Assuming that the constraint can be satisfied (i.e., the problem is linearly separable), the optimal solution is found at,

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \tau_t(\boldsymbol{f}(y_i, \boldsymbol{x}_i) - \boldsymbol{f}(\hat{y}, \boldsymbol{x}_i)) \tag{5.11}$$

$$\tau_t = \frac{\ell(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i)}{||\boldsymbol{f}(x_i, y_i) - \boldsymbol{f}(x_i, \hat{y})||^2}, \tag{5.12}$$

---

[2]In the support vector machine (without slack variables), we choose the smallest magnitude weights that satisfy the constraint of zero hinge loss. Pegasos is an online algorithm for training SVMs (Shwartz et al., 2007); it is similar to Passive-Aggressive.

where again $\hat{y}$ is the best scoring y according to $\boldsymbol{\theta}_t$. This solution can be obtained by introducing τ_t as a Lagrange multiplier for the constraint in (5.10).

If the data is not lienarly separable, there will be instances for which we can't meet this constraint. To deal with this, we introduce a "slack" variable ξ_i. We use the slack variable to trade off between the constraint (having a large margin) and the objective (having a small change in $\boldsymbol{\theta}$). The tradeoff is controlled by a parameter C.

$$\min w \frac{1}{2}||\boldsymbol{\theta} - \boldsymbol{\theta}_t||^2 + C\xi_t \tag{5.13}$$
$$s.t.\ell_{\text{hinge}}(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i) \leq \xi_t, \xi_t \geq 0$$

The solution to 5.13 is,

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \tau_t(\boldsymbol{f}(y_i, \boldsymbol{x}_i) - \boldsymbol{f}(\hat{y}, \boldsymbol{x}_i)) \tag{5.14}$$
$$\tau_t = \min\left(C, \frac{\ell(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i)}{||\boldsymbol{f}(\boldsymbol{x}_i, y_i) - \boldsymbol{f}(\boldsymbol{x}_i, \hat{y})||^2}\right), \tag{5.15}$$

- If C is 0, then infinite slack is permitted, and the weights will never change.

- As $C \to \infty$, no slack is permitted, and the optimization is identical to equation 5.10 and 5.12.

This algorithm is called "Passive-Aggressive" (PA; Crammer et al., 2006), because it is passive when the margin constraint is satisfied, but it aggressively changes the weights to satisfy the constraints if necessary.[3]

- PA is error-driven like the perceptron, but is more stable to violations of separability, like the averaged perceptron.

- PA allows more explicit control than the Averaged Perceptron, due to the C parameter. When C is small, we make very conservative adjustments to $\boldsymbol{\theta}$ from each instance, because the slack variables aren't very expensive. When C is large, we make large adjustments to avoid using the slack variables.

- You can also apply weight averaging to PA.

---

[3]A related algorithm without slack variables is called MIRA, for Margin-Infused Relaxed Algorithm (Crammer and Singer, 2003).

- **Support vector machines** (SVMs) are another learning algorithm based on the hinge loss (Burges, 1998), but they try to minimize the norm of the weights, rather than the norm of the change in the weights. They are typically trained in **batch** style, meaning that they have to read all the training instances in to compute each update. However, SVMs can also be trained in an online fashion (Shwartz et al., 2007). The LXMLS lab guide provides a simpler online learning algorithm, based on stochastic subgradient descent (Figueiredo et al., 2013).

Pros and cons of Perceptron and PA

- Perceptron and PA are error-driven, which means they usually do better in practice than naive Bayes.

- They are also online, which means we can learn without having our whole dataset in memory at once. NB can also be estimated online, in the sense that you can stream the data and store the counts.

- The original perceptron doesn't behave well if the data is not separable, and doesn't make it easy to control model complexity.

- All these models lack a probabilistic interpretation. Probabilities are useful because they quantify the classification certainty, allowing us to compute expected utility, and to incorporate the classifier in more complex probabilistic models.

5.4 Logistic regression

Logistic regression is error-driven like the perceptron, but probabilistic like Naive Bayes. This is useful in case we want to quantify the uncertainty about a classification decision.

Recall that NB selects weights to optimize the joint probability $p(y, \boldsymbol{x})$.

- In NB, we factor this as $p(y, \boldsymbol{x}) = p(\boldsymbol{x}|y)p(y)$.

- But we could equivalently write $p(y, \boldsymbol{x}) = p(y|\boldsymbol{x})p(\boldsymbol{x})$.

Since we always know $\boldsymbol{x}$, we really care only about $p(y|\boldsymbol{x})$. Logistic regression optimizes this directly. To do this, we have to define the probability function

differently. We define the conditional probability directly, as,

$$p(y|\boldsymbol{x}) = \frac{\exp\left(\boldsymbol{\theta}^\mathsf{T}\boldsymbol{f}(\boldsymbol{x},y)\right)}{\sum_{y'\in\mathcal{Y}}\exp\left(\boldsymbol{\theta}^\mathsf{T}\boldsymbol{f}(\boldsymbol{x},y')\right)} \tag{5.16}$$

$$\log p(y|\boldsymbol{x}) = \sum_{i=1}^{N}\boldsymbol{\theta}^\mathsf{T}\boldsymbol{f}(\boldsymbol{x}_i,y_i) - \log\sum_{y'\in\mathcal{Y}}\exp\boldsymbol{\theta}^\mathsf{T}\boldsymbol{f}(\boldsymbol{x}_i,y') \tag{5.17}$$

$$\hat{\boldsymbol{\theta}} = \arg\max_{\boldsymbol{\theta}}\sum_{i=1}^{N}\log p(y_i|\boldsymbol{x}_i;\boldsymbol{\theta}) \tag{5.18}$$

Inside the sum, we have the (additive inverse of) the **logistic loss**.

- In binary classification, we can write this as

$$\ell_{\text{logistic}}(\boldsymbol{\theta};\boldsymbol{x}_i,y_i) = -(y_i\boldsymbol{\theta}^\mathsf{T}\boldsymbol{x}_i - \log\left(1 + \exp\boldsymbol{\theta}^\mathsf{T}\boldsymbol{x}_i\right)) \tag{5.19}$$

- In multi-class classification, we have,[4]

$$\ell_{\text{logistic}}(\boldsymbol{\theta};\boldsymbol{x}_i,y_i) = -(\boldsymbol{\theta}^\mathsf{T}\boldsymbol{f}(\boldsymbol{x}_i,y_i) - \log\sum_{y'\in\mathcal{Y}}\exp\boldsymbol{\theta}^\mathsf{T}\boldsymbol{f}(\boldsymbol{x}_i,y')) \tag{5.20}$$

The logistic loss is shown in Figure 5.2. Because it is smooth and convex, we can optimize it through gradient steps:

---

[4]The log-sum-exp term is very common in machine learning. It is numerically instable because you can underflow if the inner product is small, and overflow if the inner product is large. Libraries like `scipy` contain special functions for computing `logsumexp`, but with some thought, you should be able to see how to create an implementation that is numerically stable.

Figure 5.2: Hinge, perceptron, and logistic loss functions

$$\ell = \sum_{i=1}^{N} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{x}_i, y_i) - \log \sum_{y' \in \mathcal{Y}} \exp \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{x}_i, y') \tag{5.21}$$

$$\frac{\partial \ell}{\partial \boldsymbol{\theta}} = \sum_{i=1}^{N} \boldsymbol{f}(\boldsymbol{x}_i, y_i) - \frac{\sum_{y' \in \mathcal{Y}} \exp \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{x}_i, y') \boldsymbol{f}(\boldsymbol{x}_i, y')}{\sum_{y'' \in \mathcal{Y}} \exp \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{x}_i, y'')} \tag{5.22}$$

$$= \sum_{i=1}^{N} \boldsymbol{f}(\boldsymbol{x}_i, y_i) - \sum_{y' \in \mathcal{Y}} \frac{\exp \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{x}_i, y')}{\sum_{y'' \in \mathcal{Y}} \exp \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{x}_i, y'')} \boldsymbol{f}(\boldsymbol{x}_i, y') \tag{5.23}$$

$$= \sum_{i=1}^{N} \boldsymbol{f}(\boldsymbol{x}_i, y_i) - \sum_{y' \in \mathcal{Y}} \mathrm{p}(y' | \boldsymbol{x}_i; \boldsymbol{\theta}) \boldsymbol{f}(\boldsymbol{x}_i, y') \tag{5.24}$$

$$= \sum_{i=1}^{N} \boldsymbol{f}(\boldsymbol{x}_i, y_i) - E[\boldsymbol{f}(\boldsymbol{x}_i, y)] \tag{5.25}$$

This gradient has a pleasing interpretation as the difference between the observed counts and the expected counts.[5] Compare this gradient with the percep-

---

[5] Recall that the definition of an expected value $E[f(x)] = \sum_x f(x)\mathrm{p}(x)$

tron and PA update rules.

The bias-variance tradeoff is handled by penalizing large $\boldsymbol{\theta}$ in the objective, adding a term of $\frac{\lambda}{2}||\boldsymbol{\theta}||_2^2$. This is called L2 regularization, because of the L2 norm. It can be viewed as placing a 0-mean Gaussian prior on $\boldsymbol{\theta}$.

This penalty contributes a term of $\lambda\boldsymbol{\theta}$ to the gradient, so we have,

$$\ell = \sum_{i=1}^{N} \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}_i, y_i) - \log \sum_{y' \in \mathcal{Y}} \exp \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}_i, y') + \frac{\lambda}{2}||\boldsymbol{\theta}||_2^2 \qquad (5.26)$$

$$\frac{\partial \ell}{\partial \boldsymbol{\theta}} = \sum_{i=1}^{N} \boldsymbol{f}(\boldsymbol{x}_i, y_i) - E[\boldsymbol{f}(\boldsymbol{x}_i, y)] - \lambda\boldsymbol{\theta}. \qquad (5.27)$$

Optimization

Batch optimization   In batch optimization, you keep all the data in memory and iterate over it many times.

- The logistic loss is smooth and convex, so we can find the global optimum using gradient descent. But in practice, this can be very slow.

- Second-order (Newton) optimization would incorporate the inverse Hessian. The Hessian is

$$H_{i,j} = \frac{\partial^2}{\partial w_i \partial w_j} \ell, \qquad (5.28)$$

  but this matrix is usually too big to deal with.

- In practice, people usually apply **quasi-Newton optimization**, which approximates the Hessian matrix. The specific method that is particularly popular is L-BFGS[6] NLP people usually treat L-BFGS as a black box; you will typically pass it a pointer to a function that computes the likelihood and gradient. L-BFGS is provided in `scipy.optimize`.

Online optimization   In online optimization, you consider one example (or a "mini-batch" of a few examples) at a time. *Stochastic gradient descent* makes a

---

[6]A friend of mine told me you can remember the order of the letters as "Large Big Friendly Giants." Does this help you?

stochastic online approximation to the overall gradient:

$$\boldsymbol{\theta}^{(t+1)} \leftarrow \boldsymbol{\theta}^{(t)} - \eta_t \nabla_{\boldsymbol{\theta}} \ell(\boldsymbol{\theta}^{(t)}, \boldsymbol{x}, \boldsymbol{y}) \tag{5.29}$$

$$= \boldsymbol{\theta}^{(t)} - \eta_t (\lambda \boldsymbol{\theta}^{(t)} - \sum_i^N \boldsymbol{f}(\boldsymbol{x}_i, y_i) - E[\boldsymbol{f}(\boldsymbol{x}_i, y)]) \tag{5.30}$$

$$= (1 - \lambda \eta_t) \boldsymbol{\theta}^{(t)} + \eta_t \sum_i^N \boldsymbol{f}(\boldsymbol{x}_i, y_i) - E[\boldsymbol{f}(\boldsymbol{x}_i, y)] \tag{5.31}$$

$$\approx (1 - \lambda \eta_t) \boldsymbol{\theta}^{(t)} + N \eta_t \left(\boldsymbol{f}(\boldsymbol{x}_{i(t)}, y_{i(t)}) - E[\boldsymbol{f}(\boldsymbol{x}_{i(t)}, y)] \right) \tag{5.32}$$

where η_t is the **stepsize** at iteration t, and $\langle \boldsymbol{x}_{i(t)}, y_{i(t)} \rangle$ is the instance selected at iteration t. (So here we are setting the mini-batch size equal to one.) As always, N is the total number of instances. As above, the expectation is equal to a weighted sum over the labels,

$$E[\boldsymbol{f}(\boldsymbol{x}_{i(t)}, y)] = \sum_{y' \in \mathcal{Y}} \mathrm{p}(y' | \boldsymbol{x}_{i(t)}; \boldsymbol{\theta}) \boldsymbol{f}(\boldsymbol{x}_{i(t)}, y'). \tag{5.33}$$

- Note how similar this update is to the perceptron!

- If we set $\eta_t = \eta_0 t^{-\alpha}$ for $\alpha \in [1, 2]$, we have guaranteed convergence.

- We can also just fix η_t to a small value, like 10^-3. (This is what we will do in the problem set.)

- In either case, we could tune this parameter on a development set. However, it would be acceptable to just find a value that gives a good regularized log-likelihood on the training set, since this parameter relates to the quality of the optimization, and not the generalization capability of the classifier.

- In theory, we select $\langle \boldsymbol{x}_{i(t)}, y_{i(t)} \rangle$ at random, but in practice we usually just iterate through the dataset.

- We can fold N into η and λ, so that $\eta^* = N\eta$ and $\lambda^* = \lambda \frac{\eta^*}{N}$. This gives the more compact form,

$$(1 - \lambda^* \eta_t^*) \boldsymbol{\theta}^{(t)} + \eta_t^* \left(\boldsymbol{f}(\boldsymbol{x}_{i(t)}, y_{i(t)}) - E[\boldsymbol{f}(\boldsymbol{x}_{i(t)}, y)] \right) \tag{5.34}$$

For more on stochastic gradient descent, as applied to a number of different learning algorithms, see (Zhang, 2004) and (Bottou, 1998). Murphy (2012) traces SGD to a 1978 paper by GT's own Arkadi Nemirovski (Nemirovski and Yudin, 1978). You can find several recent chapters about online optimization in the edited volume by Sra et al. (2012).

Adagrad Recent work has shown that you can often learn more quickly by using an **adaptive** step-size, which is different for every feature (Duchi et al., 2011). Specifically, in the **Adagrad** algorithm (adaptive gradient), you keep track of the sum of the squares of the gradients for each feature, and rescale the learning rate by its inverse:

$$g_t = - f(x_i, y_i) + \sum_{y' \in \mathcal{Y}} p(y' \mid x_i) f(x_i, y_i) + \lambda \theta \tag{5.35}$$

$$\theta_j^{(t+1)} \leftarrow \theta_j^{(t)} - \frac{\eta}{\sqrt{\sum_{t'=1}^{t} g_{t,j}^2}} g_{t,j}, \tag{5.36}$$

where j iterates over features in $f(x, y)$. The effect of this is that features with consistently large gradients are updated more slowly. Another way to view this update is that rare features are taken more seriously, since their sum of squared gradients will be smaller. Adagrad seems to require less careful tuning of η, and Dyer (2014) reports that $\eta = 1$ works for a wide range of problems.

Note that the Adagrad update can apply to any smooth loss function, including the hinge loss defined in Equation 5.9.

Names

Logistic regression is so named because in the binary case where $y \in \{0, 1\}$, we are performing a regression of x against y, after passing the inner product $\theta^\mathsf{T} x$ through a logistic transformation. You could always do a linear regression, but this would ignore the fact that the y is limited to a few values.

- Logistic regression is also called **maximum conditional likelihood** (MCL), because it maximizes... the conditional likelihood $p(y \mid x)$.

- Logistic regression can be viewed as part of a larger family, called **generalized linear models**. If you use R, you are probably familiar with `glmnet`.

- Logistic regression is also called **maximum entropy**, especially in the earlier NLP literature (Berger et al., 1996). This is due to an alternative formulation, which tries to find the maximum entropy probability function that satisfies moment-matching constraints.

The moment matching constraints specify that the empirical counts of each label-feature pair should match the expected counts:

$$\forall j, \sum_{i=1}^{N} f_j(\boldsymbol{x}_i, y_i) = \sum_{i=1}^{N} \sum_{y \in \mathcal{Y}} p(y \mid \boldsymbol{x}_i; \boldsymbol{\theta}) f_j(\boldsymbol{x}_i, y) \tag{5.37}$$

Note that this constraint will be met exactly when the derivative of the likelihood function (equation 5.25) is equal to zero. However, this will be true for many values of $\boldsymbol{\theta}$. Which should we choose?

The entropy of a conditional likelihood function $P(Y|X)$ is

$$H(P) = -\sum_{x \in \mathcal{X}} \tilde{p}(x) \sum_{y \in \mathcal{Y}} p(y|x) \log p(y|x), \tag{5.38}$$

where $\tilde{p}(x)$ is the *empirical probability* of x. We compute an empirical probability by summing over all the instances in training set.

If the entropy is large, this function is smooth across possible values of y; if it is small, the function is sharp. The entropy is zero if $p(y|x) = 1$ for some particular $Y = y$ and zero for everything else. By saying we want maximum-entropy classifier, we are saying we want to make the least commitments possible, while satisfying the moment-matching constraints:

$$\max_{\boldsymbol{\theta}} \quad -\sum_{\boldsymbol{x}} \tilde{p}(\boldsymbol{x}) \sum_{y} p(y|\boldsymbol{x}; \boldsymbol{\theta}) \log p(y|\boldsymbol{x}; \boldsymbol{\theta}) \tag{5.39}$$

$$s.t. \quad \forall j, \sum_{i=1}^{N} f_j(\boldsymbol{x}_i, y_i) = \sum_{i=1}^{N} \sum_{y} p(y|\boldsymbol{x}_i; \boldsymbol{\theta}) f_j(\boldsymbol{x}_i, y) \tag{5.40}$$

Now, the solution to this constrained optimization problem is identical to the maximum conditional likelihood (logistic-loss) formulation we've considered in the previous section.

This view of logistic regression is arguably a little dated, but it's useful to understand what's going on. The information-theoretic concept of entropy will pop up again a few times in the course. For a tutorial on maximum entropy, see `http://www.cs.cmu.edu/afs/cs/user/aberger/www/html/tutorial/tutorial.html`.

5.5   Summary of learning algorithms

- **Naive Bayes**. pros: easy and probabilistic. cons: arguably optimizes wrong objective; usually has poor accuracy, especially with overlapping features.

- **Perceptron and PA**. pros: easy, online, and error-driven. cons: not probabilistic. this can be bad in pipeline architectures, where the output of one system becomes the input for another.

- **Logistic regression**. pros: error-driven and probabilistic. cons: batch learning requires black-box software; hinge loss sometimes yields better accuracy than logistic loss.

What about non-linear classification?

The feature spaces that we consider in NLP are usually huge, so non-linear classification can be quite difficult. When the feature dimension V is larger than the number of instances N — often the case in NLP — you can always learn a linear classifier that will perfectly classify your training instances.[7] This makes selecting an appropriate **non-linear** classifier especially difficult. Nonetheless, there are some approaches to non-linear learning in NLP:

- You can add **features**, such as bigrams, which are non-linear combinations of other features. For example, the base feature ⟨*coffee house*⟩ will not fire unless both features ⟨*coffee*⟩ and ⟨*house*⟩ also fire.

- Another option is to apply non-linear transformations to the feature vector. Recall that the feature function $\boldsymbol{f}(\boldsymbol{x}, y)$ may be composed of a vector of word counts, padded by zeros. We can think of these word counts as basic features, and apply non-linear transformations, such as $\boldsymbol{x} \circ \boldsymbol{x}$ or $|\boldsymbol{x}|$.

- There is some work in NLP on using kernels for strings, bags-of-words, sequences, trees, etc. Kernelized learning algorithms are outside the scope of this class (Collins and Duffy, 2001; Zelenko et al., 2003). Kernel-based learning can be seen as a generalization of algorithms such k-nearest-neighbors, which classifies instances by considering the labels of the k most similar instances in the training set (Hastie et al., 2009).

---

[7]Assuming your feature matrix is full-rank.

- Boosting (Freund et al., 1999) and decision tree algorithms (Schmid, 1994) sometimes do well on NLP tasks, but they are used less frequently these days, especially as the field increasingly emphasizes big data and simple classifiers.

- More recent work has shown how **deep learning** can perform non-linear classification. One way to use deep learning in NLP is by learning word representations while jointly learning how these representations combine to classify instances (Collobert and Weston, 2008). This approach is very hot at the moment, so I will discuss it towards the end of the semester.

5.6 Summary of classifiers

So now we've talked about four different classifiers. That's it! No more classifiers in this class. Yay? Anyway, let's review.

	Naive Bayes	Logistic Regression	Perceptron	PA
Objective	Joint likelihood	Conditional likelihood	0-1 loss	Hinge loss
estimation	$\max \sum_i \log \mathrm{p}(\boldsymbol{x}_i, y_i)$	$\max \sum_i \log \mathrm{p}(y_i \mid \boldsymbol{x}_i)$	$\min \sum_i \delta(y_i, \hat{y})$	$\sum_i [1 - \gamma(\boldsymbol{\theta}; \boldsymbol{x}_i, y_i)]_+$
tuning	$\theta_{ij} = \frac{c(x_i, y=j)+\alpha}{c(y=j)+V\alpha}$ smoothing α	$\frac{\partial \mathcal{L}}{\partial \boldsymbol{\theta}} = \sum_i \boldsymbol{f}(\boldsymbol{x}_i, y_i) - E[\boldsymbol{f}(\boldsymbol{x}_i, y)]$ regularizer $\lambda \|\boldsymbol{\theta}\|_2^2$	$\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} + \boldsymbol{f}(\boldsymbol{x}_i, y_i) - \boldsymbol{f}(\boldsymbol{x}_i, \hat{y})$ weight averaging	$\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} + \tau_t (\boldsymbol{f}(\boldsymbol{x}_i, y_i) - \boldsymbol{f}(\boldsymbol{x}_i, \hat{y}))$ slack penalty C
complexity	$\mathcal{O}(NV)$	$\mathcal{O}(NVT)$	$\mathcal{O}(NVT)$	$\mathcal{O}(NVT)$
easy?	very	not really	yes	yes
probabilities?	yes	yes	no	no
features?	no	yes	yes	yes

Table 5.1: Comparison of classifiers. N = number of examples, V = number of features, T = number of instances.

Chapter 6

Word-sense disambiguation

Todo: add notes about WSD here

Chapter 7

Learning without supervision

So far we've assumed the following setup:

- A **training set** where you get observations x_i and labels y_i

- A **test set** where you only get observations x_i

What if you never get labels y_i?
For example, you get a bunch of text, and you suspect that there are at least two different meanings for the word *concern*.[1]
    The immediate context includes two groups of words:

- services, produces, banking, pharmaceutical, energy, electronics

- about, said, that, over, in, with, had

    Suppose we plot each instance of *concern* on a graph

- x-axis is the density of words in group 1

- y-axis is the density of words in group 2

Two blobs might emerge. These blobs would correspond to two different sense of *concern*.

- But in reality, we don't know the word groupings in advance.

---

[1] example from Pedersen and Bruce (1997)

- We have to try to apply the same idea in a very high dimensional space, where every word gets its own dimension (and most dimensions are irrelevant!)

- Or we have to automatically find a low-dimensional projection. More on that much later in the course.

Here's a related scenario:

- You look at thousands of news articles from today

- Plot them on a graph of *Miley* vs *Syria*

- Three clumps emerge (Miley, Syria, others)

- Those clumps correspond to natural document classes

- Again, in reality this is a hugely high-dimensional graph

So these examples show that we can find structure in data, even without labels.

7.1 K-means clustering

You might know about classic clustering algorithms like K-means. These algorithms are iterative:

1. Guess the location of cluster centers.

2. Assign each point to the nearest center.

3. Re-estimate the centers as the mean of the assigned points.

4. Goto 2.

This is an algorithm for finding coherent "blobs" of documents.
There is a variant called "soft k-means."

- Instead of assigning each point x_i to a specific cluster z_i

- You assign it a distribution over clusters $q_i(z_i)$

We're now going to explore a more principled, statistical version of soft K-means, called EM clustering.

By understanding the statistical principles underlying the algorithm, we can extend it in a number of cool ways.

7.2 The Expectation-Maximization Algorithm

Let's go back to the Naive Bayes model:

$$\log p(\boldsymbol{x}, \boldsymbol{y}; \boldsymbol{\phi}, \mu) = \sum_i \log p(\boldsymbol{x}_i | y_i; \boldsymbol{\phi}) P(y_i; \mu)$$

For example, $\boldsymbol{x}$ can describe the documents that we see today, and $\boldsymbol{y}$ can correspond to their labels. But suppose we never observe y_i? Can we still do something?

Since we don't know $\boldsymbol{y}$, let's marginalize it:

$$\log p(\boldsymbol{x}) = \log \sum_{\boldsymbol{y}} p(\boldsymbol{x} | \boldsymbol{y}; \boldsymbol{\phi}) p(\boldsymbol{y}; \mu) \tag{7.1}$$

$$= \log \sum_{\boldsymbol{y}} \prod_i p(\boldsymbol{x}_i | y_i; \boldsymbol{\phi}) p(y_i; \mu) \tag{7.2}$$

$$= \sum_i \log \sum_{y_i} p(\boldsymbol{x}_i | y_i; \boldsymbol{\phi}) p(y_i; \mu) \tag{7.3}$$

Now we introduce an auxiliary variable $\boldsymbol{q}_i$, for each y_i. We have the usual constraints: $\sum_y q_i(y) = 1$ and $\forall y, q_i(y) \geq 0$. In other words, q_i defines a probability distribution over Y, for each instance i.

Now since $\frac{q_i(y)}{q_i(y)} = 1$,

$$\log p(\boldsymbol{x}) = \sum_i \log \sum_{y_i} p(\boldsymbol{x}_i | y_i; \boldsymbol{\phi}) p(y_i; \mu) \frac{q_i(y)}{q_i(y)}$$

$$= \sum_i \log E_q[\frac{p(\boldsymbol{x}_i | y; \boldsymbol{\phi}) p(y; \mu)}{q_i(y)}],$$

by the definition of expectation. (Note that E_q just means the expectation under the distribution q.)

Now we apply *Jensen's inequality*. Jensen's equality says that because $\log$ is concave, we can push it inside the expectation, and obtain a lower bound.

$$\log p(\boldsymbol{x}) \geq \sum_i E_q[\log \frac{p(\boldsymbol{x}_i | y; \boldsymbol{\phi}) p(y_i; \mu)}{q_i(y)}]$$

$$\mathcal{J} = \sum_i E_q[\log p(\boldsymbol{x}_i | y; \boldsymbol{\phi})] + E_q[\log p(y; \mu)] - E_q[q_i(y)]$$

By maximizing $\mathcal{J}$, we are maximizing a lower bound on the joint log-likelihood $\log p(\boldsymbol{x})$.

Now, $\mathcal{J}$ is a function of two arguments:

- the distributions $q_i(\boldsymbol{y})$ for each i

- the parameters μ and ϕ

We'll optimize with respect to each of these in turn, holding the other one fixed.

The E-step

First, we expand the expectation in the lower bound as:

$$\mathcal{J} = \sum_i E_q[\log p(\boldsymbol{x}_i|y; \boldsymbol{\phi})] + E_q[\log p(y; \mu)] - E_q[q_i(y)]$$

$$= \sum_i \sum_y q_i(y) \left(\log p(\boldsymbol{x}_i|Y_i = y; \boldsymbol{\phi}) + \log p(y; \mu) - \log q_i(y)\right)$$

As in relative frequency estimation of Naive Bayes, we need to add a Lagrange multiplier to ensure $\sum_y q_i(y) = 1$, so

$$\mathcal{J} = \sum_i \sum_y q_i(y) \left(\log p(\boldsymbol{x}_i|Y_i = y; \boldsymbol{\phi}) + \log p(y; \mu) - \log q_i(y)\right) + \lambda_i(1 - \sum_y q_i(y))$$

$$\frac{\partial \mathcal{J}}{\partial q_i(y)} = \log p(\boldsymbol{x}_i|Y_i = y; \boldsymbol{\phi}) + \log p(y; \boldsymbol{\theta}) - \log q_i(y) - 1 - \lambda_i$$

$$\log q_i(y) = \log p(\boldsymbol{x}_i|Y_i = y; \boldsymbol{\phi}) + \log p(y; \mu) - 1 - \lambda_i$$

$$q_i(y) \propto p(\boldsymbol{x}_i|Y_i = y; \boldsymbol{\phi})p(y; \mu)$$

$$\propto p(\boldsymbol{x}_i, y; \boldsymbol{\phi}, \mu)$$

$$q_i(y) = \frac{p(\boldsymbol{x}_i, y; \boldsymbol{\phi}, \mu)}{\sum_{y'} p(\boldsymbol{x}_i, y'; \boldsymbol{\phi}, \mu)}$$

$$= P(Y_i = y|\boldsymbol{x}_i; \boldsymbol{\theta}, \phi)$$

After normalizing, each $q_i(y)$ – which is the soft distribution over clusters for data $\boldsymbol{x}_i$ – is set to the conditional probability $P(y_i|\boldsymbol{x}_i)$ under the current parameters μ, ϕ.

This is called the E-step, or "expectation step," because it is derived from updating the expected likelihood under $q(\boldsymbol{y})$.

The M-step

Next, we hold $q(\boldsymbol{y})$ fixed and update the parameters. Let's do ϕ, which parametrizes $\mathrm{p}(\boldsymbol{x}|y)$. Again, we start by adding Lagrange multipliers to the lower bound,

$$\mathcal{J} = \sum_i \sum_y q_i(y) \left(\log \mathrm{p}(\boldsymbol{x}_i|Y_i = y; \boldsymbol{\phi}) + \log \mathrm{p}(y; \mu) - \log q_i(y)\right) + \sum_y \lambda_y (1 - \sum_j \phi_{y,j})$$

$$\frac{\partial \mathcal{J}}{\partial \phi_{y,j}} = \sum_i q_i(y) \frac{x_{i,j}}{\phi_{y,j}} - \lambda_y$$

$$\lambda_h \phi_{y,j} = \sum_i q_i(y) x_{i,j}$$

$$\phi_{y,j} = \frac{\sum_i q_i(y) x_{i,j}}{\sum_{j'} \sum_i q_i(y) x_{i,j'}} = \frac{E_q[\text{count}(y, j)]}{E_q[\text{count}(y)]}$$

So $\boldsymbol{\phi}_y$ is now equal to the relative frequency estimate of the **expected counts** under the distribution $q(y)$.

- As in supervised Naïve Bayes, we can apply smoothing to add α to all these counts

- The update for μ is identical: $\mu_y \propto \sum_i q_i(y)$, the expected proportion of cluster $Y = y$. If needed, we can add smoothing here too.

- So, everything in the M-step is just like Naive Bayes, except we used expected counts rather than observed counts.

Coordinate ascent

Algorithms that alternate between updating various subsets of the parameters are called "coordinate-ascent" algorithms.

The objective function $\mathcal{J}$ is **biconvex**, meaning that it is separately convex in $q(\boldsymbol{y})$ and $\langle \mu, \boldsymbol{\phi} \rangle$, but it is not jointly convex.

- Each step is guaranteed not to decrease $\mathcal{J}$

- This is called hill-climbing: you never go down.

- Specifically, EM is guaranteed to converge to a **local optima** – a point which is as good or better than any of its immediate neighbors. But there may be many such points.

- But the overall procedure is **not** guaranteed to find a global maximum.

- This means that initialization is important: where you start can determine where you finish.

- This is not true in most of the supervised learning algorithms that we have considered, such as logistic regression; in that case, we are optimizing $\log \mathrm{p}(\boldsymbol{y}|x; \boldsymbol{\theta})$, which is defined so as to be convex with respect to the parameter $\boldsymbol{\theta}$. This means that for logistic regression (and many other supervised learning algorithms), we don't need to worry about initialization, because it won't affect our ultimate solution: we are guaranteed to reach the global minimum.

7.3 Applications of EM

EM is not really an "algorithm" like, say, quicksort. Rather, it's a framework for learning with missing data. The recipe for using EM on a problem of interest to you is something like this:

- Introduce latent variables $\boldsymbol{z}$, such that it's easy to write the probability $P(\mathcal{D}, \boldsymbol{z})$, where $\mathcal{D}$ is your observed data, and easy to estimate the associated parameters.

- Derive the E-step updates for $q(\boldsymbol{z})$, which is typically factored as $q(\boldsymbol{z}) = \prod_i q_{z_i}(z_i)$.

Some applications of this basic setup are presented here.

Word sense clustering

In the "demos" folder, you can find a demonstration of expectation-maximization for word sense clustering. I assume we know that there are two senses, and that the senses can be distinguished by the contextual information in the document. The basic framework is identical to the clustering model of EM as presented above.

Semi-supervised learning

Nigam et al. (2000) offer another application of EM: **semi-supervised learning**. They apply this idea to document classification in the classic "20 Newsgroup" dataset.

- In this setting, we have labels for some of the instances, $\langle \boldsymbol{x}^{(\ell)}, \boldsymbol{y}^{(\ell)} \rangle$, but not for others, $\langle \boldsymbol{x}^{(u)} \rangle$.

- Can unlabeled data improve learning?

We will choose parameters to maximize the joint likelihood,

$$\log \mathrm{p}(\boldsymbol{x}^{(\ell)}, \boldsymbol{x}^{(u)}, \boldsymbol{y}^{(\ell)}) = \log \mathrm{p}(\boldsymbol{x}^{(\ell)}, \boldsymbol{y}^{(\ell)}) + \log \mathrm{p}(\boldsymbol{x}^{(u)}) \tag{7.4}$$

- We treat the labels of the unlabeled documents as missing data. In the E-step we impute $q(y)$ for the unlabeled documents only.

- The M-step computes estimates of μ and ϕ from the sum of the observed counts from $\langle \boldsymbol{x}^{(\ell)}, \boldsymbol{y}^{(\ell)} \rangle$ and the expected counts from $\langle \boldsymbol{x}^{(u)} \rangle$ and $q(\boldsymbol{y})$.

- We can further parametrize this approach by weighting the unlabeled documents by a scalar λ, which is a tuning parameter.

Multi-component modeling

- One of the classes in 20 newsgroups is `comp.sys.mac.hardware`.

- Suppose that there are two kinds of posts: reviews of new hardware, and question-answer posts about hardware problems.

- The language in these **components** of the mac.hardware class might have little in common.

- So we might do better if we model these components separately.

We can envision a new generative process here:

- For each document i,

  - draw the label $y_i \sim \text{Categorical}(\theta)$
  - draw the component $z_i | y_i \sim \text{Categorical}(\psi_{y_i})$
  - draw the vector of counts $\boldsymbol{x}_i | z_i \sim \text{Multinomial}(\phi_{z_i})$

Our labeled data includes $\langle \boldsymbol{x}_i, y_i \rangle$, but not z_i, so this is another case of missing data.

$$
\begin{aligned}
\mathrm{p}(\boldsymbol{x}_i, y_i) &= \sum_z \mathrm{p}(\boldsymbol{x}_i, y_i, z) \\
&= \mathrm{p}(\boldsymbol{x}_i | z; \boldsymbol{\phi}) \mathrm{p}(z | y_i; \psi) \mathrm{p}(y_i; \mu)
\end{aligned}
$$

Again, we can apply EM

- We need a distribution over the missing data, $q_i(z)$. This is updated during the E-step.

- During the m-step, we compute:

$$
\begin{aligned}
\psi_{y,z} &= \frac{E_q[\text{count}(y, z)]}{\sum_{z'} E_q[\text{count}(y, z')]} \\
\phi_{j,y,z} &= \frac{E_q[\text{count}(z, j)]}{\sum_{j'} E_q[\text{count}(z, j')]}
\end{aligned}
$$

- Suppose we assume each class y is associated with K components, $\mathcal{Z}_y$. We can add a constraint to the E-step so that $q_i(z) = 0$ if $z \notin \mathcal{Z}_y \wedge Y_i = y$.

Chapter 8

Language models

A **language model** is used to compute the probability of a sequence of text. Why would we want to do this? Thus far, we have considered problems where text is the **input**, and we want to select an output, such as a document class or a word sense. But in many of the most prominent problems in language technology, text itself is the output:

- machine translation
- speech recognition
- summarization

As we will soon see, we can produce more **fluent** text output by computing the probability of the text.

Specifically, suppose we have a vocabulary of word types

$$\mathcal{V} = \{aardvark, abacus, \ldots, zither\} \tag{8.1}$$

Given a sequence of word tokens $w_1, w_2, \ldots, w_M$, with $w_i \in \mathcal{V}$, we would like to compute the probability $p(w_1, w_2, \ldots, w_M)$. We will do this in a data-driven way, assuming we have a **corpus** of text.

- For now, we'll assume that the vocabulary $\mathcal{V}$ covers all the word tokens that we will ever see. Of course, we can enforce this by allocating a special token ♠ for unknown words. However, this might not be a great solution, as we will see later.

- Language models typically make an independence assumption across sentences, $p(\boldsymbol{s}_1, \boldsymbol{s}_2, \ldots) = \prod_j p(\boldsymbol{s}_j)$, where each sentence $\boldsymbol{s}_j = [\boldsymbol{w}_1, \boldsymbol{w}_2, \ldots, \boldsymbol{w}_{N_j}]$.

So for our purposes, it is sufficient to compute the probability of sentences. The justification for this assumption is that the probability of words that are not in the same sentence don't depend on each other too much. Clearly this isn't true: once I mention *Manuel Noriega* once in a document, I'm far more likely to mention him again (Church, 2000). But the dependencies between words within a sentence are usually even stronger, and are more relevant to the fluency considerations inherent in applications such as translation and speech recognition (which are typically evaluated at the sentence level anyway).

So how can we compute the probability of a sentence? The simplest idea would be to apply a **relative frequency estimator**:

$$\text{p}(\textit{Computers are useless, they can only give you answers}) \tag{8.2}$$

$$= \frac{\text{count}(\textit{Computers are useless, they can only give you answers})}{\text{count(all sentences ever spoken)}} \tag{8.3}$$

It's useful to think about this estimator in terms of bias and variance.

- In the theoretical limit of infinite data, it might work. But in practice, we are asking for accurate counts over an infinite number of events, since sentences can be arbitrarily long.

- Even if we set an aggressive upper bound of, say, $n = 20$, the number of possible sentences is $\#|\mathcal{V}|^{20}$. A small vocabularly for English would have $\#|\mathcal{V}| = 10^4$, so we would have 10^{80} possible sentences.

- Clearly, this estimator is extremely data-hungry. We need to introduce bias to have a chance of making reliable estimates.

Are language models meaningful?   What are the probabilities of the following two sentences?

- *Colorless green ideas sleep furiously*

- *Furiously sleep ideas green colorless*

Noam Chomsky used this pair of examples to argue that the probability of a sentence is a meaningless concept:

- Any English speaker can tell that the first sentence is grammatical but the second sentence is not.

- Yet neither sentence, nor their substrings, had ever appeared at the time that Chomsky wrote this article (they have appeared lots since then).

- Thus, he argued, empirical probabilities can't distinguish grammatical from ungrammatical sentences.

Pereira (2000) showed that by identifying *classes* of words (e.g., noun, verb, adjective, adverb — but not necessarily these grammatical categories), it is easy to show that the first sentence is more probable than the second. We will talk about class-based language models later.

Are language models useful? Suppose we want to translate a sentence from Spanish:

- *El cafe negro me gusta mucho*.

- Word-for-word: *The coffee black me pleases much*.

- But a good language model of English will tell us:

$$P(\textit{The coffee black me pleases much}) < P(\textit{I like black coffee a lot}) \qquad (8.4)$$

- How can we use this fact?

Warren Weaver on translation as decoding:

> When I look at an article in Russian, I say: 'This is really written in English, but it has been coded in some strange symbols. I will now proceed to decode.'

This motivates a generative model (like Naive Bayes!):

- English sentence $\boldsymbol{w}^{(e)}$ generated from language model $\mathrm{p}_e(\boldsymbol{w}^{(e)})$

- Spanish sentence $\boldsymbol{w}^{(s)}$ generated from noisy channel $\mathrm{p}_{s|e}(\boldsymbol{w}^{(s)}|\boldsymbol{w}^{(e)})$

(picture)
Then the **decoding** problem is: $\max_{\boldsymbol{w}^{(e)}} \mathrm{p}(\boldsymbol{w}^{(e)}|\boldsymbol{w}^{(s)}) \propto \mathrm{p}(\boldsymbol{w}^{(s)}, \boldsymbol{w}^{(e)}) = \mathrm{p}(\boldsymbol{w}^{(e)})\mathrm{p}(\boldsymbol{w}^{(s)}|\boldsymbol{w}^{(e)})$

- The **translation model** is $p(\boldsymbol{w}^{(s)}|\boldsymbol{w}^{(e)})$. This ensures the **adequacy** of the translation.

- The **language model** is $p(\boldsymbol{w}^{(e)})$. This ensures the **fluency** of the translation.

What else can we model with a noisy channel?

- Speech recognition (original = words; encoded = sound)

- Spelling correction (original = well-spelled text; encoded = text with spelling mistakes)

- Part of speech tagging (original = tags; encoded = words)

- Parsing (original = parse tree; encoded = words)

- ...

The noisy channel model allows us to decompose NLP systems into two parts:

- The translation model, which we need labeled data to estimate.

- The language model, which we need only *unlabeled* data to estimate.

Since there is always more unlabeled data, this means we can improve NLP systems just by improving $p_e(\boldsymbol{w})$.

8.1 N-gram language models

We began with the relative frequency estimator,

$$p(\textit{Computers are useless, they can only give you answers}) \tag{8.5}$$

$$= \frac{\text{count}(\textit{Computers are useless, they can only give you answers})}{\text{count}(\text{all sentences ever spoken})} \tag{8.6}$$

We'll define the probability of a sentence as the probability of the words (in order): $p(\boldsymbol{w}) = p(w_1, w_2, \ldots, w_M)$. We can apply the chain rule:

$$
\begin{aligned}
p(\boldsymbol{w}) =& p(w_1, w_2, \ldots, w_M) \\
=& p(w_1)p(w_2 \mid w_1)p(w_3 \mid w_2, w_1) \ldots p(w_M \mid w_{M-1}, \ldots, w_1)
\end{aligned}
$$

Each element in the product is the probability of a word given all its predecessors. We can think of this as a *word prediction* task: *Computers are [BLANK]*. The relative frequency estimate:

$$p(\textit{useless}|\textit{computers are}) = \frac{\text{count}(\textit{computers are useless})}{\sum_x \text{count}(\textit{computers are } x)} = \frac{\text{count}(\textit{computers are useless})}{\text{count}(\textit{computers are})}$$

Note that we haven't made any approximations yet, and we could have applied the chain rule in reverse order, $p(\boldsymbol{w}) = p(w_M)p(w_{M-1}|w_M)\ldots$, or in any other order. But this means that we also haven't really improved anything either: to compute the conditional probability $P(W_M \mid W_{M-1}, W_{M-2}, \ldots)$, we need to model $\#|\mathcal{V}|^{N-1}$, with $\#|\mathcal{V}|$ events. We can't even **store** this probability distribution, let alone reliably estimate it.

N-gram models

N-gram models make a simple approximation: condition on only the past $n-1$ words.

$$p(w_m \mid w_{m-1}\ldots w_1) \approx P(w_m \mid w_{m-1}, \ldots, w_{m-n+1})$$

This means that the probability of a sentence $\boldsymbol{w}$ can be computed as

$$p(w_1, \ldots, w_M) \approx \prod_m p(w_m \mid w_{m-1}, \ldots, w_{m-n+1})$$

- To compute the probability of a whole sentence, it's convenient to pad the beginning and end with special symbols $\Diamond$ and $\square$. Then the bigram ($n = 2$) approximation to the probability of *I like black coffee* is:

$$p(I \mid \Diamond)p(\textit{like} \mid I)p(\textit{black} \mid \textit{like})p(\textit{coffee} \mid \textit{black})p(\square \mid \textit{coffee}) \qquad (8.7)$$

- In this model, we have to estimate and store the probability of only $\#|\mathcal{V}|^n$ events. A very common choice is a trigram model, in which $n = 3$.

- The n-gram probabilities can be determined by relative frequency estimation,

$$p(w|u, v) = \frac{\text{count}(u, v, w)}{\text{count}(u, v)} = \frac{\text{count}(u, v, w)}{\sum_{w'} \text{count}(u, v, w')} \qquad (8.8)$$

There could be too problems with an n-gram language model:

- n **is too small**. In this case, we are missing important linguistic context. Consider the following sentences:

    - Gorillas *always like to groom* **THEIR** *friends.*
    - *The* computer *that's on the 3rd floor of our office building* **CRASHED***.*

    The bolded words depend crucially on their predecessors in italics: *their* depends on knowing that *gorillas* is plural, and *crashed* depends on knowing that the subject is a *computer*. The resulting model would offer probabilities that are too low for these sentences, and too high for sentences that fail basic linguistic tests like number agreement.

- n **is too big**. In this case, we can't make good estimates of the n-gram parameters from our dataset. See the slides for some examples of this.

- These two problems point to another **bias/variance** tradeoff. Can you see how it works?

- In reality, we often have **both** problems! Language is full of long-range dependencies, and datasets are small.

We will seek approaches to keep n large, while still making low-variance estimates of the underlying parameters. To do this, we will introduce a different sort of bias: **smoothing**. But before we talk about that, let's consider how we can evaluate language models.

8.2   Evaluating language models

- Because language models are typically components of larger systems (language modeling is not really an application itself), we would prefer **extrinsic evaluation**: does the LM help the task (translation or whatever). But this is often hard to do, and depends on details of the overall system which may be irrelevant to language modeling.

- **Intrinsic evaluation** is task-neutral. Better performance on intrinsic metrics may be expected to improve extrinsic metrics across a variety of tasks (unless we are over-optimizing the intrinsic metric).

Held-out likelihood

A popular intrinsic metric is the **held-out likelihood**.

- We obtain a test corpus, and compute the (log) probability according to our model. It is crucial that the words in this corpus were not used in estimating the model itself.

- A good model should assign high probability to this held-out data.

- Specifically, we compute

$$\ell(\boldsymbol{w}) = \sum_i \sum_m \log \mathrm{p}(w_m^{(i)}|w_{m-1}^{(i)}, \ldots, w_{m-n+1}^{(i)}), \tag{8.9}$$

  for all sentences $\boldsymbol{w}^{(i)}$ in the held-out corpus.

Perplexity

Perplexity is a transformation of the held-out likelihood, into an information-theoretic quantity. Specifically, we compute

$$PP(\boldsymbol{w}) = 2^{-\frac{\ell(\boldsymbol{w})}{M}}, \tag{8.10}$$

where M is the total number of tokens in the held-out corpus.

- After this transformation, we now prefer lower values. In the limit, we obtain probability 1 for our held-out corpus, with $PP = 2^{-\log 1} = 1$.

- Assume a uniform, unigram model in which $P(s_i) = \frac{1}{V}$ for all V words in the vocabulary. Then,

$$PP(\boldsymbol{w}) = \left[\left(\frac{1}{V} \right)^M \right]^{-\frac{1}{M}}$$
$$= \left(\frac{1}{V} \right)^{-1} = V$$

- We can think of perplexity as the *weighted branching factor* at each word in the sentence.

  - If we have solved the word prediction problem perfectly, $PP(\boldsymbol{w}) = 1$, because there is only one possible choice.

- If we have only a uniform model that assigns equal probability to every word, $PP(\boldsymbol{w}) = V$.

- Most models fall somewhere in between.

- Here's how you remember: lower perplexity is better, because you are less perplexed.

Example On 38M tokens of WSJ, $V \approx 20K$, (Jurafsky and Martin, 2009, page 97) obtain these perplexities on a 1.5M token test set.

- Unigram: 962

- Bigram: 170

- Trigram: 109

Will it keep going down? See slides from (Manning and Schütze, 1999).

Information theory*

Perplexity is very closely related to the concept of entropy, the expected value of the information contained in each word.

$$H(P) = -\sum_{\boldsymbol{w}} p(\boldsymbol{w}) \log p(\boldsymbol{w}) \tag{8.11}$$

The true entropy of English (or any real language) is unknown. Claude Shannon, one of the founders of information theory, wanted to compute upper and lower bounds. He would read passages of 15 characters to his wife, and ask her to guess the next character, recording the number of guesses it took for her to get the correct answer. As a fluent speaker of English, his wife could provide a reasonably tight bound on the number of guesses needed per character. **Question: is this an upper bound or a lower bound?**

Cross-entropy is a relationship between two probability distributions, the true one $P(W)$ and an estimate $Q(W)$.

$$\begin{aligned}
H(P,Q) &= E_P[\log Q]\\
&= -\sum_{\boldsymbol{w}} \mathrm{p}(\boldsymbol{w}) \log q(\boldsymbol{w})\\
&= \sum_{\boldsymbol{w}} \mathrm{p}(\boldsymbol{w}) \log \frac{\mathrm{p}(\boldsymbol{w})}{q(\boldsymbol{w})} - \mathrm{p}(\boldsymbol{w}) \log \mathrm{p}(\boldsymbol{w})\\
&= D_{KL}(P||Q) + H(Q)
\end{aligned}$$

So the cross-entropy is the KL-divergence between P and Q – a non-symmetric distance measure between distributions, which we will see again later in the course – plus the entropy of P. Since P is the language itself, we can only control Q, and minimizing the cross-entropy is equivalent to minimizing the KL-divergence.

We do not have access to the true $P(W)$, just a sequence $\boldsymbol{w} = \{w_1, w_2, \ldots, \}$, which is sampled from $P(W)$. In the limit, the length of $\boldsymbol{w}$ is infinite, so we have,

$$\begin{aligned}
H(P,Q) &= -\sum_{\boldsymbol{w}} \mathrm{p}(\boldsymbol{w}) \log q(vw)\\
&= -\lim_{M\to\infty} \frac{1}{M} \log q(\boldsymbol{w})\\
&\approx -\frac{1}{M} \log q(\boldsymbol{w})\\
PP(S) &= 2^{-\frac{1}{M} \log q(vw)}
\end{aligned}$$

A good language model has low cross-entropy with $P(W)$, and thus low perplexity.

Further aside : A related topic in psycholinguistics is the "constant entropy rate hypothesis," also called the "uniform information density hypothesis." The hypothesis is that speakers should prefer linguistic choices that convey a uniform amount of information over time (Jaeger, 2010). Some evidence:

- Speakers shorten predictable words, lengthen unpredictable ones

- High-entropy sentences take longer to read

- Syntactic reductions (e.g., *I'm* versus *I am*) are more likely when the reducible word contains less information.

8.3   Smoothing and discounting

We want to estimate $P(W)$ from sparse statistics, avoiding $p(\boldsymbol{w}) = 0$.

Laplace/Lidstone smoothing

Simplest idea: just add "pseudo-counts"

$$p_{\text{Laplace}}(w \mid v) = \frac{\text{count}(v, w) + \alpha}{\sum_{w'} \text{count}(v, w') + V\alpha} \tag{8.12}$$

Anything that we add to the numerator (α) must also appear in the dominator ($V\alpha$). We can capture this with the concept of **effective counts**:

$$c_i^* = (c_i + \alpha)\frac{N}{N + V\alpha}$$

The **discount** for each n-gram is:

$$d_i = \frac{c_i^*}{c_i} = \frac{(c_i + \alpha)}{c_i}\frac{N}{(N + \alpha)}$$

- In general, this is called Lidstone smoothing

- When $\alpha = 1$, we are doing Laplace smoothing

- When $\alpha = 0.5$, we are following Jeffreys-Perks law

- Manning and Schütze (1999) offer more insight on the justifications for Jeffreys-Perks smoothing

Discounting and backoff

Discounting "borrows" probability mass from observed n-grams and redistributes it.

- In Lidstone smoothing, we borrow probability mass by increasing the denominator of the relative frequency estimates, and redistribute it by increasing the numerator for all n-grams.

- Instead, we could borrow the same amount of probability mass from all observed counts, and redistribute it among only the unobserved counts. This is called **absolute discounting.**

- For example, if we set an absolute discount $d = 0.1$ in a trigram model, we get: $p(w|\textit{denied the}) =$

word	counts c	effective counts c^*	unsmoothed probability	smoothed probability
allegations	3	2.9	0.429	0.414
reports	2	1.9	0.286	0.271
claims	1	0.9	0.143	0.129
request	1	0.9	0.143	0.129
charges	0	0.2	0.000	0.029
benefits	0	0.2	0.000	0.029
. . .				

- We need not redistribute the probability mass equally. Instead, we can **back-off** to a lower-order language model.

- In other words: if you have trigrams, use trigrams; if you don't have trigrams, use bigrams; if you don't even have bigrams, use unigrams. (And what if you don't even have unigrams?). This is called **Katz backoff**.

$$c^*(u, v) = c(u, v) - d$$

$$\mathrm{P_{backoff}}(v \mid u) = \begin{cases} \frac{c^*(u,v)}{c(u)} & \text{if } c(u, v) > 0 \\ \alpha(u) \times \frac{\mathrm{P_{backoff}}(v)}{\sum_{v':c(u,v')=0} \mathrm{P_{backoff}}(v')} & \text{if } c(u, v) = 0 \end{cases}$$

Typically we can set d to minimize perplexity on a development set.

Interpolation

An alternative to this discounting scheme is to do interpolation: the probability of a word in context is a weighted sum of its probabilities across progressively shorter contexts.

Instead of choosing a single n-gram order, we can take the weighted average:

$$\begin{aligned} \mathrm{P_{Interpolation}}(w|u, v) &= \lambda_1 \mathrm{p}_1^*(w|u, v) \\ &+ \lambda_2 \mathrm{p}_2^*(w|u) \\ &+ \lambda_3 \mathrm{p}_1^*(w) \end{aligned}$$

- p_k^* is the maximum likelihood estimate (MLE) of a k-gram model

- Constraint: $\sum_z \lambda_z = 1$

- We can tune λ on heldout data...

- Or we can use **expectation maximization**!

EM for interpolation   We can add a latent variable z_m, indicating the order of the n-gram that generated word w_m. Generative story:

- For each word m

  - Draw $z_m \sim \text{Categorical}(\lambda(w_m))$
  - Draw $w_m \sim p_{z_m}^*(w_m | s_{m-1}, \dots s_{m-z_m})$

As always we have two quantities of interest in our EM application:

- The parameters, λ.

- Our beliefs about the latent variables. Let $q_m(z)$ be our degree of belief that word token w_m was generated from a n-gram of order z.

Having defined these quantities, we can derive EM updates:

- **E-step**: $q_m(z) = p(z|w_{1:m}) = \frac{p_z^*(w_m|w_{m-1},\dots,w_{m-z+1})}{\sum_{z'} p_{z'}^*(w_m|w_{m-1},\dots w_{m-z'+1})} p(z'|\lambda(w_m))$

- **M-step**: $\lambda(w)_z = \frac{E_q[\text{count}(W=w, Z=z)]}{\sum_{z'} E_q[\text{count}(W=w, Z=z')]}$

By running the EM algorithm, we can obtain a good estimate of λ, which we can then use for unseen data. It should be clear how we can extend this approach to trigrams and beyond; Collins (2013) offers more details.

Kneser-ney smoothing

Kneser-ney smoothing also incorporates discounting, but redistributes the resulting probability mass in a different way. Consider the example:
   I recently visited

- *Francisco?*

- *Duluth?*

Key idea: some words are more **versatile** than others.

- Suppose $p^*(Francisco) > p^*(Duluth)$, and $c(visited\ Francisco) = c(visited\ Duluth) = 0$.

- We would still guess that $p(visited\ Duluth) > P(visited\ Francisco)$, because *Duluth* is a more versatile word.

We define the Kneser-Ney bigram probability as

$$
p_{KN}(v|u) = \begin{cases} \frac{\text{count}(u,v)-d}{\text{count}(u)}, & \text{count}(u,v) > 0 \\ \alpha(u)p_{\text{continuation}}(v), & \text{otherwise} \end{cases}
$$

$$
p_{\text{continuation}}(v) = \frac{\#|u : \text{count}(u,v) > 0|}{\sum_{v'} \#|u' : \text{count}(u',v') > 0|}
$$

- We reserve probability mass using absolute discounting d.

- The *continuation probability* $p_{\text{continuation}}(u)$ is proportional to the number of observed contexts in which u appears.

- As in Katz backoff, $\alpha(v)$ makes the probabilities sum to 1

- In practice, interpolation works a little better than backoff

$$
p_{KN}(v|u) = \frac{\text{count}(u,v)-d}{\text{count}(u)} + \lambda(u)p_{\text{continuation}}(v) \tag{8.13}
$$

- This idea of counting contexts may seem heuristic, but actually there is a cool justification from Bayesian nonparametrics (Teh, 2006).

8.4 Other types of Language Models

Interpolated Kneser-Ney is pretty close to state-of-the-art. But there are some interesting other types of language models, and they apply ideas that we have already learned.

Mixed-order n-gram models

Saul and Pereira (1997) described a "mixed-order" n-gram model, where you condition on multiple bigram contexts, skipping over intermediate words:

$$p(w_m|w_{m-1}, \ldots, w_{m-n+1}) = \sum_k \lambda_k(w_{m-k})\tilde{p}(w_m|w_m - k) \prod_j^{k-1} [1 - \lambda_j(w_{m-j})] \quad (8.14)$$

- This is an **interpolated** model, because we are taking the weighted average over a bunch of bigram probabilities.

- Note that the interpolation weight depends on the context word, $\lambda_k(w_{m-k})$. This means that some words can prefer certain dependency lengths — for example, adjectives might prefer short dependencies, since they tend to affect adjacent nouns, while verbs might prefer longer dependencies, since they can affect indirect objects that are further away.

- The final product ensures that the weights in any particular context must add up to one: each λ_k is taking a slice of the probability mass that has already been used by the earlier contexts $j < k$.

- The parameters $\lambda_k(w)$ can be estimated by expectation maximization, just like in the interpolated N-gram model above.

Class-based language models

The reason we need smoothing is because the trigram probability model $p(w|u, v)$ has a huge number of parameters. Let's simplify:

$$p_{\text{class}}(w|v) = \sum_z P(w|z; \theta)P(z|v; \phi),$$

where $z \in [1, K], K \ll V$.

We get a bigram probability using $2VK$ parameters instead of V^2.

We could use EM to estimate θ and ϕ (Saul and Pereira, 1997).

- The latent variable is the class z, so the e-step updates $q_m(z)$

- The parameters are θ and ϕ, which can be updated in the M-step.

But this is usually too slow, so there are approximate algorithms, like "exchange clustering" (Brown et al 1992), which assigns each word type to a single class.

Discriminative language models

- Or we could just train a model to predict $p(w_m|w_{m-1}, w_{m-2}, \ldots)$ directly.

- We might be able to use arbitrary features of the history to model long-range dependencies.

- Algorithms such as perceptron and logistic regression have been considered (Rosenfeld, 1996; Roark et al., 2007)

- Currently, "neural probabilistic language models" are attracting a lot of interest. The log-bilinear model (Mnih and Hinton, 2008) looks like this:

$$p_\theta^h(w) = \frac{\exp(s_\theta(w, h))}{\sum_{w'} \exp(s_\theta(w', h))}$$
$$s_\theta(w, h) = \hat{\boldsymbol{q}}_h^\mathsf{T} \boldsymbol{q}_w + b_w,$$

where h is the history context, $\hat{\boldsymbol{q}}_h$ is a latent description of the history, $\boldsymbol{q}_w$ is a latent description of the word, and b_w is an offset. The history context can be computed from the words themselves, as $\hat{\boldsymbol{q}}_h = \sum_i^{m-1} C_i \boldsymbol{q}_i$, where the matrix C_i is applied to context position i. All parameters can be estimated to directly maximize the probability of a corpus, using gradient ascent.

- Recent work has focused on efficiently training such models, with increasingly convincing results on large training sets (Mikolov et al., 2011).

Chapter 9

Morphology

[1] So far we have been focusing on NLP at the word level. Today we go **inside of words**.

We've already hinted at a morphological problem by introducing the idea of **lemmas**, where *serve/served/serving* all have the lemma *serve*.

From the perspective of document classification, these multiple forms may just seem like an annoyance, which we can get rid of by lemmatization or stemming (more on this later).

But morphology conveys information which might be crucial for some applications:

- Information retrieval

    - With a query like *bagel*, we want to get hits for *bagels*.
    - Same for *corpus/corpora*, *goose/geese*.
    - But we don't always want all the inflected forms. For example, a query for *Apple* may not want hits for *apples*

- Time. Morphology often indicates when events happen. For example, in French:

J'achete un velo	I buy a bicycle (now)
J'acheterai un velo	I will buy a bicycle
J'achetais un velo	I was buying a bicycle
J'ai acheté un velo	I bought a bicycle
J'acheterais un velo	I would buy a bicycle

---

[1]This chapter is pretty rough; better to see Chapter 2 of (Bender, 2013).

- Causality. Consider the difference between the Spanish examples:

*Si tu vas a GT, tu **seras** rica*	If you go to GT, you will be rich
*Si tu vas a GT, tu **eres** rica.*	If you go to GT, you are rich

- Lexical semantics: suppose *antichrist* is not in your sentiment dictionary. Do you think it is a positive or negative word?

In addition to recognizing morphology, there are applications in which we need to produce it.

- Translation: *you (pl) are smart → Ustedes son inteligentes* vs *Tu eres inteligente*

- Text generation: (has-property you-pl smart) → *ustedes son inteligentes*

Morphology, Orthography, and Phonology

- **Morphology** describes how meaning is constructed from combining affixes. For example, it is a morphological fact of English that adding the affix *+S* to many nouns creates a plural.

$$berry + \text{PLURAL} \rightarrow berry\text{+}s$$

- **Orthography** specifically relates to writing.
  For example,

$$berry\text{+}s \rightarrow berries$$

  is an orthographic rule. We have lots of these in English, which is one reason English spelling is difficult.

- Morphological rules also include stem changes, such as *goose* + PLURAL → *geese*.

- **Phonology** describes how sounds combine. For example, the different pronunciations of the final *s* in *cats* (s) and *dogs* (z) follow from a phonological rule (example (25) in the Bender text, page 30).

- In English, morphologically distinct words may be pronounced differently even when they are spelled the same, and this can reflect morphological differences. *read*+PRESENT vs. *read*+PAST.

- Conversely, morphological variants may be spelled differently even when they sound the same, like *The Champions' league* vs *The Champion's league* vs *The Champions league*.

Productivity

One idea for dealing with morphology is to build a morphologically aware dictionary:

- Map each **surface form** to its underlying **lemma**

- Include meaning of morphology: tense, number, animacy, posession, etc.

- Then when you encounter a surface form, just look it up.

duck	*duck*/N+SG
ducks	*duck*/N+PL
duck	*duck*/VB+PRESENT
ducks	*duck*/VBZ+PRESENT

Will this work? Besides the problem of ambiguity, still another problem is that morphology is **productive**, meaning that it applies to new words. If you only know the words *Google* or *iPad*, you can immediately understand their inflected forms.

- Have you Googled that yet?

- I have owned three iPads.

Derivational morphology (more on this later) is productive in another way: you can produce new words by applying morphological changes to existing words. hyper+un+desire+able+ity

In some languages, derivational morphology can create extremely complicated words. The J&M textbook has a fun example from Turkish:

In the homework, you'll see examples from Swahili, which also has complex morphology. A dictionary of all possible surface forms in such languages would be gargantuan. So instead of building a static dictionary, we will try to model the underlying morphological and orthographic rules.

9.1 Morphemes

Two broad classes: **stems** and **affixes**.

- Intuitively, stems are the "main" part of meaning, affixes are the modifiers

A Turkish word

uygarlaştıramadıklarımızdanmışsınızcasına
uygar_laş_tır_ama_dık_lar_ımız_dan_mış_sınız_casına

"as if you are among those whom we were not able to civilize (=cause to become civilized)"

uygar: *civilized*
_laş: *become*
_tır: *cause somebody to do something*
_ama: *not able*
_dık: past participle
_lar: plural
_ımız: 1st person plural possessive (our)
_dan: *among* (ablative case)
_mış: past
_sınız: 2nd person plural (you)   *K. Oflazer pc to J&M*

Figure 9.1: From (Jurafsky and Martin, 2009)

- Typically, **stems** can appear on their own (they are **free**) and affixes cannot (they are **bound**).

- Types of affixes:

  - **Prefixes**: *un+learn*, *pre+view*.

    * These examples are derivational. English has few inflectional prefixes, but other languages have many.
    * For example, in Swahili: *u-na-kata* versus *u-me-kata* distinguishes *you are cutting* versus *you have cut*. *na* and *me* are prefixes, *kata* is the root.
  - **Suffixes**: *I learn+ed, She learn+s, three apple+s, four fox+es*.
  - **Circumfixes** go around the stem.
    * None in English.
    * German has a circumfix for the past participle: *sagen* (say) → *ge+sag+t* (said)
    * French negation can be seen as a circumfix: *Je mange+NEG→Je ne mange pas*. (I do not eat).

(c) Jacob Eisenstein 2014-2015. Work in progress.

∗ More generally, morphemes can be non-continuous, as in the Hebrew example (7) in the Bender reading (page 12).

(7)

Root	Pattern	Part of Speech	Phonological Form	Orthographic Form	Gloss
ktb	CaCaC	(v)	katav	כתב	'wrote'
ktb	hiCCiC	(v)	hixtiv	הכתיב	'dictated'
ktb	miCCaC	(n)	mixtav	מכתב	'a letter'
ktb	CCaC	(n)	ktav	כתב	'writing, alphabet'

[heb]

In this example, the root *ktb* (related to writing) is combined with patterns that indicate where to insert vowels to produce different parts-of-speech and meanings.

– **Infixes** go inside the stem.

* Tagalog: *hingi+AGENT→h+um+ingi*
* Lakhota: *m'ani* (he walks), *ma-wá-ni* (I walk). The *wá* marks agreement with a first-person singular subject; it is an infix for this word, although it is a prefix in other words.
* English: *absolutely+fucking→absofuckinglutely*, but **absfuckingsolutely* arguably doesn't work.

9.2 Types of morphology

• **Inflection** creates different forms of a single word:

– tense: *to be, being, I am, you are, he is, I was*

– number: *book, books*

– case: *he, his, her, they, them, their*

• **Derivation** creates new words:
grace → disgrace → disgraceful → disgracefully

• **Cliticization** combines *Georgia+'s* into *Georgia's*; the possessive clitic *'s* is syntactically independent but phonologically dependent.

– Pronouns appear as clitics in French, e.g., *j'accuse* (I accuse), as does negation *Je n'accuse personne* (I don't accuse anyone).

– Another example is from Hebrew: *l'shana tova* (literally for year good, meaning happy new year); the preposition *for* appears as a clitic.

- **Compounding** combines two words in a new word:
  cream → ice cream → ice cream cone → ice cream cone bakery

- **Portmanteaus** combine words, truncating one or both.
  smoke + fog → smog
  glass + asshole → glasshole

- Word formation is *productive*: new words are subject to all of these processes

Inflectional morphology

Inflectional morphology adds information about words. English has a very simple system of inflectional morphology, compared to many languages.

Affix	Syntactic/semantic effect	Examples
-s	NUMBER: plural	*cats*
-'s	possessive	*cat's*
-s	TENSE: present, SUBJ: 3sg	*jumps*
-ed	TENSE: past	*jumped*
-ed/-en	ASPECT: perfective	*eaten*
-ing	ASPECT: progressive	*jumping*
-er	comparative	*smaller*
-est	superlative	*smallest*

Figure 9.2: From (Bender, 2013)

- English nouns are marked for number and possession; many language also mark nouns for **case**, which is the syntactic role that the noun plays in the sentence.

– In English, we do distinguish the case of some pronouns:

  ∗ *He (*NOMINATIVE*) gave her (*OBLIQUE*) his (*GENITIVE*) guitar.*
  ∗ *She gave him her guitar.*
  ∗ *I gave you our guitar.*
  ∗ *You gave me your guitar.*

  Specifically, we distinguish the **nominative** case of personal pronouns (except for 2nd-person), and the **genitive** case; all other uses are the **oblique** case.

– Other languages – such as Latin, Russian, Sanskrit, and Tamil, mark the case of all nouns. These languages have additional cases, such as dative (indirect object), accusative (direct object), and vocative (address).

– In German, noun is not inflected for case, but the articles and adjectives are:

  ∗ *Der alte Mann gab dem kleinen Affen die grosse Banane* (based on example 49 from Bender)
  ∗ The old man (NOM) gave the little monkey (DATIVE) the big banana (ACCUSATIVE)
  ∗ Notice how *der*, *dem*, and *die* all mean the, but carry the case marking.
  ∗ The adjectives are also marked for case.

– Many languages – such as Romance languages – mark the gender and number of nouns by inflecting the article and adjective. e.g., Spanish:

  ∗ *El coche rojo pasó la luz roja*: the red car ran the red light
  ∗ *Los coches rojos pasó las luces rojas*: the red cars ran the red lights
  ∗ Note that the article and adjective must **agree** for the sentence to be grammatical.
  ∗ In English, demonstrative determiners mark number, *this book* vs *these books*. In English, the determiner and noun must agree, e.g. **this books*.

– Gender is not necessarily binary.

  ∗ English pronouns include neuter *it*; German, Sanskrit, and Latin do this for all nouns.
  ∗ Danish and Dutch distinguish **neuter** from **common** gender
  ∗ Other languages distinguish **animate** and **inanimate**

- – Number if not necessarily binary.

  - ∗ Many languages, such as Arabic and Sanskrit, include a special **dual** number for two. English has residual traces of the dual number, with *both* vs *all* and *either* vs *any*.

  - ∗ Some Austronesian languages have a **trial** number, for groups of three.

  - ∗ Some languages, including Arabic, have a **paucal** number, for small groups.

- English verbs are inflected for tense and number distinguishing past (*I ate*), present (*I eat*), and 3rd-person singular (*She eats*). They are also inflected for aspect, distinguishing perfective (*I had eaten*) and progressive (*I am eating*). Note that the perfective and the past tense are identical for regular verbs, e.g. *we had talked*, *we talked*.

  - – Many languages (e.g., Chinese and Indonesian), do not mark tense with morphology. Indonesian uses time signals.

Saya makan apel	I eat an apple
Saya sedang makan apel	I am eating an apple
Saya telah makan apel	I already ate an apple
Saya akan makan apel	I will eat an apple

  - – Romance languages distinguish many more tenses than English with morphology.

    - ∗ Spanish has multiple past tenses: **preterite** and **imperfect**, distinguishing, e.g. *I ate onions yesterday* from *I ate onions every day*. These are distinguished by morphology: *comí cebollas ayer*, *comía cebollas cada día*.

    - ∗ Spanish and French have endings for conditional (*comería cebollas*) and future (*comeré cebollas*)

    - ∗ All of these are marked with time signals in English; future can also be marked this way in French and Spanish, e.g. *voy a comer cebollas*.

  - – Romance also have separate forms for every combination of number and person.

    - ∗ *(yo hablo / tu hablas / ella habla / nosotros hablamos / vosotros hableis / ellas hablan)*

    - ∗ *(je parle / tu parles / elle parle / nous parlons / vous parlez / ils parlent)*

∗ In Spanish, they they eliminate pronouns (pro-drop) in cases where the morphology makes it clear (unless they want to add emphasis). Chinese is also a pro-drop language?

∗ This doesn't happen in French, maybe because many different spellings (*parle/parles/parlent*) sound the same.

∗ In English, we only distiniguish 3rd-person singular.

- Adjectives in English mark comparative and superlative (*taller*, *tallest*). As we have seen, they can mark gender and number in other languages.

- Other things can be marked with affixes, such as **evidentiality** – how the speaker came to know the information. In Eastern Pomo (a California language), there are verb suffixes for four evidential categories:

-ink'e	nonvisual sensory
-ine	inferential
-le	hearsay
-ya	direct knowledge

Example (41) from Bender (2013) shows evidentiary marking in Turkish, *Ahmet geldi* (Ahmet came, witnessed by the speaker) vs *Ahmet gelmiş* (not witnessed by the speaker)

The **index of synthesis** measures the ratio of the number of morphemes in a given text to the number of words. Languages with complex morphology are called **synthetic**; languages with simple morphology are called **isolating** or **analytic**. English is relatively, but not extremely, analytic.

An approximation of the index of synthesis is the type-token ratio. Can you see why? If you count the number of unique surface forms in 10K *parallel* sentences from Europarl, you get:

- English: 16k word types

- French: 22k

- German: 32k

- Finnish: 55k

Language	Index of synthesis
Vietnamese	1.06
Yoruba	1.09
English	1.68
Old English	2.12
Swahili	2.55
Turkish	2.86
Russian	3.33
Inuit (Eskimo)	3.72

Figure 9.3: From (Bender, 2013)

Derivational Morphology

Derivational morphology is a way to create new words and change part-of-speech.

- **nominalization**

    - *V + -ation: computerization*
    - *V + -er: walker*
    - *Adj + -ness: fussiness*
    - *Adj + -ity: obesity*

- **negation**: *undo, unseen, misnomer*

- **adjectivization**: *V + -able : doable, thinkable, N + -al : tonal, national, N + -ous: famous, glamorous*

- **abverbization**: *ADJ + -ily: clumsily*

- **lots more**: *rewrite, phallocentrism, ...*

You can create totally new words this way.
word → wordify → wordification → wordificationism → antiwordificiationism → hyperantiwordificationism

Irregularities

English morphology contains a lot of irregularities: *know/knew/known*, *foot/feet*, *go/went*.. if you're not a native speaker, learning these was probably a pain in the neck.

- the good news is, there are fewer of these all the time! for example, the past tense of *show* used to be *shew*, like *know/knew* (the past participle is still *shown*).

- the bad news is, the most common words will be the last to change (if ever).

Attaching affixes can cause orthographic and phonological changes:

- walk + ed = walked, but frame+ed = framed, emit+ed = emitted, easy + ier = easier

- this is usually due to phonetic or orthographic *constraints*

- *optimality theory* is an approach to systematizing such interacting constraints. There's a lot of research on finite state models of optimality theory, but you'll have to take a linguistics course for that Karttunen and Beesley (2005).

Chapter 10

Finite-state automata

Finite-state automata are a powerful formalism for representing a subset of formal languages, the **regular** languages. As we will see, this formalism can also be used as a building block for an incredibly wide range of methods for manipulating natural language too (Mohri et al., 2002). This chapter will especially focus on **morphology**, which concerns how words are built out of smaller units. For a good reference on morphology for natural language processing, see (Bender, 2013).

Knight and May (2009) show how finite-state automata can be composed together to create impressive applications. They start with one such application — transliteration — and explain how it works. Here, we'll build the formalism from the ground up, starting with finite-state acceptors, then adding weights, and then adding transduction, finally arriving at the same sorts of applications.

10.1   Automata and languages

Basics of the formalism   :

- An alphabet Σ is a set of symbols

- A string ω is a sequence of symbols.
  The empty string ϵ contains zero symbols.

- A language $L \subseteq \Sigma^*$ is a set of strings.

An automaton is an abstract model of a computer which reads an input string, and either accepts or rejects it.

Chomsky Hierarchy   Every automaton defines a language. Different automata define different classes of languages. The Chomsky Hierarchy:

- Finite-state automata define **regular** languages

- Pushdown automata define **context-free** languages

- Turing machines define **recursively-enumerable** languages

Finite-state automata   A finite-state automaton $M = \langle Q, \Sigma, q_0, F, \delta \rangle$ consists of:

- A finite set of states $Q = \{q_0, q_1, \ldots, q_n\}$

- A finite alphabet Σ of input symbols

- A start state $q_0 \in Q$

- A set of final states $F \subseteq Q$

- A transition function δ

Determinism

- In a deterministic (D)FSA, $\delta : Q \times \Sigma \rightarrow Q$.

- In a nondeterministic (N)FSA, $\delta : Q \times \Sigma \rightarrow 2^Q$

- We can determinize any NFSA using the powerset construction, but the number of states in the resulting DFSA may be 2^n.

- Any **regular expression** can be converted into an NFSA, and thus into a DFSA.

The English Dictionary as an FSA   We can build a simple "chain" FSA which accepts any single word. So, we can define the English dictionary with an FSA. However, we can make this FSA much more compact. (see slides)

- Begin by taking the **union** of all of the chain FSAs by defining epsilon transitions (that is, transitions which do not consume an input symbol) from the start state to chain FSAs for each word (5303 states / 5302 arcs using a 850 word dictionary of "basic English")

- Eliminate the epsilon transitions by pushing the first letter to the front (4454 states / 4453 arcs)

- **Determinize** (2609 / 2608)

- **Minimize** (744 / 1535). The cost of minimizing an acyclic FSA is $O(E)$. This data structure is called a trie.

Operations We've now talked about three operations: union, determinization and minimization. Other important operations are:

intersection : only accept strings in both FSAs

negation only accept strings not accepted by FSA M

concatenation . accept strings of the form $s = [s_1 s_2]$, where $s_1 \in M_1$ and $s_2 \in M_2$

FSAs are closed under all these operations, meaning that resulting automaton is still an FSA (and therefore still defines a regular language).

10.2 FSAs for Morphology

Now for some morphology. Suppose that we want to write a program that accepts words that could **possibly** be constructed in accordance with English derivational morphology, but none of the impossible ones:

- *grace, graceful, gracefully*

- *disgrace, disgraceful, disgracefully, ...*

- *Google,Googler,Googleology,...*

- **gracelyful, *disungracefully, ...*

We could just make a list, and then take the union of the list using ϵ-transitions.
The list would get very long, and it would not account for productivity (our ability to make new words like *antiwordificationist*). So let's try to use finite state machines instead. Our FSA will have to encode rules about morpheme ordering, called *morphotactics*.
Let's start with some examples:

- *grace*: $q_0 \rightarrow_{\text{stem}} q_1$

- *dis-grace*: $q_0 \rightarrow_{\text{prefix}} q_1 \rightarrow_{\text{stem}} q_2$

- *grace-ful*: $q_0 \rightarrow_{\text{stem}} q_1 \rightarrow_{\text{suffix}} q_2$

- *dis-grace-ful*: $q_0 \rightarrow_{\text{prefix}} q_1 \rightarrow_{\text{stem}} q_2 \rightarrow_{\text{suffix}} q_3$

Can we generalize these examples?

Figure 10.1: I can't find the attribution for this figure right now, sorry! I think it's either from Julia Hockenmaier's slides, or from Jurafsky and Martin (2009).

- This example abstracts away important details, like why *wordificate* is preferred to **wordifycate*. But this rule is part of English **orthography** (spelling), not **morphology**. "Two-level morphology" is an approach to integrating such orthographic transformations in a finite-state framework (Karttunen and Beesley, 2001).

- It also misses a key point: sometimes we have choices, and not all choices are considered to be equally good by fluent speakers.

    – Google counts:
        * *superfast*: 70M; *ultrafast*: 16M; *hyperfast*: 350K; *megafast*: 87K
        * *suckitude*: 426K; *suckiness*: 378K
        * *nonobvious*: 1.1M; *unobvious*: 826K; *disobvious*: 5K

- Rather than asking whether a word is **acceptable**, we might like to ask how acceptable it is.

- But finite state acceptors gives us no way to express *preferences* among technically valid choices.

- We'll need to augment the formalism for this.

10.3 Weighted Finite State Automata

A weighted finite-state automaton $M = \langle Q, \Sigma, \pi, \xi, \delta \rangle$ consists of:

- A finite set of states $Q = \{q_0, q_1, \ldots, q_n\}$

- A finite alphabet Σ of input symbols

- Initial weight function, $\pi : Q \to \mathbb{R}$

- Final weight function $\xi : Q \to \mathbb{R}$

- A transition function $\delta : Q \times \Sigma \times Q \to \mathbb{R}$

We have added a weight function that scores every possible transition.

- We can score any path through the WFSA by the sum of the weights.

- Arcs that we don't draw have infinite cost.

- The shortest-path algorithm can find the minimum-cost path for accepting a given string in $O(V \log V + E)$.

Applications of WFSAs

We can use WFSAs to score derivational morphology as suggested above. But let's start with a simpler example:

Edit distance . We can build an edit distance machine for any word. Here's one way to do this (there are others):

- Charge 0 for "correct" symbols and rightward moves

- Charge 1 for self-transitions (insertions)

Figure 10.2: From (Knight and May, 2009)

- Charge 1 for rightward epsilon transitions (deletions)

- Charge 2 for "incorrect" symbols and rightward moves (substitutions)

- Charge ∞ for everything else

The total edit distance is the *sum* of costs across the best path through machine.

Probabilistic models   For probabilistic models, we make the path costs equal to the likelihood:

$$\delta(q_1, s, q_2) = \mathrm{p}(s, q_2 \mid q_1) \tag{10.1}$$

This enables probabilistic models, such as N-gram language models.

- A unigram language model is just one state, with V edges.

- A bigram language model will have V states, with V^2 edges.

Knight and May (2009) show how to do an interpolated bigram/unigram language model using a WFSA. (Last year I wrote a note that I had found a better way, with only $V + 3$ states rather than $2V + 4$. But now I can't find my solution!)

- Recall that an interpolated bigram language model is

$$\hat{\mathrm{p}}(v|u) = \lambda \mathrm{p}_2(v|u) + (1 - \lambda)\mathrm{p}_1(v), \tag{10.2}$$

with p̂ indicating the interpolated probability, p_2 indicating the bigram probability, and p_1 indicating the unigram probability.

- Unlike the basic n-gram language models, our interpolated model has non-determinism: do we choose the bigram context or the unigram context?

- What should happen to the scores as we encounter a non-deterministic choice?

- For a sequence *a,b,a*, we want the final path score to be

$$\psi(a,b,a) = (\lambda p_2(a|*) + (1-\lambda)p_1(a))$$
$$\times (\lambda p_2(b|a) + (1-\lambda)p_1(a))$$
$$\times (\lambda p_2(b|a) + (1-\lambda)p_1(b))$$

- So we could multiply along each step, and add probabilities across non-deterministic choices.

- With log-probabilities, we would add along each step, and use the log sum, $\log(e^a + e^b)$, to compute the score for non-deterministic branchings.

10.4 Semirings

We have now seen three examples: an acceptor for derivational morphology, and weighted acceptors for edit distance and language modeling. Several things are different across these examples.

- Scoring

  - In the derivational morphology FSA, we wanted a boolean "score": is the input a valid word or not?

  - In the edit distance WFSA, we wanted a numerical (integer) score, with lower being better.

  - In the interpolated language model, we wanted a numerical (real) score, with higher being better.

- Nondeterminism

  - In the derivational morphology FSA, we accept if there is any path to a terminating state.

  – In the edit distance WFSA, we want the score of the single best path.

  – In the interpolated language model, we want to sum over non-deterministic choices.

• How can we combine all of these possibilities into a single formalism? The answer is semiring notation.

Formal definition

A semiring is a system $(\mathbb{K}, \oplus, \otimes, \overline{0}, \overline{1})$

• $\mathbb{K}$ is the set of possible values, e.g. $\{\mathbb{R}_+ \cup \infty\}$, the non-negative reals union with infinity

• $\oplus$ is an addition operator

• $\otimes$ is a multiplication operator

• $\overline{0}$ is the additive identity

• $\overline{1}$ is the multiplicative identity

A semiring must meet the following requirements:

• $(a \oplus b) \oplus c = a \oplus (b \oplus c)$, $(\overline{0} \oplus a) = a$, $a \oplus b = b \oplus a$

• $(a \otimes b) \otimes c = a \otimes (b \otimes c)$, $a \otimes \overline{1} = \overline{1} \otimes a = a$

• $a \otimes (b \oplus c) = (a \otimes b) \oplus (a \otimes c)$, $(a \oplus b) \otimes c = (a \otimes c) \oplus (b \otimes c)$

• $a \otimes \overline{0} = 0 \otimes \overline{a} = \overline{0}$

Semirings of interest   :
  where $\oplus_{\log}(a, b)$ is defined as $\log(e^a + e^b)$.
  Semirings allow us to compute a more general notion of the "shortest path" for a WFSA.

• Our initial score is $\overline{1}$

• When we take a step, we use $\otimes$ to combine the score for the step with the running total.

• When nondeterminism lets us take multiple possible steps, we combine their scores using $\oplus$.

Name	$\mathbb{K}$	$\oplus$	$\otimes$	$\bar{0}$	$\bar{1}$	Applications
Boolean	$\{0,1\}$	$\vee$	$\wedge$	0	1	identical to an unweighted FSA
Probability	$\mathbb{R}_+$	$+$	$\times$	0	1	sum of probabilities of all paths
Log-probability	$\mathbb{R} \cup -\infty \cup \infty$	$\oplus_{\log}$	$+$	$-\infty$	0	log marginal probability
Tropical	$\mathbb{R} \cup -\infty \cup \infty$	min	$+$	∞	0	best single path

Example Let's see how this works out for our language model example.

$$
\begin{aligned}
score(\{a,b,a\}) = \bar{1} \otimes &\left(\lambda \otimes \mathrm{p}_2(a|*) \oplus (1-\lambda) \otimes \mathrm{p}_1(a)\right) \\
\otimes &\left(\lambda \otimes \mathrm{p}_2(b|a) \oplus (1-\lambda) \otimes \mathrm{p}_1(b)\right) \\
\otimes &\left(\lambda \otimes \mathrm{p}_2(a|b) \oplus (1-\lambda) \otimes \mathrm{p}_1(a)\right)
\end{aligned}
$$

Now if we plug in the **probability semiring**, we get

$$
\begin{aligned}
score(\{a,b,a\}) = 1 \times &\left(\lambda \mathrm{p}_2(a|*) + (1-\lambda)\mathrm{p}_1(a)\right) \\
\times &\left(\lambda \mathrm{p}_2(b|a) + (1-\lambda)\mathrm{p}_1(b)\right) \\
\times &\left(\lambda \mathrm{p}_2(a|b) + (1-\lambda)\mathrm{p}_1(a)\right)
\end{aligned}
$$

But if we plug in the **log probability semiring**, we need the edge weights to be equal to $\log \mathrm{p}_1$, $\log \mathrm{p}_2$, $\log \lambda$, and $\log(1-\lambda)$. Then we get:

$$
\begin{aligned}
score(\{a,b,a\}) = {}& 0 + \log\left(\exp(\log\lambda + \log\mathrm{p}_2(a|*)) + \exp(\log(1-\lambda) + \log\mathrm{p}_1(a))\right) \\
&+ \log\left(\exp(\log\lambda + \log\mathrm{p}_2(b|a)) + \exp(\log(1-\lambda) + \log\mathrm{p}_1(b))\right) \\
&+ \log\left(\exp(\log\lambda + \log\mathrm{p}_2(a|b)) + \exp(\log(1-\lambda) + \log\mathrm{p}_1(a))\right) \\
= {}& 0 + \log(\lambda\mathrm{p}_2(a|*) + (1-\lambda)\mathrm{p}_1(a)) \\
&+ \log(\lambda\mathrm{p}_2(b|a) + (1-\lambda)\mathrm{p}_1(b)) \\
&+ \log(\lambda\mathrm{p}_2(a|b) + (1-\lambda)\mathrm{p}_1(a)),
\end{aligned}
$$

which is exactly equal to the log of the score from the probability semiring.

- The score on any specific path will be the semiring **product** of all steps along the path.

- The score of any input will be the semiring **sum** of the scores of all paths that successfully process the input.

- What happens if we use the tropical semiring?

10.5   Finite state transducers

FSAs and WFSAs apply to single strings. FSTs and WFSTs apply to pairs of string.

FSTs define **regular relations** over pairs of strings. We can think of them in a few different ways:

- **Recognizer**: accepts string pairs iff they are in the relation

- **Translator**: reads an input, produces an output

Formally, a finite-state transducer $M = \langle Q, \Sigma, \Delta, q_0, F, \delta, \sigma \rangle$ consists of:

- A finite set of states $Q = \{q_0, q_1, \ldots, q_n\}$

- Finite alphabets Σ for input symbols and Δ for output symbols

- Initial state $q_0 \in Q$ and final states $F \subseteq Q$

- A state transition function $\delta : \langle Q \times \Sigma^* \rangle \to 2^Q$

- A string transition function $\sigma : \langle Q \times \Sigma^* \rangle \to 2^{\Delta^*}$

Unlike NFSAs, not all NFSTs can be determinized. However, special subsets of NFSTs called **subsequential** transducers can be determinized efficiently (see 3.4.1 in Jurafsky and Martin (2009)).

We can build some simple NLP systems directly from FSTs.

- A zeroth-order translation system could be made from a single state and a set of self-transitions: $Q_0 \xrightarrow[el]{the} Q_0, Q_0 \xrightarrow[los]{the} Q_0, Q_0 \xrightarrow[libro]{book} Q_0, Q_0 \xrightarrow[libros]{books} Q_0,\ldots$

- First-order translation would require a state per word in the "input" vocabulary: $Q_0 \xrightarrow[\epsilon]{the} Q_{the} \xrightarrow[los\ libros]{books} Q_0, Q_{the} \xrightarrow[el\ libro]{book} Q_0.$

- Inflectional morphology and orthography:

$$Q_0 \xrightarrow[wit]{wit} Q_{regular} \xrightarrow[+s]{+PL} Q_1$$
$$Q_0 \xrightarrow[wish]{wish} Q_{needs\text{-}e} \xrightarrow[+es]{+PL} Q_1$$

10.6 Weighted FSTs

Weights can be added to FSTs in much the same way as they are added to FSAs.

- For any pair $\langle q \in Q, s \in \Sigma^* \rangle$, we have a set of possible transitions, $\langle q \in Q, t \in \Delta^*, \omega \in \mathbb{K} \rangle$, with a weight ω in the domain defined by the semiring.

- For example, we could augment the translation transducers defined above to allow alternative possible translations for a single word.

- The same semiring operations in WFSAs apply here too.

	acceptor	transducer
unweighted	FSA: $\Sigma^* \to \{0, 1\}$	WFSA: $\Sigma^* \to \mathbb{R}$
weighted	FST: $\Sigma^* \to \Sigma^*$	WFST: $\Sigma^* \to \langle \Sigma^*, \mathbb{R} \rangle$

Example : General edit distance computer.

- $Q_0 \xrightarrow[a]{a} Q_0 : 0$

- $Q_0 \xrightarrow[\epsilon]{a} Q_0 : 1$

- $Q_0 \xrightarrow[a]{\epsilon} Q_0 : 1$

The shortest path for a pair of strings $\langle s, t \rangle$ in this transducer has a score equal to the minimum edit distance between the strings (in the tropical semiring).

We can think of each path as defining a potential **alignment** between s and t.

Operations on FSTs

- Closed under **union**

- Closed under **inversion**, which switches input and output labels.

- Closed under **projection**, because FSAs are a special case of FSTs

- Not closed under **difference**, **complementation**, and **intersection**.

- Closed under **composition**.

FST composition is the basis for implementing the noisy channel model in FSTs, and can be used to support dozens of cool applications.

Finite state composition

Suppose we have a transducer T_1 from language I_1 to O_1, and another transducer T_2 from O_1 to O_2. The composition $T_1 \circ T_2$ is an FST from I_1 to O_2.

- Unweighted definition: iff $\langle x, z \rangle \in T_1$ and $\langle z, y \rangle \in T_2$, then $\langle x, y \rangle \in T_1 \circ T_2$.

- Weighted definition:

$$(T_1 \circ T_2)(x, y) = \bigoplus_{z \in \Sigma^*} T_1(x, z) \otimes T_2(z, y) \tag{10.3}$$

- Designing algorithms for automatic FST composition is relatively straight-forward if there are no epsilon transitions; otherwise it's more challenging (Allauzen et al., 2009).

The simplest example

- $T_1 : Q_0 \xrightarrow[a]{x} Q_0, Q_0 \xrightarrow[b]{y} Q_0$

- $T_2 : Q_1 \xrightarrow{a} Q_1, Q_1 \xrightarrow{b} Q_2, Q_2 \xrightarrow{b} Q_2$

- $T_1 \circ T_2 : Q_1 \xrightarrow{x} Q_1, Q_1 \xrightarrow{y} Q_2, Q_2 \xrightarrow{y} Q_2$

For simplicity T_2 is written as a finite-state acceptor, not a transducer. Acceptors are a special case of transducers.

If we had weights, they would be combined through $\otimes$.

10.7   Applications of composition

Edit distance

Consider the general edit distance computer developed in section 10.6. It assigns scores to pairs of strings. If we compose it with an FSA for a given string (e.g., *tech*), we get a WFSA, who assigns score equal to the minimum edit distance from *tech* for the input string.

- Composing an FST with a FSA yields a FSA.

- A very useful design pattern is to build a **decoding** WFSA by composing a general-purpose WFST with an unweighted FSA representing the input.

- The best path through the resulting WFSA will be the minimum cost / maximum likelihood decoding.

Transliteration

English is written in a Roman script, but many languages are not. **Transliteration** is the problem of converting strings between scripts. It is especially important for names, which don't have agreed-upon translations.

A simple transliteration system can be implemented through the noisy-channel model.

- T_1 is an English character model, implemented as a transducer so that strings are scored as $\log p_r(c_1, c_2, \ldots, c_M)$.

- T_2 is a character-to-character transliteration model. This can be based on explicit rules,[1] or on conditional probabilities $\log p_t(c^{(f)}|c^{(r)})$.

- T_3 is an acceptor for a given string that is to be transliterated.

The machine $T_1 \circ T_2 \circ T_3$ scores English character strings based on their orthographic fluency (T_1) and adequacy (T_2).

Suppose you were given an Roman-script character model and a set of foreign-script strings, but no equivalent Roman-script strings. How would you use EM to learn a transliteration model?

Knight and May (2009) provide a more complex transliteration model, which transliterates between Roman and Katakana scripts, using a deep cascade that includes models of the underlying phonology. In their model,

Word-based translation

As we have seen, simple models of machine translation can be implemented as finite-state transducers.

- Recall the example: *the books / los libros*. Here we are interested in modeling English-to-Spanish translation, $P(S|E)$.

---

[1] http://en.wikipedia.org/wiki/Romanization_of_Russian

- Earlier we proposed a multi-state translation model to deal with the impact of pluralization on the Spanish article *los*

- If we build our translator by composing a Spanish language model $P(S)$ with an Spanish-to-English transducer $P(E|S)$, this is not necessary. This is an example of the **noisy channel model**.

Here are the specifics:

- $T1$ is a language model, implemented as a transducer, where every path inputs and outputs the same string, with a score equal to $\log p(w_1, w_2, \ldots, w_M)$. This model's responsibility is to tell us that $p(el\ libros) \ll p(los\ libros)$.

  - For example, a bigram language model would be implemented with V states, with a score of $\log P_S(W_{n+1} = i | W_n = j)$ for the edge from the state representing word j to the state representing word i.
  - We are free to use higher-order language models. Would we need a trigram language model to correctly identify *los grandes libros* as the right translation of *the great books*?

- $T2$ is a word-to-word translation model.

  - It could be a bilingual dictionary, with edge weights taking the value 0. In this case, the resulting translations will be scored only on boolean acceptability in the translation model, and on fluency according to the language model.
  - It could a probabilistic translation model, with each edge having a score equal to $\log p(w^{(e)}|w^{(s)})$, the log probability of the English word $w^{(e)}$ given the Spanish word $w^{(s)}$.

- $T3$ is an acceptor for a given foreign string $w^{(f)}$

So the translation model is $T_1 \circ T_2 \circ T_3$, which produces a WFSA, in which path are scored according to the joint log probability $\log p_{e|s}(\boldsymbol{w}^{(s)}, \boldsymbol{w}^{(e)})$. Using semiring notation, here's what happens in our example:

$$p(los\ libros, the\ books) = p_S(los|\star) \otimes p_{E|S}(the|los) \otimes p_S(libros|los) \otimes p_{E|S}(books|libros)$$

The composed FST can thus overcome its simplistic model of translation by having a more intelligent language model. This is useful, because language models can be trained without labeled data, while translation models cannot.

Structure prediction Transliteration and translation are examples of **structure prediction**, which will be a dominant theme in the remainder this course.

- Our goal for English-to-Spanish translation is to predict a sequence $\boldsymbol{w}^{(s)}$, given linguistic input $\boldsymbol{w}^{(e)}$.

- The set of possible sequences is very large — in word-to-word translation, it is V^M. Therefore, we need to score these structures in a decomposable way.

- Finite-state composition give us a way to do that. We are implicitly decomposing the score for the tuple $\langle \boldsymbol{w}^{(e)}, \boldsymbol{w}^{(s)} \rangle$ into scores for adjacent words in $\boldsymbol{w}^{(s)}$ and aligned word pairs.

- This scoring function can be written as $\boldsymbol{\theta} \boldsymbol{f}(\boldsymbol{w}^{(s)}, \boldsymbol{w}^{(e)})$, where the FST defines a specific, decomposable feature function. More on this later.

Nondeterminism While T_2 might be non-deterministic, there is no non-determinacy about **paths**: a given pair $\langle \boldsymbol{w}^{(e)}, \boldsymbol{w}^{(f)} \rangle$ can only be transduced in one way.

Knight and May (2009) introduce a more complex translation model. They relax the assumption that there is a monotonic word-to-word alignment, allowing reorderings and multiword translations.

Word segmentation

Word segmentation is a challenging problem for speech, and also for languages like Chinese, which don't explicitly segment adjacent words in orthography. We can again use a finite-state approach.

- T_1 is again a language model. We can use a simplistic model representing a dictionary, or we can use a probabilistic model.

- T_2 is a sequence of unsegmented symbols.

Suppose you are given dictionary of permissible words. How would you use EM to learn the language model?

Stemming

As discussed on Tuesday, information retrieval systems that only return exact matches aren't very useful.

- Suppose you query: *is it medically safe to kiss my cat on the lips*

- You want to get a hit even for documents like: *on the medical safety of kissing cats*.

- In morphologically complex languages like Hebrew, these differences could cause early IR systems to miss more than 90% of relevant documents Choueka (1989)!

- Stemming improves the **recall** of information retrieval systems by converting all tokens of *kissing* to *kiss*, etc.

The **Porter Stemmer** is a very popular stemming program. It is written as a set of rules, which are applied in stages, e.g.

- if the word ends in *-ing* and the preceeding part contains a vowel, delete the ending: *hopping → hopp*

- if the word ends in *-ed*, and the preceeding part contains a vowel, delete the ending: *hopped → hopp*

- Next, if the remaining part contains double letters (besides ss, ll, or zz), remove one (*hopp → hop*)

We can think of these rules as a sequence of deterministic finite state transducers:

- We can build a transducer to strip off endings like *-ing*, using two states to indicate whether we have yet seen a vowel

- Next we can build a transducer to strip off double letter endings, with exceptions for s, l, and z.

- We can **compose** these transducers into a single machine.

Morphological analysis

Recall that when we talked about morphology, there were several types of interacting rules:

1. To pluralize regular words, add an *s*; but if the word ends in *sh*, you have to add *es*.

2. To conjugate 3rd-person singular, add an *s*; but if the word ends in *sh*, you have to add *es*.

Pluralization and conjugation are different morphological systems, but once they have decided to append an s, the subsequent orthographical rules are the same.

This suggests an intermediate representation:

- *dish+*PL $\rightarrow$ *dish$\wedge$s#* $\rightarrow$ *dishes*

- *fist+*PL $\rightarrow$ *fist$\wedge$s#* $\rightarrow$ *fists*

- *fish+*PL $\rightarrow$ *fish#* $\rightarrow$ *fish*

- *fish+*3S $\rightarrow$ *fish$\wedge$s#* $\rightarrow$ *fishes*

- Special symbols for morpheme and word boundaries

- The conjugation and pluralization FSTs only need to know what affix to add, and where to add it.

- Then an orthography FST takes over, and figures out how the morphemes should be combined.

- Finite-state composition allows us to automatically build a single machine for conjugation and pluralization, incorporating both the selection of affixes and orthographic constraints.

We have only described how to *generate* the English text. But if we compose this machine with a chain FSA representing an observed string, then we obtain an FSA where the set of accepted strings reveal the acceptable morphological analyses.

- For example, an utterance like *wishes* could have been produced by *wish/N+*3S or by *wish/V+*PL; both will be accepted by the resulting FSA.

- Suppose we want to distinguish the most likely morphological analysis given the context.

  - First, we add a loop back to the beginning in the morphological FST, so that we can accept multiple words.

- Now we can build something like a language model WFST, but here we need probabilities of morphological analyses rather than just words. We might want to decompose this like

$$P(stem, affixes|context) \approx P(stem|context)P(affix|context) \quad (10.4)$$

- If we have a morphologically annotated corpus, we can make smoothed relative frequency estimates of these probabilities. If we don't, what do we do?

- The morphological analyses are missing data, so we might be able to do EM:
    * **E-step**: Decode all observed string sequences in the corpus, given the current probabilities.
    * **M-step**: Treat decoding as ground truth, update probabilities
    * Note: this is "Viterbi" EM, (a.k.a. "Hard" EM), because we are not computing probabilities over all possible decodings. We could do that using something called the expectation semiring Eisner (2002).

Context-sensitive spelling correction

Swype text entry in my old android phone does not consider context, much to my annoyance.

- I mean: *Prepare lecture for my class*

- It says: *Prepare lecture foie my class*

That's not smart. The bigram probabilities $P(lecture\ foie)$ and $P(foie\ my)$ should be ridiculously low. Once again, we can apply the noisy channel model, starting with T_1 as a language model and T_2 as a spelling model. If T_1 is a bigram language model, then the resulting machine is an FST with one state per word type.

- The cost of the transition from Q_a to Q_b on input s is the semiring product $\otimes$ of two costs:

    - The transition cost from a to b in the language model
    - The "emission" cost from s to a in the edit distance machine. Depending on exactly how we're getting our input, we could memorize these costs (for all pairs of words) and work at the token level, rather than character level. But there's no conceptual difference; alternatively you could think about character-level edit distance FSTs sitting at each state.

- Given an input sequence s, we compose

  select the spelling-corrected string t with the greatest value,

$$\hat{t} = \max_t \bigoplus_\pi s \leadsto_\pi t, \tag{10.5}$$

  where π is a path over input s and output t.

  - Why are there multiple paths from s to t? In the edit distance machine, we can always model matching input/output symbols as an insertion and a deletion.
  - Whether we care about these alternative paths depends on the semiring. In the tropical semiring, the solution is simply the minimum-cost path.

Norvig's spelling corrector

We haven't said much about where the weights come from. The bigram language model is probabilistic, and we can compute $\delta(q_s, q_t, t, t)$ as $-\log P(t|s)$. What about the edit distance model?

Before we get into that, here's an alternative approach, from Peter Norvig `http://norvig.com/spell-correct.html` (highly recommended). It may help explain how google is able to quickly and accurately spell-check your queries.

The approach can't easily be framed in exactly terms of the FST/semiring formalism, but it's close:

- Check if the word itself is in the dictionary. If so, return it.

- Else, check if any edit-distance=1 corrections is in the dictionary. If so, return the one with the highest unigram count.

- Else try corrections with edit-distance = 2.

The essay discusses extensions to bigrams a probabilistic model of errors. One way to build such a model is to look at data.

- Norvig suggests the Birbeck spelling error corpus, `http://norvig.com/spell-correct.html`.

  - Such a resource would tell us the likelihood of a word s being mispelled as t.

– But new mispellings are always possible (even inevitable)

- An alternative would be to parametrize the edit distance model with more specific probabilities: P(insertion), P(deletion), etc.

  – Maximum-likelihood estimation would compute

$$P(\text{insertion}) = \frac{\text{count(insertion)}}{\text{count(all-characters)}} \tag{10.6}$$

  – But we will probably never get a corpus that gives us these counts.
  – Can we bootstrap our way to success? Suppose we could just guess the probabilities.
    * We could find the best path for each example in the training set, and keep track of the counts of insertions and deletions in these paths.

$$\hat{\pi}_{s,t} = \arg\min_{\pi} \text{cost}(s \rightsquigarrow_{\pi} t)$$

$$\text{count(insertion)} = \sum_{\langle s,t \rangle \in \mathcal{D}} \text{count}(\text{insertion}, \hat{\pi}_{s,t})$$

    * More probabilistically, we could compute the likelihood of each path, and use these likelihood to find the *expected* counts:

$$\text{count(insertion)} = \sum_{\langle s,t \rangle \in \mathcal{D}} \sum_{\pi} \frac{P(s \rightsquigarrow_{\pi} t)}{\sum_{\pi'} P(s \rightsquigarrow_{\pi'} t)} \text{count}(\text{insertion}, \pi_{s,t})$$

    * Once we have the counts, we can go back and re-estimate the insertion probability (and all the other probabilities).

- This type of bootstrapping is another example of **expectation-maximization**.

  – We introduced EM in the simpler case of clustering.
    * Each document had a distribution over clusters $q(z)$
    * Each cluster had parameters θ.
    * The E-step computed $q(z)$, the M-step optimized θ.
  – Here, the Q distribution is over paths $Q(\pi)$.
  – In the E-step, we can compute $Q(\pi)$ given estimates of the parameters $P(\text{insertion})$.

- In the M-step, we can update the parameters $P(\text{insertion})$ from expected counts under $Q(\pi)$.

- The number of paths can be very large, often too large to store.

  - In **Viterbi EM**, we just use the best path. This can work very well.

  - Alternatively, we can use the **expectation semiring** so that the "score" of the weighted path sum includes the expected counts Eisner (2001).

    Each edge is a tuple of a probability and a vector of expected counts, $\langle p_i, p_i v_i \rangle$

    $$\langle p_i, p_i v_i \rangle \otimes \langle p_j, p_j v_j \rangle = \langle p_i p_j, p_j p_i v_i + p_i p_j v_j \rangle = \langle p_k, p_k v_k \rangle$$
    $$\langle p_i, v_i \rangle \oplus \langle p_j, v_j \rangle = \langle p_i + p_j, v_i + v_j \rangle$$

  - If we can compute the semiring shortest path ($\mathcal{O}(n^3)$ at worst), we can compute the expected counts!

Speech Recognition

Speech recognition is yet another application of finite-state methods in NLP. It takes the idea of intermediate representations even further.

We want a mapping from intended words to observed acoustics. This can be seen as a multilevel process:

- G: WFST which generates English sentences: *I went to the bank*

- L: WFST which converts each word to context-indendent phonemes (WFST. Pronunciation dictionaries have this information): *Ay w eh n t t uw ...*

- C: WFST which converts context-independent phonemes to context-dependent phonemes: *Ay w eh n t uw ...*

- A: WFST which converts context-dependent phonemes to acoustic observations

By composing the FST $G \circ L \circ C \circ A$ with an FSA representing the observed acoustics O, we obtain a single WFSA which scores proposed English sentences for the observations O.

Inference   The composed FSA would be ridiculously huge. Beam pruning is a technique for pruning away paths which are extremely unlikely, resulting in a faster, smaller FST.

Estimation   It's really hard to get labeled data for all of these levels, for example context-dependent phones. We can treat this as a hidden variable and estimate it using EM.

Software   There are mature software toolkits for working with finite state machines. OpenFST is a C++ package which I have had some experience with; it's fast and relatively well-documented. XFST and Carmel are other options.

10.8   Recap

We saw how finite-state composition can create powerful NLP systems out of simple, modular components.

- For example, we can compose a translation model (one state!) and a bigram language model (V states) to create a finite-state translation machine.

- If we compose the translation machine with a **chain FSA** representing the "input", we get a WFSA whose shortest path is the best translation of the input.

- We didn't talk about algorithms for composition, but the formal definition is

$$(T_1 \circ T_2)(x, y) = \bigoplus_z T_1(x, z) \otimes T_2(z, y). \tag{10.7}$$

  In other words, to compute the score of $\langle x, y \rangle$ in the composed machine, we take the semiring addition over all strings z, for which we compute the extension of the score of $\langle x, z \rangle$ in T_1 with the score of $\langle z, y \rangle$ in T_2.

- For example, the language model only has edges for $\langle s, s \rangle$, with score $\log P(s_i | s_{i-1})$. The translation model has edges for all $\langle s, t \rangle$, with score $\log P(t|s)$. So the composed machine must have edges for all $\langle s, t \rangle$, with score $\log P(t|s) \otimes \log P(s_i | s_{i-1})$.

- In the chain FSA, each edge takes exactly one string t_i, with score $\overline{1}$. So the composed machine is a WFSA, with edges for each possible s_i, each having score $\log P(t_i | s_i) \otimes \log P(s_i | s_{i-1}) \otimes \overline{1}$. This structure is known as a **trellis**.

Chapter 11

Part-of-speech tagging

Words can be grouped into rough classes based on syntax.

- Why is *colorless green ideas sleep furiously* more acceptable than *ideas colorless furiously green sleep*?

- Why is *teacher strikes idle children* ambiguous?

In both examples, word classes can provide an explanation.

- Word classes have strong ordering constraints:

    - J J N V R is likely
    - N J R J V is unlikely (why?)

- Ambiguity about word class leads to very different interpretations:

    - N N V N
    - N V J N (ouch!)

So clearly we have intuitions about a few parts-of-speech already: noun, verb, adjective, adverb. Jurafsky and Martin (2009) describe these as the four major **open** word classes, although apparently not all languages have all of them.

What other word classes are there?

- The Penn Treebank defined a set of 45 POS tags.[1]

---

[1] http://www.comp.leeds.ac.uk/ccalas/tagsets/upenn.html

- The Brown corpus defined a set of 87 POS tags.[2]

- Petrov et al. (2012) define a "universal" set of 12 tags.

Which is right?

Example   Let's look at some data.

> My name is Ozymandias, king of kings:
> Look on my works, ye Mighty, and despair!

The part-of-speech tags for this couplet from Ozymmandias are shown in Table 11.1.

- All tagsets distinguish basic categories like nouns, pronouns, verbs, adjectives, and punctuation.

- The Brown tagset includes a number of fine-grained distinctions:

    - specific tags for the *be*, *do*, and *have* verbs, which the other two tagsets just lump in with other verbs.

    - distinct tags for possessive determiners (*my name*) and possessive pronouns (*mine*)

    - distinct tags for the third-person singular pronouns (e.g., *it, he*) and other pronouns (e.g., *they, we, I*)

- The Universal tagset aggressively groups categories that are distinguished in the other tagsets:

    - all nouns are grouped, ignoring number and the proper/common distinction (see below)

    - all verbs are grouped, ignoring inflection

    - preposition and postpositions are grouped as adpositions

    - all punctuation is grouped

    - coordinating and subordinating conjunctions (e.g. *and* versus *that*) are grouped

---

[2]`http://www.comp.leeds.ac.uk/ccalas/tagsets/brown.html`

	Brown	PTB	Universal
My	possessive determiner (DD$)	possessive pronoun (PRP$)	pronoun (PRON)
name	noun, singular, common (NN)	NN	NOUN
is	verb "to be" 3rd person, singular (BEZ)	verb 3rd person, singular (VBZ)	VERB
Ozymandias	proper noun, singular (NP)	proper noun, singular (NNP)	NOUN
,	comma (,)	comma (,)	punctuation (.)
king	NN	NN	NOUN
of	preposition (IN)	preposition (IN)	adposition (ADP)
kings	noun, plural, common (NNS)	NNS	NOUN
:	colon (:)	mid-sentence punc (:)	.
Look	verb, base: uninflected present, imperative, or infinite (VB)	VB	VERB
on	IN	IN	ADP
my	DD$	PRP$	PRON
works	NNS	NNS	NOUN
ye	personal pronoun, nominative, non 3S (PPSS)	personal pronoun, nominative (PRP)	PRON
mighty	adjective (JJ)	JJ	adjective (ADJ)
,	comma (,)	comma (,)	punctuation (.)
and	coordinating conjunction (CC)	CC	conjunction (CONJ)
despair	VB	VB	VERB

Table 11.1: Part-of-speech annotations from three tagsets.

So which is "right"? It depends. The Brown tags can be useful for certain applications, and they may have strong tag-to-tag relations that make tagging

easier (see next chapter). But they are more expensive to annotate. The Universal tags are intended to generalize across many types of text, and should be easier to annotate.


Figure 11.1: needs attribution

11.1   Details about parts-of-speech

As usual, Bender (2013) provides a useful linguistic perspective.

- **Nouns** describe entities and concepts

  - **Proper nouns** name specific people and entities: *Georgia Tech*, *Janet*, *Buddhism*. In English, they're usually capitalized. PTB tags: NNP (singular), NNPS (plural)

  - **Common nouns** cover everything else. In English, they're often preceded by articles, e.g. *the book*, *the university*. Common nouns decompose into

    * **Count nouns** have a plural and need an article in the singular, *dogs*, *the dog*.

* **Mass nouns** don't have a plural and don't an article in the singular, *snow is cold, gas is expensive*

  – **Pronouns** refer to specific noun phrases or entities or events.

    * **Personal pronouns** refer to people or entities: *you, she, I, it, me*. PTB tag: PRP
    * **Possessive pronouns** are pronouns that indicate possession: *your, her, my, its, one's, our*. PTB tag: PRP$
    * **Wh-pronouns** are used in question forms (*Where are you going?*, WP) and as relative pronouns in forms like (*The girl who played with fire*)

  Unlike other nouns, the set of possible pronouns cannot be expanded! It is a **closed class**.

* **Verbs** describe activities, processes, and events, e.g. *eat, write, sleep*

  – The Penn Treebank differentiates verbs by morphology: VB (infinitive), VBD (past), VBG (present participle), VBN (past participle), VBP (present, non 3rd person-singular), VBZ (present 3rd person singular)

  – **modals** are a closed subclasses of verbs, such as (*should, can, will, must*). They get PTB tag MD

  – **copula** is *be* with a predicate, e.g. *she is hungry*. The Brown Tagset distinguishes copula, but PTB doesn't.

  – **auxiliary** verbs include *be, have, will* to form complex tenses, e.g. *we will have done it twice*.

    * Also includes *do* in questions and negation, e.g. *did you eat yet?*. Apparently this is from Welsh, which was spoken in England before the Anglo-Saxons invaded; *do* doesn't function this way in German.
    * The Brown corpus has special tags for HAVE and DO, but the PTB doesn't.

* **Adjectives** describe properties of entities: *antique, vast, trunkless*

  – **Attributive use**: *an antique land*

  – **Predicative use**: *the land was antique*

  – **Gradable adjectives** (*big*) have **comparative form** (*bigger*) and **superlative form** (*biggest*)

- – Can you think of an adjective that is not gradable?

- – With *big*, we can move to comparative form by adding the suffix *-est*. This is an example of agglutinative morphology. Can you think of an adjective in English where the relationship is not agglutinative? (rather, it's fusional). How about *good*, *better*, *best*?

- – PTB tags: JJ, JJR, JJS

- **Adverbs** describe properties of events.

  - – Manner: *slowly, slower, fast, hesitantly*

  - – Degree: *extremely, very, highly*

  - – Directional and locative: *here, downstairs, near*

  - – Temporal: *yesterday, Monday*

  - – Besides verbs, adverbs may also modify sentences, adjectives, or other adverbs. Apparently, the very ill man walks extremely slowly

  - – Adverbs may also be gradable. Tags: RB, RBR, RBS

- **Prepositions** are a closed class of words that can come before noun phrases, forming a prepositional phrase that relates the noun phrase to something else in the sentence.

  - – *I eat sushi WITH soy sauce*

  - – *I eat sushi WITH chopsticks*

  - – *To* gets its own tag TO, because it forms the infinitive with bare form verbs (VB), e.g. *I want to eat*

  - – Everything else is tagged IN in the PTB.

- **Coordinating conjunctions** join two elements,

  - – *vast and trunkless legs*

  - – *She eats burgers but she drinks soda*

  - – PTB tag: CC

- **Subordinating conjunctions** introduce a subordinate clause, e.g. *She thinks THAT Chomsky is wrong*. PTB tag: annotIN

(c) Jacob Eisenstein 2014-2015. Work in progress.

- **Particles** are oddball words that come with verbs and can change their meaning to a new **phrasal verb**, e.g., *come ON*, *he brushed himself OFF*, *let's check OUT that new restaurant*. They are a closed class, PTB tag RP.

- **Determiners** are a closed class of words that precede noun phrases.

  – Articles: *the, an, a*

  – Demonstratives: *this, these, that*

  – Quantifiers: *some, every, few*

  – Wh-determiners, *WHICH burger should I choose?*

  – PTB tag: DT

- **Oddballs**

  – **Existential there**, e.g. *There is no way out of here*, gets its own tag, EX.

  – So does the possessive ending *'s*, POS

  – So do numbers (CD), list items (LS), commas, and other non-alphabetic symbols.

11.2 Part of speech tagging

Part of speech tags relate to a number of other linguistic phenomena:

- Lexical semantics: *can*/V vs *can*/N, *teacher strikes children*, etc

- Pronunciation: *inSULT* vs *INsult*, *conTENT* vs *CONtent*

- Translation: *park*/v $\rightarrow$ *garer*, *park*/N $\rightarrow$ *parque*

- NP chunking: `grep {JJ | NN}* {NN | NNS}`

POS tagging is a useful preprocessing step for downstream applications.
So how can we build an automatic POS tagger?

- Observation 1: it's "easy."

  – 60% of word types have only one possible POS tag (in English).
  – If you choose the majority POS tag for each token, you get 90% right.

- Observation 2: it's not easy: a few words have a lot of possible POS tags

- *We're taking it **back**/RB*

- *The shirt off my **back**/NN*

- *Go **back**/RP where you belong*

- *If you challenge him, I'll **back**/VBP you*

- *The **back**/JJ roads are safer*

- Observation 3: 90% is not actually very good. $0.9^{10} \approx .3$, so you will only get 30% of ten-word sentences correct. Sentences have exponentially many possible POS sequences:

VBD		VB			
VBN	VBZ	VBP	VBZ		
NNP	NNS	NN	NNS	CD	NN
fed	*raises*	*interest*	*rates*	*0.5*	*percent*

Anyway, let's look at a tougher poem, Jabberwocky:

'Twas brillig, and the slithy toves
Did gyre and gimble in the wabe:
All mimsy were the borogoves,
And the mome raths outgrabe

Forget *twas*. What about *slithy*? Can you guess the POS? What about *toves*? You don't know these words. What information are you using to guess?

- Word identity: you do know that *and* is CC and *the* is DET

- **Context**

  - JJ NN is likely
  - NN JJ is unlikely

- **Morphology**

  - *-s* → noun or verb
  - *-able* → adjective (98% of the time!)
  - *-ly* → adverb
  - *un-* → adjective or verb

– (not rules, just hints)

Let's put morphology on hold for a minute.

- Suppose we have an annotated corpus, with tagged sentences, $\langle \boldsymbol{w}_{1:N_t}, \boldsymbol{y}_{1:N_t} \rangle_{1:T}$.

- Then we could estimate the likelihood of a word given a tag,

$$P(w|y) = \frac{\text{count}(w, y)}{\text{count}(y)} \tag{11.1}$$

As always, smoothing is possible...

- Given this same annotated corpus, we could also compute $P(y_n|y_{n-1})$, a sort of language model over tags.

$$P(y_n|y_{n-1}) = \frac{\text{count}(y_{n-1}, y_n)}{\text{count}(y_{n-1})} \tag{11.2}$$

- Let's combine these ideas via a **generative story**

  – For word n, draw tag $y_n \sim \text{Categorical}(\theta_{y_{n-1}})$
  – Then draw word $w_n \sim \text{Categorical}(\phi_{y_n})$

We've built a generative model that explains our observations $\boldsymbol{w}$ through a bigram generative model over the tags.

- Under this model, we can compute

$$P(\boldsymbol{y}|\boldsymbol{w}) \propto P(\boldsymbol{w}, \boldsymbol{y}) \tag{11.3}$$

$$P(\boldsymbol{w}|\boldsymbol{y})P(\boldsymbol{y}) \tag{11.4}$$

$$\prod_n^N P(w_n|y_n)P(y_n|y_{n-1}) \tag{11.5}$$

- This is a **hidden Markov model**. It's Markov because the probability of y_n depends only on y_{n-1} and not any of the previous history. It's hidden because y_n is unknown when we decode.

- We can treat this as a special case of finite state transduction. Can you see how?

Chapter 12

Sequence Labeling

In sequence labeling, we want to assign tags to words. There are many applications:

- Part-of-speech tagging: *Go*/V *to*/P *Georgia*/N *Tech*/N *next*/J *year*/N ./.

- Named entity recognition: *Go*/O *to*/O *Georgia*/B-ORG *Tech*/I-ORG *next*/B-DATE *year*/I-DATE ./O

- Phrase chunking: *Go*/B-VP *to*/B-PP *Georgia*/B-NP *Tech*/I-NP *next*/B-NP *year*/I-NP ./O

In classification, we would choose each tag independently:

$$p(y_m|w_m) \perp p(y_n|w_n), \forall m \neq n \tag{12.1}$$

In sequence labeling, we choose the sequence of tags **jointly**. Probabilistically, we might try to choose $\hat{y} = \arg\max_{y \in \mathcal{T}^M} p(y|w)$. As we will see later, we can also write this in the form of a linear predictor:

$$\hat{y} = \arg\max_{y \in \mathcal{T}^M} \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}, \boldsymbol{y}) \tag{12.2}$$

In either case, we have an immediate problem: finding the best scoring tag sequence in the set $\mathcal{T}^M$. As the notation suggests, the number of possible tag sequences is exponential in the length of the sequence. Consider part-of-speech tagging:

VBD		VB			
VBN	VBZ	VBP	VBZ		
NNP	NNS	NN	NNS	CD	NN
fed	*raises*	*interest*	*rates*	*0.5*	*percent*

Even after using a tag dictionary to restrict the set of possible tags for each word, there are thirty-six possible tag sequences for this six-word sentence. This means we will need clever algorithms to compute $\arg\max_{\boldsymbol{y}\in\mathcal{T}^M}$. We cannot enumerate all possibilities.

12.1   Hidden Markov Models

Let's first think about tagging as a probabilistic model. Specifically, we want to maximize $\mathrm{p}(\boldsymbol{y}|\boldsymbol{w}) \propto \mathrm{p}(\boldsymbol{y}, \boldsymbol{w})$, where $\boldsymbol{w}$ are words and $\boldsymbol{y}$ are tags. This is equivalent to Naive Bayes, but for sequence labeling.

As in Naive Bayes, we define the probability distribution $\mathrm{p}(\boldsymbol{w}, \boldsymbol{y})$ through a *generative story*,

- For word m, draw tag $y_m \sim \mathrm{Categorical}(\lambda_{y_{m-1}})$

- Then draw word $w_m \sim \mathrm{Categorical}(\phi_{y_m})$

Under this model, we can compute

$$\mathrm{p}(\boldsymbol{y} \mid \boldsymbol{w}) \propto \mathrm{p}(\boldsymbol{w}, \boldsymbol{y}) \tag{12.3}$$

$$\mathrm{p}_e(\boldsymbol{w} \mid \boldsymbol{y}; \phi)\mathrm{p}_t(\boldsymbol{y}; \lambda) \tag{12.4}$$

$$\prod_m^M \mathrm{p}_e(w_m \mid y_m; \phi)\mathrm{p}_t(y_m \mid y_{m-1}; \lambda) \tag{12.5}$$

- This is a **hidden Markov model**. It's Markov because the probability of y_m depends only on y_{m-1} and not any of the previous history. It's hidden because y_m is unknown when we decode.

  - The probability $\mathrm{p}_e(w_m \mid y_m; \phi)$ is the **emission probability**, since the words are treated as emissions from the tags.

  - The probability $\mathrm{p}_t(y_m \mid y_{m-1}; \lambda)$ is the **transition probability**, since it assigns probability to each possible tag-to-tag transition.

- We can describe this generative story as a graphical model. Note that although graphical models and finite-state models both use circles and arrows, the meaning is completely different.

- Our generative story assumes that the words are conditionally independent, given the tags, so
$$w_n \perp \{\boldsymbol{w}_{m \neq n}\} \mid y_n.$$
Conditional independence is not the same as independence. We do **not** have $p(w_n, w_m) = p(w_n)p(w_m)$ because the tags are related to each other. For example, suppose that (a) nouns always follow determiners, (b) *the* is always a determiner and (c) *bike* is always a noun. Then

$$
\begin{aligned}
P(W_m = \textit{the}, W_{m+1} = \textit{bike}) &= \sum_{y_{m+1}, y_m} P(W_m = \textit{the}, W_{m+1} = \textit{bike}, y_{m+1}, y_m) \\
&= \sum_{y_{m+1}, y_m} P(W_{m+1} = \textit{bike} \mid y_{m+1}, y_m, W_m = \textit{the}) \\
&\quad \times P(y_{m+1} \mid y_m, W_m = \textit{the}) P(y_m \mid W_m = \textit{the}) P(W_m = \textit{the}) \\
&= \sum_{y_{m+1}} P(W_{m+1} = \textit{bike} \mid y_{m+1}) \\
&\quad \times \sum_{y_m} P(y_{m+1} \mid y_m) P(y_m \mid W_m = \textit{the}) P(W_m = \textit{the}) \\
&= P(W_{m+1} = \textit{bike} \mid y_{m+1} = \text{NOUN}) \times 1 \times 1 \times P(W_m = \textit{the}) \\
&> P(W_{m+1} = \textit{bike}) P(W_m = \textit{the})
\end{aligned}
$$

- Another way to think about independence is that if we are told one tag, it affects all of our other tagging decisions.

  - For example, in the sentence *teacher strikes idle children*, we might choose tag sequence NN VBZ JJ NNS.
  - But if are given $y_3 = $ VBP, then suddenly $y_2 = $ VBZ looks like a bad choice because $p_T(\text{VBZ}, \text{VBP})$ is very small.
  - So we might now choose $y_2 = $ NNS.
  - This change might cascade back to y_1, etc (not in this case, but it could happen in theory)

- A classifier-based tagger, which treated the tags as IID, might ignore these dependencies, and produce a tag sequence that contained unlikely transitions like VBZ, VBP.

- A better alternative might be to tag the text from left-to-right; we could then condition on the previous tag, choosing

$$y_m = \arg\max_y \mathrm{p}_e(w_m|y_m)\mathrm{p}_t(y_m|y_{m-1}) \tag{12.6}$$

But this approach is "greedy," and can mistakenly commit to bad tagging decisions. For example, in *teacher strikes strand children*, we might initially choose $y_2 = $ VBZ, because this is more common than the noun sense of *strikes*. However, we are then stuck, because *strand* has low probability as anything but a verb, yet the verb-verb transition also has low probability. The greedy tagger is unable to recover the globally optimal sequence, NN NNS VBP NNS, without backtracking.

- This is why we need **joint inference** over $\boldsymbol{y}_{1:M}$ to find $\hat{\boldsymbol{y}} = \arg\max_{\boldsymbol{y}} \mathrm{p}(\boldsymbol{w}, \boldsymbol{y})$. The key challenge is to search over the exponential number of tag sequences efficiently.

12.2   Sequence labeling as finite-state transduction

To see whether efficient joint inference is possible, we first formulate the problem in terms of finite-state transduction.

- Transducer E has one state, and transduces from tags to words, with $\delta_{w,t}^{(e)} = \mathrm{p}_e(w|y)$.

- Transducer T has $\#|\mathcal{T}|$ states (if it's a bigram model), and transduces tags to tags, with $\delta_{y_m,y_{m-1}}^{(t)} = \mathrm{p}_t(y_m|y_{m-1})$.

- Recall the definition of finite state composition,

$$(T \circ E)(y, x) = \bigoplus_z T(y, z) \otimes E(z, x). \tag{12.7}$$

Since T only accepts identical tag pairs $\langle y, y \rangle$, we can ignore $\bigoplus$; there's only one possible $z = y$. The result of $T \circ E$ transduces tags to words, with edge

weights

$$
\begin{aligned}
\delta^{(t \circ e)}_{w, y_m, y_{m-1}} &= \delta^{(e)}_{w, y_m} \otimes \delta^{(t)}_{y_{m-1}, y_m} \\
&= \mathrm{p}(w \mid y_m) \otimes \mathrm{p}(y_m \mid y_{m-1}) \\
&= \mathrm{p}(w \mid y_m)\mathrm{p}(y_m \mid y_{m-1}) \\
&= \mathrm{p}(w, y_m \mid y_{m-1})
\end{aligned}
$$

Suppose we wanted to work with log probabilities instead. Then

$$
\begin{aligned}
\delta_{w,t} &= \log \mathrm{p}(x \mid y) \\
\delta_{y_m, y_{m-1}} &= \log \mathrm{p}(y_m \mid y_{m-1}) \\
a \otimes b &:= a + b \\
\delta_{x, y_m, y_{m-1}} &= \log \mathrm{p}(x \mid y_m) \otimes \log \mathrm{p}(y_m \mid y_{m-1}) \\
&= \log \mathrm{p}(x \mid y_m) + \log \mathrm{p}(y_m \mid y_{m-1}) \\
&= \log \mathrm{p}(x, y_m \mid y_{m-1})
\end{aligned}
$$

Can you see how many states the resulting FST will have?

- Finally, we compose with an acceptor S, which forms a chain for a sentence $w_1, \ldots w_M$.

- This composition $T \circ E \circ S$ yields a **trellis**-shaped weighted finite state acceptor (WFSA).

  - Number of columns = M, length of input.
  - Number of rows = T, number of tags.
  - Edges from states $\langle m, t_1 \rangle$ to $\langle m+1, t_2 \rangle$ have the score

$$
\delta^{(t \circ e)}_{w_{m+1}, t_2, t_1} = \delta^{(t)}_{t_2, t_1} \otimes \delta^{(e)}_{t_2, w_{m+1}} = P(Y_{m+1} = t_2 \mid Y_m = t_1) P(W_{m+1} = w_{m+1} \mid Y_{m+1} = t_2)
\tag{12.8}
$$

  - Each path in the trellis corresponds to a unique sequence of tags, $\boldsymbol{y}_{1:M}$, and every sequence of tags has a unique path. The score of the path is equal to $\mathrm{p}(\boldsymbol{w}_{1:M}, \boldsymbol{y}_{1:M})$ by construction.
  - If we define $\bigoplus = \max$ (as in the tropical semiring), then the score of the semiring shortest path is equal to $\max_{\boldsymbol{y}} \mathrm{p}(\boldsymbol{w}_{1:M}, \boldsymbol{y}_{1:M})$.

- So, can we find this score (and therefore the best path) in polynomial time?

- How expensive is it to construct the trellis?

  * Generic composition is polynomial but slower than we would like
    — it depends on the vocabulary size.
  * But since we know what the trellis is supposed to look like, we can
    just build it directly. This requires constant time per edge.
  * How big is the trellis? $\mathcal{O}(MT)$ states, $\mathcal{O}(MT^2)$ edges.

- How expensive is it find the shortest path in the trellis?:
  Generic shorest path has a time cost of $\mathcal{O}(V \log V + E) = \mathcal{O}(MT \log MT +$
  $MT^2)$ and a space cost of $\mathcal{O}(V) = \mathcal{O}(MT^2)$.

- So:

  * Building the trellis is polynomial.
  * Shortest path is polynomial.
  * Therefore, there must be a poly-time algorithm to find the best tag
    sequence, despite the apparently exponential number of paths.

12.3   The Viterbi algorithm

The Viterbi algorithm is a special-purpose best-path algorithm for FSTs in the
shape of a trellis. It has a time cost of $\mathcal{O}(MT^2)$ and a space cost of $\mathcal{O}(MT)$. (This
time cost improvement is important, because it is linear in the length of the se-
quence M, unlike the generic shortest-path algorithm, which is $M \log M$.

Based on the Markov assumption, we can decompose the likelihood recur-
sively.

$$\mathrm{p}(\boldsymbol{w}_{1:M}, \boldsymbol{y}_{1:M}) = \mathrm{p}(w_M \mid y_M)\mathrm{p}(y_M \mid y_{M-1})\mathrm{p}(\boldsymbol{w}_{1:M-1}, \boldsymbol{y}_{1:M-1})$$

- Given y_{m-1}, we can choose y_m without considering any other element of the
  history.

- Suppose we know the best path to $y_m = k$.
  The best path to $y_{m+1} = k'$ *through* $y_m = k$ must include the best path to
  $y_m = k$.

- Suppose we know the score (probability) of the best path to each $y_m = k$,
  which we write $v_m(k) = \max_{y_1 \dots y_{m-1}} \mathrm{p}(\boldsymbol{w}_{1:m}, \boldsymbol{y}_{1:m-1}, y_m = k)$.

What is the score of the best path to $y_{m+1} = k'$?

$$v_{m+1}(k') = \max_{\boldsymbol{y}_{1:m}} \mathrm{p}(\boldsymbol{w}_{1:m+1}, \boldsymbol{y}_{1:m}, y_{m+1} = k') \tag{12.9}$$

$$= \mathrm{p}_e(w_{m+1} \mid y_{m+1} = k') \max_{\boldsymbol{y}_{1:m}} P_t(Y_{m+1} = k' \mid y_m) \mathrm{p}(\boldsymbol{w}_{1:m}, \boldsymbol{y}_{1:m}) \tag{12.10}$$

$$= \mathrm{p}_e(w_{m+1} \mid y_{m+1} = k') \max_{y_m=k} P_t(Y_{m+1} = k' \mid Y_m = k) \max_{\boldsymbol{y}_{1:m-1}} \mathrm{p}(\boldsymbol{w}_{1:m}, \boldsymbol{y}_{1:m-1}, y_m = k) \tag{12.11}$$

$$= \mathrm{p}_e(w_{m+1} \mid y_{m+1} = k') \max_{y_m=k} P_t(Y_{m+1} = k' \mid Y_m = k) v_m(k) \tag{12.12}$$

The base case is $v_0(\lozenge) = 1$, with zero probability for everything else.

- We can generalize this recurrence using semiring notation:

$$v_{m+1}(k') = \delta^{(e)}_{w_{m+1}, y_{m+1}=k'} \otimes \bigoplus_k \delta^{(t)}_{k \to k'} \otimes v_m(k) \tag{12.13}$$

Then if we want to move to log-probabilities, we have

$$v_{m+1}(k') = \log \mathrm{p}_E(w_{m+1} \mid y_{m+1} = k') \otimes \bigoplus_k \log \mathrm{p}_T(k \to k') \otimes v_m(k) \tag{12.14}$$

$$= \log \mathrm{p}_E(w_{m+1} \mid y_{m+1} = k') + \max_k \log \mathrm{p}_T(k \to k') + v_m(k) \tag{12.15}$$

- We will frequently use a semiring in which the edge weights are log probabilities and $\otimes$ is addition. This is partly because addition is notationally clearer than multiplication, and because in practical settings, you will use the log probabilities to avoid underflow.

- Note that we are setting $\oplus = \max$, as in the tropical semiring. This means that the score of the best tag sequence overall is $v_M(\square)$.

- To find the best tag sequence, we just need to keep back-pointers, from $v_m(k)$ to $v_{m-1}(k')$:

$$v_{m+1}(k') = \max_k \log p_E(w_{m+1} \mid Y_{m+1} = k') + \log P_T(Y_{m+1} = k' \mid Y_m = k) + v_m(k)$$
$$\tag{12.16}$$

$$= \log p_E(w_{m+1} \mid y_{m+1} = k') + \left(\max_k \log P_T(Y_{m+1} = k' \mid Y_m = k) + v_m(k) \right)$$
$$\tag{12.17}$$

$$b_{m+1}(k') = \arg\max_k \log p_E(w_{m+1} \mid Y_{m+1} = k') + \log P_T(Y_{m+1} = k' \mid Y_m = k) + v_m(k)$$
$$\tag{12.18}$$

$$= \arg\max_k \log P_T(Y_{m+1} = k' \mid Y_m = k) + v_m(k) \tag{12.19}$$

Note that the computation of the back-pointer doesn't depend on the emission probability $p_E(w_{m+1} \mid Y_{m+1} = k')$, since Y_m is conditionally independent from w_{m+1} given Y_{m+1}.

- In the probability semiring, we had $\oplus$ as addition; in the log-probability semiring, it was log addition. What happens if we try these addition operators? We'll see in a moment.

Example

Table 12.1: $\log p_e(w \mid y)$

	they	*can*	*fish*
N	-2	-3	-3
V	-10	-1	-3

Table 12.2: $\log p_t(y_m \mid y_{m-1})$

	N	V	END
START	-1	-2	$-\infty$
N	-3	-1	-2
V	-1	-3	-2

See the slides for how the Viterbi algorithm works in this example.

12.4 The forward algorithm

In an influential survey on HMMs, Rabiner (1989) defines three problems:

- **Decoding**: find the best tags $\boldsymbol{y}$ for a sequence $\boldsymbol{w}$.

- **Likelihood**: compute the probability $\mathrm{p}(\boldsymbol{w}) = \sum_{\boldsymbol{y}} \mathrm{p}(\boldsymbol{w}, \boldsymbol{y})$

- **Learning**: given only unlabeled data $\{\boldsymbol{w}_1, \boldsymbol{w}_2, \ldots, \boldsymbol{w}_D\}$, estimate the transition and emission distributions.

The Viterbi algorithm solves the decoding problem. We'll talk about the learning problem later. Let's now consider how to compute the likelihood $\mathrm{p}(\boldsymbol{w}) = \sum_{\boldsymbol{y}} \mathrm{p}(\boldsymbol{w}, \boldsymbol{y})$.

- First, we move to a semiring where $a \oplus b = \log(e^a + e^b)$ instead of $\max$. Then

$$\alpha_{m+1}(k') = \bigoplus_k \log \mathrm{p}_E(w_{m+1} \mid Y_{m+1} = k') \otimes \log P_T(Y_{m+1} = k' \mid Y_m = k) \otimes \alpha_m(k)$$
$$(12.20)$$

$$= \log \mathrm{p}_E(w_{m+1} \mid Y_{m+1} = k') \otimes \bigoplus_k \log P_T(Y_{m+1} = k' \mid Y_m = k) \otimes \alpha_m(k)$$
$$(12.21)$$

$$= \log \mathrm{p}_E(w_{m+1} \mid Y_{m+1} = k') + \log \sum_k P_T(Y_{m+1} = k' \mid Y_m = k) \times e^{\alpha_m(k)}$$
$$(12.22)$$

$$= \log \mathrm{p}_E(w_{m+1} \mid Y_{m+1} = k') + \log \sum_k P_T(Y_{m+1} = k' \mid Y_m = k) \mathrm{p}(\boldsymbol{w}_{1:m}, Y_m = k)$$
$$(12.23)$$

$$= \log \mathrm{p}_E(w_{m+1} \mid Y_{m+1} = k') + \log P_T(Y_{m+1} = k', \boldsymbol{w}_{1:m}) \qquad (12.24)$$
$$= \log \mathrm{p}(\boldsymbol{w}_{1:m+1}, Y_{m+1} = k') \qquad (12.25)$$

- We used the inductive hypothesis in (12.23), and we used the HMM conditional independence assumptions $W_{m+1} \perp \boldsymbol{W}_{1:m} \mid Y_{m+1}$ and $Y_{m+1} \perp \boldsymbol{W}_{1:m} \mid Y_m$ in the following two steps.

- We can formalize this as an inductive proof by stating the base case,

$$\alpha_1(k) = \log \mathrm{p}_e(w_1 \mid y_1) \otimes \log P_t(Y_1 = k \mid \Diamond) = \log \mathrm{p}(w_1, Y_1 = k \mid Y_0 = \Diamond).$$
$$(12.26)$$

This is called the **forward** algorithm. The total probability of a sequence is $p(\boldsymbol{w}_{1:M}) = \alpha_M(\square)$. For a demo, see the slides.

Why solve the likelihood problem?

Why would we want to compute $p(\boldsymbol{w}_{1:M})$?

Word class language models

- Remember p(*colorless green ideas sleep furiously*)

- We don't care about the specific tags, we just want to know the probability of the utterance, so we can compare it with p(*Furiously sleep ideas green colorless*).

Comparing HMMs

- Suppose we have a few HMMs, each of which could have generated the observations.

- We want to compute the marginal likelihood of the observations given each HMM, regardless of the path.

- This approach is sometimes used in gesture recognition.

Computing marginals   : we'll soon be very interested in **marginal** probabilities $p(y_m \mid \boldsymbol{w}_{1:M})$, for all $m \leq M$. The likelihood $p(w_{1:M})$ is part of this computation.

12.5   HMM Details

Trigrams

- Can we use trigrams instead of bigrams for an HMM?

- How do we change the trellis? How big is the new trellis?

- Each cell represents a pair of tags, $\langle y_m, y_{m-1} \rangle$.

- The trellis still needs N columns, but now need T^2 rows.

- Each node can only connect to T neighbors in the next column, based on the trigram transition constraint. Number of edges = NT^3.

- We can prune very low probability edges for speed.

Estimation

In principle, we can use relative frequency estimation.

$$\lambda_{k,k'} \triangleq P_T(Y_m = k' \mid Y_{m-1} = k) = \frac{\text{count}(Y_m = k', Y_{m-1} = k)}{\text{count}(Y_{m-1} = k)}$$

$$\phi_{k,i} \triangleq P_E(W_m = i \mid Y_m = k) = \frac{\text{count}(W_m = i, Y_m = k)}{\text{count}(Y_m = k)}$$

In practice, we need smoothing and other tricks.

- The same smoothing ideas from language modeling can be applied.
  - interpolate between trigrams, bigrams, and unigrams

    $$P(y_m \mid y_{m-1}, y_{m-2}) = \lambda_2 \hat{P}_2(y_m \mid y_{m-1}, y_{m-2}) + \lambda_1 \hat{P}_1(y_m \mid y_{m-1}) + (1 - \lambda_2 - \lambda_1)\hat{P}_0(y_m)$$

  - reserve probability mass for unseen words

12.6 Tagging with features

Let's consider an example:

> 'Twas brillig, and the slithy toves
> Did gyre and gimble in the wabe:
> All mimsy were the borogoves,
> And the mome raths outgrabe

You probably don't know many of these words, yet it's not so hard to see what some of their tags should be. How do you do it?

Recall that the HMM can incorporate two sources of information:

- Word-tag probabilities, via $p_E(w_m|y_n)$.

- Local context, via $p_T(y_m|y_{m-1})$.

Local context is helpful, but the word-tag probabilities will be worthless for words like *brillig, slithy, toves, gyre*, etc.

But there are a lot of things that we're missing here. Crucially, unseen words have internal structure which can be exploited:

mimsy, 3.1415, Ke$ha

- **Orthography and morphology**. *Slithy toves* just kind of looks like JJ NNS, because of the apparent suffixes.

  - *-s* → NNS, VBZ
  - *-able* → JJ
  - *-ly* → RB
  - *un-* → JJ,RB,V
  - ...

  But morphological features are difficult to incorporate in a generative model, because they break the Naive Bayes assumption:

  $$\mathrm{p}(\textit{mimsy}, \textit{-sy} \mid \mathrm{JJ}) \neq \mathrm{p}(\textit{mimsy} \mid \mathrm{JJ})\mathrm{p}(\textit{-sy} \mid \mathrm{JJ})$$

- **Capitalization**. This is especially relevant for named entity recognition, e.g., *I bought an apple* and *I bought an Apple phone*.

- More advanced HMMs incorporate morphological, orthographic, and typographic features by creating a more complex $\mathrm{p}_E(w \mid y)$ emission probability. For example, the TNT Tagger took this approach, and is one of the best generative taggers (Brants, 2000).

In addition to word-internal features, we might want more fine-grained context. For example, in the PTB, *this* and *these* are both tagged DT. But *this* is likely to be followed by a singular noun NN, and *these* is likely to be followed by a plural noun NNS. So we might like to add word-context features to the probability $\mathrm{p}(y_m \mid y_{m-1}, w_{m-1})$.

How can we incorporate these overlapping features? The solution is to build sequence labeling models based on the perceptron and logistic regression classifiers. The first model is called **structured perceptron**, since the label space consists of structures rather than individual labels (Collins, 2002). The second model is called a **conditional random field (CRF)**, due to its relation to Markov random fields (Lafferty et al., 2001). In this model, we explicitly compute $\mathrm{p}(\boldsymbol{y} \mid \boldsymbol{w})$.

In addition to incorporating overlapping features, these models have another advantage: they are discriminative, directly maximizing the conditional probability $\mathrm{p}(\boldsymbol{y} \mid \boldsymbol{w})$, or minimizing the perceptron loss. As in standard classification, this criterion is more closely connected to the accuracy metrics that we usually care about.

12.7 Structured perceptron

Remember the perceptron update:

$$\hat{y} = \arg\max_{y} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}, y) \tag{12.27}$$

$$\boldsymbol{\theta}^{(t+1)} \leftarrow \boldsymbol{\theta}^{(t)} + \boldsymbol{f}(\boldsymbol{w}, y^*) - \boldsymbol{f}(\boldsymbol{w}, \hat{y}) \tag{12.28}$$

In sequence labeling, we have a **structured output** $\boldsymbol{y} \in \mathcal{T}(\boldsymbol{w})$. Can we still apply the perceptron rule?

$$\hat{\boldsymbol{y}} = \arg\max_{\boldsymbol{y} \in \mathcal{T}(\boldsymbol{w})} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}, \boldsymbol{y}) \tag{12.29}$$

$$\boldsymbol{\theta}^{(t+1)} \leftarrow \boldsymbol{\theta}^{(t)} + \boldsymbol{f}(\boldsymbol{w}, \boldsymbol{y}^*) - \boldsymbol{f}(\boldsymbol{w}, \hat{\boldsymbol{y}}) \tag{12.30}$$

This is called structured perceptron, because it learns to predict structured output $\boldsymbol{y}$. The problem is that $\arg\max_{y}$ must search over the entire set of structures $\mathcal{T}$. The set of permissible outputs depends on the input $\boldsymbol{w}$, and it is very large: $\mathcal{O}(K^M)$. What can we do?

Viterbi inference for structured perceptron

We will place a restriction on the scoring function $\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}, \boldsymbol{y})$, which allows us to apply the Viterbi algorithm to find $\arg\max_{\boldsymbol{y}} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}, \boldsymbol{y})$!

- Specifically, we require $\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}, \boldsymbol{y}) = \sum_{m} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}, y_m, y_{m-1}, m)$

- That is, the global score must be a **sum of local scores**.

- The local scores can consider any part of the observation, but must only consider adjacent elements in the label.

To apply Viterbi to structured perceptron, we have

$$v_m(k) = \bigoplus_{k'} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}_m(\boldsymbol{w}, k, k', m) \otimes v_{m-1}(k')$$

$$= \max_{k'} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}_m(\boldsymbol{w}, k, k', m) + v_{m-1}(k')$$

$$b_m(k) = \arg\max_{k'} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}_m(\boldsymbol{w}, k, k', m) + v_{m-1}(k')$$

Suppose we want to apply this to POS tagging? What features might we want? Here are some:

- Word-tag features, e.g. $\langle W : slithy, \text{JJ} \rangle$

- Adjacent tag-tag features, e.g. $\langle T : \text{JJ}, \text{NNS} \rangle$

- Suffix-tag features, e.g., $\langle M : \text{-}es, \text{NNS} \rangle$

- Previous-word features, e.g., $\langle P_1 : the, \text{JJ} \rangle$

- Next-word features, e.g., $\langle N_1 : slithy, \text{DT} \rangle$

- Note that we can consider arbitrarily distant words, e.g. $\langle Y_m, W_{m-15} \rangle$, because this still fits in the constraint, $\boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{w}, \boldsymbol{y}) = \sum_m \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{w}, y_m, y_{m-1}, m)$.

So to compute $v_m(k)$, we have to iterate over all $y_{m-1} = k'$,

- find the features $\boldsymbol{f}(\boldsymbol{x}_{1:M}, y_m = k, y_{m-1} = k', m)$,

- compute the inner product $\boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{x}_{1:M}, y_m = k, y_{m-1} = k', m)$,

- add it to $v_{m-1}(k')$

- take the max over all k'

This only works because of the assumption that the feature function decomposes over local parts of the sequence! If we wanted a feature that considered arbitrary parts of tag sequence, there would be no way to incorporate it into the recurrence relation.

Example

$$\boldsymbol{w} = \dots and\ the\ slithy\ toves$$
$$\boldsymbol{y} = \dots \text{CC DT JJ NNS}$$

Then we can characterize the tagging DT JJ NNS of the text the slithy toves (from Jabberwocky) in terms of the following features:

$\boldsymbol{f}(\text{the slithy toves, DT JJ NNS}) = \{ \langle W : \text{the}, \text{DT} \rangle, \langle M : \varnothing, \text{DT} \rangle, \langle T : \Diamond, \text{DT} \rangle$
$\langle W : \text{slithy}, \text{JJ} \rangle, \langle M : \text{-thy}, \text{JJ} \rangle, \langle T : \text{DT}, \text{JJ} \rangle$
$\langle W : \text{toves}, \text{NNS} \rangle, \langle M : \text{-es}, \text{NNS} \rangle, \langle T : \text{JJ}, \text{NNS} \rangle$
$\langle T : \text{NNS}, \square \rangle \}$

If this is the correct tagging, then we hope to learn a set of weights $\boldsymbol{w}$ such that $\boldsymbol{\theta}^\mathsf{T}\boldsymbol{f}$(the slithy toves, DT JJ NNS) is larger than the scores for other tag sequences, such as $\boldsymbol{\theta}^\mathsf{T}\boldsymbol{f}$(the slithy toves, DT NN VBZ).

Learning

If we define $\boldsymbol{f}(\boldsymbol{w}_{1:M}, y_m, y_{m-1}, m) = \{\langle W : w_m, y_m \rangle, \langle T : y_{m-1}, y_m \rangle\}$, then our model is identical to the HMM. If we set the weights of these features to the log of their maximum-likelihood estimates,

$$w_{\langle W:w_m,y_m \rangle} = \log \operatorname{count}(w_m, y_m) - \log \operatorname{count}(y_m)$$
$$w_{\langle T:y_{m-1},y_m \rangle} = \log \operatorname{count}(y_{m-1}, y_m) - \log \operatorname{count}(y_{m-1}),$$

then we exactly recover the HMM.

But to use more overlapping features and to get the advantages of error-driven learning, we're going to do perceptron updates. It's exactly the same as the non-structured perceptron:

- $\boldsymbol{\theta}^{(t+1)} \leftarrow \boldsymbol{\theta}^{(t)} + \boldsymbol{f}(\boldsymbol{x}, \boldsymbol{y}) - \boldsymbol{f}(\boldsymbol{x}, \hat{\boldsymbol{y}})$ is the standard update, using Viterbi to find $\hat{\boldsymbol{y}}$.

- As before, weight averaging is crucial to get good performance (Collins, 2002).

- As before, we can use Passive-Aggressive to do large-margin training (Crammer et al., 2006), computing the step size by dividing a non-negative **loss** $\ell(\boldsymbol{y}_i, \hat{\boldsymbol{y}})$ by the squared norm of the difference in the feature vectors, $||\boldsymbol{f}(\boldsymbol{y}_i, \boldsymbol{w}_i) - \boldsymbol{f}(\hat{\boldsymbol{y}}, \boldsymbol{w}_i)||^2$. A reasonable choice of loss function is the Hamming loss, which is the number of incorrect tag predictions (Taskar et al., 2003; Tsochantaridis et al., 2004).

12.8 Conditional random fields

Structured perceptron works well in practice, and you will implement in your project 2, where it should work much much better than the Hidden Markov Model.

- But sometimes we need probabilities, and SP doesn't give us that.

- The Conditional Random Field (CRF) is a probabilistic conditional model for sequence labeling.

- Just as structured perceptron is built on the perceptron classifier, conditional random fields are built on the logistic regression classifier.

$$p(y \mid \boldsymbol{x}) = \frac{e^{\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(y, \boldsymbol{x})}}{\sum_{y' \in \mathcal{Y}} e^{\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(y', \boldsymbol{x})}} \tag{12.31}$$

We can again move to structured prediction,

$$p(\boldsymbol{y} \mid \boldsymbol{w}) = \frac{e^{\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{y}, \boldsymbol{w})}}{\sum_{\boldsymbol{y}' \in \mathcal{T}(\boldsymbol{w})} e^{\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{y}', \boldsymbol{w})}}. \tag{12.32}$$

This is called a Conditional Random Field, because it models the sequence labeling task as a Markov random field, and estimates the probability of a set of variables **conditioned** on the others (as opposed to jointly, which the HMM does). We will need the same restriction on the scoring function as in Structured Perceptron: $\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}, \boldsymbol{y}) = \sum_m \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}, y_m, y_{m-1}, m)$.

An important note about CRFs is that the joint probability $p(\boldsymbol{w}_{1:M}, \boldsymbol{y}_{1:M})$ is simply the unnormalized conditional probability:

$$p(\boldsymbol{y}|\boldsymbol{w}) = \frac{p(\boldsymbol{y}, \boldsymbol{w})}{\sum_{\boldsymbol{y}'} p(\boldsymbol{y}', \boldsymbol{w})} \tag{12.33}$$

$$p(\boldsymbol{y}, \boldsymbol{w}) = e^{\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{y}, \boldsymbol{w})} \tag{12.34}$$

Aside: Maximum Entropy Markov Models (MEMMS)

Suppose we define

$$p(\boldsymbol{y} \mid \boldsymbol{w}) = \prod_m^M p(y_m \mid \boldsymbol{w}_{1:M}, y_{1:m-1}) \tag{12.35}$$

$$\approx \prod_m^M p(y_m | \boldsymbol{w}_{1:M}, y_{m-1}). \tag{12.36}$$

We can then define each local probability $p(y_m \mid \boldsymbol{w}_{1:M}, y_{m-1})$ as a logistic regression model,

$$p(y_m \mid \boldsymbol{w}_{1:M}, y_{m-1}) = \frac{\exp(\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}_{1:M}, y_m, y_{m-1}))}{\sum_{y' \in \mathcal{T}} \exp(\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}_{1:m}, y', y_{m-1}))}. \tag{12.37}$$

Recall that logistic regression is sometimes called **maximum entropy**, and observe that we are making a Markov assumption. Thus the name **Maximum Entropy Markov Model**.

Inference in the MEMM can be performed with Viterbi, choosing

$$v_m(k) = \max_{k'} P(Y_m = k \mid Y_{m-1} = k', w_{1:M}) v_{m-1}(k'). \tag{12.38}$$

The local decision model $p(y_m | w_{1:M}, y_{m-1})$ can be trained as a standard logistic regression classifier. The problem with this model is that learning to optimize individual tagging decisions is not the same as learning to produce optimal tag sequences. The local classifier is trained with the true value of y_{m-1}, not the value likely to be produced by the classifier — so, not necessarily the value that we are most likely to see in a test set tagging situation. This introduces a problem that Lafferty et al. (2001) called **label bias**. Put another way, the MEMM allows structured **prediction**, but it does not perform structured **learning**.

Decoding in CRFs

Back to CRFs! Decoding in the CRF does not depend on the denominator of $p(\boldsymbol{y} \mid \boldsymbol{w})$

$$
\begin{aligned}
\hat{\boldsymbol{y}} &= \arg\max_{\boldsymbol{y}} p(\boldsymbol{y} \mid \boldsymbol{w}) \\
&= \arg\max_{\boldsymbol{y}} \log p(\boldsymbol{y} \mid \boldsymbol{w}) \\
&= \arg\max_{\boldsymbol{y}} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{y}, \boldsymbol{w}) - \log \sum_{\boldsymbol{y}' \in \mathcal{T}(\boldsymbol{w})} e^{\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{y}', \boldsymbol{w})} \\
&= \arg\max_{\boldsymbol{y}} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{y}, \boldsymbol{w}) \\
&= \arg\max_{\boldsymbol{y}} \boldsymbol{\theta}^{\mathsf{T}} \sum_{m=0}^{M} \boldsymbol{f}(\boldsymbol{w}, y_m, y_{m-1}, m) \\
&= \arg\max_{\boldsymbol{y}} \sum_{m=0}^{M} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}, y_m, y_{m-1}, m)
\end{aligned}
$$

So we can apply Viterbi in exactly the same way that we did in the Structured Perceptron.

Learning in CRFs

Learning is a little more complicated. As with logistic regression, we need to learn weights to minimize the regularized negative log conditional probability,

$$
\begin{aligned}
\ell &= \sum_{i=1}^{N} -\log \mathrm{p}(\boldsymbol{y}_i \mid \boldsymbol{w}_i; \boldsymbol{\theta}) + \lambda ||\boldsymbol{\theta}||^2, \\
&= -\sum_{i} \boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}_i, \boldsymbol{y}_i) + \log \sum_{\boldsymbol{y}' \in \mathcal{T}(\boldsymbol{w}_i)} \exp\left(\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(\boldsymbol{w}_i, \boldsymbol{y}')\right) + \lambda ||\boldsymbol{\theta}||^2,
\end{aligned}
$$

where λ controls the amount of regularization. As in logistic regression, the gradient includes is a difference between observed and expected counts:

$$
\frac{d\ell}{d\theta_j} = \sum_{i} \mathrm{count}(\boldsymbol{w}_i, \boldsymbol{y}_i)_j - E_{\boldsymbol{y}|\boldsymbol{w}_i; \boldsymbol{\theta}}[\mathrm{count}(\boldsymbol{w}_i, \boldsymbol{y})_j] + \lambda \boldsymbol{\theta}
$$

$$
\mathrm{count}(\boldsymbol{w}_i, \boldsymbol{y}_i)_j = \sum_{m}^{M} f_j(\boldsymbol{w}_i, y_{i,m}, y_{i,m-1}, m)
$$

For example:

- If feature j is $\langle T : CC, DT \rangle$, then $\mathrm{count}(\boldsymbol{w}_i, \boldsymbol{y}_i)_j$ is the count of times DT follows CC in the sequence $\boldsymbol{y}_i$.

- If feature j is $\langle M : \text{-}thy, \text{JJ} \rangle$, then $\mathrm{count}(\boldsymbol{w}_i, \boldsymbol{y}_i)_j$ is the count of words ending in *-thy* in $\boldsymbol{w}_i$ that are tagged JJ in $\boldsymbol{y}_i$.

The expected feature counts are more difficult to compute.

$$
E_{\boldsymbol{y}|\boldsymbol{w}; \boldsymbol{\theta}}[\mathrm{count}(\boldsymbol{w}_i, \boldsymbol{y})_j] = \sum_{\boldsymbol{y} \in \mathcal{T}(\boldsymbol{w}_i)} P(\boldsymbol{y} \mid \boldsymbol{w}_i; \boldsymbol{\theta}) f_j(\boldsymbol{w}_i, \boldsymbol{y}) \tag{12.39}
$$

- This looks bad: we have to sum over an exponential number of labelings again.

- But remember that the feature function decomposition implies that

$$
f_j(\boldsymbol{w}, \boldsymbol{y}) = \sum_{m} f_j(\boldsymbol{w}, y_m, y_{m-1}, m) \tag{12.40}
$$

This means we can compute the expectation as,

$$E_{\boldsymbol{y}|\boldsymbol{w};\boldsymbol{\theta}}[\text{count}(\boldsymbol{w},\boldsymbol{y})_j] = \sum_{\boldsymbol{y}\in\mathcal{T}(\boldsymbol{w})} \text{p}(\boldsymbol{y}\mid\boldsymbol{w};\boldsymbol{\theta})f_j(\boldsymbol{w},\boldsymbol{y}) \tag{12.41}$$

$$= \sum_{\boldsymbol{y}\in\mathcal{T}(\boldsymbol{w})} \text{p}(\boldsymbol{y}\mid\boldsymbol{w};\boldsymbol{\theta})\sum_{m}^{M} f_j(\boldsymbol{w},y_m,y_{m-1},m) \tag{12.42}$$

$$= \sum_{m}^{M} \sum_{\boldsymbol{y}\in\mathcal{T}(\boldsymbol{w})} \text{p}(\boldsymbol{y}\mid\boldsymbol{w};\boldsymbol{\theta})f_j(\boldsymbol{w},y_m,y_{m-1},m) \tag{12.43}$$

$$= \sum_{m}^{M} \sum_{k,k'\in\mathcal{T}} \sum_{\boldsymbol{y}\in\mathcal{T}(\boldsymbol{w}):Y_{m-1}=k,Y_m=k'} \text{p}(\boldsymbol{y}\mid\boldsymbol{w};\boldsymbol{\theta})f_j(\boldsymbol{w},k',k,m)$$
$$\tag{12.44}$$

$$= \sum_{m}^{M} \sum_{k,k'\in\mathcal{T}} f_j(\boldsymbol{w},k',k,m) \sum_{\boldsymbol{y}\in\mathcal{T}(\boldsymbol{w}):y_{m-1}=k,y_m=k'} \text{p}(\boldsymbol{y}\mid\boldsymbol{w};\boldsymbol{\theta})$$
$$\tag{12.45}$$

$$= \sum_{m}^{M} \sum_{k,k'\in\mathcal{T}} f_j(\boldsymbol{w},k,k',m)P(Y_{m-1}=k,Y_m=k'\mid\boldsymbol{w};\boldsymbol{\theta})$$
$$\tag{12.46}$$

- The expected feature counts can be computed efficiently if we know the **marginal** probabilities $P(Y_m=k',Y_{m-1}=k\mid\boldsymbol{w};\boldsymbol{\theta})$.

- This is the probability of traversing the trellis edge $\langle m-1,k\rangle \to \langle m,k'\rangle$, conditioned on the entire observation $\boldsymbol{w}_{1:M}$. [Draw this in trellis]

- This marginal probability can be computed through the combination of two dynamic programming algorithms, the Forward and Backward algorithms.

12.9 The Forward-backward algorithm

The compute the gradient of the CRF objective with respect to the weights $\boldsymbol{\theta}$, we need marginal probabilities,

$$P(Y_m=k',Y_{m-1}=k|\boldsymbol{w}_{1:M}) = \frac{P(Y_m=k',Y_{m-1}=k,\boldsymbol{w}_{1:M})}{\text{p}(\boldsymbol{w}_{1:M})}. \tag{12.47}$$

To compute these, recall that in the CRF,

$$p(\boldsymbol{y}|\boldsymbol{w}) = \frac{\exp\left(\boldsymbol{\theta}^{\mathsf{T}}\boldsymbol{f}(\boldsymbol{w},\boldsymbol{y})\right)}{\sum_{\boldsymbol{y}'}\exp\left(\boldsymbol{\theta}^{\mathsf{T}}\boldsymbol{f}(\boldsymbol{w},\boldsymbol{y})\right)} \tag{12.48}$$

$$\propto \exp\left(\boldsymbol{\theta}^{\mathsf{T}}\boldsymbol{f}(\boldsymbol{w},\boldsymbol{y})\right) \tag{12.49}$$

$$= \exp\left(\boldsymbol{\theta}^{\mathsf{T}}\sum_{m}\boldsymbol{f}(\boldsymbol{w},y_m,y_{m-1},m)\right) \tag{12.50}$$

$$= \prod_{m}\exp\left(\boldsymbol{\theta}^{\mathsf{T}}\boldsymbol{f}(\boldsymbol{w},y_m,y_{m-1},m)\right) \tag{12.51}$$

$$= \prod_{m}\psi_m(y_m,y_{m-1}), \tag{12.52}$$

where $\psi_m(k,k') \triangleq \exp\left(\boldsymbol{\theta}^{\mathsf{T}}\boldsymbol{f}(\boldsymbol{w},Y_m=k,Y_{m-1}=k',m)\right)$. This quantity is sometimes called a **potential**. A Hidden Markov Model (HMM) can also be expressed in terms of a product of potential functions, $\psi_m(k,k') = p_E(w_m \mid Y_m = k)P_T(Y_m = k \mid Y_{m-1} = k')$.

Using this expression for the probability of a label sequence, we will compute the desired marginal probabilities in terms of three quantities:

$$Z \triangleq \sum_{\boldsymbol{y}'}\prod_{m}\psi_m(y'_m,y'_{m-1}) \tag{12.53}$$

$$\alpha_m(k) \triangleq \sum_{\boldsymbol{y}'_{1:m-1}}\psi_m(k,y_{m-1})\prod_{n}^{m-1}\psi_n(y'_n,y'_{n-1}) \tag{12.54}$$

$$\beta_m(k) \triangleq \sum_{\boldsymbol{y}'_{m+1:M}}\psi_{m+1}(y_{m+1},k)\prod_{n=m+2}^{M}\psi_n(y'_n,y'_{n-1}), \tag{12.55}$$

where Z is the **normalization constant** (also called the partition function), α is the set of forward variables (just like in the forward algorithm defined earlier), and β is the set of backward variables.

If we assume that the first tag and last tag have special values, $y_0 = \Diamond$ and

$y_M = \square$, then by definition,

$$
\begin{align}
Z &= \alpha_M(\square) \tag{12.56} \\
&= \beta_0(\Diamond) \tag{12.57} \\
&= \sum_{\boldsymbol{y}} \mathrm{p}(\boldsymbol{w}, \boldsymbol{y}) \tag{12.58} \\
&= \mathrm{p}(\boldsymbol{w}) \tag{12.59}
\end{align}
$$

The forward variables

We define the first sum as $\alpha_m(k)$, called the **forward variable**. This is just like in the forward algorithm defined earlier for Hidden Markov Models; again, we can derive a recursive expression for α as follows:

$$
\alpha_0(\Diamond) \triangleq 1 \tag{12.60}
$$

$$
\alpha_m(k) \triangleq \sum_{\boldsymbol{y}'_{1:m-1}} \psi_m(k, y'_{m-1}) \prod_{n=1}^{m-1} \psi_n(y'_n, y'_{n-1}) \tag{12.61}
$$

$$
= \sum_{Y_{m-1}=k'} \psi_m(k, k') \sum_{\boldsymbol{y}_{1:m-2}} \psi_{m-1}(k', y'_{m-2}) \prod_{n=1}^{m-2} \psi_n(y'_n, y'_{n-1}) \tag{12.62}
$$

$$
= \sum_{Y_{m-1}=k'} \psi_m(k, k') \alpha_{m-1}(k'), \tag{12.63}
$$

where we simply substitute in the recursive expression for $\alpha_{m-1}(k')$. This variable is computed from left to right, with each $\boldsymbol{\alpha}_m$ depending on $\boldsymbol{\alpha}_{m-1}$.

The backward variables

We define the second sum as $\beta_m(k)$, called the **backward variable**. We can derive a recursive expression for β as follows:

$$\beta_M(\square) \triangleq 1 \tag{12.64}$$

$$\beta_m(k) \triangleq \sum_{\boldsymbol{y}'_{m+1:M}} \psi_{m+1}(y_{m+1}, k) \prod_{n=m+2}^{M} \psi_n(y'_n, y'_{n-1}), \tag{12.65}$$

$$= \sum_{Y_{m+1}=k'} \psi_{m+1}(k', k) \sum_{\boldsymbol{y}_{m+2:M}} \psi_{m+2}(y_{m+2}, k') \prod_{n=m+3}^{M} \psi_n(y_n, y_{m-1}) \tag{12.66}$$

$$= \sum_{Y_{m+1}=k'} \psi_{m+1}(k', k)\beta_{m+1}(k') \tag{12.67}$$

This variable is computed from right to left, with each $\boldsymbol{\beta}_m$ depending on $\boldsymbol{\beta}_{m+1}$.

Computing the marginals   Now, the product $\alpha_m(k)\beta_m(k)$ is equal to

$$\alpha_m(k)\beta_m(k) = \sum_{\boldsymbol{y}_{1:m-1}} \left(\prod_{n=1}^{m-1} \psi_n(y_n, y_{n-1}) \right) \psi_m(k, y_{m-1}) \tag{12.68}$$

$$\times \sum_{\boldsymbol{y}_{m+1:M}} \psi_{m+1}(y_{m+1}, k) \left(\prod_{n=m+2}^{M} \psi_n(y_n, y_{n-1}) \right) \tag{12.69}$$

$$= \sum_{\boldsymbol{y}_{1:M}:Y_m=k} \prod_{n=1}^{M} \psi_n(y_n, y_{n-1}), \tag{12.70}$$

which is exactly equal to the sum of the unnormalized probabilities of all sequences in which $Y_m = k$. To compute the normalized probability, we simply divide by Z, so that,

$$P(Y_m = k \mid \boldsymbol{w}_{1:M}) = \frac{\alpha_m(k)\beta_m(k)}{Z} \tag{12.71}$$

To obtain the unnormalized probability of a tag-to-tag transition $\langle Y_m = k, Y_{m-1} =$

$k'\rangle$, we compute,

$$\alpha_{m-1}(k')\psi_m(k, k')\beta_m(k) = \sum_{\boldsymbol{y}_{1:m-2}} \left(\prod_{n=1}^{m-2} \psi_n(y_n, y_{n-1}) \right) \psi_{m-1}(k', y_{m-2}) \quad (12.72)$$

$$\times \psi_m(k, k') \sum_{\boldsymbol{y}_{m+1:M}} \psi_{m+1}(y_{m+1}, k) \prod_{n=m+2}^{M} \psi_n(y_n, y_{n-1})$$
$$(12.73)$$

$$= \sum_{\boldsymbol{y}_{1:M}:Y_m=k,Y_{m-1}=k'} \prod_{n=1}^{M} \psi_n(y_n, y_{n-1}), \quad (12.74)$$

which is exactly equal to the unnormalized probability of all sequences in which $Y_m = k$ **and** $Y_{m-1} = k'$. To compute the normalized probability, we again divide by Z, so that,

$$P(Y_m = k, Y_{m-1} = k' \mid \boldsymbol{w}_{1:M}) = \frac{\alpha_{m-1}(k')\psi_m(k, k')\beta_m(k)}{Z} \quad (12.75)$$

Learning in CRFs: wrapup

The overall procedure looks like logistic regression:

- Use forward-backward to compute expected feature counts under $P(\boldsymbol{y} \mid \boldsymbol{w}; \boldsymbol{\theta})$:

$$E[f_j(\boldsymbol{y}, \boldsymbol{w})] = \sum_m \sum_{k,k'} P(Y_m = k, Y_{m-1} = k'|\boldsymbol{w}_{1:M}) f_j(\boldsymbol{w}, k, k', m) \quad (12.76)$$

$$= \sum_m \sum_{k,k'} \frac{\alpha_m(k)\beta_{m+1}(k')}{Z} f_j(\boldsymbol{w}, k, k', m). \quad (12.77)$$

- Compute gradients as difference between feature counts and expected counts, $\boldsymbol{f}(\boldsymbol{y}, \boldsymbol{w}) - E[\boldsymbol{f}(\boldsymbol{y}, \boldsymbol{w})]$.

- Update $\boldsymbol{\theta}$, using quasi-newton optimization or stochastic gradient descent. The CRFsuite package implements several of these learning algorithms (http://www.chokkan.org/software/crfsuite/).

- Iterate, recomputing the expected feature counts.

12.10   Unsupervised sequence labeling

In unsupervised sequence labeling, we want to induce a Hidden Markov Model (HMM) from a corpus of unannotated text $\boldsymbol{w}_1, \boldsymbol{w}_2, \ldots, \boldsymbol{w}_N$; this is an example of the general problem of **structure induction**, which is the unsupervised version of **structure prediction**. The tags that result from unsupervised sequence labeling might be useful for some downstream task, or for better understanding the language's inherent structure; or, we might want to do probability density estimation for sequences, as in gesture or activity recognition (Mitra and Acharya, 2007). Another reason would be to do semi-supervised learning, imputing tag sequences for unlabeled data. For part-of-speech tagging, often we use a tag dictionary which lists the allowed tags for each word, simplifying the problem (Christodoulopoulos et al., 2010).

In any case, we can perform unsupervised learning by applying expectation-maximization. In the M-step, we compute the HMM parameters from expected counts:

$$P_E(W = i | Y = k) = \phi_{k,i} = \frac{E[\text{count}(W = i, Y = k)]}{E[\text{count}(Y = k)]}$$

$$P_T(Y_m = k | Y_{m-1} = k') = \lambda_{k',k} = \frac{E[\text{count}(Y_m = k, Y_{m-1} = k')]}{E[\text{count}(Y_{m-1} = k')]}$$

The expected counts are computed in the E-step, using the forward and backward variables.

$$E[\text{count}(W = i, Y = k)] = \sum_m P(Y_m = k | \boldsymbol{w}) \delta(W_m = i) \tag{12.78}$$

$$= \sum_m \frac{P(Y_m = k, \boldsymbol{w}_{1:m}) \text{p}(\boldsymbol{w}_{m+1:M} | Y_m = k)}{\text{p}(\boldsymbol{w}_{1:M})} \delta(w_m = i) \tag{12.79}$$

$$= \frac{1}{\alpha_M(\square)} \sum_m \alpha_m(k) \beta_m(k) \delta(w_m = i) \tag{12.80}$$

We use the chain rule to separate $\boldsymbol{w}_{1:m}$ and $\boldsymbol{w}_{m+1:M}$, and then use the definitions of the forward and backward variables. In the final step, we normalize by $\text{p}(\boldsymbol{w}_{1:M}) = \alpha_M(\square) = \beta_0(\lozenge)$.

$$E[\text{count}(Y_m = k, Y_{m-1} = k')] = \sum_m P(Y_m = k, Y_{m-1} = k' | \boldsymbol{w}) \tag{12.81}$$

$$\propto \sum_m P(Y_{m-1} = k', \boldsymbol{w}_{1:m-1}) P(w_{m+1:M} | Y_m = k)$$
$$\times P(w_m, Y_m = k | Y_{m-1} = k') \tag{12.82}$$

$$= \sum_m P(Y_{m-1} = k', \boldsymbol{w}_{1:m-1}) P(w_{m+1:M} | Y_m = k)$$
$$\times \text{p}(w_m | Y_m = k) P(Y_m = k | Y_{m-1} = k') \tag{12.83}$$

$$= \sum_m \alpha_{m-1}(k') \beta_m(k) \phi_{k,w_m} \lambda_{k' \to k} \tag{12.84}$$

Again, we use the chain rule to separate out $\boldsymbol{w}_{1:m-1}$ and $\boldsymbol{w}_{m+1:M}$, and use the definitions of the forward and backward variables. The final computation also includes the parameters ϕ and λ, which govern (respectively) the emission and transition properties between w_m, y_m, and y_{m-1}. Note that the derivation only shows how to compute this to a constant of proportionality; we would divide by $\text{p}(\boldsymbol{w}_{1:M})$ to go from the joint probability $P(Y_{m-1} = k', Y_m = k, \boldsymbol{w}_{1:M})$ to the desired conditional $P(Y_{m-1} = k', Y_m = k \mid \boldsymbol{w}_{1:M})$.

Linear dynamical systems

The forward-backward algorithm can be viewed as Bayesian state estimation in a discrete state space. In a continuous state-space, $y_m \in \mathbb{R}$, the equivalent algorithm is the **Kalman Smoother**. It also computes marginals $\text{p}(y_m \mid \boldsymbol{x}_{1:M})$, using a similar two-step algorithm of forward and backward passes. Instead of computing a table of values at each step ($\alpha_m(k)$ and $\beta_m(k)$), they compute a probability density function, characterized by a mean and covariance around the latent state. Connections between the Kalman Smoother and the forward-backward algorithm are elucidated by Minka (1999).

Alternative unsupervised learning methods

Expectation-maximization is just one of many techniques for structure induction. One alternative is to use a family of randomized algorithms called **Markov Chain Monte Carlo (MCMC)**. In these algorithms, we compute a marginal distribution over the latent variable $\boldsymbol{y}$ **empirically**, by drawing random samples. The randomness explains the "Monte Carlo" part of the name; typically, we employ a Markov

Chain sampling procedure, meaning that each sample is drawn from a distribution that depends only on the previous sample (and not on the entire sampling history). A simple MCMC algorithm is **Gibbs Sampling**, in which we iteratively sample each y_m conditioned on all the others (Finkel et al., 2005):

$$\mathrm{p}(y_m \mid \boldsymbol{y}_{-m}, \boldsymbol{w}_{1:M}) \propto \mathrm{p}(w_m \mid y_m)\mathrm{p}(y_m \mid \boldsymbol{y}_{-m}). \tag{12.85}$$

Gibbs Sampling has been applied to unsupervised part-of-speech tagging by Goldwater and Griffiths (2007). *Beam sampling* is a more complicated sampling algorithm, which randomly draws entire sequences $\boldsymbol{y}_{1:M}$, rather than individual tags y_m; this algorithm was applied to unsupervised part-of-speech tagging by Van Gael et al. (2009).

EM is guaranteed to find only a local optimum; MCMC algorithms will converge to the true posterior distribution $\mathrm{p}(\boldsymbol{y}_{1:M} \mid \boldsymbol{w}_{1:M})$, but this is only guaranteed in the limit of infinite samples. Recent work has explored the use of **spectral learning** for latent variable models, which use matrix and tensor decompositions to provide guaranteed convergence under mild assumptions (Hsu et al., 2012). Georgia Tech faculty Byron Boots and Le Song are among the leaders in this active area of research (Song et al., 2010).

Chapter 13

Context-free grammars

So far we've explored finite-state models, which correspond to regular languages.

- **representations**: (weighted) finite state automata

- **probabilistic models**: HMMs (as a special case), CRFs

- **algorithms**: Viterbi, Forward-Backward, $\mathcal{O}(NK^2)$ time complexity.

- **linguistic phenomena**:

  - morphology
  - language models
  - part-of-speech disambiguation
  - named entity recognition (chunking)

Is the finite state representation enough for natural language?

13.1   Is English a regular language?

Regular languages are closed under intersection:

- $K \cap L$ is the set of strings in both K and L

- $K \cap L$ is regular iff K and L are regular

How to prove English is not regular:

- Let K be the set of grammatical English sentences

- Let L be some regular language

- Show that the intersection is not regular

We're going to prove this using center embedding:

1. *The cat is fat.*

2. *The cat that the dog chased is fat.*

3. **The cat that the dog is fat.*

4. *The cat that the dog that the monkey kissed chased is fat.*

5. **The cat that the dog that the monkey chased is fat.*

Proof sketch:

- K is the set of grammatical english sentences.
  It excludes sentences (3) and (5).

- L is the regular language *the cat* $(that\ N)^+ V_t^+$ *is fat.*

- The language $L \cap K$ is *the cat* $(that\ N)^n V_t^n$ *is fat.*

It is important to understand that the issue here is not just infinite repetition or productivity; FSAs can handle productive phenomena like *the big red smelly plastic figurine*. It is specifically the center-embedding phenomenon, because this leads to the same structure as the classic $a^n b^n$ language. What do you think of this argument?

Is deep center embedding really part of English?

Karlsson (2007) searched for multiple (phrasal) center embeddings in corpora from 7 languages:

- Very few examples of double embedding

- Only 13 examples of triple embedding (none in speech)

- Zero examples of quadruple embeddings

Note that we can build an FSA to accept center-embedding up to any finite depth. Chomsky and many linguists distinguish between

- **Competence**: the fundamental abilities of the (idealized) human language processing system

- **Performance**: real utterances produced by speakers, subject to non-linguistic factors such as cognitive limitations

Even if English *as performed* is regular, the underlying generative grammar may be context-free... **or beyond**. There is a similar proof that at least some languages are not context-free! I'll post slides with this proof idea.

How much expressiveness do we need?

- Shieber (1985) makes a similar argument, showing that case agreement in Swiss-German cross-serial constructions is homomorphic to a formal language $wa^m b^n x c^m d^n y$, which is weakly non-context free. In response to the objection that all attested constructions are finite, Shieber writes:

  > Down this path lies tyranny. Acceptance of this argument opens the way to proofs of natural languages as regular, nay, **finite**.

- In practice, many real constructions are much simpler to handle in context-free rather than finite-state representations:

  > The **processor has** 10 million times fewer transistors on it than todays typical microprocessors, **runs** much more slowly, and **operates** at five times the voltage...

  - The easy way:

    $$S \rightarrow NN \ VP$$
    $$VP \rightarrow VP3S \mid VPN3S \mid \ldots$$
    $$VP3S \rightarrow VP3S, VP3S, \textit{and } VP3S \mid VBZ \mid VBZ \ NP \mid \ldots$$

  - The hard way: build an FST that basically replicates all of English grammar for VPs with 3S and non-3S subjects.

- Mainstream parsing focuses on CFGs, but there is some work on "mildly" context-sensitive grammars.

13.2   Context-Free Languages

In the Chomsky hierarchy, context-free languages (CFLs) are a strict generalization of regular languages.

regular	context-free
regular expressions	context-free grammars (CFGs)
finite-state machines	pushdown automata
paths	derivations

Context-free grammars define CFLs. They are sets of permissible *productions* which allow you to **derive** strings composed of surface symbols.

$$S \rightarrow NP\ VP_1$$
$$NP \rightarrow the\ N\ \mid NP\ \text{RELCLAUSE}$$
$$\text{RELCLAUSE} \rightarrow that\ NP\ V_t$$
$$V_t \rightarrow ate \mid chased \mid befriended \mid \ldots$$
$$N \rightarrow cat \mid dog \mid monkey \mid \ldots$$
$$VP_1 \rightarrow is\ fat$$

An important feature of CFGs is *recursion*, in which a nonterminal can be derived from itself.

More formally: a CFG is a tuple $\langle N, \Sigma, R, S \rangle$:

N a set of non-terminals
Σ a set of terminals (distinct from N)
R a set of productions, each of the form
     $A \rightarrow \beta$, where $A \in N$ and $\beta \in (\Sigma \cup N)^*$
S a designated start symbol

- Context free grammars provide rules for generating strings.

    - RHS: a non-terminal $\in N$

    - LHS: a sequence of terminals or non-terminals, $\{n, \sigma\}^*, n \in N, \sigma \in \Sigma$.

- A **derivation** t is a sequence of steps from S to a surface string $w \in \Sigma^*$, which is the yield of the derivation.

- If $\exists t : s \in \text{yield}(t)$, then w is grammatical. Equivalently, for grammar G, we say $|\mathcal{T}_G(w)| \geq 1$.

- If $\exists t, t' : w \in \text{yield}(t) \land s \in \text{yield}(t')$, then w is ambiguous. Equivalently, for grammar G, we say $|\mathcal{T}_G(w)| > 1$.

- A derivation can be viewed as trees or as bracketings:

```
                    S
          NP                VP
          PRP   VBZ              NP
          She   eats      NP          PP
                          NN     IN       NP
                          sushi  with     NNS
                                          chopsticks
```

$(_S(_{NP}(_{PRP} \textit{She})(_{VP}(_{VBZ} \textit{eats})$
$(_{NP}(_{NP}(_{NN} \textit{sushi}))(_{PP} \textit{with}(_{NP}(_{NNS} \textit{chopsticks}))))))))$

```
                    S
          NP                    VP
          PRP   VBZ    NP            PP
          She   eats   NN     IN        NP
                       sushi  with      NNS
                                        chopsticks
```

$(_S(_{NP}(_{PRP} \textit{She})(_{VP}(_{VBZ} \textit{eats})$
$(_{NP}(_{NN} \textit{sushi}))$
$(_{PP} \textit{with}(_{NP}(_{NNS} \textit{chopsticks}))))))$

Semantics   Ideally, each derivation will have a distinct semantic interpretation, and all possible interpretations will be represented in some derivation.

$(_{NP}(_{NP} \textit{Ban} (_{PP} \textit{on} (_{NP} \textit{nude dancing})))$
$(_{PP} \textit{on} (_{NP} \textit{Governor's desk})))$

$$(_{NP} \text{ Ban } (_{PP} \text{ on } (_{NP}(_{NP} \text{ nude dancing })$$
$$(_{PP} \text{ on } (_{NP} \text{ Governor's desk })))))$$

Sadly, this is not always the case.

$$(_{NP}(_{JJ} \text{ nice}) (_{JJ} \text{ little}) (_{NN} \text{ car }))$$
$$(_{NP}(_{JJ} \text{ nice}) (_{NP}(_{JJ} \text{ little}) (_{NN} \text{ car })))$$
$$(_{NP}(_{JJ} \text{ nice}) (_{NP}(_{JJ} \text{ little}) (_{NP}(_{NN} \text{ car }))))$$

13.3   Constituents

- In natural language grammars, the non-terminals should reflect syntactic categories.

- Bracketed substrings (e.g., *sushi with chopsticks*) are called **constituents**.

- There are several tests for constituency, including:

    - substitution
    - coordination
    - movement

Substitution   Constituents generated by the same non-terminal should be substitutable in many contexts:

- $(_{NP}$ *The ban* $)$ *is on the desk.*

- $(_{NP}$ *The Governor's desk* $)$ *is on the desk.*

- $(_{NP}$ *The ban on dancing on the desk* $)$ *is on the desk.*

- *$(_{PP}$ *On the desk* $)$ *is on the desk.*

A more precise test for whether a set of substrings constitute a single category is whether they can be replaced by the same pronouns.

- $(_{NP}$ *It* $)$ *is on the desk.*

What about verbs?

- *I ($_V$ gave) it to Anne.*

- *I ($_V$ taught) it to Anne.*

- *I ($_V$ gave) Anne a fish*

- **I ($_V$ taught) Anne a fish*

This suggests we need nonterminals which distinguish verbs based on the arguments they can take. The technical name for this is *subcategorization*.

Coordination Constituents generated by the same non-terminal can usually be *coordinated* using words like *and* and *or*:

- *We fought ($_{PP}$ on the hills) and ($_{PP}$ in the hedges).*

- *We fought ($_{ADVP}$ as well as we could).*

- **We fought ($_{ADVP}$ as well as we could) and ($_{PP}$ in the hedges).*

This too doesn't always work:

- *She ($_{VP}$ went) ($_{PP}$ to the store).*

- *She ($_{VP}$ came) ($_{PP}$ from the store).*

- *She (went to) and (came from) the store.*

Movement Valid constituents can be moved as a unit, preserving grammaticality.

- Passivization

  – *(The governor) banned (nude dancing on his desk)*
  – *(Nude dancing on his desk) was banned by (the governor)*

- Wh- movement

  – *(Nude dancing was banned) on (the desk).*
  – *(The desk) is where (nude dancing was banned)*

- Topicalization

  – *(He banned nude dancing) to appeal to conservatives.*
  – *To appeal to conservatives, (he banned nude dancing).*

13.4   A simple grammar of English

Noun phrases

Let's start with noun phrases:

- ***She*** *sleeps* (Pronoun)

- ***Arlo*** *sleeps* (Proper noun)

- ***Fish*** *sleep* (Mass noun)

- ***The fish*** *sleeps* (determiner + noun)

- ***The blue fish*** *sleeps* (DT + JJ + NN)

- ***The girl from Omaha*** *sleeps* (NP + PP)

- ***The student who ate 15 donuts*** *sleeps* (NP + RelClause)

So overall, we can summarize this fragment as

$$\text{NP} \rightarrow \text{PRP} \mid \text{NNP} \mid \text{DT NOM}$$
$$\text{NOM} \rightarrow \text{ADJP NOM} \mid \text{NN}$$
$$\text{NP} \rightarrow \text{NP PP} \mid \text{NP RELCLAUSE}$$

We're leaving out some detail, like pluralization and possessives, but you get the idea.

Adjectival and prepositional phrases

- *Very funny*

- *The **large, blue** fish*

- *The man **from la mancha***

$$\text{ADJP} \rightarrow \text{JJ} \mid \text{RB ADJP} \mid \text{JJ ADJP}$$
$$\text{PP} \rightarrow \text{IN NP} \mid \text{TO NP}$$

Verb phrases

- *She **sleeps***

- *She **sleeps restlessly***

- *She **sleeps at home***

- *She **eats sushi**[1]*

- *She **gives John sushi***

$$\text{VP} \rightarrow \text{V} \mid \text{VP RB} \mid \text{VP PP} \mid \text{V NP} \mid \text{V NP NP} \mid \text{V NP RB}$$

But what about **She sleeps sushi* or **She speaks John Japanese*?

- Classes of verbs can take different numbers of arguments.

- This is called **subcategorization**

$$\text{VP} \rightarrow \text{V-INTRANS} \mid \text{V-TRANS NP} \mid \text{V-DITRANS NP NP}$$
$$\text{VP} \rightarrow \text{VP RB} \mid \text{VP PP}$$

We would also need to handle modal and auxiliary verbs that allow us to create complex tenses, like *She will have eaten sushi* but not **She will have eats sushi*.

Sentences

- *She eats sushi*
$$\text{S} \rightarrow \text{NP VP}$$

- *Sometimes, she eats sushi*
$$\text{S} \rightarrow \text{ADVP S}$$

- *In Japan, she eats sushi*
$$\text{S} \rightarrow \text{PP S}$$

- What about **I eats sushi*, **She eat sushi*??
$$\text{S} \rightarrow \text{NP.3S VP.3S} \mid \text{NP.N3S VP.N3S}$$

In general, we need **features** to capture this kind of agreement.

---

[1]Sushi examples from Julia Hockenmaier

Conjunctions

- *She eats sushi and candy*

$$\text{NP} \rightarrow \text{NP } and \text{ NP}$$

- *She eats sushi and drinks soda*

$$\text{VP} \rightarrow \text{VP } and \text{ VP}$$

- *She eats sushi and he drinks soda*

$$\text{S} \rightarrow \text{S } and \text{ S}$$

- *fresh and tasty sushi*

$$\text{ADJP} \rightarrow \text{JJ } and \text{ JJ}$$

We'd need a little more cleverness to properly cover groups larger than two.

Odds and ends

- *I gave sushi to the girl **who eats sushi***. This is a relative clause,

$$\text{RELCLAUSE} \rightarrow who \text{ VP} \mid that \text{ VP}$$

- *I took sushi from the man **offering sushi***. This is a gerundive postmodifier.

$$\text{NOM} \rightarrow \text{NOM GERUNDVP}$$
$$\text{GERUNDVP} \rightarrow \text{VBZ} \mid \text{VBZ NP} \mid \text{VBZ PP} \mid \ldots$$

- ***Can** she eat sushi?* (notice it's not *eats*)

$$\text{S} \rightarrow \text{AUX NP VP}$$

- *... and many more*

13.5 Grammar design

Our goal is a grammar that avoids

- **Overgeneration**: deriving strings that are not grammatical.

- **Undergeneration**: failing to derive strings that are grammatical.

To avoid undergeneration, we would need thousands of productions.

Typically, grammars are defined in conjunction with large-scale **treebank** annotation projects.

- An annotation guideline specifies the non-terminals and how they go together.

- The annotators then apply these guidelines to data.

- The grammar rules can then be read off the data.

The Penn Treebank contains one million parsed words of Wall Street Journal text (Marcus et al., 1993).

13.6 Grammar equivalence and normal form

- Grammars are **weakly equivalent** if they generate the same strings.

- Grammars are **strongly equivalent** if they generate the same strings **and** assign the same phrase structure to each string.

- In Chomsky Normal Form (CNF), all productions are either:

$$A \to BC$$
$$A \to a$$

  - All CFGs can be converted into a weakly equivalent grammar in CNF.
  - This is very handy for parsing algorithms.

- Binarization is easy:
  group right children into new non-terminals.

- Un-binarization is important!
  people with claws is not a constituent in the original parse.

- Unary productions are best handled by modifying the algorithm.

13.7   Parsing

Parsing is the process of determining whether a sentence is in a context-free language, by searching for a legal derivation. Some possibilities:

- **Top-down**: start with the start symbol, and see if we can derive the sentence.

- **Bottom-up**: combine the observed symbols using whatever productions we can, until we reach the start symbol

- **Left-to-right**: move through the input, incrementally building a parse tree

Before we get into these different possibilities, let's see whether exhaustive search is possible. Suppose we only have one non-terminal, X, and it has binary productions

$$X \rightarrow X \ X$$
$$X \rightarrow \textit{the girl} \mid \textit{ate sushi} \mid \ldots$$

How many different ways could we parse a sentence? This is just equal to the number of binary bracketings of the words in the sentence, which is a Catalan number. Catalan numbers grow **super-exponentially** in the length of the sentence, $C_n = \frac{(2n)!}{(n+1)!n!}$.

13.8 CKY parsing

CKY is a bottom-up parsing allows us to test whether a sentence is in a context-free language, without considering all possible parses. First we form small constituents, then try to merge them into larger constituents.

Let's start with an example grammar:

$$S \rightarrow VP \ NP$$
$$NP \rightarrow NP \ PP \mid \textit{we} \mid \textit{sushi} \mid \textit{chopsticks}$$
$$PP \rightarrow P \ NP$$
$$P \rightarrow \textit{with}$$
$$VP \rightarrow VP \ NP \mid VP \ PP \mid \textit{eat}$$

Suppose we encounter the sentence *We eat sushi with chopsticks*.

- The first thing that we notice is that we can apply unary productions to obtain NP VP NP P NP

- Next, we can apply a binary production to merge the first NP VP into an S.

- Or we could merge VP NP into VP

- ... and so on

Let's systematize this. Here is the CKY algorithm:

Algorithm 1 The CKY algorithm for CFG parsing

---

for j : [1,N] **do**
  $t[j, j - 1] \leftarrow \{X \mid X \rightarrow w_j \in R\}$
  for i : [j-2, 0] **do**
    for k : [i+1, j-1] **do**
      $t[i, j] \leftarrow t[i, j] \cup \{X \mid X \rightarrow YZ \in R, Y \in t[i, k], Z \in t[k, j]\}$
    end for
  end for
end for

To hand unary transitions, we compute the *unary closure* of each non-terminal.

- e.g., if $S \rightarrow VP, VP \rightarrow V$, then add $S \rightarrow V$

- At each table entry $t[i, j]$

  - For each non-terminal $A \in t[i, j]$
    * Add all elements of the reflexive unary closure for A

$$
\begin{array}{c}
S \\
| \\
VP \\
| \\
V \\
| \\
\text{eat!}
\end{array}
$$

Complexity   What is the complexity of CKY?

- Space compexity: $\mathcal{O}(M^2|N|)$

- Time complexity: $\mathcal{O}(M^3|R|)$

- M is length of sentence,
  $|N|$ is the number of non-terminals,
  $|R|$ is the number of production rules

- But in practice... It's worse than worst-case! (Figure 13.1)

- Because longer sentences "unlock" more of the grammar.

Figure 13.1: Figure from Dan Klein's lecture slides

13.9 Ambiguity in parsing

- Syntactic ambiguity is endemic to natural language:[2]

  – Attachment ambiguity: *we eat sushi with chopsticks*,
    I shot an elephant in my pajamas.

  – Modifier scope: *southern food store*

  – Particle versus preposition: *The puppy tore up the staircase.*

  – Complement structure: *The tourists objected to the guide that they couldn't hear.*

  – Coordination scope: *"I see," said the blind man, as he picked up the hammer and saw.*

  – Multiple gap constructions: *The chicken is ready to eat*

---

[2]Examples borrowed from Dan Klein

- In morphology, we didn't just want to know which derivational forms are *legal*, we wanted to know which were likely.

- Syntactic parsing is all about choosing among the many, many legal parses for a given sentence.

Here's another example, which we've seen before:

```
              S
          /       \
        NP          VP
        |        /      \
        N       V          NP
        |       |        /     \
       Fed   raises    N        N
                       |        |
                    interest   rates
```

Figure 13.2: An example parse

- A minimal grammar permits 36 parses!

- Real-size broad coverage grammars permit millions of parses of typical sentences.

- Collins (2013) gives another example, *She announced a program to promote safety in trucks.*

Classical parsers faced a tradeoff:

- broad coverage with tons of ambiguity...

- or limited coverage in exchange for constraints on ambiguity

Consequently, deterministic parsers produced no analysis for many sentences.

Local solutions

Some ambiguity can be resolved locally:

- [*imposed* [*a ban* [*on asbestos*]]]

- [*imposed* [*a ban*][*on asbestos*]]

- Hindle and Rooth (1990) proposed a likelihood ratio test:

$$LR(v, n, p) = \frac{\mathrm{p}(p \mid v)}{\mathrm{p}(p \mid n)} = \frac{\mathrm{p}(on \mid imposed)}{\mathrm{p}(on \mid ban)}$$

  where we select VERB attachment if $LR(v, n, p) > 1$.

- But the likelihood-ratio approach ignores important information, like the phrase being attached.

  - ...[*it* [*would end* [*its venture* [*with Maserati*]]]]
  - ...[*it* [*would end* [*its venture*][*with Maserati*]]]

- The likelihood ratio gets this wrong

  - $\mathrm{p}(with \mid end) = \frac{607}{5156} = 0.118$
  - $\mathrm{p}(with \mid venture) = \frac{155}{1442} = 0.107$

  Other features (e.g., *Maserati*) argue for noun attachment. How can we add them?

Machine learning solutions   Ratnaparkhi et al. (1994) propose a maximum-entropy (logistic regression) approach:

$$P(\text{Noun attachment} \mid would\ end\ its\ venture\ with\ Maserati) =$$

$$\frac{e^{\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(would\ end\ its\ venture\ with\ Maserati)}}{1 + e^{\boldsymbol{\theta}^{\mathsf{T}} \boldsymbol{f}(would\ end\ its\ venture\ with\ Maserati)}}$$

Features include n-grams and word classes from hierarchical word clustering; accuracy is roughly 80%.

Collins and Brooks (1995) argued that attachment depends on four **heads**:

- the preposition (*with*)

- the VP attachment site (*end*)

- the NP attachment site (*venture*)

- the NP to be attached (*Maserati*)

They propose a backoff-based approach:

- First, look for counts of the tuple ⟨*with, Maserati, end, venture*⟩

- If none, try ⟨*with, Maserati, end*⟩ + ⟨*with,end,venture*⟩ + ⟨*with,Maserati, venture*⟩

- If none, try ⟨*with, Maserati*⟩ + ⟨*with,end*⟩ + ⟨*with,venture*⟩

- If none, try ⟨*with*⟩

Accuracy is roughly 84%. This approach of combining relative frequency estimation, smoothing, and backoff was very characteristic of late 1990s statistical NLP.

Beyond local solutions

Framing the problem as attachment ambiguity is limiting:

- assumes the parse is mostly done, leaving just a few attachment ambiguities to solve

- But realistic sentences have more than a few syntactic interpretations.

- Attachment decisions are interdependent:

  - *Cats scratch people with claws with knives.*
  - We may want to attach *with claws* to *scratch*.
  - But then we have nowhere to put *with knives*.

The task of statistical parsing is to produce a single analysis that resolves all syntactic ambiguities.

13.10 Statistical parsing with PCFGs

We want the parse τ that maximizes $P(\tau \mid \boldsymbol{w})$.

$$\arg\max_\tau P(\tau \mid \boldsymbol{w}) = \arg\max_\tau \frac{P(\tau, \boldsymbol{w})}{P(\boldsymbol{w})}$$

$$= \arg\max_\tau P(\tau, \boldsymbol{w})$$

$$= \arg\max_\tau P(S \mid \tau) P(\tau)$$

$$= \arg\max_{\tau : \boldsymbol{w} = \text{yield}(\tau)} P(\tau)$$

- The **yield** of a tree is the string of terminal symbols that can be read off the leaf nodes.

- The set $\{\tau : \boldsymbol{w} = \text{yield}(\tau)\}$ is exactly the set of all derivations of $\boldsymbol{w}$ in a CFG G.

PCFGs extend the CFG by adding probability to each production:

|		
S	→ NP VP	0.9
S	→ S CC S	0.1
NP	→ N	0.2
NP	→ DT N	0.3
NP	→ N NP	0.2
NP	→ JJ NP	0.2
NP	→ NP PP	0.1
VP	→ V	0.4
VP	→ V NP	0.3
VP	→ V NP NP	0.1
VP	→ VP PP	0.2
PP	→ P NP	1.0

Table 13.1: A fragment of an example probabilistic context-free grammar (PCFG)

The probabilities for all productions involving a single LHS must sum to 1:

$$\sum_\alpha P(X \to \alpha \mid X) = 1$$

The probability $p(\tau)$ is just the product of all the productions:

```
                        S
               ┌────────┴────────┐
              NP                 VP
               │         ┌───────┼────────┐
               N         V      NP        PP
               │         │       │     ┌───┴────┐
             They      eat      N      P        NP
                                │      │         │
                              sushi   with       N
                                                 │
                                             chopsticks
```

Estimation

- As in supervised HMMs, estimation is easy (for now!).

- PCFG probabilities can be estimated directly from a treebank:

$$P(\text{VP} \to \text{VP PP}) = \frac{\text{count}(\text{VP} \to \text{VP PP})}{\text{count}(\text{VP})}$$

Three basic problems for PCFGs

Let $\tau \in T$ be a derivation, $\boldsymbol{w}$ be a sentence, and λ a PCFG.

- **Decoding**: Find $\hat{\tau} = \arg\max_\tau p(\tau, \boldsymbol{w}; \lambda)$

- **Likelihood**: Find $p(\boldsymbol{w}; \lambda) = \sum_\tau p(\tau, \boldsymbol{w}; \lambda)$

- **(Unsupervised) Estimation**: Find $\arg\max_\lambda p(\boldsymbol{w}_{1\ldots N} \mid \lambda)$

These three problems are analogous to the problems identified by Rabiner (1989) for Hidden Markov Models.

	Sequences	Trees
model	HMM	PCFG
decoding	Viterbi algorithm	CKY
decoding complexity	$\mathcal{O}(M^2|K|)$	$\mathcal{O}(M^3|R|)$
likelihood	forward algorithm	inside algorithm
marginals	forward-backward	inside-outside

CKY with probabilities

It is not difficult to extend CKY to include probabilities or other weights. This is shown in Algorithm 2.

Algorithm 2 CKY with weighted productions

for j : [1,N] **do**
  for $X : tags(w_j)$ **do**
    $t[X, j, j - 1] \leftarrow P(X, w_j)$
  end for
  for i : [j-2, 0] **do**
    for $(X \rightarrow Y\ Z) \in R$ **do**
      for k : [i+1, j-1] **do**
        $t[X, i, j] \leftarrow t[X, i, j] \bigoplus (P(X \rightarrow Y\ Z) \otimes t[Y, i, k] \otimes t[Z, k + 1, j])$
      end for
    end for
  end for
end for

In the boolean semiring, we have $\oplus = \vee$ and $\otimes = \wedge$, so we recover the original CKY algorithm. In the probability semiring, we have $\oplus = +$ and $\otimes = \times$.

$$t[Y, i, k] = P(Y \rightarrow \boldsymbol{w}_{i:k}) \tag{13.1}$$

$$t[Z, k, j] = P(Z \rightarrow \boldsymbol{w}_{k+1:j}) \tag{13.2}$$

$$t[X, i, k] = \sum_{Y,Z,k} P(X \rightarrow Y, Z) P(Y \rightarrow \boldsymbol{w}_{i:k}) P(Z \rightarrow \boldsymbol{w}k + 1 : j) \tag{13.3}$$

$$= P(X \rightarrow \boldsymbol{w}_{i:j}) \tag{13.4}$$

13.11   Review of algorithms for PCFG Parsing

PCFGs score the probability of a derivation $P(\tau, \boldsymbol{s})$ as the product of all productions in τ

- CKY is a bottom-up algorithm for finding $\hat{\tau} = \arg\max_\tau \mathrm{p}(\tau, \boldsymbol{w})$.

- The **inside algorithm** finds $\mathrm{p}(\boldsymbol{w}) = \sum_\tau \mathrm{p}(\tau, \boldsymbol{w})$.

- These algorithms are related through semiring notation.

- These algorithms are efficient, but they are pretty implausible as a model of human parsing.

Shift-reduce   is a left-to-right parsing algorithm, which you may find more cognitively plausible.

- It is related to the pushdown automata representation of context-free grammars: we move through the sentence while keeping a stack with infinite depth.

- At each step, we have two choices

  - **shift** the next word on to the stack
  - **reduce** the stack by applying some production

- If we can clear all the input and end up with just S on the stack, we have parsed the sentence correctly.

- Each reduce move is a production in the derivation.

- We might train a classifier to decide between shift and reduce.

- More on this in the chapter on dependency parsing.

13.12   Improving PCFG parsing

Regardless of the parsing algorithm, pure PCFG parsing on Penn Treebank nonterminals (e.g., NP, VP) doesn't work well.

- No flexibility on PP attachment, depends only on $P(\mathrm{NP} \to \mathrm{NP}\ \mathrm{PP}) > P(\mathrm{VP} \to \mathrm{VP}\ \mathrm{PP})$

- No way to express semantic preferences or preference for right-branching or left-branching structures.

- Rare productions get subsumed by common ones.

- See slides for examples.

Modern generative parsing algorithms automatically refine these nonterminals in various ways.

- **Tree transformations**: automatically modify the parse trees — for example, Markovizing by labeling each non-terminal with its **parent**, as in NPŜ.

- **Lexicalization**: label each non-terminal with its head **word**

Tree transformations

Johnson proposed a series of transformations to PTB trees that improve parsing accuracy.
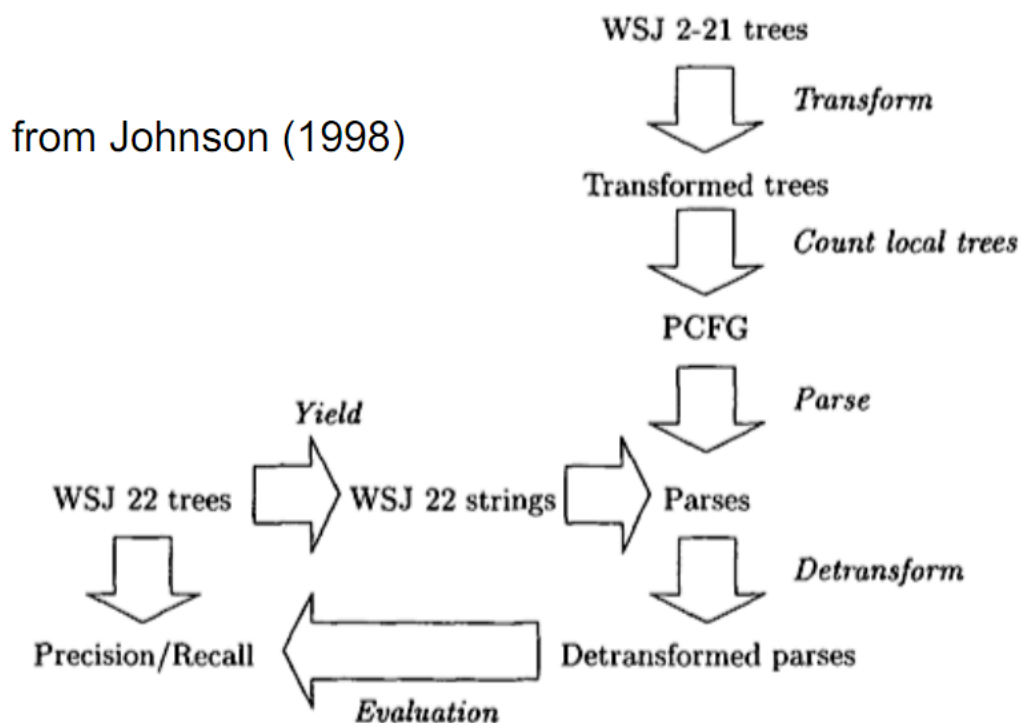
from Johnson (1998)

Flattening   Johnson (1998) proposes "flattening" nested NPs to be more like VP structures.

```
        NP                          NP
      /    \                      /  |  \
    NP      PP                  DT  NN   PP
   /  \    /  \                 |    |   /  \
  DT  NN  P   NNP              the  rain P   NNP
  |   |   |    |                        |    |
 the rain in  Spain                     in  Spain
```

Flattened rules are of course still context-free, but by reducing recursion, they allow more specific probabilities to be learned. This can eliminate the problems with rule subsumption that we saw earlier.

**Parent annotation**   The expansion of an NP is highly dependent on its parent.

$$P(NP \rightarrow NP\ PP) = 11\%$$
$$P(NP(\text{under } S) \rightarrow NP\ PP) = 9\%$$
$$P(NP(\text{under } VP) \rightarrow NP\ PP) = 23\%$$

Parent annotation: augment each non-terminal with its parent.

S
NP   VP
she   V   NP
      heard DT   NN
            the bear

S
NP-S   VP-S
she   V-VP   NP-VP
      heard DT-NP NN-NP
             the   bear

Parent annotation weakens the PCFG independence assumptions

- which could help accuracy by allowing more fine-grained distinctions

- or could hurt accuracy because of data sparseness

Overall, these transformations improve performance:

- Standard PCFG: 72% F-measure, 14,962 rules

- Parent-annotated PCFG: 80% F-measure, 22,773 rules

- In principle, parent annotation could have increased the grammar size much more dramatically, but many possible productions never occur, or are subsumed.
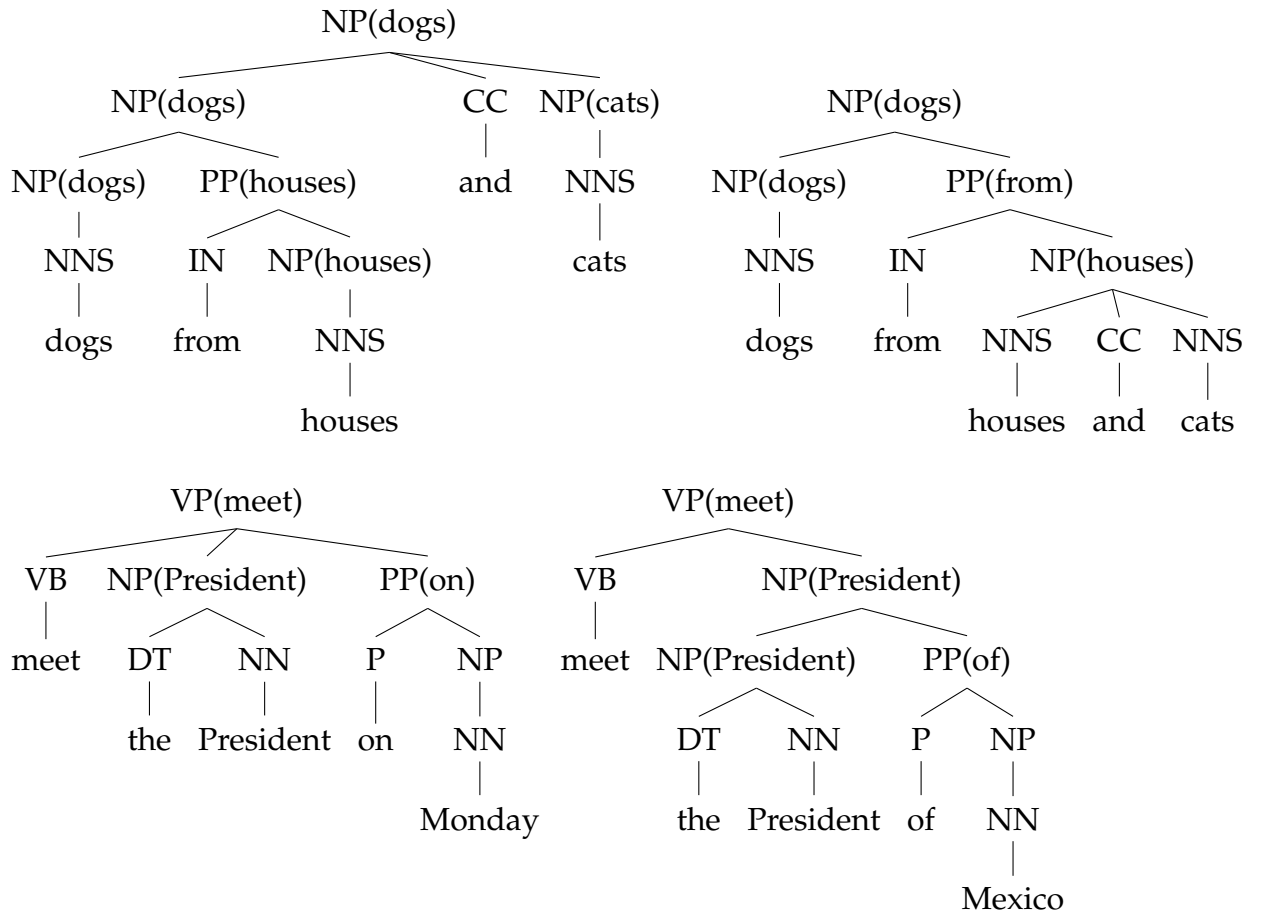
Lexicalization

A simple way to capture semantics is through the words themselves. We can annotate each non-terminal with **head** word of the phrase.

Example: coordination scope If $P(NP \rightarrow NP(\text{dogs})\ CC\ NP(\text{cats})) > (NP \rightarrow NP(\text{houses})\ CC\ NP(\text{cats}))$, we should get the right parse.

Example: PP attachment

- $P(VP(meet) \rightarrow \alpha\ PP(on)) \gg P(NP(President) \rightarrow \beta\ PP(on))$

- $P(VP(meet) \rightarrow \alpha\ PP(of)) \ll P(NP(President) \rightarrow \beta\ PP(of))$

- In plain English:

```
                              NP(dogs)
                  ┌──────────────┼──────────┐
              NP(dogs)          CC    NP(cats)              NP(dogs)
          ┌──────┴──────┐        │       │            ┌───────┴────────┐
      NP(dogs)      PP(houses)  and     NNS       NP(dogs)         PP(from)
         │          ┌───┴────┐           │           │          ┌────┴─────────┐
        NNS        IN    NP(houses)     cats        NNS        IN         NP(houses)
         │          │        │                       │          │     ┌─────┼──────┐
       dogs       from      NNS                     dogs      from   NNS   CC   NNS
                             │                                        │     │     │
                          houses                                   houses  and  cats
```

```
              VP(meet)                               VP(meet)
       ┌─────────┼─────────┐                  ┌─────────┴──────────┐
      VB   NP(President)  PP(on)             VB            NP(President)
       │     ┌────┴───┐    ┌──┴──┐            │          ┌──────┴────────┐
     meet   DT       NN    P    NP          meet   NP(President)      PP(of)
             │        │     │    │                  ┌────┴───┐       ┌───┴───┐
            the  President  on   NN                DT       NN       P      NP
                                  │                 │        │       │       │
                               Monday              the  President   of      NN
                                                                             │
                                                                          Mexico
```

– *Meeting* happens *on* things.

– *Presidents* are *of* things.

Subcategorization frames

$$P(\text{VP} \to \text{V NP NP}) = 0.00151$$
$$P(\text{VP(said)} \to \text{V(said) NP NP}) = 0.00001$$
$$P(\text{VP(gave)} \to \text{V(gave) NP NP}) = 0.01980$$

Lexicalization can capture fine-grained information that the Penn Treebank non-terminals ignore. This had a major impact on parsing accuracy, as shown in Table 13.2.

Vanilla PCFG	72%
Head-annotated PCFG (Johnson, 1998)	80%
Lexicalized PCFG (Collins, 1997, 2003; Charniak, 1997)	87-89%

Table 13.2: Penn Treebank parsing accuracies

Eugene Charniak: "To do better, it is necessary to condition probabilities on the actual words of the sentence. This makes the probabilities much tighter" (at a workshop at Johns Hopkins in 2000).

Algorithms for lexicalized parsing

Naively: we could just augment the non-terminals to include the cross-product of all PTB non-terminals and all words. But this will never work.

- Number of possible productions: $\mathcal{O}(N^3 V^3)$, where the size of the vocabulary V is $V \approx 10^5$

- Too sparse: the total amount of PTB data $= 10^6$ tokens, so we can't learn this many production probabilities.

Charniak (1997) and Collins (1997, 2003) offer practical algorithms, which each make independence assumptions to reduce the complexity.

The Charniak Parser

The Charniak (1997) parser gives a relatively straightforward way to lexicalize PCFGs.

S(rose)
NP(profits)
JJ(Corporate)   NNS(profits)   VP(rose)   VBD(rose)

- Head probabilities capture "bilexical" phenomena, like the PP attachment (*President of Mexico*) example.

- Compute the head probability: $p(w_m | t_m, w_{\rho(m)}, t_{\rho(m)})$.

  – w_m is the head word of constituent m

- – $c(w_m)$ is the **cluster** of the head word of constituent m
- – t_m is the syntactic category
- – $\rho(m)$ is the parent of node m

- In the example, we have

$$
\begin{aligned}
&\mathrm{p}(\textit{profits}|t_m = \mathrm{NP}, t_{\rho(m)} = \mathrm{S}, w_{\rho(m)} = \textit{rose}) \\
&= \frac{\mathrm{count}(w_m = \textit{profits}, t_m = \mathrm{NP}, t_{\rho(m)} = \mathrm{S}, w_{\rho(m)} = \textit{rose})}{\mathrm{count}(t_m = \mathrm{NP}, t_{\rho(m)} = \mathrm{S}, w_{\rho(m)} = \textit{rose})}
\end{aligned}
$$

- We multiply the head probability by a rule probability,

$$
\begin{aligned}
&P(\mathrm{NP} \to \mathrm{JJ\ NNS}|w_{\rho(m)} = \textit{rose}, t_m = \mathrm{NP}, t_{\rho(m)} = \mathrm{S}) \\
&= \frac{\mathrm{count}(\mathrm{NP} \to \mathrm{JJ\ NNS}, t_m = \mathrm{NP}, t_{\rho(m)} = \mathrm{S}, w_{\rho(m)} = \textit{rose})}{\mathrm{count}(t_m = \mathrm{NP}, t_{\rho(m)} = \mathrm{S}, w_{\rho(m)} = \textit{rose})}
\end{aligned}
$$

- This gives the values in the CKY table.

The rule probabilities capture phenomena like verb complement frames.

Local Tree	come	take	think	want
VP → V	9.5%	2.6%	4.6%	5.7%
VP → V NP	1.1%	32.1%	0.2%	13.9%
VP → V PP	34.5%	3.1%	7.1%	0.3%
VP → V SBAR	6.6%	0.3%	73.0%	0.2%
VP → V S	2.2%	1.3%	4.8%	70.8%
VP → V NP S	0.1%	5.7%	0.0%	0.3%
VP → V PRT NP	0.3%	5.8%	0.0%	0.0%
VP → V PRT PP	6.1%	1.5%	0.2%	0.0%

Data sparseness   The Penn Treebank is still the main dataset for syntactic analysis of English. Yet 1M words is not nearly enough data to accurately estimate lexicalized models.

- 965K constituents

- 66 examples of WHADJP

- only 6 of these aren't *how much* or *how many*

In the example above (*corporate profits rose*), the unsmoothed head probability is zero, as estimated from the PTB: there are zero counts of *profits* headed by *rose* in the treebank (hard to believe, but that's what Charniak says).

In general, bilexical counts are going to be very sparse. But the "backed-off" probabilities give a reasonable approximation. We will interpolate between them.

Smoothing the Charniak Parser   Head probability:

$$
\begin{aligned}
\hat{p}(w_m \mid t_m, w_{\rho(m)}, t_{\rho(m)}) =& \lambda_1 p_{mle}(w_m \mid t_m, w_{\rho(m)}, t_{\rho(m)}) \\
&+ \lambda_2 p_{mle}(w_m \mid t_m, \text{cluster}(w_{\rho(m)}), t_{\rho(m)}) \\
&+ \lambda_3 p_{mle}(w_m \mid t_m, t_{\rho(m)}) \\
&+ \lambda_4 p_{mle}(w_m \mid t_m)
\end{aligned}
$$

For example:

	$p(\textit{profit} \mid NP, rose, S)$	$P(\textit{corp.} \mid JJ, profit, NP)$
$p(w_m \mid t_m, w_{\rho(m)}, t_{\rho(m)})$	0	.245
$p(w_m \mid t_m, c(w_{\rho(m)}), t_{\rho(m)})$	.0035	.015
$p(w_m \mid t_m, t_{\rho(m)})$	.00063	.0053
$p(w_m \mid t_m)$	.00056	.0042

We have to tune $\lambda_1 \ldots \lambda_4$, and an equivalent set of parameters for the rule probabilities.

- The Charniak parser suffers from acute sparsity problems because it estimates the probability of entire rules.

- Another extreme would be to generate the children independently from each other.
  e.g., $P(\text{S} \to \text{NP VP}) \approx P_L(\text{S} \to \text{NP}) P_R(\text{S} \to \text{VP})$

- Collins (2003) and Charniak (2000) go for a compromise, conditioning on the parent and the head child.

The Collins Parser

- The Charniak parser focuses on lexical relationships between children and parents.

- The Collins (2003) parser focuses on relationships between adjacent children of the same parent. It decomposes each rule as,

$$
X \to L_m L_{m-1} \ldots L_1 H R_1 \ldots R_{n-1} R_n
$$

- Each L and R is a child constituent of X, and they are generated from the head H outwards.

- The outermost elements of L and R are special $\square$ symbols.

VP(dumped)

VBD(dumped) NP(sacks)   PP(into)

| |

dumped   sacks   into the river

To model this rule, we would compute:

$$\mathrm{p}(\mathrm{VP}(\textit{dumped}, \mathrm{VBD}) \to [\square, \mathrm{VBD}(\textit{dumped}, \mathrm{VBD}), \mathrm{NP}(\textit{sacks}, \mathrm{NNS}), \mathrm{PP}(\textit{into}, \mathrm{P}), \square])$$

- Here's the generative process:

  - Generate the head: $P(H \mid LHS) = P(\mathrm{VBD}(\textit{dumped}, \mathrm{VBD}) \mid \mathrm{VP}(\textit{dumped}, \mathrm{VBD}))$
  - Generate the left dependent: $P_L(\square \mid \mathrm{VP}(\textit{dumped}, \mathrm{VBD}), \mathrm{VBD}(\textit{dumped}, \mathrm{VBD}))$
  - Generate the right dependent: $P_R(\mathrm{NP}(\textit{sacks}, \mathrm{NNS}) \mid \mathrm{VP}(\textit{dumped}, \mathrm{VBD}), \mathrm{VBD}(\textit{dumped}, \mathrm{VBD}))$
  - Generate the right dependent: $P_R(\mathrm{NP}(\textit{into}, PP) \mid \mathrm{VP}(\textit{dumped}, \mathrm{VBD}), \mathrm{VBD}(\textit{dumped}, \mathrm{VBD}))$
  - Generate the right dependent: $P_R(\square \mid \mathrm{VP}(\textit{dumped}, \mathrm{VBD}), \mathrm{VBD}(\textit{dumped}, \mathrm{VBD}))$

- The rule probability is the product of these generative probabilities.

- Collins parser also conditions on a "distance" of each constituent from the head.

Smoothing the Collins Parser   Estimation is eased by factoring the rule probabilities, but smoothing is still needed.

$$
\begin{aligned}
\hat{P}(R_m(rw_m, rt_m) \mid \rho(m), hw, ht) = & \lambda_1 P_{mle}(R_m(rw_m, rt_m) \mid \rho(m), hw, ht) \\
& + \lambda_2 P_{mle}(R_m(rw_m, rt_m) \mid \rho(m), ht) \\
& + \lambda_3 P_{mle}(R_m(rw_m, rt_m) \mid \rho(m))
\end{aligned}
$$

We set λ using Witten-Bell smoothing.

Bilexical dependencies　The collins parser models bilexical dependencies in $P_{mle}(R_m(rw_m, rt_m) \mid \rho(m), hw, ht)$. Is it worth it?

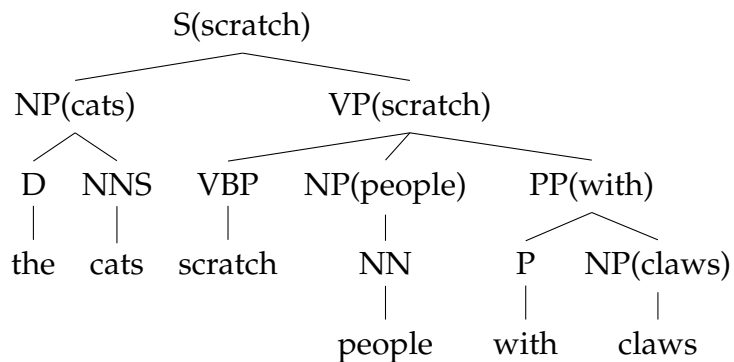| Back-off level	Number of accesses	Percentage
0	3,257,309	1.49
1	24,294,084	11.0
2	191,527,387	87.4
Total	219,078,780	100.0

- In general, bilexical probabilites are rarely available...

- ...but they are active in 29% of the rules in **top-scoring** parses.

- Still, they don't seem to play a big role in accuracy (Bikel, 2004)

13.13　Modern competitive constituent parsing

Chapter 14

Dependency Parsing

Lexicalized parsing augments the non-terminals with **head words**.

```
                    S(scratch)
            _____|_____
        NP(cats)                   VP(scratch)
         __|__              _____|_____
        D    NNS          VBP    NP(people)    PP(with)
        |     |            |         |         ___|___
       the   cats       scratch     NN        P    NP(claws)
                                    |          |        |
                                  people     with     claws
```

- A set of deterministic **head percolation** rules determine how heads move up the tree through each production.

- Some rules are pretty obvious:

    - the head of a determiner-noun constituent is the noun:

        ```
             NP(cats)
              __|__
        D(the)   NP(cats)
        ```

    - prepositional adjuncts are not heads:

        ```
                NP(people)
                 ___|___
        NP(people)   PP(with claws)
        ```

165

- – verbal predicates are the heads of sentences:

  S(scratch)

  NP(cats)   VP(scratch)

- Others are less clear-cut:

  - – the head of a conjunction is the left element:

    VP(cats)

    VP(cats)   CC(and)   VP(dogs)

  - – the head of a prepositional phrase is the preposition:

    PP(with)

    P(with)   NP(claws)

    (Alternatively, we can "collapse" out the preposition itself. This is the approach taken by the Stanford dependency parser.)

- We could argue about this stuff, but once we agree on a standard it can be applied deterministically to any parse tree.

A head-annotated parse tree defines a graph over the words in the sentence:

The cats scratch people with claws

What are the properties of this graph?

- directed

- weakly connected

- every node has one incoming edge

- (therefore) no cycles

- (therefore) a tree

What is the meaning of the edges?

- A dependency edge means there is an asymmetric syntactic relationship between the head and the modifier.

- (sometimes called nucleus/satellite, or governor/dependent, or parent/child)

- Criteria for figuring out who is the head:

  - The modifier may be optional while the head is mandatory: (*red HAT*, *EAT quickly*)
  - The head sets the syntactic category of the construction: (prepositions are the heads of prepositional phrases)
  - The head determines the morphological form of the modifier: (*los LIBROS*, *una CASA*, *these HOUSES*)

- As always, these guidelines sometimes conflict.

- Edges may be **labeled** to indicate their function:

Dependency trees tell us who did what to whom.

Expressiveness

(Unlabeled) dependency trees are less expressive than CFG derivations. That means they hide some of the ambiguity in CFGs that we may not care about. Remember the different representations for PP modification?

- Flat

- (Chomsky) adjunction

```
                        VP
              ┌──────────┴──────────┐
             VP                     PP
        ┌─────┴─────┐          ┌─────┴─────┐
       VP           PP        with a fork
    ┌───┴───┐   ┌────┴────┐
    V       NP  on the table
    │       │
   ate    dinner
```

- Two-level (PTB)

```
                    VP
        ┌───────────┼───────────┐
       VP           PP          PP
    ┌───┴───┐  ┌────┴────┐  ┌────┴────┐
    V       NP on the table with a fork
    │       │
   ate    dinner
```

These all look the same in a dependency parse:

ate dinner on the table with a fork

So if you didn't think there was any meaningful difference between these representations, you should be happy. But many kinds of CFG ambiguity remain in the dependency parse:

southern food store   southern food store

Applications

There are many; see slides for an example.

Projectivity

In projective dependency parsing, there can be no crossing edges.

- Crossing edges are rare in English

She ate a pizza yesterday which was vegetarian

- They are quite common in other languages, like Czech.

- We can build projective dependency banks from constituent treebanks, like PTB.

  - In this case, crossing edges are prohibited by construction.
  - What is the CFG analysis for the Pizza example?

- In languages where non-projectivity is common, we can annotated it directly. Prague dependency treebank: 1.5M words of Czech.

Algorithms

Let's assume that the score for a dependency parse can be **factored** across the edges:

$$\Psi(G) = \bigotimes_i \psi(g(i) \to i) \tag{14.1}$$

Then we have efficient dynamic programs for dependency parsing.

- Non-projective dependency parsing reduces to the **maximum spanning tree** problem (in directed graphs).

  - We need a tree that touches all the nodes.
  - We want to get the maximum score $\bigotimes_i \psi(g(i) \to i)$, where $g(i)$ is the head of node i.
  - Chu-Liu-Edmonds algorithm computes this in $\mathcal{O}(N^3)$. [See slides.]
  - Tarjan algorithm is $\mathcal{O}(N^2)$

- Projective dependency parsing: Naive version

– $c[i, j, h]$ is the score of the best tree spanning $i \rightarrow j$ with head h

$$c[i, j, h] = \left(\bigoplus_{k,h'} c[i, k, h] \otimes c[k, j, h'] \otimes \psi(h \rightarrow h') \right)$$

$$\oplus \left(\bigoplus_{k,h'} c[i, k, h'] \otimes c[k, j, h] \otimes \psi(h \rightarrow h') \right)$$

– The first line represents left-branching trees; the second line represents right-branching trees.

– What is the complexity? The size of the table is $\mathcal{O}(N^3)$. To fill in each node, we must consider $\mathcal{O}(N)$ possible split points and $\mathcal{O}(N)$ possible heads h' for the satellite subtree. So the time complexity is $\mathcal{O}(N^5)$.

– The Eisner algorithm reduces this to $\mathcal{O}(N^3)$, by adapting the CKY algorithm. We keep four tables instead of 1!

  * scores of **incomplete** subtrees from i to j, headed to the left
  * scores of **incomplete** subtrees from i to j, headed to the right
  * scores of **complete** subtrees from i to j, headed to the left
  * scores of **complete** subtrees from i to j, headed to the right

– Why?

  * Incomplete subtrees can subsume complete subtrees heading in the same direction.
  * Complete subtrees can combine if they are heading in opposite directions.
  * Our goal is to produce a complete tree from 0 to N.

Learning dependency parsers

– Generative: $\psi(m \rightarrow n) = \log P(h = m | n)$, estimated from Treebank
– Log-linear: $\psi(m \rightarrow n) = \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(w_m, w_n)$. Features:

  * POS tag of w_m and w_n
  * Word identity of w_m and w_n
  * Word shape (e.g., suffix, prefix)
  * Distance $(m - n)$

* ...
- Structured perceptron

$$\hat{\boldsymbol{y}} = \arg\max_{\boldsymbol{y}' \in \mathcal{T}(\boldsymbol{w})} \boldsymbol{\theta}^\mathsf{T} \boldsymbol{f}(\boldsymbol{w}, \boldsymbol{y}')$$

$$\boldsymbol{\theta}^\mathsf{T} = \boldsymbol{\theta}^\mathsf{T} + \boldsymbol{f}(\boldsymbol{w}, \boldsymbol{y}) - \boldsymbol{f}(\boldsymbol{w}, \hat{\boldsymbol{y}})$$

Same as before, but now $\arg\max_{\boldsymbol{y}' \in \mathcal{T}(\boldsymbol{x})}$ searches over dependency trees, using either MST or the Eisner algorithm.

- Conditional Random Field (CRF) a.k.a. globally-normalized conditional model

$$\frac{\partial \mathcal{L}}{\partial \boldsymbol{\theta}^\mathsf{T}} = \sum_m \boldsymbol{f}(x_{h(m)} \to w_m) - \sum_n P(n \to m | \boldsymbol{w}) \boldsymbol{f}(w_n \to w_m) \qquad (14.2)$$

We require **marginal** probabilities $P(n \to m | \boldsymbol{w})$. These can be obtained efficiently using a variant of inside-outside (for projective trees), and the matrix-tree theorem for non-projective trees (Koo et al., 2007).

Transition-based parsing

An alternative to exact global inference is transition-based parsing: making a series of local decisions. For projective dependency parsing, the algorithm is very similar to shift-reduce.

- Read the sentences left-to-right,

  - **shift**: push a word onto the stack
  - **arc-right**: make a right-facing edge
  - **arc-left**: make a left-facing edge

- **Beam search**

  - Keep a "beam" of possible hypotheses.
  - At each stage, let the k best unique hypotheses stay on the beam.
  - This way, we can recover from early mistakes

- **Learning**

- Identify the series of decisions required to produce the correct dependency parse.
- Each decision in the derivation of the correct parse is a positive instance. Each other possible decision is a negative instance.
- Huang et al. (2012) offer alternative perceptron learning rules that yield improvements when learning in the beam search setting.

A key advantage of transition-based parsing is that there is no restriction to arc-factored features; we can include any feature of the current partial parse, history of decisions, etc. It is also fast: linear time in the length of the sentence.

Bibliography

Allauzen, C., Riley, M., and Schalkwyk, J. (2009). A generalized composition algorithm for weighted finite-state transducers. In *Proceedings of Interspeech*.

Bender, E. M. (2013). *Linguistic Fundamentals for Natural Language Processing: 100 Essentials from Morphology and Syntax*, volume 6 of *Synthesis Lectures on Human Language Technologies*. Morgan & Claypool Publishers.

Berger, A. L., Pietra, V. J. D., and Pietra, S. A. D. (1996). A maximum entropy approach to natural language processing. *Computational linguistics*, 22(1):39–71.

Bikel, D. M. (2004). Intricacies of collins' parsing model. *Computational Linguistics*, 30(4):479–511.

Bottou, L. (1998). Online learning and stochastic approximations. *On-line learning in neural networks*, 17:9.

Boyd, S. and Vandenberghe, L. (2004). *Convex Optimization*. Cambridge University Press, New York, NY, USA.

Brants, T. (2000). Tnt: a statistical part-of-speech tagger. In *Proceedings of the sixth conference on Applied natural language processing*, pages 224–231. Association for Computational Linguistics.

Burges, C. J. (1998). A tutorial on support vector machines for pattern recognition. *Data mining and knowledge discovery*, 2(2):121–167.

Charniak, E. (1997). Statistical techniques for natural language parsing. *AI magazine*, 18(4):33.

Charniak, E. (2000). A maximum-entropy-inspired parser. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 132–139.

Choueka, Y. (1989). Responsa: A full-text retrieval system with linguistic processing for a 65-million word corpus of jewish heritage in hebrew. *Data Engineering*, 12(4):22–31.

Christodoulopoulos, C., Goldwater, S., and Steedman, M. (2010). Two decades of unsupervised pos induction: How far have we come? In *Proceedings of the 2010 Conference on Empirical Methods in Natural Language Processing*, pages 575–584. Association for Computational Linguistics.

Church, K. W. (2000). Empirical estimates of adaptation: the chance of two Noriegas is closer to $p/2$ than p^2. In *Proceedings of the 18th conference on Computational linguistics-Volume 1*, pages 180–186.

Collins, M. (1997). Three generative, lexicalised models for statistical parsing. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 16–23.

Collins, M. (2002). Discriminative training methods for hidden markov models: theory and experiments with perceptron algorithms. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 1–8.

Collins, M. (2003). Head-driven statistical models for natural language parsing. *Computational linguistics*, 29(4):589–637.

Collins, M. (2013). Notes on natural language processing. `http://www.cs.columbia.edu/~mcollins/notes-spring2013.html`.

Collins, M. and Brooks, J. (1995). Prepositional phrase attachment through a backed-off model. In *Workshop on Very Large Corpora*.

Collins, M. and Duffy, N. (2001). Convolution kernels for natural language. In *Advances in neural information processing systems*, pages 625–632.

Collobert, R. and Weston, J. (2008). A unified architecture for natural language processing: Deep neural networks with multitask learning. In *Proceedings of the 25th international conference on Machine learning*, pages 160–167. ACM.

Crammer, K., Dekel, O., Keshet, J., Shalev-Shwartz, S., and Singer, Y. (2006). Online passive-aggressive algorithms. *The Journal of Machine Learning Research*, 7:551–585.

Crammer, K. and Singer, Y. (2003). Ultraconservative online algorithms for multiclass problems. *The Journal of Machine Learning Research*, 3:951–991.

Duchi, J., Hazan, E., and Singer, Y. (2011). Adaptive subgradient methods for online learning and stochastic optimization. *The Journal of Machine Learning Research*, 12:2121–2159.

Dyer, C. (2014). Notes on adagrad. `www.ark.cs.cmu.edu/cdyer/adagrad.pdf`.

Eisner, J. (2001). Expectation semirings: Flexible em for learning finite-state transducers. In *Proceedings of the ESSLLI workshop on finite-state methods in NLP*.

Eisner, J. (2002). Parameter estimation for probabilistic finite-state transducers. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 1–8.

Figueiredo, M., Graça, J., Martins, A., Almeida, M., and Coelho, L. P. (2013). LXMLS lab guide. `http://lxmls.it.pt/2013/guide.pdf`.

Finkel, J. R., Grenager, T., and Manning, C. (2005). Incorporating non-local information into information extraction systems by gibbs sampling. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 363–370.

Freund, Y., Schapire, R., and Abe, N. (1999). A short introduction to boosting. *Journal-Japanese Society For Artificial Intelligence*, 14(771-780):1612.

Freund, Y. and Schapire, R. E. (1999). Large margin classification using the perceptron algorithm. *Machine learning*, 37(3):277–296.

Goldwater, S. and Griffiths, T. (2007). A fully bayesian approach to unsupervised part-of-speech tagging. In *Annual meeting-association for computational linguistics*, volume 45.

Hastie, T., Tibshirani, R., Friedman, J., Hastie, T., Friedman, J., and Tibshirani, R. (2009). *The elements of statistical learning*, volume 2. Springer.

Hindle, D. and Rooth, M. (1990). Structural ambiguity and lexical relations. In *Proceedings of the Workshop on Speech and Natural Language*.

Hsu, D., Kakade, S. M., and Zhang, T. (2012). A spectral algorithm for learning hidden markov models. *Journal of Computer and System Sciences*, 78(5):1460–1480.

Huang, L., Fayong, S., and Guo, Y. (2012). Structured perceptron with inexact search. In *Proceedings of the 2012 Conference of the North American Chapter of*

the Association for Computational Linguistics: Human Language Technologies, pages 142–151, Montréal, Canada. Association for Computational Linguistics.

Jaeger, T. F. (2010). Redundancy and reduction: Speakers manage syntactic information density. *Cognitive psychology*, 61(1):23–62.

Johnson, M. (1998). Pcfg models of linguistic tree representations. *Computational Linguistics*, 24(4):613–632.

Jurafsky, D. and Martin, J. H. (2009). *Speech and Language Processing (2nd Edition) (Prentice Hall Series in Artificial Intelligence)*. Prentice Hall, 2 edition.

Karlsson, F. (2007). Constraints on multiple center-embedding of clauses. *Journal of Linguistics*, 43(02):365–392.

Karttunen, L. and Beesley, K. R. (2001). A short history of two-level morphology. *ESSLLI-2001 Special Event titled Twenty Years of Finite-State Morphology*.

Karttunen, L. and Beesley, K. R. (2005). Twenty-five years of finite-state morphology. *Inquiries Into Words, a Festschrift for Kimmo Koskenniemi on his 60th Birthday*, pages 71–83.

Knight, K. and May, J. (2009). Applications of weighted automata in natural language processing. In *Handbook of Weighted Automata*, pages 571–596. Springer.

Koo, T., Globerson, A., Carreras, X., and Collins, M. (2007). Structured prediction models via the matrix-tree theorem. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 141–150.

Lafferty, J., McCallum, A., and Pereira, F. (2001). Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *Proceedings of the International Conference on Machine Learning (ICML)*.

Manning, C. D. and Schütze, H. (1999). *Foundations of statistical natural language processing*. MIT press.

Marcus, M. P., Marcinkiewicz, M. A., and Santorini, B. (1993). Building a large annotated corpus of English: The Penn Treebank. *Computational linguistics*, 19(2):313–330.

Mikolov, T., Deoras, A., Povey, D., Burget, L., and Cernocky, J. (2011). Strategies for training large scale neural network language models. In *Automatic Speech*

Recognition and Understanding (ASRU), 2011 IEEE Workshop on, pages 196–201. IEEE.

Minka, T. P. (1999). From hidden markov models to linear dynamical systems. In *Tech. Rep. 531, Vision and Modeling Group of Media Lab, MIT*.

Mitra, S. and Acharya, T. (2007). Gesture recognition: A survey. *Systems, Man, and Cybernetics, Part C: Applications and Reviews, IEEE Transactions on*, 37(3):311–324.

Mnih, A. and Hinton, G. E. (2008). A scalable hierarchical distributed language model. In *Neural Information Processing Systems (NIPS)*, pages 1081–1088.

Mohri, M., Pereira, F., and Riley, M. (2002). Weighted finite-state transducers in speech recognition. *Computer Speech & Language*, 16(1):69–88.

Murphy, K. P. (2012). *Machine Learning: A Probabilistic Perspective*. The MIT Press.

Nemirovski, A. and Yudin, D. (1978). On Cezari's convergence of the steepest descent method for approximating saddle points of convex-concave functions. *Soviet Math. Dokl.*, 19.

Nigam, K., McCallum, A. K., Thrun, S., and Mitchell, T. (2000). Text classification from labeled and unlabeled documents using em. *Machine learning*, 39(2-3):103–134.

Pereira, F. (2000). Formal grammar and information theory: together again? *Philosophical Transactions of the Royal Society of London. Series A: Mathematical, Physical and Engineering Sciences*, 358(1769):1239–1253.

Petrov, S., Das, D., and McDonald, R. (2012). A universal part-of-speech tagset. In *Proceedings of LREC*.

Rabiner, L. (1989). A tutorial on hidden markov models and selected applications in speech recognition. *Proceedings of the IEEE*, 77(2):257–286.

Ratnaparkhi, A., Reynar, J., and Roukos, S. (1994). A maximum entropy model for prepositional phrase attachment. In *Proceedings of the workshop on Human Language Technology*, pages 250–255. Association for Computational Linguistics.

Roark, B., Saraclar, M., and Collins, M. (2007). Discriminative¡ i¿ n¡/i¿-gram language modeling. *Computer Speech & Language*, 21(2):373–392.

Rosenfeld, R. (1996). A maximum entropy approach to adaptive statistical language modelling. *Computer Speech & Language*, 10(3):187–228.

Saul, L. and Pereira, F. (1997). Aggregate and mixed-order markov models for statistical language processing. In *emnlp*.

Schmid, H. (1994). Probabilistic part-of-speech tagging using decision trees. In *Proceedings of the international conference on new methods in language processing*, volume 12, pages 44–49. Manchester, UK.

Shieber, S. M. (1985). Evidence against the context-freeness of natural language. *Linguistics and Philosophy*, 8(3):333–343.

Shwartz, S. S., Singer, Y., and Srebro, N. (2007). Pegasos: Primal estimated sub-GrAdient SOlver for SVM. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 807–814.

Song, L., Boots, B., Siddiqi, S. M., Gordon, G. J., and Smola, A. J. (2010). Hilbert space embeddings of hidden markov models. In *Proceedings of the International Conference on Machine Learning (ICML)*, pages 991–998.

Sra, S., Nowozin, S., and Wright, S. J. (2012). *Optimization for machine learning*. MIT Press.

Taskar, B., Guestrin, C., and Koller, D. (2003). Max-margin markov networks. In *NIPS*.

Tausczik, Y. R. and Pennebaker, J. W. (2010). The psychological meaning of words: Liwc and computerized text analysis methods. *Journal of Language and Social Psychology*, 29(1):24–54.

Teh, Y. W. (2006). A hierarchical bayesian language model based on pitman-yor processes. In *Proceedings of the Association for Computational Linguistics (ACL)*, pages 985–992.

Tsochantaridis, I., Hofmann, T., Joachims, T., and Altun, Y. (2004). Support vector machine learning for interdependent and structured output spaces. In *Proceedings of the twenty-first international conference on Machine learning*, page 104. ACM.

Van Gael, J., Vlachos, A., and Ghahramani, Z. (2009). The infinite hmm for unsupervised pos tagging. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pages 678–687, Singapore.

Zelenko, D., Aone, C., and Richardella, A. (2003). Kernel methods for relation extraction. *The Journal of Machine Learning Research*, 3:1083–1106.

Zhang, T. (2004). Solving large scale linear prediction problems using stochastic gradient descent algorithms. In *Proceedings of the twenty-first international conference on Machine learning*, page 116. ACM.